Hippocratic AI
— Do No Harm —

Comment Letter: CMS-0042-NC – Health Technology Ecosystem RFI Response

Submitted by: *Hippocratic AI, Inc.|CAGE:106Y9 | UEI: G89FMSKNHAD8*
Re: Request for Information – *Health Technology Ecosystem* (CMS-0042-NC)
Date: 5 June 2025

To: **Stephanie Carlton**, Deputy Administrator of the Centers for Medicare & Medicaid Services and **Steven Posnack**, Acting Assistant Secretary for Technology Policy, Acting National Coordinator for Health Information Technology.

Thank you for the opportunity to provide comments on the ***Health Technology Ecosystem RFI (CMS-0042-NC)***. Hippocratic AI is a safety-focused, HIPAA Compliant, patient-centric generative AI company dedicated to delivering "healthcare abundance" – making high-quality care infinitely scalable and accessible. We develop non-diagnostic AI healthcare agents that engage patients via natural conversations (e.g. phone, chat) for low-risk care navigation, follow-ups, education, and administrative support. Our platform emphasizes **patient safety, prevention, and empowerment** as core design principles. We are pleased to offer perspectives from our experience as a technology vendor working with health systems, payers, and value-based care organizations across the country.

Our comments focus primarily on sections **E. Technology Vendors, Data Providers, and Networks (TD)** and **F. Value-Based Care Organizations (VB)**. We also address select questions in sections **A. Patients** (**PC-5a & TD-1**)

Delivering Healthcare Abundance – Safety-Focused and Scalable.

Recommendation: CMS should consider establishing a **Digital Health Sandbox and voluntary "Safe Harbor" certification** for non-diagnostic, safety-tested AI health agents.

Rather than traditional FDA-style approvals, this would be a **CMS-led validation program** to review digital health products for efficacy, quality, and safety in improving outcomes. Hippocratic AI strongly supports such a framework. We have demonstrated through a rigorous *Real-World Evaluation* that generative AI agents can achieve safety parity with clinicians. In over 307,000 patient-agent interactions, our AI provided correct medical advice 99.38% of the time without making a single error that could have caused severe patient harm. By engaging over *6,200 licensed clinicians* to test our AI across iterative improvements, we eliminated major safety concerns and reached human-level accuracy. A CMS sandbox could allow developers to pilot similarly safety-optimized AI tools on Medicare use cases (e.g. post-discharge follow-ups, care navigation, preventative care education) under real-world conditions. High-performing solutions that emerge from the sandbox would earn a "Safe Harbor" designation – giving providers and ACOs confidence to adopt them without fear of regulatory uncertainty.

Criteria for this safe harbor might include: robust clinical testing (e.g. better than human accuracy, no critical errors), transparency of algorithms, FHIR-based data exchange, and ongoing monitoring

for quality improvement. This voluntary pathway would **spur developer interest (TD-1)** by providing clear guidelines and a stamp of credibility, and it would **increase patient and provider trust** in AI-enabled digital health **(PC-5a)** by ensuring only **validated, safety-first tools** are scaled.

(PC-5b & PC-6) ***Promoting Prevention and Equitable Access***. To drive patient and caregiver adoption of digital health **(PC-5b)**, CMS should leverage these safe, user-friendly AI agents and incentivize their use in preventive care programs. Hippocratic AI's agents have proven effective in motivational health coaching – for example, our multilingual AI *"Ana"* conducted outreach calls that significantly boosted colorectal cancer screening rates among Spanish-speaking Medicare patients. By speaking the patient's preferred language and using *empathetic, culturally-tailored dialogue*, the AI achieved FIT-test opt-in rates 2.6× higher for Spanish speakers than for English speakers. Notably, many of these beneficiaries had limited prior experience with patient portals or apps. The key to engagement was meeting patients where they are – via a simple phone call in their own language – and building trust through empathy. The AI agent "Ana" explains the importance of preventive screening, listens to patient concerns, and helps arrange at-home test kits, mirroring the compassionate tone of a nurse. This illustrates features critical to making digital health products accessible and easy to use for Medicare beneficiaries (PC-6): multi-language support (14+ languages), voice-based interfaces, and human-like empathy. Unlike smartphone apps that some seniors may struggle to navigate, a telephone-based AI companion can reach even those with low digital literacy.

CMS could support pilot programs where such AI-driven preventive outreach is deployed at scale for chronic disease management, prevention and screenings. By **empowering individuals in their health goals** with friendly, always-available coaching, we can improve prevention and wellness across diverse populations. Early results show tangible outcomes – e.g. our AI-driven wellness calls for congestive heart failure patients saw 96% completion rates and a 4.6% reduction in 30-day readmissions (internal data), indicating fewer complications when patients receive consistent follow-up. Importantly, these agents also capture patient-reported outcomes like PHQ-9 depression screenings and Social Determinants of Health during conversations, feeding that data back to care teams to inform comprehensive holistic care plans.

In short, **safety-focused AI can be a force-multiplier for prevention and health equity**, helping achieve the *"healthcare abundance"* vision where every patient gets the guidance they need to stay healthy.

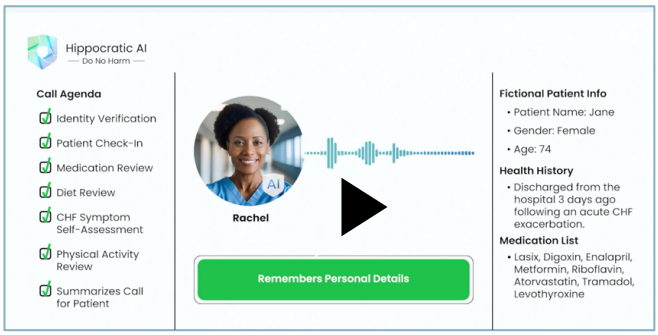

Figure 1 : A fictitious patient interacts with "Rachel," a Hippocratic AI generative agent, for a post-discharge follow-up. Such AI assistants engage patients in natural conversation, ***asking and answering health questions in real-time****. They bring a compassionate, human-like presence to digital health –* ***empowering patients*** *to voice concerns and receive guidance from home.* ***Access a demonstration via this link: https://vimeo.com/1020044938***

(PR-1) ***Reducing Burden on Providers and Care Teams.*** Generative AI health agents offer a promising solution to relieve overextended providers by handling routine patient outreach,

education, and data gathering. To encourage provider uptake of these tools, the **safety and quality must be assured** (as addressed above) and workflow integration must be seamless. **Current obstacles (PR-1a)** include provider concern about liability or accuracy of AI recommendations, and lack of time to manage yet another IT tool. We address these by keeping clinicians "in the loop" as needed. For instance, our AI agents focus strictly on non-diagnostic tasks (post-discharge check-ins, Annual Wellness Visit intakes, preventive reminders) and generate a full transcript or summary of each patient encounter for the provider's review. If any red flags or urgent issues are detected during an AI call, the system can escalate to a human nurse or physician in real-time (e.g. warm transfer to a telehealth nurse). This ensures the provider remains informed and in control of the patient's care plan – the AI is an extender, not an independent decision-maker. Providers should disclose to patients when a digital agent is assisting in their care and clarify the agent's role and limitations **(PR-1b)**. For example, at call commencement our agent introduces itself as a virtual assistant working with the care team, so the patient knows they are interacting with an AI. **Providers have a responsibility (PR-1c)** to vet that the AI tool is appropriate for the patient (e.g. only using agents that are *CMS-vetted/Safe Harbor*-certified for safety), and to monitor outcomes (reviewing transcripts or alerts). In practice, however, the burden on providers is minimal – over 95% of our AI-guided patient calls run on "auto-pilot" without needing human intervention, and the notes/transcripts are stored in the EHR for later reference rather than demanding immediate physician time. By outsourcing routine but time-intensive tasks to AI – medication check-ins, lifestyle coaching, scheduling follow-ups – clinicians can refocus on complex care. This can meaningfully reduce burnout and administrative burden. One large health system noted that our AI completed hundreds of post-discharge calls that staff would otherwise struggle to perform, yet patients rated the experience *8.95/10 in satisfaction*, comparable to human interactions. In summary, clear governance (disclosure, oversight) and proven safety outcomes will give providers confidence to delegate certain care tasks to AI, *lightening their load while keeping care quality high*.

(PR-2 & PR-4) ***Integrating AI into Clinical Workflows***. Key obstacles hindering deployment of innovative digital applications in provider workflows **(PR-2)** are lack of interoperability with clinical systems and unclear reimbursement models. To truly streamline documentation or quality reporting via AI, the agents must plug into EHRs and other health IT systems. We recommend CMS promote **standards and policies that enable third-party digital tools to perform administrative workflows (PR-4)** – for example, open scheduling APIs so an agent can book or modify appointments on a patient's behalf, or SMART on FHIR integration that allows an AI to auto-populate intake forms directly into the EHR. Today, many of these workflows (scheduling, form intake) are locked behind proprietary portals or require manual staff entry. By working with ONC to expand FHIR® API use cases to cover scheduling, referrals, and other administrative data, CMS can unlock innovation. Our AI agents, for instance, could *automatically update a patient's care plan or flag a quality measure gap in the EHR* after a call – but this is feasible only if vendor systems provide the hooks for write-back. We applaud CMS's inclusion of the Provider Access API and other new FHIR endpoints in recent interoperability rules, and encourage continued focus on data liquidity. Dual-purpose investments that enable data reuse will also **mitigate provider burden (PR-2):** when an AI gathers patient information (e.g. home blood pressure readings or social needs), sending that data via FHIR to both the clinical team and a quality reporting database means the work serves care delivery *and* reporting needs simultaneously. **CMS could further stimulate integration by funding pilot programs or challenges for EHR vendors to develop *"plug-and-play" API modules* for common workflows**

(appointments, forms, etc.), specifically to support digital health apps (including AI agents). Removing technical barriers will invite a wave of innovative tools that make providers' lives easier by automating documentation and coordination tasks that currently sap clinical productivity.

(TD-2 & PC-9) ***Data as Fuel for Innovation – Expanding CMS APIs***. Hippocratic AI leverages both clinical and claims data to tailor our agents' conversations to each patient. CMS's Blue Button 2.0 API has been a foundational resource by providing OAuth access to Medicare claims (Parts A, B, D) for patient-facing apps. However, to truly **unleash innovation (TD-2),** additional data should be made available via API beyond the current scope **(PC-9)**. We recommend CMS prioritize opening the following data through secure APIs: **(a)** ***Social risk and SDOH data*** – for example, ICD-10 Z-codes for social needs or entries from health risk assessments. These help an AI agent understand if a patient has transportation barriers, food insecurity, or other factors so it can respond with appropriate resources (or escalate to a human social worker if needed). **(b)** ***Lab results and vital signs*** – Medicare claims include procedural codes for tests but not the result values. Having API access to key lab results (e.g. A1c, lipid levels from outpatient labs) and vital signs recorded during visits would enable richer, more personalized coaching by digital tools. For instance, an AI could congratulate a diabetic patient on an improved A1c or counsel them if a value is trending poorly, using factual data pulled from claims/EHR. **(c)** ***Medication fill data including OTC usage*** – Blue Button includes Part D prescription claims; additionally incorporating known over-the-counter medications or supplements (perhaps via supplemental beneficiary surveys or integrations) could help identify potential adherence issues or interactions. **(d)** ***Provider information and appointment data*** – an API that exposes a beneficiary's upcoming appointments, or at least the ability to query their providers' scheduling systems (with consent), would allow AI assistants to help with appointment reminders and even rescheduling or finding open slots. Currently, a patient's interactions with scheduling are siloed per provider. A unified approach could be for CMS to work with provider directories and Health Information Exchanges so that a patient-authorized app can fetch their upcoming visits. In general, CMS should consider blending claims data with selected clinical and patient-generated data in its API offerings – possibly leveraging the USCDI data classes as a guide **(TD-2b)**. Our agents also connect to electronic health record systems via HL7® v2 feeds and FHIR® resources where available, but smaller providers may not have enabled those. CMS has an opportunity to fill data gaps by expanding Blue Button. Making more data available (with privacy safeguards) will not only spur developer creativity but also directly **empower patients** – they benefit when their apps and AI services have the full picture of their health **(PC-9)**. We believe adding these data types to CMS APIs, along with timely claims updates, will accelerate the next generation of highly effective digital health products.

(TD-4 & TD-5) ***Interoperability and Open Standards***. We strongly agree that open, standards-based APIs are preferable to proprietary systems **(TD-4)**. Hippocratic AI's platform is built with interoperability at its core: our *Polaris* orchestration engine communicates via FHIR® APIs(R4) to retrieve or update health records, and we conform to standards like HL7® v2 and C-CDA when exchanging information with hospital systems. All data our agents generate (call notes, patient-reported metrics, form fill etc.) can be output in standardized formats so it flows into existing IT systems. We encourage CMS and ONC to continue policies that reward use of public standards – for example, advancing the requirement that certified EHRs support bulk FHIR data export (for population health) and SMART on FHIR apps. In our experience, some providers still resort to custom CSV or flat-file exchanges because certain data (e.g. free-text notes or scanned documents) isn't exposed cleanly via API. Certification criteria could be tightened to ensure APIs

cover "all aspects of the patient's chart," **including unstructured data (TD-9c)** – perhaps through emerging standards for text or image exchange.

We support the idea of a **nationwide provider FHIR endpoint directory (TD-5)** to ease connections between systems. As a digital health vendor, finding and maintaining connections to numerous provider EHR endpoints is cumbersome. A centralized directory maintained by a trusted entity (e.g. ONC or a consortium) would help platforms like ours dynamically discover where to send data or queries for a given provider or health organization. We believe users would fund such a utility if it significantly reduces interface costs, though public funding to launch it would be welcome given the public-good nature.

We also support the idea of a **Universal Patient Identifier (UPI) (PC-14)** which would be transformative for Hippocratic AI's generative agents and for CMS's broader interoperability agenda: by giving every Medicare beneficiary a single, privacy-protected token that is recognized across hospitals, post-acute facilities, health plans, and community programs, the AI can reliably "follow" the patient from one touch-point to the next, carry forward the precise clinical context of prior conversations, and reduce incorrect mismatches when names, addresses, or portal log-ins change. For Hippocratic AI, this means each engagement—whether a post-discharge CHF check-in, a preventive cancer-screening reminder, or a social-needs assessment—can refer to the correct longitudinal data, allowing the agent to personalize guidance, prevent duplicative outreach, and maintain the high safety bar (99.38% accuracy) that our RWE-LLM framework demands. For CMS and its stakeholders, accurate patient matching via a UPI would sharply cut the $1.5 million in annual denied-claim costs per hospital and the $6 billion lost system-wide to identification errors, while raising record-matching rates that can dip to as low as 50 percent when data are shared across organizations. A UPI therefore underpins safer AI-driven care navigation, more reliable quality measurement, and smarter stewardship of taxpayer dollars—essential pillars of the digital health strategy to *Make America Healthy Again*.

Finally, Hippocratic AI supports the Trusted Exchange Framework and Common Agreement (TEFCA) network through our integration with EHR vendors. We view TEFCA's nationwide data exchange as a complement to FHIR APIs – **enabling query/retrieval of clinical records across different networks (TD-6)** and through EHR integrations, we operate within the provider and care team workflow. While still in nascent stages, TEFCA can fill gaps especially for smaller providers who aren't covered by existing networks. In sum, interoperability is critical to both innovation and patient safety, and we applaud HHS efforts (like USCDI v3 expansion and recent rules) that push the industry toward seamless, secure data exchange. We encourage CMS to continue using levers like conditions of participation and payment incentives to drive adoption of these open standards and continue to seek other innovations in industry to improve data liquidity.

(VB-1) ***Incentivizing Digital Health in Value-Based Care***. Accountable Care Organizations (ACO) and other APM (Alternative Payment Models) entities have strong motivation to keep patients healthy and out of the hospital – a perfect alignment for preventive digital health interventions. However, ACOs often operate on thin margins and may be hesitant to invest in new technology without clear ROI. CMS can encourage APM participants to leverage digital health tools by incorporating them into the value-based care incentive structure. For example, **CMS could launch a reimbursement pilot or bonus program for AI-augmented care management activities.** If an ACO uses a certified AI agent to conduct Annual Wellness Visit

(AWV) prep calls or chronic disease coaching, those touches could count as billable care management under a temporary code or as in-kind quality improvement activities. Early evidence from our deployments suggests significant payoff: one partnership saw a **4.6 percentage-point reduction in CHF readmissions** (contributing to shared savings) when post-discharge AI calls were added to standard care, and another showed **increased colorectal screening rates in underserved groups**, potentially preventing costly cancers down the line. By sharing these savings or quality gains with ACOs upfront (through model adjustments), CMS would offset the initial costs of adoption. Another incentive is providing credit in performance evaluations – e.g. if an ACO can demonstrate that a high percentage of their patients engage with digital health coaching, it could earn recognition or bonus points in Medicare Shared Savings Program (MSSP) quality scoring. **Current obstacles (VB-1)** include limited capital for technology and unclear guidelines on using such tools (some fear AI might not be allowed or could risk compliance). The Safe Harbor certification discussed earlier would alleviate compliance concerns by signaling CMS's support for these tools. In addition, CMS could facilitate bulk procurement or deployment of proven AI solutions for smaller ACOs. Similar to how CMS has offered subsidized tools for care coordination in the past, CMS could negotiate enterprise licenses for vetted AI agents and offer them to rural ACOs who might not otherwise afford them. By lowering financial and regulatory barriers, CMS will help APM entities embrace digital health solutions that **prevent complications, engage patients, and ultimately save money for Medicare** – aligning with the goal of protecting taxpayer dollars through smarter care.

(VB-2 & VB-3) ***Key Technologies for VBC Success***. We believe several technologies and capabilities are essential for value-based care organizations to thrive, many of which are enhanced by generative AI **(VB-2, VB-3)**:

- **Proactive Care Navigation & Coaching:** AI agents excel at **patient engagement between visits**. In VBC, keeping patients on track – whether it's medication adherence, diet/exercise for chronic disease, or follow-up appointments – is crucial. Our platform's use of *motivational interviewing* techniques drives behavior change by eliciting patients' own goals and barriers. Such personalized coaching can be delivered at scale via AI, ensuring *no patient falls through the cracks*. This directly integrates *patient engagement* and *risk stratification* – the AI can be programmed to focus on high-risk patients (recent discharges, rising risk scores) for more frequent outreach **(VB-2)**. It's a patient empowerment tool that complements clinical care **(VB-3a: akin to a virtual care manager).**
- **Care Planning and Coordination:** AI can support **dynamic care plans** by recording patients' progress and new issues in real time. For example, after each AI call, a summary can be sent to the care team. The care team can suggest followup and actions, which the AI can communicate back to the patient. This ensures the plan is up-to-date without a nurse manually calling. The AI can also follow standard advice guidelines to streamline cycles (e.g., "Patient reports swelling in leg; advised to elevate and call if worsens; will follow up in 24 hours to see if swelling persists). The AI essentially acts as an extra care team member who is always available. It can also trigger **patient event notifications** – e.g. if a patient indicates they visited the ER, the AI flags the ACO care coordinator immediately. These functions complement existing EHR capabilities and could be counted as meeting care coordination requirements in APMs **(VB-3a, VB-3b)**.

- **Data Integration and Quality Measurement:** Value-based models require integrating claims, clinical, and patient-reported data to get a full picture of outcomes. An AI agent that interacts with patients can collect patient-generated health data (symptoms, home readings, satisfaction scores) and feed it into the analytic pipeline. Moreover, by interfacing with claims (as discussed with Blue Button) and EHR data, the AI can help normalize and aggregate data for quality reporting. For instance, during an Annual Wellness Visit intake call, our agent can administer standardized questionnaires (PHQ-9 for depression, functional status, SDOH surveys) and output structured results that map to quality measures. This reduces the burden on providers to manually enter data for HEDIS or Stars measures – the AI has already gathered it. Having these data extraction and normalization tasks automated **(VB-3a)** means ACO staff can focus on intervening on the insights rather than chasing data.
- **Patient Attribution and Outreach:** AI tools can assist in the tedious work of patient attribution and panel management by using algorithms to identify which patients are aligned to which providers or programs and conducting outreach accordingly. For example, in a complex ACO with multiple clinics, an AI can ensure that every attributed beneficiary gets a "welcome to our care program" call and help schedule their first visit. This kind of patient empowerment outreach helps engage patients in the ACO's services early **(VB-3a: an empowerment tool and care navigation function).**

In summary, **artificial intelligence, when applied responsibly, can knit together population health analytics with personalized patient communication (VB-2)**. The result is a set of capabilities – from continuous care coaching to automated quality data collection – that dramatically augment an ACO's ability to manage health outcomes. CMS should encourage ACOs to incorporate these AI-driven capabilities by recognizing them in program requirements or offering flexibility to use them in lieu of more labor-intensive processes. Over time, as these technologies mature, they will become integral to value-based care operations, enabling providers to scale up care quality without scaling up cost.

(VB-8) ***Supplementing CEHRT: Flexibility for Innovation***. We appreciate that current CMS policy allows some flexibility in certified EHR technology (CEHRT) requirements for APMs, and we encourage even more flexibility to integrate **novel digital health tools (VB-8)**. For instance, a hospital's Conditions of Participation now require electronic patient event notifications, but ACOs could benefit from extending that requirement such that ACO care partners also receive those notifications (as the RFI notes). We recommend HHS consider broadening data-sharing mandates to cover APM entities – e.g. when a beneficiary in an ACO is admitted to any hospital, ensure the ACO's care team is notified (perhaps via the QHIN/TEFCA network or other intermediaries). This would empower our AI agents (or other tools) to react promptly, say by calling the patient post-discharge no matter where they were hospitalized. Additionally, CMS could allow substituting certain CEHRT functions with certified AI capabilities. If an APM can demonstrate that a certified AI agent performs a function (for example, patient education or care planning documentation) as effectively as a CEHRT module, that should "count" toward meeting technology use requirements. In practice, this might mean creating an alternative compliance pathway where non-EHR solutions that meet interoperability and safety criteria fulfill the tech requirement. This would prevent stifling innovation – ACOs wouldn't be forced to use only traditional EHR features if a more innovative solution exists. Overall, aligning HHS policies to support the use of cutting-edge digital health products in value-based models – through expanded

data exchange, safe harbors, and flexible tech standards – will better equip APM organizations to succeed.

Vision: Making America Healthy Again with Generative AI

In closing, Hippocratic AI envisions a future where **generative AI agents act as tireless healthcare extenders**, delivering personalized support to Medicare & Medicaid beneficiaries at scale. This technology can **create an era of healthcare abundance**: every senior who needs daily check-ins or coaching can get it, every care gap can be addressed promptly, and no one is left without guidance due to staff shortages or cost. Crucially, this is achieved safely and cost-effectively – our model has shown that an AI agent can operate at a fraction of human cost (around $0.71 per preventive outreach vs. ~$54.83 by traditional methods) while maintaining near-perfect accuracy. That translates to substantial savings for taxpayers and reinvestment in care. But beyond efficiency, it's about compassion and empowerment: these AI agents, imbued with empathy and knowledge, help patients navigate their health in a deeply supportive manner. They speak your language, literally and figuratively – from Spanish to Nepali, or from clinical jargon to plain language a layperson understands. They coach and encourage healthy behaviors, helping prevent chronic disease and illness through everyday choices. They extend the reach of physicians and nurses, so care teams can achieve more without burning out. In essence, generative AI, when developed with a Hippocratic commitment to "do no harm," becomes a powerful public health tool. It can ***prevent avoidable suffering, improve quality of life, and ensure every Medicare & Medicaid beneficiary has a helping hand in achieving their health goals.*** We urge CMS to be bold in fostering this innovation – through supportive policies, pilots, and collaborations – because the reward is a healthier America. Hippocratic AI stands ready to partner in this mission. Together, we can **Make America Healthy Again**, by harnessing safe, patient-centric AI to transform digital health for the better.

Thank you for considering our comments. We are excited about the potential of generative AI to advance CMS's goals and improve care for all patients. We welcome any further questions and look forward to ongoing collaboration.

Sincerely,

Hippocratic AI, Inc.

Contact:
W.B. "Mitch" Mitchell
CGO, Government Sector

☏ *: 650-305-3225*
✉ *: info@hippocraticai.com*

Sources Cited: https://www.wgal.com/article/pennsylvania-wellspan-health-ai-assistant-how-it-works-what-people-think/64845246 ; https://pmc.ncbi.nlm.nih.gov/articles/PMC8606361/#:~:text=The%20direct%20cost%20per%20patient,the%20absence%20of%20the%20intervention