

Public Comment on Request for Information (RFI)

Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care

Federal Register Citation: 2025-23641 (90 FR 60108)

Submitting Entity: EHY Consulting LLC

Submitted by: Edward You

Date: January 2026

Opening statement

This submission responds directly to each of the ten questions in the HHS Health Sector AI RFI and is structured to support cross-program review. Each response stands alone while also referencing related questions where system-level dependencies exist. The intent is not to recommend a single policy outcome, but to surface implementation realities, governance gaps, and second-order effects that emerge when AI systems are deployed within real clinical workflows.

A workflow-to-entity crosswalk is included to clarify how responsibilities intersect across HHS StaffDivs and OpDivs and to help avoid gaps between policy intent and operational execution.

EHY Consulting LLC respectfully submits these comments in alignment with the HHS Artificial Intelligence Strategy to support informed, non-prescriptive decision-making across the health sector. A consistent finding across AI deployments in clinical environments is that adoption and safety are determined less by model accuracy in controlled settings than by how AI systems interact with real-world clinical workflows, governance structures, data stewardship practices, procurement decisions, and learning feedback loops.

Clinical AI functions as a socio-technical system: it shapes clinician behavior, redistributes responsibility, and generates compounding learning from data produced through routine operations. Accordingly, this submission answers each RFI question directly while cross-referencing related responses to keep system-level implications visible and to inform approaches to scaling clinical AI responsibly.

HHS equities map (StaffDivs and OpDivs)

Clinical AI adoption is inherently cross-cutting. The same AI-enabled workflow can engage multiple authorities and equities across HHS. A practical way to avoid gaps is to map responsibilities to the workflow itself.

Workflow-to-entity crosswalk (illustrative)

Workflow / decision point	Primary HHS equities (examples)	What success looks like (implementation)
Clinical deployment and payment incentives	CMS (payment policy, quality measures, demonstrations)	Incentives reward measurable workflow value and outcomes, not superficial “AI presence”
Health IT, interoperability, data standards	ASTP/ONC (health IT standards, interoperability policy)	Standards enable safe integration and auditability; interoperability does not imply unconstrained secondary use
Privacy, patient rights, and enforcement	OCR (HIPAA privacy/security enforcement and related guidance)	Clear, enforceable boundaries on access, use, and disclosure; meaningful transparency and recourse
Cybersecurity and national-risk posture	CIO (enterprise security practices), ONS (national security risk and resilience processes)	Security-by-design requirements are enforceable; AI is treated as part of critical clinical infrastructure
International alignment, standards, and data-sharing regimes	OGA (international engagement, standards, agreements)	Policies reconcile domestic safeguards with cross-border standards and data-sharing realities
Contract enforcement and vendor/subcontractor controls	Office of Acquisition / contracting functions (across HHS)	RFPs and contracts require auditability, learning controls, and secure operations; flow-down to subcontractors
Safety of AI-enabled clinical products and decision support (where applicable)	FDA (device/SaMD, CDS policy)	Clear expectations for validation, change control, and post-market monitoring
Evidence generation and implementation science	AHRQ (evidence and outcomes research)	Evaluation frameworks reflect real-world workflow effects and heterogeneous

		settings
Public health integration (where AI crosses into surveillance)	CDC	Defined boundaries between clinical and public health uses; governance for secondary use
Research and biomedical data ecosystems (where clinical AI relies on research pipelines)	NIH	Strong stewardship of clinical-research interfaces and reuse pathways
Safety-net and underserved delivery settings	HRSA, IHS, SAMHSA	Adoption pathways account for resourcing constraints, equity, and operational realities

Cross-cutting note: acquisition is the control plane that converts policy intent into enforceable obligations. Interoperability, cloud use, vendor access, and model-update practices are ultimately determined by contract terms, monitoring rights, and operational controls.

Cross-cutting analytic lens: data, workflows, and learning

Practical governance test for clinical AI deployments

Use case and end-user clarity: identify the decision-maker, operator, and accountable party; specify the action taken because of the model output.

Data lifecycle boundaries: define what data enters, what leaves, retention periods, and what is used for care delivery versus model improvement.

Learning accumulation and sovereignty controls: specify whether operational data may be used for model improvement and under what constraints; treat derived artifacts (e.g., embeddings, fine-tunes, prompts, logs, telemetry) as governed outputs.

Security posture and access paths: enforce identity, least privilege, logging, incident response, and subcontractor controls; define cloud and remote access defaults with audit rights.

Bias introduction points and drift: identify where bias can enter (training data, workflow selection, labeling, clinician reliance, feedback loops) and require continuous monitoring and rollback.

Across multiple clinical domains (including high-impact areas such as precision oncology and advanced therapies), the most durable advantage—and the most consequential vulnerability—

often lies not in the AI model or the molecule, but in control over the learning system: the workflow-generated data, outcomes, failure modes, and process metadata that enable iterative optimization over time.

Therefore, it is useful to distinguish:

- Privacy: who can see my data.
- Sovereignty: who determines what my data means and how much it is worth.
- Security: who has access to my data and what they can do with it.

Many governance approaches address privacy and baseline security, but do not explicitly govern learning accumulation (how data becomes compounding capability). In clinical AI, the latter increasingly determines long-run trust, competitiveness, and resilience.

Responses to RFI questions (aligned to Federal Register ‘Specific Questions’)

Each response below restates the Federal Register question verbatim (or with minimal formatting) and provides a direct answer. Cross-references are included where useful for synthesis across HHS teams.

1. Biggest barriers to private sector innovation and adoption/use in clinical care

Clinical AI frequently crosses mismatched risk frameworks (privacy, cybersecurity, clinical safety, enterprise risk, and vendor product risk). These frameworks do not align by default, and gaps emerge at their seams unless governance is explicitly integrated.

Workflow misfit is the dominant barrier. AI systems that perform well in development can fail in practice when they add documentation burden, disrupt clinician cognitive flow, or create ambiguous responsibility at the point of care.

Key barriers include:

- Unclear accountability for AI-influenced decisions, especially when AI is adaptive or when vendor updates materially change behavior.
- Trust gaps driven by opacity, inconsistent performance across subpopulations, and unclear post-deployment monitoring.
- Data provenance and integration challenges (EHR heterogeneity, fragmented identifiers, incomplete documentation, variable labeling).
- Operational and security constraints (network segmentation, device integration, patch management, identity and access management).
- Procurement realities: contracts that do not specify auditability, model-update controls, data-use boundaries, or subcontractor constraints.

See also: Q2, Q3, Q4, Q7, Q8.

2. Regulatory, payment policy, or programmatic design changes; CFR citations where applicable

Condition participation or incentives on auditability (minimum logs, update provenance, and model-change notification).

Make data-use boundaries enforceable, clarifying whether operational data may be used for model improvement and under what authorizations.

Across regulation, payment policy, and program design, the most leverage comes from making governance and monitoring enforceable at the workflow level while avoiding incentives for superficial AI adoption.

Priority areas for HHS to consider include:

- CMS-aligned incentives that reward measurable workflow value and outcomes rather than mere AI integration (see Q6).
- Programmatic designs that require lifecycle governance (monitoring, drift management, change control) as a condition of participation in demonstrations, quality programs, or certifications (see Q4 and Q5).
- Clear, enforceable expectations for data stewardship, auditability, and vendor/subcontractor controls, implemented through acquisition and contracting requirements where feasible (see Q3 and Q5).
- Interoperability policies that distinguish data exchange for care delivery from secondary use for model training and learning accumulation (see Q8).

See also: If HHS seeks specific CFR citations, a targeted follow-on request scoped to particular programs (e.g., Medicare conditions of participation, quality reporting, certification programs) would enable more precise citations by commenters.

3. Non-medical devices: novel legal and implementation issues; HHS role

Governance and contracts should explicitly cover workflow metadata and derived artifacts (prompts, embeddings, fine-tunes, telemetry), which often represent the primary signals of value and intent rather than mere exhaust.

For non-medical device AI used in clinical care (e.g., workflow, triage, documentation, scheduling, or administrative decision support), key issues include liability allocation, indemnification, privacy/security obligations, and governance of adaptive behavior post-deployment.

Novel legal and implementation issues commonly arise from:

- Ambiguity in responsibility when AI influences clinician behavior or care pathways without being a regulated medical device.

- Contracting and indemnification gaps that fail to define change control, performance representations, or audit rights.
- Privacy/security frameworks that govern access and disclosure, but do not explicitly govern learning accumulation or model improvement using operational clinical data.
- Third-party and subcontractor risk (cloud hosting, analytics, support), especially when access paths are opaque.

See also: HHS can help by publishing model contract clauses and minimum auditability/lifecycle governance expectations usable by providers and payers, aligned across OCR, ASTP/ONC, CIO/ONS, and CMS.

4. Non-medical devices: evaluation methods pre- and post-deployment; mechanisms

Evaluate impact on clinician behavior, including automation bias, alert fatigue, de-skilling, and reliance shifts over time.

The most promising evaluation approaches combine pre-deployment validation with post-deployment monitoring that is sensitive to workflow effects, site-to-site variation, and distribution shift.

Recommended evaluation components include:

- Human-centered workflow evaluation (cognitive load, alert fatigue, override rates, time-to-task, downstream workflow changes).
- Robustness and stress testing across settings, subpopulations, and data quality conditions.
- Post-deployment drift detection and performance monitoring linked to outcomes where feasible.
- Change-control governance for model updates, including criteria for rollback.

See also: HHS can support these via contracts, grants, cooperative agreements, and prize competitions (often with AHRQ/NIH/CMS partners), prioritizing monitoring, auditability, and workflow-aware evaluation methods.

5. Support for accreditation, certification, testing, and credentialing

Establish a minimum governance checklist for certification programs, including update control, monitoring, disclosure of data use for training, and incident response.

HHS can amplify private-sector accreditation, certification, industry-driven testing, and credentialing by helping establish baseline expectations for trustworthy clinical AI that are measurable and auditable.

Support could include:

- Reference expectations for auditability (logging, traceability of inputs/outputs, update history) and minimum cybersecurity posture for clinical AI deployments.
- Credentialing frameworks for AI governance roles within health care organizations (see Q7).
- Model contract language and procurement checklists that incorporate security-by-design and lifecycle monitoring requirements (see Q3 and Q4).
- Public-private benchmarking and validation resources, including shared evaluation protocols and testbeds (see Q8 and Q10).

6. Where AI met/exceeded vs fell short; novel tools with greatest potential

Distinguish value created for patients and providers from value transferred to vendors through learning accumulation from operational data.

AI tends to meet or exceed expectations when it is tightly coupled to well-defined workflows with clear success metrics (e.g., reducing administrative burden, improving triage consistency, accelerating access, or reducing preventable harm). It tends to fall short when deployed as a generic overlay without governance, monitoring, or integration into decision rights.

High-potential tool categories include:

- Workflow-aware clinical documentation support that measurably reduces administrative burden without introducing silent errors.
- Risk stratification and triage tools with clear accountability and monitoring, especially in resource-constrained settings.
- Tools that surface quality insights by analyzing workflow metadata and outcomes (not only clinical text).
- Patient-facing navigation and care coordination supports that improve adherence and reduce friction across care settings (see Q9).

See also: To accelerate learning on 'what works,' HHS can support standardized reporting and privacy-preserving benchmarking (see Q10).

7. Organizational roles and administrative hurdles for adoption

Adoption decisions are typically shaped by clinical leadership, operational leadership, compliance/privacy, IT/security leadership, and (increasingly) AI governance committees or enterprise risk management functions.

Primary administrative hurdles include:

- Unclear ownership for AI lifecycle management and monitoring across clinical, IT, and compliance functions.
- Procurement bottlenecks and contract limitations that impede auditability or change control.
- Data access and interoperability friction across EHRs and ancillary systems (see Q8).
- Training, credentialing, and change management requirements for clinicians and staff.

- Security approvals and network integration constraints, especially where tools require cloud connectivity.

See also: This is why HHS-wide coordination across CMS, ASTP/ONC, OCR, CIO/ONS, OGA, and acquisition functions is essential.

8. Interoperability opportunities; data types, standards, and benchmarking tools

Interoperability expands the attack surface and secondary-use risk unless paired with access control, provenance, and enforceable training boundaries.

Enhanced interoperability can widen market opportunities and accelerate AI where it improves data provenance, reduces integration friction, and enables benchmarking—while still respecting governance boundaries on secondary use and learning accumulation.

Priority opportunities include:

- Standardized representations of clinical workflows and outcomes (including provenance/lineage metadata).
- Consistent patient identity and record linkage across care settings.
- Data standards and benchmarking tools that support monitoring and post-deployment evaluation (see Q4).
- Mechanisms for privacy-preserving benchmarking and federated evaluation where direct pooling is impractical.

See also: Interoperability policy should explicitly distinguish exchange for care delivery from secondary use for model training and learning accumulation.

9. Patient/caregiver desired challenges and concerns

Patients increasingly want clarity not only on who sees their data, but who benefits from derived learning and what recourse exists.

Patients and caregivers often prioritize improvements that reduce friction, shorten wait times, improve communication, and increase transparency about decisions affecting care access and quality.

Common concerns include:

- Privacy and security of sensitive health information, including secondary uses beyond direct care.
- Unclear accountability when AI influences clinical decisions or access to services.
- Bias and inequitable treatment, especially when monitoring and recourse mechanisms are weak.

- Reduced human agency or difficulty obtaining explanations, corrections, and escalation pathways.

See also: These concerns reinforce the need for lifecycle governance, auditability, and clear decision rights (see Q3–Q5).

10. AI research priorities; literature on impact and costs/benefits/transfers

Prioritize methods to measure learning accumulation and value transfer across stakeholders, and governance models that integrate privacy, sovereignty, and security operationally.

HHS research priorities that would accelerate responsible adoption include implementation science for clinical AI, workflow-aware evaluation methods, and scalable approaches to monitoring and governance.

Priority research areas include:

- Methods for post-deployment monitoring, drift detection, and change control in heterogeneous clinical settings.
- Human-centered evaluation methods that quantify workflow impact and clinician reliance over time.
- Privacy-preserving benchmarking, federated evaluation, and secure data-sharing architectures.
- Governance models for adaptive systems that integrate privacy, sovereignty, and security considerations.

See also: HHS can help curate and translate literature into practical evaluation playbooks and clarify how costs/benefits/value transfer are measured across providers, payers, vendors, and patients.

Closing

EHY Consulting LLC appreciates the opportunity to contribute to this RFI and supports HHS's efforts to advance responsible, effective, and trustworthy adoption of AI in clinical care. These comments are intended to be practical and implementable, highlighting workflow-level realities and second-order effects so that accelerated adoption does not outpace governance of safety, trust, and resilience. We welcome continued dialogue to clarify implementation challenges and governance considerations as policy and practice evolve.