

Response to Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care Docket ID: HHS-ONC-2026-0001-0001 (RIN 0955-AA13)
Submitted by: Pal Randhawa, MD, EMBA, FAANS **Email:** prandhawa@gmail.com **Capacity:** Personal Capacity **Date:** January 19, 2026

Submitter Background I am a board-certified neurosurgeon with approximately 15 years of frontline clinical experience across emergency departments, operating rooms, and intensive care units, including sustained work in Level I trauma and referral centers. In addition to my clinical role, I hold an Executive MBA (EMBA) with a focus on healthcare operations and infrastructure resilience. This dual perspective allows for the integration of tactical clinical realities with the strategic resource management required for national health security.

My experience includes firsthand observation of workforce strain, capacity degradation, and coordination failures under stress, as well as participation in collaborative and advisory efforts related to disaster preparedness and national critical infrastructure discussions.

These comments are submitted in my personal capacity. The views expressed are my own and do not represent any employer, institution, or advisory body.

Executive Summary: The Case for "Deflationary Infrastructure" The U.S. Healthcare Sector operates beyond safe margins, with the AAMC projecting a deficit of up to 86,000 physicians by 2036. This submission argues that Artificial Intelligence should be prioritized as a "Deflationary Infrastructure"—a protective capability designed to preserve the system's most expensive and fragile asset: its staffed clinical capacity.

Proposed Doctrinal Standard: We propose that HHS adopt **Workforce Readiness Telemetry (WRT)** as a critical interoperability standard.

Definition: Workforce Readiness Telemetry refers to anonymized, aggregate indicators of operational strain, cognitive load, and staffing fragility. WRT is intended for system-level planning and readiness assessment and is **explicitly not designed** for individual performance evaluation, disciplinary action, or employment decisions. It is intended to complement existing interoperability frameworks (e.g., USCDI, TEFCA), not replace them—focusing on upstream operational inputs that precede clinical events.

II. Responses to Specific RFI Questions

Question 1: Barriers to Private Sector Innovation and Adoption BLUF: Adoption is throttled by a "Readiness Bind." Hospitals lack the cognitive bandwidth to adopt tools that do not immediately reduce operational friction.

- **The Gap:** HHS currently lacks a standardized metric for "Human Capacity." Staffing is measured in "FTEs" (Full-Time Equivalents), a lagging indicator that fails to capture

burnout, moral injury, or disengagement. A facility may appear fully staffed on paper, yet remain operationally degraded due to unmeasured cognitive load.

- **The Solution:** Prioritize AI that serves as **Operational Radar**. Tools that predict attrition 30–60 days in advance allow administrators to intervene proactively, reducing the multi-million dollar annual cost of turnover. Importantly, WRT leverages existing metadata streams and minimal viable indicators, making it **low-burden** for frontline systems.
 - *Strategic Alignment:* This directly operationalizes NHSS 2023-2026 Objective 1.3 ("Strengthen the recruitment, retention, and preparedness of the HPH workforce"), moving workforce retention from a passive goal to an active, measured capability.
 - *Data Support:* While nursing turnover costs the average hospital \$5.7M annually, the loss of a single high-acuity physician costs \$500,000 to >\$1M in replacement costs alone, consistent with recent AMA and AMN Healthcare analyses.
 - *Revenue Impact:* For high-revenue specialties (e.g., Neurosurgery), a six-month vacancy results in >\$3M in lost clinical capacity. AI that preserves this capacity represents a critical form of deflationary infrastructure.

Question 3: Governance and Implementation of Non-Medical Device AI BLUF: Workflow-embedded AI falls into a "Liability Gap" that chills adoption. Frontline clinicians currently bear a disproportionate share of the risk for AI-influenced workflows.

- **The Gap:** FDA regulation covers devices; it does not cover "Algorithm-Augmented Workflow." In practice, frontline clinicians often shoulder the ambiguity and accountability for AI-influenced workflows—creating adoption hesitancy in high-stakes care.
- **Recommendation:** HHS's internal AI Governance Board should advocate for the exploration of **time-limited federal safe-harbor mechanisms** during declared staffing emergencies. If a clinician follows a federally validated AI protocol during a crisis, they should enjoy liability protection. This mechanism is essential for deploying AI in high-stakes environments without paralyzing the workforce.

Question 8: Interoperability to Accelerate AI Adoption BLUF: Interoperability must expand beyond Clinical Data (USCDI) to include Operational Telemetry.

- **The Gap:** Current interoperability standards track the patient but ignore the system.
- **The Solution:** Create a standard for sharing "Aggregate Readiness Signals" (e.g., Bed Capacity + Staffing Viability).
- **Strategic Outcome:** This enables **Regional and National Surge Prediction**. Beyond daily operations, aggregated telemetry creates a real-time heat map for mass casualty, pandemic, or geopolitical events (aligning with NDAA Section 735 readiness requirements). This allows federal planners to load-balance across the national grid before regional collapse occurs.
- **The ONC Role:** The ONC is uniquely positioned to define this operational layer in a way that preserves trust, minimizes burden, and enables downstream innovation, ensuring that readiness data becomes a national asset rather than a proprietary silo.

Question 10: Prioritized Areas of Research BLUF: Research must validate AI as a tool for Cognitive Augmentation, not just automation.

- **Priority 1: Cognitive Load Mapping.** The primary safety endpoint for clinical AI should be: Does this tool reduce time-to-decision or increase cognitive burden?
- **Priority 2: Degraded-State Validation.** AI tools must be tested for reliability under "Degraded Conditions" (e.g., cyber outages, communications latency, mass casualty surges, or geopolitical conflict). This aligns with CISA's 2025 Cybersecurity Performance Goals for critical infrastructure resilience.

Conclusion: The Aviation Parallels Successful high-reliability sectors (Aviation, Energy) do not use technology to replace the expert operator; they use it to augment situational awareness and prevent error. Healthcare AI should follow this "Human-System Teaming" doctrine. By shifting focus from "replacing doctors" to "stabilizing the grid," HHS can transform AI from a disruptive novelty into a stabilizing component of national health security.

References

1. ASPR (2023). National Health Security Strategy (NHSS) 2023-2026.
2. American Medical Association (2025). Organizational Cost of Physician Burnout.
3. NSI Nursing Solutions, Inc. (2025). 2025 National Health Care Retention & RN Staffing Report.
4. AMN Healthcare (2025). The Cost of Physician Turnover and Economic Impact.
5. The White House (2024). National Security Memorandum on Critical Infrastructure Security and Resilience (NSM-22).
6. CISA (2025). Cybersecurity Performance Goals (CPG 2.0).
7. Association of American Medical Colleges (2024). The Complexities of Physician Supply and Demand.