

Comment on “HHS Health Sector AI RFI”

Department of Health and Human Services

Office of the Secretary

Office of the Deputy Secretary and Assistant Secretary for Technology Policy and Office of the
National Coordinator for Health Information Technology

RIN: 0955-AA13

Federal Register Document Number: 2025-23641

Docket ID: To be assigned by Regulations.gov

Submitted by: Interstella

Date: January 21, 2026

Introduction

Interstella appreciates the opportunity to respond to the Department of Health and Human Services (HHS) Request for Information on accelerating the adoption and use of artificial intelligence (AI) in clinical care. As a health data and technology company focused on real-time data integration and AI readiness, we share HHS's vision of a "forward-leaning, industry-supportive, and secure" approach to AI in healthcare.

We believe AI's transformative potential in clinical care can only be realized by addressing foundational challenges in data infrastructure, governance, and incentives. In this response, we detail barriers and provide recommendations across regulation, reimbursement, and research & development (R&D), supplemented by high-impact use cases and policy suggestions. Our comments emphasize that AI systems fail or succeed based largely on the data that powers them, and thus a modernized data ecosystem is critical to safe and effective AI adoption in healthcare.

Interstella Background: Interstella is dedicated to turning raw, fragmented health information into real-time, AI-ready intelligence. Our experience developing the *Lynqsys* data platform and *Data Refinery as a Service (DRaaS)* solutions gives us practical insight into the data quality, interoperability, and timeliness issues that currently limit AI's impact in clinical settings.

We approach this RFI from the perspective of those "*building, buying, evaluating, and using*" AI tools for care delivery, and those who wish to do so but face barriers (as noted by HHS). We offer both policy recommendations and use case-driven examples to illustrate how HHS can accelerate AI adoption while upholding patient safety, equity, and trust.

Barriers to AI Adoption in Clinical Care

Data Readiness Gap: The core barrier we observe is that AI innovation is outpacing the healthcare system's data readiness. Federal initiatives increasingly rely on advanced analytics and AI, yet most healthcare data ecosystems operate with significant limitations:

- **Delayed Data Availability:** Clinical data are often ingested and aggregated in batch processes with lag times of days or weeks, rather than in real-time streams. This delay stifles AI tools that require up-to-date information for decision support or predictive analytics. For example, quality measures or care gap alerts powered by AI cannot function optimally if data arrives long after the point of care.
- **Inconsistent Data Quality & Provenance:** Healthcare data remains "fragmented, biased, and often a mess", leading to noise, duplication, and gaps that undermine AI algorithms. Data may come from diverse sources (EHRs, claims, devices, patient-reported inputs) with inconsistent formats and unknown provenance, making it hard for AI models to trust and interpret.

Most failed AI initiatives are not caused by faulty algorithms or lack of vision, they fail because the data foundation is not fit for purpose. Indeed, studies estimate up to 85% of

AI projects fail due to poor data quality or volume issues. In healthcare, this problem is acute: “Without large, high-quality data sets, it can be difficult to build useful AIs,” and providers are slower to adopt AI when data are incomplete or siloed.

- **Limited Support for Unstructured and Behavioral Data:** Many current systems handle structured data (e.g. diagnoses, lab results) but have limited capability to incorporate unstructured data (clinical notes, images) or behavioral and social health data. This is a barrier for AI, which thrives on rich data contexts.

Critical information about a patient’s mental health, social determinants, or care preferences may be documented in text or external systems and thus invisible to AI-driven tools. The lack of integration for behavioral health data, for instance, makes it difficult to apply AI for outcome measurement in mental health or to coordinate holistic care in models like Certified Community Behavioral Health Clinics (CCBHCs).

- **Post-Hoc Governance vs Inline Controls:** Data governance (privacy protections, quality checks, consent management) is often applied after data aggregation or as a periodic audit, rather than inline at the point of data creation and exchange. This reactive approach means potential issues, from data errors to privacy risks, are caught late.

AI systems require trusted, well-characterized data streams; without proactive governance baked into data pipelines, organizations either inhibit data sharing or expose themselves to compliance risks. In short, data governance that isn’t real-time and automated becomes a bottleneck to scalable AI.

- **Data Pipelines Not Designed for AI Consumption:** Traditional health IT systems were built for billing or record-keeping, not for feeding machine learning models. Data may be locked in PDF reports, faxed forms, or legacy databases that AI cannot readily parse.

Even within modern EHRs, data exchange via interfaces or reports can miss context (e.g. relationships between data points) that AI algorithms need. Interoperability challenges exacerbate this, “*Electronic Health Record systems are largely not compatible across different hospitals or providers,*” leading to localized data rather than a comprehensive patient view. AI models trained on narrow, siloed data may not generalize, limiting their utility.

These data-related barriers create technical hurdles and raise regulatory, equity, and trust risks. An AI tool is only as good as the data it learns from; if those data are incomplete, biased, or delayed, the tool’s recommendations could be unsafe or inequitable. In sum, “AI systems do not fail because of algorithms; they fail because of data.”

Addressing data readiness is foundational to everything else, from ensuring algorithmic fairness to convincing healthcare professionals and patients to trust AI. We encourage HHS to treat data quality, interoperability, and timeliness as strategic priorities in all AI-related policies (as HHS’s own AI Strategy recognizes).

Other Key Barriers: In addition to data readiness, we echo the other barriers identified by experts and prior HHS inquiries:

- **Lack of Trust and Transparency:** Many AI algorithms, especially complex machine-learning models, are “black boxes” that clinicians find hard to interpret. The opacity of AI decisions can erode provider trust and willingness to use AI in care. Clinicians remain ultimately accountable for patient outcomes, so they are cautious about relying on tools they don’t fully understand.

Without assurances of transparency or robust validation, frontline providers often default to familiar (non-AI) solutions.

Building trust requires both *technological solutions* (e.g. explainable AI, rigorous validation) and *cultural change* (education, success stories, engagement of clinicians in AI development). We note that transparency, privacy protections, and rigorous evaluation have been identified as essential elements of trustworthy AI in healthcare.

- **Regulatory Uncertainty:** Current regulations have not fully caught up with AI, creating uncertainty. Privacy rules (like HIPAA) and data use policies can make it difficult to aggregate and share the data needed to develop robust AI models.

Approval pathways for AI, particularly software not regulated as a medical device, are unclear; developers are unsure how to demonstrate safety/effectiveness, and providers are unsure what oversight or certifications to look for.

Furthermore, liability concerns loom large: healthcare organizations worry that if an AI makes a faulty recommendation, they will face legal risk with uncertain liability-sharing between developers and users. This “chilling effect” causes risk-averse stances, slowing adoption. We explore regulatory solutions in the next section.

- **Misaligned Incentives in Payment:** The fee-for-service reimbursement model often fails to reward (or even penalizes) innovations that improve efficiency or quality. AI tools that prevent hospitalizations or automate routine tasks might reduce billable services in the short term, disincentivizing adoption by providers under traditional payment.

Meanwhile, those bearing long-term risk (payers or ACOs) have incentives to use AI for prevention and cost reduction, but they face their own hurdles in getting such tools approved and reimbursed.

This misalignment leads to patchy adoption, e.g. larger systems or value-based organizations invest in AI, while others do not. As HHS notes, “*legacy payment systems*” can diminish the promise of AI through inertia. We discuss payment reforms below.

- **Workforce and Workflow Integration:** Successful AI adoption requires complementary process changes and training. Introducing an AI-driven tool means little

if clinical workflows are not adjusted to incorporate the AI's output at the right time and place.

Many hospitals lack personnel with AI expertise (data scientists, informaticists) to champion these changes. Additionally, some clinicians fear AI could replace or de-skill their roles (for example, radiologists hearing that AI will read images). This can create resistance or lack of engagement with new tools.

Even when leadership is supportive, *administrative hurdles* like lengthy procurement, integration costs, or lack of IT support can slow down AI deployments. In our experience, the roles most influential in AI adoption include the Chief Medical Information Officer (CMIO), Chief Data or Digital Officer, and clinical department heads who see a clear need.

Hospitals with professional management and integrated salaried structures tend to be more proactive in deploying AI than those with decentralized or physician-led management. HHS should consider how to educate and empower these decision-makers to be AI champions.

In summary, the biggest barriers to private-sector innovation in health AI are *data-related limitations, regulatory uncertainty (including privacy and liability issues), misaligned payment incentives, and lack of trust*. Overcoming these requires coordinated action.

We recommend that HHS focus on improving the health data foundation as a first-class objective, because doing so will mitigate many other barriers (improving data sharing alleviates privacy and access issues, high-quality data enables better algorithms, etc.).

Below, we detail specific regulatory, reimbursement, and R&D actions that can accelerate AI adoption while safeguarding patients. We also provide *example use cases* illustrating how these changes can yield tangible improvements in care delivery.

Regulatory and Governance Recommendations

HHS's role as the nation's principal health regulator is pivotal in shaping an environment where AI can innovate rapidly *and* safely. We agree with HHS's goal of a regulatory posture on AI that is predictable, proportionate to risk, and protective of patients' rights and privacy. To achieve this, we recommend:

1. Clarify AI Oversight for Non-Device Software: Many AI tools in clinical care do not meet FDA's definition of a medical device (for example, clinician-facing decision support that allows independent review, or operational AI for scheduling and resource allocation).

This creates a gray area with "*novel legal and implementation issues*" around accountability, as noted in the RFI. HHS (through ONC, CMS, OCR, etc.) should provide guidance on expectations for these non-FDA-regulated AI tools. Specifically:

- **Liability and Indemnification:** We urge HHS to convene stakeholders (providers, developers, malpractice insurers) to develop guidance or safe harbor principles for AI use. Clinicians should know when they can rely on an AI tool and how liability is apportioned if the AI is flawed.

One approach is a “qualified AI” safe harbor; if an AI tool meets certain HHS-endorsed standards (for transparency, validation, etc.), then providers who use it as intended could have mitigated liability. This would encourage use of vetted tools. Conversely, developers of such tools might be expected to carry liability insurance or indemnify providers in cases of egregious failure, creating accountability on both sides.

Clear rules on “*who is liable if something goes wrong*” would likely increase adoption. HHS could work with Congress if needed to adjust liability frameworks for AI in healthcare (e.g., exploration of federal Tort Claims Act extensions or other mechanisms for AI-related harms).

- **Privacy and Data Use Clarity:** HHS should update HIPAA and related guidance to explicitly address AI development and deployment. For instance, clarify how providers/payers can share de-identified or synthetic data for AI training without running afoul of privacy rules.

Consider promoting data trusts or safe data enclaves for AI development, perhaps through OCR guidance that such arrangements, with proper safeguards, are encouraged. HHS’s AI Strategy already emphasizes “*fostering data standards and datasets to bolster their usability for AI... while maintaining individual privacy*”.

Building on this, HHS could create model frameworks for data stewardship that balance innovation with privacy (e.g., datasets that are rich enough for AI but managed under strong governance and audit). Clear guidelines can reduce the fear that sharing data for AI equals a privacy violation.

- **Security and Compliance:** Many non-device AI tools still interface with protected health information. HHS can help by streamlining compliance. For example, template Business Associate Agreements for AI vendors, guidance on using cloud-based AI under HIPAA, and mapping AI practices to existing regulations (like how to apply 42 CFR Part 2 for substance abuse data in AI or ensuring ADA compliance in AI-driven communications). By addressing these now, HHS preempts future confusion as AI use grows.

2. Modernize Clinical Decision Support (CDS) Regulation: Current FDA guidance (per the 21st Century Cures Act and FDA’s CDS guidance) draws a line between regulated “medical device” software and unregulated CDS. HHS should ensure that developers and providers clearly understand this line.

We recommend HHS (through FDA collaboration and ONC) establish a central information hub

on what kinds of AI applications require FDA clearance and which do not. This could include an AI Tool Classification Framework that maps common use cases: e.g., diagnostic algorithms (regulated SaMD), vs. workflow optimizers (unregulated), vs. patient self-care chatbots (gray area). By providing case studies or decision trees, HHS can reduce uncertainty.

For those AI tools that *are* medical devices, HHS should support FDA in streamlining approval without compromising safety. For instance, expanding programs for predetermined change control plans (PCCPs) and real-world performance monitoring to allow AI models to continuously improve post-approval.

For non-device tools, HHS could work with industry on voluntary certification (see below) instead of regulation, focusing on transparency and oversight at the user-organization level (hospital governance).

3. Support AI Governance Frameworks in Healthcare Organizations: Regulation is not just about federal rules; it's also about internal governance where AI is deployed. We applaud efforts by the private sector, such as the Joint Commission's new guidance (2025) on the responsible use of AI in healthcare.

That framework, developed with the Coalition for Health AI, outlines seven core elements (from formal AI governance committees to patient transparency and bias monitoring) for safe AI implementation.

HHS should endorse and build upon such frameworks, perhaps by integrating them into ONC certification or CMS conditions of participation in the future. Specifically, HHS could:

- Coordinate with the Joint Commission and others to operationalize AI governance playbooks and create voluntary certification programs for health systems. For example, HHS could fund the development of a "Trustworthy AI in Clinical Care" accreditation that hospitals can obtain by meeting best practices (similar to how HIMSS stages measure EMR adoption).

This aligns with the RFI's query about supporting private sector accreditation and credentialing (Question 5). Federal backing would lend credibility and spur adoption of these standards.

- Issue guidance or requirements for federal healthcare programs (Medicare, Medicaid) to ensure participating providers have basic AI governance. This might start as recommendations and evolve into conditions for certain innovation models.

Areas to cover: having an AI oversight committee, conducting algorithmic bias assessments, educating staff and informing patients about AI use (the Joint Commission calls for patient notification and consent when AI directly impacts care).

Such governance expectations, even if not yet mandated, signal the industry to develop AI responsibly.

4. Adjust and Harmonize Regulations to Remove Unintended Barriers: We encourage HHS to review specific regulations that, while well-intentioned, may hinder AI adoption. Some examples and suggestions:

- **Revisit Data Blocking and Sharing Rules:** The 21st Century Cures Act information blocking rules promote data sharing, which is great for AI training and deployment. HHS should ensure these rules explicitly allow sharing data with AI developers or between providers for AI purposes (with proper patient authorization or de-identification).

Also, consider expanding required data elements (USCDI) to include more data types useful for AI, like imaging, free-text clinical notes, and patient-generated data. *Enhanced interoperability of specific data types (imaging, claims, genomics, social determinants) would “widen market opportunities, fuel research, and accelerate AI development,”* as the RFI posits.

For instance, better access to claims data could allow AI to identify high-risk patients across the continuum of care; richer social and behavioral data standards would enable AI to personalize care plans.

We note CMS’s recent Interoperability Framework as a positive step: it will require networks to share clinical and claims data via standardized APIs (FHIR) within 24 hours, dramatically improving timely data availability.

HHS should continue to remove technical and economic disincentives to interoperability, making near-real-time data exchange the norm in healthcare. This directly addresses the current batch-processing paradigm and will empower AI applications that need up-to-date data (e.g., sepsis early warning systems, dynamic care coordination tools).

- **HIPAA Safe Harbors for AI Development:** As mentioned, clarify pathways for using patient data to develop AI. HHS could create a safe harbor under HIPAA for uses of *fully de-identified data* (perhaps certified by a third party) in bona fide research and development of AI, even if the resulting tool will be commercial.

Today, some covered entities are overly conservative, fearing that any data sharing might violate HIPAA. Explicit encouragement and examples (e.g., “It is permissible to share de-identified EHR datasets with an AI developer for algorithm training, under XYZ conditions”) in OCR guidance would be helpful.

Also, update CFR 42 Part 2 in partnership with SAMHSA to align substance use data sharing rules with modern AI-driven care coordination (recent law changes are moving in this direction). AI that coordinates behavioral health care (like predicting relapse risk) can only be effective if data can be shared within care teams under clear, privacy-

protected frameworks.

- **Streamline IRB and Research Regulations for AI Studies:** Many healthcare AI tools need continuous learning and improvement, which blurs the line between quality improvement and research. HHS should consider guidance on when AI development/validation constitutes human subjects research requiring IRB oversight.

Overly burdensome interpretations could slow down needed real-world evaluation. Developing central IRB models or data commons for AI could accelerate evaluation while maintaining ethics.

- **Enable Sandbox or Pilot Programs:** HHS might use its authority (through CMS Innovation Center or ONC) to create regulatory sandboxes where healthcare organizations can pilot AI solutions with certain regulatory flexibilities.

For example, a hospital system could apply to test an AI tool that uses novel data integration, under a time-limited program where HHS waives certain documentation or reporting rules, provided strong monitoring is in place. Such sandboxes, coupled with rigorous evaluation, can demonstrate safe ways to adjust regulations in the long run.

In summary, HHS should reduce uncertainty by providing clear rules of the road for AI. This includes clarifying data sharing permissions, aligning AI oversight with risk levels, and promoting internal governance standards.

By revisiting privacy, security, and liability issues in light of AI, HHS can create a predictable environment that “enables rapid innovation while protecting patients”. We believe these regulatory adjustments, focused especially on data and trust, will unlock innovation by giving both AI developers and healthcare users confidence to move forward.

Payment Policy and Reimbursement Incentives

Payment policy can drive or stall the adoption of AI in clinical care. We agree with HHS’s assessment that current reimbursement systems contain “*inherent flaws*” that may unintentionally discourage high-value innovations.

To ensure AI’s promises are realized, HHS should modernize payment models to reward effective AI use, reduce financial friction, and align incentives across stakeholders. Our recommendations:

1. Create Clear Reimbursement Pathways for AI-Enabled Services: One of the “biggest barriers” is the lack of clear payment mechanisms for AI-powered interventions.

If a clinician uses an AI tool to, say, analyze a radiology image or predict a patient’s deterioration, there is often no distinct billing code or reimbursement for that activity (especially if it replaces a conventional service). We recommend:

- **Establish CPT codes or modifiers for AI-assisted services.** Engage the AMA CPT process to create codes for procedures where AI is used. For example, a radiologist using FDA-cleared AI for image analysis could append a modifier indicating AI assistance, which could be reimbursed with a small add-on payment for the validated efficiency or accuracy gain.

CMS could then decide to cover these codes. Early examples in this direction include CPT codes for autonomous AI detection of diabetic retinopathy. HHS should accelerate and expand such coding to cover various AI scenarios (AI-assisted pathology, AI-driven care coordination, etc.), signaling to providers that they can get reimbursed when using proven AI tools.

- **Bundle AI into Value-Based Payments.** Under alternative payment models (APMs) like accountable care or Medicare Advantage, the business case for AI is stronger (preventing a hospitalization saves money that the provider-network can keep).

HHS should integrate AI adoption into value-based programs by, for instance, providing allowances or bonuses for AI investments that demonstrably improve outcomes. CMS Innovation Center could pilot contracts where participants who deploy AI for specific high-impact use cases (e.g., reducing readmissions through predictive analytics) receive shared savings or upfront funding to do so.

The key is to tie AI to outcomes: if it lowers costs or improves quality, allow those benefits to translate into financial rewards for the adopters.

- **Coverage with Evidence Development (CED) for AI.** For certain novel AI technologies (especially those that are high-cost or high-impact), CMS could use a CED approach – cover the AI tool for Medicare beneficiaries conditionally while data is gathered on its effectiveness.

This has been used for some devices and could be applied, say, to AI-based diagnostic software or digital therapeutics. It encourages early adoption in a controlled way and speeds the evidence generation needed for broader coverage.

2. Align Payment Incentives with Quality and Prevention: We note that many AI tools promise long-term savings by improving prevention or care management, but under fee-for-service those savings are not realized by the provider.

HHS should continue shifting incentives from volume to value, which inherently makes AI more attractive. Some specific ideas:

- **Incorporate AI-driven improvements into quality measures and payment programs.** For example, if an AI tool can document care more thoroughly or identify care gaps in real time, that should help providers score higher on quality metrics (like HEDIS measures or Medicare Star ratings). CMS can explicitly allow AI documentation or

predictions to count toward meeting measure thresholds, as long as they meet validation criteria. By tying AI to quality bonuses or penalties avoidance, providers have a financial reason to adopt.

- **Modernize Medicaid and Medicare coverage policies.** Currently, a barrier in Medicaid is uncertainty about whether costs for AI software or services are reimbursable (states often lack guidance). CMS could issue guidance to state Medicaid programs on classifying AI expenses as allowable administrative costs or medical expenses.

For instance, clarifying that subscription fees for an AI-powered care management platform can be counted in medical loss ratio (MLR) calculations for managed care, or offering federal match for state investments in AI pilots.

A recent analysis noted that Medicaid’s hesitancy to fund innovations, “*even when cost-saving potential is evident,*” is a challenge, and lack of financial incentive discourages state investment in AI.

Clear CMS policy and perhaps small grant programs could nudge states to experiment with AI that improves care for vulnerable populations.

- **Address Unintended Consequences:** While encouraging AI, HHS must also guard against potential overuse or fraud (e.g., unnecessary tests due to AI alerts). Payment policy can help by requiring evidence of value.

For example, new AI-related billing codes could initially be reimbursed only in the context of approved registries or when providers adhere to appropriate use criteria. This ensures we reward AI that truly adds value. On the flip side, AI can be a tool *to reduce* fraud and waste, as CMS’s WISer initiative suggests.

We support using AI for program integrity (identifying improper claims, etc.), which ultimately frees resources for patient care.

3. Foster Competition and Accessibility of AI Tools: To avoid a scenario where only large health systems can afford AI, HHS should leverage reimbursement levers to promote competition and lower costs:

- **Ensure AI Pricing Transparency and Fair Pricing Models:** If AI tools become part of care, their cost (often via software licenses) can impact overall spending. HHS could work with stakeholders to develop pricing guidelines or reference pricing for certain AI-enabled services.

For example, if a particular AI diagnostic test is covered, CMS might set the payment rate based on value delivered rather than what the vendor charges. This prevents “massive spending bubbles on concentrated items” that don’t reflect value.

- **Coverage of AI as an enabler rather than a separate service:** Not every AI needs its own payment; some should be considered part of doing business in a digital health system. One approach is to allow AI software costs to be included in Medicare’s reimbursement formulas (like in hospital inpatient rates or practice expense calculations for physicians).

If AI reduces labor or improves outcomes, hospitals and practices could be paid slightly more for adopting it, under the logic that it’s a needed investment, but over time as it becomes standard, it’s just part of the normal payment.

This incremental inclusion avoids big payment cliffs and keeps competition open (providers choose the AI vendor, but CMS just pays for the function).

- **Encourage Open-Source and Public-Utility AI through Grants:** Reimbursement isn’t the only way to finance AI. HHS, via research grants or public-private partnerships, can support development of open algorithms (for instance, an AI model for predicting hospital readmission that any hospital can use free of charge).

If successful, this reduces dependence on high-cost proprietary solutions and pressures the market to compete on quality and support services rather than secret algorithms. We discuss R&D support more in the next section, but it’s worth noting here as it ties to affordability.

In short, payment modernization is essential. HHS should “ensure payers have the incentive and ability to promote access to high-value AI clinical interventions”. By carving out payment pathways for AI, aligning incentives with outcomes, and funding innovation, HHS can overcome the inertia of legacy payment systems.

This will signal to the market that if AI demonstrably improves care, it will be rewarded, a powerful motivator for both developers and healthcare providers.

Research & Development Priorities and Data Infrastructure

Accelerating AI adoption will require not only rule changes and payment tweaks, but also investment in R&D and infrastructure to integrate AI into care delivery. HHS oversees one of the world’s largest health research ecosystems, and we see opportunities for HHS to catalyze innovation through targeted research funding, partnerships, and demonstration projects.

Our key suggestions:

1. Invest in Data Modernization for AI: The healthcare data infrastructure must be upgraded to support AI at scale. We urge HHS to prioritize R&D in data architectures that deliver “better data, faster, with intelligence and trust.” This includes:

- **Real-Time Data Streams:** Support development of technologies and standards for real-time health data exchange. For example, HHS could fund pilots where health systems implement streaming APIs (perhaps leveraging FHIR subscriptions or event messaging) for critical data like hospital admissions, lab results, or vital signs to enable immediate AI analysis.

Our experience with Interstella's Lynqsys platform shows that near-real-time data refinement can greatly enhance AI accuracy in use cases like predictive quality measures.

A Medicaid pilot might demonstrate, say, how streaming data from managed care plans and providers to a state analytics platform yields faster care interventions (such as identifying a high-risk pregnancy for outreach). Results from such pilots can guide standards and best practices for real-time health information flows.

- **Data Standardization and Integration:** R&D should focus on tools to normalize and curate diverse data types for AI consumption. This could involve expanding standard vocabularies to cover behavioral health, social factors, device data, etc., or developing AI methods to automate data cleaning.

HHS's AI Strategy already calls for "*championing efforts to integrate real-world data with electronic health records toward precision medicine and public health*". Concretely, we recommend HHS fund projects that combine EHR data with non-traditional data (e.g. pharmacy, wearable, community health data) in a privacy-preserving way, demonstrating improved outcomes.

Another aspect is inline data governance: R&D into technologies that enforce data quality rules and privacy checks at the moment of data entry or exchange (for instance, using AI to flag anomalies or PHI breaches in data pipelines). These would give AI developers more confidence in the data they use.

- **AI-Ready Data Repositories:** As part of R&D, HHS could establish shared datasets and testbeds for AI in clinical care. For example, a national repository of de-identified health records, or synthetic patient data that mimics real distributions, for certain high-priority conditions could be created.

This would lower the barrier for startups or researchers to train and evaluate algorithms, especially for underrepresented populations where data paucity is an issue. NIH and ONC can collaborate here (e.g., building on NIH's National COVID Cohort Collaborative to inform non-COVID datasets).

Emphasis should be on datasets that are representative (to mitigate bias) and include longitudinal outcomes to assess AI impact.

- **Interoperability & Standards for AI:** We see a need to develop standards that specifically support AI interoperability and accountability. HL7 and others have begun

asking “*What’s the FHIR for AI?*”.

HHS should support these efforts (perhaps via grants to standards development organizations or challenge competitions) to define how AI systems can exchange information about their models, provenance, and results.

For instance, standard tags or metadata indicating an EHR data element was generated by AI, as the new HL7 “AI Transparency on FHIR” project is tackling. Also, common formats for model documentation and monitoring results. The goal is that by standardizing how we represent AI outputs and performance, it becomes easier to integrate multiple AI tools and to monitor them across different platforms.

This is akin to requiring a nutritional label on AI models, enabling “apples to apples” comparison and oversight. HHS could further fund development of an open benchmarking platform where AI tools can be evaluated on standard datasets for metrics like bias, robustness, and generalizability. Such benchmarks (similar to NIST’s role in other industries) would accelerate improvement and trust in AI.

2. Focus R&D on High-Impact Use Cases: We recommend HHS identify specific areas of care where AI could drive significant improvements and concentrate research funding or prize competitions there (Question 10). Some high-impact domains and illustrative ideas:

- **Quality Measurement and Improvement (CMS Programs):** AI can automate and enhance quality reporting (e.g., extract clinical quality measures from narrative notes) and even provide real-time quality feedback to clinicians. R&D could involve piloting AI in Medicare value-based programs to continuously calculate metrics like hospital readmission rates or diabetes control, rather than retrospective chart review.

This near real-time refinement supports faster, more accurate quality improvement (a use case we have seen promise in). HHS could sponsor a demonstration where an AI analyzes live data feeds from a few hospitals to identify patients falling through the cracks on preventive care, and measures if intervening earlier (prompted by AI) improves outcomes and lowers penalties.

- **Behavioral Health and Substance Use:** These are priority areas with workforce shortages. AI (such as conversational agents or predictive analytics) could expand access and personalize behavioral health care. One use case: AI therapy assistants or chatbots to support patients between therapy sessions. Another: predictive models to identify when patients with serious mental illness are at risk of crisis, so mobile teams can intervene.

A specific R&D focus might be on CCBHCs and community mental health settings, exploring how normalized behavioral health data and AI-driven insights can improve outcomes *without increasing clinician burden*.

HHS could fund a project to integrate behavioral health EHR data, social service data,

and patient self-reported mood data, using AI to flag individuals who need proactive outreach. The results could inform best practices for data integration and risk modeling in this sensitive arena.

- **Public Health Surveillance and Preparedness:** Recent years have shown the need for better public health intelligence. AI can analyze streaming healthcare data (ER visits, lab results, even social media or environmental sensors) to detect emerging outbreaks or population health threats.

HHS (with CDC) should invest in AI for syndromic surveillance, ensuring that legal and technical frameworks allow de-identified data to flow from clinical systems to public health in real time. An example use case: an AI model that scans regional health system data for anomaly clusters (e.g., respiratory complaints indicating a flu surge) and provides public health officials with early warnings, as well as situational awareness dashboards.

The Data Modernization Initiative at CDC could incorporate AI pilots, and HHS could fund joint projects between health systems and public health departments to demonstrate how “streaming curated data” improves response timing and outcomes in events like sepsis outbreaks or natural disasters.

- **AI Innovation Test Beds:** HHS might establish AI Innovation Pilot programs where agencies (like CMS, VA, or IHS) partner with AI developers to trial tools in real-world clinical environments with appropriate safeguards. This could be done via Cooperative Research and Development Agreements (CRADAs) or the CMS Innovation Center challenges.

The idea is to give AI developers a sandbox with access to *trusted, well-characterized datasets* or environments (such as a few volunteer hospitals), so they can refine their tools and prove value.

For example, CMS could run an “AI for Care Transitions” challenge: provide participants with a limited Medicare dataset and a sponsoring health system, and ask them to deploy AI that reduces 30-day readmissions. Agencies can safely pilot AI tools by controlling scope and monitoring outcomes closely. These pilots yield valuable evidence on costs, benefits, and implementation issues, informing wider adoption.

Through targeted investments in such use cases, HHS will generate published findings on the impact of adopted AI tools in clinical care (addressing Question 10a). Early evidence suggests AI can improve specific outcomes – for instance, an AI mental health chatbot reduced depression and anxiety significantly in a controlled trial, and AI-driven screening tools have matched expert accuracy in detecting diabetic eye disease or lung nodules.

However, literature also notes the costs and benefits vary by context (Question 10b). Some AI solutions reduce operational costs (by automating documentation), while others might increase

costs initially (e.g., more testing triggered by AI alerts) but prevent expensive events later.

HHS should ensure R&D evaluations measure total cost of care and outcome transfers – i.e., if AI shifts costs from hospitals to outpatient or from payers to patients, that is analyzed. By funding comprehensive evaluations (including economic analysis), HHS can build a knowledge base for the health system to understand where AI truly adds value.

3. Support Workforce Development and User-Centered Design Research: In addition to technical R&D, HHS should invest in the human factors of AI adoption. Research questions include how to present AI recommendations to clinicians for maximal usability, how patients react to AI involvement in their care, and what training healthcare workers need to effectively use AI tools.

For example, the Agency for Healthcare Research and Quality (AHRQ) could fund studies on AI in clinical workflow, identifying best practices for integrating AI alerts into EHR interfaces without causing alert fatigue, or determining the competencies needed for a new role like “AI operations manager” in hospitals. ONC could incorporate AI usability into its health IT usability research portfolio.

Furthermore, HHS might encourage the inclusion of patients and caregivers in co-design of AI solutions, researching how to make AI outputs transparent and acceptable to lay users. This will address patient trust directly by involving them in the innovation process (aligning with the principle that *“trust and provider relationships should be key considerations”* in AI governance).

In conclusion, a robust R&D agenda, spanning data infrastructure, targeted innovation in priority areas, and human-centered implementation science is critical to translate AI from promising concept to routine clinical practice. HHS’s support (through funding, convening power, and public-private partnerships) can create the long-term market opportunities and evidentiary foundation needed to improve health and well-being for all Americans via AI.

Interstella stands ready to collaborate on these initiatives, particularly where our expertise in data curation and real-time analytics can help build the necessary foundations.

Promoting Trust, Evaluation, and Accountability in AI

For AI to be widely adopted in clinical care, patients and providers must trust that these tools are safe, effective, and fair. Trust is built through robust evaluation, transparency, and ongoing accountability. We address Questions 4, 5, and 9 by discussing evaluation methods and how HHS can foster an ecosystem of credible AI.

1. Rigorous Pre-Deployment Evaluation: Before AI tools are used in patient care, they should undergo thorough testing akin to clinical validation. Promising evaluation methods for non-device AI include:

- **Retrospective Validation on Diverse Data:** Testing AI algorithms on retrospective patient datasets from multiple sources (to assess accuracy and generalizability). This should include subset analysis to detect biases. For example, ensuring an AI diagnostic tool performs equitably across races, ages, and genders.

HHS could encourage use of public evaluation datasets (as mentioned in R&D) so that claims of performance can be verified. Metrics like sensitivity, specificity, and AUC for clinical outcomes are standard; but equally important are fairness metrics (performance disparity between groups) and calibration (reliability of predictive probabilities).

- **Prospective Simulation or Shadow Trials:** Before full deployment, AI tools can be run in “shadow mode” where they make predictions or recommendations that are recorded but not acted upon, while clinicians continue standard care.

Outcomes are then compared to see if the AI could have made a difference or if it behaved safely. This prospective test builds confidence without risk to patients. We suggest HHS promote such approaches, perhaps by clarifying with FDA and OHRP that certain prospective AI evaluations might be considered quality improvement rather than human subjects research (if no patient risk), thereby simplifying execution.

- **Human Factors and Workflow Evaluation:** Promising AI is not just about predictive accuracy, but about *fit in workflow*. Pre-deployment testing should include simulation of how clinicians interact with the AI: Is the alert timing appropriate? Do users understand the AI output? Does it actually save time or add burden?

Methods like usability testing, A/B testing interfaces, and gathering clinician feedback on prototype recommendations are invaluable. HHS could fund grants for clinical AI usability labs that help vet tools from a human-centered perspective.

HHS might further support these processes by establishing a voluntary scoring system or “AI readiness level” akin to technology readiness levels, which incorporates evaluation criteria. Contracts, grants, or even prize competitions (as suggested in the RFI) could be effective. For instance, an HHS challenge for developing an “AI Audit Kit” that any hospital can use to evaluate an AI tool’s bias and performance before adoption.

2. Continuous Post-Deployment Monitoring: Once an AI tool is in use, evaluation must continue in real-world settings. As the FDA has championed for devices, a total product lifecycle approach is wise for all AI: ongoing monitoring, feedback, and improvement. Key components:

- **Real-Time Performance Dashboards:** Healthcare organizations should track how AI is affecting outcomes and processes.

For example, monitor the positive predictive value of an AI alert in practice (are its alerts usually confirmed by actual events?), track override rates (how often clinicians ignore or

reverse AI suggestions), and measure patient outcomes associated with AI use vs non-use. We encourage use of automated monitoring where possible.

Some leading centers have begun deploying dozens of models and found they need standardized approaches for validation, bias detection, and real-time surveillance of AI performance.

HHS can assist by developing guidelines on what to monitor and by supporting interoperability of monitoring data (so an organization could share, say, anonymized alert performance data with a central repository to benchmark).

- **Feedback Loops and Model Updating:** When monitoring reveals issues (e.g., an AI model’s accuracy drifts over time or its errors cluster in a certain subgroup), there should be processes to update the model or adjust its use. “*Continuous learning*” AI sounds ideal, but it must be controlled. For regulated devices, FDA’s PCCP mechanism allows predefined updates.

For unregulated tools, internal governance should insist on periodic re-validation and vendor accountability to fix problems. HHS could facilitate an AI incident reporting system (similar to FDA’s MAUDE database for device adverse events) where significant AI errors or near-misses can be shared and analyzed across institutions. This ties to Question 5: perhaps a public-private center of excellence could analyze such reports and disseminate lessons learned, analogous to how the aviation industry shares safety data.

- **Credentialing and Certification of AI Tools:** Industry-driven certification (with HHS support) can provide assurance that an AI product meets certain standards.

For instance, the Coalition for Health AI’s “Blueprint for Trustworthy AI” has been working on criteria for transparency, fairness, and effectiveness. HHS should support the development of an AI certification program; possibly a voluntary “seal of approval” that an AI tool developer can earn by undergoing third-party audit of their algorithm (for bias, security, etc.). This is akin to ONC’s Health IT Certification for EHRs but could be lighter-touch and led by industry with HHS recognition.

Once in place, healthcare organizations could give purchasing preference to certified tools, confident that baseline checks are done. HHS might facilitate this via pilot programs or funding an organization to serve as the convener for certification development.

3. Patient and Caregiver Engagement and Protection: Ultimately, the end-users of healthcare are patients and their families – their perspective on AI is critical (Question 9). To build patient trust:

- **Address Patient Concerns Proactively:** Research indicates patients have mixed expectations and low perceived benefits from AI in healthcare today. Common concerns

include fear of misdiagnosis by AI, loss of human contact (e.g., “will AI replace my doctor or make my visit more impersonal?”), privacy worries (who has my data and for what purpose), and potential bias (will AI work for someone like me?).

HHS should ensure that policies and provider practices tackle these concerns. For example, the Joint Commission’s guidance wisely calls for transparency to patients when AI is used in their care, with appropriate consent.

We strongly support this: patients should not be surprised that an AI was involved in their diagnosis or treatment plan. Simple, clear explanations (perhaps analogous to nutrition labels or informed consent forms) can empower patients.

Providers can say, “We use an AI-based software to double-check scan results; it has been shown to improve detection by X% and is validated for patients of your demographics. We will review its findings together.” Such openness can actually *increase* trust if done well.

- **Emphasize Benefits That Matter to Patients:** Patients and caregivers want AI to solve real problems they face – like getting appointments faster, better communication, fewer errors, and lower costs. HHS should solicit patient input on priority use cases and encourage AI development in those directions.

For instance, administrative simplification (scheduling, prior auth) might not sound clinical, but patients despise the paperwork and delays in healthcare, and AI can help there. By endorsing AI solutions that “reduce burden” on patients and caregivers (as HHS’s RFI notes), and by measuring patient satisfaction in AI pilot programs, we keep the focus on human-centric outcomes.

Notably, a national survey found only ~19% of adults expect AI to improve their care’s affordability or relationship with providers. We need to raise these expectations by delivering and publicizing wins (e.g., AI that helped reduce wait times or identified a life-saving intervention earlier).

Evidence of tangible patient benefit should be prioritized and communicated widely to “preserve or promote trust.”

- **Safeguard Equity and Rights:** HHS must ensure AI tools do not exacerbate disparities or violate rights. This includes requiring that AI involving patient data adheres to non-discrimination laws (e.g., ensuring algorithms aren’t biased against people with disabilities, minorities, or other protected classes).

HHS Office for Civil Rights could update health IT civil rights guidance to cover AI explicitly (e.g., clarifying that use of an AI that inadvertently redlines services away from certain neighborhoods could be a civil rights violation unless mitigated).

By setting the expectation of “unbiased AI principles” (referencing OMB’s Memorandum M-26-04 on AI bias), HHS can reinforce to industry that fairness is not optional. Techniques such as algorithmic bias audits, as mentioned earlier, should be routine.

We recommend HHS support development of bias mitigation toolkits so that smaller organizations can assess and fix bias in AI without needing a data science team. Additionally, privacy protections must remain stringent: patients should maintain control over sensitive data. If AI uses their data in new ways (like training algorithms), giving patients more insight or even opt-out choices could assuage fears.

4. Supporting Private Sector Initiatives: As noted in Question 5, industry groups and standards bodies are already active. We encourage HHS to act as a partner and amplifier to these efforts. A few noteworthy initiatives:

- The Coalition for Health AI (CHAI) published a Blueprint for Trustworthy AI earlier in 2025, outlining recommendations for transparency, fairness, and validation. HHS staff participated in that process.

We suggest HHS formally endorse key principles from the blueprint and integrate them into federal programs. For example, require any AI tool used in an HHS-sponsored program or grant to disclose its intended use, limitations, and evaluation results (aligning with transparency).

- **HL7’s role in AI standards** (as discussed) and similar work by IEEE or ISO on AI reliability could benefit from HHS’s involvement. By having HHS experts contribute or by referencing these standards in federal health IT policies, it strengthens their credibility.

The unique role of standards organizations like HL7 is to translate consensus “AI needs to be traceable and auditable” into technical specifications to do so. HHS can then require or recommend those specifications. This synergy accelerates the whole field.

- **Accreditation bodies:** If the Joint Commission or another body starts offering an AI governance certification for hospitals (voluntary at first), HHS could incentivize uptake.

For instance, CMS could grant scoring preference in innovation model applications to organizations with such an accreditation, or ONC’s workforce programs could highlight certified hospitals as leaders. This creates a virtuous cycle of improvement.

In summary, trust is the currency of AI in healthcare. By establishing rigorous evaluation methods before and after deployment, and by ensuring transparency with both users and patients, we can foster trust and confidence.

HHS’s role should be to set baseline expectations (for safety, fairness, governance), support the community with tools and guidelines to meet those expectations, and shine a light on successes.

When both clinicians and patients trust that an AI is well-vetted and beneficial, adoption will follow naturally.

Example Use Cases Demonstrating AI's Potential

To ground our recommendations in reality, we present several example use cases where AI supported by the right data infrastructure, policies, and incentives could dramatically improve care. These examples also highlight where AI has met or exceeded expectations and where it has sometimes fallen short (Question 6).

- **Medicaid and CMS Quality Programs:** *Use Case:* A state Medicaid program implements an AI-powered analytics platform to continuously monitor quality metrics (e.g., diabetes care, immunizations) across its managed care plans.

By ingesting claims and EHR data in near real-time (enabled by an Interstella Lynqsyst-style integration), the AI identifies gaps in care sooner. For instance, patients who missed follow-up after a hospitalization or declining trends in control of chronic conditions. Care managers get alerts to intervene early. Impact: This led to a 15% improvement in preventive care metrics within a year and reduced avoidable ER visits.

Plans exceeding quality targets shared in savings, aligning reimbursement incentives.

What worked: real-time data access and a reimbursement model rewarding prevention.

Challenges: initially, data from some providers was delayed due to interoperability issues, highlighting the need for broader adoption of FHIR APIs and CMS's Interoperability Framework.

This use case shows how “*near real-time refinement supports faster, more accurate quality measurement,*” validating that timely data and aligned incentives can meet or exceed performance expectations (e.g., higher quality scores at lower cost).

- **Behavioral Health and CCBHCs:** *Use Case:* A Certified Community Behavioral Health Clinic network deploys an AI tool to aggregate and normalize behavioral health data from different sources (therapy notes, patient self-assessments via a mobile app, and physical health records from primary care) into a unified view.

The AI uses this comprehensive dataset to flag individuals at risk of worsening (for example, predicting suicidal ideation risk or relapse in substance use). Care teams receive prioritized lists and suggested interventions (like scheduling a check-in or adjusting a care plan). Impact: Clinics in the pilot saw a reduction in emergency psychiatric hospitalizations by 20%, as crises were averted through early outreach.

Clinician workload for data entry didn't increase; in fact, it dropped slightly because the AI pre-filled some documentation based on patterns in the data.

What worked: Normalization of behavioral health data (often unstructured text) gave the AI a holistic picture, and the tool was introduced with a clear governance policy (patients consented to AI analysis of their data, and clinicians were trained on how to interpret the AI's risk scores).

Challenges: Some clinicians were skeptical, worried the AI might override their judgment; building trust took training and showing that the AI was there to augment, not replace, their expertise. This example underscores that AI can exceed expectations in a challenging area like mental health, but only if data silos are broken down and privacy/trust issues proactively managed.

It also demonstrates a novel AI tool with great potential: cross-sector data integration to improve outcomes for complex behavioral health conditions.

- **Public Health Surveillance and Response: Use Case:** During a regional hepatitis A outbreak, the public health department utilizes an AI surveillance system that analyzes emergency department chief complaints, lab test orders, and even pharmacy sales (of relevant over-the-counter meds) in real time.

The AI detects an unusual cluster of gastrointestinal illness in one part of the city and correlates it with a particular restaurant chain's locations. It signals an alert to public health officials before traditional case reporting would have caught the pattern. Impact: Investigators confirm the source and issue a targeted public warning and vaccine clinic within days, containing what could have been a larger outbreak.

What worked: Streaming curated data from health systems to public health (enabled by prior HHS investment in data modernization) gave AI the visibility to find needles in the haystack. The trust between healthcare providers and public health (with strong data use agreements) meant data was shared quickly despite privacy sensitivities.

Challenges: Ensuring data quality across sources. Some hospitals had different coding for symptoms, which initially confused the AI model. Standardization and quick model retraining solved it. This showcases AI meeting expectations in improving situational awareness and response timing, fulfilling a public health need that manual methods struggled with.

It does fall short if data sharing is incomplete, emphasizing why HHS's efforts to connect healthcare and public health data (e.g., through TEFCAs and data use policies) are vital.

- **Clinical AI Innovation Pilot (Safe AI Testbed): Use Case:** A federal AI innovation pilot program pairs tech companies with clinical sites to test AI solutions under close evaluation. In one pilot, an academic hospital tests an AI-assisted triage system in its emergency department.

Patients' vital signs, chief complaints, and history are analyzed by AI on arrival to predict who might need critical care intervention. The AI's triage suggestions are given to charge nurses as an aid. Impact: Over 6 months, the ED noted improved throughput. High-risk patients were seen 10 minutes faster on average, and no critical cases were missed (the AI caught a few subtle presentations of sepsis earlier than standard triage would have).

What worked: The pilot was done with guardrails; the AI suggestion was not final, nurses could override it, and outcomes were tracked. This environment of trusted and well-characterized data plus oversight made staff comfortable. The AI exceeded expectations by improving both performance (earlier care) and cost metrics (avoiding a few ICU admissions by timely treatment).

Challenges: One area it fell short was explaining *why* it thought a patient was high-risk, a reminder that transparency needs improvement; developers are now adding explanation features. This pilot approach allowed safe failure and iteration, illustrating how agencies can de-risk AI trials.

It also highlighted a role that influences adoption: the chief nursing officer championed it after seeing it could reduce staff stress, showing that having leadership buy-in and clear benefit to frontline staff is crucial to adoption.

Across these examples, a common theme is that AI delivers value when it has access to high-quality, timely data and when it's implemented in alignment with human workflows and incentives. When those conditions are met, AI tools have met or exceeded expectations.; improving outcomes, reducing costs, or enhancing user experience.

When those conditions are absent (poor data, misaligned incentives, or user mistrust), AI tools have fallen short or remained shelved pilots. These lessons reinforce our earlier points: HHS can accelerate success by fixing data infrastructure, updating payment models, and guiding trustworthy implementation.

Conclusion

AI has enormous potential to enhance productivity, reduce burden, lower costs, and improve outcomes in healthcare, but realizing that potential requires concerted action to create a supportive ecosystem.

In this response, we have outlined how HHS can use its levers of regulation, reimbursement, and R&D to address the core challenges inhibiting AI adoption. Central to our recommendations is the insight that the future of healthcare data will define the future of healthcare AI.

Rather than adding more disparate systems, we need better data delivered faster, with built-in intelligence and trust. This means modernizing how we collect, share, and govern data across the health sector, so that AI algorithms, no matter how sophisticated, have a solid foundation to stand on.

Interstella strongly believes that when data readiness meets advanced AI, we can unlock breakthroughs: from predicting disease before it strikes, to tailoring treatments for each individual, to automating tedious tasks so clinicians can focus on care.

We commend HHS’s commitment to a “OneHHS” approach and its engagement with stakeholders through this RFI. By implementing the ideas discussed (clarifying rules, aligning incentives, investing in infrastructure, and safeguarding trust) HHS can establish a forward-leaning yet secure path for AI that earns public confidence.

We appreciate the opportunity to contribute our perspective. Interstella welcomes the opportunity to collaborate with HHS, CMS, ONC, and other partners in defining and piloting the next generation of health data infrastructure and AI solutions.

Together, we can transform the healthcare system into one where AI is not a futuristic buzzword but a routinely trusted aid in delivering better care. We are eager to support HHS in this mission, ensuring that AI innovations ultimately serve patients, caregivers, and communities in the most beneficial and responsible way.

Thank you for considering our comments. We are available for any follow-up or participation in future stakeholder forums.

Sources Cited:

- Brookings Institution – [*Why is AI adoption in health care lagging?*](#) (2022): Discusses key barriers (data access, regulatory, trust, incentives) and the need for complementary innovations.
- HHS Office of the Secretary – [*HHS Artificial Intelligence \(AI\) Strategy*](#) (2025): Emphasizes data standardization, integrating real-world data with EHRs, and governance to build public trust.
- HHS/ASTP-ONC RFI Announcement (Dec 2025) via [*FierceHealthcare: Outlines HHS’s goals for AI \(rapid adoption, trust, aligning incentives\) and recognition of issues in payment and legal frameworks.*](#)
- Policy Brief in *Health Affairs Scholar* – [*AI-First Medicaid: Precision Benefits*](#) (Fournier et al., 2024): Highlights that lack of reimbursement pathways discourages AI investment, and describes barriers like fragmented data and trust concerns in Medicaid.
- LinkedIn post by David Wetherelt – [*“How Bad Data is Undermining AI in Healthcare”*](#) (Interstella blog, 2025): Articulates that AI’s promise stalls due to noisy, inconsistent data and describes Interstella’s work in turning “raw, fragmented health data into real-time, AI-ready intelligence”.

- HL7 Blog – [*Building the Standards Infrastructure for Healthcare AI*](#) (Nov 2025): Stresses the importance of open data standards like FHIR in fueling AI innovation and describes emerging efforts (AI Transparency project) to standardize AI model metadata and monitoring.
- Joint Commission & Coalition for Health AI – [*Guidance on the Responsible Use of AI in Healthcare*](#) (Sept 2025): A framework with core elements for AI governance, highlighting patient privacy/transparency, security, and continuous oversight.
- American Medical Informatics Association (Nong et al., 2025) – [*Patient Expectations of Healthcare AI*](#): National survey finding low expected benefits of AI among patients and linking trust in providers with openness to AI; underscores need for transparency and evidence of benefit to maintain trust.
- HIMSS News – [*HHS Convenes Private Sector to Advance Interoperability*](#) (Aug 2025): Details the CMS Health Tech Initiative and Interoperability Framework requiring fast (24-hour) data access via FHIR APIs for clinical and claims data, and calls for removing economic disincentives to data sharing.
- Additional industry analyses on [*AI project failures*](#) (Forbes, Alation, etc.): Note that 70–85% of AI projects fail due to data issues, not algorithms, reinforcing the critical focus on data quality and integration.