


ADVANCE COPY NOT FOR DISTRIBUTION



RE: HHS Health Sector AI RFI

RIN 0955-AA13

Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as part of Clinical Care

February 2, 2026

The Honorable Robert F. Kennedy, Jr.
Secretary
Department of Health and Human Services
Washington, DC

Dear Secretary Kennedy,

We appreciate the opportunity to submit comments in response to the Department of Health and Human Services Request for Information entitled “**Accelerating the Adoption and Use of Artificial Intelligence as part of Clinical Care**” (RIN 0955-AA13). Thank you for inviting public input on an issue of critical importance to the future of healthcare delivery, innovation, and patient trust in the United States.

HealthScoreAI, Inc. is a start-up consumer-focused health data and analytics company. Our leadership team brings decades of experience across healthcare operations both acute, and ambulatory, health information technology, in particular and relevant, building ambulatory electronic health record (EHR) platforms, third party administration, owning and managing Medicare Advantage networks, medical billing companies, and clinical administration. Collectively we also have been issued over forty (40) patents on the use of technology in healthcare ranging from EHR, IoT, predictive analytics, machine learning, and the use of blockchain in healthcare applications. Our perspective reflects both the promise of artificial intelligence (AI) and the risks associated with premature or misaligned deployment in clinical care settings. We strongly support innovation that improves outcomes, lowers costs, and empowers patients, while remaining cautious about policies that could unintentionally increase spending, erode trust, or substitute technology for clinical judgment.

ADVANCE COPY NOT FOR DISTRIBUTION

Enclosed below this transmittal letter are our responses to the RFI, organized section by section and clearly labeled as Comment or Question as you submitted, and our “REPLY” to align directly with the structure of the solicitation. Our comments focus on regulatory posture, reimbursement design, research & development strategy, interoperability, privacy protections, and the practical realities of AI adoption in clinical and administrative environments. We have intentionally emphasized lessons learned from prior health IT initiatives, including adoption of electronic health records (EHR), and the importance of patient centered data access, clear accountability, and market driven innovation.

We respectfully submit these comments for your consideration and would welcome continued dialogue as HHS advances its approach to artificial intelligence in the health sector. Thank you for your leadership on this important issue and for your commitment to balancing innovation with patient safety, privacy, and public trust.

Sincerely,

Noel J. Guillama-Alvarez

Noel J. Guillama-Alvarez
President
HealthScoreAI, Inc.
Wellington, Florida
nguillama@mypwer.com
<https://www.linkedin.com/in/nguillama/>

ADVANCE COPY NOT FOR DISTRIBUTION

Solicitation of Public Comments

Regulation

As the nation's principal health regulator, HHS helps shape the environment in which AI for clinical care is developed, evaluated, and deployed. HHS seeks to establish a regulatory posture on AI that is well understood, predictable, and proportionate to any risks presented to enable rapid innovation while protecting patients and the confidentiality of their identifiable health information and maintaining public trust. We seek feedback on how current HHS regulations impact AI adoption and use for clinical care.

REPLY-

We view the opportunities presented by artificial intelligence in clinical care as historically significant, but accompanied by equally significant warnings. As of December 2025, hundreds of peer reviewed studies both highlight the potential benefits of AI and raise serious concerns about its limitations and risks. While many studies demonstrate strong performance in highly controlled settings with apparent clinical validation, closer examination often shows that these systems do not perform as well as licensed medical professionals in real world clinical environments. The promise of AI is most evident in data intensive and objective domains such as medical imaging, pathology, and genomics, whereas performance degrades as clinical decision making moves beyond structured data into the more human and contextual aspects of care. A fundamental constraint is the fragmentation of health data. An estimated three thousand exabytes of clinical data in the United States are distributed across more than four hundred ONC certified ambulatory and acute care electronic health record systems. Even within standard clinical documentation, such as SOAP notes, subjective and objective information is separated, further limiting the reliability of AI driven inference. As a result, deploying AI broadly in clinical decision making today without significant safeguards presents real risk. We are particularly concerned about general purpose large language models, which can produce outputs that mix valid clinical knowledge with pseudo science or inaccurate information. Over time, this will necessitate the development and regulatory recognition of highly constrained, domain specific language models designed for clinical use. The experience with EHR adoption is instructive. Since the 2012 federal mandate, HHS has invested tens of billions of dollars, with industry investing hundreds of billions more, yet there was no measurable improvement in U.S. life expectancy from 2012 through 2019 prior to COVID. During the same period, healthcare costs increased substantially, raising legitimate questions about whether EHR adoption primarily optimized reimbursement and billing workflows rather than patient outcomes. For these reasons, we believe AI should initially be deployed primarily outside the core clinical decision making system to support

ADVANCE COPY NOT FOR DISTRIBUTION

consumers and administrative efficiency, rather than risk clouding physician judgment with opaque or poorly understood AI outputs.

Reimbursement

HHS's payment policies and programs have massive effects on how health care is delivered in the United States, often times with unintended consequences. Hypothetically, if a payer is taking financial risk for the long-term health and health costs of an individual, that payer will have an inherent incentive to promote access to the highest-value interventions for patients. Under government designed and dictated fee-for-service regimes, however, coverage and reimbursement decisions are slow. Rarely does covering new innovations reduce net spending; and waste, fraud, and abuse is difficult to prevent, often times leading to massive spending bubbles on concentrated items or services that are not commensurate with the value of such products. Given the inherent flaws in legacy payment systems, we seek to ensure that the potential promises of AI innovations are not diminished through inertia and instead such payment systems are modernized to meet the needs of a changing healthcare system. We seek feedback on payment policy changes that ensure payers have the incentive and ability to promote access to high value AI clinical interventions, foster competition among clinical care AI tool builders, and accelerate access to and affordability of AI tools for clinical care.

REPLY-

We agree that payment standards established by HHS have a profound influence on behavior across the healthcare system, often incentivizing providers and intermediaries to optimize for reimbursement rather than for better outcomes or lower costs. Historically, this dynamic has almost always resulted in higher aggregate spending. When Medicare was established in 1965, national health expenditures were approximately six percent of gross domestic product. Today, as reported by HHS actuaries, health spending exceeds eighteen percent of GDP and continues to trend upward, underscoring the structural limitations of current payment models.

We respectfully disagree with the assertion that a payer bearing long term financial risk for an individual has an inherent incentive to promote the highest value interventions. This assumption does not adequately account for the “velocity” with which patients change employers or health plans. In practice, average enrollment duration in employer sponsored coverage and Medicare Advantage plans is often five to six years. Annual benefit design changes further allow managed care organizations to indirectly shed higher risk members. As a result, the costs of underinvestment in prevention or suboptimal care are frequently externalized to future payers. In reality, the only entity that consistently bears true long term financial risk is HHS itself, which ultimately absorbs the downstream consequences of delayed care, neglect, or misaligned incentives within provider organizations and managed care plans.

ADVANCE COPY NOT FOR DISTRIBUTION

From a payment policy perspective, HHS should not reimburse for the use of AI as a discrete billable service. Reimbursing AI risks repeating the experience of prior technology mandates that optimized billing workflows without demonstrably improving patient outcomes. Instead, HHS should focus on fostering innovation by ensuring patients have meaningful access to their complete medical records in aggregate, rather than one record or encounter at a time. This would allow consumers to use carefully constrained and curated AI tools to better understand their health, engage in shared decision making, and manage chronic conditions. While AI clearly holds significant value in biomedical research, population health analytics, and drug discovery, those use cases fall largely outside the clinical reimbursement framework.

To foster competition and innovation, HHS should adopt a model similar to prior health IT certification approaches, in which AI innovators register with HHS, verify organizational identity, disclose intended data use including consumer data, and meet baseline trust and security requirements comparable to those imposed on EHR developers. HHS should set clear guardrails for safety, privacy, and accountability, while avoiding mandates that prescribe specific technical implementations. Experience with Meaningful Use demonstrates that when government mandates narrowly define requirements, vendors tend to meet the letter of the mandate without meaningful innovation. Imposing prescriptive requirements on AI risks creating a ceiling on innovation rather than enabling its full potential.

QUESTION-

Research & Development

HHS supports one of the world's largest health research ecosystems, catalyzing innovation to supplement the market. By enabling applied AI research & development, care delivery research and implementation science, as well as AI entrepreneurship in health care, we can better translate AI technologies from concept to clinical use. We seek input on ways in which HHS may invest in research & development (including public-private partnerships and cooperative research and development agreements (CRADAs)) to integrate AI in care delivery and create new, long-term market opportunities that improve the health and wellbeing of all Americans

REPLY-

We agree that HHS has the capacity to catalyze meaningful innovation through research and development, but it also has the capacity to unintentionally suppress innovation when programs are overly centralized or prescriptive. Government agencies are not well positioned to predict which technical approaches or use cases will ultimately prove transformative. In practice, many existing R&D mechanisms favor organizations that are

ADVANCE COPY NOT FOR DISTRIBUTION

adept at producing strong grant applications rather than those best positioned to deliver breakthrough innovation.

If HHS seeks to foster genuine innovation in AI, it should consider an approach more analogous to a diversified venture capital model, funding a large number of early stage efforts with relatively small awards rather than concentrating resources in a limited number of large programs. In parallel, HHS could provide access to standardized, high quality synthetic or de-ID datasets, for example a large corpus of representative electronic health record datasets, that innovators could use to explore hypotheses, develop models, and conduct back testing. This incremental and exploratory approach would allow promising ideas to emerge organically before larger investments are made.

The experience with electronic health records illustrates the risks of conflating system replacement with innovation. Over the past two decades, many hospitals have migrated from one major EHR vendor to another, such as from Meditech to Cerner to Epic. These transitions reflect vendor consolidation and rent seeking behavior rather than meaningful advances in care delivery or outcomes. AI research programs should avoid repeating this pattern by focusing on open experimentation rather than reinforcing incumbent platforms.

A particularly high value opportunity lies in expanding controlled access to de-ID data. Today, HHS and its affiliated agencies primarily possess claims and payment data, which are insufficient on their own to support robust clinical AI development. The most impactful AI applications require longitudinal clinical data, genetic information, and linked claims data to understand disease progression, treatment response, and outcomes over time. By enabling responsible access to these data types through well governed research partnerships, HHS could significantly accelerate entrepreneurial innovation and surface insights that centralized programs are unlikely to identify in advance.

Specific Questions

In addition to the general requests for information above regarding AI regulation, reimbursement, and research & development, HHS seeks input on the following specific questions:

1. What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

REPLY-

The most significant barrier to private sector innovation and adoption of AI in clinical care is that, to be practical and safe in most clinical settings, AI systems must perform

ADVANCE COPY NOT FOR DISTRIBUTION

at least as well as licensed physicians, and in many contexts demonstrably better. At present, most AI tools do not meet this threshold in real world clinical environments. Any realistic assessment must account for the fully burdened cost of AI use, not just the direct and often subsidized cost of the technology itself. This includes the time physicians spend prompting, supervising, validating, and correcting AI outputs. Physician time should be valued based on revenue generating capacity rather than salary alone, which in many specialties translates to an effective cost of approximately fifteen to thirty dollars per minute. When these opportunity costs are included, many AI tools fail to demonstrate economic value. In addition, unresolved questions around liability remain a major barrier to adoption. When AI recommendations are incorrect, it is unclear who bears responsibility for malpractice exposure, reputational harm, and potential licensure risk. Until accountability frameworks are clearly defined, providers and health systems will remain appropriately cautious in deploying AI in clinical care.

2. What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.

REPLY-

We are cautious in recommending new regulatory or payment policies related to AI, as prior experience with electronic health record mandates suggests that prescriptive approaches are unlikely to produce meaningful clinical or economic benefit. In particular, HHS should avoid establishing separate reimbursement for the use of AI. If AI is ultimately integrated into clinical care, its value proposition should be reflected in reduced costs, improved efficiency, or better outcomes, not in additional billable services. Creating payment codes for AI would raise fundamental dilemmas, such as whether a provider who uses AI should be reimbursed at a higher or lower rate than a provider who delivers the same service without AI, even when outcomes are equivalent.

In practice, AI adoption is already occurring without explicit payment incentives. Based on published medical reporting, an estimated forty to fifty percent of U.S. physicians have used tools such as OpenEvidence or similar AI enabled clinical reference systems. This underscores that clinicians will adopt tools they find useful without regulatory or reimbursement mandates. The priority role for HHS should therefore be risk mitigation rather than financial incentivization.

From a regulatory perspective, HHS should focus on strengthening protections around the use of protected health information in AI systems. In particular, HHS should clarify and enforce that identifiable patient data must not be disclosed to general purpose large language models or other AI systems that are not explicitly designed, governed, for research or clinical use only. Existing authorities under HIPAA, including 45 CFR

ADVANCE COPY NOT FOR DISTRIBUTION

Parts 160 and 164, could be revisited to provide clearer guidance on secondary data use, model training, and inference, ensuring that innovation does not come at the expense of patient privacy or public trust.

3. For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (e.g., relating to liability, indemnification, privacy, and security). What novel legal and implementation issues exist and what role, if any, should HHS play to help address them?

REPLY-

We agree with the historical approach HHS has taken toward non medical devices used by consumers, which generally relies on a consumer awareness and consumer responsibility model. In many contexts, this has been an appropriate and pragmatic posture. Our concern arises, however, when a non medical AI tool becomes affiliated with, integrated into, or operated by a designated entity under HIPAA, including a covered entity or a business associate under a business associate agreement.

Once a designated entity is involved, there is a significant legal and governance risk related to data classification. Consumer generated data that originates outside the HIPAA framework can effectively be transformed into protected health information through association with a covered entity or business associate. We believe it would be inappropriate and unsafe to allow designated entities, including EHR vendors and other HIPAA regulated actors, to move data back and forth between protected health information and so called consumer data in order to avoid regulatory obligations. This practice would undermine the integrity of HIPAA protections and create ambiguity around liability, privacy, and security responsibilities.

HHS should play a clarifying role by establishing clear guidance that once data is handled by a covered entity or business associate for clinical or operational purposes, it must be treated consistently as protected health information and remain subject to HIPAA safeguards. Clear boundaries around data status would reduce uncertainty for innovators while preserving accountability and public trust in the use of AI for clinical care.

4. For non-medical devices, what are the most promising AI evaluation methods (pre- and post deployment), metrics, robustness testing, and other workflow and human-centered evaluation methods for clinical care? Should HHS further support these processes? If so, which mechanisms would be most impactful (e.g., contracts, grants, cooperative agreements, and/or prize competitions)?

REPLY-

ADVANCE COPY NOT FOR DISTRIBUTION

As discussed in our prior responses, the most effective way to foster innovation while managing risk is to avoid overly prescriptive evaluation frameworks for non medical AI tools used in or adjacent to clinical care. Non medical devices are generally not the preferred modality for physicians when making clinical decisions, and attempts to force their integration into core clinical workflows often introduce more noise than signal. Variability in data quality, alerting logic, and model performance across consumer oriented tools can create cognitive burden for providers rather than improving care delivery.

For these reasons, HHS should be cautious in prioritizing large scale evaluation programs focused on non medical devices within clinical settings. Where evaluation is appropriate, it should emphasize real world usability, transparency of limitations, and clear boundaries around clinical reliance, rather than attempting to certify performance equivalence with medical devices. A limited, venture style funding approach could be used to support exploratory evaluation efforts, allowing innovators to test hypotheses and workflows without embedding them directly into EHR decision making.

If HHS chooses to support these processes, the most impactful mechanisms would be small, time limited grants or prize based challenges that encourage experimentation and post deployment learning, rather than long term contracts or large cooperative agreements. This approach would allow HHS to observe what delivers value in practice without prematurely standardizing evaluation methods in a rapidly evolving AI environment.

5. How can HHS best support private sector activities (e.g., accreditation, certification, industry driven testing, and credentialing) to promote innovative and effective AI use in clinical care?

REPLY-

As noted above, HHS should avoid stifling innovation through expansive accreditation, certification, or credentialing regimes that add administrative complexity without clear evidence of improved outcomes. Innovators, inventors, and entrepreneurs should not be required to assemble large compliance teams simply to engage with HHS or to explore novel AI applications. Excessive process requirements risk favoring incumbents and well resourced organizations over genuinely innovative entrants.

A more effective approach would be to apply a limited set of baseline trust requirements, similar to know your customer standards in financial services, focused on verifying organizational identity, governance, and accountability. These requirements should be sufficient to grant access to carefully governed HHS datasets, while maintaining strict adherence to existing HIPAA obligations and data protection standards. Beyond these guardrails, HHS should allow market driven testing, peer validation, and real world adoption to determine which AI tools deliver value in clinical

ADVANCE COPY NOT FOR DISTRIBUTION

care. This approach would promote innovation while preserving patient privacy and public trust.

6. Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?

REPLY-

There have been numerous reports of AI tools meeting or exceeding expectations in clinical care, particularly with respect to speed and throughput. However, when these claims are examined in detail, the performance gains often do not translate into consistent improvements in clinical quality, reliability of outcomes, or total cost reduction. Many highly positive peer reviewed studies published two or more years ago were conducted in controlled or retrospective environments that do not reflect the complexity and variability of real world clinical practice. As a result, reported performance frequently deteriorates when these tools are deployed at scale.

Publicized benchmarking studies further illustrate this gap. For example, widely cited industry sponsored evaluations have claimed that certain AI systems outperform physicians by large margins. When the underlying methodologies are reviewed, however, these studies often rely on artificial test conditions rather than real time clinical workflows, and they frequently permit AI systems to leverage knowledge sources that would not be available to a clinician at the point of care. Even under those favorable conditions, reported accuracy rates have often been closer to eighty percent, which falls short of the reliability required for unsupervised clinical decision making.

The greatest near term potential for AI lies not in replacing clinical judgment, but in supporting highly constrained and objective tasks such as image analysis, pathology screening, population level quality measurement, and administrative optimization. Novel AI tools that operate on well defined data types, produce auditable outputs, and reduce non clinical burden have the strongest potential to improve outcomes, generate new insights into quality, and reduce costs without introducing unacceptable clinical risk.

7. Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?

REPLY-

As noted in prior responses, current AI systems, particularly general purpose large language models, are not ready to replace physicians in clinical care and should not be

ADVANCE COPY NOT FOR DISTRIBUTION

treated as substitutes for clinical judgment. To the extent that AI tools are used to support clinicians, adoption decisions are primarily influenced by physicians themselves, clinical leadership, compliance and risk management teams, and health system legal counsel. These stakeholders are appropriately cautious given ongoing concerns related to patient safety, HIPAA compliance, malpractice exposure, patient acceptance, and professional reliance on AI generated outputs.

The primary administrative hurdles to adoption include unresolved accountability for errors, uncertainty around data governance, and the operational burden of integrating AI tools into existing clinical workflows. While AI may have more immediate and practical application in administrative and back office functions, HHS should remain alert to unintended consequences. For example, industry studies have reported that AI based clinical documentation tools can increase coded billing levels by as much as fifteen percent. It is not always clear whether these increases reflect more accurate documentation of care actually delivered or an expansion of documented diagnoses without corresponding clinical value.

Regardless of use case, any deployment of AI within clinical workflows must operate under strict HIPAA controls and clear governance structures. Fundamentally, AI is not yet ready for broad or unsupervised implementation in clinical environments. While future progress is likely, the timeline beyond the near term remains uncertain, and policy should reflect that uncertainty rather than assume rapid clinical readiness.

8. Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

REPLY-

Electronic health record interoperability in the United States has largely failed, and continued dialogue with large health IT vendors or major hospital systems has not meaningfully changed that reality. Structural incentives, rather than technical limitations, are the primary barrier. HHS should shift its focus away from vendor mediated interoperability and instead prioritize enforcement of HIPAA rights that allow patients direct access to their complete medical records. Patients are the most motivated stakeholders to demand portability, continuity, and meaningful use of their data.

Interoperability has stagnated for more than a decade because true data liquidity would materially weaken the economic moats of incumbent EHR vendors and reduce switching costs for health systems. If one EHR platform could efficiently and completely transfer longitudinal records to another, the financial and operational barriers to changing systems would be dramatically reduced. Today, hospitals often spend

ADVANCE COPY NOT FOR DISTRIBUTION

hundreds of millions of dollars, and in some cases more than one billion dollars, to transition data during EHR replacements. In a genuinely interoperable environment, EHR platforms would function more like regulated utilities rather than proprietary lock in systems.

Enhanced interoperability would most effectively fuel AI development if it focused on patient mediated data aggregation rather than system to system exchange alone. Longitudinal clinical records, combined with imaging, laboratory data, genomics, and claims data, made accessible through patient authorization, would create far more meaningful datasets for AI research and clinical innovation. Benchmarking tools should emphasize completeness, longitudinal continuity, and data provenance rather than simple message throughput or transactional compliance with standards.

Recent public statements by senior federal leadership have reinforced this direction, emphasizing that patients should control their own data and determine how AI tools are applied to it. HHS has a unique opportunity to accelerate AI development and restore trust by enforcing patient data access rights and allowing market forces to operate on top of truly portable health information, rather than continuing to rely on vendor centered interoperability frameworks that have repeatedly failed to deliver.

<https://youtu.be/b5p40OuTTW4?si=Alb22Tqxd2yv1RYk>

9. What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?

REPLY-

We do not believe that most patients or caregivers want AI to directly treat them. Patients overwhelmingly want to be cared for by physicians who demonstrate empathy, shared experience, and an understanding of their personal circumstances, values, and fears. Trust in care is often relational, not purely technical. In our experience as clinical administrators, many elderly patients prefer physicians who they feel understand their stage of life and lived experience. For some patients, particularly older adults or individuals from culturally tight knit communities, the introduction of AI into clinical decision making can feel impersonal or unsettling rather than reassuring.

While AI may be perceived as acceptable or even helpful by younger, more technologically fluent patients, the same tools may generate anxiety or mistrust among older populations or those with limited digital literacy. AI can simulate empathy through language and tone, but it cannot genuinely experience empathy, shared suffering, or moral responsibility. Patients and caregivers are therefore concerned that increased reliance on AI could erode the human connection that is central to healing, reduce accountability, and distance clinicians from the lived reality of their patients. Any

ADVANCE COPY NOT FOR DISTRIBUTION

adoption of AI in clinical care must respect these deeply human expectations and avoid substituting automation for relationships that patients value most.

10. Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care?

REPLY-

Please see answers above.

- a. Are there published findings about the impact of adopted AI tools and their use clinical care?
- b. How does the literature approach the costs, benefits, and transfers of using AI as part of clinical care?

Submitted by:

Noel J. Guillama-Alvarez
President

HealthScoreAI, Inc.

Wellington, Florida

nguillama@mypwer.com

<https://www.linkedin.com/in/nguillama/>