Department of Health and Human Service (HHS) in collaboration with
the Assistant Secretary for Technology Policy (ASTP) and
Office of the National Coordinator for Health Information Technology (ONC)

Request for Information (RFI)

Accelerating the Adoption and Use of Artificial Intelligence (AI) as Part of Clinical Care

Solicitation #: RIN 0955-AA13

03 February 2026

Prepared by:
BlueHalo, LLC (an AV company)
410 Jan Davis Drive, Huntsville, AL, 35806-4545

Submitted to:
Department of HHS
7033A, 330 C Street SW
Washington, DC 20201
Attention: HHS Health Sector AI RFI

Points of Contact:
Technical POC: Erin Gibson | (410) 303-2463 | erin.gibson@avinc.com
Contracting Representative: Yuko Ikedo | (301) 294-5216 | yuko.ikedo@avinc.com

Table of Contents

List of Tables

1.0 Executive Summary

Artificial Intelligence (AI) is increasingly embedded in clinical care workflows, research environments, and health-adjacent operations, as organizations seek to improve efficiency, quality, and insight, but its adoption remains uneven and constrained by systemic challenges related to governance, evaluation, and trust. We therefore see limits in AI's ability to scale safely, equitably, and sustainably in clinical and research contexts.

This response asserts that accelerating the adoption and effective use of AI in clinical care requires a shift in emphasis, from the development of additional AI models to the establishment of shared infrastructure, governance mechanisms, and evaluation frameworks that enable trust in data and AI systems. In particular, capabilities related to data readiness, interoperability, life cycle monitoring, and auditability are foundational to realizing AI's potential while safeguarding patients, clinicians, and institutions, particularly for non-medical device AI–AI embedded in workflows, documentation, decision support, data management, and operational tools that influence care without being regulated as medical devices.

Drawing on experience supporting governed, secure data and AI infrastructure in regulated health, research, as well as operational and adjacent environments, this submission offers experience-based insights into the systemic barriers impeding AI adoption and the policy levers available to alleviate them. The recommendations emphasize actions available; across regulation, research, reimbursement, and development, to standardize trust infrastructure, reduce uncertainty, and foster an ecosystem in which AI innovation can be responsibly translated into real-world clinical benefits.

The sections that follow are organized thematically to reflect the interconnected nature of these challenges, with explicit cross-references to the specific RFI questions presented.

1.1 About the Commenter

The commenter is a health and performance technology division of a large defense technology company with experience supporting complex, regulated environments through the design and operation of secure data and AI infrastructure. The company's work covers health, research, public-sector, and care-adjacent domains in which interoperability, governance, and accountability are critical to mission success.

This work emphasizes infrastructure and enablement as primary drivers of responsible AI adoption at scale. It includes facilitating evaluation, monitoring, and oversight of AI systems; supporting multi-institution and multi-stakeholder environments; and translating research capabilities into shared operational use. The perspectives offered here reflect lessons learned across these contexts and aim to inform policy and programmatic approaches that enable responsible AI adoption.

2.0 Barriers to AI Adoption in Clinical Care

Despite significant advances in AI capabilities, adoption within clinical care remains constrained by a set of interrelated, non-technical barriers. These challenges are largely structural and systemic, cutting across institutions, vendors, and care settings. Addressing them requires attention not only to individual tools, but to the broader environment in which AI is evaluated, deployed, and governed.

2.1 Fragmented Evaluation Environments

AI tools used in clinical care are often evaluated within vendor-controlled or site-specific environments, limiting independent assessment and comparison. Health systems and regulators frequently lack the ability to assess AI performance across multiple tools, populations, or deployment contexts. This fragmentation impedes informed decision-making, reinforces vendor lock-in, and constrains the ability to establish shared expectations for performance, safety, and value.

Without common evaluation frameworks or neutral environments for assessment, organizations struggle to move beyond pilot programs. Promising tools may demonstrate effectiveness under controlled conditions yet fail to generalize across institutions or patient populations. The absence of cross-vendor, cross-setting evaluation infrastructure remains a fundamental barrier to scalable AI adoption.

2.2 Lack of Real-World Performance Monitoring

Many AI deployments lack mechanisms for continuous monitoring once tools are introduced into clinical workflows. Performance is often assessed at a single point in time, prior to deployment, with limited visibility into how models behave over time as data distributions, workflows, and populations evolve.

In clinical settings, these changes are inevitable. Shifts in documentation practices, care delivery models, patient demographics, or upstream data sources can materially affect AI performance. Without ongoing monitoring and drift detection, degradation may go unnoticed, which undermines trust amongst clinicians and increases institutional risk.

2.3 Unclear Accountability and Governance

AI adoption introduces new questions of accountability that are not fully addressed by existing governance structures. When AI outputs influence clinical workflows, it is often unclear where responsibility resides among vendors, health systems, clinicians, and supporting infrastructure providers. This ambiguity creates hesitation among many adopting organizations particularly in regulated environments in which risk management and compliance are paramount concerns. In the absence of clear governance frameworks that define roles, oversight mechanisms, and escalation paths, institutions may defer adoption or limit AI use to low-impact scenarios.

2.4 Workflow Disruption and Clinician Trust

AI tools that are not well integrated into clinician workflows can impose additional cognitive and operational burden on clinicians. Even tools with strong technical performance may be resisted if they require significant changes or add many additional steps to workflows or documentation practices, introduce opaque recommendations, or lack clear mechanisms for human oversight (i.e. human-in-the-loop). This is particularly true when AI outputs exist outside of the electronic health record (EHR) environment, requiring clinicians to leave their established workflows to access or interpret them. When AI insights are not transparently integrated back into the patient record within existing clinical workflows, they are more likely to be interpreted as external impositions, rather than supportive capabilities.

2.5 Data Limitations and Bias

Clinical AI is frequently trained and evaluated using retrospective EHR data which, while essential for care delivery, often reflects local documentation practices rather than the full underlying clinical and operational context. Such data may be incomplete, biased, or unrepresentative of patient populations. Reliance on EHR data alone limits the ability to assess AI performance across diverse care settings and contexts. Moreover, EHR-centric approaches often fail to capture longitudinal, operational, and physiological factors that influence outcomes in real-world care.

These limitations contribute to bias, reduce ability to generalize, and complicate efforts to assess equity and effectiveness.

2.6 Structural Nature of the Barriers

Taken together, these challenges underscore that barriers to AI adoption in clinical care are not primarily the result of insufficient algorithms or computational capability. Rather, they reflect gaps in shared infrastructure for evaluation, monitoring, interoperability, and governance, particularly for non-medical device AI. These challenges are also evident in the growing use of consumer-grade sensors and wearable devices in care contexts, where clinicians are often presented with data without a clear understanding of how it was collected, validated, or governed, and therefore hesitate to incorporate it into clinical decision making. This further illustrates that trust in AI and data depends not only on algorithms, but on the surrounding infrastructure that establishes provenance, credibility, and accountability.

3.0 The Need for Continuous Evaluation, Monitoring, and Interoperability

Effective adoption of AI in clinical care requires moving beyond episodic validation and point-in-time approvals toward lifecycle-based evaluation, monitoring, and interoperability. For non-medical device AI performance, safety, and equity are not static properties; they evolve as data sources, workflows, and care contexts change. Without mechanisms to continuously evaluate AI behavior in real-world environments, institutions lack the visibility necessary to manage risk, sustain trust, and ensure value over time.

3.1 Pre-Deployment Versus Post-Deployment Evaluation

Current evaluation practices often emphasize pre-deployment testing conducted under controlled environments or idealized conditions. While necessary, these assessments provide limited insight into how AI tools perform once integrated into diverse clinical workflows and exposed to real-world variability. Differences in patient populations, documentation practices, staffing models, and upstream data quality can materially affect outcomes.

Post-deployment evaluation is therefore essential. Institutions need the ability to assess whether AI tools continue to meet performance expectations, behave consistently across settings, and align with clinical and organizational goals. This requires evaluation frameworks that persist beyond initial approval and are designed to operate across vendors, sites, and use cases.

3.2 Continuous Monitoring and Drift Detection

AI systems are sensitive to changes in the environments in which they operate. Over time, shifts in clinical practice, coding standards, patient demographics, or data pipelines, or the continued evolution of AI models, can introduce performance degradation or unintended effects. In the absence of continuous monitoring, such changes may remain undetected until they result in loss of confidence or other adverse outcomes.

Continuous performance monitoring enables early identification of drift, bias, and emerging risks. Monitoring mechanisms should be capable of tracking performance metrics longitudinally, comparing outcomes across populations, and automatically detecting and flagging deviations from expected behavior. For non-medical device AI, these capabilities are particularly important, as deployment often occurs outside of traditional regulatory surveillance frameworks. In many cases, these monitoring capabilities may rely on automated or AI-assisted techniques to operate at scale across complex environments.

3.3 Population- and Context-Aware Evaluation

AI tools may perform well for certain populations or settings while underperforming for others. Evaluation approaches that rely on aggregate metrics risk obscuring disparities and masking context-specific failures. To support equitable and effective use, AI evaluation must account for variation across patient populations, care setting, and operational contexts.

This requires access to diverse, representative data and the ability to stratify performance by relevant factors. Longitudinal evaluation further supports understanding how AI impacts outcomes over time, including downstream effects to workflow, clinician behavior, and patient experience.

3.4 Interoperability Beyond Data Exchange

Interoperability is a foundational requirement for AI evaluation and monitoring, but it must extend beyond the exchange of clinical data alone. Standards such as Fast Healthcare Interoperability Resources (FHIR) have significantly advanced data access and exchange; however, data movement alone does not provide sufficient context to evaluate or govern AI systems.

Effective AI interoperability must encompass additional layers, including metadata, data provenance, evaluation artifacts, and governance policies. Understanding how data is generated, transformed, and used is critical to interpreting AI outputs and assessing performance. Similarly, interoperability across evaluation and monitoring systems enables comparability and shared learning without requiring centralization of sensitive data.

Government action could help create conditions under which AI used within EHR workflows is accompanied by appropriate provenance, confidence, and governance-related metadata to support transparency and clinician trust.

3.5 Supporting Governance and Accountability through Interoperability

Interoperability across evaluation and governance layers also supports clearer accountability. When AI systems are deployed across multiple institutions or workflows, shared frameworks for monitoring and reporting enable more consistent oversight. This transparency reduces uncertainty for clinicians, administrators, and regulators, and supports coordinated responses when issues arise.

By enabling interoperable evaluation and monitoring infrastructures, HHS can help establish a common foundation upon which diverse AI tools can be responsibly deployed. Such an approach reduces fragmentation, supports trust, and accelerates adoption without constraining innovation.

3.6 Implications for Policy and Investment

The need for continuous evaluation, monitoring, and interoperability has direct implications for policy and investment decisions. Government action that supports shared evaluation frameworks, reference architectures, and interoperable monitoring capabilities can lower barriers to adoption across the health ecosystem.

Investments in these foundational capabilities complement rather than compete with private-sector innovation. By focusing on infrastructure and governance, HHS can enable a more resilient and trustworthy AI ecosystem that supports clinical care, research, and public health objectives.

4.0 Regulatory And Policy Considerations for Non-Medical Device AI

As AI becomes increasingly embedded in clinical workflows and care-adjacent operations, existing regulatory and policy frameworks face growing strain. Many current approaches were developed to govern discrete medical devices or static health IT systems and are not well suited to non-medical device AI that evolves over time, operates across organizational boundaries, and influences care indirectly through workflows, prioritization, and information presentation.

Addressing these gaps does not require new regulation for every AI application. Rather, it requires policy clarity that aligns oversight, accountability, and incentives with the realities of how AI is developed, deployed, and maintained in clinical environments.

4.1 Gaps in Accountability, Liability, and Governance

Non-medical device AI often sits outside of traditional regulatory pathways, which creates ambiguity around accountability where systems underperform or cause unintended consequences. Responsibility may be distributed among model developers, data providers, systems integrators, health systems, and end users, yet existing governance structures rarely define how these roles should interact or where oversight should reside.

This ambiguity discourages adoption, particularly in regulated environments where risk management is paramount. Organizations may hesitate to deploy AI broadly without clear expectations regarding monitoring, escalation, documentation, and corrective action. Clarifying governance responsibilities for non-medical device AI would reduce uncertainty and support more confident, responsible use.

4.2 Over-Focus on AI Artifacts versus Systems

Policy discussions often focus on the characteristics of individual AI artifacts, such as model architecture, training data, or performance metrics, rather than the systems in which those artifacts operate. In practice, many of the risks associated with AI arise not from the model itself, but from how it is integrated into workflows, how inputs are generated and governed, and how outputs are monitored and acted upon.

A systems-level perspective recognizes that AI performance and safety depend on data quality, interoperability, human oversight, and organizational context. Policies that emphasize system-level practices, rather than artifact level certification alone, are better aligned with the realities of non-medical device AI deployment.

4.3 The Role of Safe Harbors in Governance Practices

One potential policy lever is the use of safe harbors that recognize adherence to defined governance and evaluation practices. Rather than prescribing specific technologies or models, HHS could encourage adoption by clarifying that organizations implementing AI within approved governance frameworks—such as those supporting continuous monitoring, auditability, and human oversight—are meeting reasonable expectations for responsible use.

Such an approach incentivizes adoption of defined governance practices while preserving flexibility and innovation. It also aligns oversight with outcomes, focusing attention on whether AI systems are appropriately evaluated and managed over time, rather than on static design characteristics.

4.4 Reference Architectures and Shared Infrastructures

Reference architectures can play a critical role in translating policy into operational guidance. By articulating how evaluation, monitoring, interoperability, and governance components fit together, reference architectures provide a common foundation for institutions, vendors, and regulators.

HHS is well positioned to support the development or endorsement of reference architectures for non-medical device AI that build upon existing standards while extending beyond data exchange to include life-cycle oversight. Shared infrastructure such as evaluation environments or monitoring frameworks, can further reduce fragmentation and lower barriers to adoption without centralizing data or constraining competition.

4.5 Alignment with Incentives and Programs

Regulatory clarity alone is insufficient if incentives remain misaligned. Payment, research, and innovation programs play a significant role in shaping adoption behavior. Policies that recognize the cost and operational burden associated with evaluation, monitoring, and governance can encourage institutions to invest in these capabilities.

By aligning regulatory expectations with reimbursement models and research investments, HHS can reinforce the importance of responsible AI practices and support sustainable adoption across the health ecosystem.

4.6 Implications for Policy Development

Taken together, these considerations suggest that effective governance of non-medical device AI requires a shift toward system-oriented, lifecycle-based policy approaches. Clarity around accountability, support for reference architectures, and alignment of incentives can reduce uncertainty while enabling innovation. Such an approach allows HHS to guide AI adoption through shared expectations and enabling infrastructure, rather than through prescriptive control of individual technologies.

5.0 Role of R&D, CRADAs, and Public-Private Partnerships

Research and Development (R&D), cooperative research and development agreements (CRADAs), and other public-private partnerships play a critical role in translating AI capabilities from concept to sustained use in clinical care. For non-medical device AI, the primary challenge is not the absence of innovation, but the lack of environments and representative datasets in which evaluation, monitoring, governance, and interoperability practices can be developed, tested, and refined in real-world conditions.

5.1 Applied R&D versus Basic Research

Basic research continues to advance the state of the art in AI methods and algorithms. However, many of the barriers to adoption identified in this response reside downstream of algorithmic innovation. Applied R&D is needed to address questions of new applications, implementation, integration, and lifecycle management, including how AI systems interact with workflows, data quality processes, and governance structures over time.

Recent advances in generative AI (GenAI) further underscore the need for robust evaluation, monitoring, and governance practices, as these tools are increasingly integrated into clinical and operational workflows. Experience across health, research, and operational environments suggests that these applications introduce new considerations around reliability, drift, and oversight that cannot be addressed through model development alone.

HHS-supported applied R&D can help bridge this gap by focusing on these operational dimensions of AI use, such as evaluation methodologies, monitoring approaches, and interoperability mechanisms, rather than on model development alone. These efforts complement existing research investments while addressing practical challenges faced by adopting institutions.

5.2 Implementation Science and Real-World Validation

Implementation science provides a valuable framework for understanding how technologies perform when introduced to complex, real-world environments. For AI in clinical care, this includes assessing not only technical performance, but usability, workflow impact, equity considerations, and sustainability.

Public-private partnerships can support implementation-focused studies that examine how governance practices, monitoring strategies, and interoperability choices influence outcomes. Such work generates evidence that is directly relevant to policy development and program design and helps reduce uncertainty for health systems considering adoption.

5.3 Multi-Institution, Vendor-Neutral Partnerships

Many challenges associated with AI adoption span organizational boundaries. Single institutional efforts, while valuable, may not capture the diversity of contexts in which AI tools operate. Multi-institution partnerships enable broader evaluation across populations, workflows, and settings, improving generalizability and robustness.

Vendor-neutral partnerships are particularly important for non-medical device AI. Neutral environments allow evaluation and governance practices to be developed independently of specific products, reducing conflicts of interest and supporting fair comparison. HHS is well positioned to convene such partnerships and establish shared expectations.

5.4 Role of CRADAs and Cooperative Agreements

CRADAs and cooperative agreements provide flexible mechanisms for collaboration between government, industry, and research organizations. These instruments can support joint development of evaluation frameworks, reference architectures, and monitoring approaches while leveraging complementary expertise.

By using these mechanisms to focus on infrastructure and governance capabilities, HHS can accelerate progress without dictating specific technical solutions. Lessons learned through CRADAs can inform broader policy and programmatic decisions, creating a feedback loop between experimentation and guidance.

5.5 Implications for HHS Investment Strategy

Targeted investment in applied R&D, implementation science, and shared testbeds can yield outsized impact by addressing systemic barriers to AI adoption. Such investments amplify the value of private-sector innovation by providing the conditions necessary for responsible deployment.

By prioritizing partnerships that emphasize evaluation, monitoring, governance, and interoperability, HHS can foster an ecosystem in which AI tools are not only developed but sustained and trusted in clinical care.

6.0 Closing Recommendations

The following recommendations shown in **Table 1** summarize the themes discussed above and focus on actions HHS could take to accelerate the responsible adoption of AI in clinical care. Each recommendation emphasizes enablement rather than prescription, and is intended to reduce uncertainty, strengthen trust, and align incentives across the health ecosystem.

Table 1. Closing Recommendations

Theme	What HHS Could Do	Why it Matters	Outcomes Enabled
Establish Reference Architectures for Non-Medical Device AI	Develop or endorse reference architectures that describe how evaluation, monitoring, interoperability, and governance components fit together across the AI lifecycle for non-medical device applications.	In the absence of shared architectural guidance, institutions and vendors independently interpret expectations, leading to fragmentation and inconsistent practices. Reference architectures translate policy intent into operational clarity without mandating specific technologies.	Accelerated adoption through reduced ambiguity, improved comparability across AI deployments, and a common foundation for institutions, vendors, and regulators.
Support Continuous Evaluation and Monitoring Frameworks	Promote evaluation frameworks that extend beyond pre-deployment validation to include post-deployment monitoring, drift detection, and longitudinal assessment across populations and settings.	AI performance and risk evolve over time. Without continuous evaluation, degradation and inequities can go undetected, undermining trust and increasing institutional risk.	Sustained AI performance, earlier identification of issues, improved equity assessment and stronger confidence among clinicians and administrators.
Encourage Interoperability Across Evaluation and Governance Layers	Advance interoperability approaches that extend beyond data exchange to include metadata, data provenance, evaluation artifacts, and governance policies, building upon existing standards where appropriate.	Data interoperability alone is insufficient to support AI oversight. Shared visibility into context and lineage is essential for evaluation, accountability, and coordinated response.	Cross-institutional learning, reduced duplication of effort, and scalable oversight without centralizing sensitive data.
Clarify Governance Expectations and Enable Safe Harbors	Provide guidance that clarifies governance expectations for non-medical device AI and consider safe harbors for organizations that implement defined evaluation, monitoring, and oversight practices.	Uncertainty around accountability and liability discourages adoption. Safe harbors tied to clearly defined governance and oversight practices will reduce perceived risk and incentivize responsible adoption without constraining innovation.	Increased confidence among adopting institutions, alignment of incentives, and broader participation in responsible AI deployment.
Invest in Applied R&D and Public-Private Testbeds	Use R&D programs, CRADAs, and cooperative agreements to support applied research, implementation science, and real-world test beds focused on AI governance, evaluation, and interoperability.	Many adoption barriers can only be addressed through experimentation and shared learning in representative environments.	Evidence-based policy development, reusable infrastructure, and faster translation of AI innovation into sustained clinical benefit.

Together, these actions emphasize infrastructure, governance, and lifecycle oversight as critical enablers of AI adoption. By focusing on shared foundations rather than individual tools, HHS can accelerate innovation while maintaining public trust and protecting patients and providers.

Appendix A: Question Crosswalk

Table 2. Question Crosswalk

#	RFI Question	Section(s) Addressed
1	What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?	**Section 2.0**—Barriers To AI Adoption in Clinical Care **Section 6.0**—Closing Recommendations
2	What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.	**Section 4.0**—Regulatory And Policy Considerations for Non-Medical Device AI **Section 6.0**—Closing Recommendations
3	For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (e.g., relating to liability, indemnification, privacy, and security). What novel legal and implementation issues exist and what role, if any, should HHS play to help address them?	**Section 4.0**—Regulatory And Policy Considerations for Non-Medical Device AI
4	For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and other workflow and human-centered evaluation methods for clinical care? Should HHS further support these processes? If so, which mechanisms would be most impactful (e.g., contracts, grants, cooperative agreements, and/or prize competitions)?	**Section 3.0**—The Need for Continuous Evaluation, Monitoring, and Interoperability
5	How can HHS best support private sector activities (e.g., accreditation, certification, industry driven testing, and credentialing) to promote innovative and effective AI use in clinical care?	**Section 4.0**—Regulatory And Policy Considerations for Non-Medical Device AI **Section 5.0**—Role Of R&D, CRADAs, And Public-Private Partnerships
6	Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?	**Section 3.0**—The Need for Continuous Evaluation, Monitoring, and Interoperability
7	Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?	**Section 2.0**—Barriers To AI Adoption in Clinical Care
8	Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.	**Section 3.0**—The Need for Continuous Evaluation, Monitoring, and Interoperability
9	What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?	**Section 2.0**—Barriers To AI Adoption in Clinical Care
10	Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care? a. Are there published findings about the impact of adopted AI tools and their use clinical care? b. How does literature approach the costs, benefits, and transfers of using AI as part of clinical care?	**Section 5.0**—Role Of R&D, CRADAs, And Public-Private Partnerships