



Health and Human Services Department
*RFI: Accelerating the Adoption and Use of Artificial
Intelligence as Part of Clinical Care*
Submitted by: Ellipsis Health
Date: February 3, 2026

Executive Summary

Empathetic, Clinically Governed AI Can Reduce Healthcare Costs and Improve National Health

Artificial intelligence represents one of the most consequential opportunities of this decade to improve health outcomes while reducing the unsustainable growth of healthcare costs in the United States. As an AI-enabled care management technology company actively deployed in real-world clinical and population health settings, we submit this response to emphasize a central premise: **AI is most readily adopted and effectively used in clinical care when it is empathetic, highly engaging, and embedded within strong guardrails that prioritize patient safety, equity, and trust.**

AI as a Force Multiplier for Human Care

The most effective healthcare AI does not seek to replace clinicians or focus only on automation. Accelerating adoption of AI requires empathy to remain a foundational element of care delivery. Effective AI solutions **extend the reach of care teams**, allowing them to identify risk earlier, engage patients more frequently, and intervene before avoidable deterioration occurs. AI solutions adopted as

In practice, the highest-cost healthcare events—emergency department utilization, hospitalizations, readmissions, and unmanaged chronic disease—are often preceded by **unaddressed behavioral, physical, or psychosocial signals**. AI-enabled care management tools can consistently surface these signals earlier and reduce risk of unconscious bias vs. traditional models. Earlier detection and action enable more timely, lower-cost interventions that preserve both patient well-being and system resources.

Empathy and Engagement Are Core Safety Features

From the patient's perspective, engagement is not driven by algorithms alone, it is driven by whether an interaction feels **respectful, empathetic, understandable and supportive**. AI systems that leverage natural, conversational modalities (such as voice) and are designed around empathy and clarity achieve significantly higher engagement rates than transactional or form-based tools.

High engagement is not merely a usability benefit; it is a **patient safety imperative**. Patients who engage consistently are more likely to disclose emerging symptoms, emotional distress, medication concerns, and social barriers to care. This allows care teams to act earlier, reducing clinical risk and downstream cost. Such interactions are also more likely to initiate behavior changes that enable a trajectory of improved health.

Patient Safety Requires Clinical Guardrails, Not Just Technical Accuracy

AI used in care management must operate within **explicit clinical boundaries**. On-going review and use of rigorous quality assurance processes are required to maintain patient safety and adherence to clinical guardrails. Effective guardrails include:

- Clear definitions of AI scope and escalation thresholds
- Human-in-the-loop oversight by licensed clinicians
- Continuous performance monitoring and bias evaluation
- Structured clinical review of AI outputs and interventions
- Transparent communication to patients about AI use

When AI solutions are designed with these safeguards, they can safely operate as **adjuncts to clinical care**, improving consistency while preserving accountability.

Quality Assurance and Trust Are Prerequisites for Scale

Public trust is essential to national AI adoption. Trust is earned when AI tools demonstrate:

- Clinical validity grounded in real-world evidence
- Measurable improvements in outcomes and utilization
- Strong data privacy and security protections
- Alignment with existing care workflows and standards

HHS can accelerate responsible adoption of non-medical device AI by encouraging evaluation frameworks that assess **clinical impact, safety, equity, and operational performance**. Non-medical device AI is playing a growing role in care delivery and with stronger evaluations frameworks and governance established, adoption and value would be accelerated.

Deflating Costs Through Early, Continuous, and Equitable Care

AI-enabled care management supports the objectives of “Make America Healthy Again” by shifting the system from episodic, reactive care toward **continuous, preventive engagement**—especially for populations that are traditionally underserved or difficult to reach. By enabling:

- Earlier detection of behavioral and psychosocial risk
- More frequent, lower-cost patient touchpoints
- Better prioritization of clinical resources
- Reduced avoidable utilization

AI can materially reduce the total cost of care while improving patient experience and outcomes.

A Call for Alignment Across Regulation, Reimbursement, and R&D

To fully realize these benefits, federal policy should align incentives so that empathetic, safe, and clinically governed AI is rewarded—not penalized—for preventing harm and reducing cost. This includes:

- Modernized reimbursement that values prevention and engagement
- Clear, proportionate governance models for non-medical device clinical AI
- Investment in implementation science and public-private partnerships

With the right guardrails and incentives, adoption of AI-enabled care management solutions will strengthen the human foundations of healthcare and advance national health and fiscal sustainability.

QUESTION 1:

What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

AI-enabled care management has the potential to materially reduce healthcare costs while improving patient experience and outcomes. **A foundational element to success is AI developed to be empathetic, engaging, and governed by strong clinical guardrails.** Despite increasing interest and technological readiness, adoption across the healthcare sector remains uneven. The adoption challenges are not primarily technical. Organizations are facing **structural, operational, and incentive-driven challenges.**

The following outlines the most significant barriers to private sector innovation in AI for healthcare and its effective use in clinical care, informed by real-world deployment experience.

1. Misaligned Incentives That Reward Treatment Over Prevention

One of the most persistent barriers is the **disconnect between where AI creates value and where the healthcare system pays for care.**

Many AI tools, particularly in care management, behavioral health, and population health, deliver value by:

- Detecting risk earlier
- Preventing avoidable escalation
- Reducing utilization that would otherwise generate billable events

In fee-for-service environments, preventing a hospitalization or emergency department visit often **reduces revenue rather than creating it.** Even in value-based arrangements, savings are frequently

delayed, indirect, or difficult to attribute, making it challenging for health systems to justify near-term investment in AI.

Real-world impact:

Health systems may recognize that AI-driven outreach reduces ED visits among high-risk populations, yet still struggle to operationalize or scale these tools because the financial upside accrues slowly or to a different stakeholder (e.g., a payer rather than a provider).

2. Workforce Burnout and Change Fatigue Limit Adoption Capacity

Healthcare organizations are operating under unprecedented workforce strain. Nurses, care managers, and social workers are already managing unsustainable caseloads and administrative burden. These teams are often the primary users of care management AI, and where patient care shifting to prevention is most prevalently needed and realized.

While AI can reduce workload over time, **initial adoption requires change**, including:

- Workflow redesign
- Training and trust-building
- Governance and oversight processes

Organizations experiencing acute staffing shortages often lack the bandwidth to implement new solutions, even when one of the primary benefits is to alleviate long-term burnout and re-balance workload between administrative and clinical care needs.

Real-world impact:

AI initiatives stall not because they lack clinical value, but because frontline teams are stretched too thin to participate in pilots, feedback loops, or early optimization phases.

3. Lack of Clear Governance for Non-Medical Device AI

Many of the most impactful AI solutions in care management do not meet the definition of a regulated medical device. These tools may assess risk, guide outreach, or support clinical prioritization without making autonomous diagnoses or treatment decisions.

However, health systems often struggle with:

- Unclear liability and accountability frameworks
- Uncertainty around appropriate clinical oversight
- Inconsistent internal review processes
- Lack of accepted frameworks for AI system governance and assessment

Ambiguity about governance structure for non-medical device AI leads to **risk aversion, prolonged procurement cycles, and inconsistent decision-making** across organizations.

Real-world impact:

Two health systems evaluating the same AI tool may reach opposite conclusions. The conclusions are not due to evidence differences, but due to internal uncertainty about governance, compliance, or legal exposure.

4. Fragmented Data and Limited Interoperability

AI innovation depends on timely, high-quality data. Yet clinical, behavioral, and social data remain highly fragmented across EHRs, care management platforms, and community-based systems.

As a result:

- AI tools must rely on incomplete data
- Integration timelines are prolonged
- Value realization is delayed

This disproportionately affects care management and population health use cases, which rely on **longitudinal, cross-setting insight** rather than point-in-time clinical data alone.

Real-world impact:

Organizations may deploy AI that identifies behavioral health risk, but lack seamless pathways to connect those insights into care coordination workflows, reducing real-world impact.

5. Patient Engagement Is Undervalued Despite Being Foundational

Many AI evaluations focus heavily on technical performance metrics (e.g., accuracy, sensitivity) while de-emphasizing **engageability factors** that indicate the likelihood that patients will consistently interact with the technology long-term.

In real-world healthcare:

- Low engagement negates even the most accurate AI
- Vulnerable populations are least likely to engage with impersonal or transactional tools
- Engagement failures often surface as safety or quality issues downstream

Real-world impact:

AI tools that perform well in controlled settings may fail at scale because patients disengage, stop responding, or misunderstand interactions. The impact is lost trust and disengagement which can lead to missed shifts in care needs and lost opportunity for early intervention.

6. Public Trust and Patient Safety Concerns Slow Adoption

Patients and caregivers are increasingly aware of AI in healthcare and often express concerns related to:

- Privacy and data use
- Loss of human connection
- Fear of automation replacing clinical judgment

Without clear communication and strong safety assurances, these concerns can lead organizations to delay or limit AI use even when potential benefits are substantial.

Real-world impact:

Health systems frequently restrict AI functionality or scope not because of demonstrated harm, but because of perceived reputational or ethical risk in the absence of shared standards.

Question 1: Summary

In aggregate, the greatest barriers to private sector AI innovation and adoption are not rooted in algorithmic capability, but in:

- Misaligned financial incentives
- Workforce constraints
- Governance uncertainty
- Data fragmentation
- Underappreciation of engagement and empathy
- Lack of shared trust frameworks

Acceleration of AI adoption will be achieved by coordinated action plans to improve regulation, reimbursement, and research that allows continual innovation, and **enables safe, scalable, and patient-centered AI deployment.**

QUESTION 2:

What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited?

To unlock the full cost-deflation and quality-improvement potential of AI in clinical care, HHS should prioritize changes that **reward prevention, clarify governance, and reduce operational friction**, particularly for AI-enabled care management and non-medical device clinical AI solutions.

1. Modernize Payment Models to Reward Prevention and Engagement

A top priority is to align reimbursement with the outcomes AI is uniquely positioned to improve.

Many AI tools generate value by **preventing adverse events rather than producing billable encounters**. HHS should modernize payment policy to recognize AI-enabled activities that demonstrably reduce total cost of care.

Recommended actions:

- Expand CMS reimbursement pathways for **AI-supported assessment, monitoring, and outreach**, particularly in:
 - Chronic disease management
 - Behavioral health and comorbidity screening
 - Transitional care and post-discharge follow-up
- Encourage Medicare Advantage, Medicaid managed care, and ACO models to:
 - Treat AI-enabled engagement as a covered care management activity
 - Allow AI-generated insights to count toward quality and risk-adjustment workflows when clinically validated, but without specified human review times

Real-World Impact:

Without reimbursement alignment, AI adoption remains limited to pilot programs, even when evidence shows **reductions in ED visits, hospitalizations, and readmissions**.

2. Establish Clear, Proportionate Governance for Non-Medical-Device AI

A top priority of AI governance is to reduce uncertainty without over-regulating low-risk, high-value AI.

A significant portion of AI used in care delivery does not meet the definition of a medical device, yet still influences clinical prioritization and care coordination. Today, governance varies widely across organizations, slowing adoption.

Recommended actions:

- Issue HHS guidance clarifying:
 - Appropriate oversight models for non-medical device AI
 - Expectations for human-in-the-loop review
 - Documentation and audit standards proportional to risk
- Distinguish between:
 - AI that informs care (risk stratification, screening, engagement)
 - AI that autonomously diagnoses or treats (higher regulatory threshold)

Real-World Impact:

Clear guardrails reduce institutional risk aversion and accelerate responsible deployment. Creating clear proportionate governance for non-medical device AI in care management and population health, can improve AI adoption that directly improves cost, care, patient safety through early detection and proactive care.

3. Incentivize Clinical Quality Assurance and Post-Deployment Monitoring

Shifting evaluations from one-time approval to continuous safety and quality assurance is a top priority to ensure clinical quality is maintained long-term.

AI performance in healthcare is dynamic. Real-world safety depends not only on pre-deployment validation, but on **ongoing monitoring, bias detection, and clinical review**.

Recommended actions:

- Encourage or require AI-enabled care tools to demonstrate:
 - Continuous performance monitoring
 - Equity and bias assessment across populations
 - Continuous analysis of hallucinations and conversational drift
 - Structured escalation and override mechanisms
- Support these practices through:
 - CMS quality programs
 - Innovation Center (CMMI) models
 - ONC-aligned certification criteria

Real-World Impact:

Establishing suggested requirements for ongoing monitoring and quality checks will help increase public trust while avoiding rigid pre-market barriers that slow innovation.

4. Reduce Administrative Burden Through Standardized AI Evaluation Frameworks

Lower adoption friction for providers and health plans must be a priority to help mitigate administrative burden. Today, each health system often develops its own AI review and approval process, creating duplication, delay, and inconsistent standards.

Recommended actions:

- Support development of **standardized, industry-aligned AI evaluation frameworks** addressing:
 - Clinical validity
 - Safety and escalation logic
 - Data privacy and security
 - Workflow integration
 - Encourage voluntary adoption through federal programs rather than mandates.

Real-World Impact:

Standardization reduces procurement friction while preserving organizational autonomy and clinical judgment.

5. Use Federal Programs to De-Risk Early Adoption

HHS has the unique ability to reduce risk through targeted programs and demonstrations that will accelerate the use of AI from pilots to use at scale across the healthcare continuum.

Recommended actions:

- Expand CMMI demonstrations that explicitly include AI-enabled care management
- Support Medicaid pilots focused on high-need, high-cost populations
- Encourage public-private partnerships that evaluate:
 - Cost reduction
 - Patient safety
 - Workforce impact

Real-World Impact:

Early federal participation accelerates learning, builds evidence, and lowers barriers for broader market adoption.

6. Center Patient Safety, Empathy, and Transparency in Policy Design

Maintaining public trust while accelerating adoption is a top priority and will need to be coupled with development of policies that explicitly reinforce the use of effective AI:

- Enhances patient care without replacing human caregivers
- Is transparent to patients and caregivers
- Operates within clear clinical guardrails

Recommended actions:

- Encourage patient-facing disclosures that explain AI use in plain language
- Promote AI designs that prioritize engagement, accessibility, and equity
- Support research into patient experience and trust as core quality metrics

Real-World Impact:

Patient trust is foundational. Defining policies to help maintain public trust and focus on adoption of AI solutions designed with empathy and transparency can mitigate any risk of backlash or long-term harm.

Question 2: Summary

HHS can most effectively incentivize and accelerate AI adoption by:

- Aligning payment with prevention and engagement
- Clarifying governance for non-medical device AI
- Supporting continuous quality assurance
- Reducing administrative friction
- Using federal programs to de-risk innovation

These actions would accelerate safe, empathetic, and cost-reducing AI use in clinical care and preserve patient safety and public trust.

QUESTION 3:

For non-medical devices, what novel legal and implementation issues exist that challenge existing governance and accountability structures (e.g., liability, indemnification, privacy, and security)? What role, if any, should HHS play to help address them?

1. Ambiguity Around Accountability and Liability

One of the most significant challenges is **unclear allocation of responsibility** when AI informs clinical workflows but does not make independent medical decisions.

Key questions health systems regularly face include:

- Who is accountable for an AI-generated risk signal that is not acted upon?
- How should responsibility be shared between the technology vendor and the deploying organization?
- What constitutes appropriate clinical reliance versus over-reliance on AI insights?

In the absence of clear guidance, organizations often default to conservative positions regarding limiting AI use or requiring excessive manual review. Both tactics create inefficiencies, slow adoption, and diminish the ability to realize positive outcomes when AI is deployed.

Real-World Impact:

AI tools that could safely surface early warning signs are sometimes constrained to advisory roles with limited operational follow-through, not due to safety concerns, but due to fear of downstream liability exposure.

Potential HHS role:

HHS can help by:

- Clarifying expectations for **human-in-the-loop accountability**
- Providing guidance on reasonable reliance standards for AI that informs—but does not replace—clinical judgment
- Encouraging shared responsibility models aligned with existing care team structures

2. Indemnification and Contractual Complexity Slow Adoption

Procurement of non-medical device AI often involves prolonged negotiations around indemnification, insurance coverage, and risk allocation, particularly when AI outputs may influence care prioritization.

Because existing contracts were designed for traditional software or medical devices, they often fail to reflect:

- Continuous learning systems
- Post-deployment performance monitoring
- Shared oversight between vendor and provider

Real-World Impact:

Months-long contracting delays frequently occur even after clinical and operational alignment is achieved, slowing innovation and increasing administrative cost.

Potential HHS role:

HHS could support:

- Model contract language or best-practice guidance for AI-enabled care tools
- Federal frameworks that encourage proportionate indemnification based on AI risk level and use case

3. Privacy and Security in Continuous, Conversational AI

AI-enabled care management increasingly relies on **natural, conversational interactions** with patients, which may surface sensitive behavioral or social information not always captured in structured clinical encounters.

While HIPAA provides a strong foundation, challenges remain around:

- Appropriate data minimization
- Transparency in patient-facing AI interactions
- Secondary use of derived insights (e.g., risk scores, sentiment analysis)

Real-World Impact:

Organizations often apply overly restrictive interpretations of privacy requirements, limiting AI functionality or data integration even when patient consent and safeguards are in place.

Potential HHS role:

HHS can help by:

- Clarifying acceptable practices for conversational and derived data within HIPAA
- Encouraging transparency standards that explain AI data use in accessible language
- Supporting privacy-by-design approaches that protect patients while enabling innovation

- Clarifying regulations for outbound communications (e.g., TCPA) as it relates to AI for care management operations

4. Governance Gaps for AI That Operates Between Clinical and Operational Domains

Many non-medical device AI tools operate at the intersection of clinical care, operations, and population health. As a result, they may fall outside traditional governance bodies such as:

- Pharmacy & therapeutics committees
- Medical device review boards
- IT security councils

The misalignment with these groups can lead to:

- Fragmented oversight
- Redundant reviews
- Unclear ownership

Real-World Impact:

AI adoption stalls because no single internal body feels empowered to approve or oversee the technology holistically.

Potential HHS role:

HHS could encourage governance models that:

- Recognize AI as a cross-functional clinical capability
- Include nursing, care management, compliance, and patient safety leadership
- Emphasize ongoing oversight rather than one-time approval

5. Patient Trust, Transparency, and Informed Use

Patients are increasingly aware when AI is part of their care experience. For non-medical device AI, uncertainty exists around:

- How much disclosure is appropriate
- How to explain AI involvement without creating fear or confusion
- How to ensure patients feel supported rather than automated

Real-World Impact:

Inconsistent disclosure practices lead to variable patient trust and hesitation among care teams about how to introduce AI-enabled tools.

Potential HHS role:

HHS can support:

- Plain-language disclosure standards for AI in care delivery
- Guidance emphasizing AI as an extension of the care team—not a replacement
- Research into patient understanding and trust related to AI-supported care

Question 3: Summary

The primary legal and implementation challenges for non-medical-device AI stem from **ambiguity, not inherent risk**. When accountability, privacy expectations, and governance structures are unclear, organizations default to caution—slowing adoption even when patient safety and cost reduction would benefit.

HHS can play a critical enabling role by:

- Clarifying accountability and oversight expectations
- Supporting proportionate governance and indemnification models
- Reinforcing transparency, empathy, and patient trust
- Encouraging continuous quality assurance over rigid pre-approval
- These actions would accelerate responsible AI adoption while preserving patient safety, clinician accountability, and public confidence.

QUESTION 4:

For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and other workflow and human-centered evaluation methods for clinical care? Should HHS further support these processes, and if so, how?

For non-medical device AI used in clinical care in areas such as care management, patient engagement, risk identification, and clinical prioritization, traditional pre-market evaluation models are insufficient on their own. These tools derive much of their value from real-world deployment, continuous learning, and interaction with humans. As a result, **evaluation must be ongoing, contextual, and human-centered, and expand beyond static performance** benchmarks.

1. Pre-Deployment Evaluation: Clinical Appropriateness and Guardrails

Before deployment, the most important evaluation questions include *“Is the model accurate?”* and *“Is the model appropriate, safe, and clearly bounded in its role?”*

Promising pre-deployment evaluation methods include:

- Validation of AI outputs against clinically accepted reference standards
- Clear definition of intended use, limitations, and escalation thresholds
- Scenario-based testing to assess AI behavior under edge cases or ambiguous inputs
- Review by multidisciplinary clinical teams, including nursing, care management, and patient safety leaders

Real-World Impact:

AI systems that are clinically well-scoped and appropriately constrained are far safer than overly broad tools with unclear boundaries—even if raw accuracy metrics appear strong.

2. Post-Deployment Monitoring: Continuous Safety and Performance Assurance

Non-medical device AI performance evolves over time due to:

- Changes in patient populations
- Shifts in care delivery models
- Workflow adaptations
- Model updates

As a result, post-deployment evaluation is the cornerstone of patient safety. Key post-deployment metrics should include:

- Clinical signal detection rates and false-positive/false-negative trends
- Escalation appropriateness and timeliness
- Patient engagement and response consistency
- Disparities in performance across demographic groups
- Downstream clinical and utilization outcomes

Real-World Impact:

An AI system that identifies risk accurately but fails to engage patients consistently, or generates signals that overwhelm care teams, can introduce safety and quality concerns despite strong technical performance.

3. Human-Centered and Workflow-Based Evaluation

AI in clinical care does not operate in isolation. It operates within complex clinical workflows and human decision-making environments. Promising evaluation approaches include:

- Workflow impact assessments to understand how AI changes care team behavior
- Cognitive load and alert fatigue analysis
- Qualitative feedback from frontline clinicians and care managers
- Evaluation of how AI supports decision-making rather than overrides it

Real-World Impact:

Many AI failures in healthcare are not algorithmic failures—they are integration failures. Human-centered evaluation identifies these risks early.

4. Engagement and Empathy as Core Evaluation Metrics

For AI-enabled care management, **engagement is a leading indicator of both safety and effectiveness.**

Evaluation frameworks should explicitly measure:

- Patient participation and retention over time
- Completion rates of AI-supported interactions
- Patient-reported experience and understanding
- Willingness to disclose sensitive information
- Patient reported behavior change as a result of AI-supported interactions

Real-World Impact:

AI that fails to engage patients consistently cannot surface emerging risk reliably—creating blind spots that undermine safety and cost reduction goals. Improving health outcomes frequently starts with modifying the critical behaviors that are leading indicators, such as exercise, diet and sleep. AI agents are ideal for educating patients in these areas and helping them set goals and track their progress.

5. Robustness and Bias Testing in Real-World Conditions

Equitable AI requires evaluation beyond controlled datasets. **Promising practices include:**

- Performance monitoring across race, ethnicity, age, language, and socioeconomic status
- Stress-testing AI performance in populations with low digital literacy or limited access
- Ongoing bias audits rather than one-time assessments

Real-World Impact:

AI deployed at scale must perform safely and effectively across diverse populations, particularly those at highest risk and highest cost.

6. Role for HHS in Supporting AI Evaluation

HHS can play a critical role in accelerating adoption by **standardizing and supporting evaluation processes without imposing rigid, one-size-fits-all regulation.**

High-impact support mechanisms include:

- Development of reference evaluation frameworks for non-medical device clinical AI
- Grants and cooperative agreements focused on real-world implementation science
- CMMI demonstrations incorporating AI evaluation requirements
- Public-private partnerships to establish shared benchmarks and best practices

Real-World Impact:

Clear, consistent evaluation expectations reduce duplication, accelerate learning, and increase trust, while preserving flexibility for innovation.

Question 4: Summary

The most effective evaluation methods for non-medical-device AI emphasize:

- Clear clinical scope and guardrails
- Continuous post-deployment monitoring
- Human-centered workflow assessment
- Engagement and empathy as safety signals
- Equity and robustness in real-world use

HHS can accelerate safe, scalable AI adoption by supporting evaluation approaches that reflect **how care is actually delivered**, rather than forcing emerging tools into legacy frameworks.

QUESTION 5:

How can HHS best support private sector activities (e.g., accreditation, certification, industry-driven testing, and credentialing) to promote innovative and effective AI use in clinical care?

The private sector is already actively developing mechanisms to evaluate, certify, and govern AI used in clinical care. However, these efforts remain fragmented, inconsistent, and unevenly trusted across the healthcare ecosystem. HHS can play a catalytic role by **endorsing alignment, reducing duplication, and reinforcing patient safety and clinical accountability, without centralizing certification or stifling innovation.**

1. Endorse Principles-Based Accreditation Over Prescriptive Certification

Given the rapid evolution of AI, rigid, one-time certification models risk becoming obsolete quickly. Instead, HHS should encourage **principles-based accreditation frameworks** that emphasize ongoing safety, performance, and governance.

Key principles HHS could reinforce include:

- Clearly defined clinical scope and intended use
- Human-in-the-loop oversight and escalation pathways
- Continuous monitoring and quality assurance
- Equity and bias evaluation
- Transparent communication with patients and care teams

Real-World Impact:

Principles-based approaches allow innovation while maintaining consistent expectations around safety and accountability.

2. Support Industry-Driven Testing and Benchmarking

Private sector organizations are well positioned to develop **shared testing methodologies and benchmarks** for non-medical device AI, particularly in areas such as:

- Engagement and retention
- Workflow impact
- Bias and robustness
- Post-deployment performance

Recommended HHS role:

- Encourage convergence around common metrics and testing protocols
- Support public-private partnerships to develop benchmark datasets where appropriate
- Avoid mandating a single testing model, allowing specialization by use case

Real-World Impact:

Shared benchmarks reduce redundancy and enable more efficient procurement and evaluation across health systems.

3. Encourage Credentialing of AI Governance & Clinical Oversight, Not Just Technology

Effective AI adoption depends as much on **organizational readiness** as on model performance. HHS can support credentialing approaches that evaluate:

- Internal AI governance structures
- Clinical oversight models
- Training and competency of care teams using AI
- Processes for incident review and escalation

Real-World Impact:

AI safety failures are often governance failures, not algorithmic ones. Credentialing governance capabilities reinforces patient safety at scale and improves adoption of AI beyond the technology aspects of a solution.

4. Recognize Engagement and Empathy as Quality Signals

Accreditation and certification efforts should move beyond technical accuracy and explicitly incorporate **patient engagement and experience metrics**.

HHS can support frameworks that evaluate:

- Consistent patient participation
- Accessibility across diverse populations
- Patient-reported understanding, behavior change and trust
- Cultural and linguistic appropriateness

Real-World Impact:

AI that fails to engage patients reliably cannot deliver safe or equitable care and is not linked to technical performance. Measuring outcomes, such as earlier risk identification and preventative care, hinge on patient engagement levels and experience throughout their care journey.

5. Reduce Administrative Burden Through Mutual Recognition

One of the most significant barriers to adoption is the **repetition of AI reviews across organizations**.

Recommended HHS actions:

- Encourage mutual recognition of credible private sector accreditation programs
- Promote portability of evaluations across federal and state programs
- Discourage duplicative reviews that add cost without improving safety

Real-World Impact:

Reducing administrative friction by creating a systematic approach to portability of evaluation and accreditation programs assists organizations in making decisions supported by repeatable data using consistent review standards and help accelerates adoption while maintaining rigorous oversight.

6. Use Federal Programs to Legitimize and Scale Best Practices

HHS can reinforce private sector accreditation and credentialing by referencing them within:

- CMS Innovation Center models
- ONC interoperability and certification initiatives
- Federal procurement and grant requirements

Real-World Impact:

Federal recognition provides signal value without imposing top-down mandates. The signal value reduces costs and expedites deployment of AI in impactful areas of healthcare, such as care management.

Question 5: Summary

HHS can most effectively promote innovative and safe AI use by:

- Supporting principles-based accreditation
- Encouraging industry-led testing and benchmarks
- Elevating governance and oversight credentialing
- Recognizing engagement and empathy as safety indicators
- Reducing duplicative administrative burden

By acting as a **convener and validator**, rather than a sole certifier, HHS can accelerate responsible AI adoption while preserving patient safety, clinician trust, and market innovation.

QUESTION 6:

Where have AI tools deployed in clinical care met or exceeded performance and cost expectations, and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?

AI deployment in clinical care has produced **uneven results**. Success has been most consistent where AI is designed to augment human care, operate within clear clinical boundaries, and focus on early intervention and engagement. Conversely, AI has fallen short where it is deployed without sufficient workflow integration, patient-centered design, or governance.

Where AI Has Met or Exceeded Expectations

1. Early Risk Identification and Stratification

AI tools have performed particularly well in identifying **early behavioral, emotional, and psychosocial risk** that often precedes high-cost medical events.

Observed strengths include:

- Surfacing depression, anxiety, cognitive stress, and disengagement earlier than traditional screening
- Prioritizing high-risk patients for timely human follow-up
- Improving consistency of assessments across large populations

Real-World Impact:

Earlier identification enables lower-cost interventions and reduces avoidable utilization, particularly in value-based care models.

2. AI-Enabled Care Management and Outreach

AI tools that support **frequent, empathetic patient engagement**, especially via conversational modalities, have exceeded expectations in:

- Reaching patients who are traditionally difficult to engage
- Maintaining consistent touchpoints between clinical visits
- Supporting care teams with actionable insights rather than raw data

Real-World Impact:

These tools extend care team capacity, reduce administrative burden, and support proactive care at scale.

3. Operational Efficiency and Workforce Support

AI has delivered meaningful gains when used to:

- Automate routine documentation or screening tasks
- Reduce manual triage workload
- Support nurses and care managers in prioritizing outreach

Real-World Impact:

When thoughtfully implemented, AI reduces burnout and allows clinicians to operate at the top of their license.

Where AI Has Fallen Short**1. Low Patient Engagement in Transactional AI**

AI tools that rely on impersonal, form-based, or transactional interactions have often underperformed due to:

- Low response rates
- High drop-off over time
- Limited patient trust

Lesson learned:

Without empathy and usability, even technically sound AI fails to deliver clinical or financial value.

2. Poor Workflow Integration

AI solutions that operate outside existing clinical workflows and require duplicate systems or manual workarounds, frequently fail to scale.

Lesson learned:

Integration and change management are as critical as model performance.

3. Overreliance on Static Accuracy Metrics

Some AI tools have been deployed based on impressive accuracy metrics but failed to account for:

- Alert fatigue
- Bias in real-world populations
- Downstream operational burden

Lesson learned:

Static metrics do not predict real-world safety or impact on patient care, clinical team burden, or ability to control cost. Without proper use of data, and reliance only on accuracy metrics, AI will struggle to have wide adoption within healthcare.

Novel AI Tools With Greatest Future Potential**1. Empathetic, Conversational AI for Continuous Care**

AI that enables **natural, frequent, and supportive patient interactions** holds significant promise for:

- Behavioral health
- Chronic disease management
- Post-acute and transitional care

These tools can surface risk earlier and support patients between visits, where most deterioration occurs.

2. AI That Integrates Behavioral, Social, and Clinical Signals

Tools that synthesize **multidimensional data**, rather than siloed metrics, can:

- Provide earlier and more accurate risk signals
- Improve care prioritization
- Reduce unnecessary utilization

3. AI Focused on Engagement as a Clinical Outcome

AI designed to **measure and optimize engagement itself**, not just clinical indicators, can improve:

- Equity
- Safety
- Long-term outcomes

4. AI That Supports, Rather Than Replaces, Clinical Judgment

The most promising tools are those that:

- Enhance situational awareness
- Provide explainable insights
- Preserve clinician accountability

Question 6: Summary

AI has delivered the greatest value in clinical care when it:

- Identifies risk early
- Engages patients consistently and empathetically
- Operates within strong clinical guardrails
- Supports workforce sustainability

AI has fallen short where:

- Engagement was low
- Workflows were disrupted
- Governance was unclear

Future innovation should prioritize **empathetic engagement, continuous care, and clinician-supported decision-making** to meaningfully improve outcomes and reduce costs.

QUESTION 7:

Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?

Adoption of AI for clinical care is rarely driven by a single role or department. Instead, it reflects a **complex, multi-stakeholder decision-making process** shaped by clinical leadership, operational priorities, risk management, and financial incentives. Understanding these dynamics is critical to accelerating responsible AI adoption at scale.

Key Decision-Makers and Influencers

1. Clinical and Nursing Leadership

In practice, **nursing leaders, care management directors, and population health executives** exert outsized influence on whether AI-enabled care tools are adopted and sustained.

These leaders evaluate AI based on:

- Impact on patient safety and quality
- Alignment with care workflows
- Effect on workforce burden and burnout
- Ability to engage patients consistently

Why this matters:

AI that lacks nursing and care management buy-in, rarely moves beyond pilot phase, regardless of technical performance.

2. Physician Leadership and Medical Governance

Physician executives and medical committees play a critical role in assessing:

- Clinical appropriateness
- Risk to patient safety
- Alignment with standards of care

While physician leaders may not be day-to-day users of care management AI, their endorsement is often required for system-wide deployment.

3. Quality, Compliance, and Risk Management

Quality and compliance teams influence adoption by:

- Interpreting regulatory risk
- Evaluating patient safety implications
- Reviewing liability exposure

In the absence of clear external guidance, these teams may take conservative positions that delay or limit AI deployment.

4. Information Technology and Security Leadership

IT and security leaders assess:

- Integration complexity
- Data privacy and cybersecurity risk
- System reliability and scalability

Their support is essential, but AI adoption often stalls when evaluation criteria are optimized for traditional IT systems rather than adaptive, learning technologies.

5. Finance and Value-Based Care Leadership

Financial leaders increasingly influence AI adoption in:

- Medicare Advantage
- Medicaid managed care
- ACO and risk-bearing models

They assess whether AI:

- Reduces total cost of care
- Improves quality metrics
- Supports risk adjustment and performance benchmarks

Primary Administrative Hurdles

1. Fragmented and Redundant Review Processes

Many organizations lack a unified AI governance pathway. As a result, AI tools may be reviewed separately by:

- Clinical committees
- IT security
- Compliance
- Legal and contracting

This leads to duplication, delays, and inconsistent decision-making.

2. Lengthy Procurement and Contracting Cycles

AI procurement often triggers:

- Extended indemnification negotiations
- Custom risk assessments
- Repeated data security reviews

Even when clinical and operational value is clear, contracting timelines can exceed 6–12 months.

3. Lack of Internal Ownership

AI tools that span clinical, operational, and patient engagement domains may lack a clear internal “owner.” Without defined accountability:

- Decisions are delayed
- Governance is fragmented
- Adoption stalls

4. Change Management and Training Burden

AI adoption requires:

- Workflow redesign
- Training and trust-building
- Ongoing monitoring and feedback

Organizations under workforce strain may struggle to allocate resources to these efforts.

Question 7: Summary

AI adoption in clinical care is most strongly influenced by:

- Nursing and care management leadership
- Medical governance
- Quality and risk management
- IT and security
- Financial leadership in risk-based models

The greatest administrative barriers are not technical, but **organizational and procedural**, including fragmented governance, slow procurement, and lack of ownership. Reducing these barriers requires:

- Clear external guidance on AI governance
- Streamlined evaluation frameworks
- Recognition of nursing and care management leadership as central decision-makers

Addressing these realities will accelerate safe, effective AI adoption while preserving patient safety and clinician trust.

QUESTION 8:

Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

Interoperability is foundational to the safe, effective, and scalable use of AI in clinical care. While progress has been made in exchanging core clinical data, **many of the highest-impact AI use cases, particularly those focused on early risk detection, care management, and patient engagement, remain constrained by fragmented, incomplete, or delayed data access.**

Enhanced interoperability would significantly widen market opportunities, accelerate research, and improve patient outcomes by enabling AI to operate with a **longitudinal, person-centered view of health** rather than isolated clinical snapshots.

1. Integrating Clinical, Behavioral, and Social Data

Many of the drivers of healthcare cost and utilization sit **outside traditional medical data**. AI tools designed to identify early risk require access to a broader set of signals, including:

- Behavioral health indicators
- Medication adherence and self-reported symptoms
- Social determinants of health (SDOH)
- Engagement and communication patterns
- Patient Reported Outcome Measures

Current challenge:

These data are often stored in disconnected systems that include EHRs, care management platforms, community-based organizations, and patient-facing tools. If data is not connected through integration, AI's ability to synthesize meaningful insight is limited.

Real-World Impact:

Interoperability that enables secure, standardized exchange across these domains would dramatically improve AI's ability to detect risk early and support targeted intervention.

2. Longitudinal and Near–Real-Time Data Access

AI-enabled care management is most effective when it can detect **change over time**, not just static conditions.

Key gaps include:

- Delayed data feeds from claims and EHR systems
- Limited access to longitudinal engagement and outreach data
- Inconsistent linkage across care settings

Real-World Impact:

Improved interoperability for longitudinal data, particularly near–real-time signals, would enable AI to:

- Identify emerging deterioration sooner
- Reduce false positives
- Improve escalation accuracy

3. Standards That Support Care Management and Engagement

While FHIR has advanced clinical data exchange, many care management and engagement workflows rely on **data elements that are not yet consistently standardized or exchanged**, such as:

- Outreach attempts and response patterns
- Risk assessments and prioritization outputs
- Care plan progression and task completion
- Patient-reported experience and sentiment

Real-World Impact:

Expanding standards and implementation guidance to include care management and engagement data would unlock new AI use cases focused on prevention and population health.

4. Interoperability for Evaluation and Benchmarking

AI evaluation and trust depend on the ability to measure performance across settings and populations.

Current limitations:

- Inconsistent outcome definitions
- Limited access to comparable benchmarking datasets
- Difficulty linking AI insights to downstream utilization and outcomes

Real-World Impact:

HHS can support shared benchmarking by:

- Encouraging standardized outcome and utilization measures
- Supporting de-identified, privacy-preserving datasets for research
- Enabling cross-organizational comparison without exposing proprietary data

5. Patient-Mediated Data Exchange

Patients increasingly interact with AI-enabled tools outside traditional care settings. Interoperability should support **patient-mediated data sharing**, allowing individuals to:

- Understand how their data is used
- Share insights across care teams
- Retain agency and trust

Real-World Impact:

Patient-centered interoperability reinforces transparency, engagement, and safety—while expanding AIs reach

6. Role for HHS in Accelerating Interoperability

HHS can accelerate progress by:

- Prioritizing interoperability that supports prevention, engagement, and care management—not just billing and documentation
- Encouraging adoption of standards that reflect real-world workflows
- Supporting pilots and demonstrations that integrate clinical, behavioral, and social data
- Avoiding one-size-fits-all mandates that limit innovation

Question 8: Summary

Enhanced interoperability would most accelerate AI adoption where it enables:

- Integration of clinical, behavioral, and social data
- Longitudinal, near–real-time insight
- Standardization of care management and engagement workflows
- Shared evaluation and benchmarking

By focusing on **person-centered data exchange**, HHS can unlock AI's full potential to improve outcomes, reduce costs, and support safer, more empathetic care at scale.

QUESTION 9:

What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?

From the patient and caregiver perspective, the promise of AI is not abstract innovation—it is relief from the daily friction, gaps, and delays that characterize much of modern healthcare. Patients consistently express interest in AI that makes care **more accessible, responsive, and human**, while voicing concerns when AI feels opaque, impersonal, or unsafe.

What Patients and Caregivers Want AI to Address

1. Lack of Ongoing Support Between Visits

Patients frequently experience healthcare as episodic, with long gaps between appointments, especially those managing chronic conditions, behavioral health needs, or caregiving responsibilities.

Patient need:

- Ongoing check-ins and monitoring
- A way to signal concern before problems escalate
- Support that does not require scheduling or travel

Real-World Impact:

AI-enabled care management can provide consistent, low-burden touchpoints that surface emerging needs early and route them to appropriate human support.

2. Difficulty Being Heard and Understood

Many patients report that traditional encounters do not fully capture:

- Emotional distress
- Social barriers
- Cognitive or behavioral changes
- Subtle symptom progression

Patient needs:

- Interactions that feel empathetic and non-judgmental
- Time to explain concerns without feeling rushed

Real-World Impact:

Conversational, empathetic AI can create space for disclosure and help ensure important signals reach care teams.

3. Complexity and Fragmentation of Care

Patients and caregivers often shoulder the burden of navigating:

- Multiple providers
- Conflicting instructions
- Repeated assessments

Patients need:

- Better coordination
- Fewer redundant interactions
- Clear next steps

Real-World Impact:

AI can support care coordination by identifying risk, prioritizing outreach, and reducing duplicative tasks—when integrated thoughtfully.

4. Access Barriers and Health Equity Gaps

Rural populations, older adults, individuals with mobility limitations, and those with behavioral health needs often face disproportionate access challenges.

Patients need:

- Care that meets patients where they are during their care continuum
- Tools that are intuitive and inclusive

How AI helps:

Voice-based and low-friction AI modalities reduce reliance on screens, apps, and transportation, expanding access to support.

Patient and Caregiver Concerns About AI

1. Fear of Losing Human Connection

The most common concern is that AI will replace human care or make healthcare feel transactional.

Key concern:

“Will AI be used to avoid talking to me, rather than to support me?”

Mitigation:

AI must be positioned and designed as an extension of the care team, with clear escalation to humans and visible clinician involvement.

2. Transparency and Understanding

Patients often do not know:

- When AI is being used
- What decisions it influences
- How their information is interpreted

Key concern:

“Who is listening, and how is this information used?”

Mitigation:

Plain-language disclosures and patient education build trust and informed participation.

3. Privacy and Data Use

Patients express concern about:

- How sensitive information is stored
- Who has access
- Whether data could be misused

Mitigation:

Strong privacy protections, data minimization, and transparent policies are essential.

4. Accuracy, Bias, and Fairness

Patients worry that AI may:

- Misinterpret them
- Perform worse for certain populations
- Miss context that a human would understand

Mitigation:

Ongoing bias monitoring, human oversight, and inclusive design reduce these risks.

Empathy and Engagement: Foundational Components to Achieve AI Safety

From the patient perspective, **engagement and empathy are not optional features—they are safety mechanisms**. AI systems that patients trust and engage with consistently are more likely to surface early warning signs and less likely to miss critical information.

AI that patients avoid or distrust introduces risk, regardless of technical accuracy.

Question 9: Summary

Patients and caregivers want AI to:

- Provide support between visits
- Improve access and coordination
- Help them feel heard and understood
- Reduce burden and fragmentation

Their concerns focus on:

- Loss of human connection
- Lack of transparency
- Privacy and fairness

AI adoption that centers **empathy, engagement, and strong clinical guardrails** can address these concerns while delivering meaningful improvements in safety, outcomes, and cost.

QUESTION 10:

Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care?

A. Are there published findings about the impact of adopted AI tools and their use in clinical care?

B. How does the literature approach the costs, benefits, and transfers of using AI as part of clinical care?

To accelerate responsible AI adoption in clinical care, HHS should prioritize research that reflects **how AI is actually used in practice**—embedded in care teams, interacting with patients, and influencing outcomes over time. While foundational AI research remains important, the greatest near-term gains will come from **implementation-focused, patient-centered, and workforce-aware research agendas**.

1. Implementation Science for AI in Real-World Care Delivery

One of the largest evidence gaps is not whether AI *can* work, but **how it works when deployed at scale** across diverse settings and populations.

Priority research areas include:

- AI integration into clinical and care management workflows
- Change management and adoption dynamics among nursing and care teams
- Longitudinal performance and safety monitoring
- Organizational governance models for AI oversight

Why this matters:

Many promising AI tools fail not because of poor accuracy, but because implementation challenges prevent sustained use or impact.

2. Research on Engagement, Empathy, and Human-AI Interaction

Patient engagement is a leading indicator of both safety and effectiveness yet remains underrepresented in AI research.

Priority research areas include:

- Comparative effectiveness of conversational vs. transactional AI
- Patient trust, disclosure, and understanding in AI-supported care
- Caregiver interaction with AI-enabled tools
- Relationship between engagement consistency and clinical outcomes

Why this matters:

AI that patients do not engage with cannot improve outcomes or reduce costs, regardless of technical performance.

3. Evaluation of AI in Behavioral Health and Comorbid Conditions

Behavioral health and psychosocial factors are among the strongest drivers of healthcare utilization and cost yet are often under-detected.

Priority research areas include:

- AI for early detection of behavioral and emotional distress
- Integration of behavioral signals into medical care pathways
- Impact of AI-enabled behavioral health screening on downstream utilization

Why this matters:

Improving early identification and support in behavioral health has outsized potential to reduce total cost of care.

4. Workforce Impact and Care Team Sustainability

AI research has historically focused on patient outcomes while under examining workforce effects.

Priority research areas include:

- Impact of AI on nursing workload and burnout
- Redistribution of tasks across care teams
- Training and competency development for AI-supported care

Why this matters:

AI adoption will not scale if it increases burden or undermines workforce sustainability.

5. Equity, Bias, and Performance Across Populations

While bias evaluation is increasingly discussed, real-world evidence remains limited.

Priority research areas include:

- AI performance across races, ethnicity, age, language, and socioeconomic status
- Engagement disparities and digital inclusion
- Methods for continuous bias monitoring post-deployment

Why this matters:

Equitable AI performance is essential for safety, trust, and national adoption.

6. Cost, Benefit, and Transfer Effects of AI Adoption

The literature on AI cost impact remains fragmented and inconsistent.

Current gaps include:

- Limited analysis of avoided utilization (e.g., ED visits, hospitalizations)
- Underrepresentation of indirect benefits such as workforce efficiency
- Insufficient attention to cost and benefit transfers between stakeholders (e.g., providers vs. payers)

Priority research areas include:

- Total cost of care impact over time
- Attribution models for AI-enabled prevention
- Financial effects across different payment models
- Labor costs for equivalent tasks performed by AI

Why this matters:

Clear evidence of cost deflation is essential to align incentives and accelerate adoption.

7. Public-Private Partnerships and Pragmatic Research Models

HHS is uniquely positioned to convene stakeholders and support research that would be difficult for the private sector alone.

High-impact mechanisms include:

- Cooperative research and development agreements (CRADAs)
- CMS Innovation Center demonstrations
- Grants focused on real-world implementation and evaluation
- Pragmatic trials embedded in care delivery

Question 10: Summary

To accelerate AI adoption in clinical care, HHS should prioritize research that:

- Focuses on real-world implementation and governance
- Measures engagement, empathy, and human-AI interaction
- Addresses behavioral health and comorbid drivers of cost
- Evaluates workforce impact and sustainability
- Produces clear evidence on cost, benefit, and transfer effects
- Advances equity and public trust

By investing in **applied, patient-centered, and workforce-aware AI research**, HHS can accelerate adoption that is safe, empathetic, and capable of meaningfully improving outcomes while deflating healthcare costs.

Executive Closing Statement

Artificial intelligence has reached a point of maturity where it can meaningfully improve patient safety, strengthen care delivery, and reduce the total cost of care—*but only if it is deployed with empathy, accountability, and strong clinical governance*. Across our responses to this RFI, a consistent theme emerges: the greatest opportunity for AI in clinical care lies not in automation for its own sake, but in **augmenting human care teams, engaging patients earlier and more consistently, and preventing avoidable deterioration before it becomes costly or harmful**.

Patient Safety as the Foundation for AI Adoption

Patient safety must remain the cornerstone of AI deployment in healthcare. The most effective AI systems are those that operate within **clearly defined clinical guardrails**, incorporate human oversight, and are continuously monitored in real-world use. Safety is strengthened—not weakened—when AI is used to surface early risk, prioritize outreach, and support clinical judgment rather than replace it.

Equally important, **engagement and empathy are safety features**. AI tools that patients trust and interact with consistently are more likely to detect emerging concerns and less likely to miss critical signals. Conversely, AI that patients disengage from introduces blind spots that undermine safety, regardless of technical accuracy.

Improving Outcomes Through Early, Continuous, and Equitable Care

Across clinical settings, the highest-cost and highest-risk events are often preceded by unmet behavioral, emotional, or social needs. AI-enabled care management offers a practical mechanism to:

- Detect risk earlier
- Maintain consistent support between visits
- Improve coordination across fragmented care environments
- Extend the reach of constrained clinical workforces

When deployed responsibly, these tools improve outcomes for patients while supporting clinicians and caregivers who are increasingly stretched thin.

Reducing Total Medical Spend by Preventing Escalation

AI's most powerful economic impact lies in **preventing downstream utilization**, not optimizing already expensive interventions. Early identification of risk, proactive engagement, and targeted escalation reduce:

- Emergency department visits
- Hospital admissions and readmissions
- Avoidable complications of chronic disease

However, today's payment and regulatory structures often fail to reward prevention. Without aligned incentives, even proven AI solutions struggle to scale beyond pilots.

The Need for Federal Leadership, Support, and Incentives

HHS has a unique and critical role to play in accelerating safe and effective AI adoption. Federal leadership can:

- Align reimbursement with prevention and engagement
- Clarify governance expectations for non-medical-device AI
- Support standardized evaluation and accreditation frameworks
- Reduce administrative friction and duplication
- Invest in real-world implementation science and public-private partnerships

Importantly, federal action should **enable and coordinate private sector innovation**, not replace it. Clear guidance, proportionate oversight, and targeted incentives will allow payers, health systems, and providers to adopt AI responsibly and at scale.

A Call to Action

AI represents a rare opportunity to advance the shared goals of **patient safety, better outcomes, workforce sustainability, and cost containment**. Realizing this opportunity will require:

- Patient-centered design grounded in empathy and trust
- Strong clinical guardrails and continuous quality assurance
- Incentives that reward prevention rather than volume
- Federal policies that reduce uncertainty and accelerate adoption

With thoughtful leadership and collaboration, AI can help make healthcare safer, more humane, and more affordable, advancing the mission to improve health for all Americans while deflating the unsustainable growth of healthcare costs.