RE: HHS Health Sector AI RFI
RIN 0955-AA13
Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as part of Clinical Care
February 6, 2026

The Honorable Robert F. Kennedy, Jr.
Secretary
Department of Health and Human Services
Washington, DC

Dear Secretary Kennedy,

We appreciate the opportunity to submit comments in response to the Department of Health and Human Services Request for Information entitled "Accelerating the Adoption and Use of Artificial Intelligence as part of Clinical Care" (RIN 0955-AA13). Thank you for inviting public input on an issue of critical importance to the future of healthcare delivery and American competitiveness in clinical care AI.

As doctors, we are told to document everything from day one. If you don't document it, it didn't happen. Healthcare AI should be held to the same evidentiary standard. When I sign a clinical note, we can reconstruct everything, who wrote it, when, what changes were made, who accessed it. When AI generates that same note, we can't reconstruct anything. We don't know what model version ran, whether safety controls executed, or where the data traveled. That's the proof gap and it's the single largest barrier to deploying AI at the scale healthcare needs.

When an AI ambient scribe hallucinates a substance abuse diagnosis that was never discussed, or when a patient-facing chatbot discloses fabricated lab results, organizations cannot reconstruct what happened at the moment of inference. Current audit trails record that requests occurred, not which specific controls ran, with what parameters, or why they failed.

As a practicing physician, I recently encountered this firsthand. I was testing an ambient AI scribe for my healthcare organization. I had my appointment with the patient, then reviewed the AI output. The first time I used it, it actually generated a hallucination in the note. It stated that I had recommended lamotrigine for PTSD which was nowhere in the transcript that I was able to go back and check word for word.

The first concern was that it hallucinated and fabricated an off-label indication that, if it went into the medical record, could lead to a negative patient or physician outcome. But we also understand hallucinations will happen, we're never going to completely eliminate them.

Was there hallucination detection? What happened at that moment when the transcript went into the model? Did safety checks run? Did they fail, or were they never active?

I had no way to trace back through the system to understand where that failure occurred.But here’s the more consequential problem: without evidence infrastructure, we couldn’t learn from it. I had no way to trace the execution path. Without that evidence, I couldn’t improve the system, and my organization's rational response would be to slow or stop AI adoption entirely. That’s the real cost of the proof gap, not just risk, but forgone innovation.

This is a systemic gap across the industry. The infrastructure to capture execution traces just doesn't exist in most AI systems today. The result is a trust deficit that slows adoption across every stakeholder, physicians who can’t verify outputs, organizations that can’t satisfy their boards, and vendors who can’t close procurement cycles.

This pattern is not new. Every industry that scales a powerful technology eventually builds the evidence infrastructure to make that technology trustworthy at scale. Aviation did not debate whether to install flight data recorders, after enough incidents where investigators could not reconstruct what happened, the black box became standard equipment. Not because regulators demanded it first, but because airlines, insurers, and manufacturers recognized that the alternative was an industry that couldn’t learn from its own failures. Financial markets did not debate whether companies should maintain auditable books — Generally Accepted Accounting Principles emerged because capital markets could not function without verifiable records. Certificate Transparency for the internet did not require legislation — browser vendors simply began requiring cryptographic proof that TLS certificates were legitimately issued, and within five years it became universal infrastructure. In each case, the evidence layer was not a burden imposed on innovation. It was the infrastructure that made innovation scalable.

Healthcare AI is at precisely this inflection point. The question is not whether evidence infrastructure will become standard, it will. The question is whether American institutions build it, or import it.

Healthcare organizations want to deploy AI at scale. Many have the clinical vision and the institutional will. What they lack is the technical infrastructure to demonstrate to their boards, their insurers, their procurement committees, and their patients that safety mechanisms operated as designed for any given encounter.

Healthcare organizations are deploying AI systems that can claim compliance but cannot prove it both in pre and post deployment. The healthcare industry lacks technical infrastructure that provides standardized and immutable methods to demonstrate that safety mechanisms ran as designed when high risk AI decisions are made. Static compliance documentation, SOC 2 reports, architecture diagrams, policy attestations, these show controls *exist*, but cannot prove they *executed* for any given patient encounter.

GLACIS Technologies is a healthcare AI infrastructure company building the evidence layer for AI governance. We close the proof gap: the inability to verify that safety controls executed at the moment of inference

Our leadership team brings expertise spanning clinical medicine, FDA regulatory pathways for AI devices, healthcare compliance, and cryptographic systems engineering. We developed the first evidence-grade attestation infrastructure that enforces safety controls at runtime, provides independent third-party witnessing, and generates tamper-evident, independently verifiable receipts of safety control execution for every AI inference event.

We have developed ATLAS (Attestation and Transparency Layer for AI Systems) as an open standard for evidence-grade AI governance, with complete crosswalks to NIST AI RMF, ISO/IEC 42001, and EU AI Act requirements. ATLAS is open for adoption and co-development by any organization, a market-driven evidence framework, not a proprietary compliance tool.

Market forces are already creating demand for this infrastructure, independent of any single regulatory action. Insurance carriers are introducing AI-specific exclusions and requiring governance documentation that most organizations cannot produce. Healthcare procurement cycles for AI products routinely exceed twelve months, largely because organizations have no efficient way to verify vendor safety claims. International markets, including the EU AI Act (August 2026), are establishing evidence requirements that American AI companies will need to meet to compete globally. And state-level frameworks in Colorado (June 2026) and California (January 2027) reflect a bipartisan recognition that evidence infrastructure enables responsible deployment.

GLACIS addresses this through its safety and accountability infrastructure that generates verifiable evidence at runtime, not retrospectively. Our infrastructure enforces safety controls at run time, captures complete execution traces showing which guardrails ran, in what sequence, with pass/fail status, timestamps, and version identifiers, all cryptographically signed and independently verifiable by auditors, regulators, or legal counsel without access to vendor systems. For healthcare organizations facing "who owns the liability" questions when AI systems fail, we provide the chain-of-custody documentation that the market is demanding: procurement teams need it to evaluate vendors, boards need it to authorize deployment, and insurers need it to price risk.

The enclosed response to the RFI focuses specifically on the need for third-party verifiable enforcement and attestation infrastructure and technical standards as a priority for accelerating clinical AI adoption. We outline technical requirements via the ATLAS framework, and detail how evidence-grade systems enable healthcare organizations to enforce, attest, and demonstrate, not just claim. that AI safeguards protected patients at the moment of inference.

Our comments emphasize that accelerating AI adoption in clinical care requires not just innovation in models and applications, but parallel innovation in the evidence infrastructure that makes those systems trustworthy and accountable.

We have intentionally emphasized the distinction between system-level assurances (compliance certificates, vendor attestations, architectural documentation) and transaction-level proof (cryptographically verifiable evidence of what happened for specific patient interactions).

Three years of healthcare AI deployment have demonstrated a consistent pattern: governance frameworks without enforcement mechanisms, and compliance claims without verifiable evidence, create an adoption ceiling. Organizations that cannot prove safety cannot scale deployment — and the organizations with the greatest clinical need are often the most risk-averse.

We believe the fastest path to accelerating AI adoption in clinical care is not prescriptive regulation but evidence infrastructure that gives healthcare organizations the confidence to deploy. When a hospital can demonstrate to its board, its insurer, and its patients that AI systems operated safely with cryptographic proof, not just vendor promises — the risk calculus shifts from "we can't afford to adopt" to "we can't afford not to." That shift is what moves AI from pilot programs to production at scale.

We respectfully encourage HHS to prioritize research into open, interoperable evidentiary standards for healthcare AI, standards that enable vendors to produce independently verifiable proof that safety controls executed as designed, that PHI remained within customer boundaries, and that clinical decisions can be reconstructed with complete provenance. Such standards would represent the market-driven alternative to prescriptive regulation: organizations that can self-evidence don't need to be told how to be safe.

Thank you for your leadership on this important issue and for your commitment to ensuring that AI systems in healthcare can be trusted, audited, and proven safe.

Respectfully,

Jennifer Shannon, MD
Chief Medical Officer
GLACIS Technologies

Enclosure: GLACIS Response to HHS RFI Question #10 (AI Research Priorities for Evidence-Grade Compliance Infrastructure)

Contact Information:
GLACIS Technologies
jennifer@glacis.io
glacis.io

Expertise Areas:

- Cryptographic attestation for AI systems
- Healthcare regulatory compliance (EU AI Act, FDA, HIPAA)
- Clinical AI safety and governance
- Evidence-grade AI adoption infrastructure