

February 18, 2026

Re: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care (RIN 0955–AA13)

Dear Deputy Secretary and ASTP/ONC Leadership,

Thank you for the opportunity to provide input on this Request for Information. I am an independent healthcare technology consultant with a background in quality improvement and healthcare operations. My comments are informed by direct experience advising healthcare organizations on technology adoption, clinical workflow optimization, and regulatory compliance. I submit these comments as an individual.

AI adoption in clinical care is accelerating rapidly, and this momentum should be sustained. However, the speed of adoption must be matched with proportionate diligence to ensure that AI tools serve the best interests of patients and practitioners alike. My comments address four areas of the RFI where I believe HHS action can have the greatest impact.

Question 1: Barriers to Innovation and Adoption

Data fragmentation remains the most significant structural barrier to effective AI in clinical care. Health data continues to be siloed across delivery systems, payers, and public health entities, limiting the representativeness and quality of datasets that underpin AI models. When AI algorithms are trained on incomplete or non-representative data, the resulting tools risk perpetuating existing health disparities and producing biased clinical recommendations. Research has documented that AI bias in biomedicine can originate from insufficient data, sampling bias, and the use of health-irrelevant features (Yang et al., *Journal of Biomedical Informatics*, 2024, DOI: 10.1016/j.jbi.2024.104646).

HHS should prioritize policies that enable responsible data sharing across care settings while maintaining rigorous patient privacy protections. Breaking down data silos through expanded interoperability requirements, standardized data formats, and clear safe harbor provisions for data sharing in AI development will empower AI developers to build tools grounded in high-quality, representative datasets. High data quality is the prerequisite for high algorithmic accuracy, and high accuracy is the prerequisite for patient safety.

Additionally, small and medium-sized practices face disproportionate barriers to AI adoption due to limited capital, technical infrastructure, and administrative bandwidth. HHS should consider incentive structures (whether through grants, technical assistance programs, or favorable reimbursement policies) that lower the adoption threshold for these organizations, ensuring the benefits of clinical AI are not concentrated among large health systems.

An emerging dimension of this challenge is the growing ability of clinicians themselves to develop AI tools tailored to their practice. The increasing accessibility of open-source AI models

and low-code development platforms means that a physical therapist can now build a pose estimation tool for movement assessment, or a primary care clinic can develop a predictive scheduling model using their own operational data. These clinician-developed tools have the advantage of being built by domain experts who understand the clinical context intimately. However, most clinicians building these tools lack the regulatory and compliance expertise to evaluate whether their creation crosses the threshold into software as a medical device (SaMD) territory, or to conduct appropriate clinical validation. HHS should support this practitioner-led innovation through accessible compliance guidance, technical assistance resources, and grant programs that help clinicians build safe, validated tools while making clear that any AI tool used in clinical workflows must be developed and applied within the clinician's scope of practice and domain expertise.

Questions 3–4: Evaluation Methods, Legal Issues, and Accountability

For non-medical device AI tools used in clinical care, evaluation frameworks must reflect the clinical context in which these tools operate. Specifically, the calibration of false positive and false negative rates should be aligned to the clinical risk profile of the use case. An AI tool that assists in screening for a life-threatening condition requires a fundamentally different error tolerance than one optimizing appointment scheduling. HHS should encourage or support the development of risk-stratified evaluation standards that require AI developers to demonstrate how their models manage the trade-off between sensitivity and specificity relative to patient harm.

Liability and Accountability

A central governance challenge for clinical AI is aligning accountability with decision-making authority. I offer the following framework for the Department's consideration, built on the distinction between human-in-the-loop and autonomous AI architectures:

- **Augmented workflows (human-in-the-loop):** When an AI tool provides clinical decision support and a licensed clinician retains decision-making authority, accountability for patient outcomes should rest with the clinician, consistent with existing standards of professional practice. The AI tool is an instrument; the clinician is the decision-maker. This is no different from a clinician who misinterprets a diagnostic image or disregards a contraindication, the professional who exercises judgment bears responsibility for that judgment.
- **Autonomous workflows (no human-in-the-loop):** When an AI tool operates without clinician intervention at the point of clinical decision, accountability should be shared between the deploying organization and the AI product owner. The deploying health system bears responsibility analogous to its existing credentialing and privileging obligations: it must validate, monitor, and govern the AI tool's performance within its care environment, just as it would a newly credentialed clinician. The AI product owner

carries product liability for the tool performing as represented. This dual accountability mirrors existing institutional and professional liability structures and provides a familiar governance model for health systems.

Scope of Practice Congruency

Regardless of whether an AI tool operates in an augmented or autonomous capacity, its use and validation should be bounded by scope of practice and domain expertise. It would not be appropriate for a physical therapist to rely on a clinical decision support tool designed for neurologists, nor for a clinic to deploy a near-autonomous surgical AI system outside the oversight of appropriately credentialed surgical professionals. Validation of autonomous AI tools should follow peer-level review logic, which is analogous to the peer review processes used in quality management systems, where the evaluators possess domain expertise congruent with the tool's clinical function. HHS guidance on clinical AI governance should reinforce this principle to prevent scope mismatches that could compromise patient safety.

Transparency and Automation Bias Mitigation

Even in human-in-the-loop architectures, automation bias presents a material patient safety risk. Clinicians may defer to AI output without exercising independent judgment, particularly under time pressure or when the tool's reasoning is opaque. Addressing this risk requires action on two fronts:

1. AI developers should be expected to design tools that are transparent in their reasoning, surface confidence levels, and clearly communicate limitations, enabling meaningful clinician oversight rather than perfunctory confirmation.
2. Clinical training programs and continuing education should evolve to build competency in evaluating AI-generated recommendations critically, treating AI literacy as a core clinical skill rather than an elective technical interest.

Scoping reviews of EHR-integrated digital technologies have identified challenges including alert fatigue, workflow integration difficulties, and algorithmic bias, with a noted research gap on whether these tools address or worsen health inequities (Murthi et al., *Journal of Medical Systems*, 2024, DOI: 10.1007/s10916-024-02097-5). These findings underscore the need for post-deployment surveillance that accounts for differential performance across patient populations and for transparency requirements that support genuine clinical oversight.

Question 6: Where AI Has Met or Fallen Short of Expectations

Ambient AI documentation tools represent one of the most promising and well-evidenced applications of AI in clinical care today. A prospective study at Stanford Health Care found that ambient AI scribes produced statistically significant reductions in physician task load and burnout, with improvements in usability and perceived documentation quality (Shah et al., *JAMIA*, 2025, DOI: 10.1093/jamia/ocae295). Similarly, a quality improvement study at the

University of Kansas found that an ambient AI platform significantly improved clinician perceptions of documentation ease and reduced after-hours EHR work, with 48% of respondents reporting capacity to see an additional patient (Albrecht et al., *JAMIA Open*, 2025, DOI: 10.1093/jamiaopen/ooaf013).

A 2025 systematic review confirmed these patterns across multiple implementations, finding that 9 of 10 studies reported efficiency improvements and 7 of 10 reported positive effects on clinician wellness (Hassan et al., *Applied Clinical Informatics*, 2025, DOI: 10.1055/a-2597-2017). However, the same review noted that AI scribe accuracy varies and frequently requires manual edits, reinforcing the need for continued clinician oversight even in well-performing applications.

Reducing clinician burnout is not merely a workforce satisfaction issue, it is a patient safety imperative. Burned-out clinicians are more prone to medical errors and less able to engage meaningfully with patients. AI tools that restore clinician time and attention to direct patient care represent a high-value investment. HHS should consider how reimbursement structures and programmatic incentives can accelerate adoption of these administrative AI tools, particularly in resource-constrained settings where documentation burden is most acute.

Question 9: Patient and Caregiver Concerns

Patients and caregivers increasingly encounter AI-generated health information, both within clinical settings and through direct-to-consumer channels. A critical concern is that patients often cannot distinguish between reliable, clinically validated AI output and AI-generated content that may contain inaccuracies or hallucinations. When patients act on erroneous AI-generated medical information, whether by deferring necessary treatment or pursuing inappropriate interventions, the consequences can be serious.

HHS should consider how any AI product that provides medical information directly to patients can be held to standards commensurate with the clinical impact of that information. This does not necessarily require a single regulatory framework, but it does require clarity on accountability. Patients deserve to know whether the AI tools informing their care decisions have been validated against clinical evidence and whether a licensed professional stands behind the output.

At the same time, AI has significant potential to improve the patient experience by enabling more personalized treatment recommendations, reducing wait times through more efficient clinical operations, and supporting shared decision-making between patients and their care teams. Ensuring that clinical AI tools are built on high-quality, representative data is the foundation for delivering on this promise equitably across patient populations.

Conclusion

The pace of AI adoption in clinical care should be sustained, but it must be accompanied by deliberate action to ensure data quality, algorithmic accountability, clinician trust, and patient safety. As the evidence base for clinical AI matures, HHS should anticipate a future in which the adoption of validated AI tools becomes an expected component of the standard of care and should ensure that the regulatory and reimbursement environment supports that transition responsibly rather than reactively. Accountability structures should align with decision-making authority, with clinicians responsible for the professional judgment they exercise using AI tools, and AI product owners and deploying organizations sharing responsibility for autonomous systems. These frameworks should be grounded in existing clinical governance principles, including scope of practice, peer-level validation, and continuous quality monitoring, which are adapted for a technology-augmented care environment.

I would respectfully urge the Department to recognize that rigorous data quality standards and transparent validation processes are not impediments to rapid adoption: *they are prerequisites for the patient trust that sustains it*. Insufficient attention to governance and compliance when AI innovations enter clinical workflows acts as a silent killer of adoption: trust does not collapse in a single event but erodes incrementally, degrading patient confidence in AI-assisted care until adoption stalls and the root cause is difficult to trace. Reducing oversight prematurely risks high-profile failures that would set back public confidence in clinical AI and ultimately slow the very adoption trajectory HHS seeks to accelerate. If patients do not trust AI-assisted care, they will avoid providers who use it, withhold information from AI-enabled systems, or seek redundant second opinions, all of which undermine the cost savings and outcome improvements that clinical AI promises. Building trust through demonstrated rigor is the fastest path to durable, broad-based adoption.

HHS is uniquely positioned to use its regulatory, reimbursement, and research authorities to create an environment where evidence-based AI tools can flourish—improving outcomes, reducing burden, and making care more equitable and accessible. I appreciate the Department's commitment to seeking broad public input on these critical questions.

Respectfully submitted,

Caroline J. Morrow, PT, DPT

Board Certified Orthopedic Clinical Specialist, Board Certified Healthcare Executive, Certified Professional in Healthcare Quality

Managing Director, CJM Strategic Consulting, LLC