

COMMENT ON HHS HEALTH SECTOR AI RFI

Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care

Submitted by: David Condrey

Affiliation: Founder & CEO, WritersLogic, Inc.

Date: February 19, 2026

Contact: david@writerslogic.com

Introduction

HHS has correctly identified that public trust is a precondition for AI adoption in clinical care. This comment addresses a specific and underexamined barrier to that trust: the inability to verify who produced a clinical document after the fact.

When AI-assisted documentation becomes indistinguishable from documentation produced by a licensed clinician, accountability collapses. Malpractice attribution becomes probabilistic. Informed consent records become unverifiable. The physician signature at the bottom of a note attests presence but not authorship. Under current conditions, no technical mechanism exists to answer the question a court, a regulator, or a patient's family will eventually ask: did a human clinician actually compose this, or did they ratify output they did not write?

This is not a future problem. It is a present one that will intensify as AI-assisted documentation tools become standard across clinical settings.

1. The Provenance Gap in Clinical AI

HHS's RFI frames accountability and transparency as central values for AI adoption in clinical care. Current AI governance frameworks address accountability at the model level, through audits, validation studies, and performance benchmarks. These are necessary. They are not Sufficient.

What they do not address is accountability at the document level: the chain of evidence linking a specific clinical record to the specific human who produced it, in real time, during the act of Composition.

Existing attribution methods work backward from a finished document. Statistical classifiers attempt to distinguish AI-generated text from human-generated text based on patterns in the output. These methods produce approximations. They are inadmissible as evidence of authorship. They cannot establish, with cryptographic certainty, that a licensed clinician composed a document rather than reviewed or edited AI-generated text.

The practical consequences for HHS's stated goals are direct. Without verifiable authorship at the document level, the following cannot be reliably established:

- Whether a clinical note was authored by a licensed professional or generated by an AI system and co-signed without material human contribution
- Whether documentation submitted for Medicare or Medicaid reimbursement reflects genuine clinical judgment or automated output
- Whether a patient's informed consent record was produced through authentic human engagement or AI-assisted templating

Each of these is a liability exposure, a fraud vector, and a trust erosion point that AI adoption at scale will amplify.

2. Existing Technical Infrastructure

I am the developer of *witnessd*, an open-source cryptographic protocol for proving human authorship through keystroke dynamics and process evidence. The protocol operates at the moment of composition, not after the fact. It captures cryptographically signed evidence of human keystrokes, timing patterns, and compositional sequence in real time, producing a verifiable record that a human was present and active during the creation of a document.

The distinction that matters for HHS's framework is between forensic inference and prospective proof. Every current attribution method is forensic: it examines a finished document and infers probable authorship. *witnessd* is prospective: it captures evidence during composition that can be verified independently of the document's content.

Applied in clinical documentation contexts, this means a clinician using a *witnessd*-enabled documentation environment produces not just a signed note but a cryptographic record of the human compositional process behind it. That record is auditable, tamper-evident, and verifiable by regulators, payers, or legal parties without requiring access to AI detection tools whose reliability remains contested.

The protocol is open-source and has Internet-Drafts currently under review by the IETF SECDISPATCH working group:

- draft-condrey-rats-pop-protocol-06 — Proof of Process (PoP): Architecture and Evidence Format
- draft-condrey-rats-pop-appraisal-04 — Proof of Process (PoP): Forensic Appraisal and Security Model

Both drafts are profiled against the IETF Remote Attestation Procedures (RATS) architecture (RFC 9334). The companion reference implementation, `witnessd`, is written in Rust with SLSA Level 3 reproducible builds (<https://github.com/writerslogic/witnessd>). My published research, *On the Insecurity of Keystroke-Based AI Authorship Detection: Timing-Forgery Attacks Against Motor-Signal Verification* (arXiv:2601.17280), examines the adversarial threat landscape for authorship verification systems. I also submitted technical comments on NIST-2025-0035 this month, positioning unverified AI output as an agent security threat within the federal AI security framework.

3. Recommendations to HHS

HHS has three levers available: regulation, reimbursement, and research and development. This comment addresses each.

Regulation.

HHS should consider requiring that AI-assisted clinical documentation tools deployed in regulated care settings provide verifiable evidence of human authorship at the document level, not merely a clinician signature or attestation. The distinction between a signature and a cryptographic authorship record is the difference between assertion and evidence. Regulatory guidance that draws this distinction would incentivize the development and adoption of prospective authorship verification infrastructure.

Reimbursement.

CMS reimbursement policy already conditions payment on documentation requirements. HHS should evaluate whether existing documentation standards are adequate to distinguish AI-generated output from clinician-authored records, and whether reimbursement criteria should require verifiable evidence of human authorship for high-stakes clinical documentation categories. The absence of such standards creates a fraud vector that will grow as AI documentation tools proliferate.

Research and Development.

HHS and NIH should fund development of open standards for clinical authorship verification, with particular attention to interoperability with existing EHR systems. The technical infrastructure described here is nascent but deployable. Federal investment in this area would produce durable public infrastructure rather than proprietary solutions that fragment the market and create interoperability failures.

Conclusion

AI adoption in clinical care will accelerate regardless of regulatory posture. The question HHS can answer now is whether that acceleration happens with or without verifiable human accountability at the document level. The technical infrastructure to answer this question prospectively, at the moment of composition, rather than forensically after the fact, exists and is available for federal engagement.

I welcome the opportunity to provide additional technical detail, demonstrate the reference implementation, or discuss these considerations further.

David Condrey
Founder & CEO, WritersLogic, Inc.
david@writerslogic.com
<https://writerslogic.com>

References

- [PoP-Protocol]** D. Condrey, “Proof of Process (PoP): Architecture and Evidence Format,” Internet-Draft draft-condrey-rats-pop-protocol-06, IETF, 2026.
<https://datatracker.ietf.org/doc/draft-condrey-rats-pop-protocol/>
- [PoP-Appraisal]** D. Condrey, “Proof of Process (PoP): Forensic Appraisal and Security Model,” Internet-Draft draft-condrey-rats-pop-appraisal-04, IETF, 2026.
<https://datatracker.ietf.org/doc/draft-condrey-rats-pop-appraisal/>
- [RFC 9334]** H. Birkholz, D. Thaler, M. Richardson, N. Smith, W. Pan, “Remote ATtestation procedureS (RATS) Architecture,” RFC 9334, IETF, January 2023.
- [witnessd]** Writerslogic, Inc., “witnessd: Cryptographic Authorship Witnessing System,” Open-source implementation(Rust, SLSA Lvl 3). <https://github.com/writerslogic/witnessd>
- [Condrey-2025]** D. Condrey, “On the Insecurity of Keystroke-Based AI Authorship Detection: Timing-Forgery Attacks Against Motor-Signal Verification,” arXiv:2601.17280, 2025.
- [NIST AI 100-2e2025]** NIST, “Adversarial Machine Learning: A Taxonomy and Terminology of Attacks and Mitigations.”
- [C2PA]** Coalition for Content Provenance and Authenticity, C2PA Technical Specification, v2.2.
<https://c2pa.org/specifications/>