Bruce M. Bartels Management and Governance Center
45 Monument Rd., Suite 200
York, PA 17403

February 22, 2026

Dr. Thomas Keane
Assistant Secretary for Technology Policy/Office of the National Coordinator for Health Information Technology
Attention: Request for Information: HHS Health Sector AI RFI
Mary E. Switzer Building, Mail Stop: 7033A
330 C Street SW
Washington, DC 20201

Submitted electronically to Regulations.gov

Re: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care (RIN 0955-AA13)

Dear Dr. Keane:

WellSpan Health appreciates the opportunity to submit comments to the Assistant Secretary for Technology Policy (ASTP) and Office of the National Coordinator for Health Information Technology (ONC) on the Request for Information (RFI) on artificial intelligence (AI) adoption in clinical care. We support the Administration's ongoing commitment to tech-enabling healthcare to reduce administrative burden while improving access to high-quality care. We continue to advocate for a responsible, innovation-friendly and transparent regulatory framework for AI in healthcare, aligning with our longstanding commitment to improving patient outcomes and supporting clinicians through innovation.

WellSpan Health's vision is to reimagine healthcare through the delivery of comprehensive health and wellness solutions throughout our continuum of care. As an integrated delivery system focused on leading in value-based care, we encompass more than 2,700 employed providers, more than 250 locations, nine award-winning hospitals, home care, and a behavioral health organization serving central Pennsylvania and northern Maryland.

INTRODUCTION

The U.S. healthcare system stands at a critical juncture. After decades of fragmentation, rising costs, and uneven outcomes, we now possess the technological capabilities to fundamentally transform care delivery through market-driven solutions powered by applied AI. To realize this potential, the federal government must prioritize strong incentives — not punitive regulatory or enforcement postures — that encourage innovation rather than create fear-driven delays in development and adoption. A cohesive, multi-perspective, and nationally consistent framework is essential to ensure that all tiers of AI — from traditional machine learning to advanced generative models — can be safely and effectively implemented across the healthcare ecosystem.

To that end, a national framework of legal and regulatory incentives will support robust contracting practices and empower end-users to deploy rapidly advancing AI technologies with confidence. Together, we can build a forward-leaning, industry-supportive, and secure approach that accelerates clinical AI adoption while safeguarding patient privacy, civil rights, and civil liberties. By aligning regulation, reimbursement, and research and development investments, we can reduce the uncertainty that currently impedes innovation, strengthen public trust in emerging technologies, and ensure that AI is deployed in ways that enhance productivity, reduce operational burden, lower healthcare costs, and improve health outcomes for patients, caregivers, and communities across America.

Achieving these aims will require not only federal leadership but also a unified national vision that promotes innovation, protects patients, and empowers providers to deliver the highest-quality care.

WELLSPAN HEALTH'S SUCCESSES IN CLINICAL DEPLOYMENT OF AI

WellSpan Health has partnered with leading innovators to leverage AI in clinical workflows and improve patient outcomes.

At WellSpan Health, AI is a clinical support tool. It helps our teams deliver faster diagnoses, expand access to care and spend more time with patients while keeping safety, privacy and human judgment at the center of everything we do. In a clinical setting, early detection and intervention can improve patient outcomes and drastically reduce overall healthcare costs – for both patients and providers, and that is where we have focused AI deployment. For example, WellSpan Health uses Food and Drug Administration (FDA) cleared clinical AI to help radiologists identify and prioritize urgent findings such as stroke, pulmonary embolism, brain hemorrhage, aneurysm and spinal fractures – among other diagnoses. These AI tools analyze imaging studies and flag high-risk cases within minutes, allowing radiologists to review the most critical scans first. A radiologist expert still reads every image. The AI supports triage and prioritization; it does not replace clinical judgment by a trained clinician.

Additionally, WellSpan Health became the only health system in central Pennsylvania to deploy an AI solution to aid in breast cancer detection during screening mammograms. The AI helps identify cancers that are often difficult to detect, particularly in women with dense breast tissue or implants. Earlier detection allows WellSpan Health teams to more quickly connect patients with follow-up care, sometimes before a patient even leaves the appointment.

Below are our detailed responses to the Department of Health and Human Services' (HHS) RFI.

1. ***What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?***

The most significant barriers to AI adoption in healthcare include financial constraints, liability concerns, malpractice impact, change-management challenges, infrastructure complexity, misaligned reimbursement incentives, and patient trust. Up-front and ongoing costs for implementing, monitoring, and maintaining AI systems are substantial — especially when responsible use requires human oversight, safety reporting, and continuous performance auditing. Fear of legal exposure, such as liability for missed findings or automation errors, combined with unclear risk allocation between vendors and providers, further suppresses deployment. Finally, data accessibility and interoperability among legacy platform providers is also a major barrier.

Operationally, AI introduces new workflows, training needs, and integration dependencies that can slow adoption and create additional failure points. Patients and clinicians may also harbor skepticism — worrying that AI could diminish the clinician's role or produce inaccurate outputs —which underscores the need for thoughtful design, education, and transparent quality reporting.

To successfully enable market activity and promote healthy competition among AI developers and healthcare end-users, a national operating standard is essential. Such a standard would support the natural development of contracting norms that help providers consistently evaluate the trustworthiness, reliability, and measurable benefit of AI tools. Clear expectations for risk allocation, transparency, avoidance of "black box" data gaps, reporting requirements, and insurability would strengthen industry accountability and better protect patients who may be incidentally harmed as these technologies continue to evolve.

Finally, today's fee-for-service reimbursement structure creates a structural disincentive for adopting AI tools that prevent costly interventions or improve efficiency. Because AI's productivity gains are not rewarded under inflationary fee-for-service models, investment gravitates away from more expensive, high-value applications. In fact, any productivity gains that result from the use of AI would be punitive under the fee-for-service models, creating a self-defeating cycle to the adoption of AI. Aligning reimbursement with outcomes is essential to creating market pull for tools that genuinely improve care quality and operational efficiency. Transitioning to value-based models would allow the healthcare system to absorb AI's deflationary benefits and reinvest the savings generated by automation back into the broader ecosystem.

2. ***What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.***

HHS should prioritize payment policies that recognize AI-enabled care, paired with predictable and proportionate oversight. First, reimbursement pathways should permit separate payment for AI-based services when the tool demonstrably contributes to clinical value — for example, post-discharge telephonic follow-ups conducted effectively by an AI agent — so adoption is not penalized relative to fully human workflows. Second, HHS should pursue an outcomes-oriented pathway for AI-based healthcare services such as radiology triage or early detection models, acknowledging both the throughput gains and the earlier, lower-cost interventions that improve long-term patient outcomes. Third, maximum flexibility should be granted to providers using validated tools to avoid prescriptive constraints that would suppress capacity gains, particularly in regions facing acute workforce shortages. Finally, provisional clearances paired with a structured, stepwise pathway to full accreditation would support timely innovation while allowing evidence to accumulate without stalling deployment.

Expanding the amount of AI-supported care that is separately paid — whether through dedicated AI-related billing codes or incentive programs tied to demonstrated improvements in patient outcomes within value-based care — would accelerate responsible adoption across the healthcare system. At the same time, a federal ban on payer-mandated "AI-prohibition" clauses for revenue-cycle automation is essential. These clauses create significant inequity: while payers increasingly rely on AI for claims processing and denial management, they often restrict providers from using similar tools to improve their own administrative efficiency. Ensuring that providers can deploy AI on equal footing would promote fairness, reduce administrative burden, and support a more balanced and transparent claims ecosystem.

Regulatory activity should prioritize a set of nationally consistent policies that strengthen security, transparency, and trust while enabling fair, competitive AI adoption across the healthcare ecosystem.

First, cybersecurity policy should be incentive-based, rewarding providers for maintaining reasonable, diligent adherence to evolving best practices rather than penalizing them for sophisticated AI-driven attacks that overwhelm even well-prepared organizations.

Second, federal leadership is urgently needed to replace the emerging patchwork of fifty plus different state AI laws—many of which conflict with existing federal privacy and disclosure requirements — with a unified national framework for disclosure, consent, and privacy. A preemptive federal standard would reduce uncertainty for end-users and streamline compliance for vendors operating across multiple jurisdictions.

Third, standardized Business Associate Agreement (BAA) requirements are essential to closing the transparency gap between AI vendors and end-users. Today, vendors often use opaque contracting practices and maintain "black-box" data silos that shift disproportionate liability to providers, limiting their ability to evaluate products or foster healthy market competition.

Fourth, clear transparency and audit requirements should be developed to establish federal baselines for algorithmic explainability, banning black-box opacity, and requiring vendors to provide audit information to end-users. Such standards would promote fair competition, improve safety, and strengthen patient trust by making AI systems more understandable and accountable.

Finally, modernization of the Physician Self-Referral Law (Stark Law) and the federal Anti-Kickback Statute is needed to include safe harbors and exceptions that reduce legal uncertainty and enable responsible collaboration, particularly around data sharing and cybersecurity. Collectively, these regulatory actions would create industry-wide best practices for AI contracting, empower end-users, and support a competitive, innovative marketplace that ultimately benefits patients.

3. ***For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (e.g., relating to liability, indemnification, privacy, and security). What novel legal and implementation issues exist and what role, if any, should HHS play to help address them?***

For AI software that is not integrated with the developer's hardware, several unresolved governance challenges continue to impede responsible adoption, including liability allocation between providers and vendors, standards for indemnification, privacy and security obligations for data used in training and post-market monitoring, and accountability for ongoing model updates. HHS can play a defining role in addressing these gaps by clarifying safe-harbor conditions for providers who follow approved governance practices, encouraging shared-risk contracting models that align vendor accountability with claimed performance, and promoting standardized incident-reporting pathways that integrate AI-related errors into existing patient-safety systems without creating punitive disincentives to disclose. Our own operations already route AI-related errors through patient-safety workflows — for example, when an automated agent confirmed an appointment that was never scheduled — illustrating the need for normalized, "learn-from-events" processes rather than ad hoc responses.

Additionally, patient consent frameworks require modernization. HHS should clarify Health Insurance Portability and Accountability Act (HIPAA) obligations under 45 CFR §164.508 for AI-related disclosure and consent, including clear safe harbors. Updated guidance should address when and how patient consent is required for AI-supported analysis, the scope of opt-out rights, and expectations for continuously learning systems. Clear and consistent HIPAA requirements would support predictable patient communication based on actual clinical risk, reduce legal ambiguity for health systems, and strengthen patient trust in AI-enabled care.

4. ***For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and other workflow and human-centered evaluation methods for clinical care? Should HHS further support these processes? If so, which mechanisms would be most impactful (e.g., contracts, grants, cooperative agreements, and/or prize competitions)?***

Effective evaluation of clinical and operational AI requires coupling technical robustness measures — such as drift detection, hallucination and error monitoring, and bias assessments —with human-centered and workflow-based metrics, including clinician trust, handoff safety, and escalation efficacy. Accordingly, WellSpan Health combines supervisory AI systems that monitor for drift and trigger alerts with frontline error-intake processes embedded in our existing patient-safety infrastructure. This blended model of automated surveillance and human recognition ensures both high-fidelity detection and accountability. Prior to deployment, real-world simulations and pilot proofs-of-concept anchored to clinical key performance indicators (KPIs) are essential. Post-deployment, continuous monitoring and the use of registries to track performance, safety events, and outcome signals are critical to ensuring ongoing reliability and safety.

HHS can significantly amplify these efforts by leveraging contracts, grants, and cooperative agreements to fund multi-site evaluations, develop common registries, and build shared reference datasets. Prize competitions could further accelerate the development of transparent, rigorous robustness-testing frameworks that the entire industry can adopt.

5. ***How can HHS best support private sector activities (e.g., accreditation, certification, industry-driven testing, and credentialing) to promote innovative and effective AI use in clinical care?***

HHS should endorse and align with accreditation pathways that either recognize existing FDA labeling where applicable or establish new models that mirror FDA principles — ensuring baseline safety and effectiveness while supporting post-market surveillance. At the same time, accreditation must preserve room for innovation sandboxes and provisional scaling so developers can iterate and build evidence without halting progress. This balanced approach allows implementation to "start with either accreditation or an FDA label," while still enabling responsible exploration and evidence generation. HHS can further strengthen the ecosystem by supporting ACR/ABR/DNV/CMS-aligned frameworks, expanding registry infrastructure for outcomes tracking, and offering education and implementation supports for providers.

A key opportunity to support AI adoption in clinical care is a Deemed Status model in which HHS designates qualified accreditation organizations — similar to the Deemed Status accreditation system currently in effect at 42 CFR Part 488 — to certify AI systems that meet established safety, effectiveness, and equity standards. Healthcare organizations using HHS-recognized and certified AI tools would receive regulatory advantages, creating a more efficient and trustworthy marketplace. This approach would allow hospitals and clinics to adopt certified AI tools with confidence, reducing the need for extensive internal evaluations that many organizations — especially smaller community hospitals and rural health systems — lack the capacity to perform. The result is a democratization of access to high-quality AI, ensuring benefits are not confined to well-resourced academic medical centers.

Finally, HHS should support the development of clinical AI credentialing programs that train healthcare professionals to evaluate, implement, and oversee AI systems effectively. Building workforce expertise is essential to safe deployment, informed adoption decisions, and long-term governance of AI-enabled care.

6. ***Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?***

We have observed strong performance from assistive imaging technologies — such as triage tools and incidental but clinically significant finding detection — which deliver meaningful throughput efficiencies and, more importantly, earlier identification of clinical issues that enable lower-cost interventions and improved outcomes. Even with 100 percent human over-reads, we realized measurable FTE efficiency while improving detection rates, demonstrating that productivity and quality gains can occur simultaneously. In contrast, some conversational and agentic tools underperform when workflow integration is immature, resulting in errors or patient friction, such as failed appointment confirmations. Despite these challenges, high-impact opportunities remain in acute-care decision support, population-scale screening models (e.g., mobile low-dose CT paired with AI-based interpretation), and AI-augmented follow-up pathways that expand capacity without requiring proportional increases in staffing.

Generative AI agents represent another major area of potential. By automating routine administrative tasks — including communications, prior authorizations, documentation, and scheduling — these agents could reclaim 15–20 hours per week of clinician time, directly addressing burnout and allowing clinicians to focus their time on patient care. Medium-term, HHS can play a catalytic role in enabling widespread deployment of generative AI agents across healthcare. If implemented at scale, these tools could automate 1–2 billion hours annually of administrative work, alleviating pressures driven by the workforce crisis while contributing to reductions in administrative spending, currently estimated at $250–400 billion per year. By supporting the integration, evaluation, and responsible scaling of these technologies, HHS can accelerate measurable gains in workforce sustainability, patient experience, and system-wide productivity.

7. ***Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?***

Adoption hinges on clinical and implementation champions with demonstrated organizational points of contacts (POCs), supported by strong governance, finance, and IT partners who can execute, maintain, and build trust. Influence may be top down or bottom up, but success requires ground team engagement, prioritization across competing initiatives, and resources for ongoing monitoring. Administrative hurdles include insufficient local IT capacity, fragmented governance, unclear funding mechanisms for new responsibilities (monitoring, auditing, validation), and fear that AI could devalue clinician expertise. Aligning roles, incentives, and education — while budgeting for durable oversight — is critical.

8. ***Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.***

Enhanced interoperability — especially standardized, low-cost Fast Healthcare Interoperability Resources (FHIR) based exchanges — would reduce vendor lock in, improve secure data sharing, and enable AI models to operate on consistent data elements across settings. This consistency improves care quality (less test duplication, clearer cross system understanding of patient needs) and lowers integration cost for innovators. Benchmarks should emphasize common data types and metrics that support fair head-to-head evaluations and post market monitoring, ideally coupled with registries capturing system analytics and patient outcomes over time.

9. ***What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?***

Patients consistently raise several core concerns regarding the use of AI in their care, including privacy and data security — specifically who has access to their health information and whether AI systems might inadvertently expose sensitive data. They also worry about algorithmic bias and the possibility that AI tools could deliver inferior or inequitable care to certain populations. Transparency and explainability are equally important, as patients want to understand why an AI system recommended a particular treatment or course of action. Finally, patients need clarity around accountability: if something goes wrong, who is responsible?

To address these concerns, HHS should clarify liability frameworks so patients understand their rights when AI is involved in clinical decision-making. Clear, modernized medical malpractice structures are needed to appropriately assign responsibility across AI developers, healthcare organizations, and individual clinicians. Most fundamentally, HHS should ensure that AI adoption genuinely serves patient interests and improves patient experience and outcomes—rather than primarily benefiting healthcare organizations or technology vendors.

10. Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care?

a. Are there published findings about the impact of adopted AI tools and their use clinical care?

b. How does literature approach the costs, benefits, and transfers of using AI as part of clinical care?

Priority research should target multicenter, real-world trials that link AI use to patient level outcomes, not only process metrics — today's literature is niche and organization dependent, with limited ability to report outcomes absent robust analytics. HHS can catalyze shared registries and common analytics that capture clinical efficacy, workflow safety events, and economic endpoints across diverse sites. Cost/benefit studies should reflect both capacity gains (growth, earlier detection) and the ongoing costs of safe operation (human oversight, monitoring, vendor fees), while modeling transfers across stakeholders (providers, payers, patients). Embedding validated predictive models into clinical and drug trials can hasten evidence generation, and education research should tackle AI literacy for clinicians and the public to address perception gaps that impede adoption.

Programmatic payment pathways HHS should consider in the future include two complementary payment constructs that would materially accelerate high value adoption:

- AI “Meaningful Use” style framework: Define usage requirements, tie incentives to Medicare/Medicaid, and integrate with national standards for safety, transparency, and interoperability—starting where value-based models already track quality (e.g., mortality, readmissions, healthcare-associated infections (HAIs)) to show early, credible wins.
- AI Add on Payments in Medicare: Offer time limited supplemental payments (e.g., consistent payments over 3–5 years with step down percentages) for clinically valuable AI, with guardrails against overuse (per episode caps, documentation standards, periodic audits) and a structured transition to permanent codes or bundled models after evidence review. This approach could be set up in a manner that reduces financial barriers, rewards bias mitigation and interoperability, and curbs fraud risk—an issue the Centers for Medicare & Medicaid Services (CMS)/Center for Medicare and Medicaid Innovation (CMMI) leaders have prioritized and that can be preempted with reporting and validated partner mechanisms.

CONCLUSION

We have a historic opportunity to transform U.S. healthcare through AI-driven innovation. After decades of incremental change, we now have an unprecedented chance to reshape the healthcare system using market-driven solutions built on applied AI. In the near term, AI’s greatest value lies not simply in “hours saved,” but in expanding capacity and improving productivity —addressing workforce shortages while enhancing access, quality, and patient outcomes.

With sensible reimbursement structures, clear and proportionate oversight, robust post-deployment evaluation, and targeted research and development support, HHS can unlock this value at scale — especially in high-cost, acute-care, and imaging-intensive settings — without compromising safety, equity, or public trust. Achieving this vision requires national consistency, along with alignment between federal legislators and federal agencies around an incentivizing legal architecture rather than a punitive one. Given the complex patchwork of state-level requirements and the current power imbalances between AI vendors and clinical end-users, federal guidance should preempt conflicting state activity wherever feasible.

While other commenters may highlight different priorities, the recommendations above reflect the broad and practical concerns of providers and patients within the healthcare system. Implementing these approaches would deliver substantial benefit across the early and mid-stage development of the AI-in-healthcare ecosystem, supporting innovation while ensuring that adoption remains safe, equitable, and patient-centered.

WellSpan Health appreciates the opportunity to provide feedback on opportunities to advance the use of AI in clinical settings to improve patient care. We look forward to working with you to design sustainable policy solutions.

Sincerely,

Roxanna Gapstur
President & CEO
WellSpan Health