February 20, 2026

Via Electronic Submission
United States Department of Health and Human Services
200 Independence Avenue, S.W.,
Washington, D.C. 20201

Re: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care

SonderMind, Inc. ("SonderMind") is a comprehensive behavioral health provider dedicated to improving access to high-quality, continuous care through technology and clinical expertise. We connect individuals with licensed behavioral health providers and support those providers and clients with technology-enabled tools to improve clinical care. We appreciate the opportunity to comment on the examination of accelerating the adoption and use of artificial intelligence in clinical care.

SonderMind operates across all 50 states, serves tens of thousands of licensed therapists and prescribers, and is actively building and deploying AI tools within behavioral health clinical workflows, building all tools with our clinicians. Our AI-powered clinical documentation tool has been adopted by 80%+ of eligible providers and has generated over 100,000 notes submitted to health insurers, reducing note completion time by 80% and saving providers up to 90 minutes per day. Every AI feature on our platform is opt-in for both clients and providers, and every AI-generated output is fully editable by the clinician before finalization. SonderMind has published an AI Constitution establishing binding principles for safety, transparency, and clinical partnership in the use of AI.

<u>Q1: Biggest Barriers to Private Sector AI Innovation and Adoption in Clinical Care</u>

The barriers to AI innovation in healthcare, particularly in behavioral health, span reimbursement, provider readiness, liability, and regulatory fragmentation.

Misaligned Reimbursement Models
The legacy fee-for-service (FFS) model creates an economic penalty for innovation. AI tools designed to reduce visit frequency or provide continuous care, extending support beyond the traditional 50-minute session, effectively lower a provider's billable volume. Current CPT codes, including the psychotherapy codes (90832-90834) that govern most behavioral health billing, do not adequately capture or compensate for AI-enabled continuous care such as proactive outreach, between-session behavioral monitoring, or real-time crisis detection. While the 2026 CPT code set has introduced new codes for digital health, they remain focused on episodic data collection

rather than the ongoing, asynchronous support that defines AI's greatest potential in mental health. Value-based care (VBC) models are naturally aligned with AI's ability to improve outcomes and reduce costs, but VBC adoption in behavioral health remains significantly lower than in physical medicine, leaving innovators without a viable commercial pathway.

Provider Understanding and Comfort
Clinicians are trained to practice medicine, not to evaluate how AI is designed, validated, or deployed. Many providers still struggle to understand how AI tools are best used, assess their risks and benefits, or explain those factors clearly to patients. There are very few high-quality, practical education offerings to close this gap, which limits adoption even when the tools are effective.

Liability and Risk Assignment
It is unclear who bears responsibility when AI-assisted care leads to an adverse outcome: The clinician, the AI developer, or some combination. This ambiguity is especially problematic in behavioral health, where high-risk situations such as suicide risk assessment, duty-to-warn obligations, and safety planning carry significant legal complexity. Without clear frameworks, both providers and developers face uncertainty that discourages adoption.

Fragmented and Shifting State Regulation
The regulatory environment for clinical AI is not standardized across states and is changing rapidly. Different states impose different, and sometimes conflicting, requirements on how AI can be used, documented, and supervised in clinical settings. In practice, this means a provider using AI must simultaneously be a clinician, technologically literate about AI tools, and knowledgeable about the legal and regulatory framework in every state where they practice. This burden disproportionately affects behavioral health, where telehealth and multi-state practice are common.

SonderMind has direct experience with this challenge. Multiple states are advancing legislation that, while well-intentioned, is creating a patchwork of incompatible requirements that slows innovation and confuses providers. This year alone, more than a dozen states have introduced legislation regulating psychotherapy AI risks, many inadvertently prohibiting safety features like AI-enabled crisis detection and creating liability exposure for providers who use any AI tool. In some states, the regulatory framework is making it untenable for us to offer our AI-powered client support agent to members in that state (a tool that is deployed and available to clients across nearly every other state in the country). The result is that those members are denied access to a beneficial, clinically-guardrailed tool, not because of a safety concern, but because of regulatory ambiguity.

The absence of a coherent federal framework means that companies building responsible, clinically-validated AI tools face a state-by-state compliance burden that favors large incumbents and penalizes the smaller innovators most likely to develop breakthrough applications in underserved domains like behavioral health.

Q2: Regulatory, Payment Policy, and Programmatic Changes HHS Should Prioritize

Untreated and undertreated behavioral health conditions drive significant downstream healthcare costs, including emergency department utilization, worsened medical comorbidity, and workforce disability. AI tools that expand provider capacity and sustain patient engagement between sessions have the potential to reduce these costs, but current payment structures do not recognize or reimburse the mechanisms through which that value is delivered. The following reforms would begin to close that gap.

Asynchronous Reimbursement Pathways

CMS should establish specific payment pathways under 42 CFR Parts 414 and 425 for AI-enabled behavioral health interventions that reward continuous care models. This includes reimbursement for asynchronous engagement and automated monitoring that leads to documented clinical improvements. Current psychotherapy reimbursement under CPT codes 90832-90834 does not recognize or compensate for AI-enhanced treatment planning or between-session support, making adoption economically unviable even when clinical value is demonstrated.

Outcome-Based Incentives

HHS should expand value-based payment models under 42 CFR Part 425 that shift from rewarding volume to rewarding validated outcomes. AI tools that demonstrate a reduction in acute episodes (e.g., emergency department visits, crisis interventions) should be eligible for incentive payments that offset the reduction in traditional session volume. CMS should consider creating CMMI demonstration models specifically for behavioral health AI to generate the evidence base for broader coverage decisions.

Risk-Proportionate Oversight of AI Tools

Regulation should differentiate between validated clinical AI tools and general-purpose large language models (LLMs) that are not designed or tested for healthcare use. HHS should clarify the regulatory boundary under 21 CFR Part 860 between clinical decision support tools that qualify for exemption and those requiring premarket review. Unvalidated tools used in clinical contexts pose meaningful risks, including the potential to encourage harmful behaviors. HHS should implement oversight that is proportionate to the risk profile of the tool, ensuring that well-designed clinical AI is not subjected to the same regulatory burden as unvalidated consumer products, and vice versa.

Federal Preemption or Harmonization of State AI Regulation
As noted in Q1, the absence of consistent federal guidance has produced a patchwork of state-level AI regulations that impose conflicting requirements on companies operating across multiple states. HHS should work with Congress to establish minimum federal standards for the use of AI in clinical care that preempt or harmonize inconsistent state requirements. Without this, the companies most committed to responsible, multi-state deployment of clinical AI bear the highest compliance burden, while companies with less rigorous safety and compliance practices face fewer constraints.

Q3: Governance, Accountability, and Liability for Non-Medical Device AI

The most pressing governance challenge for non-device AI in clinical care is the misalignment between liability and control. Current frameworks risk holding clinicians fully responsible for technical failures they cannot prevent, detect, or fix, such as developer bugs, inaccurate AI-generated outputs, or vendor-side data breaches. This creates a chilling effect: providers avoid beneficial AI tools entirely rather than accept uninsurable risk.

This is not a hypothetical concern. SonderMind is actively engaged with state legislatures on proposed bills that would hold providers fully responsible for all aspects of an AI system's performance, including technical failures entirely outside the provider's control. As noted above, regulatory ambiguity around AI-powered client interactions has forced us to exclude some states entirely from our AI client support agents. These state-level dynamics illustrate why federal governance frameworks are urgently needed.

Distributed Liability Frameworks
HHS should convene a multi-stakeholder working group to develop model liability allocation frameworks that clarify responsibility based on who controls what. The provider should bear responsibility for the clinical decision to use an AI tool and the standard of care provided to the client. The AI developer should bear responsibility for the safety, accuracy, and security of the tool itself. Standard malpractice insurance covers clinical negligence, not third-party software failures, regulations that ignore this distinction create risk that is both unjust and uninsurable.

Transparency Standards
HHS should work with National Institute of Standards and Technology (NIST) and Office of the National Coordinator for Health IT (ONC) to establish tiered explainability requirements appropriate to clinical context and risk level. For behavioral health, where AI may be used in crisis detection or treatment matching, providers and patients need to understand why an AI made a particular recommendation. SonderMind's approach, where every AI-generated output is

fully editable by the clinician before finalization, represents one model for maintaining clinical oversight while leveraging AI capabilities.

Patient Consent and Disclosure
HHS should clarify HIPAA requirements under 45 CFR 164.508 regarding patient consent for AI use in clinical care. Patients should know when AI is involved in their care and have the ability to opt out. SonderMind has implemented opt-in consent for all AI features for both providers and clients, but industry-wide standards would build broader public trust.

Q5: Supporting Private Sector Activities (Accreditation, Certification, Industry-Driven Testing)

Differentiated Certification Standards
HHS should support the development of accreditation standards that distinguish between general-purpose AI and AI built for behavioral health with strong clinical standards. Tools built for healthcare with appropriate clinical guardrails should be eligible for expedited pathways to reimbursement.

Clear Accountability Mechanisms
To build public trust, HHS should promote frameworks where AI developers and deploying providers share accountability for outcomes. If an AI tool delivers negative clinical results due to inadequate design or insufficient safeguards, there must be transparent and proportionate consequences. Tying accountability to reimbursement eligibility, where tools that fail to meet safety and efficacy standards lose access to payment pathways, would create meaningful market incentives for quality.

Q6: Where AI Tools in Clinical Care Have Met or Exceeded Expectations

In SonderMind's experience deploying AI in behavioral health, the tools that have exceeded expectations are those that support providers within their existing workflows rather than attempting to replace clinical judgment.

Clinical Documentation (AI Notes)
SonderMind's AI-powered documentation tool reduces note completion time by 80%, saving providers up to 90 minutes per day. Since commercial launch in September 2025, the tool has been adopted by 80%+ of eligible providers and has generated over 100,000 notes submitted to health insurers. The tool transcribes and formats session notes with patient consent. Every AI-generated note is fully editable by the provider before finalization, maintaining clinical oversight while dramatically reducing administrative burden.

Quality Assurance
Automated chart auditing supports compliance monitoring by identifying documentation gaps and flagging potential issues that would otherwise require manual review.

Between-Session Client Support
SonderMind offers AI-enabled tools that help clients stay engaged with their therapeutic work between sessions, including session preparation features and guided journaling tools that support self-awareness and goal-setting between appointments. These features extend the reach of the therapist beyond the 50-minute session without replacing human clinical judgment. Over 15,000 unique clients have prepared for sessions using our AI-guided session preparation tool, helping them plan what they want to discuss in therapy sessions. Additionally, over 30,000 members have utilized our AI-guided Reflections as an interaction AI journaling tool, documenting and sharing insights in between sessions.

Key Insight
The pattern across successful implementations is consistent: AI tools perform best when they enhance existing provider behaviors and client/provider workflows rather than introduce fundamentally new workflows. Adoption is highest when the tool reduces friction in tasks providers already do (documentation, compliance) rather than asking them to fundamentally change how they deliver care. These efficiency gains carry broader implications: 90 fewer minutes per day on documentation represents capacity that can be returned to direct patient care or used to reduce the burnout and attrition that constrain the behavioral health workforce.

<u>Q10: Priority Areas for AI Research</u>

Safety and Harm Prevention
Research must prioritize ensuring AI tools cannot produce harmful outputs in mental health contexts. HHS should fund studies specifically aimed at testing LLMs in behavioral health scenarios to ensure they cannot be manipulated into encouraging self-harm or other dangerous behaviors. This is not a theoretical concern: general-purpose consumer AI chatbots have been involved in widely reported incidents where users in mental health crisis received responses that could exacerbate harm. Clinical-grade behavioral health AI requires safety testing at a standard that does not yet exist in the research literature.

SonderMind invests in this work directly: our clinical and machine learning teams collaborate daily on safety testing, clinical guardrail design, and ongoing monitoring of AI outputs in behavioral health contexts. We believe companies deploying AI in mental health care have an obligation to lead on safety, not wait for federal mandates. HHS can accelerate this work by funding shared research infrastructure and establishing behavioral health-specific benchmarks that the industry can adopt.

AI-Clinician Collaborative Models
HHS should prioritize research on how AI can best augment the clinician rather than replace them. This includes studies on AI's role in triaging acute needs, providing longitudinal support between clinical visits, and keeping patients engaged in their care.

Bias and Equity
Research into ensuring AI models do not perpetuate existing disparities in mental health diagnosis and treatment across diverse populations is critical, particularly given the historical underrepresentation of certain communities in behavioral health research data.

Conclusion

Behavioral health serves one in five Americans, yet it remains significantly underrepresented in health technology policy. The questions posed in this RFI are the right ones. The answers, however, will be incomplete if they are shaped primarily by the experiences of radiology, pathology, and acute care, domains where AI operates on images and structured data rather than on the complexity of human conversation and therapeutic relationship.

SonderMind offers a perspective grounded in direct operational experience: building AI tools for behavioral health, deploying them at scale, navigating a fragmented regulatory landscape, and learning what works when the goal is not just efficiency but better clinical outcomes for people seeking mental health care. The barriers we have described (misaligned reimbursement, unclear liability, and inconsistent state regulation) are solvable. The policy tools exist. HHS should include behavioral health in every AI framework it develops. The scale of the need demands it. We are ready to be a resource to HHS as this work moves forward.