



February 20, 2026

Steven Posnack
Principal Deputy Assistant Secretary for Technology Policy
Office of the National Coordinator for Health Information Technology
U.S. Department of Health and Human Services
330 C St SW, Floor 7
Washington, DC 20201

Submitted Electronically

Re: HHS Health Sector AI RFI

Dear Mr. Posnack:

Thank you for the opportunity to respond to your RFI regarding accelerating the adoption and use of artificial intelligence (AI) as part of clinical care. Epic's software has included various forms of AI for decades; currently we have over 175 use cases of generative AI released or in active development.

Through our work with customers, we have developed a deep understanding of how to speed the adoption of AI in clinical settings. For example, Epic offers the Jump Start Cloud Infrastructure program, which has resulted in customers increasing their AI use 2.4 times faster, by expanding use of existing features and adopting new ones. Drawing on this experience, we are pleased to offer several ideas for how the Department can encourage AI uptake.

We would be happy to answer any questions you may have and would welcome further collaboration on your work to promote AI adoption in clinical care.

Sincerely,

Ladd Wiley
Senior Vice President for Global Corporate
Affairs, Public Policy, and Advocacy

David Hoffert
Director of AI Policy

Getting to “On”

Generative AI differs from prior healthcare innovations in ways that can make clinical leadership reluctant to adopt and scale it. Because generative AI is probabilistic, its performance must be validated not only before deployment but also continuously over time, requiring skills and resources that may be scarce in healthcare settings. Its costs—driven by compute, storage, and model usage—are also highly variable and difficult for organizations to predict, budget, or govern using traditional software pricing models.

At the same time, AI delivers meaningful benefits: improved clinician productivity, reduced documentation burden, lower burnout, more timely clinical decision-making, fewer missed orders, improved throughput, and better patient engagement. Over time, these gains can translate into reduced staffing pressure and overtime costs, higher-quality documentation, improved access and satisfaction, and better clinical outcomes through earlier interventions.

Realizing these benefits requires healthcare leaders to move past understandable initial hesitation and turn AI features on. While familiarity and trust with new technology develop over time, we believe there are concrete steps the Department can take now to better align incentives and accelerate adoption.

Documenting ROI

One question any responsible healthcare administrator will ask before deploying new AI technology is whether it represents a sound investment. Healthcare organizations operate today on thinner margins than ever; for many, it simply is not an option to invest in technologies that will not pay off financially. We believe many applications of generative AI in healthcare do present sufficient return on investment (ROI).

The Department could help document ROI data in a way that healthcare providers can both rely and act on. For example, studies show that AI scribe technology—in which AI records and transcribes clinical interactions and then drafts corresponding notes—can reduce the time clinicians spend working after hours. Less “pajama time” for providers, however, does not by itself offset the cost of the technology. The true ROI of AI scribes appears downstream: happier clinicians are less likely to burn out or leave, reducing recruitment and onboarding costs.

The Department could address this gap by assembling a repository of AI use cases with demonstrated ROI. By conducting independent research, such a repository would validate vendor claims and highlight metrics relevant to healthcare administrators’ bottom-line decisions. Such a repository would be vendor-agnostic, focusing on use cases rather than specific implementations, to not advantage those vendors that get their claims published. Done right, it would serve as a trusted reference for organizations evaluating AI investments.

Bridging the Adoption-to-ROI Gap

Even when an AI use case has a proven ROI, initial adoption may still be financially challenging. The benefits of AI are rarely immediate and often require workflow changes, training, and iteration to be fully realized. Early phases of adoption are typically marked by uneven usage, ongoing configuration, and infrastructure costs that scale before productivity gains materialize. As a result, organizations may incur real costs before seeing offsetting savings, creating a financial mismatch that can slow or stall adoption.

One response to these challenges would be fee-for-service reimbursement for AI use, but this approach could quickly become costly and, for AI with proven ROI, unnecessary. Organizations primarily need support through the initial scaling period; once broad adoption is achieved, AI that delivers value will pay for itself.

Epic has had success spurring AI adoption among customers with our Jump Start Cloud Infrastructure Package, a program we created to simplify and support early adoption of generative AI. It covers the cloud infrastructure costs for roughly 50 AI features with a flat monthly fee based on organization size. This approach encourages responsible, broad adoption rather than narrow pilots driven by short-term cost avoidance.

The results have been impressive. Jump Start customers have demonstrated median month-over-month AI resource utilization growth of 45%, compared to 18% for non-Jump Start customers. They also adopt more new AI features—5.5 on average, compared to 4.6 over the same period. Overall, they use 17 times more AI resources than non-Jump Start customers of comparable size. These differences are not explained by organization size or early-adopter effects; in fact, Jump Start customers are slightly smaller on average. Rather, Jump Start succeeds because it aligns incentives to encourage rapid scaling. Once customers reach broad adoption, they see the value of AI—and continue using it.

The Department could achieve similar incentive alignment via AI adoption grants that cover infrastructure costs for a fixed period of time rather than a fixed amount of money. Traditional grants incentivize slow scaling to extend funding; time-based grants would encourage organizations to scale quickly, assess value, and determine which use cases merit continued investment. After the grant period ends, organizations could discontinue low-value AI while sustaining AI that delivers sufficient benefit—without ongoing government spending.

Such programs could be structured as true grants or as fixed-price arrangements similar to Jump Start. In the latter case, however, pricing should scale with organization size to avoid disadvantaging smaller or lower-resourced providers. (If desired, pricing could even be structured disproportionately to give an advantage to smaller organizations.) Properly

designed, time-based infrastructure support would help organizations escape the “pilot trap” and reach sustainable AI adoption.

Making Local Validation Possible

Another concern AI adopters rightly have is safety. Many healthcare AI use cases fall outside the oversight of the Food and Drug Administration, and even regulated products are poorly served by traditional evaluation methods. Especially for AI that is probabilistic, pre-deployment testing on generalized datasets has limited predictive value for real-world performance.

The better method for validating AI performance in healthcare is “local validation,” in which the healthcare facility, independent of the developer, evaluates how the AI performs on its own data while embedded in its own workflow. This can be done both prior to deployment, using local but retrospective data, and on a continual basis post-deployment (“ongoing monitoring”) to make sure performance does not “drift” over time. This way, AI features can be deemed safe no matter what patient acuity mix is being served by a particular healthcare organization.

Local validation also matters because AI performance can be sensitive to the workflow into which the AI is integrated. For example, a popular use case for predictive modelling forecasts the likelihood of imminent medical deterioration. The output of such a model is not binary but instead a probability. When implementing this model, a deployer must decide what probability threshold should trigger an alert for the care team to investigate further. This choice has a significant impact on the “performance” of the model. If an organization sets the alert threshold low, there will likely be more false positives; if the organization sets the threshold higher, there will likely be more false negatives. Yet the organization may have good reason to change this setting: organizations with dedicated remote monitoring teams may prefer false positives, while those who must rely only on frontline nurses managing a large caseload of patients may not be able to tolerate them.

The challenge for post-deployment validation, in healthcare at least, is that small or rural healthcare organizations often struggle to afford the necessary data science resources and expertise to perform effective monitoring. Yet because AI performance will vary from organization to organization, post-deployment validation is not something that can be outsourced in practice. Without appropriate support for “local validation” activities, AI could widen, rather than narrow, the digital divide in healthcare.

AI developers can help. In 2024, for example, Epic released an open-source tool called the AI Trust and Assurance Suite that automates the most resource-intensive aspects of post-deployment monitoring. The suite's real-time evaluations use the deployer's local data and workflows and can monitor both Epic and third-party AI features, using the deployer's choice of evaluation methodologies.

However, even with access to the Suite, many lower-resourced healthcare organizations report that staffing and time constraints remain a barrier. A targeted grant program to support local validation for small and rural providers could help address this gap.

Providing Guidance for AI Validation

Even with the tools and resources needed to conduct local validation, healthcare organizations may not know how to use them effectively. Today, there is little consensus on how to evaluate traditional predictive models, and even less for generative AI. Gold-standard guidance such as the National Institute of Standards and Technology's AI Risk Management Framework asserts the importance of ongoing monitoring but does not specify how to perform it. Vendors providing evaluation services may use proprietary metrics or allow the user to select among evaluation methodologies (as the AI Trust and Assurance Suite does); both limit comparability.

Regulators can add value by defining standard metrics and validation processes for AI performance. Such standards need not be mandatory; guidance alone would increase confidence among healthcare AI adopters and lower the barrier to widescale adoption. Clear, shared methodologies would allow organizations to focus on performing validation rather than designing it—especially if incorporated into existing tools.

Importantly, validation guidance should be specific to both the use case and the deployment context, not one-size-fits-all. The risks, benefits, and appropriate metrics for documentation assistance, clinical decision support, predictive modeling, and patient engagement tools differ substantially. Tailored guidance for common classes of AI applications would allow organizations to focus their efforts on meaningful evaluation rather than reinventing validation frameworks from scratch.

Clear federal guidance would also accelerate adoption by lowering the perceived risk of deployment. Many healthcare leaders delay or limit AI use not because they believe AI is unsafe, but because they lack confidence that they are evaluating it correctly. By clarifying expectations and establishing shared norms, the Department could help organizations move from cautious experimentation to responsible scale.

Building AI Fluency

Technology adoption depends not only on technical readiness, but also on human understanding and trust. Considerable attention has been paid to training clinicians on how to use AI tools effectively, verify outputs, and avoid automation bias. While this training is essential, a more foundational form of AI fluency is also needed to overcome barriers to adoption—particularly among non-technical stakeholders such as clinical leadership, frontline staff, and labor representatives.

In healthcare, AI is unlikely to displace workers. Persistent labor shortages across clinical roles, combined with rising demand driven by an aging population, mean that AI is increasingly necessary to sustain access to care. In practice, most healthcare AI use cases are designed to reduce administrative burden, support clinical judgment, and allow clinicians to spend more time with patients. Nonetheless, fears that AI will eliminate jobs remain widespread and increasingly surface in labor negotiations.

These concerns often stem from a lack of shared understanding about what healthcare AI does—and does not—do. Without credible, accessible education, AI adoption risks being perceived as a cost-cutting measure rather than a capacity-multiplying one. This perception can slow implementation, limit usage, or prompt organizations to constrain AI to narrow pilots rather than deploy it broadly where it could deliver the most benefit.

The federal government can play an important role in addressing this gap by supporting neutral, trusted education on AI's role in healthcare. This could include developing materials tailored to different stakeholder groups; supporting peer-to-peer learning networks among healthcare organizations; or incorporating AI fluency into existing workforce development and quality improvement programs. Such efforts would complement, rather than replace, vendor-led training by focusing on system-level understanding rather than tool-specific instruction.

Ultimately, however, the most effective way to build AI fluency is through experience. When clinicians and staff see AI reduce documentation burden, improve workflow efficiency, and support—not override—clinical judgment, skepticism gives way to trust. Policies that enable earlier, broader, and more responsible adoption of AI—such as those outlined in this response—will therefore do more to build AI fluency than education alone. By helping healthcare organizations move past prolonged pilots and into sustained use, the Department can ensure that AI's benefits are visible, tangible, and widely shared.