



February 20, 2026

To the Department of Health and Human Services:

Kooth USA welcomes the opportunity to submit comments to the above referenced advisory committee on the topic of what HHS can do to accelerate the adoption and use of AI as part of clinical care. Please find our comments attached.

For 25 years, Kooth has pioneered digital mental health with the National Health Service in the UK, helping people access care at their first moment of need.

Today, our services are available to over 20 million people, supporting hundreds of thousands of young users every year. Building on that experience, Kooth is now a core partner in California's Children and Youth Behavioral Health Initiative (CYBHI), where we've designed and launched Soluna: the nation's first statewide, free digital mental health service for 13–25 year olds. Co-designed with youth, Soluna removes traditional barriers, expands access, and integrates seamlessly with existing supports to help a generation build resilience and thrive.

For any questions or comments, please feel free to contact me.

Sincerely,

Dr Patrick Johnston

Chief Operating Officer, Kooth
pjohnston@kooth.com

1. What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

The most significant barriers to innovation in AI for health care, particularly in the digital mental health care space, include:

- Regulatory ambiguity, including uncertainty around medical device classification, creates unintended market consequences. Products with clinical evidence of health improvement face significantly higher regulatory burden, including substantial administrative overhead, certification costs, and extended approval timelines. By contrast, general purpose or wellness AI products face less regulatory scrutiny, thereby skewing the market toward adopting tools optimized for increasing user engagement rather than verifiable health outcomes or population health impact.
- Fragmented expectations across federal agencies and state laws, creating well meaning but at times contradictory expectations
- Regulators' and lawmakers' uneven understanding of "AI" as a category, for example, conflating LLM-based chatbot technologies with classical statistical or machine learning algorithms, which leads to broad, technology-blunt rules that may be designed to curb chatbots but in practice unnecessarily constrain lower-risk Machine Learning workflows and proven technologies.
- Lack of knowledge about AI or training on how to effectively apply AI tools in practice, for front-line practitioners/health care providers.
- Lack of trust from some patients or end users, driven by AI models that lack explainability or transparency, and corresponding concerns about bias, safety, and privacy.

2. What regulatory, payment policy, or programmatic design changes should HHSI prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.

HHS should prioritize policies that incentivize the use of AI to help clinicians do the parts of the job that only clinicians can do, instead of replacing those roles. That means building payment structures that reimburse tools used for chart reviews, documentation summaries, initiating assessments, or additional clinical decision support functions, so clinicians can spend more time with patients and health systems are supported in using AI to strengthen care rather than reduce it.

3. For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (e.g., relating to liability, indemnification, privacy, and security). What novel legal and implementation issues exist and what role, if any, should HHS play to help address them?

There is growing inconsistency in how AI systems are classified as medical devices based on context rather than risk. In several states, legislatures have enacted or are considering

requirements for general-purpose generative AI chatbots to implement automated risk or safety assessments. However, when functionally similar risk-detection, triage, or safeguarding algorithms are embedded within purpose-built user experiences in a health care setting, those same capabilities may be deemed to constitute a regulated medical device.

This inconsistency creates both a regulatory arbitrage problem and unresolved liability questions. Developers are incentivized to deploy general-purpose AI tools with minimal oversight, despite their ability to generate health-related advice, while avoiding more responsible, context-aware designs that would trigger medical device regulation. At the same time, existing governance frameworks provide limited clarity on liability when AI is used in clinical practice, for example, where AI tools are used beyond their intended scope in ways that effectively constitute de facto medical device use without corresponding regulatory oversight.

HHS could play an important role by clarifying accountability and shared responsibility across developers, deployers, and clinicians; by providing guidance on appropriate use, misuse, and “off-label” deployment of non-medical device AI in clinical contexts; and by promoting risk-based, functionally consistent approaches that support safer adoption while maintaining a level playing field for responsible innovators.

4. For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and other workflow and human-centered evaluation methods for clinical care? Should HHS further support these processes? If so, which mechanisms would be most impactful (E.g., contracts, grants, cooperative agreements, and/or prize competitions)?

The AI Evaluation approaches holding greatest external validity combine:

- Pre-deployment and/or pilot testing for accuracy, robustness, and failure modes.
- Post-deployment monitoring of real-world performance, bias, and unintended consequences.
- Design Research on the human-centered experience, including patient/user experience, user trust, and clinician workflow, impact, and experience.
- Outcome-oriented metrics, not solely technical performance (e.g., impact on access, engagement, equity, and clinical outcomes).

In practice, it is far more common for developers to publish pre-deployment and/or pilot metrics on technical accuracy and performance, while glossing over post-deployment monitoring, human-centered research, and outcomes. Thus HHS should support the latter set of evaluation methods, for example through grant funding for transparent reporting on those, and publication of shared, more complete evaluation approaches.

5. How can HHS best support private sector activities (e.g., accreditation, certification, industry-driven testing, and credentialing) to promote innovative and effective AI use in clinical care?

HHS can best support private sector activities by:

- **Addressing regulatory asymmetry between AI explicitly designed to provide or support mental health, and those positioned as “wellness.”** In practice, people are turning to both of these and general purpose AI chatbots for mental health-related reasons. Baseline alignment in safety, transparency, and post-deployment expectations across these tools would avoid market incentives that favor deploying less governed products rather than made-for-purpose technologies rooted in safety and effectiveness.
- **Ensure evidence expectations are aligned with how AI is actually deployed in the real-world**, including acceptance of proportionate real-world evidence. HHS can drive innovation by rewarding AI systems that enable services to reach into underserved communities, drive improvements in clinical, social, and functional outcomes, enhance user safety, and drive system-wide efficiencies, rather than products that deliver outcomes only in clinical trial evaluations and settings.
- **Clarify accountability across general purpose AI or “wellness” AI product developers, AI clinical product developers, and clinicians.** HHS guidance on liability, appropriate use, and “off-label” usage of AI in health contexts would increase credentialing efforts and ultimately support safer adoption, particularly since users may not appreciate the practical difference in use case between the various types of AI products available to them.
- **Invest in shared evaluation infrastructure and outcomes benchmarks.** HHS can accelerate private-sector testing and credentialing by supporting common evaluation frameworks and shared health outcome measures, alongside interoperability standards across anonymous longitudinal clinical data and broader population health data, particularly in behavioral and mental health. For example, HHS could establish an AI credentialing or registry program, separate from FDA Software as a Medical Device registration, to monitor real-world AI performance and require transparent reporting on safety and outcomes over time. This way, the private sector would be publicly accountable to ensure that safety, equity, and real-world impact can be measured and tracked over time.

6. Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?

AI tools have extraordinary potential to expand access to mental health care services, especially given the current shortage of mental health professionals in the United States, in which as many as 169 million people live in federally designated Mental Health Professional Shortage Areas. Through human-in-the-loop AI augmented systems, health care providers

and clinicians can scale their impact and reach while ensuring a baseline standard of care through:

- In-the-moment feedback for best-practice coaching that raises baseline quality,
- Care navigation optimization that improves continuity of care and routes folks to the right level of care
- Risk and safeguarding systems that detect crisis signals

Such continuous quality assurance, triage, and risk detection is impossible at scale, without AI.

7. Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?

Adoption is influenced not only by payers and commissioners, but also by provider organizations, whose clinical governance bodies and procurement processes play decisive roles in determining allowable AI use. In turn, their determinations are influenced by their users, patients, and constituents. Additionally, with tools designed to enhance or augment practitioners, individual clinicians often exercise professional judgment in how they rely on or interpret AI outputs, further influencing how AI is adopted in practice.

Administrative hurdles for adopting AI generally stem from a lack of understanding of “AI” technology, for example, conflating general purpose conversational AI with made-for-narrow-purpose, human-in-the-loop AI systems with strict guardrails and clinical governance. There is opportunity for HHS to push an AI educational campaign which minimally dissociates the two classes of tools, and then instill public confidence in the societal benefit of the latter.

8. Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

Enhanced interoperability would most widen market opportunities and accelerate AI development where it enables linkage across clinical care, population health, and outcomes data. Interoperability is particularly impactful where longitudinal clinical data could be anonymously linked with broader datasets, including population health and economic wellbeing trends; machine learning could then be used to support modelling of need, measurement of impact, and even forecast evolution of care demand and risk over time.

9. What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?

In the digital mental health space, youth expect AI to provide a personalized experience that adapts based on how they interact with the technology; thus it is expected that health

outcomes will be achieved through personalized and adaptive care pathways. In turn, practitioners and clinicians expect AI to be used to reduce or eliminate administrative friction, repetitive, automatable non-clinical tasks, and generally optimize for providing human-to-human care. The central concern that both youth and practitioners have is around loss of human connection; for that reason it is important that HHS promote AI use in clinical care subject to preserving the power of human connection and that human-led care is a core tenet of our health care system.

10. Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care?

HHS should prioritize research that evaluates AI tools in real-world settings, using large scale, longitudinal datasets that capture meaningful health outcomes for anonymized individuals. While synthetic data and LLM evaluations (e.g. one LLM assessing another) can be useful at early prototype stages, they are insufficient as evidence of external validity, safety, and value in clinical care. Adoption will accelerate only if research reflects how AI performance varies across populations, care settings, and contexts, and if real-world effectiveness (not accuracy from simulated cases) is treated as the evidentiary standard for deployment.