

Comment on HHS Health Sector AI RFI

Docket: HHS-ONC-2026-0001

Submitted by:

Eli Talbert, PhD

Senior Data Scientist, United States Special Operations Command

The submitter works on operational AI governance in defense decision-support contexts, with a focus on accountability, documentation, and escalation design; this comment applies those governance principles to clinical AI adoption and oversight.

RFI Questions Addressed: Q1, Q2, Q3, Q4, Q7, Q8, Q10

Executive Summary: Four Immediate Policy Actions

Regulatory uncertainty and liability ambiguity are among the most significant barriers to clinical AI adoption. The actions below would accelerate deployment by establishing a common accountability floor—reducing the duplicative institutional governance costs and defensive non-adoption that currently slows innovation. To that end, HHS should:

1. **Issue ONC interpretive guidance requiring minimum decision logging for high-impact AI-enabled clinical decision support integrated into EHR workflows.** (Responds to RFI Questions 2, 3, 8)
2. **Clarify accountability allocation expectations for non-device AI through joint ONC–FDA guidance.** (Responds to RFI Questions 2, 3)
3. **Direct CMS to incorporate accountable AI documentation standards into Conditions of Participation for hospitals deploying high-impact AI tools.** (Responds to RFI Questions 2, 7)
4. **Fund NIH/AHRQ implementation science pilots testing post-deployment human–AI interaction metrics, including override monitoring.** (Responds to RFI Questions 4, 10)

These actions are likely to reduce liability ambiguity, strengthen interoperability, and create measurable governance standards that accelerate responsible adoption, though empirical evaluation of adoption and compliance effects should accompany implementation. This comment does not propose new regulatory agencies, pre-market approval requirements for non-device AI, or blanket restrictions on deployment. It proposes minimum documentation and logging standards that reduce the regulatory uncertainty currently driving defensive non-adoption and duplicative institutional compliance costs.

For purposes of this comment, “high-impact AI” refers to AI systems that materially influence diagnostic, triage, medication, procedural, or discharge decisions, particularly in settings where clinician workload, time pressure, or financial incentives may shape reliance patterns.

I. Close the Regulatory Gap for Non-Device Clinical AI

[Responds to RFI Questions 1, 2, 3, 7]

Governance challenges in clinical AI mirror those in other high-stakes, time-pressured human-algorithm decision environments: authority boundaries, escalation design, and documentation standards determine whether hybrid systems remain accountable over time. In defense decision-support systems, these are not optional governance add-ons but prerequisites for deployment authorization. Clinical environments differ from military command structures in training standardization, liability regimes, and autonomy structures, but both domains confront the same core challenge: maintaining auditable authority boundaries in high-stakes human–algorithm collaboration.

Existing HHS frameworks already address important components of this landscape, including:

- FDA regulation of Software as a Medical Device (SaMD)
- Predetermined Change Control Plans (PCCP) governing planned device modifications
- ONC HTI-1 Decision Support Intervention (DSI) transparency requirements (45 CFR § 170.315)
- CMS documentation standards under the Conditions of Participation (42 CFR § 482.24)

For clarity, this comment distinguishes between (a) FDA-regulated device software pathways, including SaMD submissions and PCCP-governed modifications, and (b) ONC-certified health IT functionality and CMS documentation expectations applicable to EHR-integrated decision support that may not meet device definitions. The governance gap identified here primarily concerns workflow-integrated, non-device AI functionality embedded within certified health IT environments.

Within this non-device space, a structural gap remains. Current frameworks emphasize model transparency, performance reporting, and change management. They do not specify:

- How clinician override behavior must be logged
- How authority boundaries between model and clinician are operationally defined
- How escalation pathways are monitored over time

This creates ambiguity in accountability allocation and increases institutional hesitation to deploy AI at scale.

HHS can close this gap through interpretive guidance clarifying documentation and accountability expectations, using existing regulatory authority under 45 CFR § 170.315 (ONC certification criteria) and 42 CFR § 482.24 (CMS Conditions of Participation for medical record services).

II. Establish Minimum Decision Logging Standards

[Responds to RFI Questions 2, 4, 8]


HHS should establish — through ONC guidance or certification updates under 45 CFR § 170.315 — a minimum logging requirement for high-impact AI-enabled CDS.

At minimum, systems should record:

1. Timestamp of model output
2. Model recommendation or risk score
3. Clinician action (accept / modify / override)
4. Override flag with standardized rationale category
5. Escalation pathway triggered (if any)
6. Linkage to downstream order or documented outcome
7. Patient notification flag (whether the patient was informed AI contributed to the decision)

Illustrative workflow: An EHR-integrated sepsis risk model generates a high-probability alert for a patient in a medical-surgical unit. The attending clinician reviews the alert but determines that the elevated score reflects a known chronic condition, and overrides the recommendation. Under the proposed schema, the system logs the model output (Elements 1–2), the override action with a standardized rationale code (Elements 3–4), and the absence of an escalation trigger (Element 5). Downstream, the patient’s outcome is linked to the decision (Element 6). This structured record enables quality review to distinguish informed clinical judgment from unexamined dismissal—without requiring real-time intervention or additional clinician workflow steps.

Figure 1. Human–AI Accountability Loop (Minimum Viable Decision Logging)



Critically, this schema imposes minimal marginal burden. Elements 1–3 build directly on existing certified health IT audit logging capabilities required under 45 CFR § 170.315. Elements 4–7 require structured configuration within EHR workflows but not novel infrastructure or hardware investment. **The absence of a common standard imposes higher costs: each deploying institution must independently design governance documentation, negotiate liability allocation with vendors, and defend ad hoc logging practices—duplicative effort that a federal floor would eliminate.** Standardization reduces institutional compliance burden while creating the interoperable evidence base that regulators, payers, and health systems need to support scaled deployment.

Implementation Feasibility

Certified health IT must support audit logging of actions related to electronic health information. However, AI-specific decision telemetry, including structured capture of model output, override rationale, and escalation triggers, typically requires additional configuration or integration within EHR workflows. Elements 1–3 build on existing audit logging capabilities, while Elements 4–7 generally require structured configuration but not novel infrastructure. The patient notification flag would indicate whether AI involvement was disclosed in accordance with institutional policy; this proposal does not prescribe the notification workflow itself. The schema should be machine-readable and expressible via FHIR AuditEvent or equivalent interoperable standards, consistent with USCDI expansion goals.

This aligns with:

- HTI-1 transparency requirements under 45 CFR § 170.315(b)(11)
- FDA Predetermined Change Control Plan (PCCP) provisions under FD&C Act § 515C (21 U.S.C. § 360e-4), which address planned modifications and when additional premarket submissions are required; post-deployment monitoring expectations may arise under device quality system and postmarket reporting frameworks (e.g., 21 CFR Parts 803 and 820), though PCCP provisions are primarily change-control focused.
- Regulatory concern regarding superficial review or rubber-stamping of AI recommendations, consistent with HHS and OIG compliance principles on meaningful clinician oversight

Such a schema would not inhibit innovation. It would reduce uncertainty and demonstrate accountable deployment. Additionally, structured logging enables detection of disparate override or recommendation patterns across protected characteristics, supporting existing nondiscrimination obligations under federal civil rights law and strengthening patient trust in AI-assisted care.

III. Ground Oversight in Empirical Evidence

[Responds to RFI Questions 4, 10]

Automation bias in clinical decision support has been documented in peer-reviewed literature spanning over a decade, with recent findings confirming that the problem persists in production AI deployments:

- Goddard et al., Journal of the American Medical Informatics Association (2012): Systematic review establishing that automation bias is a consistent and measurable phenomenon in clinical decision support, identifying frequency, effect mediators, and mitigators.
- Lyell & Coiera, Journal of the American Medical Informatics Association (2017): A PRISMA-compliant systematic review of 40 studies that found that automation bias increases as verification complexity rises. This dynamic is particularly relevant to AI-

enabled clinical decision support, where clinicians often operate under high information load and time pressure.

- Gaube et al., NPJ Digital Medicine (2021): In a randomized study of chest radiograph interpretation, clinicians were influenced by both accurate and inaccurate advice. Diagnostic accuracy declined substantially when participants received incorrect recommendations, regardless of whether the advice was labeled as coming from an “AI system” or a “human expert.” The source label affected perceived credibility ratings but did not significantly change diagnostic performance.
- Sivaraman et al., Proceedings of the CHI Conference on Human Factors in Computing Systems (2023): Showed that clinician trust calibration in AI-based treatment recommendations is difficult to achieve, and that reliance patterns go beyond binary accept/reject decisions—reinforcing the need for structured override logging.
- Khera, Simon & Ross, JAMA (2023): Editorial accompanying Jabbour et al. empirical research; warned that even in controlled settings without time pressure, clinicians favored AI-based decision-making systems despite contradictory or clinically nonsensical information, identifying automation bias as a direct risk of harm from AI-driven CDS.
- Dratsch et al., Radiology (2023): In a study of AI-assisted mammography (BI-RADS) categorization, incorrect AI suggestions significantly impaired diagnostic performance across all experience levels. While greater expertise reduced the magnitude of the effect, even very experienced radiologists remained measurably susceptible to incorrect recommendations.

These findings collectively establish that performance metrics alone do not capture behavioral adaptation in hybrid human–AI systems. Over-reliance patterns emerge even among experienced clinicians and are amplified under conditions of high information load and complexity. Existing evidence suggests that transparency or labeling interventions alone do not reliably eliminate these effects.

Structured override monitoring enables evidence-based oversight without imposing blanket restrictions on deployment. Logging enables institutions to detect, quantify, and respond to behavioral adaptation patterns.

IV. Clarify Accountability Allocation for Non-Device AI

[Responds to RFI Questions 2, 3]

HHS should clarify expectations regarding:

- Designation of accountable organizational owner for AI deployment
- Documentation requirements for AI-assisted high-risk decisions
- Distinction between device-regulated AI and workflow-integrated non-device AI

This clarification could be implemented through:

- ONC interpretive guidance under HTI authority (45 CFR § 170.315)
- Joint FDA–ONC guidance on hybrid accountability allocation
- CMS Conditions of Participation updates referencing documentation standards (42 CFR § 482.24)

Clear guidance reduces defensive non-adoption driven by liability uncertainty. Such clarification reduces uncertainty without altering tort standards; it establishes documentation expectations that support, rather than supplant, existing malpractice and product liability frameworks.

V. Fund Research on Human–AI Interaction in Deployed Systems

[Responds to RFI Questions 4, 10]

HHS should fund applied research to:

1. Measure longitudinal override patterns in deployed systems.
2. Identify escalation pathway shifts following AI integration.
3. Develop standardized human–AI coordination metrics.
4. Pilot interoperable logging schemas across health systems.

Implementation science funding through NIH, AHRQ, or cooperative agreements (including CRADAs) would generate the evidence base necessary for proportionate regulation. Contracts, cooperative agreements, and prize competitions (as referenced in the RFI) are all viable mechanisms; cooperative agreements may be particularly effective for multi-site logging schema pilots that require cross-institutional coordination.

VI. Align Reimbursement Incentives with Documentation Standards

[Responds to RFI Question 2]

This comment focuses primarily on regulation and research levers. However, the documentation standards proposed above have natural connections to CMS payment policy that would reinforce adoption:

- CMS could condition New Technology Add-on Payments (NTAPs) or Advanced APM incentive payments on demonstrated AI governance documentation, including the logging schema described in Section II.
- Quality measure reporting under the Merit-based Incentive Payment System (MIPS) could incorporate AI-assisted decision documentation as an Improvement Activity,

creating a positive incentive for adoption of accountable logging without mandating specific technology choices.

- Aligning documentation standards across regulatory and reimbursement frameworks would reduce the compliance burden for health systems deploying AI, rather than creating parallel reporting obligations.

These reimbursement connections are noted here to signal their feasibility; detailed payment policy design is beyond the scope of this comment but warrants dedicated analysis in future rulemaking.

Conclusion

By clarifying documentation standards and encouraging minimal telemetry requirements, HHS would create conditions likely to:

- Reduce liability ambiguity
- Support responsible innovation
- Strengthen patient trust
- Provide regulators with objective oversight tools

These measures are proportionate, technically feasible, and aligned with existing regulatory architecture. Much of this could be pursued under existing authorities through interpretive guidance, certification updates, and program conditions; however, binding requirements would likely require notice-and-comment rulemaking or enforceable Conditions of Participation consistent with Administrative Procedure Act constraints.

The views expressed in this comment are those of the author and do not necessarily reflect the official policy or position of the Department of War, the U.S. Army, or the U.S. Government.

Appendix: RFI Question Crosswalk

The following table maps this comment's sections to the specific questions posed in the RFI.

RFI Q#	RFI Question (abbreviated)	Comment Section(s)
1	Biggest barriers to private sector AI innovation and adoption	Section I
2	Regulatory, payment, or programmatic changes to prioritize (with CFR citations)	Sections I, II, IV, VI
3	Novel legal and implementation issues for non-device AI	Sections I, IV
4	Most promising evaluation methods for non-device AI	Sections II, III, V
7	Roles and governing bodies influencing AI adoption; administrative hurdles	Sections I, IV
8	Enhanced interoperability to widen markets and accelerate AI	Section II (FHIR AuditEvent, USCDI)
10	AI research priorities; published findings on adopted AI tools	Sections III, V

Appendix: Key Regulatory Citations

CFR Citation	Authority	Relevance to This Comment
45 CFR § 170.315	ONC Health IT Certification Criteria	Regulatory hook for requiring decision logging in certified health IT under HTI-1 DSI criteria
45 CFR § 170.315(b)(11)	ONC DSI Transparency Requirements	Specific HTI-1 provision for decision support intervention source attributes and transparency
42 CFR § 482.24	CMS Conditions of Participation: Medical Record Services	Hook for CMS to require AI-assisted decision documentation in hospital medical records
FD&C Act § 515C (21 U.S.C. § 360e-4)	FDA Predetermined Change Control Plans	Change-control framework for planned device modifications; post-deployment monitoring expectations arise under related device quality system and postmarket reporting frameworks (21 CFR Parts 803, 820)