NEUROFLOW

Response to Health and Human Services Department Request for Information (RIN 0955-AA13)

February 23, 2026
HHS-ONC-2026-0001

Dr. Thomas Keane
Assistant Secretary for Technology Policy
Office of the National Coordinator
200 Independence Avenue, SW
Washington, DC 20201

Dear Dr. Keane,

NeuroFlow appreciates the opportunity to respond to the Department of Health and Human Services' Request for Information on accelerating the adoption and use of artificial intelligence (AI) in clinical care.

NeuroFlow develops AI-enabled clinical decision support (CDS) and population health tools deployed nationally across health systems, payers, and federal partners. Our systems operate exclusively within supervised clinical workflows. They do not provide autonomous medical advice or function as direct-to-consumer therapeutic agents. They augment licensed clinicians by supporting early detection, risk stratification, care matching, and measurement-informed care at scale.

Public discussion of AI in healthcare often centers on patient-facing chatbots and autonomous systems. These tools warrant appropriate oversight, but they differ materially from clinician-supervised CDS embedded within accountable care teams. Federal policy should distinguish between these categories to ensure regulation remains proportional to risk.

Behavioral health is a particularly urgent domain for thoughtful AI policy. Comorbid behavioral and physical health conditions drive substantial excess medical spending, yet behavioral health has long lacked infrastructure parity in federal health IT investment. AI-enabled CDS can function as clinical infrastructure, supporting workforce capacity, improving access, and enabling integration under value-based care.

Federal policy should focus on aligning incentives around measurable outcomes and total cost of care, rather than treating AI as a discrete reimbursable intervention. When payment rewards prevention and integration, adoption of high-value clinical decision support follows naturally.

To accelerate responsible adoption, HHS should:

- Differentiate between supervised clinical decision support and autonomous patient-facing AI systems, applying proportionate scrutiny to autonomy and risk.
- Align reimbursement with measurable outcomes and total cost accountability.

- Expand risk-bearing value-based models in Traditional Medicare and Medicare Advantage to create natural incentives for adoption of population health AI infrastructure.
- Strengthen behavioral health accountability within Medicare Advantage and other federal programs to drive systematic measurement-based care.
- Accelerate implementation of USCDI+ Behavioral Health to ensure structured, interoperable behavioral health data.
- Understand AI in behavioral health as infrastructure that operationalizes measurement-informed care and enables scalable integration, not as a replacement for clinicians.

Question 1: What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

The primary barriers to adoption of AI in clinical care are not technical limitations in model capability. They are structural misalignment across regulatory interpretation, data infrastructure, workflow integration, and payment incentives.

<u>Failure to Distinguish Clinician-Supervised CDS from Autonomous AI</u>

Public and policy discourse frequently conflates autonomous, patient-facing AI systems with clinician-supervised clinical decision support tools. These categories differ significantly in both risk profile and governance structure.

Clinician-facing CDS tools operate within established clinical accountability frameworks. They surface information, risk signals, or guideline-based recommendations for licensed clinicians to independently evaluate and act upon. These tools are addressed under Section 520(o)(1)(E) of the 21st Century Cures Act (21 U.S.C. §360j(o)(1)(E)), which excludes certain CDS software from device regulation when clinicians can independently review the basis for recommendations. Autonomous AI systems that independently generate therapeutic advice or make patient-facing care decisions without clinician oversight present a separate regulatory discussion entirely.

When these two categories are treated as equivalent, health systems often apply heightened scrutiny designed for autonomous systems to clinician-supervised CDS. This creates unnecessary friction and slows adoption of lower-risk infrastructure tools designed to augment professional judgment rather than replace it.

<u>Behavioral Health Infrastructure Deficits</u>

Behavioral health continues to operate with significantly weaker health IT infrastructure than physical health. The HITECH Act catalyzed widespread EHR adoption across medical specialties but largely excluded behavioral health providers. Behavioral health EHR adoption and data standardization lag significantly behind the rest of the health care system, and this exclusion has had compounding effects over time.

Behavioral health documentation is frequently unstructured. Data elements such as symptom scales, suicide risk assessments, and treatment response metrics are inconsistently captured across organizations. Exchange of behavioral health data across systems remains limited, and integrated behavioral-physical health records are incomplete in many care settings.

Clinical decision support systems rely on structured, longitudinal data to function optimally. While AI models can incorporate physical health data, utilization patterns, and medication history when behavioral health data are sparse, incomplete behavioral health documentation reduces precision and limits scalability. Without improvements to behavioral health data infrastructure, the clinical value of AI tools in this domain will remain constrained.

Workflow Integration and EHR Friction

Adoption depends on seamless integration into routine care delivery systems. Clinical decision support must align with established workflows and integrate into certified EHR platforms without disruption. In practice, technical and contractual barriers slow third-party integration and limit timely API access, delaying deployment even when the underlying tools are clinically ready.

CMS's Health Tech Ecosystem initiative appropriately recognizes these integration challenges. Continued progress toward more open, API-enabled interoperability will be critical to accelerating responsible adoption of clinician-supervised AI tools across care settings.

Misaligned Payment Incentives

Incentive alignment remains the final and perhaps most consequential structural barrier. In payment environments that reward volume, organizations that invest in prevention and population management may not see immediate financial return. Preventing a psychiatric hospitalization or reducing emergency department utilization improves outcomes and lowers total cost of care, but in fee-for-service models that value accrues to payers rather than the providers who made the investment.

Risk-bearing value-based arrangements create clear incentives to adopt infrastructure that supports early detection, measurement-informed care, and avoidance of preventable utilization. Until payment structures reward outcomes rather than activity, the organizations best positioned to deploy behavioral health AI will face a financial disincentive to do so.

The largest barriers, in short, are not algorithmic. They are regulatory clarity, behavioral health data infrastructure, workflow integration, and payment alignment. Addressing all four is necessary for responsible adoption at scale.

Question 2: What regulatory, payment policy, or programmatic design changes should HHS prioritize and why?

HHS can accelerate responsible adoption through targeted clarity and incentive alignment across four priority areas.

Reaffirm the Clinical Decision Support Exemption Under the 21st Century Cures Act

Section 520(o)(1)(E) of the Federal Food, Drug, and Cosmetic Act (21 U.S.C. §360j(o)(1)(E)) excludes certain clinician-supervised CDS tools from device regulation when they enable health care professionals to independently review the basis for recommendations. Behavioral health CDS tools that surface evidence-based treatment pathways, risk stratification insights, and treatment response predictions for clinician review fall squarely within this framework. Reaffirming this statutory boundary supports predictable innovation and reduces the procurement hesitation that slows adoption of lower-risk supervised tools.

Modernize Consent Infrastructure Under 42 CFR Part 2

42 CFR Part 2 governs the confidentiality of substance use disorder treatment records and currently prevents AI algorithms from accessing SUD data even within integrated care teams where patients have consented to coordinated treatment. This restriction is particularly consequential for behavioral health AI, because patients with co-occurring substance use and mental health conditions represent some of the highest-risk, highest-cost individuals in any population. When AI risk stratification models cannot access complete clinical records, their precision is compromised for exactly the patients who stand to benefit most. Modernized, electronic consent mechanisms that enable secure data exchange for treatment, payment, and operations would improve integration of behavioral and physical health data while preserving the privacy protections patients depend on. Comprehensive data access improves AI-supported risk stratification and enables coordinated care teams to act on complete clinical pictures rather than partial records.

Accelerate Risk-Bearing Value-Based Payment Models in Traditional Medicare and Medicare Advantage

Value-based payment models, including accountable care organizations, bundled payments, and capitated arrangements, align financial incentives around total cost of care and outcomes. Behavioral health conditions drive substantial excess medical spending when comorbid with chronic disease. Organizations accountable for total cost of care develop natural incentives to deploy AI-enabled measurement and population health infrastructure, because preventing crises and reducing avoidable utilization directly improves their financial position. Payment reform, not new technology codes, will be the primary driver of sustainable adoption. CMS should expand risk-bearing arrangements in Traditional Medicare and Medicare Advantage and ensure behavioral health outcomes are explicitly included in accountability structures.

Current time-based billing in programs like the Collaborative Care Model, billed under CPT codes 99492 through 99494, creates a structural disincentive against efficiency-enhancing technology. Organizations using AI tools to manage larger patient panels with better outcomes generate less billable time, not more. Payment structures should reward complexity management and measured clinical outcomes, including remission rates, crisis prevention, and reduced hospitalizations, rather than time spent on care management activities.

Strengthen Behavioral Health Accountability in Medicare Advantage and Medicaid

Incorporating robust behavioral health outcome measures into Medicare Advantage Star Ratings and managed care contracts would drive systematic adoption of measurement-informed care infrastructure. Relevant measures include remission benchmarks, proactive screening rates, and avoidable crisis utilization. Higher Star Ratings and associated quality bonus payments would reward plans using validated measurement systems achievable at population scale through AI-enabled infrastructure, without mandating specific technologies or creating new approval requirements.

Establish Federal Behavioral Health AI Validation Infrastructure

HHS should support representative, de-identified validation datasets and shared evaluation methodologies through ARPA-H, NIMH, and related public-private initiatives. Limitations in small, non-generalizable datasets undermine trust and slow adoption. Federally supported validation infrastructure would reduce duplicative review burdens on health systems, improve transparency in how tools are evaluated, and create common benchmarks that allow organizations to assess AI performance based on measurable evidence rather than vendor claims.

Question 3: For non-medical devices, what novel legal and implementation issues exist and what role should HHS play?

A primary implementation concern for AI-generated high-risk alerts, such as those triggered by suicide risk detection or crisis prediction models, is liability. Health systems currently lack clear expectations for what constitutes an adequate and defensible clinical response, particularly when immediate intervention is operationally unrealistic given staffing and resource constraints. A clinician who receives an AI-generated urgent alert faces a genuine question about what the standard of care requires in response. HHS can reduce this uncertainty by articulating best practices and tiered response expectations for AI-generated urgent alerts that account for real-world organizational capacity, without creating new liability exposure for organizations deploying these tools responsibly.

A related gap involves the scope of existing privacy frameworks. HIPAA's definition of covered entities does not adequately address the range of digital health and AI tools now operating in clinical contexts. Patient-facing tools and AI applications that connect to EHRs may handle sensitive behavioral health information without falling clearly within existing privacy and cybersecurity requirements. HHS should clarify how HIPAA applies to AI-enabled CDS tools operating within supervised clinical workflows and identify where gaps in coverage require updated guidance or rulemaking. Behavioral health data carries particular sensitivity, and patients and clinicians alike need assurance that the frameworks governing its use are keeping pace with the technology.

Health systems and payers also lack consistent criteria for evaluating AI solutions before procurement. Without shared standards, each organization conducts its own bespoke security, validation, and governance review. Vendors face inconsistent expectations across organizations, and health systems shoulder duplicative evaluation burdens that extend timelines and consume resources. HHS guidance on explainability, traceability, transparency, and validation would give organizations a coherent framework for accountable deployment. It would also create more consistent market expectations, reducing friction for responsible vendors and raising the floor for the field as a whole.

Question 4: For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and workflow and human-centered evaluation methods for clinical care? Should HHS further support these processes?

For clinician-supervised clinical decision support systems, evaluation should rest on four principles: explainability, traceability, transparency, and validation.

Explainability requires that model recommendations align with established evidence-based workflows and clearly surface the clinical variables that influenced each suggestion. Clinicians must understand how outputs relate to guideline-concordant care pathways so they can exercise independent judgment. Evaluation should confirm that recommendations are anchored in recognized medical standards rather than opaque or unconstrained reasoning. A clinician who cannot understand why a tool flagged a patient cannot responsibly act on that flag or override it.

Traceability ensures that every model output can be audited. Systems should maintain logs capturing relevant data inputs, decision thresholds, and model version history. Organizations can then review how recommendations were generated, support internal quality oversight, and respond to adverse events with a clear record of what the system surfaced and when.

Transparency requires disclosure of intended use, model limitations, validation evidence, and deployment context. Organizations should communicate clearly where and how models operate within clinical workflows and maintain recognized security and data protection standards. Transparency strengthens institutional trust, enables informed procurement decisions, and gives clinical governance committees the information they need to evaluate tools against their own patient populations.

Validation requires ongoing assessment of clinical accuracy, fairness across patient populations, and measurable impact on care delivery. For behavioral health applications, validation should examine outcomes such as access to care, symptom improvement, crisis identification, and avoidable utilization. Importantly, validation is not a one-time event. AI models that continue learning and adapting require ongoing monitoring to detect performance drift over time. Evaluation should remain proportional to risk and aligned with existing clinical governance structures rather than imposing entirely new oversight mechanisms.

HHS can accelerate responsible adoption by promoting these core governance principles and encouraging shared outcome benchmarks for behavioral health applications. Clear expectations around explainability, traceability, transparency, and validation will support innovation while reinforcing patient safety and public trust.

Question 5: How can HHS best support private sector activities (e.g., accreditation, certification, industry-driven testing, and credentialing) to promote innovative and effective AI use in clinical care?

HHS can best support responsible private-sector activity by setting clear baseline expectations while allowing accreditation and testing frameworks to evolve through industry and professional collaboration.

Rather than creating new technology-specific certification programs, HHS should articulate core expectations for clinician-facing clinical decision support: documentation of intended use, disclosure of validation evidence, auditability of outputs, and adherence to recognized security standards. Clear federal signals reduce procurement uncertainty without imposing duplicative regulatory structures that slow innovation or disadvantage smaller developers who lack resources for parallel compliance programs.

HHS can also strengthen industry-driven testing by supporting shared benchmarking infrastructure. De-identified, demographically representative validation datasets in high-impact domains such as behavioral health would enable voluntary comparative evaluation across tools using common outcome measures. This infrastructure should function as research support rather than certification, promoting transparency and comparability without creating a new approval regime. Datasets should intentionally include rural populations, Veterans, and high-comorbidity Medicare beneficiaries, groups that are consistently underrepresented in current validation studies and whose care needs are among the most acute.

Federal convening authority and procurement policies can reinforce these principles by rewarding transparency and measurable outcomes. Promoting clarity, comparability, and collaboration rather than additional licensing layers will foster a competitive marketplace that supports both innovation and public trust.

Question 6: Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?

AI has met or exceeded expectations when deployed within supervised clinical workflows to address defined operational and clinical gaps.

In behavioral health integration, AI-enabled referral matching has reduced time to treatment and lowered behavioral health-related inpatient utilization while improving total cost of care. When primary care providers are supported by algorithmic matching tools that account for clinical characteristics, patient preferences, and real-time network capacity, access improves and downstream crisis utilization declines.[1]

Natural language processing and language models applied to unstructured patient-reported data has demonstrated particular value in suicide risk detection and crisis response.[2] NLP systems identify risk signals that structured screening tools alone may not capture, enabling earlier clinician review and intervention when embedded in supervised workflows. In a retrospective database study, more than half of patients flagged by NLP analysis of free-text inputs would not have been identified through structured screening alone.[3] These patients were at elevated risk and would have gone undetected without passive, continuous analysis of the language they used outside of formal assessment windows. AI has extended this capability to crisis platforms as well, where call-, chat-, and text-based services face volume challenges that staffing alone cannot solve. AI-enabled triage systems identify high-risk contacts and route them to appropriate responders faster than manual review allows, reducing wait times and improving the match between contact severity and responder skill. AI has also demonstrated value in training crisis responders, using simulated scenarios to build consistent skills at a frequency and scale that traditional supervision cannot achieve. These applications do not replace human responders. They ensure that the highest-risk individuals reach a prepared clinician faster.

Population-level use cases have also shown meaningful impact. Suspect diagnosis identification and behavioral health risk stratification models help care teams identify patients with undiagnosed or under-managed conditions, particularly those with comorbid medical illness. When paired with measurement-informed care, these systems support more appropriate resource allocation and reduce avoidable emergency department utilization.[4]

Measurement-informed care represents one of the highest-value and most evidence-supported AI applications in behavioral health. MIC involves systematic use of validated assessment instruments during clinical encounters to track symptom trajectories, detect non-response, and guide treatment adjustments. Despite demonstrated clinical benefits, MIC is underutilized. This is largely because manual administration, scoring, and visualization are administratively burdensome. AI removes that barrier directly. Automated assessment administration, real-time scoring, and longitudinal visualization of patient trajectories enable clinicians to identify deterioration or treatment non-response between scheduled appointments. When AI operationalizes MIC at population scale, it transforms a proven but underutilized practice into routine infrastructure.

AI systems that operationalize evidence-based treatment pathways, by surfacing guideline-concordant options or predicting treatment response based on patient characteristics, have shown promise in reducing trial-and-error prescribing and improving symptom trajectories. These applications are particularly valuable in primary care settings where specialty behavioral health expertise is limited and clinicians may be managing behavioral health conditions without dedicated support.

Administrative support tools, including automated documentation assistance and coding validation, have produced efficiency gains that reduce clinician burden and support program integrity without altering clinical decision authority. These tools are often underappreciated as clinical infrastructure, but clinician burnout and documentation burden directly constrain the capacity of care teams to engage patients meaningfully.

Where AI tools fall short is typically in one of three contexts: when deployed outside supervised care environments, when not integrated into existing workflows, or when disconnected from measurable accountability structures. Tools that function as standalone digital interventions without clear linkage to care teams or outcome measurement have produced less consistent impact. The evidence base for direct-to-consumer therapeutic AI remains far thinner than for clinician-supervised decision support, and the governance structures required to ensure safety in autonomous patient-facing contexts are still developing.

The greatest potential lies in AI applications that operationalize measurement-informed, guideline-concordant care for patients with complex behavioral and physical comorbidity. These populations account for disproportionate healthcare spending. Systems that support early detection, risk stratification, treatment matching, and proactive adherence monitoring within risk-bearing payment models offer the clearest pathway to improved outcomes and cost containment.[1,4]

Question 7: Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?

Adoption of clinician-facing AI is typically shaped by three groups: clinical leadership, information technology leadership, and compliance.

Chief Medical Officers and clinical governance committees evaluate whether a tool aligns with evidence-based standards, supports appropriate workflows, and demonstrates meaningful clinical impact. Without credible validation evidence and a clear articulation of intended use, clinical approval rarely advances. These committees often lack shared frameworks for evaluating AI tools and must construct their own review criteria from scratch, a process that is time-consuming and inconsistent across organizations.

Chief Information Officers, information security teams, and compliance leaders control integration, data access, cybersecurity review, and regulatory risk assessment. These functions often determine the pace of adoption more than clinical enthusiasm does. In the absence of widely accepted evaluation standards for clinician-facing AI, organizations conduct bespoke security, validation, and governance reviews independently, extending procurement timelines significantly. A tool that passes clinical review may sit in IT security review for months while the organization works through questions that have no shared industry answer.

The primary hurdle is not resistance to AI itself but variability in institutional review expectations. When no shared framework exists for assessing transparency, validation evidence, and accountability structures, each organization must define for itself what constitutes sufficient review. This variability disadvantages smaller health systems and safety-net providers that lack dedicated AI governance infrastructure and cannot absorb extended procurement timelines. Clear federal guidance around baseline transparency and governance expectations would reduce that variability while preserving appropriate institutional oversight. Consistency in expectations, rather than additional regulation, would support more efficient and responsible adoption across organizations of all sizes.

Question 8: Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

Enhanced interoperability in behavioral health would materially expand safe and scalable AI deployment.

The most consequential data elements include structured symptom assessments for depression, anxiety, and PTSD, along with suicide risk indicators, longitudinal treatment response data, medication management history, crisis utilization signals, and relevant social determinants of health. When these data are computable and exchangeable across care settings, decision support systems can generate more precise and clinically meaningful insights. Today, many of these elements exist as scanned forms, free-text notes, or proprietary formats that AI systems cannot reliably access or interpret. An emerging body of evidence also suggests that passively collected data from smartphones and wearables, including mobility patterns, sleep duration, and heart rate variability, can serve as objective proxies for mental health status, detecting symptom changes and predicting deterioration before adverse events occur. As this evidence matures, interoperability infrastructure that can accommodate these signals will become increasingly important to realizing their clinical value.

USCDI+ Behavioral Health represents an important step toward standardizing these elements. When certified EHR systems support consistent capture and API-based exchange of behavioral health measures, data completeness improves and reliance on fragmented documentation declines. Reliable interoperability enhances both clinical accuracy and comparability across organizations. It also enables AI developers to build and validate tools against representative datasets rather than the proprietary, single-system data that currently shapes most behavioral health models.

Interoperability also supports benchmarking and research. When outcome measures and utilization data are structured and exchangeable, health systems and payers can more transparently assess performance, compare approaches, and participate in multi-site evaluation efforts. This is particularly important for Medicare and high-comorbidity populations, where longitudinal, cross-setting data are essential to understanding total cost impact. Single-site studies cannot answer the questions that policymakers and payers need answered about AI performance at population scale.

Enhanced interoperability does more than facilitate AI development. It strengthens care coordination, measurement-informed practice, and accountability across the entire healthcare system. The behavioral health field has waited too long for infrastructure parity. USCDI+ Behavioral Health, combined with meaningful enforcement through EHR certification and payment policy, offers a credible path forward.

Question 9: What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?

Patients and caregivers consistently identify access, coordination, and timeliness as primary challenges in behavioral health care. Long wait times, difficulty finding appropriate providers, fragmented communication between medical and behavioral health clinicians, and delayed identification of worsening symptoms remain persistent concerns. AI-enabled decision support can help address these gaps by supporting earlier detection, more appropriate care matching, and proactive outreach within supervised care teams.[1]

For individuals managing co-occurring medical and behavioral health conditions, improved coordination between providers is especially important. Patients with these conditions frequently repeat their histories to multiple providers, re-disclose traumatic experiences, and navigate disconnected systems without a clear point of accountability. Tools that help clinicians recognize untreated depression, substance use disorder, or emerging suicide risk within routine care settings can reduce preventable crises and avoidable hospital utilization.[3,4]

At the same time, patients and caregivers express understandable concerns about privacy, data security, bias, and the potential erosion of the clinician-patient relationship. Behavioral health information carries particular sensitivity. Disclosures made in therapy, risk assessments documented in clinical notes, and crisis contacts recorded in health systems represent some of the most personal information a patient shares. Many patients fear that automated systems could displace human judgment or reduce the empathy that makes behavioral health treatment effective.

These concerns reinforce the importance of deploying AI as clinician-supervised decision support rather than as autonomous clinical actors. Transparency about how data are used, clear communication about the role of AI within care teams, and strong institutional oversight are essential to maintaining trust. Patients also want assurance that AI systems have been validated for populations that look like them, that rural patients and Veterans are not receiving tools trained exclusively on urban academic medical center data, and that the therapeutic relationship remains central to care.

When AI is positioned as infrastructure that supports rather than replaces licensed professionals, it can enhance safety, access, and continuity while preserving the clinician-patient relationship that patients consistently identify as essential to their care.

Question 10: Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care?

HHS should prioritize research that evaluates how clinician-supervised AI decision support affects real-world clinical outcomes and total cost of care in integrated settings.

Particularly important are studies focused on populations with co-occurring behavioral and physical health conditions, as these patients drive disproportionate utilization and spending across Medicare and Medicaid. Research should examine whether AI-enabled early identification, risk stratification, treatment matching, and measurement-informed care reduce avoidable hospitalizations, emergency department visits, and symptom persistence over time.[4] Most existing studies measure algorithmic performance rather than patient-level outcomes. The field needs pragmatic effectiveness trials in diverse real-world settings, including VA Medical Centers, FQHCs, rural hospitals, and community mental health centers, that link AI deployment to remission rates, crisis events, and total cost of care.

Implementation research is equally important. Understanding how AI tools are embedded into primary care and behavioral health workflows, how clinicians interpret and act on recommendations, and what governance structures sustain impact over time will inform scalable adoption strategies. Research should examine what organizational characteristics predict successful implementation, how financial incentives affect adoption fidelity, and what implementation support rural and safety-net providers need to deploy these tools effectively.

HHS should also prioritize research that examines equity and generalizability. Behavioral health prevalence, access patterns, and treatment engagement vary significantly across demographic and geographic groups. Current validation datasets often underrepresent rural populations, Veterans, and communities of color. Research must evaluate whether AI-enabled tools perform consistently across diverse populations and actively correct for gaps in historical training data that may reflect underdiagnosis rather than lower prevalence.

Finally, HHS should support research on the appropriate design and oversight of ongoing performance monitoring for adaptive AI systems. Traditional device validation assumes a static intervention. AI tools that continue learning in deployment require frameworks for detecting performance drift, managing model updates, and maintaining accountability over time. This is an area where the field lacks consensus and where federal investment in methodology would benefit all stakeholders.

Question 10a: Are there published findings about the impact of adopted AI tools and their use in clinical care?

Yes. Peer-reviewed studies demonstrate measurable clinical and financial impact when AI-enabled tools are deployed within supervised care environments.

Technology-enabled referral optimization was associated with improved access to behavioral health care, reduced behavioral health-related inpatient admissions, and reductions in total medical spending in a risk-adjusted insured population.[1]

Natural language processing for suicide risk detection identified individuals with suicidal ideation who would not have been captured by structured screening alone. In a retrospective database study, more than half of flagged users may not have been identified without NLP-based analysis of free-text inputs.[3] This finding has direct implications for population health programs that rely solely on structured assessment tools to identify patients at risk.

Technology-enabled behavioral health integration was associated with reductions in emergency department utilization and improvements in depressive symptom trajectories when measurement-informed care was systematically applied.[4] These results held across primary care and OBGYN settings, suggesting the model generalizes across care contexts rather than depending on specialty behavioral health infrastructure.

Collectively, these findings indicate that clinician-facing AI tools, when integrated into care teams and tied to measurable outcomes, can improve access, reduce crisis utilization, and support cost containment.

Question 10b: How does the literature approach the costs, benefits, and transfers of using AI as part of clinical care?

The literature primarily evaluates avoided utilization, symptom improvement, and short-term changes in medical spending. Fewer studies assess long-term cost transfers across payment models or fully account for implementation costs, including technology infrastructure, staff training, workflow redesign, and ongoing maintenance.

Evidence increasingly suggests that AI-enabled behavioral health integration produces the strongest financial signal in risk-bearing or accountable care environments, where reductions in hospitalizations and emergency utilization translate directly into total cost savings.[1,4] In fee-for-service models, cost reductions often accrue to payers rather than providers, dampening the adoption incentive for clinical organizations that bear the implementation cost without capturing the financial benefit.

This misalignment between who invests and who benefits is a core structural problem that payment reform must address. Economic analyses that ignore transfer effects, specifically the question of which stakeholder captures the value generated by AI-enabled care improvement, will underestimate the barriers to adoption and overestimate the market incentives for deployment.

Future research should standardize economic evaluation methodologies, including attribution models and total cost of care analyses, to better quantify the fiscal impact of clinician-supervised AI across different reimbursement structures. Policymakers, payers, and health systems all need this evidence to make investment decisions, and the field cannot produce it without more rigorous and consistent approaches to economic evaluation.

Conclusion

Behavioral health AI used as clinical decision support does not displace licensed professionals. It operationalizes evidence-based pathways, strengthens measurement-informed care, and supports earlier identification of risk within accountable clinical environments. The regulatory architecture for these tools is already defined. The policy priority is ensuring that expectations remain proportionate to their function and risk profile.

In behavioral health, the need for this infrastructure is particularly urgent. Comorbid behavioral and physical health conditions drive disproportionate healthcare utilization and cost, especially within Medicare populations. Workforce shortages and fragmented data systems limit the ability of care teams to manage these conditions proactively. Behavioral health was excluded from the HITECH Act's investment cycle and fell further behind as a result. AI policy must not repeat that exclusion.

When aligned with risk-bearing payment models, interoperable data standards, and transparent governance principles, AI-enabled decision support can improve access, reduce avoidable utilization, and enhance accountability without compromising clinician judgment or patient trust. HHS has the opportunity to distinguish clearly among AI use cases, reinforce predictable regulatory pathways for supervised clinical decision support, and align incentives around measurable outcomes. Doing so will accelerate responsible adoption while preserving both innovation and patient safety.

We appreciate the opportunity to provide input and look forward to continued engagement as the Department advances its AI strategy. Please contact Matt Miclette, Vice President of Innovation & Policy, at matthew@neuroflow.com with any questions.

Respectfully submitted,

Matt Miclette

Matt Miclette
Vice President, Innovation & Policy
NeuroFlow

References

1. Zhu Y, Saynisch P, David G, Shatraw W, Mailloux A, Patel A, Dow T, Smith-McLallen A. Linking insured adults to behavioral health care: a cost-saving solution. Am J Manag Care. 2025 Dec;31(12):752-759. doi: 10.37765/ajmc.2025.89834. PMID: 41512258.

2. Holley, D., Daly, B., Beverly, B. et al. Evaluating Generative Pretrained Transformer (GPT) models for suicide risk assessment in synthetic patient journal entries. BMC Psychiatry 25, 753 (2025). https://doi.org/10.1186/s12888-025-07088-5

3. Hartz M, Hickey D, Acosta L, Brooks A, Holley D, Zaubler T. Using natural language processing to detect suicidal ideation and prompt urgent interventions: A retrospective database study. Innovations in Digital Health, Diagnostics, and Biomarkers. 2024;4:6-8. doi: 10.36401/IDDB-23-10.

4. Pardes A, Rene R, Chun P, Cherson M. Technology-enabled Behavioral Health Integration Decreases Emergency Department Utilization. Clin Pract Epidemiol Ment Health. 2022 Sep 20;18:e174501792208150. doi: 10.2174/17450179-v18-e2208150. PMID: 37274856; PMCID: PMC10158075.