waymark

2021 Fillmore Street Ste#1059
San Francisco, CA 94115-2708
P: (415) 212-8993
F: (866) 498-1987 (HIPAA-secure)

Department of Health and Human Services, Assistant Secretary for Technology Policy and the Office of the National Coordinator for Health Information Technology
Attention: Request for Information: HHS Health Sector AI RFI
Mary E. Switzer Building
Mail Stop: 7033A
330 C Street S.W.
Washington, DC 20201

Re: HHS Health Sector AI RFI

Dear Assistant Secretary for Technology Policy and the Office of the National Coordinator for Health Information Technology:

Waymark, Inc. ("Waymark" or "we" or "our") is a healthcare technology company deploying AI-driven population health management for Medicaid populations through managed care organization partnerships. We bring practical experience implementing reinforcement learning-based care coordination, predictive risk stratification, and automated intervention systems serving over 175,000 high-risk members across multiple states. Our AI-driven interventions have achieved a 20.4% reduction in avoidable emergency department visits, a 48.3% reduction in 30-day all-cause hospitalizations, and a $3.73:1 return on investment in operational Medicaid managed care deployments.

Waymark appreciates the opportunity to respond to HHS's Request for Information on accelerating AI adoption in clinical care.

The recommendations in this comment align with the Administration's stated priorities for AI policy and healthcare reform:

Reducing Regulatory Burden: Our recommendations for safe harbors, clear guidance, and harmonized federal standards would reduce compliance uncertainty and the duplicative state-by-state requirements that currently impede multi-state AI deployment. Regulatory clarity enables faster deployment of beneficial AI while maintaining appropriate safeguards.

Promoting American AI Leadership: The United States leads the world in healthcare AI innovation, but regulatory fragmentation, payment misalignment, and an emerging patchwork of state AI laws risk ceding this advantage to international competitors. Policy reforms enabling responsible AI deployment would strengthen American leadership in this strategically important sector.

Lowering Healthcare Costs: AI-driven population health management directly addresses healthcare cost growth by preventing avoidable acute care utilization, improving chronic disease management, and enabling more efficient resource allocation. Payment reforms that reward these outcomes would accelerate cost reduction.

Improving Healthcare Quality: AI enables personalized, proactive care that improves health outcomes for patients. Quality measure reforms and research investments would generate evidence demonstrating AI's contribution to better care.

Our response focuses on the highest-impact barriers to AI adoption and the most actionable policy solutions. We prioritize three areas where HHS action would most directly accelerate responsible AI deployment:

1. **Transitioning Medicaid payment models toward value-based arrangements that reward AI-driven outcome improvements.** Current payment structures represent the single greatest barrier to AI adoption. CMS guidance under 42 CFR Part 438 on performance-based payments for AI-driven population health management, together with model contract language for states to

adopt, would create the financial alignment necessary for managed care organizations to invest in AI systems that reduce total cost of care.

2. **Establishing regulatory safe harbors for population health AI systems, including fraud and abuse safe harbors for AI-driven member engagement.** Organizations deploying AI for care coordination face unresolved questions under the Anti-Kickback Statute and the Beneficiary Inducement Civil Monetary Penalty law. OIG guidance confirming the application of the care coordination safe harbor to AI-driven engagement optimization, and clarifying permissible AI-driven personalization of member communications, would remove the compliance uncertainty that slows adoption.
3. **Harmonizing state AI requirements for healthcare AI operating under federal programs.** An emerging patchwork of state AI laws, including Colorado's AI Act effective June 30, 2026, and pending legislation in California, Texas, Illinois, and other states, creates duplicative compliance obligations that disproportionately burden smaller innovators and delay beneficial AI deployment across state lines. Federal leadership establishing baseline standards that satisfy state-level requirements for AI operating within Medicaid managed care would provide the regulatory certainty needed for multi-state scaling.

Our comments and recommendations on specific areas of the RFI are set forth below.

I. Introduction: Waymark Background

Waymark is a healthcare technology company focused on value-based care delivery for Medicaid populations. We partner with managed care organizations to deploy AI-driven population health management systems that combine predictive analytics with community-based care delivery. Our platform serves rising risk, high-cost Medicaid members across Washington, Ohio, and Virginia, targeting reductions in avoidable acute care utilization through proactive interventions.

Waymark's AI systems are designed to augment–not replace–human clinical judgment and care delivery. Our platform supports community health workers, care coordinators, therapists and pharmacists by:

- **Prioritizing Caseloads**: AI identifies which members are at highest risk and most likely to benefit from intervention, enabling care teams to focus limited time on highest-value activities
- **Surfacing Relevant Information**: AI synthesizes claims, clinical and social determinants data to provide care team members with actionable insights without requiring manual chart review.
- **Optimizing Outreach Timing:** AI predicts optimal times and channels for member contact, improving engagement rates and reducing wasted outreach attempts
- **Automating Routine Tasks**: AI handles routine member communications, appointment reminders, and care gap notifications, freeing care team members for complex relationship based care.
- **Enabling Timely Escalations:** AI monitors member data for indicators of clinical deterioration, behavioral health crisis, or rising acute care risk, enabling care teams to escalate interventions before avoidable emergency department visits or hospitalizations occur.

Waymark's AI platform is designed to be consistent with the framework established in FDA's January 2026 final guidance on Clinical Decision Support software, which implements the Cures Act exclusion for Non-Device CDS under Section 520(o)(1)(E) of the FD&C Act. Several of Waymark's functions are operational, not clinical, and do not meet the device definition. Functions that do not surface clinical recommendations to healthcare professionals are designed to meet all four Non-Device CDS criteria. This "augment, not replace" design is central to Waymark's model.

This augmentation model addresses healthcare workforce shortages by enabling existing staff to serve more members effectively. Community health workers supported by AI can manage larger patient panels while maintaining personalized, relationship-based care. Rather than displacing healthcare workers, AI enables workforce multiplication that expands access to care coordination services.

II. Responses to Request for Public Comment

A. Regulation

We recommend that CMS issue guidance establishing voluntary best practices for validation and monitoring of AI systems used in Medicaid Managed Care that provide pre-deployment validation standards, ongoing monitoring requirements, transparency, and document human oversight requirements. We recommend that HHS establish clear standards for algorithmic fairness in clinical AI systems, including requirements for disparate impact testing, ongoing performance monitoring across demographic subgroups, and transparent reporting of model limitations.

While 42 CFR Part 438 permits MCOs to use AI for quality improvement and care coordination, there are no specific standards for validating AI system accuracy, reliability, or equity when used in Medicaid populations. AI systems used for risk stratification, care coordination, or utilization management in Medicaid should be validated using Medicaid specific data and assess performance metrics, including sensitivity and specificity. Systems should demonstrate that there are no statistically significant differences in sensitivity/specificity across demographic groups or have a documented plan to address identified disparities through algorithm refinement or targeted interventions.

MCOs using AI systems should implement continuous monitoring including performance tracking, algorithm drift detection, adverse event monitoring, and monitoring AI system performance across demographic subgroups.

In addition, MCOs should maintain and make available to the State documentation about their algorithm development and validation, vendor information, and reports about the performance of the algorithm.

AI systems used in Medicaid managed care should also include a clinical review process, quality oversight on AI system performance, the right for enrollees to decline AI interventions without penalty, and request information about how they were identified as rising risk.

These standards should be proportionate to the risk level of the AI application and avoid overly prescriptive technical requirements that would stifle innovation. Establishing industry-wide standards would create appropriate competitive safeguards and ensure baseline quality across all AI systems affecting patient care.

We recommend that HHS issue joint HHS guidance (OCR, CMS, SAMHSA, ONC) clarifying data sharing requirements for AI systems in Medicaid.

Medicaid MCOs, providers, and state agencies are uncertain about when and how they can share Medicaid enrollee data with AI vendors for:

- Algorithm training and development
- Risk stratification and prediction
- Care coordination and intervention deployment
- Quality improvement and program evaluation

HHS should clarify that State Medicaid agencies may share Medicaid claims, encounter, and enrollment data with AI vendors for the following purposes:

(1) Program Integrity and Quality Improvement

- Developing AI systems to detect fraud, waste, and abuse
- Creating predictive models for quality improvement initiatives
- Evaluating effectiveness of Medicaid programs and interventions

(2) Managed Care Oversight

- Validating MCO-reported data and performance metrics

- Assessing MCO compliance with contract requirements
- Supporting state-led delivery system reform initiatives

The guidance can include requirements for data sharing by states, which may include security requirements, restrictions on further disclosure without state approval, and compliance monitoring provisions.

We recommend that HHS update its 2012 de-identification guidance to address AI-driven re-identification risks and clarify permissible uses of de-identified and limited data sets for population health AI.

The HHS 2012 "Guidance Regarding Methods for De-identification of Protected Health Information" was issued before widespread deployment of artificial intelligence and machine learning technologies. Since then, AI techniques have demonstrated the ability to:

- Re-identify individuals from datasets previously considered de-identified
- Infer sensitive attributes from non-sensitive data patterns
- Link multiple de-identified datasets to enable identification
- Extract identifying information from unstructured text and images

This creates uncertainty for covered entities and business associates about:

- Whether data de-identified using older methods remains de-identified under current technology
- How to conduct Expert Determination analyses that account for AI re-identification risks
- When AI training and validation constitute permitted uses of de-identified or limited data sets
- What safeguards are necessary when sharing de-identified data with AI vendors

The guidance on the expert determination method (45 CFR 164.514(b)(1)) can be updated to address these areas of uncertainty.

Additionally, HHS should clarify permissible uses of limited data sets under 45 CFR 164.514(e) when used for AI purposes. For example, HHS can confirm that certain AI related activities, such as training AI systems for care coordination and case management, developing predictive models for quality improvement initiatives, and validating AI algorithm performance for population health management, constitute "healthcare operations" under 45 CFR 164.501 and may use limited data sets pursuant to data use agreements.

Any guidance document should take a risk based approach, acknowledging that safeguards should be proportionate to the risk of re-identification.

HHS took an important step in February 2024 when SAMHSA and OCR issued a final rule aligning 42 CFR Part 2 with HIPAA. The final rule allows Part 2 programs to obtain a single patient consent for all future uses and disclosures for treatment, payment, and health care operations; permits HIPAA covered entities and business associates that receive Part 2 records under this consent to redisclose those records in accordance with HIPAA; and applies the HIPAA de-identification standard under 45 CFR 164.514(b) to Part 2 records. With the compliance date of February 16, 2026 now upon us, this is an appropriate moment to address the remaining gaps that the final rule did not resolve, particularly with respect to AI-driven care coordination for substance use disorder populations. The recommendations below build on the foundation the 2024 final rule established.

We recommend that HHS issue guidance under 42 CFR Part 2 allowing AI-assisted substance use disorder treatment coordination.

The behavioral health crisis demands AI solutions that can identify individuals at risk, predict crisis events, and optimize intervention timing. Waymark's care teams include therapists delivering behavioral health interventions, and our AI systems support behavioral health care coordination. However, behavioral health AI faces unique barriers:

Data Fragmentation: Behavioral health data is often maintained in separate systems from general medical records, with limited integration into EHRs and claims databases. Mental health diagnoses may be under-coded due to stigma concerns, and substance use disorder data faces additional restrictions under 42 CFR Part 2.

Current 42 CFR Part 2 restrictions create barriers to integrating SUD patients into population health management programs, even when AI systems would improve care coordination. Given the breadth of Medicaid populations served by managed care organizations, clarifying that population-level risk stratification and care coordination prioritization can include SUD patients without requiring full Part 2 consent procedures would significantly improve care quality.

Crisis Prediction Challenges: AI systems predicting behavioral health crises (suicide risk, psychiatric hospitalization, substance use relapse) require access to sensitive data including prior mental health history, social determinants, and real-time indicators. Current data infrastructure and privacy frameworks limit development of effective crisis prediction models.

Integration with Community Resources: Effective behavioral health AI must connect predictions to actionable interventions, including crisis services, peer support, and community mental health resources. Interoperability between healthcare AI systems and community mental health infrastructure is limited.

To address the issues above, we recommend that HHS:

- Expand the 42 CFR Part 2 recommendations to address AI-assisted behavioral health care coordination more broadly, including crisis prediction and intervention optimization
- Fund development of behavioral health-specific AI benchmarking datasets that include longitudinal mental health outcomes, crisis events, and treatment response data with appropriate de-identification
- Support integration standards connecting healthcare AI systems with 988 Suicide and Crisis Lifeline, community mental health centers, and peer support networks
- Establish guidance on permissible AI applications for behavioral health risk prediction, including appropriate use of social media, wearable, and other non-traditional data sources

We recommend that HHS issue guidance clarifying: 1) the application of the care coordination safe harbor to AI-driven member engagement and outreach optimization within value-based arrangements, 2) permissible AI-driven personalization of member communications and incentives under the beneficiary inducement CMP, particularly for preventative care and care coordination activities, and 3) documentation and transparency requirements for AI systems operating within fraud and abuse safe harbors.

AI systems that personalize member outreach and engagement raise novel questions under existing fraud and abuse frameworks. Automated engagement systems, including SMS and voice calling for member outreach, may operate within care coordination arrangements. However, regulatory uncertainty exists regarding:

- Anti-Kickback Statute (42 U.S.C. § 1320a-7b(b)): When AI systems optimize member engagement strategies—including timing, channel, and messaging—questions arise about whether personalized outreach constitutes remuneration that could influence referrals or service utilization. While care coordination arrangements may qualify for the care coordination safe harbor (42 CFR § 1001.952(ee)), the application of this safe harbor to AI-driven engagement optimization has not been addressed in OIG guidance.
- Beneficiary Inducement Civil Monetary Penalty (42 U.S.C. § 1320a-7a(a)(5)): AI systems that personalize incentives, nudges, or engagement strategies for Medicaid beneficiaries must navigate the prohibition on offering remuneration to influence beneficiary selection of providers or services. The preventive care exception (42 CFR § 1003.110) may apply to certain AI-driven health promotion activities, but its scope for population health AI applications remains unclear.

Issuing guidance would reduce compliance uncertainty for organizations deploying AI to improve member engagement and care coordination outcomes.

We recommend that HHS: 1) coordinate with state regulators through the National Association of Insurance Commissioners (NAIC) and National Association of Medicaid Directors (NAMD) to develop harmonized AI governance standards for healthcare applications, 2) issue guidance establishing that compliance with federal AI transparency, validation, and monitoring standards satisfies state requirements for healthcare AI operating under federal programs, 3) consider whether ERISA-style preemption of inconsistent state AI requirements is appropriate for AI systems operating within federally-regulated healthcare programs, including Medicare Advantage and Medicaid managed care, and 4) establish a federal floor for healthcare AI governance that provides regulatory certainty while preserving state authority to address state-specific concerns.

Beyond variations in telemedicine and digital health regulations, an emerging patchwork of state AI laws creates significant compliance complexity for multi-state healthcare AI deployments. Colorado's AI Act (SB 24-205), effective June 30, 2026, imposes disclosure, impact assessment, and governance requirements on "high-risk" AI systems, which may include healthcare risk stratification and care coordination applications. Similar legislation is pending or under consideration in California, Texas, Illinois, and other states.

For organizations like Waymark operating across multiple state Medicaid programs, inconsistent state AI requirements create:

- Duplicative compliance obligations requiring state-specific impact assessments, disclosures, and governance documentation
- Conflicting technical requirements that may be impossible to satisfy simultaneously across jurisdictions
- Uncertainty about which state's law applies when AI systems process data from members in multiple states
- Increased costs that disproportionately burden smaller innovators and may delay beneficial AI deployment

Federal leadership on healthcare AI governance would reduce compliance fragmentation while ensuring consistent patient protections across state lines.

B. Reimbursement

For Medicaid specifically, we recommend that CMS encourage states to adopt managed care contract structures under 42 CFR Part 438 that include explicit performance-based payments for utilization management and health outcome improvements achieved through AI deployment.

Current payment models represent the single greatest barrier to AI adoption in clinical care. Fee-for-service reimbursement creates perverse incentives that discourage investment in AI systems that reduce unnecessary utilization, even when such reductions improve patient outcomes and lower total costs. Value-based payment models theoretically align incentives but face implementation challenges that limit AI adoption.

Payment Model Barriers:

- Fee-for-service billing codes do not accommodate AI-driven care coordination, preventive interventions, or population health management activities
- Medicaid managed care contracts often lack explicit incentive structures for AI-driven quality improvement and utilization management
- Limited budget authority for managed care organizations to invest in AI infrastructure when savings accrue to other entities in fragmented payment systems

CMS should develop and disseminate model contract language for Medicaid managed care agreements that states can adopt. Sample contract provisions should include performance based payments for demonstrable reductions in avoidable acute care utilization achieved through deployment of evidence-based AI systems for population health management. The health plan should be required to provide annual documentation of AI system validation, fairness testing, and performance monitoring to demonstrate evidence-based deployment.

Enhanced federal matching for states that adopt such value-based contract structures could accelerate this transition. CMS should consider increasing FMAP percentages for states implementing performance-based Medicaid managed care contracts that explicitly incentivize AI adoption for care quality improvement.

We recommend that CMS enhance risk adjustment methodologies in capitated payment arrangements to account for social determinants of health and rising risk trajectories, not just current diagnostic codes.

Existing high cost claimant-based risk adjustment creates gaming incentives around diagnosis capture rather than true health improvement. Risk adjustment that incorporates SDOH factors, utilization trajectories, and functional status would better align payment with actual population health needs and create stronger incentives for AI investments targeting rising-risk members before they develop costly complications.

We recommend that CMMI establish innovation waivers that allow alternative payment models for AI-driven care delivery outside traditional billing structures as well as establish a demonstration model testing AI-driven population health management.

The Center for Medicare and Medicaid Innovation should prioritize models that test population-based payments for AI-driven care coordination, with shared savings mechanisms that reward total cost of care reductions while maintaining or improving quality metrics. CMMI demonstration models would generate evidence for broader payment reform while creating pathways for AI-driven care delivery.

This model structure would create aligned incentives for AI investment while generating rigorous evidence on AI effectiveness in real-world care delivery. The demonstration model testing AI driven population health management should include the following parameters:

- Validation on diverse patient populations
- Ongoing performance monitoring with drift detection
- Disparate impact testing across demographic subgroups
- Transparency documentation meeting specified standards

C. Research & Development

While private sector innovation has produced impressive AI capabilities, specific research infrastructure investments would accelerate translation to real-world clinical settings.

We recommend that AHRQ invest in federated learning infrastructure for healthcare AI research, potentially through cooperative agreements with health systems and payers.

This infrastructure would enable multi-site model training and validation without requiring data sharing, addressing privacy concerns while improving model generalizability. Technical investments should include privacy-preserving computation methods, differential privacy implementations, and secure multi-party computation protocols. This approach would minimize breach risk while improving model performance across heterogeneous Medicaid populations. Current limitations in accessing multi-site data for model training restrict the ability to develop models that generalize across different states, health plan structures, and demographic compositions.

To achieve this infrastructure, specific technical capabilities needed include: 1) secure aggregation protocols that enable gradient sharing without exposing individual patient data; 2) differential privacy guarantees that formalize privacy protections while maintaining model accuracy; and 3) federated model architectures that accommodate heterogeneous data schemas across different EHR systems and claims databases. AHRQ should consider funding demonstration projects testing federated learning for population health AI applications including risk prediction, care gap identification, and intervention response modeling.

We recommend that ONC lead development of standardized SDOH data elements for population health AI.

Current pilot implementations of Gravity Project standards and Integrating the Healthcare Enterprise (IHE) SDOH profiles need widespread adoption to enable consistent capture and exchange of housing stability, food insecurity, transportation access, and social isolation data. SDOH standardization would improve predictive accuracy and enable better intervention targeting.

We recommend that FDA and CMS should establish real world evidence pathways for AI validation.

Traditional randomized controlled trials (RCTs) are often impractical for evaluating population health AI due to:

- Cluster-level interventions: AI systems typically operate at the organization or population level, requiring cluster randomization with associated sample size and cost implications
- Rapid iteration: AI systems improve continuously through retraining, making point-in-time RCT results quickly obsolete
- Ethical constraints: Withholding potentially beneficial AI interventions from control populations raises ethical concerns, particularly for high-risk members
- External validity: RCT conditions may not reflect real-world implementation contexts, limiting generalizability

FDA and CMS should establish real-world evidence (RWE) pathways for AI validation that:

- Accept quasi-experimental designs (difference-in-differences, regression discontinuity, instrumental variables) as valid evidence of AI effectiveness when RCTs are impractical
- Establish standards for observational study quality, including pre-registration, sensitivity analyses, and transparency requirements
- Create expedited review pathways for AI systems with strong RWE demonstrating safety and effectiveness
- Recognize continuous learning AI systems that demonstrate ongoing performance monitoring and improvement as meeting validation requirements without repeated point-in-time studies
- Align FDA device evaluation, CMS coverage determination, and quality measure development around common RWE standards

RWE pathways would accelerate evidence generation for AI systems while maintaining scientific rigor appropriate to the risk level of each application.

III. Responses to Specific Questions

Question 1: What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

The biggest barriers to private sector AI innovation in healthcare and its adoption in clinical care are 1) payment system misalignment, 2) regulatory uncertainty for non medical device AI applications, 3) data access, and 4) validation and evidence requirements for model development.

Payment System Misalignment

Fee-for-service payment creates negative ROI for AI systems that reduce unnecessary utilization. Organizations have no financial incentive to deploy AI that prevents hospitalizations or reduces testing when revenue depends on volume. Even in value-based arrangements, insufficient performance incentives and short time horizons limit investment in AI. Medicaid managed care contracts particularly lack explicit provisions rewarding AI-driven care management, creating misalignment between AI investment costs and potential savings.

Regulatory Uncertainty

Lack of clear regulatory pathways for non-medical device AI applications creates risk aversion. Organizations are uncertain whether their AI systems require FDA approval, what liability exposure they face, and what privacy safeguards apply when using de-identified data. This uncertainty delays deployment and increases compliance costs.

Specific areas of concern include what documentation standards apply for model validation and monitoring and how to manage liability when AI recommendations are mediated through human clinical judgment.

FDA's January 2026 final guidance on Clinical Decision Support software is a positive example of the regulatory clarity that accelerates innovation. By establishing clear criteria for Non-Device CDS under Section 520(o)(1)(E) of the FD&C Act and adopting a risk-proportionate enforcement discretion framework for lower-risk device CDS functions, the guidance gives developers a workable framework for determining whether their HCP-facing software requires premarket review. However, the CDS guidance does not address AI systems used for population-level risk stratification, care coordination prioritization, or member engagement optimization, functions that operate upstream of individual clinical decisions and do not fit neatly within the CDS framework. This is the specific regulatory gap where additional HHS guidance is most needed. We recommend that CMS and OIG adopt the same risk-proportionate approach FDA has taken in the CDS guidance when establishing safe harbors and best practices for population health AI operating within Medicaid managed care.

Data Access and Interoperability

Fragmented EHR systems, limited information exchange, and restricted access to social determinants data prevent development of comprehensive AI models. Claims data provides utilization history but lacks clinical detail. EHR data has clinical detail but limited standardization across systems. Social services data is siloed entirely. AI requires integrated longitudinal data that current infrastructure does not support. Multi-site model training faces particular challenges as data cannot be centralized due to privacy concerns.

Validation and Evidence Requirements

Limited access to diverse patient populations for model development and validation restricts generalizability. Randomized controlled trials are expensive and slow, while observational evaluation methodologies lack standardization. Payers and health systems demand evidence of effectiveness but are unwilling to support evaluation costs. Specific challenges include: establishing appropriate control groups for AI intervention evaluation; controlling for selection bias when AI deployment is not randomized; and attributing outcome changes to AI versus other concurrent quality improvement initiatives.

Question 2: What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.

Building on the single-consent framework for treatment, payment, and health care operations established in the 2024 Part 2 final rule, HHS should clarify the use of AI-generated outputs in substance use disorder treatment under 42 CFR Part 2. Specifically:

- HHS should confirm that substance use disorder patient records de-identified in accordance with 45 CFR 164.514(b) may be used for artificial intelligence model training, population health analytics, and care coordination without patient consent pursuant to 42 CFR 2.54.
- HHS should clarify that AI-generated outputs derived solely from de-identified data do not constitute Part 2 records unless they identify or provide a reasonable basis to identify specific patients.
- HHS should clarify that coordinating SUD treatment with AI systems can occur under a single patient consent. Specifically, HHS should clarify that artificial intelligence systems used for substance use disorder treatment coordination may access patient identifying information when a Part 2 program has obtained a single consent for treatment, payment, and healthcare operations

pursuant to 42 CFR 2.31, and that such AI systems are subject to the same use and disclosure restrictions as other recipients of Part 2 records.

- Finally, HHS should provide guidance on Business Associate Agreement requirements when AI vendors receive substance use disorder patient records from Part 2 programs that are also HIPAA covered entities, clarifying the interaction between Part 2 consent requirements and HIPAA business associate obligations.

We recommend that CMS issue guidance under 42 CFR Part 438 clarifying how states can structure Medicaid managed care capitation rates and performance-based payments to create financial alignment for MCO investments in AI-driven population health management that reduce total cost of care.

Current 42 CFR Part 438 regulations permit states to include technology costs in capitation rates and structure performance-based payments for cost reduction, but lack of specific guidance on AI, which creates uncertainty for states and MCOs about:

- How to project costs and savings from AI investments in actuarial rate certifications
- Optimal payment structures to incentivize AI adoption while maintaining actuarial soundness
- Performance metrics that capture AI-driven improvements in care coordination and prevention
- Multi-year investment frameworks when AI implementation costs precede savings realization

HHS could clarify that actuarially sound capitation rates under 42 CFR 438.4 may include 1) technology costs, 2) AI-driven intervention costs (e.g., CHW deployment based on AI risk stratification), 3) projected savings (e.g., from avoidable ER utilization or hospitalizations).

Question 3: For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (e.g., relating to liability, indemnification, privacy, and security). What novel legal and implementation challenges exist and what role, if any, should HHS play to help address them.

When AI systems provide care coordination recommendations that clinical staff follow, liability allocation between AI vendors, healthcare organizations, and individual clinicians is unclear. Standard professional liability insurance and vendor indemnification clauses were not designed for AI-mediated clinical judgment. Specific issues include: who bears liability when AI recommendations are technically accurate but lead to adverse outcomes; how to allocate responsibility when clinicians override AI recommendations that would have prevented adverse events; and what documentation standards apply to demonstrate appropriate use of AI in clinical decision-making. HHS should convene stakeholders to develop model liability frameworks that appropriately allocate risk while maintaining accountability.

Patients and clinicians reasonably expect to understand how AI systems make recommendations, yet vendors have legitimate intellectual property concerns about disclosing proprietary algorithms. Current transparency requirements under 21st Century Cures Act focus on interoperability rather than algorithmic explainability. Specific tensions include: what level of algorithm disclosure is necessary for informed consent; whether clinicians require access to model architecture and training data to appropriately utilize AI recommendations; and how to balance transparency needs with protection of competitive advantages. HHS should establish tiered transparency standards based on AI system risk level and use case, with higher-risk applications requiring more detailed disclosure.

AI systems that learn continuously from new data raise questions about when they constitute a new intervention requiring revalidation. Static models avoid this issue but sacrifice performance as patient populations and care delivery patterns evolve. Specific challenges include: what magnitude of model parameter changes triggers revalidation requirements; how to detect and manage performance degradation from distribution shift; and what governance processes should oversee continuous learning systems. HHS should establish frameworks for continuous learning AI that define acceptable drift thresholds, monitoring requirements, and revalidation triggers based on clinical risk.

Question 6: Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the

greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?

Where AI Has Met or Exceeded Expectations

Risk stratification models accurately predicting near-term acute care utilization have proven highly effective. For example, Waymark's operational deployments have generated measurable outcomes beyond model accuracy metrics. Across our Medicaid managed care partnerships in Washington and Virginia, AI-driven interventions have achieved:

- **Emergency department utilization**: 20.4% reduction in avoidable ED visits among members receiving AI-driven outreach compared to matched controls
- **Hospitalization rates**: 48.3% reduction in 30-day all-cause hospitalizations for high-risk members engaged through AI-driven care coordination
- **Total cost of care**: $3.73:1 ROI per-member-per-month savings in populations receiving AI-driven interventions, representing $253 in reduction in total medical expenditure per intervened patient per month
- **Member engagement**: 41.2% contact rate for AI-driven outreach compared to 10-30% for traditional care management approaches
- **Quality measure performance**: 11.8 percentage point improvement in HEDIS gap closure rates using AI-driven intervention prioritization[1]

These outcomes demonstrate that population health AI can deliver measurable value when deployed within value-based payment arrangements that align incentives for utilization reduction and outcome improvement. We are prepared to share detailed methodology and results with HHS upon request.

Automated engagement systems using SMS and voice calling have exceeded expectations for member reach and engagement rates. AI-driven messaging personalization improves response rates by 40-60% compared to generic outreach,[2] enabling scalable contact with high-risk members who historically had low engagement with traditional care management. This automation allows care coordinators to focus time on high-value in-person interventions rather than repetitive outreach attempts.

HEDIS gap closure prioritization using machine learning to identify members most likely to respond to specific interventions has improved quality measure performance cost-effectively. Rather than universal outreach, AI-driven interventions achieve similar completion rates with 30-40% fewer attempts, generating substantial savings in care management labor costs while improving quality measure performance.[3]

Where AI Has Fallen Short

Population health has not yet delivered consistently in two areas.

- Complex reinforcement learning systems for sequential care coordination decisions have shown promise in research settings but face implementation challenges. Model interpretability, clinical trust, and liability concerns limit adoption of black-box sequential decision support. Clinicians want transparent rationales for AI recommendations, which current reinforcement learning approaches struggle to provide. Specific challenges include explaining why particular outreach timing or intervention sequences are recommended, and communicating uncertainty bounds around predicted outcomes.
- Health related social needs prediction suffer from limited data availability and quality. While claims and EHR data support clinical outcome prediction, predicting social service needs requires data that is fragmented, unstandardized, and often unavailable. AI alone cannot compensate for fundamental data infrastructure gaps. Current SDOH screening tools capture point-in-time

---

1 Baum A, Batniji R, Ratcliffe H, DeGosztonyi M, Basu S. Supporting Rising-Risk Medicaid Patients Through Early Intervention. NEJM Catalyst Innovations in Care Delivery. 2024;5(11). DOI: 10.1056/CAT.24.0060.

2 Chacko J, Mazza K, Stathakos K, Kim D, Carlo L. Evaluating Technology-Driven Strategies for Enhancing Patient Outreach for Preventive Care: A Systematic Review. Cureus. 2025 Feb 22;17(2):e79467. doi: 10.7759/cureus.79467. PMID: 40130097; PMCID: PMC11932718.

3 Results from internal Waymark outreach campaigns.

assessments but lack longitudinal tracking of housing, food security, and transportation stability over time.

Novel AI Tools with Greatest Potential

- Causal AI systems that predict individual treatment effects rather than just population-level predictions, enabling truly personalized intervention selection based on expected benefit for each patient
- Federated learning systems enabling multi-site model training without data sharing, improving generalizability while maintaining privacy
- Multi-modal AI integrating claims, EHR, wearable, and social service data for comprehensive risk assessment that captures clinical, behavioral, and social determinants simultaneously
- Natural language processing extracting structured data from unstructured clinical notes, particularly for mental health symptoms, functional status, and patient-reported social needs

Question 7: Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?

Most Influential Roles for AI Adoption:

- Chief Medical Officers and Chief Quality Officers control clinical program priorities and have authority to mandate AI tool adoption for quality improvement initiatives. Their clinical credibility is essential for building physician trust in AI recommendations.
- Payer Medical Directors for managed care organizations influence coverage and reimbursement decisions affecting AI deployment at provider organizations. They control quality measure selection, care management program requirements, and performance-based payment structures that drive AI adoption incentives.
- Chief Financial Officers determine ROI requirements and budget availability, requiring clear evidence that AI investments will generate financial returns within acceptable time horizons, typically 2-3 years for managed care organizations.

Primary Administrative Hurdles:

- Procurement processes requiring lengthy vendor evaluation, security reviews, and contract negotiations that delay deployment by 12-18 months, particularly for health systems and state Medicaid agencies with extensive approval requirements
- Information security requirements including comprehensive risk assessments, penetration testing, and HIPAA compliance documentation that extend vendor onboarding, with particular challenges around cloud-based AI systems and data residency requirements
- Clinical governance processes requiring physician champion identification, committee approvals, and policy development before deployment, with resistance from clinicians concerned about AI replacing human judgment or increasing documentation burden
- Training and change management needs for clinical and administrative staff who must learn new workflows and technologies, with particular challenges when AI recommendations conflict with established care protocols or clinical intuition

Question 8: Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

Enhanced interoperability would significantly accelerate AI development in care coordination activities and medication adherence.

Non-clinical interventions by community health workers, care managers, and social service providers lack standardized representation in health IT systems. FHIR CarePlan and ServiceRequest resources need expansion to capture community-based interventions, member engagement activities, and social service referrals. Interoperable care coordination data would enable AI to learn from diverse intervention approaches across organizations. Current lack of standardization prevents training models on what interventions work for which patients in which contexts, forcing each organization to build models from scratch using only their own limited experience.

Real-time medication adherence data requires integration of pharmacy claims, prescription refill data, and potentially pill bottle sensors. NCPDP SCRIPT standard supports e-prescribing but lacks adherence monitoring capabilities. Integration of Surescripts medication history with EHR data would improve AI-driven adherence intervention targeting by identifying non-adherence patterns before clinical deterioration occurs.

Standardized SDOH data elements are important for population health AI long-term. Current pilot implementations of Gravity Project standards and IHE SDOH profiles need widespread adoption. Priority data elements include housing stability, food insecurity, transportation access, and social isolation. However, this represents a multi-year standardization effort with gradual benefits as adoption increases across care settings. Standards should support both structured data capture and natural language processing of unstructured SDOH assessments from community health workers.

In addition, improved data integration for dual-eligible populations would enable AI systems to address the full complexity of care needs for this vulnerable population.

Approximately 12 million Americans are dually eligible for Medicare and Medicaid, representing a disproportionate share of high-risk, high-cost members who would benefit most from AI-driven population health management. However, AI development for this population faces significant data integration barriers:

- Fragmented data sources: Medicare claims data (through CMS) and Medicaid claims data (through state agencies) are maintained in separate systems with different access mechanisms, data formats, and use restrictions
- Inconsistent identifiers: Lack of universal patient identifiers complicates linking Medicare and Medicaid records for the same beneficiary
- Regulatory complexity: Different privacy frameworks and data use agreements govern Medicare versus Medicaid data, creating uncertainty about permissible AI applications using integrated data
- D-SNP limitations: Dual Eligible Special Needs Plans have access to both Medicare and Medicaid data but face restrictions on data sharing with AI vendors and care coordination partners

To address this fragmentation, CMS should:

- Establish streamlined data use agreements enabling AI developers to access linked Medicare-Medicaid data for population health applications serving dual-eligible beneficiaries
- Develop common data models and standardized file formats for integrated Medicare-Medicaid datasets suitable for AI model development
- Issue guidance clarifying HIPAA and Medicaid confidentiality requirements for AI systems operating on integrated dual-eligible data
- Require Medicare-Medicaid Plans (MMPs) and D-SNPs to share integrated data with contracted care coordination vendors, including AI-driven population health management partners, subject to appropriate privacy protections

Question 9: What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?

Challenges Patients Want AI to Address:

- Care coordination complexity and fragmentation across multiple providers, specialists, and service organizations, particularly for patients with multiple chronic conditions
- Difficulty accessing appropriate care at the right time, particularly for behavioral health and specialist services where wait times can exceed months
- Medication management challenges including polypharmacy, side effects, and adherence barriers related to cost, understanding, or complexity of regimens

- Navigation of social service resources including food assistance, housing support, and transportation, where eligibility requirements and application processes create substantial barriers

Patient and Caregiver Concerns About AI:

- Privacy and data security concerns about who has access to health information and how it is used, particularly with cloud-based AI systems and third-party vendors
- Bias and fairness worries that AI systems might perpetuate existing disparities or discriminate based on race, socioeconomic status, disability, or other protected characteristics
- Loss of human connection and empathy in healthcare delivery, particularly for patients with complex needs requiring relationship-based care and trust-building over time
- Lack of transparency about when AI is being used and how it influences clinical decisions affecting their care, with concerns about whether clinicians are being pressured to follow AI recommendations even when clinical judgment differs

Balancing these opportunities and concerns requires patient engagement in AI governance, transparent communication about AI use, and maintaining human oversight of AI-generated recommendations. Patients generally support AI that augments rather than replaces human clinical judgment and care delivery, and value AI that helps them navigate complex healthcare and social service systems more effectively.

Question 10: Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care? a) Are there published findings about the impact of adopted AI tools and their use [in] clinical care? b) How does the literature approach the costs, benefits, and transfers of using AI as part of clinical care?

HHS should prioritize research in the following areas that directly support operational AI deployment and minimize implementation barriers: 1) federated learning and privacy-preserving computation, 2) causal inference and heterogeneous treatment effects, and 3) safety-shared reinforcement learning.

Federated Learning and Privacy-Preserving Computation:

Multi-site model training without centralized data aggregation is essential for improving AI generalizability while maintaining privacy. Research priorities include federated learning algorithms robust to heterogeneous data distributions across health plans and care settings; differential privacy implementations with acceptable privacy-utility tradeoffs for population health applications; and secure multi-party computation protocols scalable to healthcare data volumes. This research directly supports better model development while minimizing breach risk.

Causal Inference and Heterogeneous Treatment Effects:

Most current healthcare AI focuses on prediction rather than causal effect estimation. Research is needed on methods for estimating individual-level treatment effects from observational data, including instrumental variable approaches adapted for healthcare settings, double machine learning implementations, and causal forest methods. Priority applications include personalizing intervention selection based on predicted benefit for each patient rather than population-average effects.

Safety-Shaped Reinforcement Learning:

Sequential clinical decision support requires reinforcement learning approaches that explicitly incorporate safety constraints. Research is needed on conservative policy learning that avoids recommending interventions with insufficient evidence; offline reinforcement learning with distributional robustness for learning from historical data without risky online exploration; and hybrid approaches combining model-based and model-free learning. Applications include medication titration, chronic disease management, and care coordination sequencing.

Based on our review of the current literature, published literature on deployed AI in clinical care remains limited compared to research prototypes. A critical gap exists between algorithm performance in research settings versus operational effectiveness in real care delivery. Key findings include: sepsis prediction models show AUC values of 0.75-0.85 but implementation studies reveal high false positive rates leading to alert fatigue; hospital readmission prediction models achieve modest discrimination but interventions targeting high-risk patients show variable effectiveness; and natural language processing for clinical documentation demonstrates time savings but accuracy and liability concerns persist.[4] More implementation science research is needed documenting how AI performs when integrated into actual clinical workflows with all their complexity and constraints.

IV. Conclusion

Accelerating AI adoption in clinical care requires focused action on the highest-impact barriers. Based on Waymark's operational experience deploying AI-driven population health management for high-risk Medicaid populations, we have prioritized concrete policy recommendations that would most directly enable responsible innovation while protecting patient safety and privacy.

Three policy changes represent the highest priorities. First, CMS should issue guidance under 42 CFR Part 438 establishing performance-based payment structures in Medicaid managed care contracts that explicitly reward AI-driven reductions in avoidable acute care utilization, accompanied by model contract language for states to adopt and enhanced federal matching for states that implement these value-based structures. Payment misalignment remains the single greatest barrier to AI adoption, and addressing it would unlock investment in AI systems that have already demonstrated measurable improvements in outcomes and cost.

Second, HHS should establish comprehensive regulatory safe harbors for population health AI, including OIG guidance confirming the application of existing fraud and abuse safe harbors to AI-driven member engagement and care coordination optimization. Organizations deploying AI to improve care for Medicaid beneficiaries should not face unresolved legal exposure under the Anti-Kickback Statute and Beneficiary Inducement CMP for activities that demonstrably improve health outcomes.

Third, HHS should lead federal harmonization of state AI requirements for healthcare applications operating under federally regulated programs. Coordinating with NAIC and NAMD to develop common standards, and establishing that compliance with federal AI transparency, validation, and monitoring requirements satisfies state-level obligations for Medicaid AI systems, would eliminate the duplicative compliance burden that currently impedes multi-state deployment by organizations like Waymark.

On the research front, investment in federated learning infrastructure represents the highest priority, as it directly supports model development across diverse populations while minimizing breach risk. This would enable AI systems to generalize across heterogeneous care settings without requiring centralized data aggregation that creates privacy concerns.

We appreciate HHS's proactive approach to AI policy development and are prepared to contribute our operational experience and research expertise to ongoing policy development efforts. The opportunity to deploy AI for improved health outcomes, reduced costs, and more equitable care delivery is substantial, and focused policy action on these high-priority areas can accelerate responsible innovation while maintaining appropriate safeguards.

Respectfully submitted,

Rajaie Batniji
CEO, Waymark

[4] Wong A, Otles E, Donnelly JP, Krumm A, McCullough J, DeTroyer-Cooley O, Pestrue J, Phillips M, Konye J, Penoza C, Ghous M, Singh K. External Validation of a Widely Implemented Proprietary Sepsis Prediction Model in Hospitalized Patients. JAMA Intern Med. 2021 Aug 1;181(8):1065-1070. doi: 10.1001/jamainternmed.2021.2626. Erratum in: JAMA Intern Med. 2021 Aug 1;181(8):1144. doi: 10.1001/jamainternmed.2021.3907. PMID: 34152373; PMCID: PMC8218233; Kansagara D, Englander H, Salanitro A, Kagen D, Theobald C, Freeman M, Kripalani S. Risk prediction models for hospital readmission: a systematic review. JAMA. 2011 Oct 19;306(15):1688-98. doi: 10.1001/jama.2011.1515. PMID: 22009101; PMCID: PMC3603349; Aramaki E, Wakamiya S, Yada S, Nakamura Y. Natural Language Processing: from Bedside to Everywhere. Yearb Med Inform. 2022 Aug;31(1):243-253. doi: 10.1055/s-0042-1742510. Epub 2022 Jun 2. PMID: 35654422; PMCID: PMC9719781.