February 22nd, 2026

Office of the National Coordinator for Health Information Technology
Assistant Secretary for Technology Policy
Department of Health and Human Services
Mary E. Switzer Building, Mail Stop: 7033A
330 C Street SW
Washington, DC 20201

Re: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care (RIN 0955-AA13; Docket No. HHS-ONC-2026-0001)

Dear Deputy Secretary O'Neill and ASTP/ONC Leadership,

I appreciate the opportunity to respond to the Department's Request for Information on accelerating AI adoption in clinical care. I write as a practicing physician, clinical researcher, and the founder of an AI-powered clinical reasoning platform, a combination that gives me a perspective I believe is underrepresented in policy conversations about health AI.

I am a cardiology fellow in Southeast Michigan who has spent the past decade inside the clinical training pipeline, from medical school through internal medicine residency to subspecialty fellowship. I have passed board examinations with the American Board of Internal Medicine and Certification Board of Nuclear Cardiology, and continue to train full-time while simultaneously building and shipping an AI clinical reasoning platform. My comments below draw on both my experience as a clinician navigating the healthcare system's inefficiencies and my experience as a technologist building AI tools to address them.

I. The Clinical Reasoning Gap as a Root Cause of Healthcare Waste

The Department's RFI asks where AI tools would have the greatest potential to improve healthcare outcomes and reduce costs (Question 5). I want to draw attention to a root cause that is rarely discussed in AI policy: the systematic failure in teaching clinical reasoning frameworks in medical education, and its direct downstream impact on clinical waste.

During my internal medicine residency, I led quality improvement research examining echocardiogram utilization among medical residents for syncope evaluation at our institution. The findings were striking: approximately 13% of echocardiogram orders for syncope we reviewed met ACC/AHA guideline criteria. The remaining 87% represented unnecessary testing, studies that added cost, consumed sonographer time, delayed care for patients with genuine indications, and occasionally produced incidental findings that triggered further unnecessary workups.

This was not an isolated finding. Our review of the literature indicated the pattern was systemic across U.S. hospitals. The root cause was not physician laziness or incompetence. It was the absence of structured clinical reasoning frameworks at the point of care. Physicians and advance practice practitioners order echocardiograms for syncope because they were never systematically trained to reason through when an echocardiogram is and is not indicated. They default to pattern-matching and predefined EMR order sets ("syncope → cardiac workup") rather than applying guideline-based reasoning.

This example generalizes across specialties. Diagnostic errors remain the third leading cause of death in the United States. Unnecessary testing drives an estimated $200 billion in annual waste. The common thread is a gap in clinical reasoning: the cognitive process by which physicians synthesize patient data, generate differential diagnoses, and make evidence-based decisions. This gap begins in medical education and persists throughout a physician's career.

II. AI-Powered Clinical Reasoning: From Education to Decision Support

I've been building Prosivo (https://prosivo.ai) to address this gap. The platform uses large language models to guide medical students and practicing clinicians through realistic patient cases via Socratic dialogue, the same method an expert attending physician would use at the bedside. Rather than presenting static multiple-choice questions, it adapts to the learner's level, tracks their evolving differential diagnosis, and teaches structured clinical reasoning frameworks in real time.

The architecture is intentionally designed as a two-stage pipeline:

1. **Clinical Reasoning Education (live):** An AI-powered Socratic tutor that works across training levels (medical student through attending) and across health professions (MD, nursing, PA). This product requires no FDA clearance, no EHR integration, and has short institutional sales cycles.
2. **Clinical Decision Support (in development):** The same reasoning engine, deployed at the point of care, guiding clinicians through guideline-based decision frameworks in real time. This product will require regulatory clarity, EHR interoperability, and reimbursement pathways — precisely the levers this RFI addresses.

The education product is not a stepping stone. It is the enabling infrastructure. It builds the clinical reasoning dataset, validates the AI's ability to guide clinical thinking, and establishes institutional relationships that make clinical deployment possible. This architecture has direct implications for all three of the Department's focus areas.

III. Responses to Specific Questions

A. Barriers to Private Sector Innovation (Question 1)

As a physician who has built and shipped an AI clinical tool while practicing full-time, I have experienced the barriers firsthand:

- **Regulatory ambiguity for non-device AI:** The current framework draws a binary distinction between “medical device” and “not a medical device.” Clinical reasoning tools, like AI that helps clinicians think through differential diagnoses and guideline application, occupy an uncertain middle ground. They are not diagnostic devices (they do not produce a diagnosis), but they are more clinically consequential than administrative tools. The lack of a clear regulatory category for “clinical reasoning support” creates uncertainty that chills investment and slows development.
- **The physician-founder gap:** The people best positioned to identify clinical AI opportunities, such as practicing physicians, face the highest barriers to building solutions. Clinical training demands 60–80+ hour weeks with no flexibility for entrepreneurship. There is no NIH mechanism, SBIR pathway, or HHS program specifically designed to support physician-technologists who want to build clinical AI tools while maintaining their practice. The result is that most health AI companies are founded by technologists who have never practiced medicine, leading to tools that optimize for metrics that don’t map to clinical workflow.
- **Institutional procurement inertia:** Medical schools and health systems lack procurement frameworks for AI-powered education and clinical reasoning tools. Traditional purchasing categories (textbooks, simulation equipment, EHR modules) don’t accommodate AI software that sits between education and clinical decision support. HHS could accelerate adoption by establishing model procurement guidelines for clinical AI tools.

B. Evaluation Methods for Non-Device AI (Question 3)

For AI clinical reasoning tools that do not qualify as medical devices, I recommend an evaluation framework with three components:

3. **Clinical accuracy benchmarking:** Does the AI’s reasoning align with published clinical guidelines (ACC/AHA, ACS, ACEP, etc.)? This can be measured through structured case libraries with known correct reasoning pathways, similar to how board examinations are validated.
4. **Learner outcome measurement:** For education-stage AI tools, does exposure to the tool improve clinical reasoning performance on standardized assessments? This is measurable through pre/post testing, clinical skills examination scores, and longitudinal tracking of diagnostic accuracy in practice.
5. **Workflow integration assessment:** For clinical-stage tools, does the AI reduce unnecessary testing, improve guideline adherence, and integrate without disrupting clinical workflow? These metrics are measurable through claims data, order pattern analysis, and time-motion studies.

HHS should consider supporting the development of standardized clinical reasoning benchmarks, analogous to how USMLE validates medical knowledge, that AI tools can be evaluated against. This would give health systems, payers, and regulators a common framework for assessing clinical reasoning AI without requiring device-level regulatory oversight for non-device tools.

C. Payment Policy and Reimbursement (Question 2)

The current fee-for-service model creates a paradox for clinical reasoning AI: the tools that reduce unnecessary testing directly reduce the revenue that funds the institutions using them. A physician who uses an AI-guided reasoning framework to determine that a syncope patient does not need an echocardiogram generates less revenue than one who orders the test reflexively.

I recommend HHS consider the following reimbursement approaches:

- **Value-based incentives for guideline adherence:** CMS quality programs (MIPS, APMs) should include specific measures for AI-assisted guideline adherence, creating positive financial incentives for health systems that deploy clinical reasoning AI tools.
- **Reimbursement codes for AI-augmented clinical reasoning:** Just as CPT codes exist for clinical decision support consultations, HHS should explore coding pathways for AI-assisted clinical reasoning at the point of care — recognizing the cognitive work of applying guideline-based frameworks rather than simply ordering additional tests.
- **Training reimbursement pathways:** GME funding mechanisms should recognize AI-powered clinical reasoning training as a legitimate educational expense. If an AI tutor demonstrably improves residents’ diagnostic reasoning, the institutions deploying it should be able to account for this within their GME cost reporting.

D. High-Impact Areas for AI Research Investment (Question 5)

Based on my clinical research and product development experience, I believe the highest-impact areas for HHS AI research investment are:

6. **Clinical reasoning education at scale:** AI that teaches the next generation of physicians to reason through clinical problems, and not to just memorize facts, has multiplicative downstream effects. Every physician who learns structured clinical reasoning through an AI tutor carries that improved decision-making framework into decades of practice. **The return on investment is measured not in individual encounters but in career-long practice patterns.**
7. **Guideline adherence at point of care:** My echo/syncope research demonstrates that guideline non-adherence is pervasive and costly. AI tools that integrate clinical practice guidelines into the reasoning process at the point of ordering, not as alerts that clinicians dismiss, but as reasoning frameworks that clinicians engage with, represent a high-value R&D target.

8. **Cross-profession clinical reasoning infrastructure:** Clinical reasoning is not unique to physicians. Nurses, advanced practice providers, and other health professionals all make clinical decisions that benefit from structured reasoning support. AI architectures that are content-agnostic, where the reasoning engine generalizes across professions, offer dramatically better economics than profession-specific tools. HHS should invest in platforms, not point solutions.

E. Interoperability and Market Opportunities (Question 7)

Enhanced interoperability would significantly accelerate clinical reasoning AI. Specifically:

- **Standardized learner performance data:** There is no interoperability standard for clinical reasoning performance data. If a medical student uses an AI reasoning tutor in school and later becomes a resident using an AI decision support tool, there is no way to connect those datasets. A FHIR-based standard for clinical reasoning competency data would enable longitudinal tracking from education through practice.
- **EHR-embedded clinical reasoning APIs:** Current EHR systems (Epic, Cerner/Oracle) offer clinical decision support hooks, but they are optimized for simple alert-based interventions (“drug interaction detected”), not for the kind of multi-step clinical reasoning dialogue that mirrors how physicians actually think. HHS should encourage EHR vendors to develop APIs that support richer AI-clinician interaction at the point of care.

IV. A Note on the Physician-Technologist Perspective

I want to close with a broader observation. The most impactful clinical AI tools will be built by people who have practiced medicine. Not because technologists lack talent, but because clinical reasoning is a domain where the problem definition requires lived experience. You cannot design a tool to improve how physicians think through differential diagnoses unless you have personally struggled with that process, unless you have been the medical student who scored in the 85th percentile on board examinations but froze at the bedside, the resident who witnessed colleagues ordering unnecessary tests because they lacked access to reasoning frameworks, and the fellow who decided to build the tool rather than wait for someone else to.

HHS has an opportunity to not only accelerate AI adoption but to cultivate the physician-innovators who will build the most impactful tools. Programs that reduce the barriers for practicing clinicians to develop and deploy clinical AI, whether through dedicated funding mechanisms, regulatory sandboxes, or partnership frameworks with academic medical centers, would yield outsized returns relative to their cost.

V. Conclusion

The Department’s RFI arrives at a critical inflection point. Large language models have reached the capability threshold to conduct meaningful interactive clinical dialogue. The technology to teach clinical reasoning at scale and to support it at the point of care now exists. But the regulatory, reimbursement, and institutional frameworks have not caught up.

I urge the Department to:

1. **Establish a clear regulatory category for non-device clinical reasoning AI** that provides proportionate oversight without imposing device-level requirements on tools that augment clinical thinking.

2. **Create reimbursement pathways** that reward guideline adherence and clinical reasoning quality, rather than volume of tests and procedures.

3. **Invest in clinical reasoning education infrastructure** as a high-leverage R&D priority, recognizing that improved reasoning at the training stage prevents waste at the practice stage.

4. **Develop programs that support physician-technologists in building clinical AI** from inside the healthcare system, not just from outside it.

I would welcome the opportunity to discuss any of these recommendations in further detail. The intersection of clinical reasoning education and AI-powered decision support represents one of the highest-impact opportunities to reduce healthcare costs while improving patient outcomes, and it is an area where the United States can and should lead.

Respectfully submitted,

Sam Alsabti, D.O.

Cardiology Fellow
Board Certified, American Board of Internal Medicine
Founder, Prosivo (prosivo.ai)
Southeast Michigan

Submitted via *regulations.gov* in response to Docket No. HHS-ONC-2026-0001