

February 22, 2026

Department of Health and Human Services
Assistant Secretary for Technology Policy
Office of the National Coordinator for Health Information Technology
Attention: HHS Health Sector AI RFI
Mary E. Switzer Building, Mail Stop: 7033A
330 C Street SW, Washington, DC 20201

Re: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care (RIN 0955–AA13)

Dear Assistant Secretary Keane:

The People’s Health Clinic (PHC) appreciates the opportunity to respond to HHS’s Request for Information on accelerating the adoption and use of artificial intelligence in clinical care. We write as a 30-person free clinic in Park City, Utah delivering 21,500 patient encounters annually to a predominantly Spanish-speaking population, 95% of whom live below 200% of the federal poverty level. The number of unique patients served is increasing by approximately 10% year over year. We do not receive federal funding. We have no IT department, no compliance team, and no dedicated cybersecurity staff.

Despite these constraints, PHC has direct operational experience with both categories of AI now entering clinical care: **predictive AI** that we have successfully deployed, and **generative AI** for which we have built a governance framework and are preparing to evaluate. Our experience illustrates both the promise and the practical barriers that safety-net providers face in adopting AI—and suggests specific actions HHS can take to accelerate responsible adoption.

I. Where AI Has Met or Exceeded Expectations: Autonomous Diabetic Retinopathy Screening (RFI Question 6)

PHC has deployed LumineticsCore® (Digital Diagnostics), the first FDA De Novo-cleared autonomous AI diagnostic system, which detects diabetic retinopathy at the point of care without requiring specialist overread. This is a predictive AI: the same retinal images produce the same diagnostic output every time, and its performance has been validated through rigorous clinical trials with published bias mitigation data across sex, race, and ethnicity.

Results at PHC (August 2025 – January 2026):

- 175 total exams completed; 154 sufficient for diagnosis (88% diagnosability)
- 20% positivity rate—one in five patients screened had more than mild diabetic retinopathy detected, enabling immediate referral to eye care
- Median exam time: 5 minutes
- Zero protocol errors

- 100% camera utilization

For PHC’s patient population—predominantly uninsured and with limited access to specialty care—this AI tool closes a critical gap. Prior to implementing LumineticsCore, diabetic retinopathy screening required a separate specialist referral, often resulting in delays of up to a year. Many patients were unable to complete screening due to cost or transportation barriers.

Now, screening happens during the primary care visit. The 20% positivity rate confirms that this tool is catching disease that would otherwise go undetected.

This deployment demonstrates that autonomous predictive AI can work effectively in a resource-constrained, safety-net setting with minimal infrastructure. The tool integrates into existing workflow, requires no IT support beyond initial setup, and produces deterministic results that can be validated using standard statistical methods.

II. The Generative AI Challenge: Why Current Frameworks Are Insufficient (RFI Questions 4 and 5)

PHC is also preparing to evaluate generative AI tools—specifically ambient AI clinical documentation and AI-powered phone answering and triage. These tools present fundamentally different challenges from the predictive AI described above, and this distinction has significant implications for HHS policy.

Predictive AI (like LumineticsCore) is deterministic: the same input produces the same output every time. It can be validated once using standard statistical methods (sensitivity, specificity, AUC) and monitored for drift using established techniques. Generative AI is non-deterministic: the same patient encounter processed twice will produce two different clinical notes with different wording, different emphasis, and potentially different clinical facts. This means:

- Traditional accuracy metrics do not apply. There is no single “correct” output to measure against.
- Hallucination—the generation of clinically plausible but fabricated content—is an inherent failure mode, not an edge case.
- Validation cannot be a one-time event. Every output is unique and must be evaluated continuously.
- Bilingual performance introduces additional risk. A note generated from a Spanish-language encounter may omit clinical facts that would be captured in English, creating a two-tier documentation standard that worsens health disparities.

To address these challenges, PHC developed a generative AI evaluation framework that measures four dimensions: clinical completeness, clinical fidelity (with zero tolerance for hallucinated content), semantic consistency across repeated runs, and harm potential scoring. The framework includes a repeatability test protocol that processes the same encounter five times,

compares the clinical facts extracted across all runs, and requires safety-critical facts to appear in at least four of five outputs. It also includes bilingual stress testing with code-switching and dialect-specific medical terminology.

To our knowledge, no published national framework—including the CHAI/Joint Commission Responsible Use of AI guidance, NACHC’s AI action guides, or the HEAAL equity framework—addresses the specific evaluation challenges of generative AI in clinical settings. This is a critical gap. As ambient AI documentation tools proliferate across the healthcare system, the absence of evaluation methods designed for non-deterministic outputs creates patient safety risk at scale.

This gap is not limited to safety-net settings. A 2024 case study from Mass General Brigham—an integrated academic health system with a dedicated 18-person, multidisciplinary AI governance committee—selected ambient documentation as their generative AI test case and acknowledged the absence of straightforward metrics for evaluating generative AI performance in clinical settings (Saenz et al., “Establishing responsible use of AI guidelines,” *npj Digital Medicine* 7:348, 2024). Their evaluation relied primarily on clinician feedback scores and note retention rates. They also reported an inability to measure AI quality across patient demographics due to data linkage limitations. If a major academic health system with dedicated AI governance resources faces these measurement challenges, the need for standardized evaluation methods that smaller organizations can adopt is urgent.

III. Barriers to AI Adoption in Safety-Net Settings (RFI Question 1)

PHC’s experience highlights several barriers that are likely shared by free clinics, community health centers, and other resource-constrained providers across the country:

No governance templates for small organizations. The CHAI/Joint Commission guidance recommends establishing formal AI governance structures with representatives from compliance, IT, clinical programs, operations, and data privacy. PHC does not have separate departments for any of these functions. We built a governance framework that assigns these responsibilities to specific individuals within a 30-person team. HHS should support the development of governance templates scaled to organizations of different sizes—not just large health systems.

Vendor opacity on generative AI performance. When PHC evaluates AI vendors, we ask: which large language model powers your product? Has the vendor run repeatability or consistency testing? What is the performance data for non-English and low-income populations? Most vendors cannot answer these questions. HHS should consider requiring transparency in training data composition, known limitations, and performance benchmarks stratified by language and demographics—particularly for tools marketed to safety-net providers.

Cost barriers. Monthly licensing costs for AI-powered clinical documentation tools can represent a significant budget line for a free clinic. PHC cannot currently absorb per-provider-per-month pricing designed for large health systems. HHS should explore mechanisms—grants,

subsidies, or tiered pricing incentives—to ensure that the productivity benefits of clinical AI are not limited to well-resourced organizations.

Default-on vendor features. EHR vendors are beginning to activate AI features by default. PHC’s framework requires that any vendor-released AI feature pass the same three-gate evaluation process as any other tool—but this only works if the clinic is aware the feature has been activated. HHS should consider requiring EHR vendors to provide advance notice and opt-in consent before activating AI features that process patient data.

IV. Recommendations for HHS Action (RFI Questions 2, 4, and 5)

Based on our direct experience deploying predictive AI and building governance for generative AI, PHC offers the following recommendations:

1. Develop evaluation methods specific to generative AI in clinical settings. Current evaluation frameworks are designed for predictive models. HHS should fund the development of standardized evaluation methods for generative AI—including repeatability testing, hallucination detection protocols, and bilingual performance benchmarking—and make these methods freely available to providers of all sizes.

2. Require vendor transparency on AI model composition and performance. Providers cannot make informed adoption decisions without knowing which AI models power a product, how those models were trained, and how they perform across different languages and patient populations. HHS should establish minimum disclosure requirements for AI vendors marketing to healthcare organizations.

3. Support AI governance capacity in safety-net settings. The CHAI/Joint Commission playbooks under development should include specific guidance for free clinics, community health centers, and critical access hospitals. PHC would welcome the opportunity to contribute our framework and operational experience to this effort.

4. Protect patient trust through transparency requirements. For AI tools that record patient-provider conversations—particularly in settings serving immigrant and undocumented populations—clear patient notification and consent requirements are essential. HHS should provide guidance on minimum patient transparency standards for ambient AI in clinical settings.

5. Ensure reimbursement pathways do not exclude safety-net providers. PHC’s experience with LumineticsCore demonstrates that autonomous AI can deliver measurable clinical value in a free clinic setting. As HHS considers reimbursement policy for AI-enabled services, it should ensure that payment models account for uninsured and underinsured populations and do not create incentives that concentrate AI adoption only in commercially insured settings.

V. Conclusion

PHC's experience demonstrates two things simultaneously: that AI can meaningfully improve care for underserved populations, and that adopting AI responsibly requires governance capacity that most small providers do not have. HHS has an opportunity to accelerate AI adoption across the healthcare system—but only if it ensures that safety-net providers are included in that acceleration, not left behind by it.

We welcome the opportunity to discuss our experience further and to contribute to the development of evaluation methods and governance frameworks appropriate for resource-constrained clinical settings.

Respectfully submitted,

Mairi Leining, MD, MPH, FACP

Chief Executive Officer
People's Health Clinic
650 Round Valley Drive
PO Box 681558
Park City, Utah 84060

Bruce Fryer, MBA

AI Implementation Lead (pro bono)
People's Health Clinic
650 Round Valley Drive
PO Box 681558
Park City, Utah 84060