



Rubrum Advising

Request for Information

Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care
RIN 0955-AA13

Rubrum Advising appreciates the opportunity to comment on the Department of Health and Human Services (HHS) Request for Information on “Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care”. Rubrum Advising, co-founded by Lee A. Fleisher, MD, ML, and Matthew A. Fleisher, JD, MPH, advises clients on the adoption, diffusion, and reimbursement of innovative healthcare solutions and products to improve patient outcomes and reduce total costs of care within the context of regulation, safety, quality, and clinical standards.

What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

- i. *Fragmented, low-quality data that make it hard to build safe models that are applicable to a wide range of populations and situations.*
 - a. There is growing concern that non-representative data lead to the creation of inaccurate algorithms, resulting in poorer performance in certain patient groups. To ensure algorithmic safety, datasets need to be appropriately representative of their intended use population.¹
 - b. Since nationwide exchange of health information remains fragmented, developers and health systems struggle to assemble sufficiently representative, longitudinal datasets for training and external validation. Even with exchange networks in place, important clinical information often remains in clinician notes or non-standardized fields, making it difficult to extract relevant data. Moreover, companies may be reluctant to share the data sets used to train an AI algorithm, given concerns that competitors could reverse engineer the data to create their own tools.² The use of AI tools requires a high level of specificity that has not yet been fully reconciled with the realities of current documentation and record-keeping practices.
 - c. As private entities increasingly aggregate, analyze, and sell large-scale clinical datasets, the expansion of commercial data markets introduces important policy considerations related to transparency, representativeness, and access. In many cases, there is limited visibility into how these datasets are constructed, validated, and maintained, which creates challenges for both regulators and developers seeking to assess data quality, completeness, and potential bias. The growing commercialization of clinical data also raises fundamental questions regarding data stewardship and ownership, including whether proprietary rights should rest with the organizations that collect, process, and

¹ Kuhlman, C., Jackson, L. & Chunara, R. No computation without representation: avoiding data and algorithm biases through diversity. In *Proc. 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining (KDD '20)* 3593 (ACM, 2020); 10.1145/3394486.3411074

² Coleman, K. (2024). *Lowering health care costs through AI: The possibilities and barriers*. Paragon Health Institute.

distribute these datasets or with the patients whose information forms the underlying data. Clarifying these governance frameworks is critical to ensuring appropriate patient protections, promoting responsible innovation, and aligning data stewardship practices with evolving standards for privacy, consent, and data sharing in healthcare research and product development.

ii. *Clinician and patient skepticism of “black-box” tools*

- a. The “black box” nature of deep neural networks and other advanced algorithms can make it nearly impossible for clinicians to interpret how the system arrives at specific recommendations.³ While many studies report gains in AI-based diagnostic efficiency and resource allocation, accountability and liability remain pervasive concerns, especially when models lack clear rationale or interpretability.
- b. An AMA survey from November 2024 found that although physician sentiment and use of AI is trending positive, physicians continue to express concerns about the design of health AI tools and the potential for flawed AI-enabled tools to put privacy at risk, integrate poorly with EHR systems, offer incorrect conclusions or recommendations, and introduce new liability concerns. Nearly half of physicians surveyed also ranked increased oversight as the number one regulatory action that needed to happen to increase their trust in adopting AI tools.⁴
 1. Under current malpractice law, liability is determined based on the “reasonable physician under similar circumstances” standard. Regardless of whether AI was used, courts will judge the conduct of the physician. As of now, there is no doctrine assigning shared responsibility to AI systems, even when their recommendations directly influence patient care.⁵
- c. Recent surveys suggest that many patients lack confidence in health care systems’ ability to use AI responsibly and protect them from AI-related harms. One study found that most respondents reported low trust in their health care system to use AI responsibly (65.8%) and low trust that their health care system would make sure an AI tool would not harm them (57.7%).⁶ Patients also generally prefer explainable AI, even if that means a trade-off in accuracy compared with black-box models.⁷
- d. Given these ongoing concerns about accountability, liability, and trust, strengthening instructions-for-use (IFU) and implementation education can also play an important role

³ Saoudi, C. E. N., Uren, V., & Jariwala, S. (2025). Evaluating accountability, transparency, and bias in AI-assisted healthcare decision-making: A qualitative study of healthcare professionals’ perspectives in the UK. *BMC Medical Ethics*, 26, 89. <https://doi.org/10.1186/s12910-025-01243-z>

⁴ American Medical Association. (2025). *Physician sentiments around the use of AI in health care: Motivations, opportunities, risks, and use cases – Shifts from 2023 to 2024*. American Medical Association.

⁵ Mello, M. M., & Studdert, D. M. (2008). Professional Liability: Malpractice and Other Causes of Action. *The New England Journal of Medicine*, 359(19), 1936–1942; <https://carey.jhu.edu/articles/fault-lines-health-care-ai-part-two-whos-responsible-when-ai-gets-it-wrong>

⁶ Nong, P., & Platt, J. (2025). Patients’ trust in health systems to use artificial intelligence. *JAMA Network Open*, 8(2), e2460628. <https://doi.org/10.1001/jamanetworkopen.2024.60628>

⁷ Nair, M., Svedberg, P., Larsson, I., & Nygren, J. M. (2024). A comprehensive overview of barriers and strategies for AI implementation in healthcare: Mixed-method design. *PLOS ONE*, 19(8), e0305949. <https://doi.org/10.1371/journal.pone.0305949>.

in managing risk for both clinicians and developers. Clear, detailed IFUs that specify intended use, patient populations, performance limitations, and required clinician oversight steps can help encourage responsible use of AI. By setting transparent expectations for how and when an AI system should be used, and by establishing clear physician training protocols, developers can mitigate risks associated with common drivers of adverse events and malpractice claims.

1. One potential approach is to ensure robust instructions for use so that clinicians know how, and for whom, to apply or trust AI tools in the real world. This idea is outlined in a recent publication by Kev Coleman at the Paragon Health Institute, where he argues that, instead of (or in addition to) treating every AI model as a regulated medical device, regulators can sometimes rely on robust instructions for use so clinicians understand the tool's intended population, setting, and limits.⁸ That includes being explicit about when not to use or trust an algorithm, and what kind of clinical oversight is required in real-world deployment.
- e. As federal and state entities evaluate and regulate clinical AI tools, it is essential to establish clear and operational distinctions between tools that incorporate a physician in the loop and those that function with minimal or no direct clinical oversight. Tools that require an actively involved clinician embed professional accountability and real-time risk mitigation, which may justify more flexible evidentiary and monitoring approaches, while fully autonomous or near-autonomous tools that drive high-stakes clinical decisions should be subject to more stringent requirements related to validation, ongoing performance monitoring, and clear lines of responsibility. High-stakes clinical decisions should be subject to more stringent requirements related to validation, ongoing performance monitoring, and clear lines of responsibility.

iii. *Unclear pathways for CMS coverage*

- a. MedPAC has noted that software technologies are frequently bundled into existing payments and that payments for SaaS (software as a service) have to meet specific conditions (e.g. is FDA-approved, fits into a Medicare benefit category, is reasonable and necessary) for reimbursement. This can make it challenging for developers to secure predictable reimbursement, especially when they must also negotiate separately with commercial insurers and creates particular challenges for tools that primarily deliver general risk reduction, utilization avoidance, or workflow efficiencies rather than a discrete billable service.⁹
- b. These challenges are compounded by the central role of the AMA in maintaining CPT codes, which remain the primary mechanism for payment under Medicare and most commercial plans. Many AI-enabled services lack dedicated CPT codes or fit poorly within existing categories, and even when new codes are established, they may not be tied to a clear Medicare benefit category or coverage pathway, limiting providers' ability

⁸ Coleman, K. (2024). *Lowering health care costs through AI: The possibilities and barriers*. Paragon Health Institute.

⁹ Zabinski, D., Druckman, J., Gerhardt, G., & Ray, N. (2025, April 10). *Paying for software technologies in Medicare* [Presentation]. Medicare Payment Advisory Commission. <https://www.medpac.gov/wp-content/uploads/2025/04/Tab-G-Saas-April-2025.pdf>.

to bill for their use and making reimbursement unpredictable. In this context, alternative payment models such as the CMS ACCESS model, which is designed to expand technology-supported chronic care through outcome-aligned payments, offer a more promising pathway for supporting AI-enabled tools that serve high-needs patients and emphasize preventive care to avoid hospitalizations.

- c. Because AI-enabled services must still map to traditional Medicare benefit categories, they are frequently subject to bundled or ad hoc payment approaches that absorb their value into broader procedures, leaving innovators without a clear, standalone reimbursement pathway and contributing to FDA-authorized technologies failing to reach patients at scale.¹⁰

What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.

- i. *Create clear coverage for AI interventions that enable safer, more timely care rather than increasing volume*
 - a. CMS could prioritize payment models that reward reductions in avoidable admissions, better chronic disease management, fewer diagnostic errors, and improved access to care, such as proposed in the ACCESS model, which is explicitly designed to expand technology-supported chronic care using an outcome-aligned payment approach.
 - b. CMS could also build a clearer reimbursement pathway for software/AI that results in clinical improvements and should continue to publish instructions and guidelines for differentiating AI as its own clinical service, as an administrative service, or as part of an existing service.
- ii. *Physicians should not be paid more just because AI tools make them more efficient*
 - a. In the CY 2026 Physician Fee Schedule (PFS) final rule, CMS finalized a -2.5% efficiency adjustment to select services to better recognize that some services are likely to become more efficient over time, as compared to time-based services like office visits or behavioral health therapy. The adjustment targeted services like diagnostic imaging interpretation, outpatient interventions, interventional pain management, and orthopedic services, which have benefited from technological advancements and standardized workflows that reduce time and resource use, without corresponding payment adjustments.¹¹

¹⁰ Seeley, A. (2025, May 23). *How the Health Tech Investment Act could reshape Medicare reimbursement for algorithm-based services*. Morgan, Lewis & Bockius LLP. <https://www.morganlewis.com/pubs/2025/05/how-the-health-tech-investment-act-could-reshape-medicare-reimbursement-for-algorithm-based-services>

¹¹ Centers for Medicare & Medicaid Services. (2025, November 5). *CMS modernizes payment accuracy and significantly cuts spending waste while improving chronic disease management for Medicare beneficiaries* [Press release]. <https://www.cms.gov/newsroom/press-releases/cms-modernizes-payment-accuracy-significantly-cuts-spending-waste>

- b. Similarly, where AI leads to greater efficiency or streamlines the provision of care, CMS could treat this increased productivity as a deflationary effect and apply an efficiency adjustment to services that AI is expected to streamline. Aligning payments with the actual resources required to provide the AI-enabled services would ensure that payments are not higher simply due to efficiency gains, while still allowing AI tools that measurably improve outcomes or reduce downstream costs to be supported through outcome-based or model-specific payments.
- iii. *Embed AI-related safety expectations in Medicare quality programs and innovation models*
 - a. Existing and newly launched quality and payment programs, such as the Quality Payment Program (QPP) and the Merit-based Incentive Payment System (MIPS), as well as Innovation Center models like ACCESS, can incorporate AI governance where AI is used for clinical decision-making, including regular oversight, risk assessment, and training requirements tied to the safe use of AI tools.
- iv. *Use demonstration projects to test payment models that explicitly fund AI tools that prevent hospitalizations or support high-risk, high-need patients at home*
 - a. As stated above, the new CMS ACCESS model is designed to expand technology-supported chronic care using an outcome-aligned payment approach and could tie in AI-enabled tools that support care for high-needs patients and emphasize preventative care to avoid hospitalizations. We strongly support both the construct of the ACCESS model and the model's use of well-defined endpoints and transparency of results. Future demonstration projects should be structured so that payments reflect both outcomes and appropriate efficiency gains, rather than creating stand-alone fees for AI that do not change care delivery.

For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (e.g., relating to liability, indemnification, privacy, and security). What novel legal and implementation issues exist and what role, if any, should HHS play to help address them?

- i. *Require strong privacy and governance standards for patient-generated data from consumer devices used in clinical care, including safeguards against discriminatory secondary uses of sensitive information.*
 - a. Data from non-medical devices, such as wearables and fitness/wellness apps, increasingly informs clinical decision-making when patients share this information with their providers, yet it is often collected and stored outside HIPAA's regulatory framework. Once incorporated into clinical records, these data can influence care and outcomes, making robust privacy and security safeguards essential. These systems should require explicit, informed patient consent for any sale or secondary use of data and provide enhanced protections and remedies when those protections are violated. HHS should prioritize extending clear protections and governance standards to patient-generated data from consumer devices to ensure that information used in clinical care is subject to consistent privacy, security, and accountability requirements.

For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and other workflow and human-centered evaluation methods for clinical care? Should HHS further support these processes? If so, which mechanisms would be most impactful (e.g., contracts, grants, cooperative agreements, and/or prize competitions)?

- i. *Shift from pre- to post-market surveillance*
 - a. Pre-market evaluation has worked reasonably well for physical devices and traditional software with stable, predictable outputs, but it is poorly matched to AI systems whose behavior evolves with data, context, and implementation over time. Rather than relying primarily on pre-deployment testing, HHS (through FDA and ONC) should support a risk-based, post-market surveillance framework for clinical AI that emphasizes: mandatory adverse event and error reporting, outcome monitoring in real-world use, structured identification of implementation issues, and periodic performance re-evaluation. The goal should be to avoid both under-regulation that allows mass injury and over-regulation that delays access to beneficial technologies, while minimizing compliance burden and total cost of ownership for providers and developers. This framework should be targeted to higher-risk clinical use cases and AI systems whose design or training makes output unpredictability likely, so that oversight resources are focused where patient harm risk is highest.¹²
- ii. *Fund real-world evaluation via grants, cooperative agreements*
 - a. HHS could use grants, cooperative agreements, and contracts to support public-private evaluation partnerships, where developers receive de-identified post-market performance data and providers can safely continue using AI tools in higher-risk scenarios under structured monitoring arrangements.
 - b. HHS should also create a national, registry-style, federated post-market surveillance network for clinical AI, modeled in part on the FDA Sentinel Initiative’s distribution data infrastructure, to aggregate outcome data and support standardized reporting of errors, adverse events, and key performance metrics. Such an aggregated outcome data registry would enable early detection of safety issues across health systems and provide feedback loops for regulators, developers, and providers, without requiring continuous, centralized storage of identifiable patient data.

How can HHS best support private sector activities (e.g., accreditation, certification, industry-driven testing, and credentialing) to promote innovative and effective AI use in clinical care?

- i. *HHS should foster independent, industry-led evaluation and benchmarking of clinical AI tools, aligned with models like ACCESS and TEMPO, to support safe and effective innovation in clinical care*

¹² Coleman, K., & Pencina, M. J. (2025). *Targeted postmarket surveillance: The way toward responsible AI innovation in health care*. Paragon Health Institute.

- a. HHS can best support private sector activities by explicitly leveraging and extending efforts such as the ACCESS and TEMPO models to create a predictable environment for testing and deploying AI in clinical care. In addition to advancing outcomes-based, technology-enabled care, HHS should encourage independent, industry-led evaluation and benchmarking of AI tools to assess their safety and effectiveness in real-world settings. HHS can further align data, reporting, and quality standards so that evidence generated through these private testing initiatives can be relied upon for coverage, payment, and oversight decisions. Finally, HHS should provide technical assistance and targeted incentives for providers and developers who participate in rigorous, transparent evaluation and benchmarking activities, thereby promoting responsible innovation without relying on formal accreditation or certification schemes.

Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?

- I. *Prioritize development and payment for AI tools that demonstrably change care delivery in ways that improve outcomes and reduce downstream costs, rather than tools that solely increase clinician efficiency*
 - a. AI tools have met or exceeded performance and cost expectations primarily where they have demonstrably reduced downstream utilization while maintaining or improving outcomes, rather than merely increasing clinician efficiency. Examples include models that prospectively identify patients at high risk for sepsis or avoidable readmission, tools that reduce unnecessary imaging or procedures, and decision-support systems that lower complication rates or emergency department revisits; in these cases, the value proposition is clear because payers and providers can observe fewer admissions, repeat tests, or adverse events, and thus lower the total cost of care. By contrast, tools that mainly make clinicians faster or generate documentation more efficiently, without clear evidence of changed care patterns or improved outcomes, have fallen short of expectations from a coverage and payment perspective, because the financial return largely accrues to efficiency gains that purchasers can already value without separate reimbursement. Consistent with emerging payment principles for AI, the novel tools with the greatest potential are those that can rigorously demonstrate that they change clinical decisions and care pathways in ways that improve outcomes and reduce downstream costs (for example, earlier targeted intervention in high-risk chronic conditions, reduced low-value care, or fewer complications), and that can be incorporated into payment models that reward measurable improvements in total cost and quality rather than efficiency alone.

Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

- I. *Enhance interoperability for structured clinical data (LOINC, USCDI)*
 - a. LOINC-coded results already underpin many digital quality measures and FHIR US Core profiles (for example, the US Core Blood Pressure profile requires fixed LOINC codes for systolic and diastolic values), which shows how standardized codes can make data predictable, computable, and reusable for AI and quality measurement. However, misalignment between laboratory information systems and EHR data elements, in addition to the use of scanned reports rather than discrete results, still leads to missing values, nonspecific codes, and loss of clinical context that limit AI model performance. ONC/ASTP should promote a national LOINC reference implementation for diagnostic tests, tied to USCDI and FHIR implementation guides, so that laboratories, EHRs, and AI developers can reliably exchange test-level data needed for training, validation, and benchmarking.
- II. *Interoperability and fair access for AI tools under TEFCA*
 - a. To improve interoperability, ONC/ASTP should create a LOINC reference database for diagnostic tests and standardize use among diagnostic manufacturers. TEFCA and related nationwide networks are beginning to support exchange of USCDI data classes across organizations, but without attention to governance, large incumbents or QHINs could effectively gatekeep data flows needed by third-party AI developers. HHS should ensure that TEFCA implementation includes clear requirements that certified APIs and exchange frameworks provide non-discriminatory technical access to standardized data for authorized AI tools, prevent exclusive or preferential access arrangements that disadvantage smaller innovators, and propagate guidance on consistent use of SNOMED CT, ICD, and CPT for diagnoses and procedures across clinical documentation, payment, and reporting.
- III. *Making unstructured data more usable*
 - a. Much clinically rich information remains in unstructured formats, such as scanned PDFs, narrative notes, or loosely structured “other” fields, that are not consistently mapped to USCDI elements or standard terminologies, limiting their value for AI development and evaluation. HHS should support implementation guides, incentives, and technical assistance to move key data elements from scanned or free-text formats into discrete, coded fields aligned with USCDI and FHIR profiles, while encouraging tools that can safely structure legacy unstructured content without degrading accuracy.

What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?

- I. *Challenges patients and caregivers would like to see addressed:*
 - a. Access to specialist-level care for individuals with complex or rare diseases in rural hospitals
 1. Access to specialist-level care for individuals with complex or rare diseases in rural hospitals can be strengthened through appropriate use of AI-enabled tools. These

- tools can help reallocate limited clinical and operational resources toward more complex cases in resource-constrained settings, allowing clinicians to focus their time and expertise where it is most needed. They can also support the diagnosis and prediction of rare diseases in rural areas, particularly where local access to relevant specialists is limited, so that patients receive timely and appropriate referrals to higher levels of care.
- b. Help navigating complex regimens, symptoms, and appointments
 - 1. Patients and physicians report greater trust in efficiency-based and administrative applications, as well as in tools that offer explanations and are more transparent over black-box tools.
 - c. Patient and caregiver concerns:
 - 1. Loss of human review/connection, specifically as related to clinical decision-making.
 - 2. AI being used primarily to cut costs or deny care rather than improve safety and outcomes
 - i. e.g. AI-based prior authorization tools, AI replacing clinical tasks and decision-making.
 - 3. Potential to exacerbate discrimination in care delivery due to reliance on data sets that may not accurately reflect the population in which the AI is being used and may therefore not provide the most accurate outputs, leading to unreliable patient outcomes.
 - i. These risks are further compounded by a lack of adequate clinician training on use of AI, including how to interpret, validate, and appropriately apply AI-generated outputs or recommendations, increasing the likelihood that inaccurate recommendations will be used and reinforce existing gaps in care delivery.

Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care?

- I. *Real-world impact of AI on patient outcomes, safety, and trust over time*
 - a. Implement multicenter, prospective studies and post-market monitoring that track effects on outcomes, safety events, diagnostic accuracy, and perception of trust over time.
- II. *Comparative evaluations of AI-enabled versus traditional workflows on clinical burden, documentation time, and time with patients*
 - a. Support studies using standardized workload and time-use measures, including frontline clinician feedback, to determine when AI meaningfully improves care versus simply redistributing tasks.
- III. *Methods to detect and mitigate inaccuracies, drift, and unintended harms*
 - a. Establish a national outcomes-focused reporting and monitoring system for AI-enabled tools, drawing on lessons from FDA-regulated products (such as drugs and biologics) without simply replicating their exact model. This system should prioritize higher-risk AI

use cases and focus on meaningful outcome and safety signals rather than requiring uniform reporting for all tools.

- i. Building on concepts from the Sentinel Initiative, stakeholders (including health systems, manufacturers, and technology providers) could participate in an aggregated outcome data registry to monitor real-world performance, identify potential safety concerns, and support appropriate updates or modifications to AI-enabled tools over time. This type of targeted post-market surveillance can be structured to align with the incentives of clinicians, health care organizations, and developers.¹³
 - ii. An AI outcomes registry could allow for the conduct of post-market monitoring and collection of data on performance across diverse settings and populations. An aggregated registry could track and analyze large volumes of clinical and operational data to identify patterns of performance, adverse effects, errors, or inaccuracies that emerge only after deployment. The registry could also track utilization patterns associated with positive clinical outcomes, providing benchmarks for best practices that developers and health systems can adopt. Moreover, AI systems that learn over time or are periodically retrained based on new data can exhibit changes in performance that are not captured by pre-market studies. Continuous monitoring would enable early detection of adverse trends, triggering investigations or changes including updated training protocols, algorithmic recalibration, or pausing certain use cases. By documenting these, the registry would contribute to improving trust among clinicians and patients and informing regulatory and oversight decisions.
 - iii. The registry could leverage a distributed data model similar to the Sentinel Initiative's, where data partners retain control of their data behind firewalls and only aggregate, de-identified outputs are shared. This approach preserves HIPAA compliance while allowing regulators, researchers, and developers to track the performance of AI tools without centralizing identifiable patient information. The data elements and interoperability formats could also be aligned with frameworks like USCDI to improve consistency and reuse.
 - iv. Over time, a voluntary outcomes registry could serve as the foundation for a more formal national framework, expanding participation and data sources to create a comprehensive, scalable oversight infrastructure for AI-enabled clinical technologies.
- b. Are there published findings about the impact of adopted AI tools and their use in clinical care?**
- i. *Our CEO has published a framework arguing that hospitals can and should govern the use of AI through existing CMS Conditions of Participation by*

¹³ Coleman, K., & Pencina, M. J. (2025). *Targeted postmarket surveillance: The way toward responsible AI innovation in health care*. Paragon Health Institute.

*integrating AI-related safety, monitoring, and response into their current patient safety and quality infrastructure.*¹⁴

Conclusion

Rubrum Advising appreciates the Department's leadership in seeking concrete, experience-based input at this early stage of federal policymaking on clinical AI. The recommendations in this response reflect our belief that AI should be deployed in ways that measurably improve outcomes and reduce total costs of care, while relying on existing safety and quality infrastructure, clear governance expectations, and predictable reimbursement pathways to support responsible innovation. We look forward to continued engagement with HHS as it refines its regulatory, payment, and research strategies to accelerate the safe, effective use of AI in clinical care.

Sincerely,

Rubrum Advising

¹⁴ Fleisher, L. A., & Economou-Zavlanos, N. J. (2024). Artificial intelligence can be regulated using current patient safety procedures and infrastructure in hospitals. *JAMA Health Forum*, 5(6), e241369. <https://doi.org/10.1001/jamahealthforum.2024.1369>.