

Robert A. Harrington, M.D.

Stephen and Suzanne Weiss Dean
Weill Cornell Medicine

Provost for Medical Affairs
Cornell University

February 23, 2026

Via Electronic Submission

The Honorable Robert Kennedy, Jr.
Secretary
U.S. Department of Health and Human Services
200 Independence Avenue SW
Washington, DC 20201

RE: HHS Health Sector AI RFI

Dear Secretary Kennedy:

Cornell University and Weill Cornell Medicine appreciate the opportunity to provide comments on the Department of Health and Human Services (HHS) Health Sector Artificial Intelligence (AI) Request for Information (RFI) published on December 23, 2025. As a leading academic medical center embedded within a major research university, Weill Cornell Medicine both develops and deploys AI tools across the full continuum of health—from discovery and validation to real-world clinical implementation. Our comments address each of the three main areas of focus identified in the RFI: Regulation, Reimbursement, Research & Development.

We have recently launched a **Center for Early Detection and Prevention (CEDP)**, designed to fundamentally shift healthcare from a reactive, late-stage treatment paradigm toward **proactive, personalized prevention**. The Center leverages **multimodality AI** to integrate diverse longitudinal data—including advanced imaging, laboratory testing, electronic health record (EHR) data, multi-cancer early detection (MCED) assays, genomic and polygenic risk information, lifestyle assessments, and wearable/remote monitoring data—to generate individualized risk profiles across cancer, cardiovascular, neurodegenerative, and metabolic diseases.

Our experience underscores a central conclusion: **the greatest unrealized opportunity for AI in health care lies in prevention, early detection, and risk stratification**, yet current regulatory, reimbursement, and research frameworks remain largely optimized for episodic treatment of advanced disease. If the nation seeks to save lives, reduce the chronic disease burden, improve access and equity, and control long-term costs, federal policy must evolve to support prevention-focused, multimodality AI.



I. Regulation

A. Regulatory and Programmatic Changes HHS Should Prioritize

1. Establish a clear governance framework for non-medical-device clinical AI, particularly for prevention and risk stratification

Many of the most impactful AI tools used in clinical care—including documentation support, care navigation, population health analytics, and risk stratification—are not regulated as FDA medical devices. These tools increasingly inform screening intensity, diagnostic sequencing, and preventive interventions, yet organizations face uncertainty regarding acceptable evidence, monitoring expectations, and accountability.

HHS should establish a **risk-tiered assurance framework** for non-device AI used in clinical care, explicitly including prevention and early detection. Such a framework should define proportional expectations for:

- transparency regarding intended use and limitations,
- validation appropriate to downstream clinical action,
- post-deployment monitoring and drift management,
- clinical governance and accountability,
- patient-facing communication when AI influences screening or preventive decisions.

This framework should align with ONC certification expectations for decision support interventions under **45 CFR Part 170**, including transparency, reliability, robustness, fairness, safety, security, and privacy considerations.

2. Enable data liquidity for multimodality AI while protecting patient privacy

Prevention-focused AI requires longitudinal access to heterogeneous data sources that are often siloed across vendors and domains. While HIPAA and information-blocking regulations permit such access in principle, operational barriers remain substantial.

HHS should:

- Strengthen enforcement and implementation guidance under **45 CFR Part 171 (Information Blocking)** to ensure timely, standards-based access to electronic health information—including imaging, laboratory, genomic, and patient-generated data—necessary for AI development, deployment, and monitoring.
- Clarify that AI-driven risk stratification and early detection fall squarely within permissible **treatment and health care operations** under HIPAA (**45 CFR §§ 164.502, 164.506**).



- Issue updated guidance supporting privacy-preserving analytics, de-identification, and expert determination pathways (**45 CFR § 164.514**), particularly for multicenter and longitudinal prevention research.
- Streamline vendor contracting by specifying that a business associate may not deidentify and use data to train its own proprietary AI as “proper management and administration” under a business associate agreement (BAA) unless the covered entity explicitly permits it (**45 CFR § 164.504(e)(2)**)
- Establish baseline privacy protective guidance for any entities that process US sensitive health data, whether they are a HIPAA covered entity, business associate, or otherwise.

This will lower enforcement and contracting costs for covered entities and enhance patient trust if digital health technology providers by default must use data in a manner consistent with the purposes for which it was collected and subject to minimum necessary standards. (**45 CFR § 164.514(d)**).

3. Modernize regulatory posture to support prevention as a core clinical use case

Existing regulatory frameworks implicitly prioritize treatment decisions over prevention. HHS should explicitly recognize **screening optimization, early detection, and personalized prevention** as high-value clinical use cases for AI and ensure that regulatory expectations do not unintentionally discourage innovation in these areas.

B. Novel Legal and Implementation Issues for Non-Medical-Device AI

The use of AI for longitudinal risk prediction and prevention raises novel governance challenges:

- **Accountability over time:** Risk models evolve as new data streams are incorporated, raising questions about responsibility for recalibration, communication of changing risk, and oversight.
- **Action thresholds:** Determining when probabilistic risk estimates trigger screening, preventive therapy, or referral requires clinical and policy alignment.
- **Equity:** Prevention-focused AI must avoid systematically disadvantaging populations with less complete data (e.g., limited access to wearables or genomic testing).
- **Contracting and indemnification:** Health systems face challenges negotiating audit rights, update policies, and liability sharing for adaptive models.

HHS should play a convening role to:

- define best practices for governance of longitudinal, prevention-focused AI,
- develop model contractual provisions addressing auditability, updates, and incident response,
- promote minimum expectations for equity monitoring and transparency.



C. Evaluation Methods and Metrics for Non-Device AI

Pre-deployment evaluation should include:

- external validation across diverse populations,
- calibration of absolute risk estimates,
- simulation of downstream screening and prevention workflows,
- human-centered usability testing.

Post-deployment monitoring should assess:

- performance drift and subgroup equity,
- workflow impact (e.g., access, timeliness, burden),
- safety signals and override rates,
- clinical outcomes such as stage shift and avoidance of advanced disease,
- long-term economic impact.

HHS support would be most impactful through **cooperative agreements, contracts, and multicenter evaluation networks** that provide shared infrastructure for real-world evidence generation in prevention-focused AI.

D. Supporting Private-Sector Accreditation and Certification

HHS can catalyze innovation by supporting:

- voluntary **clinical AI assurance labels** documenting intended use, limitations, subgroup performance, update cadence, and monitoring commitments;
- independent testing sandboxes and neutral evaluation laboratories;
- alignment between private accreditation efforts and ONC-aligned transparency and risk management expectations under **45 CFR Part 170**.

II. Reimbursement

A. Structural Misalignment in Current Payment Systems

One of the most significant barriers to AI-enabled prevention is a profound misalignment in reimbursement. **Screening, early detection, and prevention are reimbursed at levels far below those for treating advanced disease**, including surgery, hospitalization, and high-cost therapeutics.



As a result:

- Health systems bear the cost of prevention infrastructure and AI deployment.
- Financial returns accrue downstream to payers and society.
- Incentives favor late-stage intervention over early action.

This misalignment undermines national goals to reduce mortality, chronic disease burden, and healthcare costs.

B. Payment Policy Changes HHS Should Prioritize

HHS should:

- Develop reimbursement pathways for **AI-enabled risk stratification, personalized screening, and longitudinal prevention management**.
- Use **CMMI demonstration models** to test payment approaches that reward prevention and share savings when AI reduces downstream utilization.
- Ensure that payment policy recognizes **value over time**, including avoided hospitalizations, reduced disease severity, and improved access.
- Tie payment to outcomes to avoid scaling ineffective or low value technologies by deploying a coverage with evidence development approach¹
- Create tailored – not one size fits all – payment mechanisms for AI-enabled clinical services that incentivize quality improvement. When an AI-enabled clinical service complements an existing billable diagnostic or image service, bundled payment design should be used, potentially with add-on payments for use of AI in scenarios where the AI augmentation improves quality of care and/or patient outcomes. When the AI-enabled clinical service does not have a complementary diagnostic or image service, a separate payment mechanism should be created only when there is evidence of quality improvement. This framework is delineated by Zink and colleagues²
- When clinical efficacy of an AI device or algorithm is established, CMS should consider assigning a temporary procedure code with provisional status that authorizes payment for a limited period. This strategy could help address the current catch-22 where payers and the AMA panel that approves CPT codes require evidence of meaningful utilization before approving a new code with payment, but providers are unlikely to adopt new technologies in the absence of payment. See recommendation from Razavian and colleagues³
- Categorize AI services by their functions (e.g., augment vs replace existing services, diagnostic vs predictive) and establish umbrella payment policies across functional groups. When

¹ <https://pubmed.ncbi.nlm.nih.gov/35981123/>

² <https://pubmed.ncbi.nlm.nih.gov/38805195/>

³ <https://ai.nejm.org/doi/full/10.1056/AIpc2500871/>



multiple AI products have the same function, reference pricing should be used to set the payment rate based on the cost of the most efficient product. See recommendation from Razavian and colleagues ⁴

C. Decision-Makers and Administrative Hurdles

Adoption decisions are shaped by clinical leadership, population health leaders, CMIO/CIOs, quality and compliance teams, and enterprise AI governance bodies.

Primary hurdles include:

- lack of reimbursement for prevention-focused AI,
- integration of diverse data sources (imaging, genomics, lab tests, wearables),
- costs of ongoing monitoring and governance,
- difficulty demonstrating ROI within short budget cycles.

D. Patient and Caregiver Perspectives

Patients and caregivers seek:

- earlier identification of disease risk,
- personalized screening rather than one-size-fits-all approaches,
- reduced fear and morbidity associated with late diagnoses,
- equitable access to preventive care and disease management, through AI-enhanced virtual and hybrid care.

Concerns include privacy, data use, fairness, transparency, accessibility, and accountability. HHS can address these through patient-facing transparency standards and enforceable equity monitoring expectations.

III. Research & Development

A. Barriers to Innovation and Adoption for Private Sectors

Key barriers include fragmented data ecosystems, limited shared infrastructure for longitudinal prevention research, insufficient funding for implementation science, and misaligned economic incentives. Also rapid pace of change presents challenges for decisions about investment in and scaling of AI for clinical care. Today's technology will be out of date quickly.

⁴ Ibid



B. Where AI Has Succeeded and Fallen Short

AI has succeeded when tightly integrated into workflows and when benefits are tangible and near-term. It has fallen short when models are siloed, poorly monitored, or when prevention benefits accrue outside existing payment structures.

There are several areas where clinical AI has met or exceeded expectations. For example, AI used as a second reader/triage in breast cancer screening has shown meaningful gains in detection and/or workload reduction at population scale, with emerging evidence on favorable outcomes and cost-effectiveness in Swedish programs and trials. Diabetic retinopathy (DR) autonomous screening systems have demonstrated strong diagnostic performance in real-world settings and can close screening gaps when implemented well. In multi-site evaluations, ambient scribes are associated with reduced documentation burden and improved clinician experience—often meeting the “quality of life” expectation even when hard ROI is still evolving.

At the same time, there are examples of areas where AI has fallen short. Sepsis prediction is a classic example. Some systems have shown limited external validity and/or high false-alert burdens, and real-world impact can be far less than vendor claims. Diabetic retinopathy autonomous screening illustrates this. Claims-based data suggest low real-world uptake, with reimbursement levels and implementation costs acting as practical bottlenecks. For ambient scribes, early evidence shows promise for burden/burnout, but system-level financial or throughput gains can be mixed—and there are open questions about downstream spending effects (e.g., coding intensity).

C. Interoperability as a Catalyst

Enhanced interoperability—particularly for imaging, genomics, laboratory data, wearables, and patient-reported outcomes—would dramatically accelerate development of multimodality AI. . The recent ASTP HTI-5 proposed rule⁵ on promoting FHIR-based data exchange and AI-enabled interoperability solutions that use FHIR alongside newer standards like the Model Context Protocol (MCP), may go a long way in improving interoperability for these diverse data. Standardized representations for AI-derived risk scores, audit trails, and outcomes are essential.

D. Priority Research Areas for HHS

HHS should prioritize:

- multimodality risk prediction across diseases,
- AI-enabled screening optimization and early detection,
- equity-aware prevention and management models,
- longitudinal economic analyses comparing prevention versus late-stage treatment,

⁵ <https://www.federalregister.gov/documents/2025/12/29/2025-23896/health-data-technology-and-interoperability-astponc-deregulatory-actions-to-unleash-prosperity>



- scalable governance models for lifelong risk assessment.

IV. Benefits and Challenges of Use of AI in the Clinical Space

Part of the promise of AI, making it unique among health technologies, is that the speed and cost of device/algorithm development is relatively low (e.g., as opposed to drug development) once appropriate infrastructure and access to data is achieved. HHS can facilitate the creation of the infrastructure necessary to both develop and subsequently deploy by setting regulations/standards for interoperability and the use of data in a secure and ethical manner.

A requirement for a standard-based interoperability for AI devices must apply to electronic health records (EHR), AI algorithms/platforms, and medical devices. AI can process and incorporate large volumes of data and detect signals and patterns that humans cannot. These patterns may be useful in early detection prediction, and ultimately prevention of disease across populations that extend beyond individual or even combinations of health care systems. However, when standards are lacking, the effort to acquire, normalize, and combine data sets has been and will continue to be a barrier requiring significant cost and expertise to overcome. Data and interface standards are necessary to allow AI to achieve its true potential to improve the health of patients on a population scale.

Another barrier is the lack of guidance on the use of individual patient data. AI requires individual patient data for development. In many cases, there may not be direct benefit to the individual whose data is utilized, rather for future patients or populations. The lack of guidance and understanding on the appropriate use of patient data has and will continue to result in reluctance to allow data to be used. This has resulted in a slowing or stoppage in algorithm development and improvement. By providing clear guidance on how appropriately deidentified and protected patient data may be incorporated into AI models, presumably without individual patient consent, HHS would facilitate the development of large data sets necessary to develop population-based algorithms.

HHS and other entities seeking to maximize adoption of AI in the clinical setting will encounter many other challenges. As relates to data access, quality and fragmentation, health data is currently siloed across institutions, vendors and payers. EHR data are often incomplete, biased and shaped by billing and workflow rather than clinical truth. And many AI models are trained on narrow populations, raising concerns about generalizability. And the competitive landscape in the commercial EHR space may limit innovation by smaller firms.

When examining clinical validation and evidence gaps, AI tools often demonstrate retrospective performance but lack prospective, workflow-embedded validation. And there is no universal standard for clinical AI across tasks. It is difficult to demonstrate that AI, rather than clinicians or other systems, are the cause of improved outcomes.

AI can disrupt workflow, rather than enhance it, because AI tools often add cognitive load rather than reduce it. Alerts, predictions and explanations are often misaligned with how clinicians think and act. Clinical settings require near-perfect uptime and predictable behavior, which make them vulnerable to any latency or reliability constraints.



Providers will also encounter gaps as relates to trust, accountability and liability. Clinicians are hesitant to rely on systems they cannot interrogate. An over-reliance on AI could result in errors in the clinical space that are worse than would otherwise be the case. As relates to errors, healthcare providers will worry that AI will increase malpractice risk without legal protection.

Any use of AI tools in the clinical space will require constant calibration and adjustment. Clinical practice, coding and populations change over time. And models often need site-specific retraining or recalibration. Significant resources will need to be devoted to monitoring, updating and governing AI systems.

There should be a mechanism for gathering and integrating the patient perspective in the creation and use of AI models in the clinical space. There is ample evidence that shows patient-reported outcomes (PROs) can inform clinical practice in a manner that reduces patient mortality. These observations are particularly valuable in use of AI in disease screening, early detection and prevention.

There are certain features that AI tools could or should have that could substantially improve outcomes, reveal quality gaps, and reduce costs:

- Tools that turn recommendations into reliably executed care: next-best-action suggestions tied to order sets, contraindications, care pathways, and local policies—designed to reduce missed steps and variation rather than add alerts.
- Multimodal longitudinal risk + monitoring (the “early signal” stack) AI tools that combine EHR, notes, imaging, labs and wearable/home data to detect deterioration earlier (e.g., heart failure decompensation, post-op complications) with calibrated uncertainty and actionable thresholds—plus monitoring for drift and subgroup performance.
- AI tools that continuously find missed follow-ups, delays in diagnosis, gaps in screening, unsafe medication combinations, inequities in access or outcomes by subgroup, or documentation/ordering patterns signaling process failures, with outputs routed to quality teams (not only frontline clinicians).

AI tools have great potential for cost savings across both the entire healthcare delivery system and within specific systems. Tools can be utilized that directly optimize staffing, bed management, scheduling, imaging queues, discharge timing, and supply utilization under real constraints. These can generate hard savings because they target throughput and waste. AI can make it easier to run rapid, reproducible “learning health system” studies by utilization for cohort construction, phenotype definitions, confounding-aware analyses, and pragmatic trial operations. This also means that adoption decisions can be driven by local evidence, rather than marketing.



Benchmarking tools can and should be utilized, such as open federated evaluation platforms like MedPerf⁶ for privacy-preserving, multi-site evaluation. This is critical for generalization, bias, and drift testing. Such tools enable scalable products in risk stratification, care gaps, quality measurement, and operations because vendors can deploy the same pipelines across many health systems.

We should have cross-network exchange for longitudinal outcomes and follow-up for care continuity and research. AI value often depends on outcomes observed outside one institution such as readmissions data, follow-up care, mortality and utilization at other institutions or in other systems. Network exchange makes outcome measurement and model monitoring more complete. The Trusted Exchange Framework and Common Agreement (TEFCA) is an example of a framework that enables nationwide exchange across health information networks under common “rules of the road.” TEFCA work is explicitly evolving toward enhanced exchange approaches, including FHIR-based exchange pilots.

Establishment of standards as relates to devices, wearables, and remote monitoring signals is particularly important for “closed-loop” clinical AI. The biggest outcome gains often come from continuous monitoring and intervention workflows. Interoperability is the difference between pilots and scalable programs. There are other areas where establishment of such standards could show significant improvement in health outcomes. Examples are device vitals, waveforms, activity/sleep proxies, home measurements and infusion/ventilator parameters. The result is scalable “hospital-at-home,” chronic care, and high-risk monitoring programs with measurable ROI.

V. Conclusion

If the United States is serious about saving lives, reducing the chronic disease burden, improving access, and bending the healthcare cost curve, **AI policy must move beyond treatment optimization and explicitly support prevention, early detection, and personalized risk stratification.**

Multimodality AI—supported by aligned regulation, reimbursement, and research investment—offers one of the most promising paths to achieving these goals.

Sincerely,



Robert A. Harrington, M.D.

⁶ <https://www.nature.com/articles/s42256-023-00652-2>

