

RESPONSE TO HHS REQUEST FOR INFORMATION

Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care

90 FR 60108 | Document Number 2025-23641 | Submitted February 23, 2026

Submitted by: Dennis Morrison, PhD, Eleos Health | Focus Area: Behavioral Health

Introduction

This response addresses the ten specific questions posed in HHS's RFI on accelerating AI adoption in clinical care (90 FR 60108), with an exclusive focus on behavioral health—encompassing mental health, substance use disorders (SUD), and community behavioral health organizations (CBHOs). Behavioral health occupies a unique position in the AI landscape: it faces the some of the most severe workforce shortages, the greatest documentation burden relative to encounter volume, the most significant access disparities, and clinical complexity that demands nuanced, evidence-based approaches that current AI tools are only beginning to approximate.

Behavioral health organizations, particularly Certified Community Behavioral Health Clinics (CCBHCs) and community mental health centers (CMHCs), treat the nation's most vulnerable populations—individuals with serious mental illness (SMI), co-occurring disorders, and those involved in the justice system. These organizations have chronically operated on thin margins, and the combination of Medicaid funding pressures, clinician burnout, and documentation mandates threatens their sustainability. AI tools—properly governed, evaluated, and reimbursed—represent a rare opportunity to meaningfully address each of these challenges simultaneously.

This response draws on the work of pioneering behavioral health AI companies including Eleos Health (ambient documentation and compliance AI) and Lyssn (clinical quality improvement and evidence-based practice fidelity AI), as well as the clinical research of Dr. John Torous, MD, MBI, Director of the Digital Psychiatry Division at Beth Israel Deaconess Medical Center (BIDMC) / Harvard Medical School, whose decade-plus body of work on digital mental health tools, AI chatbots, and mobile mental health applications provides essential evidence-based grounding.

Question 1: What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

Workforce Crisis and Documentation Burden

Behavioral health faces a workforce crisis of historic proportions. Clinicians in community behavioral health settings report spending 30-50% of their working hours on documentation rather than direct care. This documentation burden is a primary driver of the 90% burnout rate cited across the sector. Companies like Eleos Health have demonstrated that ambient AI documentation tools can reduce documentation time by more than 70%, with providers accepting 80% of AI-generated note suggestions and 90% of notes submitted within 24 hours—dramatically reducing late submission denials.

Yet adoption remains uneven because many behavioral health organizations, particularly rural and safety-net providers, lack the IT infrastructure, broadband connectivity, and technical staff to implement sophisticated AI platforms. The "digital divide" that Dr. John Torous has long documented in his research on digital equity and the BIDMC Digital Navigator program extends directly into organizational AI readiness.

Regulatory Fragmentation

The behavioral health regulatory environment is uniquely fragmented. Unlike general medical AI, which primarily navigates FDA and HIPAA frameworks, behavioral health AI must simultaneously address 42 CFR Part 2 restrictions on substance use disorder records, state-by-state licensure and telehealth regulations, payer-specific prior authorization requirements, and an emerging patchwork of state AI-specific legislation. Dr. Torous's 2025 50-state legislative review published in JMIR Mental Health identified widely varying approaches — from Illinois's near-ban on certain AI uses in behavioral health to Utah's disclosure-focused framework — creating a compliance environment that disadvantages smaller innovators and CBHOs.

The Part 2 landscape shifted substantially with HHS's February 2024 Final Rule (89 Fed. Reg. 12496), which implemented Section 3221 of the CARES Act by aligning significant portions of 42 CFR Part 2 with HIPAA. The rule took effect April 16, 2024, with a compliance deadline of February 16, 2026 — the week this RFI comment period closes. Key changes include: a single patient consent now covers all future uses and disclosures for treatment, payment, and healthcare operations (TPO), replacing the previous per-disclosure consent requirement; covered entities and business associates receiving Part 2 records under TPO consent may redisclose in accordance with HIPAA, without separate patient authorization; record segregation is no longer required for entities receiving Part 2 records under a single TPO consent; and enforcement authority has shifted from DOJ-only criminal penalties to OCR civil and criminal enforcement aligned with the HIPAA penalty framework.

Despite this significant alignment, important Part 2 barriers specific to behavioral health AI tools remain. The Final Rule created a new, highly protected category called "SUD Counseling Notes" — defined as a clinician's notes documenting or analyzing the content of a SUD counseling session, maintained separately from the treatment record — analogous to HIPAA's psychotherapy notes protections. SUD Counseling Notes require a separate patient consent and cannot be covered by a broad TPO consent. This new category is directly relevant to ambient AI documentation tools: when a system like Eleos Health captures and analyzes a SUD counseling session, the resulting AI-generated session analysis may constitute SUD Counseling Notes under the new definition, requiring consent workflows that current ambient AI implementations were not designed to accommodate. HHS has not yet issued guidance clarifying how ambient AI documentation systems should navigate this new category, leaving SUD treatment programs and their AI vendors in legal uncertainty.

Furthermore, Part 2 does not preempt more protective state laws, and the prohibition on using SUD records in civil, criminal, administrative, or legislative proceedings against a patient remains absolute without separate written consent or a court order — a protection that goes well beyond HIPAA and has direct implications for AI-generated SUD documentation that could be discoverable in legal proceedings. HHS should issue specific guidance addressing how ambient AI documentation systems operating in SUD treatment settings should implement the SUD Counseling Notes consent requirements, and should clarify the conditions under which AI-generated session analysis constitutes Part 2-protected records. Because some ambient listening documentation assistants require retention of audio recordings and/or transcripts of the therapy session to analyze the session, HHS should declare these recordings as non-discoverable.

A foundational barrier to AI adoption and interoperability in Community Mental Health Centers (CMHCs) serving Medicaid populations is the absence of standardized procedure coding and service definitions across state lines. While HCPCS Level II and CPT codes provide a nominal federal framework, Medicaid is a state-administered program, and individual states retain broad latitude to define covered services, establish documentation requirements, specify unit structures, and determine reimbursement rates independently. The practical consequence is that the same billing code — for example, an H-code for community-based mental health services — may carry materially different service definitions, staff qualification requirements, and clinical documentation expectations depending on the state and even the managed care payer within that state. For AI-powered clinical documentation and coding assistance tools to function reliably in CMHC environments, they must be trained, configured, and continuously updated to accommodate this patchwork of state- and payer-specific rules. This variability substantially increases the cost and complexity of AI implementation, limits the scalability of vendor solutions across state markets, and creates persistent risk of claim denials, audit exposure, and compliance failures when AI-generated outputs do not precisely align with jurisdictionally-specific requirements. HHS should consider whether federal guidance, technical assistance, or Medicaid waiver incentives could encourage greater harmonization of behavioral health service definitions and coding standards as a prerequisite for effective AI adoption in this sector.

Lack of Behavioral Health-Specific AI Training Data and Validation

General-purpose large language models (LLMs) like ChatGPT, Gemini and others perform poorly on behavioral health-specific tasks without domain-specific fine-tuning. As Dr. Torous noted in Congressional testimony (November 2025), one patient's apt observation captures the challenge: general AI tools tend to repeat the same responses 'in different ways' without demonstrating genuine clinical insight. These tools rely on sycophancy to build rapport. Lyssn, which has built its platform on 17 years of academic research and analysis of more than 4.3 million clinical sessions, and has produced over 60 peer-reviewed publications validating its approach, represents the kind of domain-specific foundation behavioral health AI requires—but that standard AI development pipelines do not automatically provide.

Eleos Health has pursued a parallel but distinct data strategy — building what it describes as the industry's largest dataset of real-world behavioral health treatment sessions to train its ambient documentation and clinical insight AI. A peer-reviewed study published in the *Journal of Medical Internet Research* reports that Eleos's machine learning research drew on 2.83 million therapist-client micro-dialogues — a dataset the lead researcher described as "unprecedented," specifically because it captures how therapy is actually delivered in real-world clinical settings rather than in the controlled conditions that have historically characterized psychotherapy research. Eleos's in-house clinical team, with over 180 combined years of behavioral health experience, continuously tests, validates, and fine-tunes its models to recognize the clinical terminology, interaction patterns, and emotional indicators that a general-purpose LLM trained on broad internet text cannot reliably detect. The platform's coverage extends across CBT, DBT, ACT, Motivational Interviewing, and the Zero Suicide Framework, with active research studies conducted in collaboration with Stanford, MIT Sloan, Vanderbilt, Harvard, Palo Alto University, Imperial College London, and Hebrew University — including novel AI-based detection of social determinants of health discussions in session content and multi-modal audio feature analysis as indicators of depression and anxiety symptom severity. Critically, this dataset was assembled under a strict privacy architecture: Eleos does not store raw session recordings or use protected health information to train its models; any data used for model improvement is fully de-identified and reviewed exclusively by a clinically licensed, security-cleared team. Together, the Lyssn and Eleos datasets represent the most significant behavioral health-specific AI training foundations in existence, and the research programs surrounding them demonstrate precisely the kind of domain-specific investment that HHS should incentivize, protect, and replicate through federal research funding.

The scarcity of large, well-labeled, consent-appropriate behavioral health datasets for AI training—particularly for specialized modalities like group therapy, peer support, and crisis intervention—limits the quality ceiling for behavioral health-specific AI.

Reimbursement Gaps

Current CMS reimbursement frameworks do not recognize AI-assisted clinical documentation or AI-enhanced quality improvement as billable services. Behavioral health organizations cannot

claim additional reimbursement for deploying AI tools that improve documentation integrity, enhance evidence-based practice fidelity, or reduce clinician turnover—even when these tools demonstrably improve patient outcomes and organizational financial sustainability. The result is that the upfront cost of AI tools falls entirely on organizations that are already operating on thin margins.

Question 2: What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.

42 CFR Part 2 Safe Harbor for AI Documentation Tools

As described in the “Regulatory Fragmentation” section above, HIPAA and 42 CFR § 2 alignment has helped providers’ significantly. However, issues remain especially regarding SUD Counseling Notes. HHS should take two targeted actions. First, issue authoritative interpretive guidance clarifying: (a) whether AI-generated session analysis in an ambient documentation tool constitutes SUD Counseling Notes under 42 CFR § 2.11; (b) the conditions under which AI vendor processing of session content constitutes a use or disclosure requiring Part 2 consent; and (c) what a compliant “separate consent” workflow for SUD Counseling Notes looks like in an ambient AI implementation. Second, establish a prospective safe harbor for AI documentation tools that process SUD session content within the treating provider’s care setting, under a defined consent architecture, where the AI-generated output remains within the treating provider’s EHR and is not disclosed to external parties — on the grounds that such use is analytically equivalent to a clinician’s own note-taking rather than a disclosure to a third party. Applicable CFR citations: 42 CFR §§ 2.11 (SUD Counseling Notes definition), 2.31 (consent requirements), 2.32 (notice accompanying disclosure).

CCBHC Prospective Payment System Recognition of AI-Enhanced Services

CMS should modify the CCBHC Prospective Payment System (PPS) methodology under the CCBHC Medicaid Demonstration program to recognize AI-assisted documentation and quality improvement tools as qualifying infrastructure investments that affect the cost report used to establish PPS rates. Under the current framework, investments in AI tools that reduce documentation time and improve compliance do not result in sustainable revenue adjustments, because their value accrues primarily through avoided costs (reduced denials, improved staff retention) rather than through billable activity. Applicable CFR citation: 42 CFR § 438.6 (managed care CCBHC provisions); 42 CFR Part 447 (payment for services).

Medicare Behavioral Health Integration and Digital Health Codes

CMS should expand CPT code coverage for AI-augmented behavioral health services. The existing Behavioral Health Integration (BHI) codes (99492-99494) and General BHI (99484) should be updated or supplemented to recognize AI-assisted care coordination, digital

phenotyping-informed treatment adjustment, and AI-enabled evidence-based practice supervision. Additionally, CMS should establish new HCPCS codes specifically for AI-assisted clinical documentation in behavioral health settings, analogous to how remote patient monitoring codes (CPT 99453-99458) created a reimbursement pathway for technology-mediated monitoring. Applicable CFR citation: 42 CFR § 410.67 (behavioral health integration services).

ONC Behavioral Health EHR Certification Criteria

ASTP/ONC should establish behavioral health-specific criteria within the Health IT Certification Program (45 CFR Part 170) that include requirements for AI transparency, bias testing on behavioral health populations, and interoperability with behavioral health data standards including SNOMED CT mental health subsets and the Gravity Project's social determinants of health (SDOH) terminology. Current certification criteria were developed primarily for acute and primary care workflows and do not adequately address behavioral health-specific documentation structures, group therapy notes, or the unique data elements relevant to SMI populations.

Question 3: For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (e.g., relating to liability, indemnification, privacy, and security). What novel legal and implementation issues exist and what role, if any, should HHS play to help address them?

In general, there are currently two broad uses of AI in behavioral healthcare: the first is the use of AI to expedite and improve a clinician's ability to complete the onerous documentation burden seen in behavioral healthcare, especially community based behavioral healthcare. When used in this manner, clinicians are being provided information to help them write clinical notes. They and they alone are the arbiters of what, if any of the suggestions provided by the AI are used in the clinical note. The second is AI being deployed without clinician oversight. AI chatbots are the most common example of this form of AI usage in behavioral health care. The fundamental difference between these two uses is that the former assists clinicians in their work but does not replace clinical judgement. The latter has no human clinical oversight.

Liability in AI-Augmented Clinical Documentation

Ambient AI documentation tools like Eleos Health's platform generate progress note suggestions that clinicians review, modify, incorporate into the EHR note and sign. For some, this has created confusion regarding liability questions: if an AI-generated note contains a clinical inaccuracy that contributes to a patient harm event, how should liability be apportioned among the parties that participated in the creation of that including the AI vendor, the EHR vendor, the clinician who approved the note, and the organization? Since these tools assist professionals only in documenting clinical interactions, current malpractice frameworks appropriately assume human authorship and human accountability for clinical documentation. HHS should reaffirm this position and work with state insurance commissioners and CMS to develop guidance clarifying that the supervising clinician retains professional responsibility for reviewed and approved AI-generated content, while establishing minimum accountability standards for AI vendors providing documentation assistance including transparency regarding model limitations, error rates, and update protocols. Some AI products allow or even require that transcripts and/or audio recordings of the clinical session and the recommendations made but not accepted by the clinician be retained for some period of time. These recordings and transcripts are similar to other quality improvement documents such as chart reviews and supervision and should be protected from legal discovery.

AI Mental Health Chatbots and Crisis Situations

Compared to AI tools that assist clinicians in documenting clinical transactions, mental health chatbots provide significantly more serious risks to end users. These products work semi-autonomously often without a "human in the loop". Dr. John Torous and others have documented

extensively the risks of AI chatbots deployed for mental health support without adequate safety guardrails. As he noted in Congressional testimony, OpenAI reported over one million users per week engaging ChatGPT with explicit indicators of suicidal planning, and numerous lawsuits have alleged AI chatbot involvement in patient suicides. The legal status of AI chatbots that provide mental health guidance but are not classified as medical devices creates a significant accountability gap. HHS should issue guidance establishing minimum safety requirements—including mandatory crisis detection, escalation protocols, and handoff to human crisis services—for AI applications marketed for mental health support, regardless of FDA classification. This would build on Dr. Torous's work benchmarking AI chatbot performance on mental health scenarios through platforms like [mindApps.ai](https://mindapps.ai). See “Real World Performance Surveillance” below for more on this topic.

Privacy in Psychotherapy and Group Therapy AI

Lyssn's AI platform analyzes session recordings for quality improvement and evidence-based practice fidelity—a valuable capability that raises novel privacy questions when applied to group therapy, where multiple clients simultaneously disclose protected health information. Unlike individual therapy, group therapy transcripts contain the PHI of potentially 10-15 individuals per session. However, like individual notes, the information not used to create the clinical notes for group therapy participants should be regarded as legally non-discoverable.

Informed Consent Standards for AI-Assisted Care

No federal standard currently exists for what behavioral health patients must be told when AI tools are used in their care. Dr. Torous and others have long advocated for transparent disclosure to patients about digital health tool use. There are two use cases for consent regarding the use of AI tools. The first is general consent which is typically given when treatment begins along with other, more traditional consents to treatment. The second is specific consent, or more exactly, the ability for a patient to withdraw previously given consent for one or more treatment sessions. AI vendors must accommodate both. HHS should establish minimum informed consent standards for behavioral health settings that address: the use of ambient listening technologies during sessions, AI analysis of session content for documentation or quality improvement, and the use of AI chatbots or digital therapeutics as adjuncts to care.

Question 4: For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and other workflow and human-centered evaluation methods for clinical care? Should HHS further support these processes? If so, which mechanisms would be most impactful (*e.g.*, contracts, grants, cooperative agreements, and/or prize competitions)?

Evidence-Based Practice Fidelity as a Core Metric

For behavioral health AI, the most clinically meaningful evaluation metric is fidelity to evidence-based practices (EBPs)—the degree to which AI-assisted clinical encounters actually improve adherence to validated treatment protocols like Motivational Interviewing (MI), Cognitive Behavioral Therapy (CBT), and Dialectical Behavioral Therapy (DBT). Lyssn's platform, built on 17 years of academic research and validated against more than 4.3 million analyzed sessions, represents the state of the art in AI-powered EBP fidelity measurement. Its approach—using natural language processing to assess therapist behaviors against validated coding systems like the Motivational Interviewing Treatment Integrity (MITI) scale—provides a model for how AI evaluation in behavioral health can be clinically grounded rather than purely technical. Using EBPs in this way presumes that managing the process *i.e.* the protocol will result in the same (or similar) outcomes achieved by the EBP when that protocol was validated. It is important to note that doing an EBP is a means to an end, not the end itself. EBPs exist to provide clinicians with a research-based way to get good clinical outcomes. They define what improvement is possible for various treatments, in a way defining a gold standard. In cases where delivering high fidelity EBPs is not possible or advisable, the outcomes demonstrated by the EBPs provide a benchmark for non-EBP interventions. That is, if a non-EBP intervention can achieve the same or better outcomes as demonstrated with the EBP for the same or less cost, the non-EBP intervention is achieving what the EBP treatment did even though it was not maintaining fidelity to the protocol. In this way, EBPs serve dual purposes: providing a detailed protocol for clinicians to follow and a clinical outcome benchmark that non-EBPs should strive to attain.

Therapeutic Alliance Preservation

The therapeutic relationship—the alliance between clinician and client—is one of the most powerful predictors of behavioral health treatment outcomes across all modalities. AI evaluation frameworks must assess whether AI tools enhance or diminish therapeutic alliance quality. Eleos Health's published research and received numerous anecdotal observations demonstrating that providers using its ambient documentation platform achieve 50% higher client engagement, 3-4x better symptom improvement, and 36% greater use of evidence-based techniques compared to baseline, suggesting that removing the documentation burden allows clinicians to be more fully present with clients. Post-deployment evaluation should systematically measure therapeutic alliance using validated tools such as the Working Alliance Inventory.

Bias and Health Equity Testing

AI systems trained primarily on data from majority populations can perform poorly for underrepresented groups—a particular concern in behavioral health, where racial and ethnic minorities, LGBTQ+ individuals, and people with limited English proficiency face well-documented care disparities. Dr. Torous's Digital Navigator program at BIDMC specifically addresses digital equity as a prerequisite for beneficial AI deployment. HHS should require that behavioral health AI tools undergo pre-deployment bias testing disaggregated by race, ethnicity, language, gender identity, and socioeconomic status, with post-deployment monitoring to detect disparate performance. Grant mechanisms through SAMHSA and NIMH would be the most appropriate vehicles for funding this research.

Real-World Performance Surveillance

Dr. Torous has consistently argued that most mental health apps and AI chatbots remain understudied, with unknown real-world risks — a position now supported by a robust independent evidence base. Linardon and colleagues conducted a systematic review of 171 clinical trials of mental health apps published in *npj Digital Medicine* (2024), finding that only 55 trials reported adverse events at all, the meta-analytic deterioration rate across app conditions was 6.7%, and those trials that did report adverse events frequently provided insufficient information to allow proper assessment of user risk. Writing in *World Psychiatry* (2025), the same research team documented that negative effects from apps can extend to symptom deterioration, onset of new symptoms, and suicidal ideation, with rates of negative effects in severe mental illness populations reaching as high as 20% in some trials — a state of evidence that, in their words, "makes it difficult to integrate apps into clinical practice safely."

The crisis response failure problem has been quantified independently. Dwyer and colleagues publishing in *Psychiatric Services* (2024), evaluated 302 mental health apps available in January 2024 using the APA's app evaluation model and found that only 15% referred users to 988, the national crisis standard, and — most strikingly — 14 apps with more than 3.5 million combined downloads provided incorrect or nonfunctional alternative crisis hotlines. This is not a theoretical risk: millions of people in potential crisis were being routed to telephone numbers that do not connect. This finding directly extends foundational work by Larsen, Nicholas, and Christensen, whose systematic review of suicide-focused apps found that none abided with best practices for providing visible crisis support information, and that the vast majority demonstrated limited ability to recognize anticipatory warning signs of suicidal crisis.

For behavioral health AI, HHS should support the development of post-deployment surveillance infrastructure analogous to FDA's MedWatch system—a mechanism for clinicians and patients to report AI-related clinical errors, safety events, and unexpected behaviors in behavioral health settings. AHRQ's existing patient safety infrastructure could serve as a model.

Question 5: How can HHS best support private sector activities (e.g., accreditation, certification, industry-driven testing, and credentialing) to promote innovative and effective AI use in clinical care?

Behavioral Health-Specific AI Certification

HHS should support the development of a behavioral health-specific AI certification program, potentially through SAMHSA or in partnership with the National Council for Mental Wellbeing. Such a program would extend beyond ONC's general health IT certification criteria to address behavioral health-specific requirements including EBP fidelity measurement, crisis protocol adequacy, 42 CFR Part 2 compliance for SUD applications, and therapeutic alliance preservation. Eleos Health's recent achievement of ISO 42001 certification (the AI-specific management system standard) demonstrates that behavioral health AI vendors are ready for rigorous certification frameworks.

SAMHSA Technology Assessment Program

SAMHSA should establish a technology assessment program analogous to the Substance Abuse Prevention and Treatment Block Grant's evidence-based practices mandate, creating a tiered recognition system for behavioral health AI tools based on evidence strength. Dr. Torous, who led the American Psychiatric Association's work group on smartphone app evaluation and developed the APA's technology evaluation framework, has created models for how this can be structured. His mindapps.org database—the largest mental health app registry—and mindApps.ai benchmarking platform provide a template for systematic, ongoing behavioral health AI evaluation that HHS could support and formalize. Other partners in this effort include The Joint Commission and CARF, two of the most used accrediting bodies in the behavioral health industry.

Academic-Industry Research Partnerships

HHS should incentivize partnerships between behavioral health AI companies and academic medical centers with established digital psychiatry programs. The BIDMC Digital Psychiatry Division's research collaborations—including studies validating Eleos Health's impact on documentation burden and clinical outcomes—demonstrate how these partnerships can produce rigorous peer-reviewed evidence that benefits the entire field. Cooperative research and development agreements (CRADAs) between NIMH-funded research centers and behavioral health AI companies would accelerate the evidence base. Eleos Health has established multiple academic partnerships with Stanford, Palo Alto University, Hebrew University and others.

Question 6: Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?

Successes: Ambient Documentation in Behavioral Health

AI-assisted clinical documentation has clearly met or exceeded performance expectations in behavioral health settings. Eleos Health's deployments across organizations including Wellpower, Jefferson Center, GRAND Mental Health, and Gaudenzia demonstrate consistent documentation time reductions exceeding 70%, with providers accepting 80% of AI-generated note suggestions. It should be noted that Eleos Health believes this 80% mark is a best practice. That is, no more than 80% of the note should be AI-generated since all notes should also include the unique voice of the clinician. These results translate directly into organizational sustainability: 90% of notes submitted within 24 hours reduces claim denials, and reduced documentation burden improves staff retention in a sector experiencing severe workforce shortages. Dr. Brandon Ward, when he was CIO at Jefferson Center, reported that providers are enthusiastic about the technology, including initial skeptics. These results are particularly significant given that behavioral health organizations serve some of the most complex patients in the healthcare system.

Successes: AI-Powered Quality Improvement and EBP Fidelity

Lyssn's AI platform has demonstrated success in scaling evidence-based practice training and quality improvement in ways that were previously impossible at the organizational level. Traditional EBP fidelity assessment required manual coding of session recordings—a process so resource-intensive that most organizations could only evaluate a tiny fraction (5-10%) of clinical encounters. Lyssn's AI enables continuous, automated fidelity monitoring across entire workforces, allowing organizations to identify training needs, target coaching resources, and demonstrate fidelity to funders and accreditors. This addresses a longstanding gap between what EBP research says is effective and what actually occurs in real-world clinical practice.

Shortfalls: AI Chatbots for Direct Mental Health Support

AI chatbots deployed directly to patients for mental health support have not met clinical expectations and present documented safety risks. Dr. Torous, whose early research was foundational in assessing the potential of digital mental health tools, has documented that most mental health chatbots remain untested in clinical populations and present unknown risks. The 2017 emergence of Woebot and similar chatbots was accompanied by enthusiasm that has been tempered by subsequent evidence: few have undergone randomized controlled trials, many fail to recognize crisis situations, and some have been associated with patient harm. The FDA's April 2024 approval of the first app to treat depression (Rejoyn for adjunctive treatment of MDD) represents a positive regulatory development, but the vast majority of the thousands of mental

health apps on the market remain unvalidated. NAMI does not endorse AI for mental health treatment and has partnered with Dr. Torous to develop independent benchmarks for AI tool safety.

Shortfalls: Interoperability and EHR Integration

AI tools in behavioral health consistently struggle with EHR integration. The behavioral health EHR market is highly fragmented, with dozens of specialized platforms (Netsmart's myAvatar and Streamline, Qualifacts/Credible, Kipu, ICANotes, and others) lacking the standardized APIs that acute care EHRs have developed under ONC's information blocking rules. Eleos Health addresses this by operating as an EHR-agnostic overlay via browser extension—a practical solution, but one that highlights the underlying interoperability problem rather than solving it. True integration requires standardized behavioral health data exchange frameworks that do not currently exist.

Question 7: Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?

Key Decision Makers

In community behavioral health organizations, AI adoption decisions typically require alignment among: the Chief Executive Officer (organizational mission and financial sustainability), the Chief Clinical Officer (clinical appropriateness, staff engagement, evidence base), the Chief Information Officer (technical infrastructure, security, EHR integration), and the Chief Medical Officer (psychiatrist or physician oversight of AI-generated clinical content). Unlike hospital systems, most CBHOs lack a dedicated Chief Digital Officer or Chief Innovation Officer, meaning AI advocacy must come from within existing leadership roles. Board-level engagement is increasingly common as AI represents both opportunity and liability.

Eleos Health's deployment experience across organizations like Jefferson Center (where the Chief Innovation Officer and President/CEO both participated actively in adoption decisions) illustrates that multi-stakeholder buy-in is essential—and that mission alignment and clinician enthusiasm are as important as technical capability.

Primary Administrative Hurdles

The most consistently reported administrative hurdles to behavioral health AI adoption include: (1) procurement complexity, including legal review of BAAs for ambient listening technologies and vendor security assessments that small organizations lack capacity to conduct; (2) state licensing and consent requirements for AI tools that capture session content, which vary widely and are changing frequently and quickly and are rarely addressed in existing technology guidance; (3) payer contracting constraints, as some Medicaid managed care contracts include provisions that could be interpreted to restrict AI-generated documentation; (4) workforce training requirements, as even well-designed AI tools require change management investment that organizations with limited training budgets struggle to fund; and (5) skepticism from clinicians who have experienced technology-driven increases in administrative burden and are appropriately cautious about new tools promising relief.

Question 8: Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

Behavioral Health-Specific Data Standards

The most significant interoperability gap affecting behavioral health AI is the absence of standardized behavioral health data exchange frameworks. The Meaningful Use standards first implemented with the passage of the Affordable Care Act started the process of interoperability, but behavioral health care organizations were immediately at a disadvantage since they were exempted from the EHR stimulus incentives available to medical care providers. Many Health Information Exchanges (HIEs) were hesitant to accept and coordinate behavioral health data because of the challenges of 42CFR Part 2 and a general hesitancy about the confidentiality requirements inherent in behavioral health information. Now, while HL7 FHIR has transformed data exchange in acute and primary care, behavioral health data—including mental status examinations, symptom rating scales, substance use assessment tools, and session notes—lacks standardized FHIR resource profiles. SAMHSA and ONC should jointly develop and mandate FHIR implementation guides specifically for behavioral health, including: standardized profiles for PHQ-9, GAD-7, Columbia Suicide Severity Rating Scale (C-SSRS), AUDIT, DAST, and other validated assessment instruments; FHIR resources for group therapy documentation; and behavioral health-specific social determinants of health data elements aligned with the Gravity Project.

Digital Phenotyping Data Standards

Dr. Torous's mindLAMP platform at BIDMC represents a leading approach to digital phenotyping—using passive smartphone sensor data (GPS, accelerometry, screen time, communication patterns) and active ecological momentary assessment to capture the lived experience of individuals with mental illness between clinical encounters. Standardized data models for digital phenotyping would enable AI systems to incorporate real-world behavioral data into clinical decision support, treatment planning, and relapse prediction. NIH's All of Us research program, in which Dr. Torous serves as an advisor for the smartphone mood study, provides a model for how such data could be collected at population scale.

Crisis and Emergency Data Exchange

Behavioral health AI crisis tools—including suicide risk prediction algorithms and crisis chatbot escalation systems—require seamless data exchange with crisis services, emergency departments, and 988 Suicide and Crisis Lifeline infrastructure. Currently, no standardized data exchange protocols exist for transmitting risk-relevant information between AI-powered behavioral health applications and crisis response systems. The 988 infrastructure buildout

represents an opportunity to establish behavioral health AI crisis data standards that could benefit both direct patient-facing AI tools and clinician-facing decision support systems.

Question 9: What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?

Desired Applications

Patients and caregivers in behavioral health settings consistently express desire for AI applications that: increase access to care (particularly for those in underserved rural and urban areas); reduce administrative friction (appointment scheduling, medication refill requests, prior authorization navigation); extend the reach of therapeutic support between sessions through evidence-based digital tools; and improve the quality and attentiveness of clinical encounters by reducing the time their clinician spends looking at a screen. The last point is particularly salient—patients notice when their clinician is distracted by documentation. Eleos Health's research and clinical testimonials frequently highlight that clients receive more engaged, screen-free care when their providers use ambient documentation AI.

Dr. Torous's research on patient perspectives on digital mental health, including his widely-cited guides to evaluating mental health apps, consistently finds that patients want tools that are evidence-based, privacy-respecting, and clinically supervised. The patient experience of AI chatbots varies widely—some find them accessible and destigmatizing, while others find them superficial and inadequate for their needs.

Patient Concerns

Patient concerns about AI in behavioral health are concentrated in three areas. First, privacy: psychotherapy is one of the most sensitive data categories, and patients fear that AI-analyzed session content could be used for insurance underwriting, employment decisions, or law enforcement purposes. The history of stigma and discrimination against people with mental illness and substance use disorders makes these fears reasonable and well-founded. Second, depersonalization: patients fear that AI will replace the human therapeutic relationship with algorithmic interaction, reducing their care to a transaction with a machine. Third, accuracy: patients worry that AI-generated documentation may misrepresent what was said or decided in a session, affecting their treatment or generating records they cannot correct.

NAMI's partnership with Dr. Torous to develop independent AI safety benchmarks—explicitly centered on lived experience perspectives—represents a model for how patient voices can be systematically incorporated into behavioral health AI evaluation. HHS should require lived experience representation on any federal advisory bodies addressing behavioral health AI.

Question 10: Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care?

a. Are there published findings about the impact of adopted AI tools and their use in clinical care?

b. How does the literature approach the costs, benefits, and transfers of using AI as part of clinical care?

10a: Published Findings on Adopted AI Tools

The behavioral health AI evidence base is growing but remains insufficient relative to the pace of deployment. Key published findings include:

- Eleos Health has produced 9 peer-reviewed studies on its ambient AI documentation platform, with more than a dozen in progress, including preliminary randomized controlled trial data suggesting improved mental health outcomes for patients whose clinicians use the platform. Findings include 50% higher client engagement, 3-4x better symptom improvement, and 36% greater use of evidence-based techniques.
- Lyssn's platform is grounded in 60+ peer-reviewed academic publications developed over 17 years of research, validating its NLP-based approach to assessing Motivational Interviewing fidelity and other EBPs. The platform has analyzed more than 4.3 million sessions, providing an unparalleled evidence base for behavioral health AI quality measurement.
- Dr. Torous and colleagues published a comprehensive review in *World Psychiatry* (2025) on the evolving field of digital mental health, covering smartphone apps, generative AI, chatbots, and virtual reality, concluding that while evidence for some modalities is accumulating, most mental health chatbots remain undertested. His team's 50-state legislative review (JMIR Mental Health, 2025) provides the most comprehensive analysis of how states are regulating AI in mental health.
- The FDA's April 2024 approval of Rejoyn (Otsuka and Alto Neuroscience's prescription digital therapeutic) for adjunctive treatment of major depressive disorder represents a regulatory milestone, demonstrating that behavioral health AI can meet the FDA's evidence standards for medical device clearance when rigorously evaluated.

10b: Costs, Benefits, and Transfers

The behavioral health AI literature approaches cost-benefit analysis with increasing sophistication, though significant gaps remain:

- Documentation AI demonstrates clear return on investment through reduced clinician hours spent on notes (>70% reduction), improved claim submission timeliness (90% of notes within 24 hours with Eleos Health), reduced late/rejected claims, and

improved staff retention. Given that behavioral health clinician turnover costs organizations an estimated \$15,000-\$30,000 per vacancy, retention benefits alone can justify AI investment within months.

- Quality improvement AI (Lyssn) generates value through improved EBP fidelity, which correlates with better patient outcomes and reduced treatment duration—ultimately reducing the per-episode cost of care. The ability to evaluate 100% of clinical encounters rather than a 5-10% sample also reduces audit risk and accreditation vulnerability.
- AI chatbots and direct-to-consumer mental health AI present a more complex cost-benefit picture. While they offer potential cost savings through reduced clinician time, Dr. Torous and colleagues have documented risks including failure to recognize crisis situations, inadequate evidence bases, and the possibility of displacing rather than supplementing human care. The costs of adverse events—including patient harm from inadequate crisis response—are difficult to quantify but potentially catastrophic.
- A critical and underappreciated transfer dynamic is the reimbursement cliff: behavioral health organizations bear the full cost of AI tools while payers capture most of the financial benefit through reduced claims (fewer denials, improved documentation integrity). CMS and state Medicaid agencies should develop value-sharing mechanisms that ensure AI-adopting providers capture a fair share of the financial value they generate.

Priority Research Investments for HHS

Based on the above analysis, HHS—through NIMH, SAMHSA, AHRQ, and NIH—should prioritize the following behavioral health AI research investments:

- Research support for AI use in crisis assessment and triage including AI tools to improve suicidal risk identification and prevention.
- Rigorous clinical trials of AI ambient documentation tools measuring both provider outcomes (burnout, satisfaction, retention) and patient outcomes (therapeutic alliance, symptom improvement, treatment completion rates).
- Development and validation of standardized behavioral health AI evaluation frameworks, building on Dr. Torous's APA technology evaluation framework and Lyssn's evidence-based practice fidelity measurement approaches.
- Safety research on AI chatbots and digital mental health tools, including prospective studies of adverse event rates, crisis response adequacy, and the impact of AI tool use on suicidal ideation and self-harm outcomes.
- Digital equity research examining differential outcomes of behavioral health AI deployment across racial, ethnic, linguistic, and socioeconomic groups, addressing the disparities documented in Dr. Torous's Digital Navigator program research.
- Implementation science research on AI adoption in CCBHCs and CMHCs specifically, examining the organizational, workforce, and regulatory factors that predict successful adoption in under-resourced safety-net settings.

Citations and References

Eleos Health. (2026). Automation and AI Solutions for Behavioral Health. <https://eleos.health>. Accessed February 2026. [Source for ambient documentation statistics: 70%+ documentation time reduction, 80% provider acceptance rate, 90% notes within 24 hours, 2x client engagement, 3-4x symptom improvement, 36% greater EBP use, 9 peer-reviewed studies.]

Lyssn. (2026). Training & Quality Improvement AI Solutions for Health & Human Services. <https://www.lyssn.io>. Accessed February 2026. [Source for: 17 years academic research foundation, 4.3M+ analyzed sessions, 60+ peer-reviewed publications, NLP-based EBP fidelity measurement, CCBHC and crisis care applications.]

Torous, J., et al. (2025). The evolving field of digital mental health: current evidence and implementation issues for smartphone apps, generative artificial intelligence, and virtual reality. *World Psychiatry*, 24(2), 156-174. doi:10.1002/wps.21299.

Torous, J. (2025, November 18). Congressional Testimony: Governing AI in Mental Health. U.S. House of Representatives, Committee on Energy and Commerce. <https://www.congress.gov/119/meeting/house/118669/witnesses/HHRG-119-IF02-Bio-TorousMDMBIJ-20251118.pdf>

Torous, J., et al. (2025). Governing AI in Mental Health: 50-State Legislative Review. *JMIR Mental Health*, 12, e80739. [Source for state-level AI regulation analysis.]

Torous, J. (Director, Digital Psychiatry Division, BIDMC/Harvard Medical School). Profile and research overview. <https://www.johntorousmd.com>. [Source for mindapps.org, mindLAMP, mindApps.ai, Digital Navigator program, APA technology evaluation framework.]

National Alliance on Mental Illness (NAMI). (December 2025/January 2026). NAMI Takes the Lead to Push for Clarity and Safety in AI Mental Health Tools. Partnership with Dr. John Torous, BIDMC. <https://www.nami.org/press-releases/nami-takes-the-lead-to-push-for-clarity-and-safety-in-ai-mental-health-tools/>

NAMI. (2025). AI and Mental Health — FAQ with Dr. John Torous. <https://www.nami.org/about-mental-illness/ai-and-mental-health/> [Source for NAMI's position that AI is not a substitute for mental health treatment.]

Torous, J. (2017, September 25). 'I'm sorry to hear that': Why training Siri to be a therapist won't be easy. [Media coverage of early AI/NLP chatbot limitations for mental health applications.]

Torous, J. (2017, December 3). 'The Woebot will see you now' — the rise of chatbot therapy. [Early analysis of AI chatbot mental health applications and evidence gaps.]

Torous, J. (2017, July 1). How Clinicians Evaluate Mental Health Apps. [Foundation for APA app evaluation framework.]

Torous, J. (2016, December 6). Smartphone health apps often miss real medical emergencies. [Source for AI/app safety concerns, failure to recognize crisis situations.]

Torous, J. (2017, December 21). Why choosing a mental health app is harder than you think. [Source for app quality and safety disparities.]

FDA. (April 4, 2024). The FDA has approved the first app to help treat depression. [Regulatory milestone for behavioral health digital therapeutics — Rejoyn approval.]

Torous, J. (2025). Digital Psychiatry Division, BIDMC. Academic Profile. <https://dci.bidmc.org/people/john-torous> [Source for APA technology evaluation leadership, All of Us smartphone mood study advisory role.]

HHS. (December 23, 2025). Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care. 90 FR 60108. Document Number 2025-23641.

Sadeh-Sharvit, S., & Hollon, S.D. (2020). Leveraging the Power of Nondisruptive Technologies to Optimize Mental Health Treatment: Case Study. *JMIR Mental Health*, 7(11), e20646. [Source for Eleos AI platform clinical validation,

2.83 million micro-dialogue dataset, and published RCT/retrospective study program.]

<https://pubmed.ncbi.nlm.nih.gov/33242025/>

Linardon, J., Fuller-Tyszkiewicz, M., Firth, J., Goldberg, S.B., Anderson, C., McClure, Z., & Torous, J. (2024). Systematic review and meta-analysis of adverse events in clinical trials of mental health apps. *npj Digital Medicine*, 7(1), Article 363. doi:10.1038/s41746-024-01388-y

Torous, J., Linardon, J., Goldberg, S.B., et al. (2025). The evolving field of digital mental health: current evidence and implementation issues for smartphone apps, generative artificial intelligence, and virtual reality. *World Psychiatry*, 24(2), 156–174. doi:10.1002/wps.21299 [also cited elsewhere in this RFI response]

Dwyer, B., et al. (2024). *Mental Health Apps and Crisis Support: Exploring the Impact of 988. Psychiatric Services*. doi:10.1176/appi.ps.20240485 [Lead author affiliated with BIDMC; published August 2024]

Larsen, M.E., Nicholas, J., & Christensen, H. (2016). A systematic assessment of smartphone tools for suicide prevention. *PLOS ONE*, 11(4), e0152285. doi:10.1371/journal.pone.0152285