

Supplemental Closing Statement

To: RFI Review Committee

From: David Bynon

Date: February 23, 2026

Subject: Supplemental Closing Comments – HHS-ONC-2026-0001

Executive Summary

In reviewing the volume and tenor of recent comments, it is evident that many stakeholders perceive a tension between responsible AI oversight and administrative burden. I respectfully submit that this tension arises not from oversight itself, but from ambiguity in how policy meaning and validation are structured.

The path to efficiency is not reduction of governance—but architectural refinement of how governance is implemented.

On Oversight and Operational Load

Oversight becomes burdensome when compliance is interpretive, manual, and reactive. It becomes efficient when policy meaning is explicit, machine-readable, and auditable by design.

As outlined in my initial submission, a deterministic semantic framework allows AI systems to retrieve from authoritative, versioned policy memory rather than infer from narrative text. This approach shifts compliance from after-the-fact review toward structured, audit-ready outputs.

Clarity reduces friction. Structure reduces interpretive drift. Deterministic retrieval reduces downstream reconciliation.

On Cost-Efficiency and Innovation

Innovation is not constrained by documentation; it is constrained by uncertainty.

By aligning AI systems with structured, machine-verifiable policy representations, organizations can reduce reconciliation costs currently borne by CMS, OMHA, and

beneficiary support operations. When oversight is embedded at the data layer, validation becomes continuous rather than episodic.

Efficiency in AI governance is achieved through standardization of semantic responsibility—not through deregulation. In the context of ongoing government efficiency initiatives such as the Department of Government Efficiency (DOGE), clarity — not reduction — in policy semantics enables scalable oversight that aligns with both organizational efficiency goals and AI governance principles.

On Data Transparency and Accountability

Concerns regarding data sovereignty are understandable in an evolving AI environment. However, structured provenance, versioning, and machine-readable definitions protect both regulated entities and the public interest.

A transparent, pre-defined semantic layer creates mutual accountability while preserving proprietary implementation. It clarifies “what is policy” without exposing “how systems are built.”

Closing Observation

As this RFI closes, I encourage the committee to consider that the future of high-impact AI governance in healthcare will depend less on the volume of reporting and more on the precision of structured policy memory.

The federal challenge is not whether to govern AI—but how to do so in a way that reduces interpretive variance, minimizes audit friction, and preserves public trust in regulated domains such as Medicare.

Respectfully,
David W. Bynon

Ref:

- HHS-ONC-2026-0001-0010 (<https://www.regulations.gov/comment/HHS-ONC-2026-0001-0010>)
- HHS-ONC-2026-0001-0059 (<https://www.regulations.gov/comment/HHS-ONC-2026-0001-0059>)

