

Department of Health and Human Services (HHS)

Office of the Secretary
RIN 0955-AA13

Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care

Feb 23, 2026

Submitted to:

via FederalRegister.Gov



722 East Market Street, Suite 102-A9 | Leesburg, VA 20176
www.code360.io | (571) 443-9331

This proposal includes data that shall not be disclosed outside the Government and shall not be duplicated, used, or disclosed-in-whole or in part-for any purpose other than to evaluate this proposal. If, however, a contract is awarded to this offer or as a result of, or in connection with, the submission of this data, the Government shall have the right to duplicate, use, or disclose the data to the extent provided in the resulting contract. This restriction does not limit the Government's right to use information contained in this data if it is obtained from another source without restriction. The data contained on each page of this offer is subject to this restriction.



23 February 2026

Submission Method: Submitted electronically via Regulations.gov.

Department of Health and Human Services

Office of the Deputy Secretary and Assistant Secretary for Technology Policy (ASTP)

Office of the National Coordinator for Health Information Technology (ONC)

(collectively, ASTP/ONC)

Attention: ASTP/ONC Representatives

Subject: Comments to RIN 0955-AA13

Dear ASTP/ONC Representatives

About Code360:



Code360 Inc. is a digital IT services company supporting federal agencies, state and local governments, and commercial sector clients. We specialize in Health IT, Data Analytics, IT Governance, Advanced Statistics, AI/ML, Solution Architecture, Cloud Engineering, Software Modernization, and Business Intelligence. As an SBA-certified 8(a) and HUBZone minority owned small business, we are committed to delivering innovative, cutting-edge IT solutions that drive impact and transformation across the public sector.



Background:

On December 4, 2025, the Department of Health and Human Services (HHS) issued the HHS Artificial Intelligence (AI) Strategy ^[1] in response to Office of Management and Budget (OMB) Memorandum 25-21. ^[2] The HHS AI Strategy represents a “OneHHS” approach to AI and is primarily focused on the Department's internal interests and ambitions to harness the transformative potential of AI and embed it into HHS's operations, while upholding patient privacy, civil rights, and civil liberties. Consistent with the President's artificial intelligence (AI) Action Plan, ^[3] recent Executive Orders ^[4] on AI, and Office of Management and Budget AI memoranda, ^[5] the Department seeks public feedback on the actions it can take to establish a forward-leaning, industry-supportive, and secure approach to accelerate the adoption and use of AI ^[6] as part of clinical care. In the past 12 months, HHS Divisions, including the Food and Drug Administration (FDA), ^[7] National Institutes of Health (NIH), ^[8] Centers for Medicare & Medicaid Services (CMS), ^[9] and ASTP/ONC ^[10] have sought public feedback on various aspects of AI and how it intersects with the Department's policy interests. In this RFI we seek concrete, experience-based feedback from those building, buying, evaluating, using, and receiving care from AI tools that are part of clinical care as well as from those who wish to do so but face barriers. Public feedback will inform HHS-wide use of three different approaches: regulation, reimbursement, and research & development. In general, HHS seeks feedback on ways in which these approaches can be most effectively applied to support the rapid adoption and use of AI in clinical care, to foster public trust and confidence in modern technology solutions, to reduce uncertainty that impedes AI innovation, and to align federal incentives so that AI is deployed in ways that enhance productivity, reduce burden, lower health care costs, and improve health outcomes for patients, caregivers, and communities.

Section A:

Partner Response — Architecture and Business Systems Engineering for Safe AI Adoption in Clinical Care (Non-Device and Device-Adjacent)

A.1 Perspective and focus

This response is provided by a small business advisory focused on **AI-adoption architecture** (enterprise/business/technical) and **business systems engineering for AI**, with emphasis on **governed, auditable, human-centered operation** of AI capabilities in regulated environments. The offering is **advisory + architecture** (no product claims; no implementation commitments), with the intent to support primes and health-sector organizations as they operationalize HHS goals for rapid adoption that maintains public trust.

HHS is seeking “concrete, experience-based feedback” to accelerate AI in clinical care using **regulation, reimbursement, and R&D**, while fostering trust, reducing uncertainty, and aligning incentives.

A.2 Core thesis: adoption bottlenecks are governance-and-operations bottlenecks, not “model” bottlenecks

Across federal and state health programs, we see a consistent pattern: the biggest adoption drag is rarely the algorithm in isolation; it is the **system around it**—including **requirements, accountability, change control, evidence, interoperability, and operational oversight**.



Three practical implications:

1. **AI is operated, not delivered.** Intelligent behavior changes over time (data drift, prompt/tool changes, workflow changes) and must be governed continuously.
2. **Regulated trust requires traceability** from intent/requirement → decision → action → evidence, including auditable approvals and governed change.
3. **Requirements for AI are not just “features.”** They must explicitly model: what is off-limits, when to escalate, how humans intervene, and what evidence must be produced.

Practical “method call-outs” (used as concepts, not marketing)

A. Structured, agent-aware requirements

A proven approach is to surface ambiguous needs early, classify requirements (including **agentive** and **policy/constraint** requirements), and link them to events/actors to preserve design traceability.

B. Capability-governed, event-driven operating model

For governed intelligent systems, model behavior on an “event spine,” define explicit approval gates, maintain audit trails for each run/case, enforce change control for prompts/tools/thresholds, and produce “evidence packs” suitable for audit sampling.

A.3 Responses to HHS Specific Questions (1–10)

1) Biggest barriers to private-sector innovation and adoption in clinical care

HHS asks: “What are the biggest barriers...?”

Barrier themes (adoption horizontal):

- **Unclear governance boundary for non-device AI** (documentation, validation, accountability expectations vary by organization and legal posture).
- **Operational risk and “silent change”** (prompt/tool updates, model drift, workflow drift) without standardized change-control patterns.
- **Evidence deficit:** organizations cannot reliably answer “what changed / who approved / what data / can we reproduce outputs,” which blocks scale-out.
- **Interoperability friction:** costly, inconsistent integration to clinical data types; uneven benchmarking; and limited “shared evaluation infrastructure.”
- **Procurement mismatch:** acquisition often funds pilots but not the operational controls (monitoring, governance, evidence, sustainment) needed for production.

Actionable HHS enablers:



- Publish a **Non-Device Clinical AI Governance Profile** (minimum “operational control set”): approvals, audit trail, traceability, monitoring, change control, escalation paths, and evidence artifacts. (These are the controls regulated organizations repeatedly need to scale safely).
- Fund reusable reference architectures for: **clinical workflow integration + evidence logging + controlled deployment** (patterns, not products).
- Create “fast-lane” adoption guidance for **low-risk, assistive AI** with clearly bounded autonomy and required safeguards.

2) Regulatory, payment, or program design changes HHS should prioritize (include CFR references)

HHS asks what changes to prioritize and requests CFR citations. HHS-AI-RFI

Regulatory priorities (reduce uncertainty without weakening protections):

- **Clarify how existing privacy/security obligations apply** to AI workflows, especially where AI vendors act as business associates and where models/logs may contain sensitive data. HIPAA’s Privacy and Security framework is already grounded in **45 CFR Part 160 and Part 164**. HHS+1
- Address adoption friction for sensitive subdomains (e.g., SUD records) by aligning practical implementation guidance with **42 CFR Part 2** while protecting patient trust. eCFR+1
- Reinforce that interoperability is not optional: reduce “AI data lock-in” by aligning AI deployment expectations with ONC’s information blocking framework in **45 CFR Part 171**. eCFR+1

Programmatic / reimbursement design (architecture lens):

- Encourage payment models that reimburse the **operational controls** that make AI safe (monitoring, human oversight time, evidence production), not just the “tool.” HHS notes reimbursement can unintentionally slow innovation and seeks payment policy modernization.
- Where AI is used in coverage/utilization processes, reinforce **individualized decision requirements** and transparency. Medicare Advantage regulations include medical-necessity determination requirements (see **42 CFR § 422.101**). GovInfo+1

Concrete “do now” changes (low regret):

- Issue **implementation guidance** (not necessarily new regulation) defining acceptable patterns for: model/prompt versioning, approval gates, run-level audit logging, and evidence packs—especially for non-device AI embedded in clinical workflow.
- Expand HHS/ONC procurement playbooks to require vendors to deliver **operational governance artifacts** (change logs, monitoring plans, escalation policies, reproducibility strategy) as first-class deliverables.

3) Novel legal and implementation issues for non-medical-device AI; HHS’s role

HHS highlights liability, indemnification, privacy, and security for non-device AI.



Novel issues (implementation reality):

- **Accountability ambiguity** when AI participates in decisions but humans remain responsible; “who is the decision-maker” is often unclear in practice.
- **Duty-to-monitor and change-control**: if the AI behavior evolves, the organization needs a defensible, repeatable oversight process.
- **Evidence expectations**: legal defensibility depends on being able to reconstruct the basis for outputs (data provenance, model/prompt/tool versions, human overrides).
- **Privacy/security risk surfaces** expand: prompts, logs, embeddings, and fine-tuning datasets can inadvertently become regulated data stores (HIPAA/Part 2 implications). HHS+1

Recommended HHS role:

- Publish a **shared vocabulary and minimum controls** for non-device AI in clinical workflow: autonomy levels, escalation triggers, evidence artifacts, and approval expectations (a “common contract” between builders, buyers, and auditors).
- Convene a cross-OpDiv working group (ASTP/ONC + CMS + NIH, as relevant) to align “trust & safety” expectations with interoperability and reimbursement incentives. HHS is explicitly taking a “OneHHS” posture and seeks to accelerate adoption while upholding privacy/civil liberties.

4) Promising evaluation methods (pre/post), metrics, robustness testing; how HHS should support

HHS asks for pre- and post-deployment evaluation methods and mechanisms.

Evaluation must be lifecycle-based (not a one-time “validation”):

- Pre-deployment: define the **operational boundaries** (allowed actions, prohibited actions, escalation rules) and test them, not just model accuracy.
- Post-deployment: continuous monitoring for **behavioral drift**, workflow misfit, error modes, and human override patterns (signals of unsafe automation).

Metrics categories that accelerate trust:

- **Clinical/workflow outcomes** (burden reduction, time-to-task completion, error rates, escalation quality).
- **Safety-of-execution metrics** (rate of off-limits outputs, near-miss frequency, override frequency, “exception raised” patterns).
- **Governance metrics** (evidence completeness score; % of runs traceable to approved versions; time-to-revoke/rollback). These map directly to auditability controls.

How HHS should support (highest leverage):

- Fund shared **test harnesses + benchmark toolchains** for clinical AI workflows (including human-centered evaluation), not just datasets.



- Use **prize competitions** for: (a) “best evidence pack / reproducibility design,” (b) “best controlled escalation UX,” and (c) “best drift detection + rollback patterns,” because these are adoption blockers across domains.
-

5) Supporting private-sector accreditation/certification/testing/credentialing

HHS asks how it can best support accreditation/certification and industry-driven testing.

Recommendation: certify the *operating controls*, not the “AI.”

A pragmatic approach is to support accreditation programs that assess whether an AI-enabled workflow has:

- explicit approval gates,
- run-level audit trails,
- governed change control for prompts/tools/thresholds,
- traceability from requirement → decision → action → evidence,
- operational UX for review/override.

HHS contribution options:

- Publish **reference checklists** and “audit-ready evidence templates” (machine-readable where possible) so accreditation is consistent across organizations.
 - Sponsor an “AI for Clinical Care Controls Catalog” aligned to privacy/security and interoperability expectations.
-

6) Where AI has met/exceeded expectations vs fallen short; most promising novel AI tools

HHS asks where AI has met expectations and where it falls short; and what novel tools have greatest potential.

Where expectations are often met (pattern-level):

- **Administrative burden reduction** (documentation assistance, summarization, routing/triage) when autonomy is bounded and human review is designed well.
- **Decision support augmentation** (surface relevant context, suggest next-best actions) when traceability and escalation are explicit.

Where tools fall short (adoption failure modes):

- “Pilot success” but **production failure** due to missing operating controls: monitoring, versioning, change approval, evidence packs, and interoperability plumbing.
- **Opaque or inconsistent behavior** caused by ungoverned prompt/tool changes (“silent production change”).
- Integration costs and data friction overwhelm value when interoperability isn’t planned early.

Most promising “novel” tools (framed as capabilities):



- **Agentic workflow orchestration with strict boundaries:** AI that can coordinate tasks across systems (e.g., gather context, draft artifacts, route for approvals) while being constrained by policy and escalation rules. The key is a governed event spine + approval gates + evidence.
 - **Evidence-producing AI:** systems designed to generate audit-ready evidence packs (inputs, versions, decisions, overrides) as a native output.
-

7) Who influences adoption inside healthcare orgs; primary admin hurdles

HHS asks which roles/bodies influence adoption and what hurdles exist.

Decision-makers and governing bodies (typical):

- Clinical leadership (CMO/CNO), Quality & Safety, Compliance/Privacy/Security, IT/EHR leadership, and Procurement—often with enterprise risk committees involved for anything that touches clinical workflow at scale.

Primary administrative hurdles (systemic):

- No shared “definition of done” for production AI (governance artifacts missing from acquisition packages).
- Fragmented accountability (clinical vs IT vs compliance) and unclear escalation/override design.
- Evidence burden: teams cannot efficiently produce the artifacts risk/compliance needs to approve and sustain AI.

What helps (architecture lever):

- Use an **agent-as-role accountability model**: explicitly assign which human or automation role fulfills each requirement and what triggers their actions.
 - Treat operational UX (review queues, explanation views, delta dashboards) as part of governance, not “nice-to-have.”
-

8) Where enhanced interoperability widens market opportunities; data types, standards, benchmarking tools

HHS asks where interoperability would widen opportunities and accelerate development (data types, standards, benchmarking).

Highest-leverage interoperability targets:

- **Clinical context + provenance:** not only “clinical facts,” but “what changed and why” (orders, results, meds, problems, notes) with time context and source attribution.
- **Workflow events:** standardized event semantics for common transitions (admission/discharge/transfer; lab-result-available; referral-created; prior-auth-requested; etc.)—because AI is event-driven in operation.
- **Quality and outcomes signals** to support post-deployment monitoring and continuous evaluation.

What HHS can do:



- Promote interoperability approaches that include **benchmarking and evaluation harnesses**, not just data exchange. (Benchmarking is explicitly requested by HHS here.)
- Encourage “composable” integrations so organizations can swap AI components without rebuilding pipelines (reduces lock-in; accelerates competition).

9) Patient/caregiver challenges to address; concerns about AI adoption in clinical care

HHS asks what patients/caregivers want addressed and what concerns they have.

Challenges patients/caregivers often want improved (system-level):

- Reduced administrative friction (scheduling, referrals, prior auth navigation, instructions clarity).
- Better continuity and coordination (the system “remembers” context across encounters).
- More transparent decision pathways (why something was recommended/denied).

Common concerns (trust & safety):

- Privacy and secondary use of sensitive data (HIPAA/Part 2 contexts). HHS+1
- Bias, unequal performance across populations, and lack of recourse when AI is wrong.
- Opaqueness: “who is accountable” and “how do I appeal or correct the record.”

Adoption recommendation:

- Require that AI-enabled workflows expose: (a) **what the AI did**, (b) **what the human decided**, (c) **what data was used at a high level**, and (d) **how to contest/correct**—and log those interactions as part of the evidence trail.

10) AI research areas HHS should prioritize; published findings; costs/benefits/transfers

HHS asks what AI research to prioritize and how literature approaches impacts and economics.

Priority research gaps that unlock adoption (and align to ONC + CMS interests):

1. **Operational governance science for AI in clinical workflow:** measurement, auditability, safety-of-execution, and how to design escalation/override UX that clinicians actually use.
2. **Interoperability + benchmarking infrastructure:** shared evaluation harnesses, standardized task definitions, and reproducible “evidence packs” as first-class research outputs.
3. **Implementation science for AI:** what organization/role structures and operating models reliably scale beyond pilots (cross-site generalizability).
4. **Economics of operational controls:** quantify the costs and benefits of monitoring, evidence, and governance (so reimbursement and procurement can finance the controls that make AI safe).



Mechanisms likely to be most impactful:

- Cooperative agreements and contracts for **reference architectures and open evaluation toolchains**.
- Prize competitions for **governance artifacts** (evidence packs, reproducibility, drift detection/rollback), because these are reusable across many clinical applications and reduce adoption friction.

A.4 Closing: what we can contribute in a prime-led response

Consistent with HHS's call for experience-based input and a secure, industry-supportive approach to accelerate AI adoption in clinical care, this partner's contribution is to help primes and health-sector customers:

- design **enterprise AI adoption architectures** aligned to trust & safety,
- build **governed operating models** (approvals, evidence, change control, monitoring),
- engineer **agentive requirements** that make AI behavior bounded, testable, and auditable.