Department of Health and Human Services
Office of the Secretary
HHS Health Sector AI RFI
RIN 0955-AA13

Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care

HHS Health Sector AI RFI – Specific Question Response

February 23, 2026

Provided to:

Department of Health and Human Services
Office of the Secretary

Provided by:

Bright Star Solutions LLC
Andrew Burk
Business Development Lead
3309 Runnymede PL NW
Washington, DC 20015
aburk@bright-star.com
www.bright-star.com

SBA 8(a) Certified
8(a) STARS III Contract Number: 47QTCB22D0055
UEI: GMB8FVTPZLD3
Top Secret Facilities Clearance

Table of Contents

1 Organizational Perspective and Why Our Experience Transfers to Clinical Care

Bright Star Solutions LLC is a Washington, DC–based SBA 8(a) small business that has delivered secure, mission-critical AI, analytics, automation, and digital modernization solutions across high-stakes federal environments since 2016. We support organizations where public trust is the product, where mistakes create real-world harm, and where compliance is not an afterthought. Those are the same core conditions driving AI adoption barriers in clinical care.

Our work has required turning AI from a “lab concept” into an operational capability under real constraints: identity and access management, data protection, auditability, change management, cross-organization governance, and measurable outcomes. This is why we believe our federal AI delivery model maps directly to healthcare AI, even though we have not yet delivered production systems within HHS.

A key difference between public sector AI and commercial AI is that success is not just accuracy. Success is **institutional trust** at scale. In our experience, AI adoption succeeds when the government can answer five questions confidently:

- ✓ Who approved this and under what policy?
- ✓ What data did it touch and why?
- ✓ What decisions did it influence and how?
- ✓ How do we monitor performance, bias, drift, and safety over time?
- ✓ What happens when it breaks or causes harm?

Those are the same questions HHS needs to be able to answer for AI in clinical care. Bright Star’s core strength is building AI programs that have answers to those questions built into the design, delivery, and operations.

Specialist Capital and Institutional Advisory Contributor

ZXV Group is an independent investment and strategic advisory firm focused on capital allocation, institutional governance, and risk architecture in complex, regulated environments. The firm provides decision-focused advisory support to senior leaders in public and private sectors when technology adoption intersects with regulatory constraint, accountability requirements, and capital commitment decisions.

ZXV Group’s contribution to this response centers on how economic incentives, reimbursement structures, liability allocation, and governance clarity shape the pace and durability of artificial intelligence adoption in clinical settings. Our perspective complements enterprise modernization and secure AI deployment capabilities by grounding adoption in financial, institutional, and risk realities.

Relevant Qualifications

- ✓ Advisory participation in strategic transaction contexts valued at over $500M
- ✓ Capital allocation and portfolio advisory exceeding $200M within regulated environments
- ✓ Governance and oversight advisory within federal and government-facing systems
- ✓ Independent research on technology adoption, institutional change, and capital formation
- ✓ Experience aligning technology deployment with compliance, audit, and accountability frameworks

ZXV Group operates as a principal-led advisory practice. In the context of this RFI, ZXV Group contributes structured analysis regarding incentive alignment, risk containment, regulatory clarity, and capital formation mechanisms that influence responsible AI deployment in healthcare. Bright Star Solutions and ZXV Group are pleased to provide HHS with this RFI providing our expertise in accelerating AI adoption in the clinical care industry.

2 Specific Question #1

What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

Barrier A: Fragmented data access and unclear data permissioning
Healthcare AI requires data from EHRs, labs, imaging systems, pharmacy records, claims, and patient-generated sources. Even when data exists, innovators face a problem we have repeatedly seen in federal AI environments: data is owned by diverse groups, governed by different interpretations, and stored in different formats. The result is uncertainty, slow access, and uneven data quality.

On our U.S. Citizenship and Immigration Services automation work, we supported workflows that required ingesting data from multiple legacy systems with different owners, formats, and sensitivity constraints. Progress did not accelerate until we established a practical model: explicit data-use boundaries, a documented data flow, access control, and full auditability. Once governance and permissioning were unblocked, automation scaled reliably and securely.

How this maps to HHS:
Clinical AI cannot scale without a similar model. HHS can accelerate innovation by helping normalize how AI teams define permissible use, data lineage, and governance artifacts for training, testing, and monitoring.

Recommendation:
HHS should sponsor and publish reference “AI data access patterns” for clinical care that show compliant approaches for 1) model development and 2) validation and evaluation 3) post-deployment monitoring and drift detection. This reduces uncertainty for innovators and health systems without weakening privacy protections.

Barrier B: Governance immaturity and unclear accountability
Healthcare AI adoption stalls when organizations treat AI like a static software feature rather than continuous capability. AI is a lifecycle. It needs owners, controls, monitoring, and escalation paths.

At U.S. Department of State, Bright Star supported enterprise AI tools and adoption programs where the biggest challenge was not building the tool. The challenge was establishing the governance, training, support model, and communications needed to sustain adoption responsibly. Our enterprise-scale generative AI work required policies, use-case boundaries, and stakeholder alignment to safely expand adoption.

How this maps to HHS:
Clinical AI requires explicit governance: model approval processes, clinical oversight, patient safety mechanisms, and defined accountability for monitoring and incident response. HHS can accelerate adoption by promoting governance norms that require 1) named accountable roles for AI operations (clinical owner and technical owner), 2) minimum monitoring standards, 3) mandatory incident reporting mechanisms, 4) clear “human in the loop” decision boundaries where appropriate.

Barrier C: Security and privacy implementation uncertainty, especially for cloud and modern AI
Most organizations understand privacy laws exist. What they struggle with is how to implement modern AI architectures compliantly at scale.

In DHS environments, uncertainty around security implementation is often the slowest gating factor. Bright Star's delivery model addresses this by starting with security architecture and compliance as first-order design constraints and implementing auditable pipelines, role-based access, and monitoring from day one.

How this maps to HHS:
Healthcare AI adoption slows when vendors and providers are unsure how to operationalize compliant AI pipelines, especially for continuous model improvement and post-deployment monitoring. HHS should provide more practical guidance and reference architectures for secure AI operations, including: model supply chain integrity, logging and auditability standards, AI incident response best practices, and access control patterns for clinical settings.

Barrier D: Misaligned financial incentives and ROI ambiguity
Many clinical AI tools create value by reducing burden and improving outcomes, but payment structures do not consistently reward those gains. When ROI is unclear, adoption slows.

In federal programs, adoption accelerates when performance metrics are tied to operational outcomes, not just delivery of features. For example, adoption metrics, time savings, throughput improvements, and user satisfaction were critical to scaling enterprise AI tools in sensitive environments.

How this maps to HHS:
Clinical AI needs "value recognition mechanisms," including reimbursement models that reward burden reduction, error reduction, improved clinical throughput, and improved access.

From a capital and institutional perspective, the principal barrier to healthcare AI adoption is not model performance, but incentive alignment.

Health systems operate under constrained margins and complex reimbursement frameworks. AI adoption scales when it demonstrably:

- ✓ Improves contribution margin per unit of care
- ✓ Reduces denial or compliance risk
- ✓ Lowers administrative burden
- ✓ Improves throughput within fixed-cost environments
- ✓ Reduces measurable liability exposure

Absent economic alignment, AI remains confined to pilot programs. In regulated sectors, uncertainty in liability allocation, monitoring obligations, and reimbursement pathways increases perceived downside risk and suppresses capital deployment. HHS can accelerate adoption by reducing ambiguity in accountability structures and by supporting reimbursement models that recognize operational and quality gains, not only novel technical outputs.

3 Specific Question #2

What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.

We recommend changes in three areas: regulatory clarity, incentive alignment, and shared infrastructure.

Priority A: Reduce uncertainty through operational clarity
AI adoption slows when organizations lack predictable guidance on how existing rules apply across the AI lifecycle. A key government lever is not always new regulation. It is making the existing rules easier to apply consistently.

Bright Star crossover story:
In federal AI deployments, uncertainty was reduced by producing standardized governance artifacts: decision logs, data flow diagrams, usage boundaries, and monitoring reports. Standardization enabled scale.

HHS opportunity:
HHS can accelerate adoption by defining minimum standard "AI governance documentation" that buyers can request and vendors can produce consistently.

Priority B: Incentivize AI that improves outcomes and reduces burden, not just AI that generates novel outputs
HHS reimbursement and program design is one of the most powerful levers for adoption. Where incentives reward volume and not value, adoption will lag.

Bright Star crossover story:
When enterprise AI tools delivered demonstrable time savings and productivity gains, adoption rose rapidly. The difference was measurement discipline and governance mechanisms that demonstrated impact.

HHS opportunity:
HHS can support incentive models that reward: 1) reduced administrative burden, 2) improved throughput and access, 3) improved quality metrics, 4) reduced preventable utilization, and 5) paired with required monitoring and transparency.

Priority C: Invest in shared benchmarks, evaluation frameworks, and public-private pilots
One of the most effective accelerants is shared infrastructure that lowers the cost of responsible adoption.

Bright Star crossover story:
In federal settings, shared platforms, repeatable intake processes, and standardized evaluation approaches enabled scale. For example, enterprise intake and feasibility assessment mechanisms created a repeatable pathway for innovation.

HHS opportunity:
HHS can find: 1) benchmark datasets and test harnesses, 2) reference architectures for monitoring,
3) implementation of science pilots in real clinical environments, 4) scalable evaluation templates and reporting standards.

Regulatory modernization should prioritize operational clarity over additional rulemaking volume.
Institutions require predictable guidance on:

- ✓ Liability demarcation
- ✓ Documentation standards

- ✓ Monitoring expectations
- ✓ Audit pathways
- ✓ Escalation and override authority

In capital-intensive environments, clarity lowers risk perception and enables deployment.

HHS can support adoption through:

- ✓ Standardized AI governance documentation templates
- ✓ Safe harbor frameworks for AI-assisted decision support
- ✓ Regulatory sandboxes tied to measurable outcomes
- ✓ Reimbursement pilots that reward burden reduction and quality improvement

From a capital formation standpoint, policy that reduces downside exposure is a stronger accelerator than policy that increases speed alone.

4 Specific Question #3

For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (e.g., relating to liability, indemnification, privacy, and security). What novel legal and implementation issues exist and what role, if any, should HHS play to help address them?

Non-medical device AI used in clinical care introduces legal and implementation issues that challenge governance structures even when the tool is not regulated as a medical device.

Issue A: Accountability ambiguity and liability boundary confusion
Many AI tools influence clinical workflows even if they do not “make the decision.” When a tool influences decisions, accountability must be explicit.

Bright Star crossover story:
In federal AI, we treated this as a governance problem, not a legal mystery. We implemented: 1) explicit “decision support only” language where appropriate, 2) audit logs showing what data was used and what output was provided, 3) escalation paths for suspected errors, and 4) documented human review requirements for high-risk use cases.

HHS role:
HHS can help by publishing model governance guidance and procurement language that clearly defines: 1) responsibility for monitoring and remediation, 2) expectations for auditability and transparency, 3) requirements for incident reporting and corrective actions.

Issue B: Data privacy, security, and model supply chain risks.
Non-medical AI tools still process protected health information and create risk through third-party services, dependencies, and opaque models.

Bright Star crossover story:
Federal AI delivery requires supply chain risk thinking: what dependencies exist, what vendors touch sensitive data, where are logs stored, who can access, and how do you revoke.

HHS role:
HHS should encourage standardized AI security and supply chain assurance documentation, especially for tools integrated into clinical workflows.

Non-medical AI used in clinical settings introduces accountability ambiguity.

Institutions require explicit governance structures that define:
- ✓ Decision authority
- ✓ Human override boundaries
- ✓ Monitoring responsibilities
- ✓ Incident reporting and remediation pathways

Liability ambiguity increases adoption hesitation even where technology is sound.

Additionally, model supply chain risk, third-party dependencies, and audit traceability must be incorporated into procurement frameworks.

HHS can assist by publishing standardized procurement language and governance expectations that clarify responsibility for:

- ✓ Model monitoring
- ✓ Drift detection
- ✓ Data lineage documentation
- ✓ Incident remediation

Clear accountability reduces perceived institutional risk.

5 Specific Question #4

For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and other workflow and human-centered evaluation methods for clinical care? Should HHS further support these processes? If so, which mechanisms would be most impactful (e.g., contracts, grants, cooperative agreements, and/or prize competitions)?

Bright Star strongly believes evaluation must be treated as a lifecycle function, not a one-time test.

Pre-deployment evaluation should include:

- ✓ clinical validity and utility testing aligned to workflow realities
- ✓ subgroup performance evaluation (equity and fairness)
- ✓ robustness testing under real-world data quality variance
- ✓ usability testing focused on clinician cognitive load and workflow fit
- ✓ clear documentation of intended use and limitations

Bright Star crossover story:
In federal generative AI deployments, the most important early evaluation was not “does it generate text,” but whether it was safe, policy-aligned, and operationally useful under real constraints. We used structured stakeholder review, controlled rollouts, feedback loops, and monitoring dashboards to validate readiness.

Post-deployment evaluation must include:

- ✓ drift detection
- ✓ usage analytics and adoption patterns
- ✓ outcome tracking tied to intervention pathways
- ✓ error reporting and safety incident mechanisms
- ✓ periodic model and prompt updates with controlled release

HHS role and best mechanisms:
HHS should support these processes through:

- ✓ contracts for evaluation infrastructure
- ✓ grants for implementation science
- ✓ cooperative agreements that build shared test harnesses
- ✓ prize competitions focused on monitoring, drift detection, and human-centered evaluation

Evaluation must be treated as a lifecycle function.

Pre-deployment evaluation should include:

- ✓ Clinical utility and workflow alignment
- ✓ Subgroup and equity analysis
- ✓ Robustness testing under real-world data variability
- ✓ Documentation of intended use and limitations

Post-deployment evaluation should include:

- ✓ Drift detection mechanisms
- ✓ Usage analytics
- ✓ Outcome tracking tied to intervention pathways
- ✓ Documented update protocols

From an institutional perspective, evaluation infrastructure is capital infrastructure. HHS can accelerate adoption by funding:

- ✓ Shared test harnesses
- ✓ Benchmark datasets
- ✓ Implementation science pilots
- ✓ Monitoring and audit frameworks

Lifecycle governance increases executive confidence and improves capital discipline.

6 Specific Question #5

How can HHS best support private sector activities (e.g., accreditation, certification, industry-driven testing, and credentialing) to promote innovative and effective AI use in clinical care?

HHS should support certification systems that validate operational readiness, not just documentation completeness.

Bright Star's view:
What healthcare needs is a norm that "AI is certifiable as an operational system." That includes governance, monitoring, security, and support.

Bright Star crossover story:
In federal environments, systems achieve trust when they have repeatable assurance artifacts, consistent review, and transparent operational processes. The closer clinical AI assurance resembles disciplined government assurance models, the faster adoption can scale.

Recommendation:
HHS should support frameworks that certify:

- ✓ monitoring readiness
- ✓ auditability

- ✓ governance structures
- ✓ security patterns
- ✓ human factors evaluation without creating fragmented competing certification burdens.

Certification systems should validate operational readiness, not only documentation completeness.

Effective certification frameworks should evaluate:

- ✓ Governance structures
- ✓ Monitoring readiness
- ✓ Auditability
- ✓ Security posture
- ✓ Human factors integration

Certification aligned with institutional risk management increases trust and lowers procurement friction.

7 Specific Question #6

Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?

We can speak from a cross-sector operational AI lens: tools succeed when they remove burden and integrate into workflows. They fail when they are bolted on, untrusted, or unsupported.

Where AI typically exceeds expectations:

- ✓ administrative burden reduction
- ✓ documentation and summarization
- ✓ workflow triage and routing support

- ✓ population-level analytics tied to clear interventions

Bright Star crossover story:
At USCIS, RPA and AI components succeeded when they were embedded into process pathways with monitoring, auditability, and human oversight. Adoption rose when time savings were measurable and error rates were managed.

Where AI typically falls short:

- ✓ tools that are not integrated into workflows
- ✓ tools that change clinician processes without training and change management
- ✓ tools with opaque outputs that clinicians do not trust
- ✓ tools without monitoring and drift management

Novel AI tools with high potential:

- ✓ care navigation copilots that reduce access friction
- ✓ prior authorization and claims automation with transparency
- ✓ predictive models tied to intervention pathways, not just prediction
- ✓ patient communication tools that increase comprehension and engagement while preserving privacy and consent

Across sectors, AI tools meet expectations when they reduce burden and integrate into workflows. They fall short when deployed without governance, monitoring, or workflow alignment.

High-performing use cases include:

- ✓ Administrative automation
- ✓ Documentation support
- ✓ Workflow triage
- ✓ Population-level analytics tied to defined interventions

Underperforming deployments typically lack:

- ✓ Clear ROI measurement
- ✓ Monitoring discipline
- ✓ Workflow integration
- ✓ Defined accountability

AI success correlates with governance maturity.

8 Specific Question #7

Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?

AI adoption is a governance decision, not a technology purchase.

Key decision makers include clinical leadership, informatics, IT security, compliance, procurement, legal risk management, quality and safety governance, and executive leadership.

Primary administrative hurdles:

- ✓ procurement cycles and contracting risk
- ✓ unclear ROI and evidence thresholds
- ✓ legal uncertainty and accountability concerns
- ✓ integration complexity with legacy systems
- ✓ workforce and training capacity

Bright Star crossover story:

At DOS, scaling AI required coordinated buy-in from leadership, security, policy stakeholders, and end users. The adoption model worked because we treated change management, training, and governance as primary deliverables, not secondary tasks.

AI adoption as stated is a governance decision.

Key actors include:

- ✓ Clinical leadership
- ✓ Informatics
- ✓ Compliance and legal
- ✓ Procurement
- ✓ Executive leadership

Administrative hurdles include:

- ✓ Procurement cycles
- ✓ Liability ambiguity
- ✓ ROI uncertainty
- ✓ Integration complexity
- ✓ Workforce capacity

Adoption improves when change management, governance, and executive framing are treated as primary deliverables rather than secondary tasks.

9 Specific Question #8

Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

Interoperability is the growth engine for clinical AI. HHS can unlock major market opportunity by supporting standardization and accessible benchmarking across:

- ✓ clinical longitudinal data
- ✓ imaging
- ✓ labs and pathology
- ✓ claims and utilization
- ✓ social determinants
- ✓ patient-reported outcomes

Bright Star crossover story:
At USCIS, integration across legacy systems enabled scalable automation and analytics. Once data exchange patterns were standardized, capabilities scaled quickly. Healthcare needs the same interoperability enablement, aligned to data standards and real-world benchmarking infrastructure.

Interoperability challenges are frequently governance and incentive challenges rather than purely technical constraints.

HHS can accelerate adoption by supporting:

- ✓ Standardized reporting formats
- ✓ Structured data-sharing agreements
- ✓ Benchmarking infrastructure
- ✓ Incentives for audit traceability

Interoperability lowers capital barriers and expands responsible market opportunity.

10 Specific Question #9

What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?

Patients and caregivers want AI to reduce friction and increase clarity, but they fear loss of control.

Top desired improvements:

- ✓ reduced administrative burden
- ✓ faster access and scheduling
- ✓ clearer explanations and care plans
- ✓ improved coordination across providers
- ✓ proactive identification of risks

Top concerns:

- ✓ privacy and data misuse
- ✓ opaque decision making
- ✓ bias and inequitable outcomes
- ✓ reduced human interaction
- ✓ inability to appeal or challenge AI-influenced decisions

Bright Star crossover story:
Public trust is central in immigration and diplomacy systems. We have learned that transparency and agency matter: clear explanations, auditability, and feedback channels are essential to trust. Clinical AI must build those features in from the start.
Patients and caregivers seek reduced friction, improved clarity, and preserved agency.

Top desired outcomes:

- ✓ Reduced administrative burden
- ✓ Faster access
- ✓ Clear explanations
- ✓ Improved coordination

Primary concerns:

- ✓ Privacy
- ✓ Bias
- ✓ Opaque decision-making
- ✓ Loss of human oversight

Trust requires transparency, appeal mechanisms, and clear communication pathways.

11 Specific Question #10

Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care?
a. Are there published findings about the impact of adopted AI tools and their use clinical care?
b. How does the literature approach the costs, benefits, and transfers of using AI as part of clinical care?

Yes. HHS should prioritize research that closes the gap between “model performance” and “clinical reality.”

Priority research areas:

- ✓ Implementation science for AI in clinical workflows
- ✓ Long-term monitoring and drift detection
- ✓ Bias and subgroup performance evaluation
- ✓ Human factors and clinician cognitive load
- ✓ Privacy-preserving analytics and federated methods
- ✓ Economic impact measurement tied to real delivery models
- ✓ Benchmarking infrastructure and shared evaluation harnesses

10a: Are there published findings about the impact of adopted AI tools in clinical care?

There is a growing body of research, but impact varies widely depending on workflow integration, monitoring, and governance maturity. Tools that are integrated, measured, and continuously managed show stronger outcomes than tools deployed as stand-alone pilots.

Bright Star’s federal experience supports the same principle: adoption and impact are driven by governance, training, monitoring, and operational integration.

10b: How does literature approach costs, benefits, and transfers?

The most credible approaches account for:

- ✓ direct operational savings
- ✓ downstream cost avoidance
- ✓ workforce burden reduction
- ✓ quality and safety outcomes
- ✓ infrastructure and monitoring costs
- ✓ change management and training investments

AI “benefit claims” fail when they ignore adoption and lifecycle costs.

Bright Star’s approach to ROI emphasizes measurable operational outcomes and lifecycle sustainment requirements, rather than one-time implementation success.

Research should focus on economic and institutional drivers of adoption, including:

- ✓ Contribution-margin analysis
- ✓ Liability exposure mapping
- ✓ Monitoring cost modeling
- ✓ Implementation science for workflow integration
- ✓ Governance maturity benchmarks

Adoption speed will be determined less by incremental model improvements and more by clarity in economic incentives and institutional accountability.

12 Closing

Bright Star believes HHS has an opportunity to accelerate clinical AI adoption through a practical model: reducing uncertainty, aligning incentives, and funding shared evaluation infrastructure. Our public sector AI delivery story correlates strongly with HHS’s needs because it is built on trust, governance, security, and operational accountability. Those are the same foundations required for safe, scalable AI in clinical care.

Responsible AI adoption in clinical care requires alignment among regulatory clarity, economic incentives, governance discipline, and institutional culture.

Bright Star contributes enterprise modernization and secure AI delivery expertise. ZXV Group contributes independent capital and institutional analysis focused on improving decision quality before large-scale commitments are made. Together, this combined perspective supports HHS in designing policy environments where safe, accountable, and economically sustainable AI deployment can scale.