



515 KING STREET, ALEXANDRIA VA 22314

MEMORANDUM

To: Hon. Robert F. Kennedy, Jr., Secretary, United States Department of Health and Human Services

From: Andrew Langer, Director, Center for Regulatory Freedom, CPAC Foundation

Date: February 23, 2026

Re: Department of Health and Human Services Information Collection Request, “HHS Health Sector AI RFI - Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care”, Docket #HHS-ONC-2026-001, Fed. Reg. 2025-23641, Published January 5, 2026

Below are comments of the American Conservative Union Foundation's (d/b/a. Conservative Political Action Coalition Foundation) (hereinafter “CPAC Foundation”) Center for Regulatory Freedom (hereinafter “CRF”), in response to the Department of Health and Human Services Information Collection Request, “HHS Health Sector AI RFI - Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care”, Docket #HHS-ONC-2026-001, Fed. Reg. 2025-23641, published January 5, 2026.

CRF is a project of the CPAC Foundation, a non-profit, non-partisan 501(c)(3) research and education foundation. Our mission is to inject a common-sense perspective into the regulatory process, to ensure that the risks and costs of regulations are fully based on sound scientific and economic evidence, and to ensure that the voices, interests, and freedoms of Americans, and especially of small businesses, are fully represented in the regulatory process and debates. Finally, we work to ensure that regulatory proposals address real problems, that the proposals serve to ameliorate those problems, and, perhaps most importantly, that those proposals do not, in fact, make public policy problems worse.

Introduction

The Center for Regulatory Freedom (CRF) supports responsible, innovation-driven reform in health care delivery. Advances in data analytics, digital infrastructure, and artificial intelligence present meaningful opportunities to modernize clinical workflows, reduce inefficiencies, and

strengthen the overall performance of the health system. Sound regulatory policy should enable beneficial innovation while preserving the safeguards necessary to protect patients and sustain trust in medical institutions.

Artificial intelligence, in particular, holds substantial promise. Properly deployed, AI systems can reduce administrative burden on clinicians, streamline documentation requirements, assist with longitudinal chart review, identify care gaps, and surface clinically relevant patterns that might otherwise be overlooked in complex patient histories. In an era where physicians face escalating paperwork obligations and time constraints, intelligent decision-support tools can improve efficiency and allow clinicians to focus more fully on direct patient care.

At the same time, medical care is uniquely sensitive among regulated domains. Errors in this context are not merely financial or reputational; they may produce irreversible physical harm. Patients are often in vulnerable states, lack specialized knowledge, and must rely on professional expertise. Moreover, health care involves profound information asymmetries between providers and patients, as well as fiduciary obligations that distinguish it from most other commercial environments. For these reasons, the governance framework for AI in medical care must be more disciplined than general-purpose software regulation.

A further structural reality must inform HHS's approach: the longstanding value of longitudinal primary care relationships. Continuity with a primary care physician strengthens clinical accuracy, improves medication management, reduces duplicative testing, and enhances early detection of emerging conditions. A physician who knows a patient's history—both medical and contextual—is better positioned to make sound decisions than one encountering that patient for the first time without adequate background.

However, evolving care models are reshaping how patients access services. Younger cohorts increasingly rely on telehealth platforms, urgent care centers, and episodic or on-demand primary care arrangements. While these models expand access and convenience, they also introduce fragmentation risks. In such an environment, informational continuity becomes as important as relational continuity. AI-assisted clinical briefing tools—designed to synthesize longitudinal patient information at the point of care—may help bridge that gap, provided they are appropriately governed.

The central policy question, therefore, is not whether AI should be used in medical care, but how to preserve the safety and quality benefits of traditional primary care relationships while accommodating technological advancement and evolving delivery models. The goal should be to strengthen physician decision-making through structured, transparent information support—not to displace clinical judgment or automate high-stakes decisions without accountability.

CRF's position is that HHS should adopt a calibrated, risk-tiered governance framework for AI in medical care. Permissionless innovation should remain the default posture for low-risk, assistive functions that operate under clear transparency and human-review safeguards. At the same time, heightened evidentiary and oversight requirements should apply where AI systems materially influence diagnosis, triage, treatment selection, or other consequential clinical decisions. Such a balanced framework will promote innovation, preserve patient safety, and ensure that artificial intelligence serves as a tool to reinforce—rather than erode—the integrity of the physician-patient relationship.

Executive Summary

Artificial intelligence is rapidly entering the clinical environment. Properly structured, these technologies can enhance physician efficiency, reduce administrative burden, and improve patient care. Improperly governed, however, they risk amplifying fragmentation, obscuring accountability, and introducing preventable safety concerns. HHS faces a pivotal policy choice: how to enable innovation in health care AI while preserving patient safety, physician judgment, and system integrity.

CRF submits that the appropriate framework is neither blanket precaution nor unbounded permissionless deployment. Instead, HHS should adopt a calibrated, risk-tiered, lifecycle governance model that preserves innovation for low-risk assistive functions while applying proportionate safeguards where AI materially influences clinical outcomes.

Substantively:

- **Continuity of care is a measurable quality and safety asset.** Longitudinal primary care relationships reduce fragmentation, improve medication management, and enhance clinical accuracy by preserving contextual patient knowledge.
- **Care delivery models are evolving.** Younger cohorts increasingly rely on episodic, telehealth, and on-demand services, creating structural information discontinuity even as access expands.
- **AI-powered “clinical briefing” tools offer a potential bridge.** Properly designed systems can synthesize longitudinal patient data into concise, structured summaries that enhance physician situational awareness at the point of care.
- **Material risks must be addressed.** These include hallucinated or fabricated data, omission of critical contraindications, automation bias, privacy vulnerabilities, model drift, and over-reliance that weakens independent clinical reasoning.
- **A binary “precautionary vs. permissionless” framework is ill-suited to health care.** Medical AI presents heterogeneous risk profiles that require calibrated governance rather than categorical approval or prohibition.
- **CRF recommends a risk-tiered, lifecycle-based governance model:**
 - Permissionless deployment for low-risk assistive functions, subject to baseline transparency, provenance, and human-review safeguards.
 - Escalating oversight triggers for tools that materially influence diagnosis, triage, treatment, or other consequential clinical decisions.
 - Strong accountability, auditability, incident reporting, and post-deployment performance monitoring to address model drift and real-world safety signals.
- **HHS should harmonize governance across CMS, ONC, and FDA authorities.** Fragmented or duplicative regulatory approaches will undermine clarity, slow responsible innovation, and increase compliance burdens without commensurate safety gains.

A disciplined, risk-tiered approach will allow AI to reinforce informational continuity in an era of evolving care models, strengthen the physician-patient relationship, and improve affordability through efficiency gains—while maintaining clear lines of accountability and patient safety protections.

I. CONTINUITY OF CARE AS A QUALITY AND SAFETY FOUNDATION

Longitudinal primary care relationships function as a structural safeguard within the health care system. A primary care physician who follows a patient over time accumulates contextual knowledge that cannot be fully captured in isolated clinical notes. This continuity reduces fragmentation in decision-making by anchoring diagnosis, treatment, and referral decisions in a coherent understanding of the patient's medical trajectory. In this sense, continuity is not merely a matter of convenience; it is an operational mechanism for improving safety and clinical precision.

A sustained physician–patient relationship strengthens diagnostic accuracy and therapeutic calibration. Familiarity with prior medication responses, historical lab trends, chronic condition progression, and patterns of symptom presentation enables clinicians to interpret new complaints in light of established baselines. Subtle deviations that might appear insignificant in a single encounter may carry greater meaning when viewed longitudinally. Likewise, knowledge of prior adverse reactions or treatment failures helps prevent repetition of ineffective or harmful interventions.

Continuity also extends beyond strictly clinical data. Primary care physicians often develop awareness of social determinants, behavioral health factors, family history nuances, and lifestyle constraints that influence treatment adherence and outcomes. Understanding a patient's employment schedule, caregiving responsibilities, transportation access, or financial limitations can materially shape care planning. These contextual variables frequently determine whether a treatment recommendation is feasible, sustainable, and safe.

When continuity erodes, the system incurs measurable inefficiencies and risks. Fragmented care environments increase the likelihood of duplicative testing, inconsistent medication management, and missed preventive screenings. Episodic providers who lack access to complete patient histories may order repeat diagnostics out of caution, inadvertently raising costs and patient burden. More concerning, incomplete knowledge of prior contraindications or drug interactions can introduce avoidable safety hazards.

Discontinuity also weakens follow-through. When responsibility for care is diffuse, accountability for monitoring results, adjusting medications, or ensuring specialist coordination can become unclear. Missed follow-ups, unmanaged chronic conditions, and delayed interventions become more likely in the absence of a central coordinating clinician. Continuity therefore functions as a stabilizing force in an otherwise complex and decentralized system.

At the same time, care delivery patterns are evolving. Younger cohorts increasingly utilize telehealth platforms, urgent care centers, retail clinics, and other on-demand service models. These models expand access and convenience, and they may lower certain barriers to entry for patients who might otherwise defer care. However, they inherently reduce relational continuity and increase the likelihood that individual encounters occur without the benefit of long-term contextual knowledge.

In this environment, policy should prioritize preservation of informational continuity even where relational continuity declines. If patients increasingly move between providers or platforms, the system must ensure that longitudinal medical information remains coherent, accessible, and

clinically actionable. The objective should not be to resist innovation in care access models, but to mitigate the fragmentation risks that accompany them.

Artificial intelligence, properly structured, may serve as a continuity amplifier rather than a substitute for physician judgment. AI-powered clinical briefing tools can synthesize longitudinal data into structured summaries that restore situational awareness at the point of care. However, these tools must operate in support of—rather than in place of—the primary care physician’s reasoning. Continuity ultimately resides in accountable human judgment; technology should strengthen that judgment by improving information clarity, not displace it.

II. THE “AI CLINICAL BRIEFING” MODEL: OPPORTUNITY AND FUNCTIONAL ROLE

The specific use case CRF supports is the AI-powered “clinical briefing” tool deployed immediately prior to a patient encounter. These systems ingest structured and unstructured data from electronic health records—including problem lists, medication histories, laboratory trends, imaging results, prior encounter notes, discharge summaries, and referral documentation—and synthesize that information into a concise, structured pre-visit summary for the treating clinician. The purpose is to reduce cognitive and time burdens associated with manual chart review while preserving access to the underlying data.

Modern EHR environments often contain years of accumulated information presented in fragmented, duplicative, or poorly organized formats. Clinicians operating under time constraints may face dozens or even hundreds of pages of notes for a complex patient. AI summarization tools can compress that volume into an actionable briefing that highlights relevant history, recent changes, unresolved issues, and potential risk factors. When properly designed, such tools enhance situational awareness without displacing clinical discretion.

It is essential to distinguish between three distinct functional categories. First, administrative summarization tools primarily organize and condense information already present in the record. They may extract key data points, compile timelines, or generate structured summaries, but they do not introduce new diagnostic or treatment recommendations. These tools function as informational aids rather than decision engines.

Second, clinical insight–highlighting tools go beyond simple summarization by identifying patterns or flagging potential concerns. For example, they may highlight medication interactions, overdue screenings, abnormal lab trends, or inconsistencies in documentation. While still assistive, these systems begin to shape clinical attention and therefore carry incrementally greater influence over decision-making.

Third, diagnostic or treatment recommendation engines actively suggest courses of action, risk scores, triage determinations, or therapeutic pathways. These systems materially influence clinical judgment and may, in some contexts, function with a degree of autonomy. They operate at a higher risk tier and require correspondingly stronger validation and oversight.

The “AI clinical briefing” model contemplated here is primarily focused on the first two categories—administrative summarization and bounded insight highlighting—rather than autonomous diagnostic decision-making. The goal is to improve information clarity, not to

substitute algorithmic judgment for physician reasoning. The tool’s core function is compression and organization, not independent clinical authority.

Properly designed briefing systems can materially reduce chart-review burden. By surfacing medication changes, prior adverse events, hospitalizations, and pending follow-ups in a structured format, they allow clinicians to enter the patient encounter better prepared and more focused. This improved preparation may enhance both efficiency and quality, particularly in environments where patients move between providers or care settings.

Importantly, the objective is information compression rather than clinical replacement. The physician remains responsible for interpreting the summarized information, validating its relevance, and making ultimate care decisions. The tool should present data with traceability and transparency, enabling clinicians to drill down into underlying records as needed. At no point should the system obscure the source of its conclusions or operate beyond physician oversight.

Integration into physician workflow is critical. The briefing must be delivered in a manner that supports pre-visit preparation without disrupting established clinical processes. It should be available within the existing EHR environment, clearly labeled as assistive, and accompanied by appropriate disclaimers regarding scope and limitations. The system should facilitate review and correction, not create friction or encourage passive acceptance.

When appropriately bounded—limited to summarization and structured highlighting, operating under explicit physician authority, and supported by transparency safeguards—the AI clinical briefing model represents a low-to-moderate risk application of artificial intelligence in health care. Its potential to restore informational continuity in fragmented care environments justifies a governance posture that enables deployment while maintaining disciplined guardrails. Properly calibrated policy can allow these tools to strengthen physician performance without compromising safety or accountability.

III. RISK CATEGORIES AND FAILURE MODES IN CLINICAL AI

Artificial intelligence in clinical environments presents heterogeneous and context-dependent risks. A disciplined regulatory framework must begin with a clear understanding of distinct failure modes rather than treating “AI risk” as a monolithic concept. The safety profile of an assistive chart summarization tool differs materially from that of an autonomous diagnostic engine, and policy must reflect those distinctions. The following categories outline the principal risk domains relevant to AI-enabled clinical briefing systems and related applications.

A careful delineation of these risks is not an argument against innovation. Rather, it is a prerequisite to proportionate governance. Identifying how systems fail—technically, operationally, and behaviorally—enables regulators to design guardrails that preserve benefits while mitigating foreseeable harm.

A. Hallucination Risk

Hallucination risk refers to the generation of fabricated or unsupported information. In a clinical context, this may include nonexistent diagnoses, incorrect medication lists, spurious allergy notations, or inaccurate descriptions of prior medical events. Even a single fabricated data point can distort clinical reasoning if it appears authoritative.

The legal and patient safety implications are substantial. A hallucinated allergy could improperly constrain treatment options; a fabricated prior diagnosis could lead to unnecessary testing or inappropriate referrals; an inaccurate medication entry could introduce dangerous drug interactions. Because clinicians often operate under time pressure, they may not immediately detect subtle inaccuracies embedded within a polished summary.

Unlike simple typographical errors, hallucinations carry an epistemic risk: they present as coherent and plausible statements. This makes them more difficult to detect through casual review. In high-acuity settings, even minor inaccuracies may cascade into consequential errors.

Mitigation requires design features such as source traceability, confidence indicators, and explicit differentiation between extracted data and inferred conclusions. Systems must make it operationally easy for clinicians to verify claims against the underlying record. Transparency is the primary countermeasure to hallucination risk.

B. Omission Risk

Omission risk arises when a system fails to surface critical information that is present in the underlying record. In many clinical scenarios, omission may pose a greater danger than fabrication. A missing notation regarding a prior adverse drug reaction, anticoagulant use, or cardiac event can have immediate and serious consequences.

Clinical decision-making often depends on identifying contraindications or outlier events within large datasets. If a summarization system compresses information in a way that excludes such signals, clinicians may proceed under incomplete assumptions. The brevity that enhances efficiency can inadvertently obscure nuance.

Omission risk is especially pronounced in longitudinal records containing years of documentation across multiple providers and systems. Data fragmentation, inconsistent coding, and unstructured narrative notes can make it difficult for algorithms to detect relevant entries, increasing the likelihood that critical information is not highlighted.

Effective mitigation requires structured completeness checks, explicit acknowledgment of data sources reviewed, and clear indicators when records are incomplete or unavailable. Systems should err on the side of signaling uncertainty rather than presenting a falsely comprehensive summary.

C. Automation Bias

Automation bias refers to the human tendency to over-rely on automated outputs, particularly when those outputs are presented in a clear and authoritative format. In time-constrained clinical environments, physicians may unconsciously defer to summarized or highlighted information rather than independently interrogating the full record.

Interface design can exacerbate this tendency. Clean dashboards, risk scores, and concise bullet-point summaries can create a perception of comprehensiveness even when the system is assistive rather than determinative. The more seamless the integration, the greater the risk that clinicians internalize AI outputs as definitive.

Over time, excessive reliance on automated summaries may degrade independent clinical reasoning. If physicians become habituated to accepting compressed narratives without verification, diagnostic vigilance may decline. This dynamic is particularly concerning in training environments where less-experienced clinicians may treat AI output as authoritative guidance.

Mitigating automation bias requires deliberate workflow design. Systems should clearly indicate their assistive nature, facilitate easy access to underlying documentation, and incorporate mechanisms for clinician review and correction. Accountability should remain firmly anchored in human decision-makers.

D. Distribution Shift and Model Drift

Clinical environments are dynamic. Patient demographics change, treatment standards evolve, new medications enter the market, and coding practices are revised. Distribution shift occurs when the data environment in which a model operates diverges from the environment in which it was trained.

Over time, such shifts can degrade performance. A model trained on historical patterns may fail to recognize emerging conditions or updated therapeutic standards. In summarization contexts, evolving documentation practices may alter how information is recorded, reducing the system's ability to extract relevant content accurately.

Model drift may occur gradually and invisibly. Performance degradation can accumulate without obvious failure events, especially in assistive systems where outputs are not directly benchmarked against clear ground truth metrics. This creates a latent safety risk.

Addressing distribution shift requires lifecycle governance. Ongoing monitoring, periodic performance reassessment, and structured update protocols are essential. AI in health care cannot be treated as static software; it must be subject to adaptive oversight consistent with the dynamic nature of clinical practice.

E. Privacy and Security Risks

AI clinical systems necessarily process sensitive protected health information. This creates heightened privacy and security obligations. Unauthorized access, data breaches,

or inadequate encryption protocols can expose patients to identity theft, discrimination, or reputational harm.

Beyond traditional cybersecurity concerns, AI systems introduce risks related to secondary use of data. Training processes, model improvement cycles, or third-party integrations may expand data exposure beyond the original purpose for which information was collected. Without clear guardrails, mission creep can erode patient trust.

Large-scale data aggregation also increases the attractiveness of systems as cyber targets. The concentration of longitudinal health records within AI-enabled platforms amplifies the consequences of security failures.

Mitigation requires strict adherence to purpose limitation, data minimization, access controls, and audit logging. Governance structures must clearly define who may access data, for what purposes, and under what safeguards. Privacy and security are not peripheral concerns but foundational conditions for sustained adoption.

In evaluating these risk categories, it is essential to recognize that magnitude varies significantly by use case and degree of autonomy. A bounded assistive summarization tool presents a different risk profile than an autonomous diagnostic or treatment engine. Accordingly, governance should scale with material clinical influence rather than applying uniform constraints across heterogeneous technologies. A calibrated framework that distinguishes between low-risk assistance and high-impact autonomy will better protect patients while preserving the benefits of innovation.

IV. PERMISSIONLESS INNOVATION AND THE PRECAUTIONARY PRINCIPLE IN MEDICAL AI

Debates over emerging technologies often collapse into a false binary: either innovation should proceed without regulatory friction, or precaution should dominate to prevent potential harm. In the context of medical AI, neither extreme is appropriate. A purely precautionary posture that demands exhaustive pre-deployment proof for all AI-enabled tools would freeze low-risk innovation and delay beneficial improvements in efficiency and care coordination. Conversely, an unbounded permissionless approach would expose patients to preventable risks in high-stakes clinical environments.

CRF rejects precautionary paralysis where assistive technologies pose minimal and reversible risk. Low-risk tools that organize information, summarize records, or highlight existing data—when deployed under clear transparency and human-review safeguards—should not be subjected to disproportionate regulatory barriers. Excessive friction in these domains would preserve administrative inefficiencies and undermine opportunities to improve informational continuity.

At the same time, CRF rejects the notion that medical AI should be deployed without calibrated oversight in high-risk contexts. Where AI systems materially influence diagnostic pathways, triage decisions, or treatment selection, the stakes rise significantly. In such circumstances, the consequences of error are not abstract; they may directly affect patient morbidity or mortality. Regulatory restraint must not become regulatory abdication.

Accordingly, CRF proposes a hybrid standard. Permissionless innovation should remain the default posture for reversible, assistive, and non-autonomous functions. These include bounded summarization tools and informational aids that operate under physician supervision and do not independently initiate or execute clinical actions. Such tools should be permitted to deploy subject to baseline transparency, traceability, and accountability safeguards.

Escalating safeguards should apply where AI systems exert what CRF defines as “Material Clinical Influence.” Material Clinical Influence exists when AI outputs could reasonably alter diagnosis, triage, or treatment decisions. In these contexts, validation, performance benchmarking, monitoring, and clear delineation of intended use become essential to protect patient safety.

CRF further defines “Autonomous Action” as AI-driven execution of clinical decisions or administrative determinations without contemporaneous physician verification. Autonomous routing of patients, automated treatment adjustments, or coverage determinations made without human review represent a qualitatively different risk tier. It is autonomy—not mere assistance—that should trigger heightened oversight, evidentiary requirements, and structured accountability mechanisms.

This distinction between assistance and autonomy is critical. Many AI-enabled tools enhance human judgment without displacing it. When clinicians retain authority, review outputs, and exercise independent reasoning, the risk profile remains bounded. When AI systems act independently or create de facto determinative outcomes, oversight must increase proportionally.

Finally, innovation discipline should be understood as a trust-building mechanism rather than a constraint on progress. Clear standards that distinguish low-risk assistive tools from high-impact autonomous systems provide certainty to developers, clinicians, and patients alike. A calibrated framework that protects safety while enabling responsible experimentation will strengthen public confidence in medical AI and promote sustainable technological advancement within the health care system.

V. RISK-TIERED GOVERNANCE FRAMEWORK FOR HHS

A durable policy framework for AI in medical care must move beyond abstract principles and provide operational clarity. HHS should adopt a risk-tiered governance architecture that scales oversight to functional impact and autonomy level. Such an approach promotes innovation where risk is bounded while ensuring that systems with material clinical consequences are subject to proportionate safeguards. The objective is not to regulate “AI” as a monolith, but to regulate specific use cases according to their real-world risk profile.

A tiered model also provides regulatory certainty. Developers, health systems, and clinicians benefit from knowing in advance which category their tools fall into and what compliance expectations attach to each tier. Clear delineation reduces ambiguity, avoids over-deterrence of low-risk innovation, and concentrates oversight resources where patient safety implications are most significant.

Tier 1: Assistive / Informational Tools

Tier 1 encompasses assistive and informational systems that support clinicians without independently influencing or executing clinical decisions. These include documentation drafting tools, chart summarization systems, and automated care gap reminders. Their primary function is organizational or clerical efficiency rather than diagnostic or therapeutic determination.

These systems operate within physician workflow and under direct human supervision. They do not independently initiate clinical action, nor do they generate prescriptive treatment recommendations. While they may influence clinician awareness, they remain fundamentally assistive and reversible in impact.

Baseline governance requirements should apply uniformly across Tier 1 tools. Provenance transparency must be embedded by default, enabling clinicians to trace summarized content back to underlying source records. Systems should clearly disclose which data sources were included and whether external or incomplete records were excluded.

In addition, Tier 1 tools should be clearly labeled as assistive within the user interface, avoiding any implication of autonomous authority. Audit logging should capture system usage and clinician review activity to preserve accountability and enable post hoc evaluation if necessary. These safeguards are sufficient for bounded, low-risk applications.

Tier 2: Decision-Influencing Tools

Tier 2 includes systems that meaningfully influence clinical judgment, even if they do not execute decisions autonomously. Examples include risk scoring tools, triage prioritization systems, and diagnostic suggestion engines. These systems can shape physician attention and alter decision pathways, thereby elevating their risk profile.

While physicians retain ultimate authority, Tier 2 tools may exert Material Clinical Influence. Their outputs can reasonably affect diagnosis, referral patterns, medication choices, or patient prioritization. As such, they warrant stronger evidentiary and governance requirements than purely assistive tools.

Context-specific validation should be required prior to deployment. Performance benchmarking must demonstrate that the system operates reliably within the intended patient population and clinical setting. Error tolerance thresholds should be explicitly defined, particularly where outputs may affect high-acuity or vulnerable populations.

Tier 2 systems should also incorporate structured physician attestation workflows. Clinicians should confirm review of AI-generated outputs before finalizing clinical decisions. This reinforces human accountability and mitigates automation bias. Ongoing performance monitoring should detect drift or emerging safety concerns over time.

Tier 3: Autonomous or High-Impact Systems

Tier 3 encompasses autonomous or high-impact systems capable of executing or effectively determining outcomes without contemporaneous physician verification. Examples include independent triage routing, automated authorization or coverage determinations, and direct execution of clinical actions such as dosage adjustments or treatment pathway activation.

These systems present the highest risk profile because autonomy introduces the potential for harm without immediate human correction. Even if designed to optimize efficiency, autonomous systems may create systemic errors that scale rapidly if not properly governed. Accordingly, Tier 3 tools require the most rigorous oversight architecture.

A pre-deployment safety case should be mandatory for Tier 3 systems. This includes clearly defined intended use, validation within the target population, explicit performance metrics, and documented risk mitigation strategies. Regulators should require demonstration that the system performs at least as well as established clinical standards under realistic operating conditions.

Post-deployment obligations must include ongoing performance monitoring, structured reporting of adverse events, and periodic reassessment of model accuracy. Independent review pathways should allow clinicians or affected parties to challenge system outputs. Override and appeals mechanisms must be easily accessible and functionally meaningful to ensure that autonomy does not displace accountability.

A risk-tiered governance framework enables HHS to balance innovation with safety in a principled and operationally coherent manner. By scaling oversight to functional impact and autonomy level, the Department can preserve permissionless deployment for low-risk assistive tools while imposing proportionate safeguards where patient safety implications intensify. This architecture promotes clarity, accountability, and sustained public trust in AI-enabled medical care.

VI. HUMAN ACCOUNTABILITY AND WORKFLOW INTEGRATION

The introduction of AI into clinical environments must not obscure a foundational principle: physicians remain the ultimate decision-makers in patient care. No deployment model should dilute professional accountability or create ambiguity regarding responsibility for diagnosis, treatment, or follow-up. AI systems may inform and support clinical judgment, but they must not displace the physician's fiduciary and ethical obligations to the patient.

Preserving physician primacy requires deliberate interface design. User interfaces should encourage active review rather than passive acceptance of AI-generated outputs. Visual presentation should clearly indicate that summaries, risk scores, or highlighted insights are assistive tools. The system should prompt engagement—such as confirmation steps or expandable evidence views—rather than defaulting to silent incorporation into the clinical record.

Traceability is a core safeguard. Clinicians must be able to click through from any summarized statement or highlighted alert directly to the underlying source data within the electronic health record. This traceability allows physicians to validate accuracy, assess context, and identify potential omissions. Without easy access to primary data, clinicians may be forced to rely on compressed summaries without sufficient verification.

Accountability mechanisms should include documentation of physician corrections or overrides. When a clinician identifies an inaccuracy, clarifies context, or rejects a suggested insight, that interaction should be logged in a structured manner. This documentation not only preserves legal clarity but also creates an opportunity for iterative system improvement. Correction pathways should be efficient and integrated into workflow rather than burdensome.

Continuous feedback loops are essential to responsible AI deployment. Health systems and developers should establish processes through which clinician feedback is aggregated, analyzed, and used to refine model performance. Patterns of recurrent correction or misunderstanding may signal systemic issues requiring recalibration. Governance should treat clinician interaction as an input to ongoing quality assurance rather than a one-time review.

Workflow integration must be calibrated to avoid cognitive overload. AI tools should reduce friction, not create additional documentation tasks that undermine efficiency gains. Human oversight should be meaningful but streamlined, reinforcing thoughtful engagement without imposing redundant administrative steps. Effective integration requires close alignment with real-world clinical processes.

Ultimately, AI should augment clinical reasoning, not erode it. The objective is to enhance situational awareness, improve informational clarity, and reduce administrative distraction so that physicians can exercise independent judgment more effectively. A governance framework that embeds transparency, traceability, and correction mechanisms into workflow will preserve physician authority while enabling the benefits of technological advancement.

VII. POST-DEPLOYMENT MONITORING AND ADAPTIVE GOVERNANCE

Artificial intelligence systems in clinical environments cannot be treated as static products. Unlike traditional medical devices or fixed software applications, AI models evolve over time and operate within dynamic care settings. Patient populations shift, documentation practices change, and standards of care advance. For these reasons, governance must extend beyond initial deployment and incorporate structured, lifecycle-based oversight.

HHS should require periodic reassessment of model performance, calibrated to risk tier. Assistive tools may warrant routine evaluation at defined intervals, while systems exerting Material Clinical Influence should undergo more frequent and rigorous review. Performance reassessment should examine accuracy, error rates, false-positive and false-negative patterns, and any observable degradation relative to initial validation benchmarks.

Significant adverse events plausibly linked to AI outputs should be subject to structured reporting requirements. Where an AI-generated summary, recommendation, or autonomous action contributes to patient harm or near-miss events, health systems should document and

report those occurrences through appropriate oversight channels. Transparent reporting promotes accountability and facilitates system-wide learning rather than isolated corrective action.

HHS should also encourage structured data collection to evaluate real-world impact. Deployment should not rest solely on pre-launch validation; instead, agencies should promote longitudinal assessment of clinical outcomes, workflow efficiency, patient safety indicators, and potential disparities. Real-world evidence can inform whether systems are delivering promised benefits or introducing unintended consequences.

Clear update protocols are essential to responsible lifecycle governance. When models are modified—whether through retraining, parameter adjustment, or expanded scope—those updates should be documented and, where appropriate, revalidated. Governance should distinguish between minor performance refinements and material functional changes that alter risk profiles. Structured version control and defined approval pathways enhance clarity and stability.

Transparency regarding model revisions strengthens trust among clinicians and patients. Health systems should maintain accessible documentation describing what changes were made, why they were implemented, and how performance was assessed. Silent or opaque updates undermine accountability and may obscure emerging safety signals.

Finally, HHS should consider sunset or formal review triggers for higher-risk deployments. Autonomous or high-impact systems should not operate indefinitely without structured reevaluation. Periodic renewal, contingent upon demonstrated safety and performance, reinforces adaptive governance. By embedding monitoring, transparency, and reassessment into policy architecture, HHS can ensure that innovation remains aligned with patient safety over time rather than relying solely on initial deployment assumptions.

VIII. FEDERAL COORDINATION AND REGULATORY CLARITY

The effectiveness of any AI governance framework depends not only on substantive standards but also on institutional coherence. Within HHS, multiple components exercise jurisdiction over aspects of health care technology, including payment policy, health IT certification, and medical device regulation. Without deliberate coordination, overlapping authorities may produce fragmented guidance, duplicative compliance obligations, and regulatory uncertainty that discourages responsible innovation.

Alignment between CMS, ONC, and FDA is particularly important in the context of AI-enabled clinical tools. CMS influences adoption through reimbursement structures and participation requirements. ONC shapes health IT architecture and certification standards. FDA oversees certain software functions that qualify as regulated medical devices. A coherent federal approach requires these components to operate from a shared risk-tiered framework rather than issuing isolated or inconsistent directives.

Clear boundaries between clinical decision support functions and device-regulated systems are essential. Developers and health systems must understand when a tool is treated as assistive CDS and when it crosses into device territory requiring additional regulatory processes. Ambiguity in these boundaries can chill innovation at the margins and create compliance risk that exceeds the actual safety concern.

HHS should also work to avoid duplicative or conflicting compliance regimes. If transparency requirements under health IT certification diverge from expectations under payment policy or device oversight, stakeholders may face redundant reporting and validation burdens. Harmonized standards reduce administrative cost while maintaining safety protections.

Interoperability and data portability are foundational to safe AI deployment. Clinical briefing tools and related systems depend on accurate, comprehensive data drawn from multiple sources. Federal policy should promote consistent data standards, secure exchange mechanisms, and portability across platforms. Fragmented data environments undermine both performance and patient safety.

Consistency in transparency expectations across federal programs is equally important. If CMS requires certain disclosures for AI-enabled tools used in Medicare settings, those expectations should align with ONC certification requirements and FDA guidance where applicable. A unified transparency baseline strengthens predictability and fosters trust among clinicians and patients.

Coordinated federal policy also improves market signaling. Developers are more likely to invest in compliant, safety-focused systems when regulatory expectations are clear and stable. Conversely, shifting or inconsistent standards encourage minimal compliance strategies and reactive adjustments rather than proactive safety design.

By harmonizing oversight across HHS components, clarifying jurisdictional boundaries, and aligning transparency and interoperability standards, the Department can reduce fragmentation while preserving robust safeguards. Regulatory clarity is not merely an administrative convenience; it is a precondition for safe, scalable, and trustworthy deployment of AI in medical care.

IX. PRESERVING PATIENT TRUST AND ACCESS

Patient trust is the foundation upon which all health care innovation rests. Technological advancement that undermines confidence in clinical integrity will ultimately fail, regardless of its theoretical efficiency gains. HHS must therefore evaluate AI governance not only through a safety lens, but also through its effect on patient trust, system stability, and equitable access to care.

Excessive regulatory uncertainty can deter beneficial innovation. When developers and health systems lack clarity regarding compliance expectations, risk thresholds, or jurisdictional boundaries, investment slows and responsible experimentation is deferred. In health care, where administrative burden is already substantial, additional ambiguity may discourage deployment of tools that could reduce clinician workload and improve care coordination.

At the same time, poorly governed AI systems can erode trust even more quickly. High-profile failures, opaque decision-making, or unexplained clinical outcomes may lead patients to question whether technology has displaced professional judgment. Trust once lost is difficult to restore. A governance framework that anticipates foreseeable risks and embeds transparency mechanisms helps prevent reactive regulatory backlash driven by isolated but visible incidents.

Trust is also closely linked to equity. If AI systems are deployed without adequate validation across diverse patient populations, they risk reinforcing disparities. Performance gaps across demographic groups can magnify existing inequities in diagnosis, treatment, or access. Disciplined governance that requires performance monitoring and correction supports broader, more equitable adoption.

Regulatory clarity encourages responsible private-sector investment. Developers are more likely to allocate capital toward safety-conscious, interoperable systems when standards are predictable and proportionate. Clear tiered requirements allow firms to design compliance pathways from the outset rather than retrofitting systems after deployment. This stability supports long-term innovation rather than short-term opportunism.

Calibrated oversight also promotes affordability. Assistive AI tools that reduce chart-review time, streamline documentation, and minimize duplicative testing can lower administrative overhead across the health system. When governance is proportionate to risk, low-impact tools can deploy without unnecessary friction, enabling efficiency gains that translate into cost containment.

Conversely, overregulation of low-risk assistive functions may entrench existing inefficiencies. If health systems are prevented from adopting technologies that meaningfully reduce clinician burden, administrative costs remain embedded in the delivery system and are ultimately borne by patients and taxpayers. A risk-tiered approach ensures that safeguards do not inadvertently block cost-saving innovation.

In sum, preserving patient trust and expanding access are not in tension with innovation; they are dependent upon disciplined governance. Clear standards, proportional oversight, and transparent accountability mechanisms foster confidence among patients, clinicians, and developers alike. By calibrating regulatory expectations to actual risk, HHS can promote responsible innovation that enhances affordability, broadens access, and strengthens trust in the evolving health care system.

Conclusion

The Center for Regulatory Freedom strongly supports innovation that strengthens—rather than fragments—the primary care foundation of the American health system. Technological advancement should reinforce the physician-patient relationship, reduce unnecessary administrative burden, and improve clinical precision. Artificial intelligence, if deployed responsibly, can serve these ends by enhancing physician awareness and restoring coherence to increasingly complex medical records.

As care delivery models evolve toward greater use of episodic and on-demand services, informational continuity becomes the critical bridge between traditional longitudinal primary care relationships and emerging access pathways. Even where relational continuity declines, the system must preserve the integrity of longitudinal patient information. AI-powered clinical briefing tools offer a means of compressing and organizing that information so physicians can enter encounters informed and prepared.

However, AI must remain a tool of augmentation—not substitution. Properly bounded systems can amplify physician insight and reduce fragmentation. Poorly governed systems risk obscuring accountability, introducing error, and undermining public trust. The policy challenge is therefore not whether to enable AI in health care, but how to calibrate governance to preserve safety and clinical integrity.

CRF urges HHS to adopt a risk-tiered, lifecycle-based framework that distinguishes assistive tools from materially influential or autonomous systems. Permissionless innovation should remain the default for low-risk, reversible functions operating under clear transparency safeguards. At the same time, heightened validation, monitoring, and accountability requirements should apply where AI materially influences diagnosis, triage, or treatment decisions.

A clear, harmonized, and forward-looking regulatory architecture will support responsible innovation, encourage private-sector investment, and protect patients. Alignment across CMS, ONC, and FDA authorities will reduce fragmentation and provide the predictability necessary for sustained technological development. Calibrated oversight can improve efficiency, strengthen affordability, and preserve confidence in clinical decision-making.

CRF remains committed to disciplined, economically grounded, patient-centered regulatory reform. By adopting a balanced framework that protects safety while enabling progress, HHS can ensure that artificial intelligence serves as a continuity-enhancing force in American health care—advancing innovation without compromising the integrity of primary care.

Sincerely,

A handwritten signature in black ink, reading "Andrew M. Langer". The signature is fluid and cursive, with the first name "Andrew" and last name "Langer" being more prominent than the middle initial "M".

Andrew M. Langer
Director
CPAC Foundation Center for Regulatory Freedom