Cleveland Clinic

Melissa Myers, JD, MPA
Vice President, Government Relations

February 23, 2026

The Honorable Thomas Keane, MD, MBA
Assistant Secretary for Technology Policy
Office of the National Coordinator for Health Information Technology
Department of Health and Human Services
Mary E. Switzer Building, Mail Stop 7033A
330 C Street SW
Washington, DC 20201

RE: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as part of Clinical Care (RIN 0955-AA13)

Submitted electronically via: Regulations.gov

Dear Assistant Secretary Keane:

Cleveland Clinic is a not-for-profit, integrated healthcare system dedicated to patient-centered care, teaching, and research. With a footprint in Northeast Ohio, Florida, and Nevada, Cleveland Clinic Health System operates 23 hospitals, including a main campus near downtown Cleveland, with more than 6,700 beds and 280 outpatient locations. Cleveland Clinic employs over 5,700 physicians and researchers, and 16,800 nurses. Last year, our system had 15.7 million patient encounters, including 14.1 million outpatient visits and 333,000 hospital admissions and observations.

Cleveland Clinic appreciates the opportunity to provide feedback to the Department of Health and Human Services (HHS) and the Assistant Secretary for Technology Policy and Office of the National Coordinator (ASTP/ONC) regarding the adoption and use of artificial intelligence (AI) in the delivery of clinical care. This RFI represents a practical next step in translating federal AI strategies into policy that directly shapes real-world clinical adoption.

Cleveland Clinic recognizes the current and critical inflection point for clinical AI, particularly in primary care and other high-volume, lower-acuity settings. Advances in AI models and workflow automation are expanding the range of clinical and operational tasks that can be supported through software, while the U.S. health care system continues to face persistent access constraints, rising costs, and clinician burnout. From Cleveland Clinic's perspective, the greatest near-term opportunity for clinical AI lies in end-to-end enablement of care delivery rather than isolated point solutions. When deployed with appropriate oversight, AI can support intake, triage, routine care, documentation, navigation, and follow-up in ways that expand access, improve continuity, and reduce provider burden. Within this context, Cleveland Clinic's perspectives are detailed below.

<u>Section I: Framing the Opportunity for Clinical AI in Care Delivery</u>

Regulation

To establish trust and encourage adoption, clear standards must be established to define how AI is used in the delivery of patient care and expected outcomes. A regulatory framework that effectively advances the adoption of AI in clinical care while setting protections for patient safety and investment should encompass:

- **Stronger interoperability frameworks and enforcement** across the health care ecosystem, particularly to support transitions of care across settings and health systems.

- **A clear, risk-tiered clinical AI framework** to evaluate AI tools, including those that operate without a human in the loop, distinct from traditional medical device pathways. Core regulatory elements that enable safe scaling should include:

    - Clearly defined pilot guidelines with escalation and oversight protocols and explicit triggers for clinician review, override, and intervention.
    - Scaled commercialization only after demonstrating performance against defined quality benchmarks.
    - Minimum quality standards requiring performance comparable to representative physician benchmarks in real-world clinical practice, evaluated through pilot data and independent physician review using like-for-like clinical scenarios.
    - Continuous performance monitoring to detect degradation and prevent model drift, with adverse events investigated using processes comparable to existing medical device oversight.

- **Expedited review pathways** for new clinical AI innovations that meet defined risk and quality thresholds.

- **Data ownership and stewardship protections** that support provider investment in data quality, governance, and responsible reuse.

Reimbursement

Payment approaches must recognize and appropriately compensate physicians and providers for the investment, risk, and ongoing innovation required to deliver AI-enabled care. This includes not only clinician time, but also the costs of implementation, workflow redesign, data governance, and clinical oversight. While multiple reimbursement and compensation models may be viable, payment policy should remain flexible rather than constraining innovation to a single structure. Aligning incentives so that providers are rewarded for responsible deployment of AI in clinical practice is essential to driving adoption, improving access, and sustaining high-quality care. This can be achieved in the following ways:

- **Telehealth parity with in-person services**, recognizing that in-person escalation should be the final step of care rather than the default starting point.

- **Payment for services provided or augmented by AI without a human in the loop**, reflecting the clinical value delivered and the oversight responsibilities borne by physicians and providers.

- **Engagement with commercial payers to promote coverage for clinically validated AI tools**, reducing variability and uncertainty that can otherwise limit adoption.

Research and Development
National investment in applied AI research and development is crucial to building the evidence base to support the safe and effective use of AI in clinical care. Key strategies to foster innovation, ensure effective integration, and maintain leadership in the development and deployment of clinical AI within the United States should include:

- **Non-dilutive grants for innovators** building clinical AI applications, alongside **provider pilot grants** to support real-world testing and implementation.

- **Federal investment in applied research, validation, and workforce readiness**, accelerating translation of promising AI tools from concept to routine clinical use.

- **Protections for U.S.-developed clinical AI intellectual property and data**, promoting fair competition and long-term leadership in responsible clinical AI.

Finally, interoperability and data stewardship are foundational across regulation, reimbursement, and research. From a regulatory perspective, interoperability standards enable longitudinal data access and integration across systems. From a reimbursement perspective, providers incur significant costs and labor to extract, clean, and organize clinical data, and adoption depends on whether payment and incentive structures support that investment. From a research and development perspective, interoperable, well-governed data is essential for training, validation, and scaling of clinical AI. Policies that standardize access while supporting provider investment and stewardship are therefore critical to aligning incentives, improving data quality, and sustaining innovation in clinical AI.

Section II: Responses to RFI Questions

Question 1: What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

Several interrelated barriers limit private sector innovation and adoption of clinical AI. To begin with, regulatory frameworks remain fragmented and insufficiently predictable, with no clear AI-specific pathway to support rapid clinical adoption and safe scaling. This uncertainty complicates investment decisions and slows deployment of clinically promising tools.

Compounding regulatory uncertainty, reimbursement incentives are misaligned. Under current payment policies, providers are often not reimbursed for care delivered with or augmented by AI, making it difficult to justify sustained investment. Given already-thin hospital margins, AI initiatives that improve quality or access but lack a clear revenue or cost-offset pathway are challenging for health systems to sustain at scale.

Beyond payment, data-related challenges remain substantial. Clinical data are frequently fragmented, inconsistent, or of variable quality, and interoperability across systems, settings, and vendors remains limited. Providers appropriately prioritize patient privacy, security, and responsible data governance. However, the significant cost and labor required to meet these standards are not consistently supported by aligned incentives, which can limit providers' ability to invest in high-quality, privacy-preserving data preparation. This, in turn, constrains data quality, slows equitable adoption of clinical AI, and increases the risk that models are trained or deployed without sufficient real-world validation.

At the same time, implementation barriers are high. Scaling clinical AI requires meaningful up-front investment in infrastructure, engineering, governance, and change management. Integrating AI into existing clinical workflows, including in federated analytics environments, can be complex and resource-intensive without adequate institutional support.

Finally, adoption depends on trust and workforce readiness. Limited transparency, explainability, and confidence in AI outputs can deter clinician use, while gaps in AI literacy, training, and organizational change management further constrain uptake. Taken together, these barriers mean that even technically strong AI tools often struggle to move from pilots to sustained, system-wide adoption.

Question 2: What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.

To address the interlocking barriers described above, HHS should prioritize policy changes that reduce regulatory ambiguity, align incentives, and enable responsible scaling of clinically valuable AI. From a regulatory perspective, HHS should establish a clinical AI-specific, risk-tiered regulatory framework distinct from traditional medical device pathways. This framework should clarify evaluation and monitoring expectations across HIPAA, HITECH, CMS, and ASTP/ONC; harmonize agency requirements; and provide predictable pathways for approval, updating, and scaling of AI tools, including those that operate without a human in the loop.

In parallel, reimbursement policy should be modernized to support adoption of high-value clinical AI. Under current fee-for-service structures, payment often lags innovation, discouraging investment in tools that improve care but do not immediately generate revenue or reduce utilization. HHS and CMS should therefore advance payment approaches that recognize clinically valuable AI-enabled care while allowing flexibility across reimbursement models.

Programmatic support can further reduce adoption friction. HHS should invest in federal pilot programs and provider grants that defray the cost and complexity of adoption, including workflow integration, system integration, and local validation. HHS should also address contractual barriers to responsible AI use, such as overly restrictive “no-AI” clauses embedded in clinical systems.

Finally, HHS should advance national data infrastructure and interoperability, paired with incentives that support provider participation in data reuse. Standardized data access, benchmarking infrastructure, and bias-mitigation resources are essential to scaling AI safely, while policy should recognize the significant cost and labor required to prepare clinical data for AI development and deployment.

Question 3: For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (e.g., relating to liability, indemnification, privacy, and security). What novel legal and implementation issues exist and what role, if any, should HHS play to help address them?

Non-medical-device clinical AI raises governance, legal, and implementation issues not fully addressed by existing frameworks. In practice, AI capabilities are increasingly embedded within software platforms and clinical systems without consistent disclosure, while responsibility for model behavior, updates, and performance is often poorly defined across vendors and providers. Contractual restrictions can further limit responsible quality and safety use of AI in care delivery.

To address these gaps, HHS should establish vendor disclosure requirements, promote standardized contract language, and clarify shared accountability across developers and deploying organizations. In particular, HHS should define governance and liability expectations for AI vendors supplying models, agents, or AI-enabled features to health systems, including responsibilities for ongoing monitoring and rapid remediation of performance degradation.

Liability and indemnification remain central challenges. Because AI outputs frequently influence but do not dictate clinical decisions, ambiguity persists around responsibility when harm occurs. HHS should therefore define how liability applies to assistive versus more autonomous AI and establish safe harbors for providers and vendors that meet baseline governance, monitoring, and documentation standards.

AI also introduces novel privacy and security risks, including inference-based re-identification, derivative data generation, and multimodal processing not fully contemplated by existing HIPAA guidance. Accordingly, HHS should update guidance to clarify standards for use of protected health information (PHI) in ongoing model training and updating, including when health-system and patient consent is required.

Finally, large language models (LLM) introduce distinct risks. At present, there are no widely adopted, scalable methods to evaluate, monitor, and govern LLM performance in clinical settings. HHS should prioritize development of national standards and tools for LLM evaluation, monitoring, and safety that can be implemented consistently across health systems.

Question 4: For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and other workflow and human-centered evaluation methods for clinical care? Should HHS further support these processes? If so, which mechanisms would be most impactful (e.g., contracts, grants, cooperative agreements, and/or prize competitions)?

Effective evaluation of non-medical-device AI requires a full lifecycle approach integrating pre-deployment validation, post-deployment monitoring, and human-centered workflow assessment.

Prior to deployment, evaluation should include standardized documentation such as model cards describing intended use, data sources, assumptions, and limitations; champion-challenger testing under identical conditions; robustness and stress testing, including subgroup and bias analyses; and transportability testing across sites, populations, and workflows. Workflow-integrated usability testing with clinicians is essential to ensure interpretability, reduce cognitive load, and support safe decision-making.

Once deployed, continuous quality monitoring becomes critical. This includes drift detection, recalibration triggers, performance thresholds tied to clinical outcomes, and criteria for model pause, retraining, or retirement. Auditability and traceability, including immutable logs of data access, model versions, and inference outputs, are necessary to support accountability and incident investigation. Shadow-mode trials and structured qualitative feedback loops help identify unintended consequences before full activation.

HHS should support these processes through nationally aligned evaluation frameworks. The most impactful mechanisms include contracts and cooperative agreements to build shared evaluation infrastructure; competitive grants for provider-led pilots; prize competitions to incentivize novel evaluation and explainability methods; and national standards for documentation, monitoring, and risk-tier classification of non-device AI.

Question 5: How can HHS best support private sector activities (e.g., accreditation, certification, industry driven testing, and credentialing) to promote innovative and effective AI use in clinical care?

HHS can best support private-sector innovation by strengthening the ecosystem for accreditation, certification, and independent testing. As a starting point, HHS should set clear national baseline expectations for safety, transparency, implementation, monitoring, and bias mitigation so that private accrediting bodies can align around a consistent floor.

Building on that foundation, HHS should enable independent, trustworthy testing by supporting neutral testbeds, including universities, consortia, and public-private partnerships, that evaluate robustness, fairness, drift, and workflow fit using realistic clinical scenarios. HHS should also recognize and incentivize credentialing and training programs for clinicians, analysts, and governance teams, including voluntary HHS-recognized certifications for organizations demonstrating mature AI governance.

Targeted funding mechanisms, including grants, cooperative agreements, and prize challenges, can help scale effective validation tools, testing frameworks, and credentialing programs. Finally, convening standards bodies, health systems, developers, and certification organizations can reduce fragmentation and promote interoperable, industry-wide approaches to accreditation and assurance.

Question 6: Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?

AI has met or exceeded expectations in domains where data are relatively structured, workflows are standardized, and model outputs map clearly to clinical or operational decisions. For example, sepsis and early-warning models have improved early recognition and outcomes when tightly integrated into workflows with clear escalation pathways. In diagnostics, AI applications regulated as medical devices have increased throughput and improved consistency in radiology. Provider enablement tools such as ambient documentation have reduced administrative burden and improved clinician productivity when deployed with appropriate safeguards.

Operational applications have also performed well. AI-enabled forecasting of admissions, discharges, surgical flow, bed needs, and staffing has generated measurable efficiency and cost gains. Revenue cycle management tools supporting prior authorization and billing workflows have reduced administrative friction in persistent pain points for providers.

At the same time, AI has fallen short where foundational conditions are weak. Fragmented data environments, or environments where the meaning of clinical data changes over time or across settings, degrade model performance and erode clinician trust. Tools deployed without clearly defined success metrics, action pathways, or monitoring plans often fail to deliver sustained value. Clinical event forecasting models illustrate this challenge, as performance metrics such as area under the receiver operating characteristic (AUROC) can overestimate real-world utility while masking downstream burden from false alerts when guardrails are insufficient.

The most promising next generation of AI tools are those that span care episodes, integrate multimodal data, and deliver actionable insights embedded directly into clinical workflows. This includes both autonomous and AI-augmented care models. In many cases, significant value is created not necessarily by fully replacing clinician encounters, but by AI systems that perform intake, synthesize patient history and signals, stratify risk, surface relevant evidence, and prepare care plans in advance of a clinician interaction. These tools can materially improve visit efficiency, diagnostic accuracy, and clinician experience, while preserving appropriate human oversight. Trajectory-focused models that predict deterioration, recovery, or functional decline across settings further extend this value by supporting proactive, rather than reactive, care delivery.

In parallel, there are substantial opportunities in AI-enabled care management and operational intelligence. Tools that support transitions of care, longitudinal follow-up, adherence, and patient education address persistent gaps that drive avoidable utilization and poor outcomes under episodic care models. Administrative and operational applications are similarly critical to scalability and return on investment, including AI-enabled prior authorization, revenue cycle workflows, and forecasting of admissions, discharges, bed needs, and staffing. Across these cases, Cleveland Clinic's experience suggests that AI delivers the greatest impact when success metrics, caregiver action pathways, and monitoring plans are clearly defined, and when tools are deployed with guardrails that translate predictions into reliable, workflow-appropriate action.

Question 7: Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?

Adoption of clinical AI is shaped by a broad set of stakeholders. Senior clinical leaders, enterprise AI governance bodies, and boards of directors influence decisions related to patient safety and institutional risk. Legal, compliance, privacy, and cybersecurity leaders shape acceptable deployment models and contractual terms. Data governance teams determine data readiness and monitoring capabilities, while operational leaders assess workflow impact and scalability. Frontline physicians exert substantial influence through their trust and day-to-day use of AI tools.

Administrative hurdles include fragmented governance with unclear decision rights, inconsistent data quality, privacy and contractual constraints, and limited monitoring infrastructure. Integration into electronic health records (EHRs) and clinical workflows remains technically complex, particularly in federated environments. Gaps in AI literacy, training, and change management further slow adoption. Overly centralized, executive-driven deployment can create risk when implementation outpaces governance and clinical engagement.

Question 8: Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

Enhanced interoperability would have outsized impact by unlocking high-value data types currently siloed or non-computable, including unstructured notes, imaging, pathology, genomics, and real-time device data. Stronger adoption and enforcement of shared data standards such as FHIR read/write, HL7, SNOMED, LOINC, and OMOP would improve portability of features, labels, and clinical logic across systems and vendors.

Enterprise-grade semantic and operational interoperability across EHRs, operational data platforms, quality measurement systems, and external data sources is essential to moving AI beyond local pilots. Standardized system-to-system integration points, including common APIs and data models, would reduce integration burden and widen market access.

Interoperable benchmarking and monitoring infrastructure, including shared datasets, standardized outcome definitions, performance metrics, model cards, and drift-detection tools, would support transparent evaluation, cross-institution research, and continuous improvement. Together, these advances would reduce friction, fuel applied research, and accelerate development of AI that is portable, verifiable, and workflow ready.

Question 9: What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?

AI provides a means to improve coordination, reduce delays, personalize care, and relieve administrative and caregiving burden, such as repetitive data entry, scheduling difficulties, and gaps in follow-up across care settings. AI can provide welcome relief by handling routine or burdensome tasks that detract from time spent on direct human interaction and care.

At the same time, reservations persist regarding privacy, fairness, transparency, reliability, and preservation of human-centered care. Concerns include whether AI-enabled care meets appropriate quality standards, whether errors or bias will be detected and corrected, and whether patients are adequately informed about AI use. Clear guardrails, transparency expectations, and accountability standards are essential to sustaining trust.

Question 10: Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care? Are there published findings about the impact of adopted AI tools and their use in clinical care? How does the literature approach the costs, benefits, and transfers of using AI as part of clinical care?

HHS should prioritize AI research that strengthens the evidence base for safe and effective clinical adoption. Primary care is a priority, where AI can expand access, improve continuity, and reduce provider burden through triage, routine care support, and proactive follow-up. Transition of care is another high-impact area, where AI can improve coordination and follow-up across fragmented settings.

The literature demonstrates meaningful benefits of AI when deployed with appropriate oversight, including improved screening and triage, faster response times, reduced administrative burden, and more consistent care delivery. Evidence also supports AI-enabled asynchronous functions such as messaging, intake, and scheduling, as well as ambient tools that improve documentation quality and clinician efficiency.

At the same time, research often under-characterizes real-world performance, equity impacts, and full lifecycle costs. HHS should support research on bias mitigation, generalizability across populations and settings, and economic evaluation of how costs, benefits, and value are transferred across clinicians, health systems, and technology vendors as AI redistributes work across the care delivery ecosystem.

Thank you for conducting a thoughtful process that allows us to provide input on such important issues and for your consideration of this feedback. Should you need any further information, please contact me at myersm21@ccf.org.

Sincerely,

Melissa Myers, JD, MPA
Vice President, Government Relations