

Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care

I. Executive Summary

ALIGNMT AI is an AI governance platform provider that works with health systems, payers, regulatory bodies, and AI developers to operationalize AI governance frameworks, risk assessment, and accountability structures for healthcare AI. We offer these comments from the perspective of practitioners who see daily the gap between the theoretical promise of healthcare AI and the governance and oversight realities on the ground.

We are broadly supportive of HHS's ambition to accelerate AI adoption in clinical care. We are equally convinced that acceleration without robust governance is not a shortcut - it is a liability. The history of health IT is replete with examples of rapidly deployed technology that generated unintended harm, eroded clinician trust, and ultimately set back adoption rather than advancing it.

Our comments focus primarily on two areas where we believe HHS can take the most consequential action:

- Standards and oversight: Establishing a coherent, risk-proportionate, and enforceable governance framework for clinical AI - particularly for the large and growing category of non-medical device AI tools - that reduces uncertainty for developers while providing meaningful accountability for patients.
- Patient safety and equity: Ensuring that the imperative to accelerate AI adoption does not come at the cost of the populations most likely to be harmed by poorly governed AI, including patients from underserved communities, those with disabilities, and those with limited health literacy.

We urge HHS to resist the false choice between speed and safety. The best governance frameworks are enabling, not obstructive - they create the predictability and trust that sustainable AI adoption requires.

Our core recommendation: HHS should establish a tiered, risk-proportionate oversight framework for clinical AI that is coherent across FDA, ONC, and CMS - backed by quality and equity requirements, transparency standards, and meaningful post-market surveillance - and that applies to all AI tools used in consequential clinical decisions, not only those that meet the FDA's definition of a medical device.

II. Barriers to Innovation (Question 1)

From our experience advising health systems and developers, the most consequential barriers to responsible AI adoption in clinical care are governance-related, not technical:

Regulatory Fragmentation and Ambiguity

The fundamental governance problem facing clinical AI in the United States is jurisdictional incoherence and incompleteness. FDA regulates AI that meets the definition of Software as a Medical Device (SaMD) under 21 C.F.R. Part 880. ONC shapes health IT through certification requirements under 45 C.F.R. Part 170. CMS influences behavior through payment policy. Each operates on a different timeline, with different requirements, and for different subsets of the AI ecosystem.

The result is that a large and clinically consequential category of AI tools - non-device clinical decision support (CDS) that does not meet FDA's criteria under 21 U.S.C. §399d - operates in a near-regulatory vacuum. These tools may analyze patient records, generate differential diagnoses, recommend treatment pathways, flag patients for escalation, guide provider or patient queries, suggest care plans, or provide clinical reference information. This category also includes AI tools offered directly to patients and consumers — whether through patient portals, mobile health applications, or direct-to-consumer platforms — that generate symptom assessments, triage recommendations, care navigation guidance, or health coaching, and tools offered by providers to patients as an extension of the clinical encounter. These tools occupy the same regulatory vacuum as provider-facing non-device CDS, despite often operating with less clinical oversight and reaching populations with greater health literacy variability. They are used in life-affecting decisions. Yet they face no mandatory evaluation requirements, no performance standards, no transparency obligations, and no post-deployment surveillance or performance measurement requirements.

This is not primarily a barrier to innovation. It is a patient safety gap that also, secondarily, creates business risk for health systems that deploy these tools and reputational risk for developers whose products underperform in ways they did not anticipate.

Liability Ambiguity Creates Perverse Incentives

The current liability environment simultaneously discourages adoption by cautious health systems or clinicians and fails to adequately incentivize responsible development by AI vendors. When a clinical AI tool contributes to a harmful outcome, there is no settled legal framework governing how responsibility is allocated across the developer, the deploying health system, and the treating clinician. This ambiguity causes legal departments to counsel non-adoption, even when the clinical evidence for a tool is strong - and it means that developers face limited accountability when their tools cause harm.

Clear governance frameworks - including defined standards of care for AI deployment, transparency requirements, and appropriate safe harbors for compliant adopters - would benefit both sides of this equation.

Inequitable Access to Governance Infrastructure

Smaller health systems, rural providers, federally qualified health centers (FQHCs), community health centers (CHCs), physician and clinician practice groups, generally lack the internal capacity to conduct meaningful AI governance: they cannot independently validate vendor claims, audit models for quality or bias, or establish robust post-deployment surveillance. This creates an equity paradox: the organizations that serve the most vulnerable patients are least equipped to protect those patients from poorly performing AI - and least able to capture AI's benefits from a position of informed decision-making.

III. Regulatory Priorities (Question 2)

We recommend HHS treat the following regulatory actions as the highest-priority interventions for accelerating responsible AI adoption:

3.1 Establish a Unified, Tiered Federal Framework for Clinical AI Governance

The most impactful single action HHS could take is to develop - in collaboration across FDA, ONC, and CMS - a unified risk-tiered framework that applies consistently to all AI tools used in consequential clinical decisions, regardless of whether those tools qualify as FDA-regulated medical devices.

We recommend a three-tier structure for classifying healthcare AI tools based on their purpose of use for risk management:

- Tier 1 — Autonomous or high-stakes decision AI: Tools that make or substantially determine clinical decisions with limited clinician override capacity. These should be subject to the most rigorous evaluation requirements, including prospective clinical validation, mandatory demographic performance disaggregation, and active post-market surveillance.
- Tier 2 — Clinical decision support with clinician review: Tools that generate recommendations reviewed and acted upon by a clinician. These should require pre-deployment validation documentation, transparency disclosures, and periodic performance reporting. Patient and consumer-facing AI tools — including those offered through providers via patient portals or care management platforms and those offered directly to consumers outside a clinical relationship — require explicit placement within this framework. Where such tools generate clinical recommendations, triage decisions, or care guidance that a patient acts upon without an intervening clinician review, they should be classified at minimum as Tier 2 and evaluated accordingly, notwithstanding the absence of a traditional provider-patient workflow. Direct-to-consumer tools that operate fully outside any provider relationship warrant additional consideration given the reduced safety net available when recommendations are misunderstood or misapplied.

- Tier 3 — Administrative and workflow AI: Tools that do not directly affect clinical decision-making (e.g., scheduling, coding, documentation assistance). These warrant lighter oversight focused primarily on data privacy and security.

This framework would not require new statutory authority for many of its elements - it could be implemented through ONC certification criteria (45 C.F.R. §170.315), CMS conditions of participation, requirements for AI developers under the CMS Health Tech Ecosystem, and HHS-wide guidance. It would also provide the clarity that developers and health systems need to plan for compliance.

3.2 Mandatory Transparency and Disclosure Standards for Tier 1 and Tier 2 Tools

Health systems currently evaluate AI vendors with almost no standardized information required from the developer. Vendors make performance claims using their own metrics, their own test sets, and their own definitions. There is no equivalent of the FDA's Summary of Safety and Effectiveness Data (SSED) for non-device clinical AI.

HHS should require, as a condition of use in any health systems, hospitals, physicians, or clinicians funded by or in receipt of federal funds related to the provision of designated health services as defined by 42 CFR 411.351, that developers of Tier 1 and Tier 2 clinical AI tools provide standardized transparency documentation including:

- Training data characteristics: source, date range, demographic composition, and any known gaps or biases.
- Validation methodology: whether validation was prospective or retrospective, internal or external, and the characteristics of the validation population.
- Performance metrics: disaggregated by race, ethnicity, sex, age, primary language, and insurance status at minimum.
- Known limitations and contraindications: specific patient populations or clinical contexts where the tool has demonstrated degraded performance.
- Update and versioning history: what has changed in each model version and how changes have been validated.

This standard should be developed in collaboration with NIST (building on the AI Risk Management Framework), standards development organizations (HL7, IEEE), and industry, and should be codified as an ONC certification criterion under 45 C.F.R. §170.315.

3.3 Post-Market Surveillance Requirements

The FDA's real-world performance monitoring framework for AI/ML-based SaMD represents a sophisticated and practical model for post-market surveillance. HHS should develop an analogous framework for non-device clinical AI that requires health systems and developers to monitor and report on:

- Model drift: changes in AI tool performance over time as patient populations, clinical practices, or data inputs change.

- Disparate impact monitoring: ongoing tracking of whether performance gaps across demographic subgroups are emerging, widening, or narrowing post-deployment.
- Adverse event reporting: a mechanism analogous to MedWatch by which clinicians and patients can report suspected AI-related clinical errors.

This surveillance infrastructure should be administered through a multi-agency accountability structure that reflects the actual reach of federal authority over the AI developer community. ASTP/ONC is a logical administrative home for surveillance reporting — building on existing infrastructure such as the Health IT Feedback and Inquiry Portal — and reporting obligations should be linked to any standing an HIT developer holds under ONC's certification program. However, ONC's reach is limited to certified HIT developers, and certification itself is voluntary except where it conditions provider participation in Promoting Interoperability programs. Pure AI developers and direct-to-consumer health IT vendors have little practical motivation to seek ONC certification and therefore fall outside that accountability structure. To achieve the reach necessary for meaningful post-market surveillance across the full AI developer community, CMS should establish parallel reporting obligations tied to participation in CMS's Health IT Ecosystem as a recognized vendor or developer — conditions that carry direct market consequences for a broader class of AI developers than ONC certification alone. Additionally, HHS should engage OIG as an enforcement backstop, recognizing that OIG holds investigative and enforcement authority over certified HIT developers in matters involving Information Blocking and fraud, and that this authority could extend meaningfully to material misrepresentations in AI performance disclosures. No single agency has sufficient reach to cover the full landscape; a coordinated ONC-CMS-OIG framework is necessary to ensure that surveillance obligations have genuine teeth across the developer community.

3.4 HIPAA Modernization for Responsible AI Development

The current HIPAA Privacy Rule (45 C.F.R. Parts 160 and 164) was designed before large-scale machine learning was foreseeable. Its application to AI training and validation creates genuine compliance uncertainty. HHS should issue guidance clarifying:

- That training and validating AI models on de-identified data meeting the standards of 45 C.F.R. §164.514 does not constitute impermissible use or disclosure, and that a business associate covered entity's AI governance program constitutes a legitimate "healthcare operations" function under §164.501.
- That federated learning approaches — in which model training occurs at the site of the data without raw data leaving institutional control — satisfy the Privacy and Security Rule's requirements for minimum necessary and authorized use requirements when implemented with appropriate technical security and administrative security controls.
- That patients must be meaningfully informed when AI plays a material role in their diagnosis or treatment, consistent with the principle of informed consent, and that this notification requirement applies regardless of whether the AI tool is FDA-regulated.

IV. Novel Legal and Implementation Issues (Question 3)

From our governance advisory practice, we identify the following as the most consequential unresolved legal and implementation challenges for non-device clinical AI:

4.1 The Accountability Gap

The central legal problem with non-medical-device clinical AI is the accountability gap: there is no party who is clearly responsible for ensuring the tool performs safely across the full range of patients and clinical contexts in which it is deployed. Developers disclaim liability through contractual terms of use. Health systems and clinicians rely on vendor performance representations they lack the expertise to verify. Clinicians are expected to exercise independent judgment but are rarely given the information they need to calibrate their trust in an AI tool. This problem is made materially worse by ASTP/ONC's proposed HTI-5 rule, which removes requirements for exactly this kind of performance and transparency information for developers certifying to 170.315(b)(11) — apparently on the mistaken assumption that such information is not useful to end users in their clinical workflow. Transparency documentation about model training, validation methodology, known limitations, and demographic performance is not primarily a point-of-care reference tool for end users; it is essential infrastructure for health systems and clinicians as buyers and implementers of clinical AI — necessary for responsible evaluation, procurement, governance, and ongoing oversight. Removing these requirements from the certification framework at the moment when clinical AI deployment is accelerating most rapidly compounds the accountability gap rather than addressing it, and HHS should treat the HTI-5 proposal on this point as a cautionary example of how regulatory action can inadvertently erode the governance foundation that responsible AI adoption requires.

HHS should take the following steps to close this gap:

- Establish, through CMS Conditions of Participation (42 C.F.R. Parts 482 and 485), that accredited hospitals and health systems bear primary accountability for ensuring that AI tools deployed in clinical care have been appropriately evaluated, monitored, and documented. This places responsibility with the deploying institution — which has the most context about its patient population — rather than exclusively with developers.
- Clarify, through OIG guidance modeled on existing safe harbor provisions at 42 C.F.R. §1001.952, that health systems that deploy AI tools in compliance with HHS-established governance standards have an affirmative defense in regulatory proceedings arising from AI-related adverse events. This creates a compliance incentive without eliminating accountability.
- Require AI developers as part of condition of use in any health systems, hospitals, physicians, or clinicians funded by or in receipt of federal funds related to the provision of designated health services as defined by 42 CFR 411.351 to provide contractual representations regarding model performance, update notification, and post-deployment support, with HHS enforcement authority for material misrepresentations.
- Direct ASTP/ONC to retain and strengthen — rather than reduce — the transparency and risk management documentation requirements currently embedded in the certification program under 170.315(b)(11) for predictive decision support interventions. As discussed above, the proposed HTI-5 rule would eliminate requirements that are

foundational to responsible AI procurement, evaluation, and governance by health systems and clinicians. These requirements should be preserved as a floor, not removed, and should be treated as a governance asset of the certification program rather than an administrative burden on developers.

- HHS should further require that any HIT or AI developer seeking recognition as a participating vendor or developer within CMS's Health IT Ecosystem meet equivalent transparency and risk management documentation standards as a condition of that participation — ensuring that developers who have not pursued or do not hold ONC certification are nonetheless subject to comparable disclosure obligations when they seek access to the federal market through CMS-recognized channels.

4.2 Informed Consent and Patient Notification

There is currently no federal requirement that patients be informed when AI plays a material role in their clinical care. This is a meaningful gap: patients have a right to understand how decisions about their care are made, and that right is not contingent on whether a tool is FDA-regulated. HHS should work with OCR to develop a practical patient notification standard — not a consent form, but a clear notification obligation — that applies to Tier 1 and Tier 2 AI tools.

A distinct but related issue arises with patient and consumer-facing AI tools, where the patient or consumer is the direct user rather than the subject of an AI-assisted provider decision. In these contexts, disclosure obligations run directly to the individual using the tool, and must address not only that AI is involved but the nature of its limitations, the absence of a clinical relationship, and the appropriate role of the tool's output relative to professional medical advice. HHS should clarify that notification standards apply with equal force to consumer-facing tools offered through CMS's Health IT ecosystem, including those offered by providers to patients and those offered directly to consumers through app marketplaces or digital health platforms.

4.3 Civil Rights and Algorithmic Discrimination

Section 1557 of the Affordable Care Act (42 U.S.C. §18116) prohibits discrimination in HHS-funded health programs on the basis of race, color, national origin, sex, age, and disability. AI tools that systematically underperforms for protected groups — due to non-representative training data, proxy variables that correlate with protected characteristics, or differential deployment across patient populations — may constitute discriminatory practices under this statute.

The HHS Office for Civil Rights (OCR) should issue enforcement guidance clarifying that health systems have a proactive obligation to evaluate AI tools for disparate impact on protected groups before deployment and on an ongoing basis. This is not a new legal standard — it is an application of existing law to a new technology context. Clear guidance would provide health systems with the framework to comply and would create meaningful accountability for discriminatory AI deployment.

Governance Principle: Accountability for AI in clinical care must be distributed across the development, deployment, and monitoring lifecycle — not concentrated at any single point. HHS governance frameworks should reflect this by placing complementary obligations on developers, deploying health systems, and ongoing monitoring processes.

V. AI Evaluation Methods (Question 4)

Evaluation of non-device clinical AI should occur across three phases: pre-procurement, pre-deployment, and post-deployment. HHS can play a meaningful role in supporting all three.

5.1 Pre-Procurement Evaluation

Health systems and providers, particularly smaller organizations and physician/clinician practices, lack standardized tools for evaluating AI vendor claims. HHS should support the development of:

- A publicly available AI procurement checklist and evaluation rubric, developed by ASTP/ONC in collaboration with CMS under their Health Tech Ecosystem and health system associations, that provides health systems with a minimum standard of due diligence to apply before contracting with AI vendors.
- A national AI tool registry (modeled on ClinicalTrials.gov and administered by NIH or ASTP/ONC) where developers register clinical AI tools, submit transparency documentation, and post validation evidence. This allows health systems, researchers, and payers to compare tools on a common informational basis.

5.2 Pre-Deployment Validation

Before deploying a clinical AI tool in a live clinical environment, providers should conduct local validation on a representative sample of their own patient population. HHS should:

- Develop and publish standardized validation protocols for common clinical AI use cases (sepsis prediction, imaging analysis, medication management, etc.) that providers can apply without needing to design evaluations from scratch.
- Require, as a condition of use in CMS-funded programs, that Tier 1 AI tools undergo prospective shadow mode validation — in which the AI generates recommendations that are tracked but not acted upon — for a defined period before live deployment, with pre-specified performance thresholds as go/no-go criteria.
- Fund, through AHRQ grants, independent evaluation support for FQHCs, rural health clinics and independent clinician groups, and safety-net providers that lack internal validation capacity.

5.3 Post-Deployment Monitoring

Post-deployment monitoring is the most underinvested phase of AI governance. HHS should encourage that providers deploying Tier 1 and Tier 2 AI tools maintain and periodically review:

- Model performance dashboards tracking accuracy, alert acceptance rates, clinician override rates, and demographic performance disaggregation on an ongoing basis.
- Structured incident reporting processes for suspected AI-related clinical errors, with a defined escalation pathway to the AI developer and, for serious adverse events, to HHS. This reporting regime should be modeled after the Patient Safety Organization (PSO) program established under the Patient Safety and Quality Improvement Act of 2005 and administered by AHRQ, which provides a proven federal framework for the voluntary collection and analysis of patient safety event data reported by healthcare providers, with legal protections for reported information that encourage candid disclosure. Under that model, AHRQ certifies and lists PSOs that receive provider-reported safety data, aggregate it through standardized Common Formats, and contribute non-identifiable data to the Network of Patient Safety Databases (NPSD) for systemic learning. HHS should direct AHRQ to extend the PSO program's Common Formats and reporting infrastructure to expressly capture AI-related adverse events and near-misses as a discrete event category, enabling health systems to report suspected AI-related clinical errors through existing PSO relationships with the same legal protections that apply to other patient safety work product — while ensuring that serious adverse events trigger a defined escalation pathway to the AI developer and to HHS for regulatory follow-up. This approach leverages existing infrastructure rather than creating a parallel reporting bureaucracy, and it connects AI safety surveillance to the broader patient safety learning system in a way that can generate actionable evidence at scale.
- Annual re-validation against updated local data to detect model drift.

VI. Patient Safety and Equity (Questions 3, 9)

We want to be direct: AI governance that does not explicitly center on patient safety and health equity is incomplete governance. The populations most likely to be harmed by poorly governed AI - patients from racial and ethnic minority communities, patients with disabilities, patients with limited English proficiency, patients in underserved rural and urban settings — are also the populations least likely to have advocates within health systems' AI procurement and governance processes.

6.1 Mandatory Equity Evaluation as a Governance Standard

Patient and consumer-facing AI tools present a compounded equity risk. These tools are increasingly marketed to and used by patients managing chronic conditions, navigating care access, or seeking health information — populations that disproportionately include individuals with limited health literacy, limited English proficiency, and limited access to professional clinical guidance to contextualize AI-generated recommendations. HHS equity standards should explicitly apply to direct-to-consumer clinical AI and to tools offered by providers through patient-facing channels, with particular attention to whether these tools have been validated for the literacy levels, languages, and access contexts of the populations likely to use them.

HHS should establish, as a federal standard applicable to all AI tools used in Tier 1 and Tier 2 clinical decisions in healthcare providers in receipt of federal funds for provision of designated health services" that are overseen by the FDA as SaMD, by ASTP/ONC as certified products to 170.315(b)(11), or as are provided by AI developers participating in CMS's Health Tech Ecosystem, a mandatory equity evaluation requirement that includes:

- Pre-deployment demographic performance disaggregation: vendors must provide performance data broken down by race, ethnicity, sex, age, primary language, and insurance status. Tools that cannot provide this data should not be deployed in federally funded settings until they can.
- Local equity audit: health systems must evaluate whether AI tool performance in their own patient population reflects the performance claimed by the vendor, with particular attention to any disparities for demographic subgroups overrepresented in their patient panels.
- Ongoing disparate impact monitoring: post-deployment monitoring must include demographic disaggregation, with a defined threshold for performance disparity that triggers mandatory review and corrective action.

This is not a burdensome standard — it is a baseline that any responsible developer should already be meeting. Making it mandatory ensures that the market cannot compete on the basis of ignoring equity.

6.2 Structural Inclusion of Patient and Community Voices

HHS AI governance processes — both at the federal policy level and in the local governance frameworks that health systems establish — currently lack meaningful patient and community representation. This is a structural deficit that leads to governance frameworks designed primarily by and for the interests of developers and deploying institutions.

We recommend:

- HHS establish a permanent Patient and Caregiver AI Advisory Council with a formal advisory role to ASTP/ONC, FDA's Digital Health Center of Excellence, and CMS's Innovation Center. Council membership should reflect the diversity of the patient population, with particular representation for communities historically underserved by health care.
- ASTP/ONC incorporate community representation requirements into the AI governance standards it develops for certified health systems — analogous to community benefit requirements for nonprofit hospitals — recognizing that communities affected by AI deployment have a legitimate interest in how those tools are governed. These community representation requirements should extend beyond the health systems that deploy AI tools to the developers who create them. HHS should require that any AI developer seeking or maintaining recognition as a participating vendor or developer within CMS's Health IT Ecosystem demonstrate meaningful engagement with the

communities their tools are designed to serve — including communities defined by race, ethnicity, primary language, disability status, and geography — as a condition of that participation. This is particularly important for developers whose tools are deployed in safety-net, rural, or underserved settings, where the gap between the populations represented in model development and the populations affected by model deployment is most likely to be consequential. Community engagement in this context does not require a prescribed structure, but should at minimum include documented processes for incorporating community input into tool design, validation, and ongoing governance — with HHS establishing minimum standards for what constitutes meaningful engagement as distinct from nominal compliance.

6.3 Equitable Access to AI Governance Capacity

The equity problem in clinical AI is not only about whether AI tools perform equitably for diverse patients. It is also about whether the organizations that serve those patients have the capacity to govern AI responsibly. Safety-net providers, FQHCs, rural health clinics, and Critical Access Hospitals disproportionately serve vulnerable populations and are disproportionately under-resourced for AI governance.

HHS should address this through:

- Dedicated AHRQ and HRSA funding for AI governance capacity-building at FQHCs, rural health clinics, and safety-net providers — including training, shared evaluation resources, and access to standardized procurement tools.
- A shared services model, potentially administered through regional health information exchanges (HIEs) or ASTP/ONC's Trusted Exchange Framework and Common Agreement (TEFCA) infrastructure, through which smaller organizations can access AI governance and validation support without bearing the full cost independently.
- CMS reimbursement consideration for AI governance activities — including model validation, equity auditing, and post-deployment monitoring — as legitimate quality improvement expenditures under existing value-based care frameworks.

VII. Interoperability for Governance (Question 8)

From a governance perspective, interoperability is not only a technical enabler for AI — it is also the foundation for meaningful AI accountability. Tools that operate on fragmented, inaccessible, or non-standardized data are impossible to audit, validate, or monitor effectively.

HHS should prioritize the following interoperability investments to support AI governance:

- Standardized FHIR-based AI model cards: HL7 should develop and ONC should adopt as federally recognized interoperability standards FHIR implementation guides for AI model documentation — standardized, machine-readable disclosures of model training data, validation evidence, performance metrics, and intended use — that can be embedded in health system AI registries and queried by governance tools.

- Demographic data standardization: The ability to conduct meaningful equity audits depends on the availability of standardized, structured race, ethnicity, language, and disability data in EHR systems. HHS should strengthen requirements for collection and exchange of this data under ONC certification criteria and ensure USCDI includes SDOH and demographic variables sufficient for equity auditing.
- Federated audit infrastructure: ASTP/ONC, AHRQ and CMS should invest in the development of federated audit tools that allow regulators, researchers, and accreditation bodies to evaluate AI tool performance across health systems without requiring raw patient data to leave institutional control. This is technically feasible and would dramatically improve post-market surveillance capacity.

VIII. Supporting Private Sector Governance Activities (Question 5)

Private sector governance activities - including accreditation, certification, and industry-driven standards - have an important role to play, but HHS should be clear-eyed about their limitations. Voluntary industry standards, unmoored from regulatory requirements, tend toward the least-common-denominator. For high-risk clinical AI, self-regulation is insufficient.

We recommend a complementary model:

- HHS-anchored minimum standards: ASTP/ONC should establish minimum governance standards for Tier 1 and Tier 2 clinical AI as certification criteria under 45 C.F.R. §170.315. These represent the regulatory floor.
- Accreditation as a market signal above the floor: HHS should recognize and support independent accreditation programs — from bodies such as the Joint Commission, NCQA, or a new AI-specific accreditation body — that evaluate health systems' AI governance practices against standards above the regulatory minimum. Accreditation status should be visible in CMS quality reporting and usable by payers in network contracting decisions. Beyond its role as a market signal, HHS should give serious consideration to making demonstrated AI governance accreditation a Condition of Participation (CoP) under 42 C.F.R. Parts 482 and 485 for hospitals and health systems that deploy Tier 1 or Tier 2 clinical AI tools in their care delivery operations. Establishing accreditation or equivalent governance standards as a CoP would transform AI governance from a voluntary best practice into a baseline accountability requirement for participation in Medicare and Medicaid — giving it the same structural weight that existing CoPs carry for infection control, patient rights, and quality assurance. HHS should initiate rulemaking to explore this pathway, with particular attention to how accreditation standards could be designed to be achievable by smaller and safety-net health systems without creating compliance burdens that effectively disadvantage the organizations that serve the most vulnerable patient populations. In the interim, prior to any CoP rulemaking, accreditation status should function as a meaningful market signal

through its visibility in CMS quality reporting and its usability by payers in network contracting and value-based care arrangement decisions.

- Developer certification programs: ASTP/ONC or a designated third party should operate a voluntary AI developer certification program that evaluates clinical AI tools against published transparency and equity standards. Certification should not be required as a standalone matter but should confer meaningful and concrete market obligations and advantages. Specifically, certification — or demonstrated compliance with equivalent transparency and equity standards — should be established as a condition of participation in CMS's Health IT Ecosystem as a recognized AI vendor or developer, ensuring that the federal market is not accessible to developers who have not met minimum governance standards. Beyond ecosystem participation, certification should be tied to the use of health information technology in CMS reimbursement pathways that already carry HIT requirements — including CMS Innovation Center models that require the use of Certified Electronic Health Record Technology (CEHRT) as a condition of participation, and CMS Advanced Alternative Payment Models (Advanced APMs) under the Quality Payment Program — such that AI tools used by providers operating under these models must meet the same certification or equivalent standards as a condition of their use in those payment arrangements. This approach gives certification genuine regulatory consequence rather than treating it as a purely voluntary market signal. This certification program should explicitly extend to developers of patient and consumer-facing clinical AI tools, including those distributed through app marketplaces or offered directly to consumers, recognizing that these developers are a material part of CMS's Health IT Ecosystem and that the absence of a clinical intermediary makes governance standards more — not less — important for protecting consumers. Preferred status in federal procurement should apply to certified developers across all of these categories.

IX. Research and Development Priorities (Question 10)

We recommend HHS prioritize the following research investments through NIH, AHRQ, and ASTP/ONC:

- Governance science: Rigorous research on what AI governance frameworks and practices actually work — what evaluation methods predict real-world safety, which organizational governance structures improve outcomes, and how to design accountability mechanisms that do not inadvertently impede beneficial adoption. This field is nascent, and federal investment would have high leverage.
- Equity research: Studies specifically examining AI performance disparities across demographic subgroups, the mechanisms by which bias enters clinical AI systems, and the effectiveness of bias mitigation techniques in real-world clinical settings. NIH should create a dedicated funding mechanism for clinical AI equity research.
- Implementation science: Research on the organizational, workforce, and cultural factors that determine whether clinically effective and responsibly governed AI is actually adopted and used appropriately in practice — with particular focus on community health settings.

- Longitudinal outcomes research: Long-term studies evaluating the population health effects of widespread clinical AI adoption, including effects on health equity, total cost of care, and the patient-clinician relationship. Without this evidence base, the field cannot make credible claims about AI's societal value.

X. Conclusion

AI has genuine and significant potential to improve clinical care. Realizing that potential requires governance frameworks that are rigorous enough to protect patients — particularly the most vulnerable — while predictable enough to support responsible investment and innovation.

The United States does not currently have such a framework. The choice facing HHS is not between acceleration and safety: it is between building the governance infrastructure now, in a deliberate and coherent way, or encountering it reactively, after patient harm has eroded public trust in ways that would be far more damaging to AI adoption in the long run.

We urge HHS to invest in the governance foundation that durable AI adoption requires: a unified, tiered oversight framework; mandatory transparency and equity standards; meaningful patient voice; and post-market surveillance with teeth. These are not obstacles to innovation. They are the conditions under which innovation can earn and sustain public trust.

We appreciate the opportunity to submit these comments and welcome further dialogue with HHS as it develops its approach.