February 23, 2026

Assistant Secretary for Technology Policy (ASTP)
and Office of the National Coordinator for Health IT (ONC)
U.S. Department of Health and Human Services
Washington, DC

Ashley Cairns, MBA
Founder and Chief Executive Officer
Keywell.ai
Austin, TX

RE: Response to HHS Health Sector AI RFI (RIN 0955-AA13)

I am pleased to submit this response to HHS' RFI on accelerating the adoption and use of artificial intelligence as part of clinical care. I am the Founder and CEO of Keywell.ai, a tech-enabled services firm that provides analytics and AI solutions in the health and human services sector. Our clients include payers, clinicians and providers, state agencies, and private-sector health care organizations. I have 16 years of experience in healthcare analytics and AI, including leading data and AI modernization efforts for Texas HHSC, California OSI, and national payer and provider organizations.

<u>Response to Question 1: Barriers to Adoption</u>

The Department correctly identifies that payers bearing financial risk have an inherent incentive to promote access to high-value AI clinical interventions — but translating that incentive into practice requires a governance infrastructure that many regional and mid-market payers are structurally unable to build on their own. Unlike large national carriers with dedicated data science organizations and compliance teams, smaller payers lack the in-house capacity to evaluate AI vendor claims, implement HIPAA-compliant large language model workflows, navigate HTI-1 Decision Support Intervention disclosure requirements, or sustain the continuous performance monitoring that responsible AI deployment demands. Payment policy reforms designed to reward value-based AI adoption will accelerate the intended outcomes only if accessible enterprise governance frameworks are available to organizations that cannot amortize these implementation costs across large member populations. HHS should therefore consider incentive structures — such as enhanced risk corridor adjustments or quality bonus provisions — that explicitly account for the disproportionate implementation burden borne by smaller, resource-constrained payers, and should actively support the development of shared governance infrastructure that enables those payers to act on the incentives the Department seeks to create.

<u>Response to Question 3: Legal and Implementation Risks</u>

Legal and implementation risks in AI adoption are not hypothetical: they represent active compliance burdens that organizations deploying non-device AI in clinical care must navigate, currently without clear federal guidance. HIPAA's minimum necessary use standard and six-year audit retention requirements impose architectural constraints on LLM context windows and logging pipelines that most general-purpose AI platforms were not designed to accommodate; HTI-1's Decision Support Intervention framework creates classification uncertainty for tools that provide clinical insight without explicit treatment directives; FDA's medical device boundary remains difficult to operationalize for risk stratification and clinical documentation applications; and the 250 health AI-related bills introduced across 34 states in 2025 alone have produced overlapping and sometimes contradictory transparency, prior authorization, and human-in-the-loop mandates at the state level. The Pieces Technologies enforcement action in Texas — which required post-hoc disclosure of error rates, training data sources, and known limitations following overstated performance claims — illustrates that the absence of harmonized federal standards creates both patient safety risk and material litigation exposure. HHS can meaningfully reduce this fragmentation by issuing unified guidance that maps federal AI governance expectations to established frameworks such as the NIST AI Risk Management Framework, HITRUST, ISO 42001, and the CMS AI playbook, giving organizations a credible, auditable compliance pathway and reducing the due diligence overhead that currently deters smaller market participants from deploying clinical AI at all.

Response to Question 8: Interoperability

The most consequential barrier to scalable clinical AI is not connectivity between systems but semantic consistency within them. AI models trained on claims or electronic health record data that applies different coding conventions, condition hierarchies, and risk stratification logic across organizations will produce inconsistent outputs, a problem that compounds when models are deployed across payer-provider partnerships or multi-state populations. Standardized healthcare ontologies that provide consistent definitions for clinical concepts, cost categories, and quality measures form the necessary semantic foundation for AI tools that can be developed, validated, and deployed reliably across heterogeneous data environments. Federal support for ontology standardization — through ONC's data standards program, CMS value-based care reporting requirements, or ASTP's interoperability rulemaking — would reduce duplicative data engineering work, enable more reproducible AI validation studies, and materially lower barriers to entry healthcare organizations that stand to benefit most from clinical AI adoption but currently lack the data infrastructure to participate.

Respectfully submitted,

/s/ Ashley Cairns

Ashley Cairns, MBA
Founder and CEO
Keywell.ai