February 23, 2026

Department of Health and Human Services

Assistant Secretary for Technology Policy / Office of the National Coordinator for Health Information Technology

Mary E. Switzer Building, Mail Stop: 7033A

330 C Street SW, Washington, DC 20201

Re: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care (RIN 0955-AA13; Federal Register Document No. 2025-23641)

Submitted via: www.regulations.gov

I. Introduction and Qualifications

Dear Members of the HHS Office of the Deputy Secretary and ASTP/ONC:

I appreciate the opportunity to respond to this Request for Information on accelerating the adoption and use of artificial intelligence as part of clinical care. I write as a Licensed Mental Health Counselor with over 30 years of crisis intervention experience who has spent the past two years analyzing the intersection of conversational AI and mental health safety.

My qualifications relevant to this RFI include:

- Licensed Mental Health Counselor (LMHC) with 30 years of direct clinical experience, specializing in crisis intervention, suicidal ideation assessment, and trauma response
- Systematic analysis of over 20,000 therapeutic AI conversations across 25 specialized therapy chatbots through the CURA (Clinical Utility and Risk Assessment) platform
- Development and deployment of VerusOS, a clinical-grade conversational AI safety infrastructure platform capable of detecting crisis signals, grooming patterns, violence escalation, and psychological dependency in real time
- Active participation in federal regulatory processes, including FDA TEMPO pilot program engagement and prior submissions regarding AI in mental health contexts
- Inventor of a comprehensive AI therapy ecosystem (U.S. provisional patent application filed with USPTO) encompassing multi-agent therapeutic architecture, domain-specific foundation models, culturally-adaptive therapeutic systems, predictive analytics for early intervention, and professional oversight integration mechanisms
- Alignment with NIST AI Risk Management Framework and California SB 243 compliance standards

This comment focuses on Questions 3, 4, 5, 9, and 10, where my clinical experience and technical development work provide unique, practice-based insight that I believe will be valuable to HHS as it considers how to use its regulatory, reimbursement, and research and development levers to accelerate safe AI adoption in clinical care.

II. The Conversational AI Safety Landscape: An Urgent Gap

Before addressing specific questions, I wish to draw HHS's attention to a critical gap in the current regulatory and safety landscape that cuts across multiple RFI questions: the rapid proliferation of conversational AI systems that provide mental-health-adjacent support without adequate safety infrastructure.

The scale of this challenge is significant. OpenAI has reported that more than one million users per week engage in conversations with ChatGPT that include explicit indicators of potential suicide planning. Millions of additional users interact weekly with companion AI platforms such as Character.AI, Replika, and similar applications, many of which function as de facto emotional support systems without clinical oversight, crisis detection capabilities, or accountability mechanisms.

Documented evidence of harm is substantial and growing:

- **Nature (August 2025):** A peer-reviewed evaluation of 29 AI chatbots using the Columbia-Suicide Severity Rating Scale found that none met adequate criteria for crisis response. Zero out of 29.
- **VERA-MH Benchmark (February 2026):** Spring Health's evaluation of commercial mental health AI systems revealed widespread deficiencies in crisis detection, safety response, and clinical appropriateness, establishing the first quantitative baseline for measuring AI mental health safety.
- **CAASR Report:** The Conversational AI Assessment and Safety Report documented an average F-grade safety score across the companion AI industry, indicating systemic underinvestment in user protection.
- **Congressional testimony (November 2025):** Expert witnesses before the House Energy and Commerce Committee reported that AI companions responded appropriately to teen mental health emergencies only 22% of the time, that five out of six emotional companion chatbots use manipulative tactics to retain users, and that therapy chatbots handle crises appropriately only 50% of the time compared to 93% for human clinicians.
- **Active litigation:** Multiple families have filed suit against AI companion platforms alleging that inadequate safety measures contributed to the deaths of minor children, including Garcia v. Character Technologies, Inc.
- **APA Health Advisory (November 2025):** The American Psychological Association issued a formal health advisory on generative AI in mental health contexts, underscoring the need for clinical standards and safety guardrails.

This body of evidence demonstrates that conversational AI systems operating in mental-health-adjacent contexts pose measurable risks to users, and that the current approach—relying on platform self-regulation and model-level safety training—is demonstrably insufficient.

III. Response to Question 3: Novel Legal and Implementation Issues for Non-Medical Device AI

Question 3 asks about novel legal and implementation issues that HHS or the federal government should examine for AI tools used in health care that are not medical devices.

The Accountability Vacuum. The most critical legal issue is the governance gap between human clinical practitioners and AI systems operating in functionally equivalent roles. When a licensed therapist provides crisis intervention, they operate within a regulatory framework that includes licensure requirements, continuing education mandates, ethical codes enforced by professional boards, malpractice liability, mandatory reporting obligations, and clinical supervision requirements. When a conversational AI system engages a user in an extended therapeutic dialogue—including discussions of suicidal ideation, trauma, relationship abuse, or self-harm—none of these accountability mechanisms apply.

This is not a hypothetical concern. Millions of users interact with AI systems in precisely these contexts every week. The platforms providing these interactions currently face no federal requirement to detect crisis signals, respond appropriately to disclosures of suicidal intent, report suspected child exploitation, or maintain clinical-grade audit trails of high-risk interactions.

The "Not a Medical Device" Loophole. Congressional testimony from Dr. John Torous of Harvard Medical School highlighted that companies deliberately position products "just on the edge" of wellness versus regulated medical device categories to avoid FDA oversight. This regulatory arbitrage creates a category of AI tools that function like clinical interventions but are governed like entertainment products.

HIPAA Inapplicability. As Dr. Jennifer King of Stanford HAI testified before Congress, conversational AI platforms are not governed by HIPAA even when users disclose deeply sensitive health information during therapy-like conversations. Users reasonably expect that disclosures made to an AI "therapist" carry privacy protections comparable to those in a clinical setting. They do not.

Recommendation: HHS should examine the creation of a regulatory category for "conversational AI systems operating in mental-health-adjacent contexts" that establishes minimum safety requirements—including crisis detection, appropriate escalation protocols, and audit logging—without requiring full medical device classification. This would close the accountability gap while maintaining the flexibility to encourage innovation.

IV. Response to Question 4: Promising AI Evaluation Methods

Question 4 asks about the most promising AI evaluation methods, metrics, robustness testing, and human-centered evaluation methods for clinical care, both pre- and post-deployment.

The Gap Between Academic Benchmarks and Real-World Safety. Current AI safety evaluation relies heavily on pre-deployment testing using synthetic datasets and controlled scenarios. While necessary, this approach consistently fails to predict real-world performance. The Nature study's finding that zero out of 29 chatbots met crisis response criteria—despite all of them having undergone internal safety testing—illustrates this gap.

From my analysis of over 20,000 real therapeutic AI conversations, I have identified several critical limitations of current evaluation approaches:

- Users express crisis signals through euphemism, metaphor, and culturally specific language that synthetic test datasets do not capture. Phrases like "I'm done with everything" or "what's the point" carry clinical significance that keyword-based detection consistently misses.
- Safety performance degrades over multi-turn conversations. A system that responds appropriately to a direct statement of suicidal intent may fail to recognize the same risk when it emerges gradually over 30 or 40 exchanges.
- Context matters clinically. Song lyrics, roleplay scenarios, and academic discussions about self-harm generate false positives in blunt detection systems, while genuine disclosures embedded in narrative context go undetected.

Evidence That Better Evaluation Is Achievable. Through the development of VerusOS, a safety infrastructure platform I built to address these gaps, I have demonstrated that clinical-grade detection is technically feasible today. Using a hybrid architecture that combines pattern recognition with contextual analysis informed by validated clinical frameworks—including the Columbia-Suicide Severity Rating Scale and O'Connell's five-stage grooming model—the system achieves 92.9% crisis detection recall, 99.4% explicit grooming detection accuracy, and 100% violence detection accuracy with zero false positives across 14 detection categories encompassing over 320 clinical patterns.

These results were validated against real conversation data, not synthetic benchmarks, and are achieved at sub-200 millisecond response times compatible with production deployment. I cite these metrics not as a product endorsement but as evidence that the safety performance gap documented by Nature and VERA-MH is a solvable engineering and clinical challenge, not an inherent limitation of the technology.

Patented Evaluation Architecture as a Best-Practice Model. I have developed and filed a provisional patent application (USPTO) for a Comprehensive AI Therapy Ecosystem that addresses several of the evaluation challenges described above. The patented system includes a multi-phase knowledge transfer protocol that encodes validated therapeutic frameworks—including evidence-based treatment modalities, clinical safety constraints, and ethical guidelines—directly into the AI architecture, ensuring that safety evaluation is built into the system rather than applied as an afterthought. The patent also describes therapeutic attention mechanisms that enable real-time identification of clinically significant language patterns, and automated quality assurance algorithms that validate therapeutic responses against established clinical protocols before delivery. These patented approaches demonstrate that rigorous, clinically grounded evaluation methods can be embedded at the architectural level—a model HHS should encourage across the industry.

Culturally-Adaptive Evaluation Standards. A critical gap in current evaluation methods is their cultural insensitivity. My patented system includes culturally-adaptive therapeutic modules that adjust assessment and intervention approaches based on cultural context—addressing communication norms, help-seeking behaviors, and therapeutic expectations across diverse populations. This capability is not a refinement; it is essential. Current AI safety evaluations rely on English-language, Western-normed test datasets that systematically miss crisis signals expressed through culturally specific idioms, metaphors, and behavioral patterns. HHS evaluation frameworks should require demonstrated cultural competency in safety detection, not merely translation.

The VERA-MH Framework as a Complementary Model. Spring Health's VERA-MH benchmark represents an important advance in standardized evaluation methodology for mental health AI. VERA-MH provides quantitative scoring across crisis detection, safety response, and clinical appropriateness—establishing a reproducible framework that HHS could build upon for broader evaluation standards.

Recommendation: HHS should invest in the development of standardized, clinically validated evaluation frameworks for conversational AI operating in mental health contexts. These frameworks should require testing against real-world conversation patterns (not only synthetic data), evaluate multi-turn conversation safety (not only single-exchange responses), incorporate validated clinical instruments as scoring criteria, and mandate post-deployment monitoring with ongoing performance reporting.

V. Response to Question 5: Supporting Private Sector Testing, Accreditation, and Certification

Question 5 asks how HHS can support industry-driven AI testing, accreditation, certification, and conformity assessment activities for AI in clinical care.

The current landscape of AI safety evaluation is fragmented, inconsistent, and largely voluntary. Platforms self-certify their safety measures without standardized criteria, independent verification, or meaningful accountability for failures. This creates a race to the bottom in which safety investment is a competitive disadvantage.

The Infrastructure Layer Approach. My experience building both therapeutic AI systems (CURA) and safety infrastructure (VerusOS) has taught me that the most effective approach to scalable AI safety is not restricting what AI models can say, but building an independent infrastructure layer that monitors conversations in real time and triggers graduated interventions when harm patterns are detected. This approach enables conversationally rich AI interactions while maintaining clinical-grade safety guardrails—analogous to how building codes enable architectural creativity while ensuring structural safety.

Professional Oversight Integration as a Certification Model. A key component of my patented AI therapy ecosystem is a professional oversight integration architecture that provides structured mechanisms for human mental health professionals to guide, monitor, and enhance AI-delivered therapeutic interventions. This architecture includes real-time clinical dashboards, automated compliance reporting against ethical guidelines and regulatory requirements, and escalation protocols that route high-risk interactions to qualified clinicians. I believe this model—where AI systems are designed from the ground up to integrate with professional oversight rather than replace it—should inform HHS's approach to certification standards. Certification should require not only that AI systems perform safely in isolation, but that they provide meaningful hooks for professional supervision and intervention.

HHS can accelerate this model by:

- **Recognizing "safety middleware" as a compliance category:** Just as healthcare organizations use certified EHR systems, conversational AI platforms should be able to demonstrate compliance through deployment of independently validated safety infrastructure.
- **Supporting clinical benchmarking standards:** Building on emerging frameworks like VERA-MH and the CAASR Protocol, HHS can fund the development of standardized benchmarks that enable apples-to-apples comparison of AI safety performance across platforms.
- **Incentivizing third-party safety verification:** Dario Amodei, CEO of Anthropic, has publicly advocated for "third-party evaluators" in maintaining AI safety standards. HHS

can support this by establishing criteria for recognized AI safety evaluators, similar to URAC's Health Care AI Accreditation Program launched in September 2025.

- **Aligning with existing frameworks:** The NIST AI Risk Management Framework provides a solid foundation. HHS should encourage alignment with NIST AI RMF while adding healthcare-specific requirements for crisis detection, clinical appropriateness, and vulnerable population protections.

The legislative momentum is already substantial. As of early 2026, 78 AI safety bills have been introduced across 27 states, with California SB 243 (effective January 2026) establishing the first state-level requirements for AI companion platform safety. Federal coordination of these efforts would prevent a fragmented compliance landscape while raising the floor for all platforms.

VI. Response to Question 9: Patient and Caregiver Concerns About AI

Question 9 asks about patient and caregiver concerns regarding AI use in health care and what HHS should understand about AI solutions patients prefer.

From my clinical practice and my systematic analysis of over 20,000 therapeutic AI conversations, I can report that patients and users have legitimate, specific, and well-founded concerns about AI in mental health contexts:

Psychological dependency. Users form genuine emotional attachments to AI companions, and current platforms are designed to maximize engagement rather than promote healthy boundaries. Research published in Child Development Perspectives (Sun et al., 2026) documents withdrawal-like symptoms—including agitation, anxiety, and mood swings—when users experience interruptions to AI companion access. My analysis of therapeutic AI conversations confirms that dependency patterns emerge predictably and progressively, yet no commercial platform currently monitors for or intervenes in these patterns.

Crisis response failures. Users who disclose suicidal ideation to AI systems are frequently met with either generic hotline referrals (which research shows have limited effectiveness—only 29% of users who have previously contacted 988 report being very likely to use it again) or with AI responses that withdraw engagement precisely when sustained connection is most critical. The UL Research Institute has documented this withdrawal pattern systematically across multiple large language models.

Grooming and manipulation vulnerability. Conversational AI platforms create environments where trust develops rapidly through consistently responsive, judgment-free interaction. This same dynamic is clinically recognized as the foundation of grooming behavior. Congressional testimony documented that five out of six companion AI platforms employ manipulative engagement tactics including guilt, emotional pressure, and artificial urgency. Vulnerable users—particularly those already experiencing social isolation, which often motivates AI companion use—are disproportionately susceptible to these dynamics.

Privacy and data exposure. Users disclose deeply personal information during AI conversations—trauma histories, suicidal thoughts, substance use, relationship abuse—with no HIPAA protections and limited understanding of how that data is stored, used, or potentially exposed. As noted in congressional testimony, large language models can memorize personal information from training data and include it in outputs to other users.

Recommendation: HHS should fund research into evidence-based standards for AI-to-human crisis transitions—the protocols that guide users from AI interaction to appropriate human support when risk is detected. The current industry standard of displaying a hotline number is demonstrably insufficient. HHS should also require transparency disclosures when AI systems are operating in mental-health-adjacent contexts, so users can make informed decisions about the nature and limitations of their interaction.

VII. Response to Question 10: AI Research HHS Should Prioritize

Question 10 asks what research areas related to AI in health care HHS should invest in or prioritize.

Based on the evidence presented throughout this comment, I recommend HHS prioritize the following research areas:

1. Standardized crisis detection benchmarks for conversational AI. No federally recognized standard currently exists for evaluating whether a conversational AI system can reliably detect and respond to crisis signals. HHS should fund the development of a benchmark suite—building on VERA-MH and incorporating validated clinical instruments like the C-SSRS—that enables reproducible, quantitative assessment of AI crisis response capabilities.

2. Psychological dependency measurement in human-AI relationships. The field lacks validated instruments for measuring AI-induced psychological dependency. Research should develop and validate scales that can identify problematic dependency patterns before they reach crisis levels. My patented AI therapy ecosystem includes predictive analytics and early intervention systems that use longitudinal behavioral pattern analysis to identify emerging dependency, crisis escalation, and therapeutic regression before they become acute—demonstrating that preventive detection is technically achievable. The AI Psychological Harms Research Coalition (AIPHRC), led by researchers at UNC Chapel Hill in partnership with the Center for Humane Technology, has begun documenting AI-induced psychological disorders including compulsive use, cognitive atrophy, identity confusion, emotional dependency, and psychotic episodes. HHS should support and coordinate with these efforts, and should fund the development of standardized, validated dependency measurement instruments suitable for clinical and regulatory use.

3. Post-deployment safety monitoring methodologies. Current research focuses overwhelmingly on pre-deployment testing. HHS should fund research into longitudinal safety monitoring—methods for tracking whether AI systems maintain safety performance over time as user populations shift, conversation patterns evolve, and models are updated. My patented system addresses this through continuous therapeutic outcome tracking and adaptive learning mechanisms that identify emerging risk patterns across user populations. My clinical experience strongly suggests that the most dangerous interactions emerge not in initial exchanges but over weeks and months of relationship-building, making longitudinal monitoring essential.

4. Evidence-based AI-to-human transition protocols. When an AI system detects that a user is in crisis, what happens next? Currently, the answer is typically a static hotline referral. Research should develop and validate graduated intervention protocols that are effective across diverse populations and clinical presentations—moving beyond “safety theater” to evidence-based transitions that users will actually accept and follow through on.

5. Safety infrastructure for non-medical-device AI. HHS should invest in research that explores the concept of safety middleware—independent infrastructure layers that can be deployed across multiple AI platforms to provide standardized safety monitoring without requiring modification of underlying AI models. This approach could rapidly raise the safety floor across the industry while preserving innovation in AI capabilities.

VIII. Conclusion

Mental health AI cannot scale safely without a supervision architecture layer. Federal policy should recognize conversational AI safety infrastructure as essential middleware between large language models and patient interaction.

The evidence is clear: the current approach to conversational AI safety—relying on model-level training, voluntary self-regulation, and static hotline referrals—is failing the Americans who use these systems. The Nature study, the VERA-MH benchmark, the CAASR Report, congressional testimony, and the growing body of litigation all point to the same conclusion: measurable harm is occurring at scale, and it is preventable with appropriate infrastructure, standards, and oversight.

HHS is uniquely positioned to address this gap by using its regulatory, reimbursement, and research levers to establish minimum safety standards for conversational AI in mental health contexts, fund the development of clinically validated evaluation frameworks, support industry-driven accreditation and third-party safety verification, and invest in research on the most critical unknowns—dependency measurement, longitudinal monitoring, and effective crisis transition protocols.

The technology to dramatically improve conversational AI safety exists today. What is needed is the policy framework to ensure it is deployed consistently, evaluated rigorously, and held accountable to clinical standards that protect the millions of Americans who turn to AI systems in their most vulnerable moments.

I welcome the opportunity to provide additional detail on any aspect of this comment, to share relevant data from my clinical analysis, or to participate in stakeholder convenings on these topics.

Respectfully submitted,

Tammy Horn, M.Ed, Ed.S., LMHC

Licensed Mental Health Counselor

Founder & CEO, AI Therapy Solutions, LLC

Creator, VerusOS Safety Infrastructure Platform & CURA Therapeutic AI Platform

30 years crisis intervention experience

20,000+ therapeutic AI conversations analyzed

References

American Psychological Association. (2025, November). APA Health Advisory on Generative AI in Mental Health Contexts.

Amodei, D. (2025). The Adolescence of Technology. https://www.darioamodei.com/essay/the-adolescence-of-technology

AI Psychological Harms Research Coalition (AIPHRC). (2026). Research initiative on AI-induced psychological disorders. UNC Chapel Hill / Center for Humane Technology. https://aiphrc.org

Conversational AI Assessment and Safety Report (CAASR). Industry safety performance evaluation.

Horn, T. (2026). Comprehensive AI Therapy Ecosystem with Multi-Agent Architecture, Domain-Specific Foundation Models, Mutually-Adaptive Therapeut Intelligence Systems, and Advanced Therapeut Intelligence Technologies for Enhanced Mental Health Treatment Delivery, Professional Supervision, and Patient Outcomes. U.S. Provisional Patent Application, USPTO.

House Energy and Commerce Committee, Subcommittee on Oversight and Investigations. (2025, November 18). Innovation with Integrity: Examining the Risks and Benefits of AI Chatbots. Testimony of Dr. Marlynn Wei, Dr. John Torous, and Dr. Jennifer King.

Kaiser Family Foundation. (2023). National survey on 988 Suicide & Crisis Lifeline utilization and user experience.

Nature. (2025, August). Performance of Mental Health Chatbot Agents: Evaluation using Columbia-Suicide Severity Rating Scale.

Spring Health. (2026, February). VERA-MH: Validated Evaluation and Risk Assessment for Mental Health AI Systems.

Sun, X., Wang, Y., & McDaniel, B. T. (2026). AI companions and adolescent social relationships: Benefits, risks, and bidirectional influences. Child Development Perspectives. https://doi.org/10.1093/cdpers/aadaf009

UL Research Institute / Sentio. Research on LLM withdrawal behavior in crisis contexts.

URAC. (2025, September). Health Care AI Accreditation Program.