

February 23, 2025

Innovaccer, Inc.
201 Mission Street, Ste. 2900
San Francisco, CA 94105

Department of Health and Human Services
Office of the Secretary
Assistant Secretary for Technology Policy and Office of the National Coordinator for Health
Information Technology (ASTP/ONC)
200 Independence Avenue SW,
Washington, DC 20201

Submitted electronically via Regulations.gov

Re: RIN 0955-AA13, HHS Health Sector AI Request for Information

Dear Secretary Kennedy and Assistant Secretary Keane,

Innovaccer, Inc. (Innovaccer) is honored to contribute our perspective to the Request for Information (RFI) issued by your offices to seek public comment on what HHS can do to accelerate the adoption and use of AI as part of clinical care.

Innovaccer submits these comments as a healthcare AI platform company with direct operating experience across the deployment challenges this RFI addresses. Our [Gravity platform](#), a unified enterprise data and AI orchestration system, connects data from 100+ electronic health records and automating workflows spanning revenue cycle, prior authorization, population health, care management, and clinical documentation. Since our founding in 2014, we have invested more than \$500 million in healthcare-specific AI infrastructure, scaled a team of more than 750 AI/ML and data engineers working exclusively in the healthcare context. Our track record was recently validated by KLAS ratings above 90% across three categories (provider data platform, payer data platform, and CRM). We offer these comments not as advocates for AI in the abstract, but as practitioners accountable for AI performance inside the health systems, providers, payers, and public sector agencies that serve Americans.

Our responses are organized around a thesis we believe should shape HHS's approach throughout this RFI: the U.S. health sector faces compounding crises that will determine its financial sustainability and mission capacity for the coming generation. On the clinical side, an aging population, a rapid rise in chronic disease prevalence, and increasingly complex and expensive pharmacologic therapies and medical devices are driving care demands that a fragmented system, already constrained by primary care and nursing shortages, was not designed to absorb. On the financial side, labor represents 50 to 70 percent of total operating cost for most health systems and is growing faster than revenue. Clinical complexity is directly amplifying these costs. AI automation applied systematically to the workflows where these twin pressures are most acute, including revenue cycle, prior authorization, clinical documentation, scheduling, and utilization management, is not incremental efficiency. It is the mechanism by which health systems sustain both financial viability and clinical quality simultaneously. This is what Innovaccer's Gravity platform is built to deliver, and it is why we describe the goal as better margins for stronger mission.

The administrative and operational AI driving this transformation is not separate from the clinical AI this RFI addresses. Investment in one accelerates the other, because both depend on the same underlying data foundation, governance infrastructure, and interoperability standards. HHS has a distinct opportunity through this RFI to recognize enterprise AI infrastructure as a shared prerequisite for both, and to design its regulatory, reimbursement, and R&D policies to accelerate healthcare's automation transformation rather than treating its operational and clinical dimensions as separate problems requiring separate policy frameworks, respecting that the business of healthcare needs to align with the clinical effectiveness of healthcare.

A final framing consideration that should inform HHS's approach throughout this RFI: clinical AI must be held to a clear standard of clinical benefit balanced against cost. As general-purpose foundation model platforms seek to expand into healthcare through health records aggregation and participation in government programs, HHS should ask a question that is straightforward but consequential: what should clinical AI not do? It should not optimize for engagement over outcomes. It should not serve as a channel for pharmaceutical or device advertising. It should not aggregate patient data in ways that serve platform growth rather than patient care, hard lessons learned in other domains and in previous generations of healthcare IT development. These are not hypothetical risks. When general-purpose consumer AI platforms that derive revenue from advertising gain access to clinical care environments, the structural conflict between commercial incentive and clinical integrity does not disappear because the interface looks clinical. Purpose-built healthcare AI, operated by companies whose only business is healthcare improvement and who are accountable to health system operators rather than to advertisers or platform engagement metrics, is the architecture most likely to maintain that standard. HHS's governance frameworks should reflect the difference between general-purpose AI adapted for healthcare and healthcare AI built from the ground up and should create incentive structures accordingly.

Responses to selected questions from the RFI follow:

Question 1: Barriers to Private Sector Innovation

The pressure driving healthcare AI adoption has two interconnected roots. The first is clinical: an aging population, a rapid rise in chronic disease prevalence, and increasingly complex and expensive pharmacologic therapies and medical devices are compounding the demand on health systems writ large at the same time that primary care and nursing shortages are limiting capacity to meet it. Care is simultaneously becoming more complex and more fragmented, with patients navigating more conditions, more specialists, and more care transitions than the system was designed to support.

The second is financial: labor represents 50 to 70 percent of total operating cost for most health systems, and those costs are growing faster than revenue. These two pressures are not independent; the clinical complexity driving higher acuity and longer care episodes is a direct contributor to the labor cost crisis, and the workforce shortages constraining clinical capacity make the financial pressure harder to absorb. AI automation applied to the workflows where this combined pressure is most acute is not incremental efficiency; it is the mechanism by which health systems sustain both financial viability and clinical quality simultaneously.

We have identified the following barriers to private sector innovation:

First, fragmented data and governance infrastructure. Health systems attempting to adopt AI encounter data scattered across siloed EHR, billing, scheduling, clinical knowledge repositories

and HR systems; inconsistent data quality and standards; and no unified governance layer to manage AI deployment safely at scale. Organizations testing AI tools in isolation, such as a point-solution tool for revenue cycle, another for facilitating prior authorization, another for coding and appropriate documentation, face governance overhead that multiplies with each addition, creating vendor sprawl that adds risk rather than reducing it. A unified data and AI governance platform that standardizes data once and applies governance consistently across use cases dramatically reduces this overhead and accelerates time to value and orchestrates the data inputs and outputs of AI solutions facilitating embedding into clinical workflows.

Second, data access restrictions imposed by EHR vendors. Innovaccer addresses these in detail in our forthcoming HTI-5 comments: restrictions on automated API-based data access, inconsistent FHIR implementations, and prohibitive fees collectively prevent health systems from building the unified data foundation that operational AI requires. Finalizing HTI-5's information blocking provisions will directly address this barrier.

Third, workflow integration gaps. AI tools that require staff to switch to a separate system achieve minimal adoption regardless of their technical quality. But the solution is not simply routing outputs back into the EHR. That framing preserves the EHR as the permanent center of the workflow universe and limits what AI can accomplish. The more powerful architectural model, and the one Innovaccer's Gravity platform is built around, treats the EHR as one system of record within a broader AI orchestration layer that coordinates workflows across clinical, financial, and operational systems simultaneously. HHS should recognize this class of enterprise AI orchestration platforms as a distinct category of health IT; one that transcends any single EHR and requires interoperability standards that support bidirectional workflow coordination across systems, not just data retrieval from them.

Question 2: Regulatory/Payment/Program Changes to Prioritize

Innovaccer recommends that HHS prioritize the following regulatory, payment, and programmatic changes, with a focus on unlocking operational AI automation.

Regulatory changes to prioritize:

First, finalize and enforce HTI-5's information blocking provisions and updates, as detailed in our forthcoming HTI-5 comments. These changes are prerequisite to the unified data access that operational AI requires.

Second, coordinate with FDA to issue joint HHS-FDA guidance on the boundary between clinical decision support exclusions under 21 U.S.C. §360j(o) and SaMD classification.

Third, consider a regulatory safe harbor for healthcare AI tools whose developers can demonstrate rigorous pre-deployment validation against representative patient populations, documented model development and testing methodology, prospective post-deployment performance monitoring, and transparent disclosure of model limitations, intended use, and known failure modes. Anchoring safe harbor criteria on the rigor of the development and validation process, rather than on deployment platform type or scale, rewards responsible innovation at any developer size, avoids creating structural advantages for incumbent technology vendors, and creates durable incentives for the field to invest in better AI development practices.

Fourth, establish HHS guidance under 45 C.F.R. §164.514 and §164.506 clarifying HIPAA's Privacy Rule applicability to AI training data use, including Business Associate Agreement requirements for locally hosted AI model architectures.

Question 3: Novel Legal and Implementation Issues (Non-Medical Devices)

Non-medical device AI deployed in healthcare operations raises governance, liability, and implementation challenges that existing federal frameworks do not fully address. Innovaccer's Gravity platform was designed specifically to address several of these challenges at the enterprise level, and our experience provides direct insight into where we believe federal action would create the most value.

Governance fragmentation is a major challenge. Health systems adopting multiple AI tools and point solutions face an accumulating stack of governance obligations with no unified framework for managing them: each tool requires its own security review, training data governance assessment, and ongoing monitoring process. Gravity addresses this by providing a single governed AI layer across use cases. HHS can reduce this fragmentation burden not by building a new federal compliance apparatus, but by articulating the governance principles including data standards, audit requirements, and oversight mechanisms, against which enterprise platforms can be formally recognized. As we describe in our response to Question 5, that recognition is the mechanism: HHS defines the standard; the private sector builds to it and earns the safe harbor.

Question 5: Supporting Private Sector Activities

The most effective near-term action HHS can take is to articulate governance principles that define what responsible AI deployment in healthcare looks like. What HHS should explicitly avoid is translating those principles into a mandatory federal certification program for non-device healthcare AI. Mandatory certification creates compliance overhead that disadvantages smaller developers, produces technical standards that become outdated faster than they can be revised, and adds a bureaucratic layer to a sector already burdened by HIPAA, ONC certification, and existing quality programs.

The appropriate federal role, in short, is to define the standard and provide the incentive. The private sector, through AI orchestration platforms, industry consortia, and professional society frameworks, builds the governance content that meets those principles. This division of labor produces faster and more durable accountability than a federal program could, and it mirrors how effective voluntary frameworks in adjacent domains have achieved broad market adoption without regulatory mandate.

Question 6: Where AI Has/Hasn't Met Expectations

Based on Innovaccer's experience deploying Gravity across the healthcare enterprise, we offer the following evidence-based assessment of AI performance.

In early results, Innovaccer's AI-powered financial performance platform, [Flow](#), has generated measurable cost savings through automation of manual prior authorization processes. These use cases succeed because ROI is concrete: fewer denials, lower labor cost per claim, faster cash collection. AI that automates high-volume, rules-based administrative work and routes only exceptions to human reviewers consistently performs well. Unified enterprise data intelligence, including AI-assisted analytics allowing finance, operations, and clinical teams to query across

consolidated enterprise data has also exceeded expectations by replacing slow, labor-intensive manual reporting cycles. The common success factor across all high-performing deployments is integration: AI that surfaces results and triggers actions within existing workflows, without requiring staff to switch systems, achieves dramatically higher adoption than tools that require workflow changes.

AI has fallen short of expectations in fragmented, point-tool deployments where governance is absent or inconsistent. Organizations adopting multiple AI tools without a unified governance and data layer report that governance overhead grows faster than the value those tools deliver: each tool requires its own security review, its own data pipeline, and its own monitoring process. The cumulative effect is vendor sprawl that creates risk. In widely reported cases in our industry, real-time sepsis prediction systems have generated significant alert fatigue due to poor local calibration. The pattern of failure is clear: ungoverned, non-integrated AI fails regardless of model quality.

High-potential novel applications include: end-to-end revenue cycle automation where AI manages the full workflow from charge capture through denial resolution; agentic AI that autonomously handles routine operational tasks while escalating complex exceptions to humans; conversational enterprise search enabling any team member to query across clinical, financial, and operational data without analyst support; and AI-powered prior authorization that replaces a labor-intensive, error-prone process with a governed, auditable automated workflow.

Question 7: Decision Makers and Administrative Hurdles

AI adoption decisions in healthcare increasingly involve two parallel tracks that must both converge before a purchase is made. The first is the business track: the CEO, CFO, and COO must see a credible path to measurable margin improvement, which is typically quantified as reductions in labor cost, administrative overhead, or revenue leakage. In the current environment, where labor represents 50 to 70 percent of health system operating costs and margins are under sustained pressure, the business case for operational AI is both urgent and increasingly legible to executive leadership. What delays the business track is difficulty quantifying ROI with confidence; and what accelerates it is evidence-based projections tied to specific workflow automation outcomes.

The second is the governance and technical track: the CIO, CISO, and an increasingly formalized AI governance committee must evaluate security architecture, data governance, EHR integration, and AI accountability practices before approving any new platform. This track has lengthened significantly over the past two years as health systems have become more structured in how they approve AI. Health systems with mature AI governance frameworks move through this track significantly faster than those creating evaluation criteria from scratch. Gravity's design directly addresses this track: centralized controls, built-in auditability, and locally hosted model options reduce the security review burden substantially, which is why governance-ready platform design has become a meaningful competitive differentiator.

Question 8: Interoperability Opportunities

Enhanced interoperability is the data foundation that effective healthcare AI requires. Innovaccer's Gravity platform operates as an enterprise data fabric, bringing together clinical, financial, operational, and workforce data into a unified layer that supports AI automation across the organization. Our experience shows clearly where current interoperability gaps create the greatest friction.

First, enterprise data integration beyond clinical FHIR. AI in healthcare is most powerful when it can reason across the full enterprise data picture: clinical data from the EHR, financial data from billing and revenue cycle systems, operational data from scheduling and workforce systems, and claims data from payers. Current interoperability standards focus almost exclusively on clinical data exchange (FHIR, USCDI), leaving operational and financial data without the same focus on standardized exchange mechanisms. This benefits incumbents and disadvantages newer, more innovative players.

Second, FHIR API consistency and write-back for workflow-embedded AI. FHIR APIs enable AI to access clinical data, but inconsistent implementation across EHR platforms, and the near-universal absence of standardized write-back support, prevents AI from completing operational workflows that require returning results to the EHR or triggering downstream actions. HHS should establish rigorous FHIR conformance testing requirements and mandate write-back API support, directly enabled by HTI-5's removal of the third-party modification condition.

Thank you again for this opportunity to respond to the RFI. Our team is available to discuss any of our responses in more detail or provide demos of our Gravity platform or any of our AI-first solutions that comprise our healthcare automation stack.

Sincerely,

Abhinav Shashank
CEO and Co-Founder
Innovaccer

David Nace, MD
Chief Medical Officer

Anil Jain, MD, FACP, FACHDM
Chief Innovation Officer

Lisa Bari
Vice President, Policy and Partnerships