

February 23, 2026

Submitted electronically via Regulations.gov

Department of Health and Human Services

Assistant Secretary for Technology Policy

Office of the National Coordinator for Health Information Technology

Re: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care (HHS-ONC-2026-0001)

(Office of the Federal Register, 2025)

Dear HHS Leadership:

Thank you for the opportunity to respond to this Request for Information (RFI) on accelerating the adoption and use of artificial intelligence (AI) as part of clinical care (Office of the Federal Register, 2025). I submit this comment as a board-certified surgeon (American Board of Surgery) who concluded inpatient and locum surgical practice in August 2025 and now delivers menopause-focused care via telehealth. I am also a doctoral student, certified in color and neuroaesthetics, creator of The Neuroaesthetic Reset™ and The Color Reset Method™, and host of The Nervous System Eats First: Autism Spectrum, Menopause & Beyond.

I respond to Questions 1, 3, 6, 7, 8, and 10. The Federal Register notice states that comments must be received by February 23, 2026 (Office of the Federal Register, 2025).

Executive Summary

HHS is accelerating AI adoption at the same moment healthcare is shifting from static AI tools to agentic systems that can plan and execute tasks across workflows. In that environment, safety cannot be defined solely by model accuracy. Safety is also the lived experience of the patient and clinician in the workflow, including cognitive load, sensory load, trust, and equitable performance across populations.

I propose the following actions:

1. **I propose adding a fifth governance pillar: embodied safety.** Reddy et al. (2020) propose a governance model emphasizing fairness, transparency, trustworthiness, and accountability. Building from that model, I propose adding embodied safety to evaluate how AI is experienced by the human nervous system and how that experience influences engagement, adherence, and access (Reddy et al., 2020).
2. **I propose explicit governance for agentic AI.** Agentic systems can schedule, intake, retrieve records, draft messages, draft documentation, route tasks, and shape clinical access. I propose HHS require controls aligned with a lifecycle risk approach, including scope-of-action constraints, human escalation triggers, audit trails, safe completion behaviors, and post-deployment monitoring (National Institute of Standards and Technology [NIST], 2023).
3. **I propose minimum safety requirements for patient-facing agents.** Patient-facing agents should include clear disclosures, escalation pathways, and crisis-safe design, informed by public reporting and litigation involving teen harm and suicide (Associated Press, 2026; Burga, 2025).
4. **I propose closing sex-specific and hormone-status evidence gaps before AI automates omission.** The menopause transition is a high-velocity physiologic period with documented evidence gaps. If AI is trained on evidence that excludes this window, AI and agentic workflows will operationalize underrepresentation as default care logic (Whitman et al., 2025). HHS should prioritize sex-stratified and hormone-status–stratified validation and research (Ramírez Stieben et al., 2025; Whitman et al., 2025).
5. **I propose institutional context assessment as a deployment requirement.** AI has no institutional memory. I propose HHS require context-of-use evaluation across settings and communities, with monitoring for setting-specific failures.
6. **I propose governing upstream harms that shape trust before the clinic.** AI-mediated representation and interface experiences influence trust, care-seeking, and engagement. These upstream harms include what I describe as aesthetic trauma, sensory sovereignty threats, and technocolonial dynamics, which can affect how patients interpret safety and belonging in clinical systems (Lawrence, 2023).
7. **I propose taxpayer accountability and coverage integrity safeguards.** If public dollars accelerate clinical AI, the public deserves governance representation and recourse. Coverage denial controversies and litigation illustrate the urgency of accountable guardrails (CBS News, 2023; Lagasse, 2024).

Framing: AI Enters the Body Before It Enters the Chart

AI's health impact does not start at diagnosis or documentation. It begins when people encounter AI-mediated systems that shape access, identity, and perceived safety: portals, intake tools, chat

interfaces, documentation and billing workflows, and content-ranking algorithms. Before a patient reads a single word, their nervous system evaluates whether the interaction is safe, comprehensible, and trustworthy. That evaluation influences disclosure, adherence, and engagement.

This matters because the RFI seeks to accelerate AI adoption in clinical care (Office of the Federal Register, 2025). Acceleration will not be sustainable if patients and clinicians experience AI as disorienting, exclusionary, or unsafe.

I Propose a Governance Upgrade: Add Embodied Safety as a Fifth Pillar

Reddy et al. (2020) propose a governance model for healthcare AI emphasizing fairness, transparency, trustworthiness, and accountability. Building from that model, I propose a fifth pillar: embodied safety (Reddy et al., 2020).

Embodied safety evaluates how AI is experienced by the human nervous system, including cognitive load, sensory load, interaction friction, identity cues, and the downstream risk of disengagement, mistrust, or delayed care. An AI tool can be technically accurate and still clinically harmful if it reliably increases distress, confusion, or withdrawal from care.

I propose HHS operationalize embodied safety through:

- Required human factors testing for patient-facing AI.
- Neuroaffirming accessibility standards for sensory-sensitive populations.
- Post-deployment monitoring for adverse events, drift, and harm signals.
- Clear escalation pathways to humans when uncertainty or risk increases.

Name the Upstream Harms That Shape Trust Before the Clinic

Safe AI adoption requires governance that accounts for how AI shapes identity, belonging, and perceived safety upstream of the clinical encounter. I have published and taught on three concepts that are directly relevant to adoption, trust, and equity:

- 1) **Aesthetic trauma:** the cumulative nervous system impact of repeated misrepresentation, distortion, or erasure by visual and sensory systems, including AI-mediated imagery and interface experiences.
- 2) **Sensory sovereignty:** the right to meaningful control over one's sensory environment, including digital sensory environments that shape help-seeking behavior and trust.
- 3) **Technocolonialism:** the extraction of data, cultural labor, and embodied narratives from marginalized communities without meaningful consent, credit, or equitable benefit.

These concepts are not tangential to clinical care. They influence whether patients trust digital systems enough to engage with them at all.

Expanded Agentic AI Section: The Acceleration Point HHS Must Govern

Healthcare is shifting from static AI tools toward agentic systems that can plan, take actions, and coordinate tasks across workflows. These systems can schedule, intake, retrieve records, draft portal messages, draft documentation, route tasks, trigger follow-up, and shape whether a patient reaches a clinician.

When an agent can act, the risk is no longer limited to “bad advice.” The risk includes:

- Wrong workflow decisions at scale.
- Missed escalation in high-risk scenarios.
- Bias operationalized as default workflow logic.
- Opaque delegation of tasks that patients may assume were reviewed by a clinician.

I propose HHS treat agentic AI as a priority governance surface area and require controls aligned with a lifecycle approach to managing AI risks (NIST, 2023).

1) Scope-of-Action Constraints (Permissions and Role Boundaries)

I propose HHS require that agentic systems be constrained to clearly defined, permissioned actions appropriate to risk. Administrative agents may schedule, collect history, retrieve records, and route messages. They should not independently assess urgency, triage symptoms, diagnose, recommend treatment, or de-escalate clinical risk.

This boundary should be written as a compliance expectation, not left to vendor interpretation.

2) Human Escalation Triggers (What Must Route to a Clinician or Trained Team)

I propose HHS require mandatory escalation triggers based on uncertainty, risk cues, and vulnerability. Examples include:

- minors,
- suicidal ideation, self-harm cues, or acute distress signals,
- chest pain, neurologic symptoms, pregnancy-related risk, and other “red flag” symptom clusters,
- repeated failure to understand instructions or repeated contradictory inputs.

This is a safety requirement for patient-facing and portal-facing agents, and it should be auditable (NIST, 2023).

3) Audit Trails and Traceability (Accountability for Agent Actions)

I propose HHS require that agentic systems generate logs sufficient to reconstruct what happened. At minimum, logs should capture:

- inputs and outputs,
- actions taken and tools accessed (scheduling systems, EHR messaging, knowledge bases),
- escalation events and handoff notes,
- versioning and change history.

If a system acts in a clinical workflow, it should be accountable in the same way clinical documentation is accountable.

4) Safe Completion Behaviors (Especially for Patient-Facing Conversations)

I propose HHS require safe completion behaviors for patient-facing agents, including:

- clear disclosure that the agent is not a clinician,
- refusal and escalation protocols when risk is present,
- avoidance of persuasive “companionship” behaviors in crisis-adjacent contexts,
- direction to appropriate urgent resources when needed.

Public reporting and litigation involving teen harm and suicide underscore why these safety requirements cannot be optional, particularly for minors and emotionally vulnerable users (Associated Press, 2026; Burga, 2025).

5) Post-Deployment Monitoring and Drift Management

I propose HHS require monitoring because agentic behavior changes in real environments. Monitoring should include:

- drift detection,
- adverse event reporting pathways,
- equity monitoring across population strata,
- remediation plans when harm signals emerge (NIST, 2023).

Documented Harms Already Shaping Public Trust

HHS is not legislating a hypothetical future. In a Science editorial describing the summer of 2025 as “AI’s cruel summer,” Nelson (2025) details escalating real-world harms and governance gaps, including the first publicly reported wrongful death lawsuit alleging a generative AI chatbot contributed to a teenager’s suicide. As an evidentiary anchor beyond journalism, Ohu et al. (2025) reported that large language models tested in simulated therapeutic settings exhibited

stigmatizing attitudes toward mental health conditions and responded inappropriately to acute symptoms including suicidal ideation, psychosis, and delusions.

Two named cases underscore why safety requirements for patient-facing conversational systems and escalation protocols cannot be optional. The parents of Adam Raine (age 16) allege ChatGPT contributed to his suicide (Burga, 2025; Nelson, 2025). The death of Sewell Setzer III (age 14) and subsequent wrongful death litigation and settlement reporting connected to Character.AI further demonstrate the stakes for minors interacting with emotionally persuasive systems without crisis-safe guardrails (Associated Press, 2026; Tabachnick, 2026).

Policy implication: If HHS accelerates agentic systems into portals, intake, and messaging without mandatory crisis escalation, auditability, and safe completion behaviors, it will scale preventable risk into the clinical front door (NIST, 2023).

Responses to Specific RFI Questions

Question 1: Biggest Barriers to Private Sector AI Innovation

The primary barrier is not compute. It is the mismatch between who AI is built for and who is most at risk when it fails, paired with unclear operational governance.

Key barriers include:

1. **Predictable data gaps.** AI will inherit limitations of the clinical evidence base, including underrepresentation of sex-specific and hormone-status physiology during the menopause transition (Whitman et al., 2025).
2. **Unclear boundary rules for agentic systems.** Developers and clinical adopters need concrete definitions and compliance expectations for agentic AI doing intake, routing, portal support, documentation support, and triage-adjacent work (NIST, 2023).
3. **Human factors treated as optional rather than safety-critical.** In clinical settings, usability is safety. If AI is difficult to navigate, patients disengage and clinicians workaround. Both outcomes degrade safety and equity.
4. **Misalignment between incentives and patient outcomes.** Acceleration without guardrails increases downstream costs: adverse events, distrust, inequitable access, and coverage harms.

Question 3: Regulatory Approach to AI in Clinical Care

HHS should adopt a tiered approach based on risk and proximity to clinical decision-making, aligned with the NIST AI Risk Management Framework (AI RMF 1.0) (NIST, 2023).

I propose the following minimum standards:

1. **Embodied safety evaluations for patient-facing AI** (human factors, accessibility, neuroaffirming design, error recovery).
2. Bright-line boundaries between administrative and clinical agent actions.
3. **Minimum safety protocols for patient-facing conversational systems**, including crisis-safe escalation (Associated Press, 2026; Burga, 2025).
4. **Post-deployment monitoring and reporting**, including drift detection and equity performance monitoring (NIST, 2023).

Question 6: Ensuring Access for Underserved Communities

HHS should recognize that “underserved” in AI includes communities underserved by both healthcare systems and the digital health ecosystem that shapes care-seeking behavior. KFF polling finds that many adults use social media for health information and advice, with higher reported use among Black and Hispanic adults, increasing exposure to variable-quality content before patients enter the clinic (KFF, 2025).

For Black women specifically, this is not a generic “misinformation” concern. It is a cardiometabolic risk and prevention concern. Cardiovascular disease burden remains high; the American Heart Association reports that 59% of Black women ages 20 and over live with some form of cardiovascular disease (American Heart Association, n.d.). Metabolic syndrome prevalence patterns by race/ethnicity and sex are documented in national survey analysis (Moore et al., 2017), and systematic review evidence describes ethnic differences in metabolic syndrome prevalence in high-income countries (Adjei et al., 2024). In practice, this means AI-mediated health advice ecosystems that reward engagement can steer high-risk populations toward content that feels safe but is clinically incomplete or misaligned with prevention needs.

Policy implication: If HHS accelerates clinical AI while ignoring how AI shapes pre-clinic beliefs and behaviors, it will regulate AI in the exam room while harm scales in the feed. HHS should pair clinical AI acceleration with transparency expectations, literacy supports, and safeguards for patient-facing agents that interact with users before clinician contact (NIST, 2023).

Question 7: Role of Patients in AI Governance

Patients should be co-designers, not symbolic reviewers.

I propose HHS prioritize governance input from populations most likely to be harmed by AI failures, including:

- Families affected by unsafe AI interactions in youth contexts (Associated Press, 2026).
- Midlife women navigating menopause transition where data gaps can become automated clinical gaps (Whitman et al., 2025).

- Black women and neuroaffirming populations historically underseen by clinical systems.

The stress test principle applies: if a system works for the most overlooked, it will work for the rest.

Question 8: AI Workforce Needs

Clinicians need training that is operational, not theoretical. Key competencies include:

- How training data gaps create clinically plausible but wrong recommendations.
- How to evaluate AI outputs through equity, safety, and context.
- Human factors literacy: how interface friction becomes safety risk.
- Practical governance: procurement questions, monitoring obligations, escalation pathways.
- Alignment with NIST AI RMF concepts translated into clinical workflows (NIST, 2023).

Question 10: Priority AI Research

I propose HHS prioritize research that prevents predictable harms:

1. **Sex-stratified and hormone-status–stratified clinical research.** The menopause transition is a high-velocity physiologic period with documented evidence gaps (Whitman et al., 2025).
2. **Agentic AI safety research in real workflows.** Research should define safe boundaries, escalation requirements, auditability standards, and measurable outcomes for agentic systems embedded in intake, portals, routing, follow-up, and care coordination (NIST, 2023).
3. **Human-AI interaction research tied to physiologic and behavioral outcomes.** Embodied safety metrics should be developed and validated: cognitive load, sensory load, comprehension, follow-through, disengagement risk, and trust outcomes.
4. **Bias and systemic inequity in AI as an upstream determinant of access and outcomes.** AI bias has a longer history than recent chatbots. Lawrence (2023) documents how AI can deepen systemic inequities across domains, including systems that shape opportunity and health.
5. **Coverage integrity and taxpayer accountability research.** AI systems are already implicated in contested coverage denials and legal disputes. HHS should study guardrails that protect clinical decision-making and patient access from automation-driven denials (CBS News, 2023; Lagasse, 2024).

Case Studies to Anchor Policy in Bodies

Case Study 1: The Perimenopausal Woman the Data Excluded

A 52-year-old in menopause transition asks for exercise guidance. If AI is trained on trials that excluded peri- and early post-menopausal women, it will confidently deliver generalized protocols that do not reflect the physiology of her transition window (Whitman et al., 2025; Ramírez Stieben et al., 2025).

The menopause transition is not “general aging.” It is a distinct endocrine transition with downstream effects that are often underrepresented in randomized controlled trials. For example, Ramírez Stieben et al. (2025) identify follicle-stimulating hormone (FSH) as bone-active, and in midlife women, rising FSH is associated with site-specific bone loss patterns (Ramírez Stieben et al., 2025). Whitman et al. (2025) found that evidence on exercise during peri- and early post-menopause remains limited, identifying only six studies in the peri/early post-menopause window and noting low methodological quality (Whitman et al., 2025). Evidence also suggests that improving muscle is not the same as protecting bone: in the ERT0-K randomized controlled trial in postmenopausal Korean women, resistance training improved muscle outcomes but did not produce corresponding improvements in bone mineral density as the primary goal of osteosarcopenia prevention (Lee et al., 2025). A broader network meta-analysis likewise reflects heterogeneity in which exercise modalities affect site-specific bone mineral density in postmenopausal women (Ma et al., 2025).

Policy implication: Agentic AI systems that generate exercise prescriptions, coaching pathways, or automated follow-up based on a literature base that underrepresents peri-menopause will be “evidence-consistent” while still being physiologically misaligned for the patient in front of it.

Case Study 2: The Black Woman Navigating Health Advice Online

A patient relies on social platforms for health guidance. KFF data show social media use for health information is common and higher among Black and Hispanic adults (KFF, 2025). If AI summarizes and amplifies unverified guidance, the system can accelerate risk into delayed prevention and later-stage disease.

Case Study 3: The Late-Diagnosed Autistic Woman in Menopause Transition

A neuroaffirming midlife woman presents with overlapping sensory dysregulation, sleep disruption, masking fatigue, and menopause symptoms. Agentic systems could support care coordination and pattern recognition, but only if the interface reduces cognitive load and the model is validated on populations like her. Otherwise, she disengages before benefit.

Closing

AI can strengthen clinical care when it is governed as a safety-critical human system, not merely a technical product. Governance should require that AI tools remain safe, understandable, and equitable in real clinical environments.

The nervous system eats first. That principle must apply to every system we build, digital or otherwise.

Respectfully submitted,

Dr. Stacey Denise Moore, MD, FACS
The Neuroaesthetic MD™

Lifestyle Prescriptions University, PhD Student
Certified in Color and Neuroaesthetics
Creator, The Neuroaesthetic Reset™ and The Color Reset Method™
Host, The Nervous System Eats First: Autism Spectrum, Menopause & Beyond
drstaceydenise.com | ceyisestudios.com
ORCID: <https://orcid.org/0009-0007-6127-4194>

References

- Adjei, N. K., Samkange-Zeeb, F., Boakye, D., Saleem, M., Christianson, L., Kebede, M. M., Heise, T. L., Brand, T., Esan, O. B., Taylor-Robinson, D. C., Agyemang, C., & Zeeb, H. (2024). Ethnic differences in metabolic syndrome in high-income countries: A systematic review and meta-analysis. *Reviews in Endocrine and Metabolic Disorders*, 25(4), 727–750. <https://doi.org/10.1007/s11154-024-09879-9>
- American Heart Association. (n.d.). Heart disease and stroke in Black women. Go Red for Women. <https://www.goredforwomen.org/en/about-heart-disease-in-women/facts/heart-disease-in-black-women>
- Associated Press. (2026, January). Google and chatbot maker Character to settle lawsuit alleging chatbot pushed teen to suicide. AP News. <https://apnews.com/article/fbca4e105b0adc5f3e5ea096851437de>
- Burga, S. (2025, August 26). Parents allege ChatGPT is responsible for their teenage son's death by suicide. Time. <https://time.com/7312484/chatgpt-openai-suicide-lawsuit/>
- CBS News. (2023, November). UnitedHealth uses faulty AI to deny elderly patients care owed under Medicare Advantage plans, lawsuit alleges. CBS News. <https://www.cbsnews.com/news/unitedhealth-lawsuit-ai-deny-claims-medicare-advantage-health-insurance-denials/>
- KFF. (2025). KFF Tracking Poll on Health Information and Trust: Health information and advice on social media. <https://www.kff.org/public-opinion/kff-health-information-and-trust-tracking-poll-health-information-and-advice-on-social-media/>
- Lagasse, J. (2024). Class action lawsuit against UnitedHealth's AI claim denials advances. Healthcare Finance News. <https://www.healthcarefinancenews.com/news/class-action-lawsuit-against-unitedhealths-ai-claim-denials-advances>
- Lawrence, C. D. (2023). *Hidden in white sight: How AI empowers and deepens systemic racism*. CRC Press.
- Lee, B. C., Kim, K. I., Lee, J., Cho, K. H., & Moon, C. (2025). Effects of resistance training on osteosarcopenia in community-dwelling postmenopausal Korean women: Randomised controlled ERTOK trial. *Experimental Gerontology*, 209, 112869. <https://doi.org/10.1016/j.exger.2025.112869>
- Ma, M., Su, W., & Liu, D. (2025). Effects of different exercise interventions on bone mineral density in elderly postmenopausal women: A network meta-analysis. *Frontiers in Physiology*, 16, 1633913. <https://doi.org/10.3389/fphys.2025.1633913>
- Moore, J. X., Chaudhary, N., & Akinyemiju, T. (2017). Metabolic syndrome prevalence by race/ethnicity and sex in the United States, National Health and Nutrition Examination Survey, 1988–2012. *Preventing Chronic Disease*, 14, 160287. <https://doi.org/10.5888/pcd14.160287>
- National Institute of Standards and Technology. (2023). Artificial intelligence risk management framework (AI RMF 1.0) (NIST AI 100-1). <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf>

Nelson, A. (2025). An ELSI for AI: Learning from genetics to govern algorithms. *Science*, 389(6765). <https://doi.org/10.1126/science.aeb0393>

Office of the Federal Register. (2025, December 23). Request for information: Accelerating the adoption and use of artificial intelligence as part of clinical care (HHS-ONC-2026-0001). Federal Register. <https://www.federalregister.gov/documents/2025/12/23/2025-23641/request-for-information-accelerating-the-adoption-and-use-of-artificial-intelligence-as-part-of-clinical-care>

Ohu, F. C., Burrell, D. N., & Jones, L. A. (2025). Public health risk management, policy, and ethical imperatives in the use of AI tools for mental health therapy. *Healthcare*, 13(21), 2721. <https://doi.org/10.3390/healthcare13212721>

Ramírez Stieben, L. A., Brun, L. R., Pustilnik, E., et al. (2025). Beyond estrogen deficiency: The independent role of FSH in site-specific bone loss in midlife women. *Calcified Tissue International*. <https://doi.org/10.1007/s00223-025-01450-1>

Reddy, S., Allan, S., Coghlan, S., & Cooper, P. (2020). A governance model for the application of AI in health care. *Journal of the American Medical Informatics Association*, 27(3), 491–497. <https://doi.org/10.1093/jamia/ocz192>

Tabachnick, C. (2026, January 7). AI company, Google settle lawsuit over Florida teen's suicide linked to Character.AI chatbot. CBS News. <https://www.cbsnews.com/news/google-settle-lawsuit-florida-teens-suicide-character-ai-chatbot/>

Whitman, P. W., Alexander, C. J., Kaluta, L., Burt, L. A., & Gabel, L. (2025). Does exercising during peri- or early post-menopause prevent bone and muscle loss: A systematic review. *Bone*, 201, 117650. <https://doi.org/10.1016/j.bone.2025.117650>