Assistant Secretary for Technology Policy
Office of the National Coordinator for Health Information Technology
Mary E. Switzer Building, Mail Stop 7033A
330 C Street SW
Washington, DC 20201

Submitted via regulations.gov

RE: RIN 0955-AA13 Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care

Dear Deputy Assistant Secretary Posnack,

The National Health Law Program (NHeLP) is a public interest law firm working to advance access to quality health care and protect the legal rights of low-income and under-served people. NHeLP has a long history of advocacy regarding the use of automated decision-making systems (ADS), including algorithms and other forms of artificial intelligence (AI) to determine health care eligibility and services.[1] We are particularly familiar with how the lack of transparency and accountability of many ADS impacts individuals' ability to understand and fight care denials or reductions. We have also seen how poor transparency in the ADS development and implementation process obscures opportunities for evaluation of errors or bias in these systems. We have also seen repeated instances of harm from ADS without individuals understanding ADS has been used or how to seek remedy for that harm. We appreciate the opportunity to provide information about HHS's approach to the use of AI as part of clinical care.[2] Our long advocacy history around these systems supports our position that the use of AI in health care is an opportunity to improve care, but there must be protections for people in the process that reflect the magnitude of potential harm.

NHeLP provides some overarching comments that apply to the RFI's queries generally, before answering questions 6, 9, and 10.a.

Innovation Can Be Supported While Providing Appropriate Protections.

Efforts to foster innovation in AI in clinical settings must be balanced against its real and potential harms. When safeguards are fully enforced, they can support innovation by helping to promote awareness and identify issues, including discrimination and unintended impacts that cause harm. Research into AI in healthcare has repeatedly shown that many tools incorporate bias, either because the statistical analyses or other prediction mechanisms the AI is based upon or trained with incorporates bias or lacks a full scope of data.[3] Research also shows that not all populations are typically represented in clinical data and this affects the reliability and bias of AI systems.[4] Our response to questions 6, 9, and 10.a of the RFI provide more information on this research. Overall, studies that examine whether bias exists in clinical algorithms consistently call for greater transparency and accountability around the use of these tools throughout their lifecycle because of the risk of bias and the limitations of the research that typically underlies these tools.[5] HHS must

[1] *See generally* Nat'l Health Law Program, Fairness in Automated Decision-Making Systems, https://healthlaw.org/algorithms/; Elizabeth Edwards & David Machledt, Nat'l Health Law Program, *Principles for Fairer, More Responsive Automated Decision-Making Systems* (May 15, 2023), https://healthlaw.org/resource/principles-for-fairer-more-responsive-automated-decision-making-systems/ [hereinafter NHeLP ADS Principles].

[2] HHS, Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care, 90 Fed. Reg. 60108 (Dec. 23, 2025), https://www.govinfo.gov/content/pkg/FR-2025-12-23/pdf/2025-23641.pdf.

[3] *See, e.g.,* HHS, AHRQ, IMPACT OF HEALTHCARE ON RACIAL AND ETHNIC DISPARITIES IN HEALTH AND HEALTHCARE (2023, addendum 2024), https://effectivehealthcare.ahrq.gov/sites/default/files/related_files/cer-268-racial-disparities-health-healthcare-addendum.pdf (finding algorithms have been shown to perpetuate, exacerbate, and sometimes reduce racial and ethnic disparities and important mitigation steps include increased transparency in development and implementation, promoting awareness, including in patients, of potential risk, and further research before widespread implementation).

[4] *Id.*; NHeLP ADS Principles (discussing history of bias in ADS); *see also* Abdullah Alsaleh, *The Influence of Artificial Intelligence on Individuals with Disabilities,* ACTA PSYCHOLOGICA (Feb. 2026), https://www.sciencedirect.com/science/article/pii/S0001691825013241 (discussing literature review of AI research including disability and noting the limited user diversity of much of the AI, the limitations of the AI tools, and issues such as ethical and privacy concerns); Ariana Aboulafia, CDT & Henry Claypool, AAPD, *Building A Disability-Inclusive AI Ecosystem A Cross-Disability, Cross-Systems Analysis Of Best Practices* (Mar. 11, 2025), https://cdt.org/insights/building-a-disability-inclusive-ai-ecosystem-a-cross-disability-cross-systems-analysis-of-best-practices/; *see generally* New York City Bar, *The Impact of The Use of Artificial Intelligence on People with Disabilities* (June 2025), https://www.nycbar.org/wp-content/uploads/2025/06/20221483-AnalysisAIDisabilities.pdf.

[5] *See* note 4, *supra.* The research cited repeatedly cites the need for more transparency and accountability in the use of these tools, both so clinicians understand the limitations and so that people can identify both when such tools are being used and can identify errors.

incorporate significant protections and accountability mechanisms into any regulatory approach.

Meaningful Transparency and Accountability Measures are Necessary for Public Trust in Clinical AI.

At a very basic level, when it is in use, all parties using the AI tool or whom are affected by the use of an AI tool must have sufficient information that an AI tool is in use, any limitations or biases of the tool, enough information to understand how the decision was made so that it can be challenged, and mechanisms to correct any wrongs made by the tool at an individual or larger population level. These basic, well-accepted standards are not new concepts for AI creators and will not hinder innovation, but will help foster comfortability with the use of AI and are critically important in the identification of problems. All of this is foundational to public trust in AI, which in turn fosters innovation.

Transparency alone is insufficient; HHS must provide strong oversight and investigation. Transparency is critical, but it puts the onus on individuals impacted by the AI to understand it, have the expertise to challenge it, and have the resources to do so. HHS can take lessons learned from mental health parity compliance that simple requirements of transparency may not be sufficient if there is not enforcement of those requirements and when greater expertise may be required to identify improper bias.[6]

HHS has existing enforcement tools that can help regulate AI tools. AI does not necessarily create novel legal issues with respect to existing civil rights laws like Section 504 of the Rehabilitation Act, the Americans with Disabilities Act, and Section 1557 of the Affordable Care Act (incorporating Title VI, Title IX, and The Age Discrimination Act of 1975). These laws and rights all still apply to various aspects of healthcare. Merely incorporating AI into how decisions are made does not insulate that decision or its impact from civil rights enforcement. However, the government's oversight and expertise in these areas must be strengthened as increased investigation and accountability is needed, especially as significant information and expertise imbalance will exist between those who develop and employ AI (payers and providers) and those who are harmed by it (patients and their families).

If the implementation of AI continues to occur with a great deal of secrecy and limited ability by users to identify problems, any public trust in the use of AI will erode quickly and thus stifle innovation. If HHS truly wants to support the promise of AI to improve health

[6] For a discussion of the history of mental health parity legislation, requirements of plans, how plans employed techniques that denied services in new ways, and subsequent legislation and enforcement activities, including supporting sources, see, e.g., Amicus Brief of Nat'l Health L. Program & The Kennedy Forum in E.W. v. Health Net, *available at* https://healthlaw.org/resource/amicus-brief-of-nhelp-and-the-kennedy-forum-in-e-w-v health/?fbclid=IwAR3sBc20UO2ZqOhO1GbjHGuJJeboMAQ1T1lmH6KIynOiIasEBoC5muxFwM.

care, it must focus on the choices that will protect individuals and thus public trust in these systems. Part of such efforts must include significant data transparency and accountability mechanisms. States are currently taking the lead on AI regulation and should be allowed to continue to develop these mechanisms. AI accountability is not an easy problem to solve, but very traditional approaches regarding transparency and accountability, including investigation by federal agencies under existing laws are important pieces to protecting individuals from harm and overall public trust in the use of AI.[7]

RFI Questions:

6. Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?

Studies on effective use of AI and positive beneficiary outcomes do not exist. Instead, studies have shown that AI tools deployed in clinical care have fallen short in many ways, most notably because of data issues. These issues include incompatible, inconsistent, outdated and incomplete data; biased data; and the failure of systems to communicate with one another (interoperability failures). These studies indicate that as HHS approaches supporting the use of AI in clinical settings, it must provide protections that address these well-known and identified issues. Each of these issues can generate harm to individuals, but can do so at speed and scale that has not been experienced before. Therefore, in order to protect individuals from harm and facilitate the long-term promise of AI innovation, HHS must ensure that it takes all due care in setting forth criteria and oversight in these early years of AI use.

Flawed Data. Poor design, faulty or missing data, and other implementation issues can cause AI tools to fail outright or fail to anticipate unexpected circumstances. Resulting service and coverage losses may affect whole populations, and often fall disproportionately on marginalized groups. According to the American Medical Association:

> [t]he risk of generative AI 'hallucinating,' 'confabulating,' or otherwise fabricating information in response to a user-generated query has been well documented. Notably, tools such as ChatGPT have shown a not-uncommon tendency to falsify references cited in response to these queries. Generative AI tools have demonstrated the ability to generate fraudulent scientific/medical literature. They

[7] *See, e.g.*, CHAI Patient Survey Report, COALITION FOR HEALTH AI 3(survey finding that an average of 51% of respondents said AI use makes them trust health care less, with just 12% saying it made them trust health care more), Paige Nong & Jodyn Platt, *Patients' Trust in Health Systems to Use Artificial Intelligence*, 8 JAMA Network Open 2 (Feb. 2025) (finding that nearly 60% of U.S. adults do not trust their health care systems to ensure that they will not be harmed by an AI tool, and that roughly 65% do not trust their health care systems to use AI responsibly).

> are also capable of plagiarizing, falsifying, or misrepresenting data in ways that could compromise research integrity. Additionally, retracted papers may have the ability to continue to impact the content generated by LLM-based tools, potentially leading to dissemination or inaccurate or otherwise discredited information.[8]

The scale of data involved in these technologies and the lack of transparency also poses many problems and "may obscure the patterns from human analysts."[9] Machine learning can detect patterns, but data volume does not automatically correct for bias, missing information, or misrepresentation. At the same time, some private entities (and developers) get exclusive rights to large databases which limits transparency and independent validation.

Incompatible Data and Interoperability Failures. There are major structural barriers to workflow integration and consistent data-sharing across the entire health care system. These data issues can lead to errors, inaccuracies, or limitations in AI systems. According to the GAO:

> Data are often siloed in different systems, such as medical imaging archival systems, diagnostic systems, EHRs, electronic prescribing tools, and insurance databases. The various siloes store data in varying formats that cannot easily be reconciled and aggregated even if the data could be shared. The more systems involved, the more difficult it can be to aggregate the data.[10]

One review of EHR systems found that "there are no clinical decision support AI tools using EHR data that have successfully scaled across a health care system."[11]

This is especially problematic in the context of Electronic Health Records (EHR). EHR standards usually vary not just across vendors, but also within the same hospital systems, leading to "disjointed, incomplete records" that negatively impact the accuracy of AI performance.[12] Issues that GAO described 10 years ago related specifically to Electronic Health Records (EHR) have yet to be resolved. These include "(1) insufficiencies in health data standards, (2) variation in state privacy rules, (3) accurately matching patients' health records, (4) costs associated with interoperability, and (5) the need for governance and

---

[8] Am. Med. Assoc., *AMA Principles for Augmented Intelligence Development, Deployment, and Use* (Nov. 12, 2024), https://www.ama-assn.org/system/files/ama-ai-principles.pdf; Rami Hatem, Briana Simmons, Joseph Thornton, *Chatbot Confabulations Are Not Hallucinations*, JAMA INTERN. MED. 183(19):1177 (2023), doi:10.1001/jamainternmed.2023.4231.
[9] Kun-Hsing Yu et al., *Artificial Intelligence in Healthcare*, 2 NAT. BIOMED. ENG'G 719 (2018).
[10] U.S. Gov't Accountability Off., *Artificial Intelligence in Health Care: Benefits and Challenges of Machine Learning Technologies for Medical Diagnostics*, GAO-21-7SP, at 21 (2020).
[11] Mark Sendak et al., *A Path for Translation of Machine Learning Products into Healthcare Delivery*, 11 EMJ INNOV. 43 (2020).
[12] John D. Halamka et al., *Early Experiences with Personal Health Records*, 15 J. AM. MED. INFORMATICS ASS'N 1 (2008).

trust among entities."[13] While CMS interoperability rules have made some progress in detailing and requiring integrated systems, they are still not completely implemented (and won't be until 2027).[14] Until this actually happens, AI systems will continue to rely on partial or inconsistent records.

Even when unified, AI vendors often have access only to structured electronic health record data about an individual, which commonly excludes information about social determinants of health, functional limitations, caregiver supports, or housing instability. As Health Affairs noted, these omissions can cause AI tools to underperform when predicting care needs for historically marginalized populations.[15]

Data used in AI, and AI, scales bias. Racial bias is also well-documented across AI tools in health care settings.[16] In one case, because an algorithm used cost as a proxy for health, Black patients were systematically under-identified for high-risk care management because the tool associated higher health care costs with sicker patients.[17] Other documented cases include biased transplant prioritization, unsafe clinical recommendations, and chatbot agents providing inaccurate medical advice.[18]

Bias across payers and settings is a specific problem in programs like Medicaid. Because Medicaid serves people with historically underserved identities, the risks of AI-driven clinical care decisions are especially serious. These tools are likely to rely on historical insurance data from across the system (not just Medicaid) - data that reflects long-standing inequities and bias. When that data trains automated systems, it can quickly lead to care denials at scale. Additionally, when AI is deployed in lower-resource settings without the same resources, recommendations can increase patient harm. As the GAO reported,

> [i]f an AI tool is developed using data from a high-resource hospital and later applied to a community-based hospital with a different patient population, the tool may not necessarily be safe for the new population or perform as effectively in the new location.[19]

---

[13] GAO-21-7SP, *supra* note 10.
[14] Patient Access and Interoperability Final Rule, 42 C.F.R. §§ 422, 431, 438, 457; 45 C.F.R. § 156.
[15] Michelle M. Mello et al., *The AI Arms Race in Health Insurance Utilization Review*, 45 HEALTH AFF. 6 (2026), https://www.healthaffairs.org/doi/10.1377/hlthaff.2025.00897.
[16] AHRQ, *supra* note 3; Patrick Tighe et al., AHRQ, *Artificial Intelligence and Patient Safety: Promise and Challenges*, PSNet, Agency for Healthcare Research & Quality (2024); Kerstin Denecke et al., *Artificial Intelligence and Patient Safety: Promise and Challenges* (2024); Aziz Sheikh et al., *Artificial Intelligence and Health Inequities*, 52 LANCET REG'L HEALTH-EUR. (2024).
[17] "Although these two variables are correlated—sicker patients generally need more care—disparities in health care access… mean that on average African American patients have lower medical costs at the same level of health as Caucasian patients." Ziad Obermeyer at al., *Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations*, 366 SCIENCE 447 (2019).
[18] Kerstin Denecke et al., AHRQ, *Artificial Intelligence and Patient Safety: Promise and Challenges* (2024).
[19] GAO-21-7SP, *supra* note 10, at 27.

Electronic health data often do not reflect specific populations, like Medicaid beneficiaries. "High-resource hospitals often were early adopters of EHR systems and thus may have larger volumes of high-quality electronic data that now can be used to develop and train AI tools. However, the data from such hospitals may underrepresent some patient populations."[20] For example, rare diseases and low-prevalence conditions are systemically excluded.[21] As a result, AI tools often perform poorest where clinical risk is highest.

Even where studies "document the accuracy of AI tools, accuracy does not necessarily represent clinical efficacy."[22] We have very, very few studies of beneficiary outcomes (19), compared with thousands of studies about AI performance.[23] "Overall, among thousands of studies [about the impact of AI decision-making on patient outcomes]…, only five without a high risk of bias studied patient-relevant outcomes, on the basis of an externally validated algorithm" and "[o]nly two showed at least one benefit, in depression treatment and with regard to end-of-life care." There are no studies on long-term patient outcomes.[24] Other reports and studies have found widespread instances of harm: AI hallucinations, delayed transplants for Black patients, inaccurate health information, and promotion of disordered eating.[25] In one example, "Physician-graders found the AI model made mistakes when describing images and explaining how its decision-making led to the correct answer."[26]

And the studies showing positive impacts are mostly flawed - either they are not validated or they do not reflect current medical study standards.[27]

---

[20] W. Nicholson Price II, *Medical AI and Contextual Bias*, 33 HARV. J.L. & TECH. 65 (2019); Vindell Washington et al., *The HiTECH Era and the Path Forward*, 377 NEW ENG. J. MED. 907 (2017)
[21] Cong. Rsch. Serv., Artificial Intelligence in Health Care, R48319 (2024).
[22] *Id.*
[23] Christian Wilhelm et al., *Benefits and Harms Associated with the Use of AI-Related Tools in Public Health*, 10 LANCET PUB. HEALTH (2025), https://www.thelancet.com/journals/lanepe/article/PIIS2666-7762(24)00314-4/fulltext.
[24] *Id.*
[25] Kerstin Denecke et al., *The Unexepected Harms of Artificial Intelligence in Healthcare: Reflections on Four Real-World Cases*, Vol. 325 STUDIES IN HEALTH TECH. & INFORMATICS (2025), https://doi.org/10.3233/SHTI250219.
[26] Nat'l Insts. of Health, NIH Findings Shed Light on Risks and Benefits of Integrating AI into Medical Decision-Making (2023).
[27] "[R]ecent developments more often enable implementation without causal evidence on benefits of algorithmic technologies, which is at odds with established requirements to medical devices. As with other medical technologies before, health authorities' approvals should be tied to patient-related evidence." Wilhelm, *supra* note 23. Yet, "few AI tools have been evaluated by RCTs [randomized clinical trials]" which is the gold standard of medical review. Derek C. Angus et al., *AI, Health, and Health Care Today and Tomorrow: The JAMA Summit Report on Artificial Intelligence*, 334 JAMA 1650 (2025), https://doi.org/10.1001/jama.2025.18490. Frequently, the tests of accuracy are not even complete. "[T]ools would need to be subjected to comparison on the same independent test set that is representative of the target population, but such comparisons are not

What we do know is that AI produces high rates of care denial, in some cases 16 times higher than is typical, but "[n]o studies have compared rates of denials or wrongful denials (those reversed on appeal) in reviews with and without AI, making it difficult to disentangle potential causes of rising denial rates or assess the impacts of AI use."[28] This is especially true in utilization management contexts, like prior authorization.[29] Despite all of these well-studied issues with data quality and denials, identifying and correcting errors is shifted onto those with the least access to information and advocacy: patients.

9. What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption/use of AI in clinical care?

Trust starts with transparency.[30] Providers, patients, and caregivers need to know when and how AI tools are used, and how these tools influence clinical decisions or coverage determinations. According to the AMA, "Poor system design and inadequate workflow considerations during implementation can also lead to alert fatigue and mistrust of the system among patients and healthcare providers, undermining the intended benefits of AI."[31] Before deployment of AI tools, national standards for transparency and disclosure are essential, along with systems that collect and analyze disaggregated data to evaluate how AI affects specific patient populations and outcomes.[32] Transparency about data sources is also critical to identify and correct bias, and federal oversight should incorporate existing bias-assessment frameworks to ensure AI does not exacerbate inequities.[33]

At the same time, patients and caregivers express deep concern about opaque and time-consuming appeals processes, particularly when AI is used in utilization management or coverage decisions. AI-driven denials are often difficult to understand, with unclear clinical criteria and insufficient explanations for denials. Simply asserting, or even requiring, that a

---

done." C.J. Kelly et al., *Key Challenges for Delivering Clinical Impact with Artificial Intelligence*, 17 BMC MED 195 (2019).

[28] Mello, *supra* note 15.

[29] Jennifer Lubell, *How AI Is Leading to More Prior Authorization Denials*, Am. Med. Ass'n (Mar. 10, 2025), https://www.ama-assn.org/practice-management/prior-authorization/how-ai-leading-more-prior-authorization-denials.

[30] Tighe, *supra* note 16.

[31] *Id.*

[32] NHeLP ADS Principles, *supra* note 1; NAIC Model Bulletin (Dec. 4, 2023), https://content.naic.org/sites/default/files/cmte-h-big-data-artificial-intelligence-wg-ai-model-bulletin.pdf.pdf; Angus, *supra* note 27.

[33] R. Agarawal et al., *Addressing Algorithmic Bias and the Perpetuation of Health Inequities: An AI Bias Aware Framework*, HEALTH POLICY & TECH. (Mar. 2023), https://www.sciencedirect.com/science/article/abs/pii/S2211883722001095.

"human is in the loop" is not enough.[34] Meaningful safeguards require truly independent medical judgment (absent any financial incentives), full disclosure of all clinical standards used in the decision-making process, and clear notice to patients about the specific reasons that triggered a denial. The appeal processes must also work, and in a timely way. And there must be further accountability mechanisms from regulators to address both individual and systemic issues. Without these protections, appeals become procedural hurdles rather than genuine opportunities to correct errors or appropriate denials.

Automation bias and financial incentives are equally concerning. Clinicians may defer to AI-generated recommendations even when those recommendations conflict with their judgment, particularly when those clinicians lack training on AI systems' limitations or are unaware of how the tool generates outputs. Patients worry that AI tools, especially those used by payers, are designed in ways that favor cost containment over medical necessity. Uncertainty about accountability compounds these risks. Stakeholders remain confused about who bears ultimate responsibility for AI-influenced decisions and whether they have meaningful protection if/when systems fail to meet standards.
Finally, beneficiaries and caregivers need robust regulatory protections, oversight, and enforcement. This includes timely appeals and grievance processes grounded in independent medical judgment. It also includes strong data protections and compliance with privacy laws. And it includes effective oversight, like the authority to investigate AI-related harms and impose penalties on bad actors, as well as enforcement of existing protections in use cases like prior authorization. Without enforceable standards, transparency, and accountability, patients and the general public fear that AI adoption will entrench existing problems in the health care system, rather than address them.

10. Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care?

A. Are there published findings about the impact of adopted AI tools and their use in clinical care?

Published findings about the impact of AI tools on health care are widespread, and they broadly support the same conclusion: these tools can do great harm to impacted individuals and must be used with care, if at all. Some types of tools and use have greater risk than others. Although there are no universally agreed upon mechanisms for monitoring AI tools, there are thru lines that say transparency, accountability, data collection, and other robust guardrails should be part of any AI use, especially for higher risk uses like health care decisions.

---

[34] *See, e.g.,* Hannah Quay-de la Vallee & Kevin De Liban, CDT & Benefits Tech Advocacy Hub, *Rethinking the Loop: Encircling Public Benefits AI with Human Oversight* 11-12 (Sept. 2025), https://cdt.org/wp-content/uploads/2025/09/2025-09-22-Humans-in-the-Loop-CDT-Civic-Tech-report-final.pdf (discussing the limitations of humans in the loop as an accountability mechanism for AI systems).

Existing research shows that AI and other automated systems have real risks to their use that can cause great harm. Discriminatory bias can be baked into these tools; in the health care context, this innate bias can exacerbate health disparities and impact patient outcomes.[35] AI tools also frequently get things wrong.[36] For instance, an AI tool used to predict a given patient's care needs may fail to consider all factors pertinent to their needs; a tool used to predict the success of a prior authorization or reimbursement request may deprioritize claims similar to those that have failed in the past without performing an independent assessment of their merit, reinforcing payers' bad behavior and leaving patients with little recourse.[37] AI tools and other algorithms engage in opaque decision-making, making it unreasonably difficult to challenge adverse decisions.[38] The humans

---

[35] *See generally* AHRQ, *supra* note 3 (collecting research on racial bias in clinical algorithms); Elizabeth Edwards, David Machledt, Skyler Rosellini, & Liz McCaman, *NHeLP Comments on AHRQ Request for Information on the Use of Clinical Algorithms That Have the Potential to Introduce Racial/Ethnic Bias Into Healthcare Delivery* (May 4, 2021), https://healthlaw.org/resource/nhelp-ahrq-comments/ [hereinafter "NHeLP AHRQ Comments"] (discussing various types of bias in clinical algorithms); *see also, e.g.*, Mahmud Omar et al., *Sociodemographic Biases in Medical Decision Making by Large Language Models*, 31 NATURE MED. 1873 (Apr. 2025), https://www.nature.com/articles/s41591-025-03626-6 (disparate pain management recommendations from an LLM for LGBTQIA+ individuals); Crystal T. Chang et al., *Evaluating Anti-LGBTQIA+ Medical Bias in Large Language Models*, 4 PLOS DIG. HEALTH 9, e0001001 (Sept. 8, 2025) https://journals.plos.org/digitalhealth/article?id=10.1371/journal.pdig.0001001 (prompts describing clinical scenarios with "LGBTQIA+ identity terms" elicited more severe bias from LLMs); Melissa D. McCradden et al., Hannah E. Knight et al., *Challenging Racism in the Use of Health Data*, 3:3 THE LANCET E144 (Feb. 3, 2021), https://www.thelancet.com/journals/landig/article/PIIS2589-7500(21)00019-4/fulltext (racial bias in the use of health data generally); *Ethical Limitations of Algorithmic Fairness in Health Care Machine Learning*, 2 THE LANCET E222 (May 2020), https://www.thelancet.com/journals/landig/article/PIIS2589-7500(20)30065-0/fulltext; Ziad Obermeyer et al., *Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations*, 366 SCIENCE 447 (2019), https://escholarship.org/uc/item/6h92v832#main (racial bias in population health algorithm); Darshali A. Vyas et al., *Challenging the Use of the Vaginal Birth After Cesarean Section Calculator*, 29-3 WOMEN'S HEALTH ISSUES 201 (2019), https://www.whijournal.com/article/S1049-3867(19)30098-2/pdf; Ibram X. Kendi, *There is No Such Thing as Race in Health-care Algorithms*, 1:8 THE LANCET E375 (Dec. 2019), https://www.thelancet.com/journals/landig/article/PIIS2589-7500(19)30201-8/fulltext (history of racial bias in health care can be perpetuated and amplified by AI tools); Foad Hamidi et al., *Gender Recognition or Gender Reductionism? The Social Implications of Automatic Gender Recognition Systems*, CHI HUM. FACTORS IN COMPUTING SYS. (Apr. 2018) (impact of "Automatic Gender Recognition" (AGR) algorithms on transgender individuals); Anupam Datta et al., *Proxy Discrimination in Data Driven Systems* (2017), https://arxiv.org/pdf/1707.08120.pdf ("Machine learnt systems inherit biases against protected classes, historically disparaged groups, from training data," which can be difficult to see because "they rely on subtle correlations discovered by training algorithms").

[36] *See, e.g.*, Qiao Jin et al., *Hidden Flaws Behind Expert-Level Accuracy of Multimodal GPT-4 Vision in Medicine*, 7 NPJ DIG. MED. 190 (2024) (GPT-4V was comparably accurate to human physicians when asked multiple-choice questions about imaging it had reviewed, but its rationales were often flawed).

[37] *See* Mello, *supra* note 15.

[38] *See* NHeLP AHRQ Comments, *supra* note 34, at 3 ("The rise of ADS has exacerbated the problem of 'black box' decision-making. This has multiple impacts. A person denied care may simply be told

responsible for operating and monitoring these tools are also fallible. Automation bias – the tendency of humans to accept a computer-generated solution as correct irrespective of its accuracy and ignore non-automated contradictory information – is a significant risk and can have catastrophic consequences if not addressed and mitigated.[39] Finally, public mistrust of AI in health care remains high, impacting its utility.[40]

AI tools in health care certainly have the potential to be helpful. They can automate simple tasks and reduce administrative burdens, "free[ing] up human capacity," allowing providers and agency staff to serve people with "dignity and efficiency."[41] But safeguards are necessary to realize the full potential of these tools while minimizing their potential harm. Additional testing and data collection is imperative. AI tools should be studied for their impact on patient outcomes, construed broadly to ensure that all aspects of the health care experience are meaningfully improved without harm. Any strategy to deploy these tools must include robust guardrails and protections.[42] In sum: moving too fast magnifies the already-significant risk of harm posed by health care-related AI applications. Innovation is often a net positive, but innovation for innovation's sake should not be the focus when it comes to developing and expanding health-related AI.

---

the results of the ADS without any information about what about their condition triggered the denial, which may deter appealing the denial") (citations omitted).

[39] *See, e.g.*, Giuseppe Romeo & Daniela Conti, *Exploring Automation Bias in Human-AI Collaboration: A Review and Implications for Explainable AI,* 41 AI & SOC'Y 259, 262 (Jul. 2025); Lauren Kahn et al., *AI Safety and Automation Bias: The Downside of Human-in-the-Loop* 6, CTR. FOR SEC. & EMERGING TECH. (Nov. 2024) (commenting that the "challenge of automation bias has only grown with the introduction of progressively more sophisticated AI-enabled systems and tools across different application areas" and noting failure to heed AI warnings); *see also* Mello et al., *supra* note 15, at 9 (research team conducting ethical assessments of AI tools at Stanford Health Care found that some health system staff "could not explain in even basic terms how generative AI creates output, did not know that AI could be biased, and could not name any ways in which the tools might have performance problems"); Moustafa Abdelwanis et al., *Exploring the Risks of Automation Bias in Healthcare Artificial Intelligence Applications: A Bowtie Analysis*, 5 J. SAFETY SCI. & RESILIENCE 460, 463 (Dec. 2024) (commenting that automation bias poses a substantial risk to patients' lives through medical errors such as misdiagnosis and incorrect prescriptions).

[40] *See* note 7, *supra.*

[41] *See* Danny Mintz & Jennifer Thom, *Identifying Use Cases for AI in Medicaid Delivery*, CODE FOR AMERICA (Jan. 21, 2026), https://codeforamerica.org/news/identifying-use-cases-for-ai-in-medicaid-delivery/.

[42] For instance, any adopter of an AI tool should conduct pilot studies before the tool is implemented and develop "evaluation criteria and success metrics" to ensure the tool is effective and does not amplify existing health inequity. *See* Sophia Fox-Sowell, *Can AI Help States Prepare for Medicaid Changes Coming in 2027?*, STATESCOOP (Jan. 22, 2026), https://statescoop.com/states-medicaid-changes-hr1-ai/. And HHS should develop a multifaceted accountability framework including *meaningful* human-in-the-loop oversight and radical transparency about the tool itself and how it works (including known weaknesses and risks, and how vendors will monitor the tool's performance); *see, e.g.*, CHAI (noting that similar strategies will help improve public trust in health care AI tools).

Conclusion

For the reasons stated above, we urge HHS to balance carefully the support of innovation with the needed protections for people affected by AI in clinical decision making. We direct HHS to the research we have cited and made available through active links, the full text of which should be part of the record for this RFI. We appreciate the opportunity to provide comment on this issue. If you have any questions, please feel free to contact Elizabeth Edwards, edwards@healthlaw.org, Brit Vanneman, vanneman@healthlaw.org, or Shandra Hartly, hartly@healthlaw.org.

Sincerely,

Elizabeth G. Taylor
Executive Director