AMERICAN PSYCHIATRIC ASSOCIATION

800 Maine Avenue, S.W.
Suite 900
Washington, D.C. 20024

Thomas Keane, MD, MBA
Assistant Secretary for Technology Policy
National Coordinator for Health Information Technology
Department of Health and Human Services
Mary E. Switzer Building
330 C Street SW
Washington, DC 20201

Re: HHS Health Sector AI RFI [RIN 0955-AA13

Dear Assistant Secretary Keane,

The American Psychiatric Association (APA) appreciates the opportunity to respond to ONC/ASTP's Request for Information on the use of artificial intelligence in clinical care. Psychiatrists approach AI cautiously because behavioral health information is uniquely sensitive, and patient trust is foundational to effective treatment. Accordingly, APA's comments emphasize privacy, security, transparency, and evidence. While AI may offer promise, particularly in reducing administrative burden, adoption in psychiatric practice should be guided by rigorous, real-world evaluation, clear accountability, and safeguards that prevent secondary use or unintended disclosure of mental health data. The recommendations to certain questions posed by HHS below focus on practical HHS actions that support responsible innovation without weakening existing protections.

Barriers to private sector innovation in AI

Behavioral health clinicians are approaching AI cautiously, and the biggest barriers to private sector innovation and adoption are not a lack of tools but a lack of trusted, consistent guardrails, leaving patients and providers without clear transparency into when AI is used, how it performs, and what happens to highly sensitive mental health data, including whether it is used without awareness or consent or reused for other purposes. To reduce these barriers while preserving innovation, HHS should require or at minimum strongly incentivize, independent healthcare AI accreditation as a scalable way to set risk-proportionate expectations for quality, transparency, safety, and privacy, so developers have a clear pathway to trust and providers have confidence that accredited tools meet baseline protections.

Regulatory changes to incentivize the effective use of AI

HHS can accelerate the safe and effective integration of AI into clinical care by prioritizing targeted regulatory, payment, and programmatic reforms that align innovation with patient safety, transparency, and accountability. Additionally, APA recommends:

- Stronger HIPAA-era rules for modern data flows:
  Modernize privacy and security safeguards to address contemporary data ecosystems, including AI model training, secondary data use, and vendor access to protected health information (PHI). HHS should require clearer limits on data reuse and model training, enhanced transparency for patients, and stronger accountability standards for third party vendors handling PHI. Strengthened data governance would enhance patient safety, protect privacy, and build public trust which is critical prerequisites for widespread AI adoption.
- Payment for safety work:
  Establish reimbursement mechanisms that recognize clinician time spent reviewing, validating, and monitoring AI-generated outputs, as well as payment pathways for evidence-based digital and AI-supported care. Without reimbursement, safety oversight becomes uncompensated labor, creating disincentives for responsible implementation. Payment reform would align financial incentives with safe, high quality AI integration.
- Baseline requirements for "clinical AI" transparency:
  Develop national baseline standards requiring disclosure of model purpose, intended use population, performance characteristics, known limitations, drift monitoring processes, and expectations for human oversight. Clear transparency requirements would reduce unsafe deployment, support informed clinical decision-making, and ensure that AI tools are used within validated boundaries.
- National adverse-event reporting for AI:
  Create a standardized, nationwide mechanism to report AI-related harms and near misses, similar to pharmacovigilance or medical device reporting systems. A centralized reporting infrastructure would enable early detection of safety signals, promote shared learning across health systems, and support continuous improvement of AI tools.

Regulations or policies to revisit

HHS can modernize privacy protections, clearer accountability structures, and strengthen federal standards to ensure that AI in clinical care, particularly in psychiatric settings, is implemented safely, transparently, and with strong patient trust:

- HIPAA Privacy/Security (OCR):
  - Strengthen and clarify HIPAA guardrails: Provide clear guidance under existing HIPAA regulations that limits AI to purpose-bound uses, prohibits secondary use of health data without explicit patient authorization, and requires enforceable safeguards (data minimization, strong access controls, and auditable logs) across the full AI workflow.
  - Limit secondary use for "product improvement," requiring robust audit logs, and tightening modern de-identification expectations to minimize re-identification risk.
- 42 CFR Part 2 (SAMHSA): Update Part 2 implementation guidance so behavioral health such as SUD information can support safe AI-enabled care only with granular, revocable consent, clear redisclosure limits, and enforceable segmentation and access controls across systems.
- ONC/ASTP Interoperability, Certification & Information Blocking: Require certified health IT to support privacy-preserving, least-privilege data access (including segmentation tags, provenance, and purpose-of-use), and clarify that patients can obtain AI-influenced clinical content in the record without exposing sensitive psychotherapy notes beyond existing protections.

- Do not roll back certification transparency requirements for AI-enabled health IT including standardized model documentation (“model cards”) and instead strengthen them so psychiatric providers can evaluate intended use, data handling/retention, limitations, and subgroup performance before adoption.

Novel legal and implementation issues

The primary legal and implementation challenges of non-device AI in clinical care include unclear liability when harm occurs, uncertainty about discoverability of AI-generated content in the medical record, and insufficient transparency and consent frameworks particularly when sensitive psychiatric information is involved. APA recommends establishing clear, system-level standards for the governance of non-device AI used in clinical care. Federal agencies, in collaboration with developers and health care organizations, should clarify liability frameworks, documentation expectations, and discoverability rules for AI-generated content.

Addressing novel legal and implementation issues

HHS should play a convening and standard-setting role that reduces uncertainty without lowering privacy or safety protections:

- Unified guidance: HHS should publish clear rules for non-device AI use in health care, including human review and escalation.
- Behavioral health privacy: HHS should limit psychiatric data use to purpose-bound, minimal access with default non-retention and no secondary use without explicit authorization.
- Keep certification transparency: HHS should not roll back ONC certification transparency (including model cards) and should strengthen it for limits, updates, and data practices.

Promising AI evaluation methods

The most promising evaluation approaches for non-medical AI in clinical care combine rigorous pre-deployment validation with continuous post-deployment monitoring (drift detection, adverse-event reporting, etc.), alongside human-centered assessments of usability, clinician reliance, documentation accuracy, and impact on patient safety and equity.

- HHS should align evaluation methods for non-medical device AI used in clinical care to existing evaluation methods such as NIST AI Risk Management Framework which emphasize (1) governing with clear accountability, documented intended use, change control, and incident response; (2) mapping the clinical context and foreseeable harms (including behavioral health privacy and high-risk workflows like crisis triage); (3) measuring performance using pre-deployment test sets and real-world “shadow mode” that quantify safety-critical errors, subgroup impacts, usability/workload, and privacy/security leakage; and (4) managing risk post-deployment through continuous monitoring for drift, trigger-based audits for high-risk content, transparent update notices, and rollback thresholds tied to patient safety and trust.

HHS support of evaluation methods

APA recommends funding shared evaluation infrastructure and finance privacy-preserving, multisite “AI evaluation networks” (secure enclaves/federated evaluation) so providers can test tools without centralizing sensitive psychiatric data.

Administrative hurdles to the adoption of AI

While administrative complexity is often cited as a barrier to AI adoption in clinical care, the more fundamental challenge is trust. Clinicians and patients must have confidence that AI tools are accurate, safe, privacy-protective, and aligned with clinical judgment before they will meaningfully integrate them into care.

- Model bias: without real-world evidence on error modes, drift, and subgroup performance, reimbursing clinical AI would prematurely scale tools that can have disparate impacts on protected and clinically vulnerable groups. HHS should condition any payment pathway on transparent accountability expectations plus rigorous, stratified outcome evaluation that demonstrates meaningful clinical improvement and equitable performance in routine practice.

- Privacy uncertainty slows interest: If it's not crystal clear how sensitive notes and messages are protected (and not reused), clinicians prefer to wait.
- Lack of evidence for implementation into practice: Until there are robust, real-world evidence of benefit in behavioral health settings and without unintended harm, enthusiasm remains cautious.
- Fear of clinical harm: clinicians worry about tone, stigma, mislabeling, and overconfidence in summaries or "risk" outputs that could steer care incorrectly.
- Accountability is unclear: If an AI-generated summary is wrong or a crisis message is mishandled, clinicians want to know who is responsible, what are the steps to correct, and what protections exist.
- Workflow disruption is a concern: Even helpful tools can add steps (reviewing, correcting, documenting), and clinicians do not want new burdens disguised as "automation."
- Patient acceptance is unknown: Providers expect some patients to object to AI involvement in sensitive communications, and clinics don't want to jeopardize engagement.
- The bar is higher in mental health: Because the data are deeply personal and harms can be serious, psychiatrists and their patients expect stronger guardrails than in many other specialties.

Areas of AI research that HHS should prioritize

There are ways HHS can accelerate the adoption of AI by developing evidence that is lacking today including:

- Pragmatic clinical trials for AI-in-workflow: Fund multisite pragmatic trials (not just retrospective accuracy studies) that test AI's impact on outcomes, safety events, clinician workload, and patient experience.
- Real-world effectiveness in behavioral health: Pragmatic studies showing where AI helps and does not in psychiatry including documentation burden, access, continuity, and outcomes, using credible comparators.
- Head-to-head comparative effectiveness: Support studies comparing AI-assisted care versus best existing tools (human scribes, templates, standard triage protocols), not just "AI vs. nothing."

Thank you for your review and consideration of these comments. If you have any questions or would like to discuss any of these comments further, please contact Zuhal Haidari (zhaidari@psych.org), Deputy Director, Digital Health.