



February 23, 2026

The Honorable Thomas Keane, MD, MBA
Assistant Secretary for Technology Policy
National Coordinator for Health Information Technology
U.S. Department of Health and Human Services
Mary E. Switzer Building
330 C Street, SW
Washington, DC 20201

RE: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as part of Clinical Care (RIN 0955-AA13)

Dear Assistant Secretary Keane:

Amazon.com, Inc. (Amazon) appreciates the opportunity to comment on opportunities to accelerate the use of artificial intelligence (AI) to improve care delivery, empower individuals, and save lives. We support the Assistant Secretary for Technology Policy/Office of the National Coordinator for Health Information Technology (ASTP/ONC) mission to encourage the safe adoption of AI in clinical care.

Over the last 25 years, Amazon has developed and deployed state-of-the-art AI and machine learning models to power customer experiences. This includes customer-facing services and internal operations that personalize the shopping experience on Amazon.com, the AI-powered robots that optimize order fulfillment in our warehouses, and Amazon Web Services (AWS) building generative AI capabilities that transform customer experiences in industries across the globe.

In healthcare, AI is increasingly driving the ability for patients, physicians, and scientists to use the information they have to improve health outcomes. Through use of AI, patients can get insights early to make decisions on where and how to get care; healthcare providers can spend more time directly caring for their patients—not behind a monitor; and scientists can surface insights on the root causes of disease to then develop new innovations to save and improve lives.

Amazon Pharmacy and One Medical are leveraging AI to transform healthcare delivery, reduce administrative burden, and make getting health care easier for patients. Amazon Pharmacy uses AI to process prescriptions 90% faster by automatically structuring prescription data, predicting medication inventory needs, providing real-time price estimates without manual insurance entry, and enhancing customer support through AI-powered documentation analysis. Beyond operational efficiency, AI-powered clinical decision support can improve patient safety by identifying potential adverse drug interactions, predicting patient deterioration for early intervention, and accelerating diagnosis of time-critical conditions through AI-enhanced imaging triage.

One Medical is deploying generative AI to handle routine administrative tasks such as reviewing health records, identifying care gaps, routing messages, and automating workflows—giving providers more time to focus on patient care rather than paperwork. For patients, One Medical is using AI to make accessing health care even simpler, more highly personalized, and more actionable through Health AI that can answer health questions, book appointments, and help manage medications. One Medical's Health AI assistant provides 24/7 personalized health guidance, can explain lab results, book

appointments, and manage medications with HIPAA-compliant security. Health AI was developed with One Medical providers providing clinical guidance across the entire development lifecycle, embedding safety principles from conception through launch. We are committed to using advanced technology to make healthcare more personal, accessible, and human-centered while respecting patient privacy and enhancing clinical decision-making.

We are dedicated to making our customers' and patients' lives better while also establishing and implementing the necessary safeguards to protect them. Amazon has a long-standing track record of protecting customer data. For example, One Medical and Amazon Pharmacy protect protected health information (PHI) with HIPAA-compliant privacy and security practices to keep information safe and secure. All our services go through rigorous security testing, and we maintain strict administrative, physical, and technical safeguards to protect our systems and customer information, such as through encryption technology and strict controls over who can access records.

Through AWS, we support healthcare providers, insurers, electronic health record (EHR) vendors, and other organizations to leverage AI to generate efficiencies and improve patient outcomes. AWS offers purpose-built healthcare AI services, including HealthScribe to support clinical documentation, Health Lake to enable data interoperability, and HealthImaging for faster analytics of medical images. AWS also recently announced agentic AI capabilities in Amazon Connect—which automates routine patient interactions like appointment scheduling that represent 50% of call volume. We also developed an open-source Healthcare and Life Sciences Agentic AI toolkit, which includes starter agents built on Amazon Bedrock. And, we recently released Amazon Nova Forge, which empowers organizations to build custom frontier AI models tailored to their specific needs while maintaining data privacy.

We encourage ASTP/ONC to take several steps to further accelerate the benefits of AI adoption in care and to work with innovation-forward states testing novel AI technologies within their regulatory jurisdictions.

Considerations for Oversight of AI/ML in Healthcare

When evaluating the role of government in overseeing the use of AI/ML in healthcare, we encourage you to consider several factors prior to the development of any new policies.

First, ASTP/ONC should consider whether the robust policies already in place have gaps unique to AI. For example, the Food and Drug Administration (FDA) is already well suited with its existing authorities to oversee the use of AI for regulated products, including clinical decision support tools that meet the definition of medical devices. In fact, as of late 2025, FDA had already cleared or approved more than 1350 devices with AI/ML components—dating back to 1995. ASTP/ONC should only focus new policies on AI-specific gaps in existing authorities. We also caution that new policies that duplicate existing authorities can raise complications and burdens that do not generate any further risk mitigation or patient care benefits. Instead, duplicative regulation may slow deployment of health technologies that could improve the quality of care. Therefore, we recommend ASTP/ONC publish a formal mapping of existing regulations to AI use cases in clinical care, identifying only those gaps that require new policy intervention. This would prevent duplicative requirements that could delay deployment of beneficial technologies.

Second, Amazon supports a risk-based approach to AI/ML oversight and appropriate transparency for patients and clinicians. Recognizing that AI applications are heterogeneous, a focus on high-risk activity helps to ensure that potential harms are addressed without stifling innovation or impeding access to

benefits from low-risk uses of AI. In the healthcare context, risk-based approaches are already widely adopted—with policies varying based on the specific use cases and effect of the technology on patient care (e.g., as medical devices are already regulated based on risk). We appreciate ASTP/ONC's recent proposed rule that would rescind the Decision Support Intervention transparency requirements as they required disclosures of source attribute information regardless of the risk associated with the technology.

Third, when promoting transparency to support patient and clinician trust in AI-based tools, we encourage ASTP/ONC to first evaluate the methods that AI developers already use to disclose information on models—such as AWS AI Service Cards that provide appropriate disclosure about model capabilities, limitations, and intended uses. When fostering transparency, we caution against the disclosure of proprietary information, such as training data, training recipes, or model architecture, which are confidential commercial secrets with significant competitive value. If regulators require model developers to disclose information about such data, it may foreclose the developer's ability to claim trade secret protection under the law. These concerns are only amplified if organizations seek to use larger models, like foundation models, which are trained on a central corpus of internet scale data and are intended to be used by many customers across many use cases. Training data disclosure requirements are therefore impractical and not well suited to evaluating appropriateness for use. Training datasets vary based on many factors, including size, demographic makeup (the number and type of groups), quality (amount of confounding variation, fit for purpose), the types and quality of labeling, and other factors. Simply because a model has been trained on certain data does not mean it will perform as needed in a specific use case.

Finally, we encourage ASTP/ONC to align with international standards—many of which are currently in development—to provide consistency, clarity, and thoroughness to support the effective and safe development and use of AI in healthcare. For example, ISO 42001 is a management system standard that provides guidelines for managing AI systems within organizations. This standard emphasizes a commitment to responsible AI practices, encouraging organizations to adopt controls specific to their AI systems, fostering global interoperability, and setting a foundation for the development and deployment of responsible AI. Amazon is committed to the responsible development of AI, as demonstrated by AWS being the first major cloud service provider to announce ISO 42001 accredited certification for several of its services.

Apply new policies to applications, not specific technologies

Many uses of AI provide enhancements to existing healthcare workflows and processes. For example, clinicians already use diverse approaches to document care, including via: direct entry by clinicians; template-based documentation tools; in-room medical scribes; remote scribes; and transcription technologies based on AI. Each approach carries different risks and benefits; applying regulatory requirements on only one of those modes could impede innovation and adoption.

Instead, when evaluating policy guidance or regulations on an aspect of clinical care, ASTP/ONC should consider existing quality and safety requirements applicable to that use case. Taking an approach grounded in the use case—rather than specific tools, such as AI—would avoid singling-out certain technologies and stifle innovation in their development and application.

Focus on testing in deployed settings

As mentioned above, regulatory focus on training data does not provide actionable information for patients, clinicians, and deployers. Especially for high-risk healthcare use cases, rather than focusing on

training data, testing AI-based tools in the intended environment with data reflecting the intended population is the scientifically appropriate and industry-standard mechanism that provides the best indication of performance.

There are some parallels between AI and EHR systems. With EHRs, the performance and safe use of the system can vary based on differences in users, patient populations, the prevalence of other technologies, workflows, and many other factors. That's why, with EHRs, testing in the intended environment is a core component of evaluating the safety and usability of the system.

The same applies with AI, where deployers—such as a health system—should test AI-based technologies in their environments to evaluate whether the tool works appropriately and—where needed—apply risk mitigations. ASTP/ONC should consider guidance to health systems on how to assess the performance of technologies in their intended settings, much like how the agency led the development of the Safety Assurance Factors for Electronic Health Record Resilience (SAFER) Guides for evaluation of EHR performance in healthcare organizations.

Encourage greater data interoperability to improve care, support AI

EHR systems, claims databases, and other technologies increasingly house and make available information to empower patients and improve care. When coupled with AI, this information can help: identify treatments for unmet needs; equip patients and clinicians with personalized recommendations through decision support tools; and otherwise generate insights to better understand the safety and quality of care.

Yet, the data currently accessible via standard application programming interfaces (APIs) is not yet fully sufficient to support many improvements to the patient experience and enhanced care. For example, the versions of the US Core Data for Interoperability (USCDI) currently required for support by certified EHRs lacks access to cancer care data, genomic data elements links to download medical images—including X-Rays and MRI scans. API-based access to these data elements, when used with AI-based technologies, can enable enhanced patient-specific guidance that factors in each individual's unique clinical situation. AI integration with EHR systems can transform clinical workflows by providing real-time decision support, automating documentation, and enabling seamless data exchange through FHIR-based APIs. AWS HealthLake provides infrastructure for organizations to meet ASTP/ONC and CMS interoperability requirements. We encourage ASTP/ONC to make these—and other—data elements available via standard APIs to accelerate the utility of AI-powered tools in clinical care.

Relatedly, with appropriate patient authorization and in compliance with HIPAA, healthcare organizations have sought to use protected health information (PHI) to develop AI tools to support the treatment, payment, and operations of facilities. Health systems have also sought to use de-identified data to improve the capabilities of AI-based technologies, however, the de-identification of data may make the information less useful. Through AWS Clean Rooms, healthcare organizations can securely analyze and collaborate on their collective datasets—without sharing or copying one another's underlying data. This technology, when coupled with AI, can further enable the development of health-specific technologies without disclosure of the underlying, sensitive data. This service can allow organizations to discover, evaluate, and access patient cohort data in a privacy-preserving manner and without any data movement.

Thank you for considering our recommendations, and we look forward to serving as a resource to ASTP on accelerating the ability for AI to improve patient care.