

RESPONSE TO HHS REQUEST FOR INFORMATION (RFI)

Accelerating the Adoption and Use of Artificial Intelligence (AI) as Part of Clinical Care

Submitted by:

Acoer, Inc.

Jim Nasr, Chief Executive Officer

jnasr@acoer.com | acoer.com | healthdataexplorer.io

Atlanta, Georgia

Docket:

HHS-RIN-0955-AA13

Submission Date:

February 23, 2026

About Acoer, Inc.

Founded by a former Chief Software Architect at US Centers for Disease Control (CDC) and developers of many large-scale open health platforms, Acoer, Inc. is a health technology company headquartered in Atlanta, Georgia. Acoer specializes in transforming complex, unstructured public health and clinical data into actionable intelligence. Our flagship platform, Health Data Explorer (healthdataexplorer.io), is deployed across medical examiner offices, forensic agencies, and public health departments to modernize medicolegal death investigation through private AI-driven analytics.

Acoer has built and deployed systems that directly address the United States opioid and overdose crisis, working in partnership with programs supported by the CDC Foundation and institutions aligned with the Bloomberg Philanthropies Overdose Prevention Initiative. Our work converts unstructured autopsy narratives, toxicology reports, and investigative documentation into structured, cross-case intelligence. We enable earlier detection of emerging drug combinations, identification of environmental risk clusters, and coordinated system-level prevention response.

Acoer is a recognized innovator in privacy-preserving AI architecture, data sovereignty, FHIR-native interoperability, and cryptographic data governance. We use distributed ledger technology (specifically the Hedera network) to create tamper-evident, immutable audit trails that provide verifiable proof of what data was used, when, and under what governance conditions. These capabilities are all in production today, not prospective development or research concepts.

High-level summary of our capabilities:

Health Data Explorer Data analytics and reporting for public health	Forensic Intelligence Platform Private AI technologies for forensic data	Cryptographic Audit & Data Verification Blockchain-based data sovereignty and governance	FHIR-Native Interoperability HTI-5 aligned, data sovereignty-first architecture
---	--	--	---

Executive Summary

Three Imperatives for HHS AI Strategy

1. Reduce Preventable Deaths: harness AI as a real-time early warning system against overdose, emerging pathogens, and chronic disease onset.
2. Deploy AI Safely: anchor AI trust in cryptographic verifiability, data sovereignty, and accountability.
3. Enable Intelligent Syndromic Surveillance: close the gap between mortality data, clinical systems, and prevention response.

More than 70,000 Americans die every year from drug overdoses. This is a crisis costing the United States an estimated \$1.5 trillion annually in healthcare, lost productivity, and criminal

justice expenditures. The majority of these deaths are preventable, and the intelligence needed to prevent them exists today, often buried in unstructured autopsy narratives, toxicology reports, and emergency department records that no human workforce can analyze fast enough to act on.

This is the problem Acoer has spent years solving in production. We do not submit this response as a research proposal or a theoretical framework. We submit it as practitioners who have deployed AI systems inside medical examiner offices, forensic agencies, and public health departments. These are systems that today convert months of manual cross-case review into near-real-time intelligence, detect emerging drug combinations before they appear in structured reporting systems, and generate cryptographically verifiable proof of every data action taken.

AI will transform clinical care and public health when three conditions are met: (1) when trust is engineered into systems through verifiable governance rather than asserted through policy language; (2) when reimbursement and funding incentives reward demonstrated prevention outcomes rather than technology adoption theater; and (3) when interoperability frameworks evolve to connect the upstream intelligence of mortality and forensic data with the downstream response of clinical care.

Secretary Kennedy's vision to Make America Healthy Again demands action on preventable death. AI is the most powerful tool available to accelerate that mission. This submission provides concrete, implementation-informed recommendations aligned to HHS's regulatory, reimbursement, and R&D authorities, grounded in Acoer's direct experience deploying AI-driven prevention intelligence at scale.

The Strategic Imperative: AI for Preventable Death Reduction

The United States overdose crisis represents one of the most significant preventable death challenges in modern American history. In the past decade, more Americans have died from drug overdoses than were killed in all of World War II. Fentanyl and synthetic opioids (increasingly adulterated with xylazine, nitazenes, and novel synthetic compounds) evolve faster than any structured surveillance system can track. The window between a new substance combination appearing on the street and its identification in public health reporting can be months. People die in that window.

AI can close that window. But only if we design the right data infrastructure, governance frameworks, and incentive structures to make it possible.

Acoer's work demonstrates the model. Medical examiner/coroner offices generate the most outcome-validated health intelligence in the American health system: confirmed causes of death, toxicological findings, scene-level context, and demographic data. However, that intelligence is locked in unstructured, siloed files and systems. Our private AI systems extract, structure, and cross-reference that intelligence in near-real-time, surfacing early warning signals that can inform Emergency Department protocols, harm reduction resource deployment, and public health prevention strategy.

Current Gap	Private AI-Enabled Solution
• Substance combinations evolve faster than ICD codes capture them • Scene-level risk context buried in narrative case files • Weeks-to-months lag between death events and actionable intelligence • Fragmented data across medical examiner/coroner offices, emergency departments, and public health agencies	• Natural language extraction from unstructured autopsy and tox narratives • Cross-case pattern detection surfacing emerging drug combinations • Near-real-time syndromic surveillance intelligence • Open standards interoperability with clinical care workflows

This is not a theoretical future. Acoer's Forensic Intelligence Platform, including the Smart Toxicology module deployed in production with medical examiner offices, already performs structured extraction from unstructured reports, cross-case substance combination tracking, and early identification of emerging overdose patterns. The technology works. What is needed now is federal policy that scales it.

I. Regulation

A. Risk-Based Governance Clarity to Accelerate Responsible Adoption

Question 1: *What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?*

Acoer's direct experience deploying the Health Data Explorer and Forensic Intelligence Platform in public health and clinical-adjacent environments identifies five foundational barriers:

- **Regulatory ambiguity for non-medical device AI:** Most AI deployed in public health and prevention settings does not meet the definition of a medical device, yet governance expectations are unclear. This creates risk aversion among CIOs, compliance officers, and agency legal teams who cannot confidently deploy systems without defined accountability frameworks.
- **Unverifiable AI lineage:** AI systems in clinical and public health settings currently cannot demonstrate provenance of training data, consent status, model versioning history, or post-deployment performance drift. This makes institutional trust impossible to engineer and audit.
- **Data sovereignty concerns in sensitive environments:** Medical examiner offices, behavioral health agencies, and county health departments are deeply cautious about any AI architecture requiring centralized data pooling or external model training. The inability to guarantee local data control is a primary adoption barrier in the public sector.
- **Fragmented interoperability:** Forensic, clinical, and public health systems operate in separate data silos. AI-generated intelligence from mortality data cannot flow into emergency department workflows or syndromic surveillance platforms without standardized, trust-anchored data exchange.
- **Misaligned procurement pathways:** Innovation procurement pathways for public health agencies and clinical-adjacent public sector organizations are far less developed than those for hospital systems, creating structural disadvantages for responsible AI deployment in the jurisdictions that need it most.

Question 2: *What regulatory, payment policy, or programmatic design changes should HHS prioritize?*

HHS should prioritize the following specific regulatory changes:

1. Define a Minimum Verifiable AI Baseline Under 45 CFR Part 170

ASTP/ONC should amend 45 CFR Part 170 (Health IT Certification) to require AI systems used in certified health IT to meet a minimum verifiable governance standard, including:

- Versioned model metadata with tamper-evident storage
- Dataset lineage and provenance disclosures tied to consent status
- Cryptographically signed audit logs for inference events

- Documented post-deployment performance monitoring protocols

Acoer's Health Data Explorer and Forensic Intelligence platforms implement these capabilities today using the Hedera distributed ledger network to create immutable, cryptographically verifiable records of every data access event, AI inference, and governance action. These capabilities should become the baseline expectation for health AI systems, not exceptional features of forward-leaning vendors.

2. Principles-Based Guidance for Non-Medical Device AI Under 45 CFR Part 164

HHS should issue formal guidance clarifying governance expectations for AI systems that operate in clinical and public health settings but do not meet FDA medical device definitions. This guidance should:

- Define proportionate oversight requirements based on risk level
- Establish auditability and documentation standards
- Clarify HIPAA applicability to AI inference processes (45 CFR Part 164)
- Define human-in-the-loop requirements for population-level decision support

3. Support Privacy-Preserving AI Architecture Standards

HHS should formally recognize and incentivize AI architectures that:

- Operate within dedicated, jurisdiction-specific infrastructure without cross-jurisdiction data pooling
- Maintain full local data sovereignty and explicit consent tracking
- Provide transparent, auditable extraction and inference logic
- Support third-party validation of model governance claims

Question 3: *What novel legal and implementation issues exist and what role should HHS play?*

Three legal and accountability gaps require immediate HHS attention:

- **Liability ambiguity for AI-assisted prevention decisions:** When a clinical system or public health agency acts on AI-generated pattern intelligence (e.g., an emerging drug combination identified from mortality data) and that action prevents or fails to prevent harm, existing frameworks under 42 CFR Part 482 do not clearly allocate responsibility between the AI system, its operator, and the clinician or public health official who acted. HHS should convene a multi-stakeholder working group to define accountability frameworks for AI-assisted prevention decisions
- **Cross-jurisdiction data sovereignty in multi-agency intelligence sharing:** AI systems such as Acoer's Forensic Intelligence Platform that aggregate signals across medical examiner jurisdictions, EDs, and public health departments require explicit frameworks for who owns derived intelligence, who can act on it, and how consent flows across entities. HHS should model these frameworks through TECCA governance extensions.

- **Absence of AI incident reporting analogous to adverse event reporting:** There is currently no standardized mechanism to report AI system failures in clinical or public health settings. HHS should establish a voluntary (and ultimately mandatory) AI incident reporting framework capturing: model version, data context, workflow integration point, and human-AI interaction pattern. This would materially improve safety learning without chilling innovation.

B. Supporting Private Sector Innovation

Question 5: *How can HHS best support private sector activities to promote innovative and effective AI use in clinical care?*

HHS should establish a tiered, risk-proportionate AI assurance framework that:

- Creates a voluntary certification pathway for health AI vendors demonstrating verifiable governance practices
- Establishes public-private partnerships to develop open benchmarking datasets for mortality intelligence, syndromic surveillance, and overdose pattern detection
- Funds cooperative agreements with health AI vendors operating in underserved public health environments (county health departments, tribal health authorities, medical examiner offices) to accelerate deployment of responsible private AI platforms in the jurisdictions with the greatest preventable death burden
- Establishes safe harbor provisions for organizations that implement verifiable AI governance standards, reducing liability uncertainty that currently deters adoption

Question 7: *Which roles have the most influence on AI adoption and what are the primary administrative hurdles?*

In public health and clinical-adjacent settings, Acoer's experience identifies the following as primary adoption gatekeepers:

- Chief Medical Examiners and Forensic Directors: control access to the most valuable outcome-validated health data in the system, but have minimal IT procurement support
- State and County Public Health Officers: primary decision-makers for prevention strategy, but operate under constrained budgets and procurement processes designed for traditional IT, not AI
- Healthcare System CIOs and CMIOs: strong influence but face compliance teams that cannot approve systems without verifiable governance documentation
- State Medicaid Directors: increasingly relevant as prevention AI demonstrates downstream cost avoidance

The primary administrative hurdles are: (1) absence of standard AI governance requirements that give legal teams a basis for approval; (2) HIPAA ambiguity for AI inference on public health data; (3) procurement vehicle constraints that favor established vendors over innovative small companies; and (4) the requirement that AI tools demonstrate ROI in reimbursement terms that do not yet capture prevention value.

II. Reimbursement

A. Aligning Payment Policy to Reward Prevention Outcomes

Question 2: *What payment policy changes should HHS prioritize?*

Under current fee-for-service frameworks governed by 42 CFR Parts 410 and 414, AI that prevents overdose deaths, detects emerging drug threats, or reduces unnecessary Emergency Department utilization generates no direct reimbursement signal. This is a fundamental misalignment: the most valuable AI applications in public health (specifically those that prevent costs and deaths before they occur) are precisely the ones that receive no payment reward.

Acoer recommends the following specific payment policy changes:

1. Fund Prevention AI Demonstration Pilots through CMMI

CMS should authorize CMMI to fund 3-5 year demonstration projects measuring the impact of AI-enabled mortality intelligence and syndromic surveillance on downstream healthcare costs. Eligible outcomes for reimbursement credits should include:

- Reductions in overdose-related Emergency Department visits
- Faster deployment of harm reduction resources (naloxone distribution, treatment referrals) in response to AI-identified emerging substance threats
- Reduction in time-to-detection for emerging drug combinations (from months to days)
- Measurable reduction in overdose mortality in geographies where AI-enabled prevention intelligence is deployed

The Economic Argument for Prevention AI

A single prevented opioid overdose death represents an estimated \$1-2 million in avoided lifetime productivity loss, healthcare costs, and downstream criminal justice expenditure. AI systems that cost hundreds of thousands of dollars annually to operate can prevent hundreds of deaths. The return on investment is not marginal. It is transformational. Existing payment models simply do not capture this value.

2. Create AI Outcome Modifiers in Value-Based Care Arrangements

For Accountable Care Organizations, Medicare Advantage plans, and Medicaid managed care entities, CMS should establish AI-specific outcome modifiers that provide financial credit for documented AI-assisted prevention outcomes. Critically, all outcome claims must be auditable, not self-attested. This is where Acoer's cryptographic governance infrastructure provides unique value: our systems generate tamper-evident proof of what AI actions were taken, what data was used, and what outcomes were observed, generating the kind of independently verifiable outcome attribution that outcome modifier programs require.

3. Include Forensic and Mortality Intelligence in Population Health Payment Models

HHS should revise the definition of "eligible data sources" in population health payment models to include structured mortality and forensic intelligence, enabling providers and health plans to receive credit for incorporating upstream prevention signals into clinical workflows.

Question 6: *Where have AI tools met or exceeded expectations, and what novel tools have the greatest potential?*

Acoer's direct deployment experience provides the following assessment:

Where AI Has Demonstrated Clear Value

- Natural language processing for unstructured forensic and clinical narratives: Acoer's Forensic Intelligence Platform demonstrates that private AI consistently outperforms manual review in speed, consistency, and cross-case correlation, converting months of work into hours without sacrificing accuracy.
- Cross-case pattern detection in overdose intelligence: AI systems that aggregate and cross-reference individual case data identify emerging substance trends 4-8 weeks earlier than traditional structured reporting systems.
- Administrative burden reduction: AI-driven pre-population of reporting fields, automated coding assistance, and documentation support demonstrably reduce administrative time for forensic and clinical staff.

Where AI Has Fallen Short or Faces Challenges

- Trust deficit in clinical adoption: AI tools deployed without verifiable governance documentation face persistent resistance from clinical staff, compliance teams, and legal counsel, regardless of technical performance.
- Integration complexity: AI tools that require deep EHR integration face months-to-years implementation timelines that erode their value in rapidly evolving public health contexts.
- Lack of outcome feedback loops: Most deployed AI systems do not have mechanisms to measure whether their outputs led to different decisions or better outcomes, making it impossible to demonstrate or improve value over time.

Novel Tools with Greatest Potential to Reduce Costs and Save Lives

- AI-powered syndromic surveillance integration: Systems that aggregate signals from mortality data, ED visits, EMS runs, and harm reduction encounters to provide a unified, near-real-time early warning intelligence layer for emerging public health threats.
- Predictive overdose risk scoring at the individual patient level: AI models trained on mortality intelligence ("what led to death in confirmed cases") that can identify high-risk individuals in ED and primary care settings for targeted intervention.
- Substance combination early detection: AI systems that identify novel polysubstance combinations in toxicology data before they are represented in ICD coding systems, enabling faster public health response.

III. Research & Development

A. Recognizing Mortality Intelligence as Prevention Infrastructure

Question 10: *What specific areas of AI research should HHS prioritize to accelerate AI adoption in clinical care?*

Mortality data is the most reliable, outcome-validated intelligence in the American health system. Unlike clinical diagnostic data, mortality data reflects confirmed outcomes. Unlike survey data, it captures what actually happened, not what was reported. Yet it is systematically excluded from AI-enabled prevention strategy. HHS should treat structured mortality intelligence as a foundational layer of national health AI infrastructure.

HHS R&D investment priorities should include:

- **Natural language processing for forensic and clinical narratives:** Investment in private AI and NLP models specifically trained and validated on autopsy narratives, toxicology reports, and death investigation documentation. These differ substantially from clinical encounter notes and require domain-specific training, validation, and evaluation frameworks. Acoer's Forensic Intelligence Platform represents a deployable reference implementation for this class of system.
- **Cross-case pattern detection infrastructure:** Development of federated AI architectures that enable cross-jurisdiction pattern detection without requiring centralized data pooling, thereby preserving data sovereignty while enabling population-level intelligence.
- **Syndromic surveillance AI integration:** Research and implementation science for integrating AI-derived mortality intelligence into existing CDC BioSense Platform syndromic surveillance workflows, enabling real-time early warning signals to inform public health response.
- **Predictive prevention AI for high-risk populations:** Development and validation of AI models that use mortality intelligence to identify and intervene with living individuals at highest risk of overdose death, before the next acute event.
- **Verifiable AI governance standards:** Investment in standards development for cryptographic AI audit trails, model versioning, and inference provenance documentation, to enable interoperable, independently verifiable AI governance across health system boundaries.

Regarding published evidence (Q10a): Acoer's deployment experience is consistent with emerging literature demonstrating that AI-assisted NLP of clinical and forensic narratives achieves extraction accuracy comparable to expert human review at dramatically higher speed. The gap between the evidence base and current deployment is not evidence quality. The real constraint is governance framework maturity and interoperability infrastructure.

Regarding cost-benefit analysis (Q10b): Published literature consistently finds that AI systems in high-volume, pattern-intensive health contexts (radiology, pathology, population surveillance) generate positive ROI when total cost of care effects are captured over multi-year horizons. The primary gap in current cost-benefit models is the failure to capture upstream prevention value:

specifically the deaths and hospitalizations that did not occur because early warning intelligence triggered an intervention.

B. Interoperability as Prevention Infrastructure

Question 8: *Where would enhanced interoperability widen market opportunities, fuel research, and accelerate AI for clinical care?*

Acoer is an active implementer of FHIR-native interoperability and a strong supporter of HTI-5 policy direction. Our CEO, Jim Nasr, serves in a leadership role on the Coalition for Decentralized Healthcare and Life Sciences, an industry working group advancing open, verifiable, community-governed research models, reflecting Acoer's commitment to the broader ecosystem of transparent health data governance. From our deployment experience, interoperability gaps most limiting to AI-enabled prevention include:

Priority Interoperability Gaps

Data Type	Gap and Recommendation
Overdose toxicology findings	No standardized FHIR representation for multi-substance toxicology findings. HHS should fund development of FHIR extensions for overdose substance combinations, enabling mortality intelligence to flow into ED and population health workflows.
ME case management systems	ME case management systems are largely excluded from health IT certification and TEFCA participation. HHS should extend TEFCA participation eligibility to medicolegal death investigation systems and establish FHIR-based data exchange standards for mortality case data.
AI model metadata	No standardized schema for AI model provenance metadata in FHIR. HHS should fund development of FHIR extensions for AI model version, training data provenance, and inference audit trails, to enable interoperable AI governance documentation across system boundaries.
Syndromic surveillance feeds	AI-derived mortality signals cannot currently flow into CDC BioSense Platform in structured, machine-readable format. HHS should define standardized APIs for mortality-derived early warning signals to integrate with national syndromic surveillance infrastructure.
Computable consent	AI systems that use patient data across jurisdictions require machine-readable consent provenance aligned with 45 CFR Part 164. HHS should fund development of computable consent standards that travel with data through FHIR exchanges, enabling AI systems to verify consent status at inference time.

C. Public-Private Partnerships and Demonstration Pilots

Question 4: *What are the most promising AI evaluation methods and what mechanisms would be most impactful?*

Acoer recommends HHS fund demonstration pilots through cooperative agreements and CRADAs with a dual evaluation framework:

Pre-Deployment Evaluation Standards

- Structured extraction accuracy testing validated against confirmed toxicological and forensic findings (ground truth from completed death investigations)
- Expert domain validation by licensed forensic pathologists and toxicologists
- Reproducibility, consistency, and population-representativeness testing across demographic and geographic subgroups
- Cryptographic baseline documentation of model version, training dataset, and validation methodology before deployment

Post-Deployment Evaluation Framework

- Reduction in time-to-detection for emerging substance combinations (measured against structured reporting lag)
- Accuracy of early warning signals validated against subsequent confirmed death data
- Clinical workflow integration effectiveness: did AI-derived intelligence actually change ED protocols, harm reduction deployment, or prescribing patterns?
- Downstream outcome measurement: overdose ED visits, EMS naloxone deployments, and mortality in jurisdictions with and without AI-enabled surveillance

The most impactful mechanisms for HHS investment would be: cooperative agreements with state and county public health departments for deployment of AI-enabled mortality intelligence platforms; CRADAs with technology vendors (like Acoer) to co-develop evaluation frameworks and open benchmarking datasets; and prize competitions for demonstrated prevention outcomes rather than technology features.

D. Leadership in Decentralized Science and Open Health Research

Acoer's commitment to verifiable, transparent, community-grounded research extends beyond commercial product deployment. Acoer CEO Jim Nasr serves in a leadership role on the Coalition for Decentralized Healthcare and Life Sciences, an industry working group advancing the principles of decentralized science (DeSci), an emerging movement that applies blockchain-based governance, open data standards, and token-enabled coordination to scientific research, peer review, and clinical validation.

DeSci is directly relevant to HHS's AI strategy because it provides an operational model for the open, verifiable, community-accountable research infrastructure that AI governance requires. Key principles:

- **Open provenance:** Research data, methods, and findings are published with cryptographically verifiable provenance, so any stakeholder (whether a regulator, patient, or peer researcher) can independently verify what was done and what was found.
- **Community-driven validation:** Citizen science and open peer review platforms enable distributed validation of research findings, broadening the evidence base beyond traditional academic and industry gatekeepers.
- **Data sovereignty:** DeSci frameworks ensure research participants retain ownership of their data and can verify how it is used in research protocols, a critical requirement for sensitive health data.
- **Transparent incentive alignment:** Token-based coordination mechanisms can align researcher, patient, and public health incentives around measurable health outcomes rather than publication metrics or commercial returns.

DeSci in Action: Open Cancer Research on ResearchHub

Acoer is an active contributor to an open science project on ResearchHub (researchhub.com/fund/4459) investigating the Wim Hof Method (WHM), a structured cold exposure and breathwork practice, as a novel complementary approach for cancer patients. This citizen-scientist pilot applies the principles of open, verifiable, community-funded research to a clinical question that traditional research funding mechanisms have been slow to address. It demonstrates that rigorous, outcome-driven health research can be conducted transparently, with public accountability, at a fraction of the cost of conventional clinical trials.

HHS Implications: DeSci as a Model for AI Research Governance

HHS should consider how DeSci principles can inform the governance of AI research in clinical and public health settings:

- **Open benchmarking datasets:** HHS-funded AI benchmarking datasets for overdose pattern detection, syndromic surveillance, and clinical decision support should be published with DeSci-style open provenance, enabling independent validation and continuous improvement by the research community.

- **Community validation of AI models:** Before AI models are deployed in clinical settings, DeSci-inspired community validation processes (involving clinicians, patients, and independent researchers) could provide a more robust and trusted validation pathway than proprietary testing alone.
- **Open science funding mechanisms:** HHS should explore whether platforms like ResearchHub and other DeSci infrastructure can serve as channels for distributing small, milestone-based research grants for prevention-focused AI validation studies, reducing the administrative burden of traditional grant mechanisms while broadening participation.
- **Citizen science for health surveillance:** DeSci frameworks that enable patient communities to contribute their own data to surveillance and research efforts, with full transparency and consent controls, could significantly expand the data available for AI-driven prevention research.

IV. Additional Question Responses

Patient and Caregiver Perspectives

Question 9: *What challenges do patients and caregivers wish to see addressed by AI in clinical care, and what concerns do they have?*

Acoer's work with families affected by overdose deaths provides a distinctive perspective on patient and caregiver priorities:

What Patients and Caregivers Want AI to Deliver

- **Prevention before death:** Families who have lost loved ones to overdose consistently ask why the warning signs were not detected earlier. They want AI systems that identify high-risk individuals in clinical encounters and trigger earlier, more aggressive intervention.
- **Transparency about what happened:** Families navigating the aftermath of an overdose death want clear, comprehensive information about what substance combinations were involved, what medical interventions were attempted, and what systemic factors contributed. AI that can answer these questions from existing data provides immense value.
- **Fewer administrative barriers to care:** Patients with substance use disorders consistently encounter administrative and insurance barriers to treatment access. AI that can automate prior authorization, reduce treatment matching latency, and predict which interventions are most likely to be effective for a given individual can meaningfully reduce these barriers.
- **Privacy and control:** Patients consistently express concern about AI systems that aggregate their health data without their knowledge or consent, and about the use of sensitive behavioral health information in employment, insurance, or law enforcement contexts.

What Concerns Patients and Caregivers Have About AI

- **Invisible automation replacing human judgment:** Patients and families fear that AI-driven clinical decisions will reduce the quality of human attention they receive from clinicians.

- Data misuse: Particularly in behavioral health and substance use contexts, patients fear that AI-accessible data will be used against them in ways they did not anticipate or consent to.
- Algorithmic bias: Community members in high-overdose-burden communities, which are disproportionately communities of color, are acutely concerned about AI systems that could perpetuate or amplify existing inequities in how care is provided and resources are allocated.

Addressing these concerns requires governance systems where consent is explicit, inspectable, and enforceable, not buried in terms of service. Acoer's computable consent and cryptographic audit capabilities provide a model for how AI governance can be made visible and verifiable to patients, not just institutions.

V. Summary of Specific Recommendations

Lever	Recommendation	Authority / CFR
Regulation	Amend 45 CFR Part 170 to require verifiable AI governance baseline (model provenance, cryptographic audit logs, inference documentation) for certified health IT	45 CFR Part 170 (ONC Health IT Certification)
Regulation	Issue principles-based guidance for non-medical device AI in clinical and public health settings with risk-proportionate oversight standards	45 CFR Part 164; 42 CFR Part 482
Regulation	Establish standardized AI incident reporting framework analogous to MedWatch adverse event reporting	21 CFR Part 803; new guidance authority
Regulation	Extend TECCA participation eligibility to ME/coroner case management systems	45 CFR Part 170; TECCA governing framework
Reimbursement	Fund CMMI demonstration pilots measuring AI-enabled prevention outcomes: overdose deaths prevented, ED utilization reduced	42 CFR Part 425; Section 1115A authority
Reimbursement	Establish AI outcome modifiers in value-based care arrangements with cryptographically verifiable attribution requirements	42 CFR Parts 410, 414, 425
Reimbursement	Include structured mortality intelligence as eligible data source in population health payment models	42 CFR Part 425; MSSP program rules
R&D	Fund cooperative agreements for public-private deployment of AI-enabled mortality intelligence in high-burden jurisdictions	42 U.S.C. § 300hh; PHSA Section 301
R&D	Fund FHIR extension development for overdose toxicology, AI model metadata, and computable consent provenance	45 CFR Part 170; ONC Standards Advisory
R&D	Develop open benchmarking datasets for overdose pattern detection and syndromic surveillance AI validation	NIST AI RMF alignment; CMMI research authority

R&D	Integrate AI-derived mortality signals into CDC BioSense Platform for real-time national syndromic surveillance	<i>Public Health Service Act; CDC cooperative agreement authority</i>
-----------------------	--	--

Conclusion

AI will not save lives until it is trusted, funded, and connected to the systems where lives are at risk. The United States loses more than 70,000 people to drug overdoses every year. A significant proportion of those deaths are preceded by signals, seen in clinical encounters, in emergency runs, in the unstructured text of toxicology reports, that AI can detect, correlate, and act on, if the governance, interoperability, and incentive infrastructure exists to support it.

Acoer stands at the intersection of this challenge and this opportunity. Our Health Data Explorer and Forensic Intelligence Platform are deployed today in medical examiner offices and public health agencies, converting unstructured mortality data into actionable prevention intelligence using private AI that keeps sensitive data under local jurisdiction control. We use cryptographic technologies to make every AI action verifiable and every data governance decision auditable. We are aligned with FHIR and HTI-5 interoperability standards. We have built the privacy-preserving architecture that sensitive public health environments require.

What we need from HHS is the governance clarity, the interoperability infrastructure, and the reimbursement alignment to scale what works. The technology exists. The evidence is accumulating. The political will to Make America Healthy Again is present. The moment to act is now.

We recommend HHS prioritize three immediate actions:

- **Establish a verifiable AI baseline:** Amend 45 CFR Part 170 to require cryptographically verifiable AI governance documentation as a condition of health IT certification. Make trust in AI systems a requirement, not an aspiration.
- **Fund prevention AI demonstration pilots:** Use CMMI authority and cooperative agreement mechanisms to fund 3-5 year demonstrations of private AI-enabled mortality intelligence and syndromic surveillance, with clear outcome metrics and attribution standards, prioritizing architectures that preserve data sovereignty.
- **Extend interoperability to the full prevention data ecosystem:** Bring ME/coroner systems into TEFCA, fund FHIR extensions for mortality and AI metadata, and create APIs that allow AI-derived prevention intelligence to flow into national syndromic surveillance infrastructure.

Acoer welcomes the opportunity to demonstrate our deployed systems, share our evaluation data, and partner with HHS in designing the governance frameworks and pilot programs that will accelerate responsible AI adoption across clinical and public health settings. Every day that AI-enabled prevention intelligence remains siloed, fragmented, or insufficiently trusted to act on is a day when preventable deaths continue.

We appreciate the opportunity to contribute to this policy dialogue and stand ready to support HHS in turning this RFI's insights into action.

Respectfully submitted,

Jim Nasr

Chief Executive Officer, Acoer, Inc.
jnasr@acoer.com

Acoer, Inc.

Atlanta, Georgia
Acoer.com
HealthDataExplorer.io

February 23, 2026