

**Department of Health and Human Services (HHS)
Office of the Deputy Secretary and Assistant Secretary for Technology
Policy and Office of the National Coordinator for Health Information
Technology (ASTP/ONC)**

**Accelerating the Adoption and Use of
Artificial Intelligence as Part of Clinical Care
RFI Document #2025-23641**

Submitted to: Steven Posnack, via Regulations.gov
Date of Submission: February 23, 2026

Prepared by:



Mettler
SOLUTIONS

Mettler Solutions, LLC
An SBA Certified 8(a) Business



POC: Kishore Metla (Chief Executive Officer)

Phone: (443) 545-3852 | **Fax:** (443) 545-3401 | **Email:** kishore.metla@mettles.com
10320 Little Patuxent Pkwy, Suite 233-02, Columbia, MD 21044 | **Website:** www.mettles.com
SAM UEI: CURWPHLP9PK9 | **DUNS:** 079346709 | **FEIN:** 27-4746158 | **Cage Code:** 7NKX3
GSA MAS: 47QTC A21D004Y | **GSA STARS III:** 47QTC B21D0023

Table of Contents

1	Executive Summary & Perspective	1
2	Responses to Specific Questions	2
2.1	Question 1: Barriers to Private Sector Innovation	2
2.2	Question 2: Regulatory and Payment Policy Levers	2
2.3	Question 3: Novel Legal and Implementation Issues	2
2.4	Question 4: Evaluation Methods, Metrics, and Robustness	3
2.5	Question 5: Accreditation and Certification	3
2.6	Question 6: Performance and Limitations	4
2.7	Question 7: Governance and Decision Making	4
2.8	Question 8: Interoperability	4
2.9	Question 9: Patient and Caregiver Priorities and Concerns	4
2.10	Question 10: Research Priorities and Evidence	4
3	Conclusion	5

1 Executive Summary & Perspective

Mettle Solutions respectfully submits this response to the Department of Health and Human Services (HHS). We respond not as an AI developer ("Builder"), but as a specialized **Evaluator** and **Builder of Assurance Tools** that support the safe adoption of AI in clinical care. Our perspective is informed by working across diverse care settings, from large academic health systems to resource-constrained community practices, with AI tools that are often classified as non-medical device technologies.

Regulatory Impact on AI Adoption: The Validation Gap

Current HHS regulatory frameworks appropriately emphasize patient safety, privacy, and public trust. However, for many clinical AI tools, particularly those that fall outside traditional medical device classifications—regulatory expectations for **post-deployment validation and ongoing performance monitoring remain unclear or implicit**. This lack of clarity contributes to what we describe as a "Validation Gap": uncertainty among providers regarding how AI models perform on their **local patient populations over time**.

In practice, validation is often retrospective, static, and disconnected from real-world clinical workflows. Existing regulations do not consistently specify how health care organizations should generate local evidence, monitor model drift, or demonstrate continued safety and effectiveness once an AI tool is deployed. As a result, providers face two undesirable outcomes:

- 1) **blind adoption**, which introduces patient safety and liability risks, or
- 2) **adoption paralysis**, where otherwise valuable AI tools are delayed or rejected due to regulatory ambiguity and unclear accountability.

This uncertainty disproportionately affects smaller and rural practices, which often lack the resources to interpret evolving regulatory expectations or to independently design robust validation processes.

An Independent Safety Layer as a Regulatory Enabler

To support a regulatory posture that is predictable, proportionate to risk, and enabling of innovation, we advocate for an **Independent Safety Layer** approach that distinguishes between building AI and assuring AI. This approach aligns with HHS's goals of protecting patients and maintaining public trust while allowing innovation to proceed more rapidly.

Our response outlines two implementation pathways to promote equitable adoption:

- **Enterprise Track:** "Shadow mode" validation prior to clinical use, combined with continuous post-deployment surveillance ("sidecar" monitoring) to ensure models perform safely across changing populations and conditions.
- **Resource-Constrained Track:** Mettle Solutions is currently developing a Local AI Evaluator (under a contract with ASTP) which is a lightweight, no server backend application that enables resource constrained healthcare providers to evaluate the accuracy metrics of AI tools.

2 Responses to Specific Questions

2.1 Question 1: Barriers to Private Sector Innovation

What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

The single greatest impediment is **uncertainty**. Healthcare organizations operate in a zero-failure environment. Providers cannot confidently adopt AI because they lack reliable tools to answer the fundamental question: *"Will this model work for my patients, in my workflow, today?"*

Key observed barriers include:

- **The "Local Validity" Gap:** A model validated in an urban academic center often degrades when deployed in a rural clinic due to differences in demographics and workflow.
- **Undefined Liability:** Providers fear that relying on AI advice creates new liability exposure, particularly when models "hallucinate" or drift.
- **The "Resource Gap":** While large systems can hire data science teams to evaluate tools, small practices cannot. This process risks creating a "two-tier" health system where AI benefits are restricted to well-resourced centers.

2.2 Question 2: Regulatory and Payment Policy Levers

What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.

We are open to collaborating with HHS to determine the best solution.

2.3 Question 3: Novel Legal and Implementation Issues

For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (e.g., relating to liability, indemnification, privacy, and security). What novel legal and implementation issues exist and what role, if any, should HHS play to help address them?

1) Liability & Standard-of-Care Ambiguity

When "non-medical device" AI errors, it is unclear whether the fault lies with the clinician (malpractice) or the vendor (product defect). This ambiguity stalls procurement and adoption.

Proposed Solution: "Gated Decision" Architecture

We recommend HHS endorse "Gated Decision" Architectures (or Safety Interlocks) as a standard governance control. In this framework, an AI prediction does not reach the clinician unless it passes a real-time safety check (e.g., Is the patient already on treatment? Is the confidence score above threshold?).

- **HHS Role:** Recognize that providers utilizing "Gated" architectures are exercising due diligence. This shifts reliance from "gut feeling" to auditable safety protocols.

2) Indemnification & Cryptographic Evidence

Indemnification is difficult because post-event forensic analysis is often impossible. Parties cannot prove what the AI "saw" or "thought" at the moment of care.

- **HHS Role:** Encourage the adoption of **Cryptographic Evidence Chains**. By digitally signing every validation event (input, prediction, safety check, user action), providers and vendors create an immutable "Audit Package" that can factually resolve disputes and support indemnification clauses.

2.4 Question 4: Evaluation Methods, Metrics, and Robustness

For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and other workflow and human-centered evaluation methods for clinical care? Should HHS further support these processes? If so, which mechanisms would be most impactful (e.g., contracts, grants, cooperative agreements, and/or prize competitions)?

Validation is not a one-time event; it is a continuous lifecycle. We propose the following standard methods:

A. Pre-Deployment: "Shadow Mode" Validation

The most effective evaluation method is Shadow Mode, where the AI runs silently alongside clinical workflows. It captures real-time data and generates predictions without alerting the clinician.

B. Post-Deployment: The "Auditor Sidecar"

To address "robustness testing", An "**Auditor Sidecar**" sits parallel to the AI, continuously monitoring for "Drift" (statistical changes in input data) and "Degradation" (drops in confidence or accuracy).

HHS Support Mechanism:

HHS should use Contracts and Cooperative Agreements to fund "Clinical AI Innovation Sandboxes." These programs would subsidize the cost of third-party validation (Shadow Mode infrastructure) for participating health systems, allowing them to test AI tools risk-free before purchase.

2.5 Question 5: Accreditation and Certification

How can HHS best support private sector activities (e.g., accreditation, certification, industry-driven testing, and credentialing) to promote innovative and effective AI use in clinical care?

Certification must be based on **local**, ongoing evidence, not just static vendor claims. The entity that evaluates AI models in production cannot be the same as the developers of those models.

Recommendation: HHS should support an "AI Assurance Evidence Package" standard.

Rather than a one-time seal of approval, certification should require a dynamic Evidence Package containing:

- 1) **Local Shadow Mode Results:** How the model performed on *this* site's data.
- 2) **Monitoring Plan:** Defined thresholds for drift and safety signal detection.
- 3) **Incident Response Protocol:** Pre-agreed pathways for when the model fails.

2.6 Question 6: Performance and Limitations

Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?

We are open to collaborating with HHS to determine the best solution.

2.7 Question 7: Governance and Decision Making

Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?

We are open to collaborating with HHS to determine the best solution.

2.8 Question 8: Interoperability

Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

Interoperability is usually focused on patient data (USCDI). To accelerate AI adoption, we need **Interoperability of Evaluation Evidence.**

The Opportunity:

If HHS promotes a standard for "AI Evaluation Event Logs", then independent Evaluators can audit any model from any vendor.

- **Data Standards Needed:** A schema that standardizes:
 - *Input Context* (what the model saw)
 - *Prediction & Confidence* (what the model said)
 - *Outcome Linkage* (what actually happened to the patient)

2.9 Question 9: Patient and Caregiver Priorities and Concerns

What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?

We are open to collaborating with HHS to determine the best solution.

2.10 Question 10: Research Priorities and Evidence

Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care?

- a. Are there published findings about the impact of adopted AI tools and their use clinical care?*
- b. How does literature approach the costs, benefits, and transfers of using AI as part of clinical care?*

We are open to collaborating with HHS to determine the best solution.

3 Conclusion

The barriers to effective clinical AI adoption are not primarily technical; they are rooted in uncertainty. Health care organizations lack clear, scalable mechanisms to verify how AI tools perform in their specific clinical environments over time. This uncertainty affects regulatory confidence, liability management, procurement decisions, and ultimately patient trust.

Throughout this response, we have described how an ecosystem of **independent, post-deployment validation and continuous assurance** can address these challenges in a manner that is predictable, proportionate to risk, and equitable across care settings. Approaches such as enterprise-scale Shadow Mode validation, standardized evaluation evidence, and auditable safety controls offer HHS practical levers to reduce adoption friction while maintaining strong patient protections.

By clarifying expectations for ongoing performance verification, supporting independent assurance mechanisms, and enabling interoperable evaluation evidence, HHS can help shift the clinical AI market from reliance on static claims to reliance on real-world performance data. This shift would strengthen public trust, support responsible innovation, and ensure that the benefits of AI in clinical care are accessible to providers and patients across the U.S. health system.