

February 23, 2026

Thomas Keane, MD
Assistant Secretary for Technology Policy and Director, Office of the National
Coordinator for Health Information Technology,
Mary E. Switzer Building
Mail Stop: 7033A
330 C Street SW, Washington, DC 20201

Attention: HHS Health Sector AI RFI

Dear Dr. Keane,

Below are my comments, which I provide in response to the HHS Health Sector AI.

Thank you for the opportunity to provide input into the ASTP's thinking and decision making about shaping our approach to incorporating AI into clinical care.

Sincerely,

Anna McCollister
Patient Advocate
Technology Company Founder and Consultant
Member, HITAC

8. Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

Make IAS Possible in the Real World

AI has remarkable potential and some nascent demonstrated ability to support patients and providers in providing care guidance and assistance, filling in major gaps in physician-patient communications opportunities, information about issues that are outside of the provider's specialties and personalizing insights based on an individual's data. One of the biggest hurdles in the development and deployment of apps that can

bridge the gap between and for patients and providers is the inability to access patient data through IAS on existing data exchange systems, such as Carequality and TEFCA. As we have known for years, attempts to get the data are more often than not thwarted by providers who claim they cannot confirm that the patient requesting the data is in fact the patient represented in the requested data. Out of fear of a potential accidental disclosure of personal health data, providers stop the data from being exchanged. As a result, the apps cannot access the data they need to help fuel the AI-generated support. Patients are left waiting. Many have chosen to manually download their data from their various portals and then upload the data into consumer LLMs, which have no protections for privacy. Meanwhile, data brokers use well developed “digital finger printing” algorithms to knit together data from individuals with remarkable accuracy and completeness with zero regulatory oversight or limitations.

HHS’ Office of Civil Rights has resisted requests to provide guidance that could help alleviate this issue. While OCR’s perspective is important - we absolutely must protect against data breaches - there are ways this could be addressed. OCR could and should develop a HITRUST-esque approach, whereby they define a gold standard for proof of identification. These algorithms have been developed and used in other contexts. They could work here, as well. Once the standards are in place, OCR should grant any company who demonstrates they met those standards with a legal “safe harbor”, protecting them against any enforcement actions and fines.

This kind of policy leadership is essential to make IAS real. Getting patients access to their data for use in AI tools can be an essential part of not just self care management, but also for clinical care. We cannot rely on physicians and provider systems to be the gatekeepers for healthcare-focused AI tools. Patients must be an equal party to your considerations. Enabling real patient data access and incorporation of their personal data into AI tools will enhance clinical care, promote primary and secondary prevention and improve disease outcomes.

Research Exchange Purpose

As a patient with complex chronic disease, it is essential that we enable and facilitate data liquidity for research purposes, allowing researchers to better understand disease patterns in the real-world setting. ASTP and TEFCA can play an important role in enhancing this by creating a Research Exchange Purpose through TEFCA.

One of the tensions that exists within TEFCA is the lack of an authorized exchange purpose for research. From the outset, we have known that companies and researchers would seek to access data through the exchange to enable secondary research of integrated, aggregated health data. As ASTP worked to create and build TEFCA, they rightfully prioritized the most essential exchange purposes. Research was not among them. To get the system off the ground, that was the right call. However, at this point in TEFCA's development, we need to create a legitimate mechanism through which authorized exchange can happen for research purposes. In so doing, we will proactively alleviate some of the pressure on QHINs to police the use of the treatment exchange purpose. At the same time, we need to be very clear about what research is in scope and what is not in scope, creating stiff penalties for companies that misuse the research purpose to harvest the data for monetization and other unintended uses. In establishing the research exchange purpose, we need to set requirements for those who use that exchange purpose, such as mandating patient consent and requiring transparency about the specific ways the data was used, publicly releasing a list of the ways the data was used, such as reports, summaries, publications or abstracts.

9. What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?

Improve Data Granularity

The clinical record for any patient is inherently incomplete. The EHR captures data taken during clinical encounters and laboratory testing. The data contained within the patient's record represent discrete elements related to health that are easily measured in a clinical setting.

By their nature, these clinical visits are episodic. They are snapshots in time, without context for what has happened and why since the last clinical encounter. Any clinician and patient knows the data are not a representation of the patient's full health status and the things that contribute to their ongoing health management or disease progression. For some diseases or status checks, that might be OK. But for chronic disease management, especially for data-intensive chronic disease, such as diabetes, hypertension, congestive heart failure and cardiovascular disease, among many others,

basing the patient's treatment decisions solely on the data contained in their EHR would be unwise. The doctor needs more context, more data.

The good news is, much of this data exists in machine-readable, standardized formats that could be incorporated into the EHR. Despite efforts over many years to encourage incorporation of this data into certified EHRs, thus far, ONC/ASTP has not made it a requirement.

In traditional clinical care, the gaps in data can be augmented by physician and patient discussions or record exchange. It's time consuming, clunky and not ideal, but it helps guide the physician in truly understanding the patient's status, concerns and needs. With this context, the physician can then review the data and make choices about treatment. But the context is absolutely essential for getting the best chance to pick the treatment and care course that will work for that patient.

As we move into the era of AI at full speed, I am greatly concerned about introducing AI tools into clinical decision making without requiring them to have context beyond the EHR. As a patient with type 1 diabetes, I have more than once overheard physicians making insulin dosing decisions in the ER that would have caused me significant harm if implemented. Fortunately, to date, I have been sufficiently conscious to be able to advise them that I am very insulin sensitive and correct their dosing strategy, which very likely saved my life. The physicians were basing their dosing decisions on well vetted algorithms developed by expert clinicians who relied on data from well designed studies and years of clinical care. And for me, those guidelines could have been deadly. The data about my insulin sensitivity factor is coded into my insulin pump through an app that doses me very precisely in response to changes in glucose.

Imagine I'm in the ER in two years. Decisions about medication and dosing that are well established in clinical guidelines have been incorporated into an AI-driven tool that automatically doses patients in diabetic ketoacidosis without supervision from a physician or consultation with a patient who may be unconscious. If the data from the patient's pump and CGM are missing, that algorithm is essentially blind dosing a lethal product. It has no idea what it doesn't know. Insulin can easily kill patients if too much is given.

Now imagine, my pump and my artificial pancreas app data are incorporated into the EHR or is readily accessible and can be incorporated into the ER's AI tool. In this case, there's a very good chance the AI would do a better job than the ER doctor who may know about diabetes but is not a specialized endocrinologist. In this instance, AI could save the patient's life.

The most important element in this hypothetical isn't the AI. That part is doable now. What's missing is the granular data that actually gives sufficient context, history and more sensitive measurements that will enable more accurate, timely and safe care. This is one narrow example of how AI could go drastically wrong or be heroically helpful I can cite many others from my own personal care and treatment regimens, and I am one person.

Patients live in a world of risk. We measure it and weigh it daily, sometimes hourly. When it is transparent and we are well enough, we can make decisions. When decisions are automated, the risks change. In some ways they may be mitigated, but if not done wisely, the risks to patients will be multiplied. This can be fixed in part by more granular data. If we do not address this issue, AI is likely to do more harm than good. We must enable patient-generated health data into the EHR in a structured, machine-readable manner.

AI in Payment Decisions

As mentioned above, the data within the clinical record fails to represent many critical factors that need to be considered when assessing a patient's health, understanding their disease and choosing a treatment that is appropriate. As patients and providers have experienced for many years, payment decisions made based on these data elements alone miss the context required to choose the treatment that works best for patients. Insurers tend to focus on very specific clinical measures represented in the patient's data. These measures are often considered without truly understanding the patient's full health profile, lifestyle, personal needs or biological idiosyncrasies. Often the measures chosen by insurance companies to determine if patients can access a treatment or care option are based on randomized controlled trials of patient populations who have very different characteristics than the patient in question. As a result, insurers choose to take the most restrictive, conservative approach referenced in a study and apply those conservative approaches in their treatment access decisions. The result is denial of access. The patient and provider are forced to jump

through multiple hoops to appeal, a process that for patients is often required to happen through a letter sent via the postal service.

There are many things about this process that could benefit from AI, but not if it's used to automate the current decisionmaking algorithms. If the current algorithms are automated, the process becomes further calcified, extracted from human decisionmaking. Judging from horror stories of deployments to date, nearly impossible to appeal.

If HHS deploys or allows insurance companies to deploy AI-based access decisions, they must require that the AI have a more complete data set and context. In addition, there needs to be clear mandates around appeals processes, timeframes for decisions and transparency about decision makers, rather than only offering a PO Box for an appeals department. These denials create remarkable burden for patients and caregivers, costing them countless hours of time, anguish, pain and frustration. Many give up, which results in worsening health and ultimately - especially for Medicare - higher costs of care.

10. Are there specific areas of AI research that HHS should prioritize to accelerate the adoption of AI as part of clinical care?

Individualized Benefits Determination

On more than one occasion, I have been told by prior heads of CMMI that it is impossible to create a mechanism for doing nuanced, individualized decisions about patient treatment access at scale. In response, I replied that we can't do that yet, but we should figure out a way to make it happen. We have incredible data. AI has reached new levels of sophistication. AI that is informed with granular data, combined with patient and provider health goals could be much better at making decisions about enabling access to care that works for that specific patient in promoting their very specific health and care needs.

This will take work, but it is doable. And, frankly, I think it's the only way that we can truly enable personalized medicine, tertiary prevention and healthspan, especially

within the context of Medicare. The existing “once size fits all” treatment algorithms produce predictable results. The treatments work for some, work OK for others, don’t work for many and often, for a subgroup of patients, cause harm. Making “one size fits all” decisions about access may have made sense 20 years ago, but it doesn’t in the age of AI.

I encourage HHS/CMS to invest in research into the development of AI tools that can create personalized, context and goal-driven reimbursement decisions at scale. It would be a bold, future-oriented approach to updating our reimbursement and access decisions in ways that improve health, encourage more patient engagement in measuring and maintaining aspects of their health and ultimately reducing the substantial taxpayer burden of poorly managed chronic disease.