



Christi J. Braun, Esq.
christi.braun@ikshealth.com

February 23, 2026

Via Federal eRulemaking Portal:

Assistant Secretary for Technology Policy
Office of the National Coordinator for Health Information Technology
Attn: HHS Health Sector AI RFI
Mary E. Switzer Building
Mail Stop: 7033A
330 C Street SW
Washington, DC 20201

Re: HHS Health Sector Artificial Intelligence Request for Information

Dear Assistant Secretary for Technology Policy,

On December 4, 2025, the Department of Health and Human Services (HHS) issued the HHS Artificial Intelligence (AI) Strategy in response to OMB Memorandum 25-21. As the Department seeks public feedback to establish a forward-leaning, industry-supportive approach to AI in clinical care, Inventurus Knowledge Solutions, Inc. (d/b/a IKS Health) is pleased to provide our experience-based feedback.

IKS Health is a healthcare solutions partner of health systems and physician organizations with nearly two decades of experience supporting clinical, operational, and revenue cycle performance across the care continuum. We improve revenue performance, enhance patient access and engagement, relieve administrative burden for care teams, and advance sustainable care delivery under both fee-for-service and value-based payment models. Our leadership team includes former health system executives and practicing clinicians with substantial experience in healthcare operations, clinical workflow design, and health data governance. This operational perspective informs how we develop, evaluate, and deploy Artificial Intelligence (AI) within complex delivery environments, integrating advanced automation with structured human oversight to promote safety, accountability, and demonstrable impact. Our ultimate goal, which AI makes possible, is to eliminate the “chores” that increasingly burden medical professionals, allowing them to focus on their core mission: care delivery.

In response to HHS’s Request for Information regarding AI and the American healthcare system, we submit the following feedback to inform HHS’s strategy across regulation, reimbursement and research and development. Our recommendations center on dismantling data silos, modernizing regulatory frameworks, and aligning payment policies to ensure that AI serves as a tool for clinical excellence.

1. What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

While issues of ethics and liability are frequently debated, the single greatest barrier to the adoption and scaling of AI is the persistent lack of access to clinical data stored in electronic health records (EHRs), practice management systems, remote patient monitoring, and population health technologies. AI presents a unique opportunity to fundamentally re-design the EHR experience itself, shifting away from

transactional billing and siloed documentation systems that prioritize physician time for documentation over patient care. By integrating AI into platforms that orchestrate relationships and automate non-core tasks, we can allow physicians to return to their primary role, but this evolution is entirely dependent on AI's ability to read from and write to the EHR without vendor-imposed friction.

For third-party applications to communicate with an EHR, the vendor must provide a functional application programming interface (API), which acts as a secure “doorway” between the EHR and the applications. Without a functional API, the EHR acts as a “walled garden,” trapping data inside and requiring manual, “click-heavy” human intervention. For an AI tool to effectively support a physician's workflow, a seamless API connection is required for reading from and writing to the EHR.

IKS Health's ambient AI scribe solutions—Scribble Now, Scribble Swift, and Scribble Pro—illustrate this technical necessity. Through an API, the Scribble app can instantly query the patient's medical history, current medications, and previous diagnoses. For ambient scribing, reading the longitudinal record ensures the AI understands the difference between a new acute symptom and a long-standing chronic condition, significantly reducing the risk of “hallucinations.” Once the Scribble app generates a clinical note, it must have the ability to write back to the EHR, or the clinician will be forced to copy and paste text, which maintains, rather than eliminates, the documentation chore. A robust, standardized API allows the Scribble app to write the clinical notes—in real time—directly into the correct EHR fields for the clinician to review and finalize.

The challenge is not just the existence of APIs, but their variability. Many EHR vendors utilize proprietary, rather than standardized, APIs that force app developers to build unique connections for every hospital system and physician organization, a process that is cost-prohibitive for smaller innovators. The Health Level Seven (HL7®) Fast Healthcare Interoperability Resources (FHIR®) standard was intended to be the universal language of health data, allowing vendors to securely share information through standard HTTP operations and enabling health care providers and systems to get a more complete view of their patients and business processes. FHIR provides the necessary structure that allows AI to navigate through vast datasets consistently and efficiently. By utilizing the same “RESTful” architecture that enables global entities like Google and Amazon to exchange data over the internet, FHIR brings healthcare into the modern digital era.

Despite the progress intended by the 21st Century Cures Act and the Cures Act Final Rule, which included the adoption a suite of FHIR-based requirements in the Health IT Certification Program and mandated that certified EHRs support FHIR APIs by the end of 2022, we are currently seeing a period of lax compliance regarding the Office of the National Coordinator for Health IT (ONC) requirements. Many EHR vendors have implemented FHIR inconsistently or in ways that maintain proprietary silos. By enforcing a standardized FHIR API, the Assistant Secretary for Technology Policy (ASTP)/ONC can ensure that an AI developer can build a tool once and deploy it across the entire healthcare ecosystem. This “plug-and-play” capability is the only way to scale AI tools that can meaningfully reduce clinician burnout and lower costs across the American healthcare system.

In the Health Data, Technology, and Interoperability: ASTP/ONC Deregulatory Actions to Unleash Prosperity (HTI-5) Proposed Rule, ASTP/ONC recognizes that “[t]he transition to a FHIR-based API ecosystem is a strategic imperative”. While the proposed rule removes non-FHIR-based certification criteria, it takes no steps to require support for more updated criteria. ASTP/ONC must mandate and enforce specific technical standards that ensure data is not only available but “AI-ready.” Specifically, ASTP/ONC should require full adherence to the HL7® FHIR® Bulk Data Access (Flat FHIR) specification (v1.0.1 or higher) and HL7®'s FHIR US Core Implementation Guide (v5.0.1 or higher). These standards enable the standardized export of large-scale population datasets using NDJSON (Newline Delimited JSON), which

is essential for training robust AI models. Without this “bulk” capability, AI is fed data one patient at a time, which is insufficient for identifying systemic care gaps or performing population-level analytics.

We support ASTP/ONC’s proposed revisions in HTI-5 to Part 171 to clarify that interfering with AI accessing EHI is information blocking. In particular, we agree that modifying the definitions of “access,” “exchange,” and “use” under 45 CFR §171.102 is necessary and appropriate and support retiring the “third party seeking modification use” condition of the Infeasibility Exception (45 CFR §171.204(a)(3)). We also agree with the removal of the Trusted Exchange Framework and Common Agreement™ (TEFCA™) Manner Exception (45 CFR §171.403).

We note, however, that HHS cannot stop at rule making. It needs to prosecute information blocking and levy meaningful fines.

2. What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.

To effectively incentivize AI adoption, HHS must address the misaligned regulatory and payment frameworks that currently penalize innovation. Our existing systems—built for a manual, encounter-based era—frequently obstruct the data-rich, longitudinal nature of AI-driven medicine.

A primary bottleneck is the “minimum necessary” standard (45 CFR § 164.502(b)) within the HIPAA Privacy Rule, which is often interpreted to require significant restrictions on data use and access. Clarifying this standard is essential because modern clinical AI models require massive, high-fidelity datasets to ensure clinical safety and statistical reliability. HHS could issue guidance specifying that, for the purposes of training and validating clinical AI, a comprehensive longitudinal view of the patient is the baseline requirement. We recommend that HHS issue guidance that “bulk” access is the actual “minimum necessary” to achieve a safe and effective AI tool or revise 45 CFR § 164.506 to clarify that “health care operations” includes AI validation and includes the right to access comprehensive, longitudinal datasets.

Furthermore, the requirement for specific patient consent for uses beyond treatment, payment, or operations (45 CFR § 164.508) creates friction for vital “secondary” uses like AI safety and bias monitoring. HHS could establish a new “AI-Ready De-identification Standard” refining the existing safe harbor methods under 45 CFR § 164.514 to allow for the use of “Limited Data Sets” that retain critical temporal and geographic information that are currently restricted but vital for AI to track disease progression.

Scaling AI also requires moving away from manual data extraction. As noted above, ASTP/ONC should mandate HL7® FHIR® Bulk Data Access (v1.0.1+) to facilitate the export of population-level datasets in NDJSON format. This should be supported by the US Core Implementation Guide (v5.0.1+), ensuring an AI model built for one health system works in another without reconfiguration. Additionally, the SMART on FHIR Framework must be the standard for “plug-and-play” AI apps, allowing tools like ambient listening to integrate directly into the physician workflow without vendor-specific barriers.

From a reimbursement standpoint, we must avoid the “Telehealth Trap,” where productivity gains from technology are used as a justification for immediate fee schedule cuts. When telehealth scaled, providers faced automatic reimbursement decreases under the assumption that cost-savings should be passed directly to the payers. HHS should prioritize a “Value-Preservation” payment model. If an AI tool reduces documentation time, CMS should not respond by simply lowering the Work Relative Value Units

(RVUs) for those services. We need a commitment that productivity gains realized through AI will be shared with the providers who invest in these technologies.

The most profound shift required is a move away from our current per-encounter payment model toward population-based and outcome-based payments, such as via 42 CFR Part 414, Subpart O (the Quality Payment Program). As long as the financial “engine” of a hospital is built on the volume of visits, AI tools that *prevent* visits will be seen as a financial liability. We recommend 42 CFR §414.1375 (Quality Category) be updated to explicitly permit and reward “digital quality measures” (dQMs) that pull directly from AI-filtered data, reducing the manual “chart-chase” for staff and physicians. By aligning technical standards with a refreshed MIPS incentive structure, we can ensure that AI restores the clinician-patient relationship and shifts the focus toward lasting health outcomes.

3. For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (e.g., relating to liability, indemnification, privacy, and security). What novel legal and implementation issues exist and what role, if any, should HHS play to help address them?

While the question specifically focuses on “non-medical devices,” we want to acknowledge that a significant implementation barrier and administrative hurdle that stalls deployment of AI remains the “gray zone” of what constitutes “Software as a Medical Device” (SaMD) subject to FDA regulation under 21 CFR § 860. While the Cures Act excluded some clinical decision support software from the definition of a medical device, current ambiguity makes it unclear if a novel AI application, particularly one involving predictive analytics or providing a differential diagnosis based on pattern recognition, is considered SaMD. This “gray zone” creates a culture of risk aversion, as health systems fear the legal and compliance ramifications of deploying unregulated software that may later be categorized as a medical device. HHS should clarify whether, if software output “drives” clinical action and the logic remains “opaque”—meaning the physician cannot fully verify the underlying reasoning—it should be classified as SaMD.

Even when a tool requires regulatory oversight, premarket review cycles are incompatible with the rapid iterative nature of AI software. A model that is clinically sound today may be based on data that becomes obsolete during a twelve-month review period. Without a modernized, fast-track “process validation” pathway at the FDA, regulatory bottlenecks will continue to stifle innovative care.

In addressing the governance and accountability of non-medical device AI, HHS should adopt a policy of “Augmented Intelligence,” where the physician remains the “final click” and ultimate authority. However, to prevent “automation bias,” the government must provide specific guardrails that ensure clinicians do not over-rely on AI outputs.

If we are to have widespread adoption of AI, accountability cannot rest solely on the clinician. If a clinician follows the standard of care but is misled by an AI “hallucination,” the liability should shift. We need a dual-accountability framework where physicians are responsible for clinical application, while AI developers are responsible for structural failures, such as data-drift or hallucinations that a reasonable clinician could not have detected. To codify this, HHS should update 21 CFR §820 (Quality System Regulation) to include “AI hallucination monitoring” and “bias testing” as mandatory design controls. HHS should also clarify 42 CFR § 482 (Conditions of Participation) to require that hospitals establish a “Clinical AI Governance Committee.”

4. For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and other workflow and human-centered evaluation methods for clinical care? Should HHS further support these processes? If so, which mechanisms would be most impactful (e.g., contracts, grants, cooperative agreements, and/or prize competitions)?

HHS must move beyond static validation and adopt a “Human-in-the-Loop” (HITL) philosophy. This approach prioritizes real-world performance over laboratory accuracy. By treating AI evaluation as a continuous clinical audit, we can implement a lifecycle approach where pre-deployment adversarial testing identifies “hallucination” risks and post-deployment automated discrepancy checks flag outputs that diverge from human consensus. Shifting focus to human-centered metrics, such as “Intervention Accuracy,” allows us to measure how often an expert agrees with the AI.

HHS should leverage the authorities granted by the One Big Beautiful Bill Act (Public Law 119-21) to catalyze this innovation. Specifically, the Rural Health Transformation (RHT) Program provides a prime opportunity to fund “AI-Human Teams” that help rural generalists practice at the top of their license. These cooperative agreements can be supplemented by prize competitions for clinical reasoning, modeled after the successful ARPA-H PARADIGM program, to incentivize “rugged” AI tools capable of maintaining accuracy in resource-constrained environments. Additionally, utilizing the framework of the AI in Medical Education Act, HRSA can provide educational grants to train the next generation of physicians to oversee these tools, specifically to close care disparities in minority communities.

The most impactful move HHS can make is a CMMI Demonstration Project that shifts to outcome-aligned payments based on a patient’s current health status and longitudinal outcomes. Rewarding the prevention of acute events through AI’s predictive power aligns the financial success of a hospital or physician with the actual health of the community. This strategy directly supports Pillars 3 and 5 of the HHS AI Strategy (December 2025).

5. How can HHS best support private sector activities (e.g., accreditation, certification, industry-driven testing, and credentialing) to promote innovative and effective AI use in clinical care?

To foster a robust ecosystem for clinical AI, HHS must move toward a policy of active enablement. As noted above, the ONC Health IT Certification Program (45 CFR Part 170) should prioritize data accessibility as a core requirement for certification.

To bridge the gap between private innovation and public safety, HHS could establish a Federally Qualified National Patient Database. By leveraging HL7® FHIR® standards, this repository could provide de-identified, population-level datasets that preserve clinical integrity while protecting individual privacy. Access to this database should be offered to developers and researchers through a subsidized, affordable fee structure, ensuring that American ingenuity is not sidelined by the prohibitive costs of data acquisition currently controlled by a few dominant firms.

As the market for clinical AI grows, both providers and patients need a reliable way to verify the credibility of new tools. Rather than a rigid federal mandate, HHS should facilitate a Public-Private Partnership (PPP) to create a voluntary certification program for AI application developers. This model of co-regulation would allow industry leaders to participate in setting performance and safety standards, incorporating globally recognized frameworks, like ISO/IEC 42001 for AI Management Systems. Such a certification would serve as a “seal of trust,” signaling that a vendor adheres to rigorous self-regulatory standards and ethical design principles.

Finally, we must establish a clear legal baseline for the digital age: EHR companies are stewards of data, not owners. HHS should issue definitive guidance or update 45 CFR §164 to clarify that patients

own their data. Under this principle, EHR vendors must be prohibited from claiming intellectual property rights over any AI product derived from patient data processed within their software. We need to ensure that medical progress remains a shared human benefit.

6. Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?

We are still in the nascent stages of AI deployment, so evidence of exceptional performance is currently limited to specific, controlled-use cases rather than systemic healthcare transformation. Despite this early stage, AI shows significant promise in disease prevention, closing persistent gaps in care, and reducing hospital readmissions. Perhaps most immediately, AI is effectively reducing "pajama time"—the time physicians spend after-hours on documentation—by automating clinical scribing and streamlining data retrieval for quality reporting.

We have found that, where well-governed data is highly available with full clinical context, AI is remarkably cost-effective to develop and implement. Conversely, AI has consistently underperformed when developed as a "proof-of-concept" on limited or siloed data, as performance often collapses upon introduction to a live clinical workflow.

The greatest potential for novel AI lies in its ability to act as a "great equalizer" for care delivery. AI could function like a clinical "Siri" or "Alexa" to triage patients to the appropriate level of care, significantly reducing the burden on overcrowded emergency departments. Additionally, AI tools with access to longitudinal records for large populations can provide treatment insights that are currently impossible for humans to synthesize. This is vital for rare diseases, where AI can identify "patients like me" and guide physicians toward specific tests and treatments, reducing both the time to diagnosis and the cost of care.

7. Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?

In addressing the governance of AI within healthcare organizations, it is critical to distinguish between those who champion the technology and those who authorize its procurement. Within most health systems, physician-led committees—often structured as Clinical Technology Councils or AI Oversight Boards—exert the most significant influence on adoption. These clinicians are the primary arbiters of a tool's clinical validity, safety, and workflow compatibility. While hospital executives and Chief Information Officers typically hold the ultimate "purchasing" authority, they rarely move forward without a formal endorsement from these physician leaders, who ensure the technology serves as augmented intelligence, rather than an administrative burden.

8. Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

As explained in the answer to Question 1, interoperability, enabled through implementation of HL7® FHIR®, is the indispensable bridge that allows AI to move beyond theoretical analysis into meaningful, frontline care. For AI to truly transform the clinical landscape, it must have the capability to read directly from the EHR, capture the patient's current state in real-time, and execute the "chores," such as writing care notes, supporting the creation of complex plans of care, and managing prescription refills.

By automating these administrative burdens, we allow physicians to return their focus to the patient, restoring the human element of medicine.

Furthermore, to fuel research and large-scale model development, HL7® FHIR® Bulk Data Access (v1.0.1+) must be made universally available. Accessing data patient-by-patient is insufficient for the training and benchmarking of sophisticated AI; we require the ability to export population-level datasets to identify trends and validate clinical interventions at scale.

Concluding Recommendations

To facilitate innovation in AI, ASTP/ONC must enforce a standardized FHIR API and mandate and enforce specific technical standards like HL7® FHIR® Bulk Data Access (v1.0.1 or higher) and the FHIR® US Core Implementation Guide (v5.0.1 or higher). Further, HHS must issue guidance on the “minimum necessary” standard and ensure access to bulk data is permissible under HIPAA. By aligning technical standards with modernized payment and regulatory structures, HHS can ensure AI acts as a catalyst for clinician wellness and superior patient care.

We appreciate the opportunity to respond to this Request for Information and stand ready to answer any questions. Given the complexities of integrating AI with healthcare workflows, IKS Health welcomes the opportunity to provide your staff with a technical briefing. We are also happy to facilitate a follow-up conversation with your policy team to discuss how a dual-accountability framework can be implemented to protect both patients and clinicians.

Sincerely,

A handwritten signature in black ink, appearing to read "Christi J. Braun", with a stylized flourish extending to the right.

CHRISTI J. BRAUN
Chief Legal & Compliance Officer