

February 23, 2026

Department of Health and Human Services
Assistant Secretary for Technology Policy and the Office of the National Coordinator for Health
Information Technology
Attention: Request for Information
HHS Health Sector AI RFI
Mary E. Switzer Building
Mail Stop: 7033A
330 C Street SW
Washington, DC 20201

To: Department of Health and Human Services,

Re: Request for Information (RFI)—Accelerating the Adoption and Use of Artificial Intelligence as Part
of Clinical Care (RIN 0955-AA13; 90 FR 60108)

Thank you for the opportunity to comment on the Department of Health and Human Services' (HHS) efforts to accelerate the adoption and use of artificial intelligence (AI) as part of clinical care. We are at a critical juncture in ensuring that we draw a clear balance between fostering innovation without sacrificing safety and quality and unduly shifting liability to one side of the deployer/developer paradigm. As both a developer and a deployer of healthcare AI, Johns Hopkins Medicine (JHM) understands the need for this balance and commends HHS and the Assistant Secretary for Technology Policy/Office of the National Coordinator for Health Information Technology (ASTP / ONC) leadership¹ in seeking feedback from the industry on these topics. Below, we provide comments on the areas highlighted in the RFI (Regulation, Reimbursement, and Research), as well as provide specific recommendations aimed at ensuring accelerated adoption of artificial intelligence in clinical care in a safe, reliable, and innovative manner, without overburdening healthcare systems.

1. Regulation

a. Deregulation of Developers Has a Paradoxical Chilling Effect on Clinical AI Adoption.

i. Overview:

¹ Although this RFI originates from ASTP/ONC, we understand that HHS has adopted a "One HHS" approach to AI, as articulated in the December 2025 HHS AI Strategy. Given this coordinated, department-wide framework, our comments address regulatory considerations that span all relevant HHS operating divisions—including the FDA, CMS, and ASTP/ONC—where AI oversight intersects with clinical care delivery. We believe this cross-cutting perspective is both necessary and consistent with HHS's stated goal of providing a unified, coherent approach to regulation of healthcare AI.

JHM agrees that AI has enormous potential to transform healthcare delivery and is committed to developing and deploying AI in ways that are innovative, compliant with current regulations, and safe and beneficial for our patients. However, HHS's focus on a path of deregulation for AI developers will unintentionally *chill* AI adoption by healthcare AI deployers, rather than spur it. Without a trusted framework to evaluate new AI tools, providers will gravitate towards engaging with known vendors for AI tools, rather than assuming the risk of newer, unknown, start-up companies that have an unproven track record and potentially shorter-term financial viability. To balance these priorities and encourage adoption, we recommend a streamlined, less burdensome approach to regulating clinical AI tools—one that is expedient but does not sacrifice accountability, and that appropriately allocates responsibility among developers, deployers, and regulators.

ii. Explanation

In an environment where AI tools enter the market without meaningful federal oversight or validation, the regulatory burden has been shifted onto the healthcare provider. The resultant fear of liability and concerns for patient safety creates a chilling effect and adds significant administrative burden, requiring individual systems to vet tools on their own, rather than at a centralized level. Specifically, healthcare organizations looking to deploy these tools assume the entire burden of validation testing, bias assessment, performance monitoring, and post-deployment surveillance—functions that would otherwise be performed or overseen by regulatory bodies or developers, and that are nearly impossible to operationalize within a healthcare organization at this time. This burden falls particularly heavily on rural and community care organizations. Ultimately, clinicians and patients bear the risk when AI tools underperform, produce biased outputs, or cause patient harm, and with limited recourse against developers who have often disclaimed liability through contractual terms, providers will be hesitant to adopt the tools.

In short: deregulation of developers has become de facto delegation of such responsibilities to deployers. This asymmetry in responsibility creates a powerful chilling effect on AI adoption. The cost, expertise, and time required to conduct rigorous internal evaluations of each AI tool—evaluations that the tech industry is best equipped to perform—becomes a barrier to adoption. Moreover, in the absence of reasonable federal regulations, states have started developing their own laws and regulations, creating inconsistent approaches and making adoption of AI for multi-state providers challenging.

We urge HHS to recognize that thoughtful regulation is an enabler of innovation and a driver of adoption—not an impediment to innovation. A dependable, well-understood federal regulatory framework that appropriately allocates responsibility will provide healthcare organizations and providers with the confidence to adopt AI tools, accelerate deployment timelines, and ultimately deliver the benefits of AI to patients.

iii. Recommendations

We encourage HHS to adopt a regulatory framework that is proportional to the potential for patient harm, increasing oversight for high-risk applications and enabling faster adoption for lower-risk tools. Notably, “proportionate” should not mean “absent”—even low-risk AI tools must meet baseline requirements that provide healthcare organizations and providers with actionable information for deployment decisions. To achieve this, we propose a two-phase approach:

- Phase 1: Expedited Pre-Deployment Review with Transparency Requirements

For clinical AI tools, rather than lengthy pre-market approval processes, to encourage innovation, HHS could establish an expedited review pathway that focuses on transparency and documentation requirements, minimum safety thresholds, and preliminary conditional clearance with clearly defined conditions and monitoring obligations. As part of the pre-deployment review, HHS should require disclosure of training data characteristics, known limitations, and performance metrics across diverse patient populations to enable clinicians to make safe and appropriate clinical decisions in real world settings. Without transparency requirements, health care systems are less likely to rapidly deploy innovative technologies.

- Phase 2: Robust Post-Deployment Monitoring

The unique characteristics of AI—including the potential for model drift, performance degradation over time, and differential performance across patient populations—make post-deployment monitoring essential. We recommend that HHS establish or enhance regulatory mechanisms for mandatory adverse event reporting, real-world performance surveillance, algorithmic auditing authority, expedited withdrawal or modification authority, communication requirements to deployers and end-users, and transparency in post-market findings that vests post-deployment monitoring responsibilities in AI developers. AI developers are best able to modify the AI to accommodate any identified issues. Alternatively, responsibility for such post-deployment monitoring could rest with regulators to ensure transparency and consistent application across all deployments.

This two-phase approach preserves the speed-to-market benefits sought by the Administration and positions AI developers, and not healthcare delivery organizations, to leverage their expertise and bear responsibility for demonstrating baseline safety and fitness for purpose. Importantly, expedited review would provide healthcare organizations with a regulatory signal they can rely upon, dramatically reducing internal governance bottlenecks and accelerating adoption. Additionally, adding the post-deployment phase aligns with the RFI's recognition that "pre- and post-deployment" evaluation methods are critical for clinical AI. By focusing regulatory attention on the post-deployment phase, HHS can enable faster market entry while maintaining robust accountability throughout the product lifecycle. This approach also better allocates responsibility among developers, deployers, and clinician end users, ensuring liability matches degrees of control. Safe harbor provisions for health systems that follow established governance and monitoring protocols could accelerate adoption while maintaining accountability.

b. Existing Regulatory Frameworks Are Ambiguous and Hard to Implement - Chilling Development and Adoption

With respect to existing regulatory frameworks, we appreciate recent and ongoing efforts from the Food and Drug Administration (FDA) to further clarify what constitutes clinical decision support (CDS) as a device and when FDA intends to exercise its CDS enforcement discretion. However, the nuanced-fact-dependent analysis required under the guidance becomes a significant impediment to development and deployment of CDS. Rather than offering fact-specific examples, the guidance could be strengthened by identifying the broader qualities and principles distinguishing CDS devices from other support tools. Such an approach would address potential inconsistent application and interpretation across the industry.

The extensive internal review processes that healthcare organizations will develop to address the complex guidance will delay deployment and result in inconsistency in the quality of patient care across the industry. This will be exacerbated in rural and community care settings with deployers that are less resourced. **We recommend a regulatory framework that is easier to understand, easier to apply to different fact-patterns, and that would lead to consistent application and create confidence in the efficacy and safety of tools being deployed.**

2. Reimbursement

We support HHS and the Centers for Medicare and Medicaid Services (CMS) in the expansion of pathways that support reimbursement for outcomes. AI has the potential to produce cost savings across the care continuum. These benefits, however, will materialize over time and often deliver savings that may inadvertently accrue more to payors, rather than the providers and health systems that bear the cost and effort of AI procurement and implementation.

Reimbursement is critical for the adoption of AI tools in healthcare. To ensure equitable and widespread adoption, we urge HHS and CMS to pursue targeted funding and value-based reimbursement initiatives, supporting providers as they implement, monitor, and govern AI tools. This balanced approach will prevent unintended fee-for-service consequences, protect core clinical services, and promote financial sustainability as AI-generated savings and improved outcomes are realized over time.

With a balanced framework outlined above, implementation and infrastructure funding opportunities could help leverage trusted care delivery entities to expand the use of AI, increase cost-savings across the care ecosystem, and remain financially viable as the future benefits of these tools manifest over time.

3. Research and Development

To drive accelerated adoption of AI in clinical care, investment will be needed in key research areas. The costs of clinical AI are at risk of multiplying if reproducible learnings about the real-world effectiveness of these tools, on both clinical quality, safety, outcomes, cannot be accelerated. Every health system independently incurs significant human capital (expanding health technology and engineering expertise) and infrastructure costs to deploy and monitor clinical AI safely. These duplicative efforts slow adoption and limit our collective ability to learn what works. **Increased funding opportunities from HHS, the National Institutes of Health (NIH), the Agency for Healthcare Research and Quality (AHRQ) and others to accelerate implementation science in clinical AI would help both advance these implementation initiatives and share reproducible lessons across the healthcare ecosystem.**

Research should prioritize understanding human-AI interactions, which will have direct implications for future regulatory and reimbursement structures.

To support payment evolution, the U.S. should fund health economics research on the cost of AI procurement and implementation (including infrastructure, training, validation, and ongoing monitoring), downstream impacts on healthcare utilization, clinical outcomes, and costs, time horizon for cost reduction returns, and allocations of those costs and benefits throughout the healthcare continuum.

Clinical AI will be core to our healthcare future, which also means investments are needed in undergraduate, graduate, and continuing medical education. Academic medical centers are positioned at the center of this educational continuum. AI-in-education funding should support both research and development of novel approaches to clinical training. Investment in clinical AI tools will yield minimal return unless we educate, upskill, and reskill our current and future workforce. Training, validating, and monitoring clinical intelligence—historically human, and now both human and artificial—is core to the mission of academic medicine.

Finally, we believe high quality, validated healthcare data is one of the most important assets for accelerating clinical AI development. However, that data does not exist at zero cost – especially unstructured data, which represents the vast majority of the underutilized healthcare datasets. This data represents clinical expertise and health system investment in centralized infrastructure for flexible use including privacy protecting tools. **We encourage funding opportunities that will accelerate expansion of safe, secure, and federated data platforms to support AI tool research and development.**

Conclusion

Johns Hopkins Medicine appreciates the opportunity to comment on this critical effort to accelerate clinical AI adoption in a safe and responsible manner. As a large academic health system, we are committed to harnessing AI's potential to improve patient outcomes, reduce clinician burden, and address the workforce challenges facing healthcare today. At the same time, we recognize our responsibility to ensure that AI is deployed safely, consistently, and in ways that maintain the trust of the patients and communities we and our fellow health systems serve.

We stand ready to partner with HHS in this important work and welcome further dialogue on how academic health systems can contribute to the responsible advancement of AI in clinical care.

Sincerely,

/s/ Deanna Hanisch

Deanna Hanisch, MS
Vice President and Chief Information Officer
The Johns Hopkins Health System Corporation

A handwritten signature in black ink, appearing to read 'B. Hasselfeld', enclosed within a large, loopy oval shape.

Brian Hasselfeld, MD
Executive Medical Director, Digital Health and Innovation
Johns Hopkins Medicine