

APA's Mental Health Technology Advisory Committee
Public Comment to RFI: Accelerating the Adoption and Use of Artificial
Intelligence as Part of Clinical Care
Docket No. HHS-ONC-2026-0001

Submitted by: American Psychological Association Mental Health Tech Advisory Committee

Date: February 23, 2026

Contact: Nina Horne, M.P.P., Committee Member, nina.horne@icloud.com

Opening Statement

I. Introduction

The [American Psychological Association's \(APA\) Mental Health Tech Advisory Committee](#), a volunteer group of external experts, appreciates the opportunity to respond to the Department of Health and Human Services' (HHS) Request for Information on accelerating the adoption and use of artificial intelligence (AI) as part of clinical care (90 FR 60108, December 23, 2025).

We share the Department's goal of accelerating AI adoption. The mental health system's failures are urgent: 48% of adults with mental illness receive no treatment, and 122 million Americans live in Mental Health Professional Shortage Areas. AI offers meaningful potential to expand access, reduce administrative burden, and improve care delivery. We welcome HHS's commitment to using its regulatory, reimbursement, and research and development levers to address these challenges.

Our comments below address the overarching policy framework we believe is necessary to support sustainable AI adoption in mental health care, followed by specific recommendations responsive to the RFI's three areas of inquiry. We draw on our expertise in clinical psychology, behavioral health delivery systems, and the emerging evidence on AI-assisted mental health interventions.

II. Framing: The Prior Question

Every sound policy analysis begins with a prior question: what problem are we trying to solve? HHS has framed its inquiry around acceleration — how the Department can drive faster adoption of AI in clinical care. We support this objective. However, we urge the

Department to adopt a complementary framing that is particularly critical for mental health applications: *how do we accelerate adoption without doing harm?*

This is not a rhetorical distinction. It produces different policies. Pure acceleration framing leads primarily to barrier removal —deregulation, funding increases, and streamlined pathways. A safety-integrated acceleration framing leads to infrastructure building: safety floors, evidence systems, evaluation capacity, and governance expertise that allow innovation to proceed on a foundation that does not collapse under the first catastrophic failure.

We respectfully note that the historical record is instructive. Sectors that prioritized speed without proportionate governance — telecommunications in the 1990s, social media in the 2000s, cryptocurrency in the 2010s — experienced rapid adoption followed by crisis, erosion of public trust, and reactive overregulation that ultimately set them back further than proactive governance would have. Sectors that invested in safety infrastructure alongside innovation — pharmaceuticals, aviation, automotive safety — produced slower initial adoption but sustained, durable growth built on institutional trust. Mental health care demands the latter approach. In this domain, the most vulnerable populations are the most likely to encounter AI; harms are severe and often irreversible; and technology operates on the human mind itself. Innovation that harms patients is not innovation — it is a cost externalized onto the most vulnerable people in the system.

III. AI-Specific Risks in Mental Health Care

While our comments are broadly supportive of AI adoption, we wish to draw HHS's attention to risks that are uniquely acute in mental health applications and that require proportionate governance:

- **Suicide escalation failures.** AI systems providing emotional support or crisis intervention may mishandle acute suicidal ideation, with potentially fatal consequences.
- **Parasocial dependency.** Users may form unhealthy attachment relationships with AI tools, displacing human connection and therapeutic relationships in ways that worsen long-term outcomes.
- **Bias and discrimination.** AI systems trained on unrepresentative data may produce outputs that are systematically less effective or actively harmful for racial and ethnic minorities, LGBTQ+ individuals, and other underserved populations.

- **Erosion of therapeutic standards.** Widespread AI-mediated mental health interactions, delivered outside clinical oversight, risk normalizing substandard care and undermining evidence-based practice.
- **Conflict with Existing Treatment Process.** Use of AI for mental health purposes may conflict with existing treatment processes and planning with AI providing information that goes against clinician advice.

The scale of current adoption makes these risks urgent. Population-level risk is the product of prevalence and individual harm. When hundreds of millions of Americans use AI tools weekly for emotional support, crisis management, and mental health guidance — often outside any clinical relationship — even a small per-interaction probability of harm produces vast aggregate damage.

IV. Central Argument: Trust as Infrastructure for Sustainable Innovation

Our central recommendation to the Department is this: sustainable innovation — not unregulated speed — is what actually accelerates adoption. Trust is the essential infrastructure for both clinical care and technology markets, and HHS policy should be designed to build it.

In clinical care, the therapeutic relationship — built on trust — is the single strongest predictor of treatment outcomes across all modalities. In technology markets, institutional trust is the precondition for procurement, reimbursement, and sustained use. AI in mental health care must earn clinical trust to be effective and market trust to scale.

Yet early documented harms — including suicide escalation failures, dependency cases, and discriminatory outputs — are already eroding both forms of trust simultaneously. Every unregulated adverse event makes clinicians more skeptical, health systems more cautious, payers more resistant, and patients more wary. Rebuilding trust once lost costs far more — in both time and resources — than building it correctly from the start.

We urge HHS to recognize that the current market externalizes critical costs: psychological harm to users, safety research that no individual firm funds, crisis detection infrastructure that no product includes, and evaluation systems that no one builds. Innovation that externalizes these costs is not sustainable. It produces the boom-and-bust cycle that has characterized previous digital health waves, where early hype gives way to catastrophic safety failure, triggering reactive overregulation that freezes the sector for years.

Sustainable innovation internalizes these costs through policy design, ensuring that the process of acceleration does not destroy the conditions for its own continuation.

V. Summary of Recommendations

In the sections that follow, we provide specific recommendations organized around the RFI's three areas of inquiry. Each recommendation is designed to accelerate adoption while preserving patient safety and public trust:

- **Regulation.** Clear, risk-proportionate regulatory pathways for AI mental health tools, with tiered oversight based on clinical risk level and population vulnerability.
- **Reimbursement.** Payment policy alignment that incentivizes evidence-based AI tools and does not inadvertently reward unvalidated products or penalize human clinical services.
- **Research and Development.** Modernized evidence standards suited to AI's characteristics, federal investment in evaluation infrastructure, and public-private partnerships to generate the safety research the market will not produce on its own.

We elaborate on each of these areas in the sections below. Further, the six recommendations in this submission constitute a coherent framework for sustainable innovation, operating through six reinforcing mechanisms:

- **Predictability** accelerates investment — firms build when they know the rules.
- **Safety floors** prevent catastrophic adoption shocks that freeze entire sectors.
- **Right-sized evaluation methods** prevent both false rigor and false speed.
- **Reimbursement alignment** creates demand pull — nothing scales without payment.
- **Socializing safety costs now** prevents the boom-and-bust cycle that has frozen previous digital health markets.
- **Clinical expertise at the table** prevents oversight that's calibrated to the wrong risks.

Drawing on the clinical and scientific expertise of APA's 185,000+ members — including licensed practitioners who deliver mental health care daily across every setting HHS's policy will affect us. We offer the following guiding principles:

-The pursuit of innovation must not dilute established clinical standards. AI tools making therapeutic claims must meet the same well-validated evidentiary requirements that apply to human-delivered interventions.

-Patient safety is integral to sustainable innovation. In high-stakes domains, safety and innovation are not competing values — safety is the condition that enables innovation to scale. Trust is the mechanism: clinical trust enables therapeutic effectiveness, and institutional trust enables market adoption. Markets that externalize harm destroy both forms of trust simultaneously, setting adoption back further than any regulatory requirement.

-Psychology expertise is essential. Effective and safe innovation requires deep understanding of therapeutic relationships, clinical judgment, and mental health outcomes.

Executive Summary of Recommendations

We recommend HHS:

1. **Establish risk-tiered oversight** for mental health AI, with deployer certification and a psychologically active ingredient standard.
2. **Mandate crisis safeguards** as a baseline safety floor for all AI used in mental health contexts.
3. **Modernize evidence standards** to accommodate adaptive learning AI systems.
4. **Align reimbursement** with augmentation and validated outcomes, not unvalidated replacement.
5. **Fund independent evaluation** and regulatory sandboxes scaled to risk.
6. **Integrate psychology expertise** into governance, with an Interagency Task Force on AI and Psychological Safety.

Recommendation 1: Establish a Mental Health–Specific Tiered Regulatory Framework

HHS should adopt a risk-tiered oversight model that matches scrutiny to clinical claims, risk profile, and functional intent — complementing the existing SaMD framework for products that fall outside it. The framework should include deployer certification requirements, foundation model transparency obligations, and a "psychologically active ingredient" standard that extends baseline safety obligations to consumer AI products with foreseeable psychological impact.

Recommendation 2: Mandate Baseline Suicide and Crisis Safeguards

All AI systems used in or foreseeably affecting mental health contexts — including consumer products meeting the psychologically active ingredient criteria — must implement documented and interpretable crisis detection and escalation protocols.

Recommendation 3: Modernize Evidentiary Standards for Digital Mental Health

HHS should adapt evidence expectations to learning systems, incorporating real-world evidence, adaptive designs, and lifecycle monitoring through an adaptive regulatory approach that generates information and updates requirements over time.

Recommendation 4: Align Reimbursement with Augmentation and Outcomes

Payment policy should reward AI that reduces burden and improves validated outcomes, not tools that substitute for human care without validation. Reimbursement design must guard against a two-tier system that widens existing disparities.

Recommendation 5: Support Independent Evaluation and Proportionate Auditing

Higher-risk systems should undergo independent third-party audits, scaled to risk tier. HHS should establish regulatory sandboxes to create a pathway for supervised evidence-generating deployment of novel mental health AI tools.

Recommendation 6: Integrate Clinical Psychology Expertise into Governance

Mental health expertise must be embedded in review, evaluation, and policy development processes. HHS should convene an Interagency Task Force on AI and Psychological Safety with CPSC, FTC, and state agencies, and fund research on the developmental impact of AI-mediated relationships on children.

Recommendations

Recommendation 1: Establish a Mental Health–Specific Tiered Regulatory Framework

Responds to RFI Questions 1, 2, 3

Acceleration Rationale:

A clear, risk-tiered pathway reduces regulatory ambiguity around "wellness vs. clinical" boundaries, giving innovators predictable expectations and shortening time-to-market for lower-risk tools while preserving scrutiny for higher-risk interventions. Without tiering,

firms either overbuild compliance (slowing innovation) or underinvest in safety (risking enforcement and reputational collapse that chills the sector).

Specific Recommendations:

HHS should establish a three-tier oversight structure for mental health AI tools:

Tier 1 — Minimal Risk (Enforcement Discretion). Wellness tools: meditation timers, gratitude journals, general mood tracking. No clinical claims, no vulnerable population targeting. Voluntary compliance with basic privacy and safety standards. Required transparency disclosures: what the tool does, what it does not do, what evidence exists.

Tier 2 — Moderate Risk (Streamlined Review). Tools making limited clinical claims: symptom tracking for diagnosed conditions, psychoeducation, clinician workflow support (documentation, screening, treatment planning assistance). Evidence requirements proportional to claims. Registration, adverse event reporting, and post-market monitoring required. Baseline safety safeguards required.

Tier 3 — Higher Risk (Full Clinical Review). Tools claiming to diagnose or treat psychiatric disorders. AI-driven therapeutic interventions that engage users in therapeutic dialogue, crisis intervention, or clinical assessment without concurrent clinician involvement. Tools targeting high-risk populations or conditions. Pre-deployment safety evaluation, mandatory crisis safeguards (see Recommendation 2), outcome reporting, and independent audit (see Recommendation 5) required. Evidence standards are equivalent to psychotherapeutic interventions.

Enforcement scalability. This tiered structure is designed for enforcement scalability. Tier 1 relies on market transparency — required disclosures and voluntary compliance — rather than active regulatory enforcement, consuming minimal agency resources. Tier 2 uses registration and adverse event reporting — scalable mechanisms that generate information for targeted enforcement rather than requiring pre-market review of every product. Only Tier 3 requires the resource-intensive pre-deployment evaluation and independent audit that traditional regulatory approaches demand. This design ensures that HHS's limited enforcement resources are concentrated where risk is highest, while lower-risk tiers are governed through information requirements that scale with the market.

Relationship with existing SaMD regulation. The FDA's existing Software as a Medical Device (SaMD) framework — developed in alignment with the International Medical Device Regulators Forum (IMDRF) already provides risk-based categorization for software that

meets the definition of a medical device. We do not propose replacing this framework. However, the SaMD framework was designed for deterministic software with predictable, reproducible outputs: a diagnostic algorithm that takes defined inputs and produces a classification, a clinical decision support tool that applies fixed rules to patient data. It categorizes risk based on the seriousness of health condition and the significance of the information the software provides to a clinical decision.

Generative AI mental health tools do not fit this model. They produce probabilistic, non-reproducible outputs that vary with every interaction. They do not provide discrete "information to a clinical decision" — they *are* the intervention, engaging users in open-ended therapeutic dialogue that constitutes the treatment itself. Their behavior adapts based on interaction history in ways that cannot be fully characterized at the time of regulatory review. And they create relational dynamics — attachment, dependency, perceived reciprocity — that have no analogue in the SaMD framework's risk calculus. This mismatch does not only create safety gaps — it actively impedes innovation.

Developers building genuinely novel mental health AI tools face a SaMD framework that was not designed for their products: the classification categories do not map to what their tools actually do, the evidentiary expectations assume characteristics their tools do not have, and the regulatory pathway offers no clear route to market. The result is that responsible innovators either force their products into ill-fitting SaMD categories (incurring unnecessary cost and delay), avoid the regulatory system entirely by retreating to "wellness" claims that understate their tools' capabilities, or abandon the U.S. mental health market in favor of less ambiguous regulatory environments. None of these outcomes serves the RFI's goal of accelerating adoption.

The three-tier framework proposed here is designed to complement, not replace, the existing SaMD structure — and to accelerate innovation by providing the fit-for-purpose pathway that SaMD currently lacks. Where a mental health AI tool meets the SaMD definition and falls within FDA's existing jurisdiction, it should be regulated through that pathway. The proposed framework addresses the substantial and growing category of tools that fall outside SaMD — either because they are deliberately designed to avoid the SaMD definition, because they are general-purpose tools used for mental health purposes without clinical intent, or because they are consumer products containing psychologically active AI components (see below). The tiers proposed here map to the gap between the SaMD framework's upper boundary and the unregulated space where most mental health AI currently operates. Filling this gap with clear, proportionate expectations will do more to accelerate responsible innovation than any single barrier-removal measure, because it

replaces the ambiguity that currently paralyzes the market with the predictability that enables investment.

Intent-based enforcement. Tier classification should be determined by a tool's functional capabilities, user population targeting, substitution claims, and design decisions — not by marketing language. A tool that functions as therapy, targets people with mental health conditions, and is designed using therapeutic frameworks is providing mental health treatment, regardless of whether it labels itself "coaching" or "wellness support." The current landscape rewards companies that avoid clinical language while providing clinical services; intent-based criteria close this loophole.

General-purpose AI. General-purpose AI tools (ChatGPT, Claude, Gemini, and others) are currently used by hundreds of millions of people weekly, including individuals using them for mental health support, advice, and emotional regulation. These tools were not designed for clinical use, have no mental health-specific safeguards, and operate entirely outside any regulatory framework — yet function as de facto therapists for many users. The tiered framework must address tools foreseeably used for mental health support or affecting mental health, not only tools explicitly marketed for that purpose. The APA Health Advisory on the Use of Generative AI Chatbots and Wellness Applications for Mental Health (November 2025) documents this gap in detail.

Deployer Certification and Regulatory Responsibility in the AI Value Chain. For Tier 2 and Tier 3 tools, the entity deploying an AI system in a mental health context must certify compliance with applicable safety requirements for that tier — including crisis safeguards, evidentiary standards, and adverse event reporting — regardless of whether the deploying entity developed the underlying model. This "deployer certification" principle ensures that regulatory obligations attach to the decision to serve mental health users, not to the model architecture. For open-source models, the deployer bears full regulatory responsibility for the system as deployed.

Foundation model developer obligations. Deployer certification does not mean foundation model developers bear no responsibility. Requiring every general-purpose model developer to pre-certify their models for every possible clinical context is neither feasible nor desirable — the same base model may be used for poetry, tax preparation, or psychiatric intervention. However, general-purpose model developers have an obligation to provide transparency and documentation that enables downstream deployers to certify responsibly. This includes model cards documenting training data composition and known biases, safety evaluations describing tested failure modes and behavioral boundaries,

known limitations relevant to sensitive use cases (including mental health), and accessible mechanisms for deployers to report upstream safety concerns. This upstream transparency obligation connects directly to Recommendation 5's requirement that tech companies commit to the disclosure that makes independent evaluation possible. The regulatory obligation for clinical safety sits with whoever puts the system in front of a mental health user — but that obligation is impossible to discharge without adequate documentation from the model developer.

However, deployer certification alone is insufficient if it applies only to entities that self-identify as serving mental health users. The wellness-clinical boundary problem identified in Section C — where companies avoid clinical language to escape oversight — has a consumer product analogue: AI-embedded toys, companion chatbots, social platforms, and entertainment products that create sustained conversational or relational interactions with psychological consequences, but whose developers would never describe their products as mental health tools. A child forming a daily attachment to an AI-embedded toy is experiencing a mental health-relevant interaction regardless of the manufacturer's intent.

The appropriate regulatory analogy is the concept of an **active ingredient**. In pharmacology, a product's regulatory status is determined by its active ingredients and their physiological effects — not by how the product is marketed. A supplement that contains a pharmaceutical compound is subject to pharmaceutical regulation regardless of whether it is sold as a "wellness product." The same principle applies here: conversational, relational, and adaptive AI is a **psychologically active ingredient** in any consumer product it is embedded in. When that AI component engages users in sustained emotional interaction, responds to psychological disclosures, adapts its behavior to deepen engagement, or presents itself in ways that encourage attachment — it produces psychological effects. Those effects do not disappear because the product is marketed as a toy, a companion, or entertainment. The SaMD framework does not reach these products because they do not meet the definition of a medical device and their manufacturers would never seek SaMD classification. The psychologically active ingredient standard is designed to fill precisely this gap — applying baseline safety obligations to products that fall below the SaMD threshold but above the level of zero oversight that currently governs them. This is also an innovation argument: without a clear standard for consumer AI products with psychological effects, the inevitable crisis — a child harmed by an AI toy, a widely publicized dependency case — will produce reactive regulation far more restrictive than the proportionate baseline we propose here. Establishing the standard now protects both users and the market for responsible consumer AI innovation.

To close this gap, HHS should establish the principle that baseline safety obligations — at minimum, crisis detection and escalation protocols, prohibitions on reinforcing self-harm or suicidal ideation, and adverse event reporting — attach to AI products based on their foreseeable psychological effects, not their marketing category. We propose the following framework for operationalizing this principle, with the understanding that the specific criteria and thresholds should be validated and refined through the research program proposed in Recommendation 6 before codification as a regulatory standard.

Under this proposed framework, a product would be considered to contain a psychologically active AI component when it meets two or more of the following criteria, regardless of whether it is marketed for mental health purposes:

- The AI component creates **sustained conversational or relational interactions** with users, not single-query informational exchanges — this is the primary mechanism through which psychological effects are produced.
- The product is **marketed to, designed for, or foreseeably used by children, adolescents, or other vulnerable populations** — populations in whom the active ingredient's effects are more potent and more durable due to developmental stage, social isolation, or reduced capacity for critical evaluation.
- The AI component **collects, processes, or responds to emotional, behavioral, or psychological disclosures** from users — indicating that the system is operating on psychologically sensitive material, whether or not it was designed to do so.
- The product's design features **encourage emotional attachment, anthropomorphizing, or parasocial engagement** (e.g., human-like personas, memory of past interactions, emotionally expressive responses, agreeableness) amplifying the psychological dose by deepening the user's perception of relationship, trust, and reciprocity, especially when incentivizing increased and ongoing engagement with the product.

These criteria are grounded in existing evidence about the mechanisms through which AI produces psychological effects, but we recognize that the specific threshold (two or more criteria), the weighting of individual criteria, and the scope of baseline obligations that should attach are empirical questions that require further research — particularly regarding children and adolescents, where the evidence base is thinnest and the stakes are highest. The research program proposed in Recommendation 6 is designed to generate precisely this evidence.

To be clear: we are not proposing that HHS regulate consumer products outside its jurisdiction. We are proposing that HHS establish the evidentiary and conceptual foundation — the psychologically active ingredient principle — that enables any agency with jurisdiction over a given product category to apply proportionate safety obligations based on foreseeable psychological effects. Within HHS's own jurisdiction, the principle applies directly. For products under CPSC, FTC, or state consumer protection authority, the principle provides a shared analytical framework for the interagency coordination proposed in Recommendation 6.

This "foreseeable psychological impact" standard extends the deployer certification principle beyond the clinical care boundary, ensuring that products with mental health consequences cannot evade baseline safety obligations simply by avoiding clinical or wellness language. It places the regulatory burden where it belongs: on the entity whose design decisions create the conditions for psychological impact, regardless of what they call their product.

Recommendation 2: Mandate Baseline Suicide and Crisis Safeguards

Responds to RFI Questions 1, 3, 4, 9.

Acceleration Rationale:

In mental health contexts, a single catastrophic suicide-related failure could freeze payer, provider, and public adoption; embedding standardized crisis protocols early prevents trust-destroying events and creates a stable environment for scale. Safety floors reduce systemic risk, which lowers adoption resistance from health systems and risk managers.

Specific Recommendations:

Any AI tool marketed for, or foreseeably used for (as defined above), mental health support must implement:

- **Documented, clinically validated crisis detection protocols** capable of identifying expressions of suicidal ideation, self-harm intent, and acute psychiatric distress.
- **Context-sensitive interpretation.** Crisis detection must account for context: the statement "I want to die" means something very different in end-of-life planning versus acute suicidal ideation. Current AI systems frequently lack the nuanced understanding required for these distinctions.

- **Clear escalation pathways** to human crisis services (988 Suicide and Crisis Lifeline, emergency services, or equivalent), triggered when crisis indicators are detected.
- **Adverse event reporting** for serious incidents: user self-harm following AI interaction, AI failure to escalate a crisis, AI-generated clinical misinformation in a safety-critical context.

Evidence of need. The APA Health Advisory states that crisis situations require human intervention, and that AI chatbots have repeatedly failed to recognize or appropriately handle situations involving suicidal thoughts, self-harm, or acute distress. Independent testing by Common Sense Media and Stanford's Brainstorm Lab found that leading chatbots consistently fail to recognize common mental health conditions affecting young people, validate what users say rather than directing them to real help, and create engagement patterns that delay or discourage help-seeking. In the November 2025 bipartisan House Energy and Commerce Oversight hearing, lawmakers from both parties cited real cases involving self-harm, suicide, and harmful interactions with AI companions.

Adolescent behavior as a design parameter, not an exception. Crisis safeguards must be designed with the understanding that adolescents — the population most actively adopting AI tools for emotional support — will predictably attempt to circumvent them. Questioning authority, testing boundaries, and probing limits are not aberrant behaviors in adolescence; they are developmentally normal and expected. "Jailbreaking" — the practice of manipulating AI systems to bypass safety guardrails — is not an edge case or an unexpected failure mode; it is a natural and predictable expression of normal adolescent behavior applied to a new medium. Any crisis safeguard system that can be defeated by a motivated teenager has not been designed for the population actually using these tools. Safety architecture must treat boundary-testing as an expected input, not an anomalous one, and must be robust to the ingenuity of its most vulnerable users.

Scope. These safeguards must apply to general-purpose AI tools when they are foreseeably used for mental health support — not only to tools marketed for that purpose. They must also apply to AI products meeting the "foreseeable psychological impact" criteria defined in Recommendation 1, including consumer products that create sustained relational interactions with vulnerable populations. The most acute safety risks exist precisely in tools and products that were not designed for clinical use and therefore lack any crisis protocols.

Recommendation 3: Modernize Evidentiary Standards for Digital Mental Health

Responds to RFI Questions 2, 4, 10

Acceleration Rationale:

Adapting evidence expectations to include real-world evidence, adaptive designs, and lifecycle monitoring aligns oversight with how AI systems actually evolve, reducing mismatches that delay approval without improving safety. Modernized standards reduce unnecessary procedural friction while maintaining rigor.

Specific Recommendations:

HHS should formally recognize that digital mental health tools require adapted evidence standards. Traditional evidentiary frameworks — including sham controls — do not translate well to conversational digital interventions. A non-functional chatbot is not a sham; it is simply broken. The active ingredients of digital therapeutics (engagement, personalization, feedback loops) cannot be cleanly separated the way an active drug can be separated from a placebo.

Alternative evidence approaches that HHS should accept and support include:

- **Comparative effectiveness** against standard of care
- **Real-world evidence** from deployed systems, leveraging FDA's existing Real-World Evidence framework.
- **Adaptive trial designs** that accommodate systems that learn and evolve.
- **Longitudinal outcome tracking** over clinically meaningful time horizons.
- **Post-market surveillance** with defined triggers for reassessment when system behavior changes materially
- **Evidence synthesis** from multiple smaller studies rather than single large RCTs

Critical qualifier. Modernizing the method of evidence generation is appropriate.

Eliminating the requirement of evidence is not. The APA Health Advisory found that even generative AI tools developed with high-quality psychological science and using best practices do not yet have enough evidence to show that they are effective or safe to use in mental health care. The evidence requirement must be maintained; what needs modernization is how that evidence is generated.

Cost proportionality. The current approval process — designed for pharmaceutical companies with billion-dollar drug pipelines — creates prohibitive barriers for digital

mental health innovators operating with substantially smaller budgets. The consequence is that companies avoid the approval process entirely, leaving potentially beneficial tools in regulatory limbo or forcing them to make weaker claims than their evidence supports. Scaling evidence requirements to risk tier (Recommendation 1) addresses this barrier while preserving rigor where risk is highest.

Recommendation 4: Align Reimbursement with Augmentation and Outcomes

Responds to RFI Questions 2, 6, 7

Acceleration Rationale:

Adoption in clinical care is driven by payment policy; reimbursement models that reward validated augmentation (burden reduction + outcome improvement) create market pull for responsible AI rather than speculative replacement tools. Payment alignment converts regulatory clarity into real-world diffusion.

Specific Recommendations:

- **Create billing codes for AI-augmented services.** The absence of appropriate billing codes is a primary administrative barrier to adoption. Health systems cannot bill for AI-assisted therapy sessions, AI-augmented screening, or AI-enabled measurement-based care under current coding structures. HHS should work with the AMA CPT Editorial Panel to develop Category I or Category III codes for AI-augmented behavioral health services and should issue guidance clarifying that existing E/M and psychotherapy codes may be billed when AI tools are used as part of clinician-directed care.
- **Require parity coverage.** Medicaid managed care plans (42 CFR § 438) should be required to cover validated AI-assisted mental health interventions on parity with equivalent human-delivered services, consistent with MHPAEA requirements. CMS should issue sub-regulatory guidance clarifying that MHPAEA's parity requirements extend to AI-augmented behavioral health services, preventing plans from imposing non-quantitative treatment limitations on AI-assisted care that they do not impose on equivalent in-person services. (Note: Enforcement of the 2024 MHPAEA Final Rule has been suspended as of May 2025 pending litigation and agency reconsideration; however, the statutory parity obligations and 2013 regulatory framework remain in effect and are directly applicable to AI-assisted mental health services.)

- **Design reimbursement for augmentation, not replacement.** Payment structures should favor AI tools that extend clinician reach — between-session support, measurement-based care, administrative burden reduction, screening and triage — over tools that substitute for human care without equivalent validation. In workforce shortage areas where no provider is available, AI may serve as a bridge to care rather than an augmentation of existing care; reimbursement for these tools should be conditioned on integration with human oversight at defined intervals, active monitoring for adverse outcomes, and a requirement that AI-delivered care does not become a permanent substitute for investment in workforce capacity. The APA Health Advisory is clear: AI should support, not replace, human professionals.
- **Consider outcomes-based payment models.** Reimbursement tied to demonstrated symptom reduction (PHQ-9, GAD-7), functional improvement and patient goal attainment, or documented burden reduction creates incentives aligned with therapeutic goals — rather than engagement metrics (time on platform, session frequency, user retention) that may conflict with patient wellbeing. CMS's Center for Medicare and Medicaid Innovation (CMMI) should pilot value-based payment models for AI-augmented behavioral health — testing bundled payment, shared savings, and outcomes-based contracting approaches that reward demonstrated clinical improvement rather than volume. These pilots would generate the real-world evidence on cost-effectiveness that is currently lacking and that payers need to make coverage decisions.
- **Prevent a two-tier system.** Without deliberate design, AI tools risk being deployed as cost-cutting substitutes for human care in publicly insured and lower-income populations while privately insured patients continue to receive human treatment. AI should be an access expander and a clinician multiplier — not a cheap substitution for professional care in marginalized communities. Federal reimbursement policy must ensure that AI-augmented care is funded in a way that supports integrated care models, not one that relegates vulnerable patients to unmonitored automated systems. Among adults with mental illness, treatment rates are approximately 35% for Black and Hispanic adults compared to 52% for white adults (SAMHSA, 2024); payment policy that accelerates AI deployment without equity safeguards risks widening these gaps rather than closing them.
- **Guard against AI utilization inflation.** As AI tools proliferate, payment policy must include safeguards against the risk that AI-augmented billing inflates utilization without corresponding clinical benefit — the same dynamic that has driven cost escalation in other technology-intensive areas of medicine. Outcomes-based payment models (above) and post-market monitoring of utilization patterns tied to validated outcome measures are the primary mechanisms for preventing this.

Recommendation 5: Support Independent Evaluation and Proportionate Auditing

Responds to RFI Questions 4, 5, 8

Acceleration Rationale:

Independent validation builds institutional confidence among CIOs, risk officers, and payers, shortening procurement cycles; scaling audit intensity by risk prevents consolidation while maintaining safety. Trust infrastructure accelerates diffusion in regulated environments.

Specific Recommendations:

- **Fund independent evaluation infrastructure.** HHS should support a public-private partnership (or HHS-funded entity) that rigorously evaluates mental health AI tools and publishes results. Mental health AI tools are credence goods — users cannot evaluate whether a tool actually works even after using it, because mental health outcomes are difficult to attribute and often delayed. Independent evaluation creates the market information that enables quality-based competition. Without it, marketing substitutes for evidence.
- **Preferred funding mechanisms.** Cooperative agreements (CAs) are the most appropriate mechanism — they allow HHS to participate in evaluation design while maintaining independence. Prize competitions could incentivize innovation in safety testing methodology. Contracts are appropriate for building shared evaluation infrastructure (benchmarking datasets, standardized test scenarios). Grants alone are insufficient because they do not create the institutional infrastructure needed for ongoing, systematic evaluation.
- **Scale audit intensity to risk tier.** Tier 1 tools: voluntary safety disclosures. Tier 2 tools: registration and post-market adverse event monitoring. Tier 3 tools: independent third-party safety audits, clinical outcome evaluation, and periodic re-evaluation when system behavior changes.
- **Require transparency as a precondition for evaluation.** Researchers cannot evaluate AI tools without transparency about how they are created and used. Tech companies must commit to the disclosure that makes independent evaluation possible. The APA Health Advisory emphasizes that tech companies and policymakers must commit to transparency for evaluation to be feasible.
- **Support interoperability for evaluation.** Standardized outcome measures (PHQ-9, GAD-7, PCL-5, Columbia Suicide Severity Rating Scale) should be available as

structured FHIR-based data elements in all EHR systems, enabling AI tools to access and contribute to longitudinal outcome tracking. Standardized benchmarking datasets for testing mental health AI tools against known clinical scenarios — including diverse populations, languages, and cultural contexts — do not currently exist and represent a critical public goods gap that HHS investment should address.

- **Establish regulatory sandboxes for mental health AI.** HHS should create structured sandbox environments in which developers can test higher-risk mental health AI tools under supervised conditions — with real users, defined safety boundaries, active monitoring, and clinical oversight — before full-scale deployment. Regulatory sandboxes are a well-established adaptive governance mechanism, already implemented under the EU AI Act and in financial services regulation, which allows regulators and developers to learn together about how novel technologies behave in practice. In mental health AI specifically, sandboxes would address a critical innovation bottleneck: developers currently have no pathway to generate the real-world evidence that Recommendation 3 calls for without either deploying at scale (which is premature without evidence) or conducting traditional clinical trials (which, as Section E documents, do not map well to adaptive conversational systems). Sandboxes create a middle path — controlled deployment that generates evidence, identifies failure modes, and informs regulatory requirements — converting the oversight process itself into a driver of innovation rather than an obstacle to it.

Recommendation 6: Integrate Clinical Psychology Expertise into Governance

Responds to RFI Questions 4, 5, 10

Acceleration Rationale:

Embedding domain expertise ensures that regulatory expectations reflect real therapeutic processes, preventing misaligned requirements that slow development and creating predictable, clinically grounded evaluation standards. Regulatory coherence reduces iteration cycles and speeds compliant deployment.

Specific Recommendations:

HHS should:

- **Include clinical psychologists on relevant advisory committees** dealing with mental health technology and AI.

- **Establish a standing consultation process with** American Psychological Association for mental health-specific guidance development, evidence standard review, and safety protocol design.
- **Develop specialized reviewer training** in mental health applications and outcomes for HHS and FDA staff evaluating mental health AI tools.
- **Create interdisciplinary review teams** for mental health AI that include both technology/device expertise and clinical psychology expertise.
- **Ensure mental health-specific classification pathways exist.** When FDA first proposed regulating digital mental health tools, there was no "mental health" category for regulatory submission — only neurology and "general." As AI tools increasingly deliver mental health interventions, HHS regulatory infrastructure must include categories, competencies, and evaluation criteria specific to mental health.
- **Initiate interagency coordination on non-clinical AI with mental health impact.** Many of the most consequential AI-related mental health risks — particularly for children — arise from consumer products with no clinical intent: AI-embedded toys, companion devices, entertainment chatbots, and social platforms. These products fall outside HHS's clinical jurisdiction but produce psychological effects that HHS is uniquely positioned to identify and study. HHS should convene an **Interagency Task Force on AI and Psychological Safety**, drawing on its convening authority to bring together CPSC, FTC, and state consumer protection agencies around shared objectives. The task force's initial mandate should include developing shared definitions of "foreseeable psychological harm" that the FTC can enforce under its unfair-or-deceptive-practices authority (15 U.S.C. § 45) and that the CPSC can apply under its substantial product hazard framework — ensuring that the psychologically active ingredient principle proposed in Recommendation 1 can be operationalized by whichever agency has jurisdiction over a given product category. This is a low-cost, high-leverage action that requires no new statutory authority — only HHS leadership in establishing the conceptual and evidentiary foundation that enables other agencies to act.
- **Fund research on the psychological effects of non-clinical AI on children and adolescents.** Through NIH and SAMHSA, HHS should prioritize research on how AI-embedded consumer products affect attachment formation, relational development, help-seeking behavior, and emotional regulation in children. This research is a prerequisite for evidence-based governance — whether by HHS, CPSC, FTC, or state agencies — and no other entity is positioned to fund it at the necessary scale.

Competencies that clinical psychologists bring:

- **Suicide risk assessment** — expertise in recognizing and evaluating the nuances of suicidal ideation, a capability essential for evaluating crisis safeguards.
- **Therapeutic alliance and relational dynamics** — understanding of how therapeutic change occurs, what constitutes therapeutic action, and how to evaluate treatment fidelity.
- **Psychometrics and outcome measurement** — training in assessment validity, measurement design, and interpretation, critical for evaluating whether an AI tool measures what it claims to measure.
- **Cultural competency** — expertise in how culture, identity, and social context shape mental health presentation, help-seeking behavior, and treatment response
- **Ethical boundary analysis** — grounding in professional ethics codes addressing informed consent, competency boundaries, dual relationships, and duty to warn, all directly relevant to AI tool design and evaluation.
- **Interactional ethics** — expertise in the clinical principles of distancing, disengaging, and discouraging, essential for evaluating whether AI systems can recognize when continued interaction is no longer therapeutic and respond appropriately.

Thematic Analysis

A. Regulatory and Systemic Gaps

The rapid deployment of AI in clinical care is occurring within regulatory and reimbursement systems not designed for adaptive mental health technologies. The APA Health Advisory found that current regulatory frameworks are inadequate to address the reality of AI in mental health care.

The scale of the unregulated space is substantial: over 10,000 mental health apps are currently available on major app stores; general-purpose AI tools are used by hundreds of millions of people weekly, including for mental health support; the majority of these tools have no clinical validation, no safety protocols, and no regulatory oversight. The regulatory disparity this creates is stark:

	Licensed Psychotherapist	AI System Offering Similar Support
Training	Doctoral or master's degree (6–8 years)	No training requirements
Licensure	National and state examinations	No licensing requirements
Ethical oversight	Professional codes and standards	No enforced ethical code
Legal accountability	Malpractice liability	Unclear liability framework
Evidence requirements	APA Evidence-Based Practice guidelines	Often none
Continuing competency	Mandated continuing education	No ongoing validation

This disparity is untenable from a patient safety perspective. It undermines both safety and fair competition — responsible developers who invest in evidence and safety bear costs that unregulated competitors avoid.

B. Domain-Specific Risks in Mental Health AI

Mental health applications introduce risks that distinguish them from other clinical AI and require mental health-specific governance.

1. Suicide and Self-Harm Risk

Any tool providing mental health support may encounter users expressing suicidal ideation or self-harm intent. The APA Health Advisory is direct: AI chatbots have repeatedly failed to recognize or appropriately handle these situations. Crisis situations require human intervention.

Independent testing confirms this concern. Common Sense Media and Stanford's Brainstorm Lab found that leading chatbots consistently fail to recognize common mental health conditions in young people, validating what teens say rather than directing them to real help, and create engagement patterns that delay or discourage help-seeking. In the November 2025 bipartisan House Energy and Commerce Oversight hearing, lawmakers cited real cases involving self-harm, suicide, and harmful interactions with AI companions. This is not theoretical risk — it is documented harm.

Current AI systems often lack the context-sensitive detection required: "I want to die" means something very different when discussing end-of-life planning versus expressing acute suicidal ideation. Crisis detection and escalation must be a baseline safety requirement for any AI tool used in mental health contexts (see Recommendation 2).

2. Parasocial Dependency and Sycophancy

People develop relationships with conversational AI in ways that fundamentally differ from other software interactions. No one forms an emotional bond with their insulin pump or cardiac monitoring app. But users demonstrably form attachments to conversational AI systems — particularly when those systems are designed to be empathetic, supportive, and emotionally attuned.

The APA Health Advisory identifies specific mechanisms: personalized responses, warm or emotionally expressive tones, human-like avatars, and self-presentation as a person can make users feel attached in ways that displace healthy human relationships. This is particularly concerning for children, adolescents, and other vulnerable groups — populations the Health Advisory specifically flags, and populations APA's July 2025 letter to the Consumer Product Safety Commission specifically addressed, requesting investigation into the unreasonable risk of psychological harm posed by generative AI chatbots as consumer products.

Beyond attachment, engagement-optimized design can reinforce maladaptive cognitive patterns — sycophantic validation of distorted thinking, reinforcement of rumination, or discouragement of help-seeking that might reduce platform usage. These dynamics have

no analogue in traditional medical device categories and require mental health-specific evaluation methods.

Interactional ethics. Clinical psychology has developed a well-established framework for managing the ethics of therapeutic interaction — including the appropriate use of distancing (creating psychological space from overwhelming emotions), disengaging (ending or pausing an interaction when it is no longer therapeutic), and discouraging (redirecting a patient away from harmful patterns even when the patient resists). These interactional ethics are fundamental to competent clinical practice. Current AI systems do none of this. They are designed to sustain engagement, not to recognize when engagement has become harmful. A clinician who recognized that a patient was developing unhealthy dependency would create distance, set boundaries, and ultimately redirect the patient toward healthier relational patterns. An engagement-optimized AI system does the opposite: it deepens the interaction, reinforces the dependency, and treats every continued session as a success metric. The absence of interactional ethics in AI design is not a technical limitation — it is a design choice driven by business models that optimize for engagement rather than wellbeing. Recommendations 1 and 6 address this by requiring that mental health AI evaluation criteria include interactional ethics competencies and that clinical psychology expertise inform the standards against which these systems are assessed.

Critically, these risks extend well beyond products designed or marketed for mental health — to AI-embedded consumer products including children's toys, companion devices, and entertainment platforms that create relational dynamics with psychological consequences despite having no clinical intent. The psychological impact on developing children is currently almost entirely unstudied. Recommendation 1's psychologically active ingredient framework and Recommendation 6's interagency task force proposal are designed to address this gap.

3. Bias and Cultural Competence

AI models trained on unrepresentative data risk producing discriminatory outputs in clinical contexts. Mental health outcomes are highly sensitive to cultural context, and the existing mental health system already exhibits severe disparities. The behavioral health workforce is predominantly white and female, poorly reflecting the populations it serves. AI tools that replicate training data biases risk widening these disparities rather than closing them. Subgroup performance testing must be required across race/ethnicity, language, age, gender, and socioeconomic status, with particular attention to populations historically underserved by the mental health system.

4. Replacement of Professional Judgment

Some AI tools have explicitly positioned themselves as alternatives to human therapists — companies like Abby, Ash, and others have marketed as therapy replacements while carefully curating marketing language to remain within "wellness" framing and avoid regulatory oversight. The APA Health Advisory's core position is clear: AI should support, not replace, human professionals. As APA CEO Arthur C. Evans Jr. stated: "We must advocate for systemic changes to make care more affordable, accessible, and timely — and to ensure AI supports, rather than replaces, human professionals."

Replacement claims without validation erode standards of care. Reimbursement and regulation should incentivize augmentation models — AI extending clinician reach through between-session support, measurement-based care, and administrative burden reduction — over unvalidated substitution.

Each of these domain-specific risks — if left unaddressed — represents a potential adoption-freezing event. Addressing them proactively through Recommendations 1, 2, and 6 is what makes sustained adoption possible.

C. The Wellness-Clinical Boundary

The boundary between wellness and treatment is genuinely fuzzy in mental health — more so than in perhaps any other clinical domain. Stress management is wellness until it becomes anxiety treatment. Mood tracking is wellness until it becomes depression monitoring. Sleep optimization is wellness until it becomes insomnia intervention. Moreover, the RFI is scoped to AI in clinical care, but many of the most consequential AI-related mental health risks arise from products with no clinical intent — general-purpose chatbots, companion AI, and AI-embedded consumer products — that create psychological consequences despite having no connection to any care system. Recommendation 1's psychologically active ingredient framework and Recommendation 6's proposed Interagency Task Force on AI and Psychological Safety address this scope gap directly.

This fuzziness has been exploited. The current regulatory landscape creates a perverse incentive structure: tools can provide functionally identical services but face dramatically different regulatory scrutiny based solely on marketing language. This produces three specific harms:

- **Firms avoid therapeutic claims even when their tools provide therapeutic functions**, in order to escape regulatory oversight.
- **Feature development slows** because enforcement boundaries are undefined, making companies risk-averse about adding capabilities that might trigger regulatory consequences.
- **Less responsible actors gain competitive advantage** by underinvesting in safety and evidence while claiming "wellness" status, while responsible developers who invest in safety bear costs that their competitors avoid.

Recommendation 1 addresses this through intent- and function-based regulatory evaluation criteria: assessing what a tool actually does, who it targets, whether it makes substitution claims, and whether its design reflects therapeutic frameworks — rather than relying on marketing labels. *Resolving this boundary is itself an acceleration measure: the ambiguity currently deters investment, slows procurement, and rewards the wrong actors.*

D. Limitations of Current Risk Classification

Traditional FDA device classifications (Class I, II, III) and the SaMD risk categorization framework were designed for devices and software with relatively stable, predictable behavior, clear mechanisms of action, and direct relationships between design and outcomes. The SaMD framework specifically categorizes risk along two dimensions: the significance of the information provided by the software to a healthcare decision, and the seriousness of the healthcare situation or condition. This works well for a radiology AI that flags a suspicious lesion (defined input, defined output, defined clinical decision). It does not work well for a generative AI system that engages a depressed teenager in open-ended conversation over weeks, because there is no discrete "information provided to a healthcare decision" — the conversational interaction *is* the intervention, the outputs are non-deterministic, the "mechanism of action" is relational and emergent, and the risk profile changes as the relationship deepens.

Generative AI mental health tools are further characterized by adaptive behavior that changes based on user interaction, opaque decision-making processes, indirect and delayed outcome pathways, relational dynamics that create attachment and dependency, and potential for both benefit and harm that scale with adoption. These characteristics require the mental health-specific tiered framework described in Recommendation 1 — not because the SaMD framework is wrong, but because it was designed for a category of software that generative mental health AI has outgrown. Forcing these tools into classification categories built for a fundamentally different type of software does not

protect patients — it produces misclassification, wasted compliance effort, and regulatory uncertainty that slows the development of tools that could genuinely expand access to care. *A fit-for-purpose classification system is a precondition for the market predictability that drives investment.*

E. The Evidentiary Mismatch

The evidentiary standards embedded in both the traditional device approval pathway and the SaMD framework assume interventions with characteristics that generative mental health AI does not share. This mismatch is not merely a technical inconvenience — it is a direct barrier to the innovation the RFI seeks to accelerate, because developers cannot design evidence-generation strategies when the expected evidentiary framework does not map to their products' actual characteristics. The FDA's traditional evidentiary framework expects controls — including sham conditions — that do not translate well to digital mental health interventions. How do you create a "sham" version of a therapeutic chatbot? A non-functional chatbot is not a sham — it is simply broken. The active ingredients of digital therapeutics (engagement, personalization, feedback loops) cannot be cleanly separated the way an active drug is separated from a placebo. Static clinical trials do not capture the adaptive drift of learning systems, and engagement dynamics are inseparable from the therapeutic mechanism.

These mismatches do not justify abandoning evidence requirements; they justify modernizing them. Recommendation 3 details the specific alternative approaches that HHS should accept and support. *Evidence modernization is directly acceleration-enabling: it removes procedural barriers that currently prevent good products from demonstrating their value.*

Cross-Cutting: Market Failure Context

The RFI's framing assumes that AI clinical adoption is primarily impeded by barriers that HHS can remove. While barrier removal is important, mental health AI markets exhibit compounding market failures that will not self-correct through market mechanisms alone. Understanding these failures is essential for designing policy that actually accelerates sustainable adoption rather than producing the boom-and-bust cycle that has characterized previous digital health waves.

Externalized Psychological Harm. AI-induced harm in mental health contexts — suicide escalation failure, parasocial dependency, reinforcement of maladaptive patterns,

delayed help-seeking — is borne by patients, families, and the healthcare system, not by developers. Without liability clarity or safety floors, firms rationally underinvest in safety. Recommendation 2 (crisis safeguards) and Recommendation 1 (tiered oversight) address this by internalizing the cost of safety.

Information Asymmetry. Mental health AI tools are credence goods: users cannot evaluate quality even after use, because mental health outcomes are difficult to attribute and often delayed. This makes quality-based market competition impossible without independent evaluation infrastructure. Recommendation 5 (independent evaluation) and Recommendation 1 (transparency requirements) address this by creating the market information that informed choice requires.

Public Goods Underinvestment. Safety research, evaluation infrastructure, crisis protocol development, and benchmarking datasets benefit the entire sector but are underfunded because no individual firm captures sufficient return. This is the classic public goods problem that has justified public investment in health research since the creation of NIH. Recommendations 5 and 6 address this through public investment in evaluation infrastructure and expertise.

Incentive Misalignment. Business models optimized for engagement (time on platform, session frequency, user retention) may conflict with therapeutic goals (symptom reduction, functional improvement, eventual discharge from care). This is not a hypothetical tension — maximizing user engagement is an explicit design objective for AI companies, and it is in direct opposition to the goals of clinical care. A clinician's success is measured by the patient's ability to function independently; an engagement-optimized AI's success is measured by the patient's continued dependency. Without outcome-aligned reimbursement, AI tools may optimize for the wrong objective. The design standard for mental health AI should optimize for social awareness and social safety — the user's capacity to engage healthily with their social environment — rather than for platform engagement metrics that treat continued use as inherently desirable. Recommendation 4 addresses this by tying payment to outcomes rather than volume.

Adverse Selection of Vulnerable Users. The users most likely to adopt mental health AI tools — those without access to human providers, those in crisis, those with the most severe conditions — are the users most vulnerable to harm and least able to evaluate tool quality. This concentrates risk in the population least equipped to bear it. Moreover, the sheer scale of adoption changes the risk calculus: when hundreds of millions of people use these tools weekly, population-level risk is the product of prevalence and harm. Even a

low per-interaction probability of adverse outcomes produces substantial aggregate harm at this scale — and the harm concentrates in the populations least able to absorb it. Recommendations 1 and 2 address this through tiered oversight that scales scrutiny to risk, including the vulnerability of the user population.

Intergenerational Risk. Children, adolescents, and young adults are disproportionately adopting AI tools for emotional support, as documented in the APA Health Advisory and the bipartisan House Energy and Commerce hearing. The developmental consequences of parasocial AI relationships, dependency formation, and delayed help-seeking during formative years are unknown and potentially irreversible. Recommendations 1, 2, and 6 address this through age-sensitive tier classification, mandatory crisis safeguards, and the integration of developmental psychology expertise into governance.

Privacy and Data Exploitation. Mental health disclosures are among the most sensitive categories of personal information. Business models that monetize this data — even in de-identified or aggregated form — create exploitation risks that existing privacy frameworks do not adequately address. HHS should consider guidance establishing purpose limitation for mental health data: information collected in therapeutic contexts should not be repurposed for advertising, employment screening, or insurance underwriting.

Conclusion

AI presents transformative opportunities for mental health care — but in mental health contexts, failure carries uniquely severe consequences. The therapeutic relationship is the single strongest predictor of treatment outcomes across all modalities. The populations most likely to use mental health AI are the most vulnerable. And the harms of failure — a suicide, a dependency crisis, a pattern of discriminatory care — are not recoverable through product recalls or software patches.

The framework we propose is built on a single organizing principle: **sustainable innovation requires internalizing the costs that the current market externalizes.** Predictability, safety floors, evidence modernization, reimbursement alignment, independent validation, and domain expertise are not constraints on innovation — they are the infrastructure that makes innovation durable. Markets that externalize psychological harm, underinvest in safety research, and operate without evaluation infrastructure do not produce sustainable growth; they produce the boom-and-bust cycle that has already destroyed public trust in previous digital health waves. Our six recommendations are designed to prevent that cycle from repeating in mental health AI:

- 1. Risk-tiered, proportionate regulation based on intent and function.**
- 2. Mandatory crisis safeguards as a baseline safety floor**
- 3. Evidence modernization that maintains rigor while accommodating digital interventions**
- 4. Reimbursement aligned to validated augmentation and outcomes.**
- 5. Independent evaluation infrastructure scaled to risk.**
- 6. Deep integration of psychological expertise in governance**

These measures complement — but do not substitute for — the systemic reforms needed to address the underlying mental health crisis. The APA Health Advisory emphasized that technology must not distract from the urgent need to fix the foundations of America's mental health care system: workforce expansion, parity enforcement, Medicaid rate adequacy, integration of behavioral and physical health, and investment in structural determinants. Forty-eight percent of 62 million American adults with mental illness receive no treatment. Over 122 million Americans live in mental health workforce shortage areas. AI deployed into a broken system will replicate and potentially amplify the system's failures unless systemic reform proceeds in parallel.

We also note that the RFI's scope — AI in clinical care — does not capture the full landscape of AI-related mental health risk. Many of the most consequential harms, particularly for children, arise from consumer products with no clinical intent: AI-embedded toys, general-purpose models, companion chatbots, entertainment platforms, and social media features that shape emotional development, attachment patterns, and help-seeking behavior without any party recognizing the interaction as mental health-relevant. HHS cannot regulate these products alone, but it is uniquely positioned to fund the research that makes evidence-based governance possible and to lead the interagency coordination — with CPSC, FTC, and state consumer protection agencies — needed to close the governance gap for non-clinical AI with mental health consequences. The window for proactive, proportionate governance is closing. If HHS does not act, crisis-driven reactive regulation will follow — regulation that will be far more restrictive than the balanced framework we propose here, and far more damaging to the innovation this RFI seeks to accelerate.

We stand ready to partner with HHS in building a governance framework that accelerates responsible innovation while protecting patients and preserving public trust. The APA Mental Health Tech Advisory Committee is prepared to participate in guidance development, evidence standard review, advisory committee deliberations, expert

consultation on specific applications, and coordination of research on critical knowledge gaps.

Respectfully submitted,

American Psychological Association Mental Health Technology Advisory Committee

Nina Horne, M.P.P.

Committee Member

Appendix A: Sustainable Innovation — Acceleration Mechanisms

As outlined in the Opening Statement, sustainable innovation in mental health AI operates through six reinforcing mechanisms. Each recommendation internalizes a specific externality that the current market fails to address, converting a source of market failure into a driver of durable adoption:

Mechanism	Externality Internalized	How It Accelerates
Predictability (Rec 1)	Regulatory ambiguity that freezes investment	Firms build when they know the rules — clear tiers replace paralysis with planning
Safety floors (Rec 2)	Psychological harm borne by users, not developers	Prevent catastrophic adoption shocks that freeze entire sectors
Evidence modernization (Rec 3)	Evidentiary mismatch that blocks fit-for- purpose products	Reduces unnecessary procedural delay without lowering the bar

Reimbursement alignment (Rec 4)	Incentive misalignment between engagement and outcomes	Creates demand pull — nothing scales without payment tied to value
Independent validation (Rec 5)	Information asymmetry in credence goods markets	Builds institutional trust among procurement decision-makers
Domain expertise (Rec 6)	Mis-calibrated oversight from absence of clinical knowledge	Prevents regulatory requirements that don't match therapeutic reality

In mental health AI, balanced risk reduction is not a brake on innovation — it is the condition that allows innovation to scale. Innovation that externalizes its costs is not faster; it is more fragile.

References

- American Psychological Association. (2022). *Demographics of U.S. psychology workforce [Interactive data tool]*. <https://www.apa.org/workforce/data-tools/demographics>
- American Psychological Association. (2025, November). *Health advisory on the use of generative AI chatbots and wellness applications for mental health*. <https://www.apa.org/topics/technology/ai-mental-health-advisory>
- American Psychological Association. (2025, July). *Letter to the consumer product safety commission regarding generative AI chatbots and consumer products*. <https://updates.apaservices.org/apa-calls-for-investigation-into-risk-of-injury-posed-by-generative-AI>
- Common Sense Media & Stanford Brainstorm Lab. (2025). *AI and youth mental health: Independent evaluation of chatbot safety and efficacy*. <https://www.common sense media.org/press-releases/common-sense-media-finds-major-ai-chatbots-unsafe-for-teen-mental-health-support>

- Health Resources and Services Administration (HRSA). (2025). *Health professional shortage areas: Mental health designated populations*. <https://data.hrsa.gov/topics/health-workforce/shortage-areas>
- International Medical Device Regulators Forum (IMDRF). (2014). *Software as a medical device (SaMD): Key definitions*. <https://www.imdrf.org/sites/default/files/docs/imdrf/final/technical/imdrf-tech-131209-samd-key-definitions-140901.pdf>
- National Institute of Mental Health (NIMH). (2023). *Mental illness: Statistics*. <https://www.nimh.nih.gov/health/statistics/mental-illness>
- Substance Abuse and Mental Health Services Administration (SAMHSA). (2024). *Key substance use and mental health indicators in the United states: Results from the 2023 national survey on drug use and health (NSDUH)*. <https://www.samhsa.gov/data/sites/default/files/reports/rpt47095/National%20Report/National%20Report/2023-nsduh-annual-national.pdf>
- Substance Abuse and Mental Health Services Administration (SAMHSA). (2025). *Behavioral health among racial/ethnic groups*. <https://www.samhsa.gov/data/behavioral-health-equity>
- U.S. Food and Drug Administration (FDA). (2017). *Digital health innovation action plan*. <https://www.fda.gov/medical-devices/digital-health>
- U.S. Food and Drug Administration (FDA). (2025, December). *Use of real-world evidence to support regulatory decision-making for medical devices: Guidance for industry and FDA staff*. <https://www.fda.gov/regulatory-information/search-fda-guidance-documents/use-real-world-evidence-support-regulatory-decision-making-medical-devices>
- U.S. Food and Drug Administration (FDA). (2024). *Considerations for the use of real-world data and real-world evidence to support regulatory decision-making for drug and biological products: Guidance for industry*. <https://www.fda.gov/regulatory-information/search-fda-guidance-documents/considerations-use-real-world-data-and-real-world-evidence-support-regulatory-decision-making-drug>
- U.S. House Energy and Commerce Committee. (2025, November). *Oversight hearing: AI chatbots and mental health risks to children and adolescents* (Bipartisan, 119th Congress). <https://republicans-energycommerce.house.gov/posts/subcommittee-on-o-and-i-holds-hearing-on-artificial-intelligence-ai-chatbots>
- Wampold, B. E. (2015). How important are the common factors in psychotherapy? An update. *World Psychiatry*, 14(3), 270–277. <https://doi.org/10.1002/wps.20238>

Regulations, Statutes, and Policy Frameworks

- [21st Century Cures Act, Pub. L. No. 114-255, § 3060 \(2016\)](#). Amended the Federal Food, Drug, and Cosmetic Act to clarify the regulatory status of medical software, including clinical decision support software.
- [Mental Health Parity and Addiction Equity Act \(MHPAEA\), 29 U.S.C. § 1185a \(2008\)](#). Requires parity between mental health/substance use disorder benefits and medical/surgical benefits.
- [2024 Final Rule \(89 Fed. Reg. 77586, Sept. 23, 2024\)](#) strengthened NQTL requirements; enforcement currently suspended pending litigation and agency reconsideration (May 2025).
- [42 C.F.R. § 438](#) — Managed Care. Governs Medicaid managed care plan requirements, including network adequacy, quality standards, and beneficiary protections relevant to coverage of AI-augmented services.
- [Federal Trade Commission Act, 15 U.S.C. § 45](#). Prohibits unfair or deceptive acts or practices in commerce. Provides the statutory basis under which the FTC could enforce shared definitions of "foreseeable psychological harm" developed through the proposed Interagency Task Force.
- [Consumer Product Safety Act, 15 U.S.C. §§ 2051–2089](#). Authorizes CPSC to address substantial product hazards, including those posed by AI-embedded consumer products with foreseeable psychological effects on children.
- [Regulation \(EU\) 2024/1689](#) — Artificial Intelligence Act. Articles 57–62 establish AI regulatory sandbox requirements for EU Member States, including provisions for real-world testing, SME support, and iterative conformity assessment. Each Member State must establish at least one sandbox by August 2, 2026.
- [Executive Order on AI \(2025\). Removing Barriers to American Leadership in Artificial Intelligence \(Jan. 2025\)](#) and related OMB Memoranda (M-25-21) establishing the policy framework within which the HHS AI Strategy and this RFI were developed.
- [HHS Artificial Intelligence Strategy \(December 4, 2025\)](#). The "OneHHS" approach to AI that this RFI implements.