

In response to the Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care (Published Document: 2025-23641 (90 FR 60108)) by the US Department of Health and Human Services, Office of the Secretary (RIN 0955-AA13), [Reveleer](#) is submitting the following information for consideration to the specific questions outlined in the published document.

Please direct any questions to the responses below to alan.tam@reveleer.com.

1. Biggest barriers to AI innovation and adoption in clinical care

From Reveleer's work with health plans and risk-bearing providers in value-based care, the primary barriers are:

- Fragmented data and limited interoperability. Despite interoperability policies, clinical and administrative data remain scattered across an array of systems EHRs, claims platforms, HIEs, point solutions, and document repositories, with critical signal buried in unstructured notes. There is some adoption of transmission standards through standard FHIR based APIs but this only covers part of the data needs required for AI to comprehensively and reliably gather the context for an accurate response. Given the variety of formats driving the source data, when AI is deployed on this landscape without first standardizing and normalizing the inputs, it tends to amplify gaps, increase hallucination risk, and produce inconsistent outputs that clinicians do not trust.

- Lack of an AI-ready evidence layer. Many organizations still depend on partial claims feeds, limited EHR connections, and manual chart retrieval, preventing the creation of longitudinal, provenance-managed member records that can support risk adjustment, quality, member eligibility, and care management consistently across the care continuum.
- Insufficient governance, explainability, and auditability. Leaders are wary of “black box” AI, hallucinations, bias, and opaque decision paths—especially where AI touches coding, quality reporting, or point-of-care decision support. Many offerings lack strong MLOps, traceability to source documents, or clear human-in-the-loop controls. Explainable AI requires new approaches like our “clinical corpus” model that delivers reliable medical evidences in a language familiar to clinicians and auditors to reduce the cognitive load driving administrative burden. By using this approach, evidence is clear and defensible.
- Misaligned incentives and unclear measurement. Value-based contracts increasingly emphasize digital quality, equity, and downside risk, but organizations lack standardized outcome and ROI frameworks to compare AI-assisted versus human-only workflows. This uncertainty slows investment and scale. By using incentive management solutions like our Care Gap Manager, value alignment streamlines the coordination between payer and provider to focus on patient outcomes leading to lower costs for all.
- Workflow disruption and provider abrasion. AI tools that require new portals, double documentation, or low-precision suspect lists create extra work and “audit anxiety,” leading clinicians and coders to disengage even when underlying models are accurate. AI prediction and recommendation engines like our Hybrid AI, are built to support the placement of high value clinical next best

actions in the right interface at the right time. EMRs must allow these tools into their ecosystems so that the richest and most reliable insight can drive the best patient outcome.

- Regulatory uncertainty and overlapping oversight. Rapidly evolving rules around interoperability, PHI use, and expectations for transparency and safety create perceived legal and compliance risk, particularly for non-device clinical AI that does not fit neatly into existing FDA, CMS, HHS, and TECCA paradigms.

As noted above, the Reveleer EVE framework –a predictive, explainable AI layer operating over a unified evidence foundation–is designed to mitigate these barriers, but they remain systemic obstacles to broader private-sector innovation and adoption.

2. Regulatory, payment, and programmatic changes HHS should prioritize

Reveleer recommends that HHS prioritize the following changes to accelerate effective, safe AI use in clinical care, especially for non-device AI used in risk adjustment and quality workflows.

2.1. Codify expectations for explainable, evidence-linked AI

- Establish industry-wide guidance that clinical AI used in coding, quality measurement, and decision support must provide traceable links from each recommendation back to specific clinical evidence (e.g., note, lab, report) and maintain reproducible audit trails.
- Encourage architectures that explicitly separate evidence extraction (which may use generative or agentic techniques) from deterministic clinical reasoning to establish a clear boundary separating AI from diagnoses generation or authoritative narrative

content but instead organizes existing evidence and suggestions for human review. This is core to Reveleer's design of EVE Hybrid AI, which combines two sequenced services—agentic evidence extraction and clinician-authored rules in a human managed clinical corpus—to produce audit-ready suspects.

Possible regulatory anchors (illustrative):

- Under 42 CFR part 422 (Medicare Advantage) and 42 CFR part 438 (Medicaid managed care), clarify that AI-assisted risk adjustment and quality workflows are acceptable when diagnoses and quality elements meet existing documentation standards, maintain traceability to source records, and are validated by human experts.
- Through ONC's health IT rulemaking, align deregulatory steps on certification with explicit expectations around standards-based APIs, provenance metadata, and logging sufficient to support explainable AI.

2.2. Align reimbursement and program design with AI-enabled workflows

- Recognize AI-assisted digital abstraction, prospective risk suspecting, and real-time quality gap closure as acceptable methods of documentation and reporting under CMS quality programs (e.g., eQMs, Star Ratings, HEDIS-aligned measures), so long as evidence and audit standards are met.
- Create value-based payment model options (e.g., new Medicare Advantage and ACO model tests through CMMI) that explicitly reward investment in AI-enabled evidence layers and prospective risk adjustment—rather than only retrospective chart review—by linking shared savings or bonuses to improvements in risk accuracy, quality performance, and avoidable utilization attributable to

AI-enabled workflows. Then manage those incentives through certified VBC incentive management applications like Reveleer Care Gap Manager which is designed specifically to guide providers and plans through annual VBC programs to ensure both parties understand and agree on the program performance and outcome each year.

Specific references (non-exhaustive and illustrative):

- 42 CFR part 422 Subpart G (Payments to Medicare Advantage Organizations): clarify that AI-assisted coding and suspecting are permissible when diagnoses conform to existing medical record documentation rules and remain subject to RADV and other audit requirements.
- 42 CFR parts 412 and 413: acknowledge AI-enabled documentation support as a form of clinical decision support and workflow optimization, without treating these tools as autonomous decision makers.

2.3. Strengthen, but simplify, interoperability

- Maintain and refine requirements that payers and providers expose standardized APIs (e.g., FHIR-based patient access and provider directory APIs under CMS Interoperability and Patient Access rules) while avoiding overly prescriptive certification that impedes innovation. Emphasize outcome-oriented metrics such as latency, uptime, and data completeness that directly enable safe AI.
- Encourage payer and provider adoption of interoperable evidence layers that support reuse of curated clinical data across risk, quality, and care management programs, not just passive “warehousing.”

- Require EHRs to make it easy for third party AI solutions to provide enriched evidence, suggestions, and potential diagnosis to relevant places within the provider workflow to enable choice of vendor for the best clinical insights and guidance.

2.4. Create unified expectations for non-device AI oversight inside HHS

- Issue cross-agency guidance clarifying the roles of ONC, CMS, OCR, OIG, and FDA for non-device AI systems that support coding, quality, and decision support. This would allow organizations to build AI governance programs anchored in a single, coherent set of expectations instead of anticipating overlapping or conflicting requirements.

These changes would directly support platforms like EVE, which already implement evidence-linked, human-in-the-loop AI within HIPAA- and HITRUST-aligned governance frameworks.

3. Novel legal and implementation issues, and HHS's role

For non-device AI used in clinical workflows, Reveleer sees several emerging legal and implementation issues:

- Attribution of liability in human-AI collaboration. When AI surfaces a suspected diagnosis or quality gap that a clinician accepts or rejects, questions arise about responsibility if subsequent audits or outcomes suggest that the AI contributed to over- or under-coding, or to missed care. Provide clear guidance for how accepted recommendations should be logged in systems and rules as to the significance a single instance can have on the continuous tuning of a model.

- Use of PHI for training and cross-client learning. Organizations must navigate HIPAA, state privacy laws, and contract terms when using PHI to train, fine-tune, or adapt AI models. The line between de-identified and identifiable data, and between client-specific and shared models, is often unclear in practice. Expand on minimum necessary guidance to include the use of AI and explicit consent by the patient for this information to be used.
- Black-box models and due process. When AI models are opaque, plans and providers may struggle to demonstrate to regulators and courts how coding or care decisions were reached, raising due-process and fairness concerns in payment disputes or audits. Establish rules or guidance blocking the use of black box models in clinical use cases.
- Indemnification and vendor risk allocation. Contracts increasingly need to allocate risk where AI-assisted workflows contribute to errors in submissions or clinical decisions, especially when both vendor AI design and client configuration choices interact.
- Dynamic model behavior and “drift.” Continuous learning and retraining can change AI behavior in ways that are hard to document and explain to auditors absent robust MLOps, versioning, and change-management practices.

HHS can help by:

- Publishing model contractual clauses or high-level guidance that clarifies shared responsibility among vendors, payers, and providers in AI-assisted workflows, emphasizing human final decision-making and traceability to source evidence.
- Clarifying under HIPAA that PHI may be used for training and validation of non-device clinical AI subject to strict purpose

limitation, de-identification where feasible, and robust access controls, and by encouraging standardized language in business associate agreements to cover AI use.

- Encouraging or requiring versioned records of datasets, model artifacts, and configurations used in AI-assisted clinical workflows, so that AI behavior is reproducible and auditable over time.
- Recognizing human-in-the-loop, evidence-linked AI as a distinct category from fully autonomous systems, with commensurate expectations for liability and oversight. This aligns with EVE's design, which treats AI as clinical decision support rather than as a replacement for clinician judgment.

4. Evaluation methods, metrics, robustness testing, and human-centered evaluation

For non-device AI integrated into clinical workflows, Reveleer uses and recommends a layered evaluation approach.

4.1. Technical and robustness evaluation

Pre- and post-deployment, models should be evaluated on:

- Precision, recall, and accuracy for task-specific outputs (e.g., suspected diagnoses, chart prioritization, quality data element extraction), segmented by population, contract type, and data source. This should include standard definitions for selected benchmark metrics to prevent misinterpretation of model quality in the market.

- Robustness and drift, via continuous monitoring for performance degradation, shifts in input distributions, and outlier confidence scores, with automated triggers for investigation and retraining.
- Fairness, via segmented performance analysis across demographic and clinical subgroups to detect systematic differences in sensitivity, specificity, or error rates.

4.2. Workflow and operational evaluation

Because AI's value is realized in workflow, Reveleer tracks:

- Charts reviewed per coder per day, charts eliminated from need for first-pass review, and reductions in time per chart in AI-assisted versus non-assisted workflows.
- Noise rate (the amount of inaccurate or inappropriate suggestions presented to the provider) and provider acceptance (the amount of accepted recommendations by the provider) are the two key metrics driving effective AI-driven suspecting. EVE Hybrid AI, for example, has shown up to three-fold noise reduction compared with legacy NLP in prospective risk adjustment, which directly improves provider engagement and reduces abrasion.
- Cycle times for risk and quality programs, including speed from data ingestion to action at the point of care.

4.3. Clinical, financial, and user-experience evaluation

Reveleer's measurement framework covers four domains:

- Clinical outcomes, including condition detection and documentation rates, adherence to evidence-based care pathways, and avoidable utilization.

- Financial performance, including RAF accuracy, Star Ratings, medical loss ratio, and ROI for AI-enabled programs (e.g., increased value per chart, reduced administrative costs).
- Operational efficiency, including abstraction throughput, error rates, and end-to-end cycle time for key value-based workflows.
- User experience and AI trust, including coder and clinician satisfaction, agreement rates with AI recommendations, and adoption analytics.

HHS should support these processes by:

- Funding multi-site evaluation consortia and pragmatic trials that use standardized measurement frameworks to compare AI-assisted and human-only processes in risk adjustment and quality.
- Supporting prize competitions and cooperative agreements for tools that provide open benchmarking harnesses, synthetic test sets, and measurement standards for non-device AI in clinical care.
- Using federal contracts and grant conditions to require evidence-linked, explainable AI and continuous monitoring in HHS-funded pilots, normalizing these evaluation practices.

5. Supporting private-sector accreditation, certification, and testing

To promote responsible innovation while avoiding duplicative regulation, HHS could:

- Endorse or co-develop voluntary AI assurance frameworks for non-device clinical AI that emphasize transparency, evidence linkage, bias monitoring, and human oversight—analogue to HITRUST and SOC 2 for security and privacy.

- Partner with industry on “trustmarks” for AI-ready clinical data and interoperable evidence layers, signaling that organizations meet standards for data quality, provenance tracking, and standardized APIs necessary for safe AI.
- Support independent testing labs or collaboratives that benchmark AI tools against standardized test cases (including synthetic data) for risk adjustment, quality abstraction, and care gap detection, with results published in a transparent registry.
- Leverage CMS contracts and quality programs to reward participation in such accreditation schemes—for example, by recognizing accredited AI workflows as meeting certain documentation sufficiency criteria for quality programs, while retaining audit rights.

These actions would reinforce the direction Reveleer has already taken with EVE—emphasizing explainable, audit-ready outputs and rigorous governance—while giving the broader market clear reference points for trustworthy AI.

6. Where AI has met or fallen short; high-impact future tools

6.1. Areas where AI has met or exceeded expectations

Reveleer and its clients have observed strong performance where AI is:

- Powering prospective risk adjustment and clinical intelligence, enabling more accurate risk stratification, earlier detection of undiagnosed conditions, and improved provider documentation at the point of care. Organizations using such tools have seen better provider engagement and measurable RAF improvements.

- Supporting digital quality abstraction and care gap closure, increasing throughput and precision in quality measures and improving Star Ratings through better gap targeting and closure.
- Improving operational efficiency and staffing resilience, allowing organizations facing workforce shortages to process more charts per coder with higher yield, effectively increasing coding and abstraction capacity without proportional staffing increases.
- Delivering favorable cost-effectiveness, as broader literature shows AI can reduce screening costs, unnecessary procedures, and readmissions when seamlessly integrated into clinical workflows.

6.2. Areas where AI has fallen short

Common shortfalls include:

- Noisy, non-explainable suspect lists. Legacy NLP systems often produce high volumes of low-specificity suspects without clear evidence links, overwhelming providers and eroding trust. This has been a key driver of provider abrasion.
- Workflow misalignment and low adoption. Tools that sit outside core EHR workflows, require extra log-ins, or lack role-specific UX tend to underperform regardless of model accuracy.
- Insufficient measurement and drift management. Many implementations lack rigorous pre/post evaluation and ongoing monitoring, obscuring true impact and allowing subtle performance degradation over time.

6.3. High-potential future tools

Reveleer believes the most impactful future non-device AI tools will:



- Operate over unified, multimodal evidence layers, combining structured data, unstructured content, and in the future images and signals, to deliver richer, context-aware clinical intelligence to help clinicians deliver the right care at the right moment from the right suggestion.
- Use hybrid, agentic evidence extraction with deterministic clinical reasoning for a growing number of clinical workflows, as EVE Hybrid AI does today for prospective risk adjustment, enabling audit-ready, configurable suspects with drastically reduced noise and rapid adaptation to new risk models and documentation standards.
- Embed digital quality measurement directly into care pathways rather than treating quality as a standalone reporting workflow, surfacing in-encounter opportunities to close gaps while maintaining measure-aligned data and provenance for eQMs and Stars.
- Combine ambient AI and prospective AI, where ambient tools reduce point-of-care documentation burden and EVE-like engines orchestrate pre- and post-encounter planning and panel management across the year. Leverage ambient data as a data set in prospective AI to add another data dimension to recommendations leading to a virtuous cycle of ever improving care.

7. Organizational roles and administrative hurdles

In Reveleer's experience, AI adoption decisions in clinical care are typically driven by:

- Chief Medical Officers and clinical leadership, who assess clinical validity, workflow fit, and provider trust.
- Chief Information Officers and IT/data leaders, who own data integration, interoperability, and infrastructure readiness.
- Value-based care and risk/quality program leaders, who are accountable for RAF, Star Ratings, and total cost of care and must see clear ROI and compliance alignment.

The primary administrative hurdles include:

- Fragmented governance across clinical, IT, compliance, and finance, which slows decisions and leads to inconsistent AI policies.
- Lack of structured AI change-management programs (stakeholder mapping, communication, training, escalation paths), which can cause confusion about roles, responsibilities, and safety oversight.
- Procurement and contracting processes that are not adapted to AI (e.g., difficulty evaluating explainability, data usage terms, and model lifecycle commitments).
- Unclear internal measurement standards, making it hard to prioritize AI investments or justify scaling pilots.

Reveleer has found that formal AI steering committees—combining clinical, operational, IT, and compliance leadership—and clear, outcomes-based measurement frameworks are critical to overcoming these hurdles and enabling durable adoption of an enterprise AI engine like EVE.

8. Interoperability opportunities



Enhanced interoperability would materially widen AI market opportunities and accelerate development in several areas:

- FHIR-based, near-real-time exchange of clinical data across payers and providers, beyond minimum USCDI, so AI engines can access timely, rich evidence for risk adjustment, quality, and care management rather than relying on delayed batch files and manual chart retrieval.
- Standardized representation of data provenance and evidence links, including where each data element came from, how it was transformed, and which documents support a given diagnosis or quality data element. This directly supports explainable AI and auditability.
- Shared, de-identified benchmarking datasets for risk adjustment, quality measures, and care gap detection, enabling vendors and health systems to benchmark AI performance and robustness fairly.
- Interoperable “evidence layers” across payers and providers, with standard interfaces and vocabularies for longitudinal records that can be reused across retrospective and prospective risk, quality, and utilization management.
- Open EHRs to support AI-based insights tools to power the best intelligence within the provider workflow. Closed or pay-to-play EHR systems limit choice for the providers to the AI initiatives within the EHR infrastructure and in parallel contributes to the provider abrasion when new innovative tools are activated but must operate outside of the primary clinical workflow tools.

The Reveleer platform already uses interoperable clinical repositories and standardized APIs to support EVE across risk adjustment and quality solutions; strengthening national interoperability in these

areas would significantly expand the impact and safety of similar AI engines.

9. Patient and caregiver challenges and concerns

9.1. Challenges patients and caregivers want AI to address

From the vantage point of value-based care operations, patients and caregivers most want AI to help with:

- Fragmented care and incomplete information. Patients want clinicians to see the full picture of their health across settings to ensure the care and advice they are getting considers their health conditions and needs uniquely. AI operating over unified evidence layers can help ensure important conditions and care gaps are not overlooked.
- Delays and access barriers. AI-enabled risk stratification and care gap identification can prioritize high-risk patients for outreach, ensuring timelier care and reducing avoidable hospitalizations. But knowing that healthcare is local, these systems need to account for regional factors (medical personnel, geographic isolation, socioeconomic status, and population demographics) and physician-specific approaches to care.
- Administrative burden and confusion. Patients benefit when AI reduces clinician documentation burden and back-office inefficiencies, freeing more time for communication and shared, personalized decision-making.
- Health equity and tailored care. AI can help identify disparities in access and outcomes and support more targeted interventions,

provided fairness is monitored and addressed. As above, need to account for regional factors (medical personnel, geographic isolation, socioeconomic status, and population demographics) and physician-specific approaches to care.

9.2. Concerns about AI in clinical care

Common concerns include:

- Loss of human judgment and relationship. Patients and caregivers worry that AI may replace, rather than augment, clinician judgment. Reveleer deliberately positions EVE as clinical decision support, with non-binding recommendations and required human validation.
- Privacy, data security, and secondary use. There is concern that AI requires extensive sharing and repurposing of sensitive health data. Reveleer addresses this with HIPAA- and HITRUST-aligned controls, de-identification or redaction where appropriate, and strict purpose limitations on data use.
- Bias and fairness. Patients may fear that AI will perpetuate or worsen inequities. Continuous fairness monitoring and transparent performance metrics by subgroup are essential to mitigate this risk.
- Opacity and lack of recourse. Without explainability and clear appeal mechanisms, patients may feel that algorithmic decisions are unchallengeable. Evidence-linked AI and human-in-the-loop workflows provide a foundation for explainable decisions and remediation.

HHS can support trust by emphasizing evidence-linked, human-supervised AI and promoting patient-facing transparency about when and how AI is used in care decisions.

10. Priority AI research areas; literature on costs, benefits, and transfers

10.1. Priority research areas

Reveleer recommends that HHS prioritize research in:

- Hybrid AI architectures for non-device clinical decision support. Comparative studies should evaluate hybrid agentic-plus-deterministic approaches (like EVE Hybrid AI) against pure generative or black-box models in terms of accuracy, auditability, provider acceptance, and regulatory compliance in risk adjustment and quality workflows.
- AI-enabled evidence layers and continuous learning systems. Studies should examine how unified, provenance-rich evidence layers coupled with continuous-learning AI engines affect outcomes, costs, and resilience of value-based care programs over multi-year horizons.
- Digital quality measurement embedded in care. Research should develop and validate methods for AI to maintain measure-aligned quality data and surface in-encounter opportunities to close gaps, instead of treating quality reporting as an after-the-fact exercise.
- Human-AI collaboration and workflow design. Empirical work should focus on UX patterns, alerting strategies, and feedback loops that maximize adoption, trust, and safety in AI-augmented coding and clinical workflows.
- Equity and fairness in AI-enabled value-based care. Research should investigate how AI affects care across different populations and

payment models, including potential unintended consequences in risk selection and resource allocation.

10.2. Published findings on impact, costs, benefits, and transfers

The emerging literature on AI in clinical care shows:

- Across multiple clinical domains (e.g., screening, chronic disease management, imaging), AI interventions have demonstrated strong cost-effectiveness or cost-saving potential, often via reduced unnecessary procedures, improved early detection, and fewer readmissions.
- Evaluations of AI-powered prospective risk and care management have found improvements in documentation completeness, reduced readmissions, and better alignment of reimbursement with true disease burden, while lowering administrative costs through automation of chart review and suspecting.
- Economic evaluations typically combine clinical outcomes (sensitivity, specificity, readmission rates, adverse events) with utilization changes (length of stay, number of procedures, outpatient versus inpatient mix) to derive budget impact and ROI. Implementation costs such as integration, training, and model maintenance are explicitly accounted for, along with indirect effects on workflow and provider satisfaction.

HHS-funded research can extend this evidence base by focusing specifically on non-device AI in value-based care, using standardized frameworks to ensure that comparisons of AI-assisted and traditional workflows are fair, transparent, and actionable for regulators and program designers.

