

February 23, 2026

James O'Neill  
Deputy Secretary  
U.S. Department of Health and Human Services  
200 Independence Ave. SW  
Washington, DC 20201

Thomas Keane, MD  
Assistant Secretary for Technology Policy ("ASTP")  
U.S. Department of Health & Human Services  
200 Independence Ave. SW  
Washington, DC 20201  
*By electronic submission*

*Re: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care; RIN 0955-AA13*

Dear Assistant Secretary Keane and Deputy Secretary O'Neill:

I'm grateful for the opportunity to comment on the Department of Health and Human Services (HHS), in collaboration with the Office of the Deputy Secretary and Assistant Secretary for Technology Policy/Office of the National Coordinator for Health Information Technology (ASTP/ONC), Request for Information on Accelerating the Adoption and Use of Artificial Intelligence (AI) as part of Clinical Care from the patient, carepartner, and caregiver perspective as we work towards building safe, effective, collaborative, and accessible AI in clinical care for all patients and families.

From the perspective of a board-certified patient advocate, the primary carepartner to two disabled adults, and a patient myself, I have experienced firsthand, personally and professionally, both the promise and the failures of our healthcare system. In 2008, I was misdiagnosed with advanced lymphoma. Because I knew to request my medical records and seek a second opinion, months of testing, invasive procedures, endless imaging, labs, video-assisted thoracic surgery, and waking up in recovery after surgery in isolation for suspected tuberculosis ultimately revealed that I did not have cancer, but rather a fungal infection acquired on my honeymoon that required no treatment, only monitoring. Looking back, I can't help but wonder if an AI powered software to read my radiology could have provided a differential diagnosis that could have prevented months invasive and traumatic testing and procedures as well as unnecessary, costly care. That experience shaped my understanding of what is at stake when information, interpretation, and decision-making tools fail, or succeed, in clinical care.

### **Executive summary**

AI is already shaping clinical decisions across the health system, yet the patients whose lives those decisions affect are largely excluded from visibility, participation, and

oversight. The greatest barriers to responsible private-sector innovation are not solely technical or regulatory but rather they are structural: unclear accountability, misaligned incentives, reimbursement gaps, workforce readiness challenges, privacy concerns, and insufficient mechanisms for real-world validation and patient partnership. Patients face a paradox: AI is influencing our care while we often cannot access, understand, or question the outputs used to make decisions about us.

Large language models and AI powered tools are transforming how patients, carepartners, and caregivers engage with our health information, interpret diagnoses, and navigate daily clinical realities while waiting weeks or months for appointments. Our use of AI represents patient engagement, empowerment, and personalized operationalization of care decisions in real time. Like any tool, these technologies carry both promise and risk, which is precisely why patient, carepartner, and caregiver voices must be embedded not only in listening exercises but in co-creation, co-development, governance, and funded leadership roles that shape how AI is designed and deployed.

HHS is uniquely positioned to close this trust and governance gap by aligning policy, payment, evaluation, and transparency expectations so innovation scales alongside safeguards. Designed, implemented, and governed intentionally with patients as partners, AI can improve outcomes, reduce costs, strengthen trust, and enable clinicians to focus on what matters most: caring for people. Healthcare does not occur only within the four walls of medicine, and neither should the governance, oversight, or design of the technologies that increasingly shape it.

## **1. What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?**

After having the privilege of serving on numerous multi-stakeholder efforts and taskforces on AI in health care, including the National Academy of Medicine's AI Code of Conduct project and the Digital Medicine Society's AI Playbook work, the most commonly discussed barriers to private sector innovation that I am hearing from various stakeholders in AI include concerns about regulatory uncertainty, implementation challenges, lack of reimbursement frameworks, ownership of liability for consequent harm, privacy concerns, and lack of appropriate workforce and workforce training to support assessment.

## **2. What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.**

As we carefully define what effective use of AI in clinical care looks like, there are concerns and problems defined by the patient, carepartner, and caregiver stakeholder voice that need to be considered as we examine regulatory, payment policy, and programmatic changes the HHS should prioritize to incentivize collaborative, trusted use of AI in the healthcare ecosystem.

AI outputs and clinical decision support have been successfully used in patient care for well over the last decade. Outputs in the form of, for example, risk prediction scores, diagnostic support outputs, treatment optimization recommendations, early warning alerts, medication safety intelligence, and the likes of prognostic and trajectory modelling are routinely used to make informed decisions about patient care.

Building trust and transparency in AI is one of the cornerstone conversation topics in US and global healthcare sectors. Ironically, patients and families have essentially little knowledge of and no access to any of the AI intelligence that runs behind the scenes of our care. AI derived outputs that are used to make decisions about an individual's care or coordination of care are not documented in our medical records. Not having access to AI derived outputs has a broad spectrum of consequences, including but not limited to:

- robbing us of opportunities to participate in shared decision-making in our or our loved one's care using intelligence from AI derived outputs.
- depriving us of clarity about the impact of use of AI in our care on the costs of our care, and anecdotally, leading to growing concerns about our out-of-pocket costs increasing as a result of the deployment AI powered technologies, such as ambient listening.
- preventing us from learning about the transformative power of AI in our care and outcomes and the celebrating positive impacts AI powered tools can bring; downstream effects include denying us the opportunities to become trusted sources of information in evangelizing and educating our families, friends, places of worship, and communities about the positive impacts of AI.
- shielding us from the inconsistencies, hallucinations, errors, and deviations of AI outputs that are experienced by our physicians and care teams, preventing us from fully appreciating the art of medicine; downstream effects include denying us the opportunities to become trusted sources of information in evangelizing and educating our families, friends, places of worship, and communities about the potential harms and negative impacts of AI.
- denying us the opportunities to inspect and curate the outputs and intelligence generated for potential errors; patients, carepartners, and caregivers can help assist in reducing "garbage in and garbage out" concerns.

- being routinely excluded, overlooked, and denied opportunities to participate as partners in the co-design, implementation, and integration of AI tools and technologies that are going to impact our lives.

As we examine what HHS regulations, policies, or programs could be revisited to augment use of AI in our clinical care, I strongly recommend the following considerations.

AI-derived outputs used to make decisions about an individual's care or coordination of care should be recognized as EHI (45 CFR § 171.102) and hence the designated record set (45 CFR § 164.501) and consequently, the Right of Access (45 CFR § 164.524) should apply. Individuals should have a right to access AI outputs as part of our medical record. Measures should be taken to include access to AI outputs as in scope for considerations of information blocking (45 CFR Part 171).

Programmatically speaking, patients, carepartners, and caregivers can help advance the effective use of AI in our care. Listening to patient voices is not enough. Including patients on stages to talk about our experiences with AI is not enough. Budgeting for patient stakeholder voice partnership and leadership is necessary as we explore specific changes that can be implemented programmatically at scale. Currently there are zero federally patient designated AI governance, engagement and strategy focused leadership roles. Innovation means doing things a new way. The time is now to fund the roles that are needed to do the work.

At minimum, patients need to know where AI is used in our care, when it mattered, how we can interpret it, how it impacts our costs of care, and how to operationalize on the knowledge obtained. It should be documented as to how our doctors adjudicate disputes, i.e., model output stands but input corrected, output rejected because of patient history, output recalculated to include missing patient records.

It is well documented that errors in our medical record are common. The downstream effects of these rampant errors in our records contribute to outputs that could cause harm. Patients, carepartners, and caregivers should have a right to not only access AI outputs used in our care and amend errors in our records, but also address errors in potential AI outputs used in our care. A feedback loop is missing that can help support post-deployment monitoring. As a patient, where can I go to report an error on an AI tool that was implemented in my care? Who do I report that to at my healthcare system or cancer center? I've yet to find a contact or point person or feedback portal.

Standard of care fails or in some cases isn't even available for patients living with life-altering, life-limiting conditions, rare disease, or multiple comorbidities. There are

structured, standardized pathways, i.e., Right to Try or Expanded Access (compassionate use), where patients who have exhausted all approved treatment options can make an informed decision to try innovative treatment approaches that have yet to be FDA approved when all standard of care has been exhausted or no standard of care is available. We currently do not have comparable pathways for AI or digital health for patients with serious or immediately life-limiting conditions. Innovation can not only move from top-down. Patients and families are playing proactive roles in our care and many are pleading for access to various innovations to manage their condition only to be told by their health system, physician, cancer center, “we don’t do it that way here.” I strongly encourage colleagues at FDA to explore Right to Try and/ Expanded Access type pathways for patients to access promising AI powered innovations who are unable to participate in a clinical trial, or where a clinical trial may not (yet) be available, to access certain investigational digital health treatment or precision medicine options.

We are currently witnessing a metamorphosis of the doctor-patient relationship. Programmatically, doctors and care providers need training for engaging with patients, carepartners, and caregivers who are utilizing AI to drive our informed decision making. There needs to be programming to also best support and educate patients, our carepartners and caregivers, on how to best use AI in our care. These necessary updates in education and training must be co-designed with the patient and family stakeholder voice. Individuals who share their lived experience and experience to co-create and reimagine physician and nursing training as well as co-creating credible educational resources to level the playing field for all patients should be fairly compensated using models such as the National Health Council’s Fair Market Value calculator or frameworks from PCORI.

We need to compensate and reimburse for the human work and the art of medicine that ensures AI is safe and effective, including oversight, patient education, and shared decision-making. We want to be able to discuss AI in our care with the doctors and care teams we trust with our care. We need to reimagine how we virtually, digitally, and asynchronously engage in our care. Patients need safe, trusted, secure spaces to collaborate with our physicians, care teams, community health workers, social workers, etc., in LLMs. I can invite a patient, carepartner, or caregiver to collaborate in a shared google doc, but we could be more efficient if I could collaborate within an LLM shared working space, where everyone can see my prompting, the outputs, the interrogations, the feedback, the hallucinations and errors instead of me sending screenshots in group texts. ChatGPT has a feature to start a group chat that in its current state, is useless and misleadingly named. I am unable to invite my fellow patient advocates, patients,

carepartners, caregivers, or colleagues in medicine into my deep dives on any given day of patient advocacy work. The time is now to dream bigger.

America's AI Action plan mentions sandboxes for research and testing. Using patient and care team-led sandboxes to test payment policies that reimburse clinician time to validate and interpret results, exploring the impact of shared decision-making conversations where AI outputs guide choices. Multi-disciplinary reviews, such as tumor-board style governance for predictive tools, could be implemented as another sandbox use case. Payment should be paired with minimal documentation standards, i.e., what output, how it impacted care, shared decision-making, and impact on outcomes.

**3. For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (e.g., relating to liability, indemnification, privacy, and security). What novel legal and implementation issues exist and what role, if any, should HHS play to help address them?**

Patients, carepartners, and caregivers are encountering a new category of risk in AI-enabled clinical care: decisions affecting our care may be influenced by systems we cannot see, question, understand, or collaborate with. The core issue is not whether AI can be useful, but whether governance has kept pace with deployment. Current legal and operational frameworks do not clearly define responsibility when AI contributes to clinical decision-making, how errors should be investigated, or how patients can challenge outputs derived from incorrect or incomplete data. Privacy concerns are equally urgent. AI development and deployment rely on extensive health data flows that patients often cannot easily track or meaningfully understand, creating legitimate concerns about misuse, secondary use, and security vulnerabilities. HHS is uniquely positioned to close these governance gaps by setting consistent national expectations for transparency, documentation, and accountability when AI informs clinical care. Patients should be able to know when AI influenced a decision, understand the basis for its output, and access mechanisms to question or correct it when necessary. Federal leadership is essential to ensure that innovation scales alongside safeguards so that advances in clinical AI strengthen patient trust rather than undermine it.

**4. For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and other workflow and human- centered evaluation methods for clinical care? Should HHS further support these processes? If so, which mechanisms would be most impactful (e.g., contracts, grants, cooperative agreements, and/or prize competitions)?**

From the patient and care partner perspective, HHS should ensure that non-device AI used in clinical care is evaluated across its lifecycle for real-world safety, equity, transparency, and usability, and technical accuracy through standardized pre-deployment validation, post-deployment monitoring, and robustness testing. Patient-partnered sandbox environments described above can be impactful and should be strongly considered. There should be clear disclosure in clinical workflows, especially in patient-facing disclosure when AI informs care decisions. CMS should align payment and quality programs to support validated, monitored AI and the clinician time required for safe use, oversight, patient education, and shared decision-making.

HHS could support real-world AI evaluation networks with patients as funded partners by prioritizing this work within existing appropriations and authorities. HHS could prioritize cooperative agreements and grants to build real-world evaluation networks with patients as funded partners and co-creators, supplemented by targeted contracts and prize competitions to accelerate techquity, explainability, and safety testing. CMS demonstration and innovation models could fund real-world testing through CMMI pilots, alternative payment models, i.e. ACCESS, and quality-reporting incentives to move evaluations real-world practice. If patients are to be meaningfully placed in charge of our care, federal policy must support us with the tools to evaluate the technologies shaping clinical decisions about us and our lives. A dedicated funding stream should be established for patient-driven AI evaluation networks to ensure empowerment is paired with evidence, accountability, and safety.

HHS should align AI policy with established federal patient-engagement frameworks (e.g., PCORI, FDA PFDD, AHRQ models) that treat patients, carepartners, and caregivers as end-users and partners in design, evaluation, and oversight, while fairly compensating individuals for their time, expertise, and lived experiences.

## **5. How can HHS best support private sector activities (e.g., accreditation, certification, industry-driven testing, and credentialing) to promote innovative and effective AI use in clinical care?**

From the patient and care partner perspective, HHS can best support private-sector accreditation, certification, testing, and credentialing efforts for clinical AI by establishing clear national expectations for transparency, safety monitoring, and real-world performance while requiring meaningful patient, carepartner, and caregiver involvement in design, evaluation, implementation, and governance. HHS should align payment, certification, and procurement policies to incentivize adoption of independent validation programs that assess not only technical accuracy but also techquity, usability, patient

comprehension, and metrics that include patient perspectives on value. Federal agencies can further accelerate responsible innovation by supporting public-private-patient testing collaboratives and publishing model evaluation criteria that embed patient stakeholder participation as a core standard. Trustworthy clinical AI is built not only on strong algorithms, but on structured inclusion of the patient, carepartner, and caregiver voices it is designed to serve.

**6. Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? What kinds of novel AI tools would have the greatest potential to improve health care outcomes, give new insights on quality, and help reduce costs?**

AI tools have demonstrated the greatest real-world impact when they address access gaps, reduce administrative burden, or enable earlier clinical action within patient care workflows. Some examples of AI tools that were deployed and met or exceeded in clinical care, performance, and cost expectations include:

- **Autonomous diabetic retinopathy screening in primary care**  
FDA-authorized autonomous AI systems for diabetic retinopathy have demonstrated accurate detection of referable disease in primary care settings, enabling point-of-care identification and referral without requiring immediate specialist interpretation, which can help close screening gaps where specialty access is limited.
- **Ambient clinical documentation tools**  
Ambient listening has been documented to demonstrate meaningful reductions in clinician documentation time and administrative burden as well as improving clinician well-being.
- **Hospital deterioration prediction models**  
AI-driven early-warning models analyzing EHR data in real time have reduced adverse events by identifying clinical decline earlier and enabling earlier clinical intervention.
- **4. AI clinical decision support for unnecessary testing (SmartAlert system)**  
Significant volume reductions of repeat laboratory testing demonstrating cost savings through eliminating low-value care.

With respect to novel tools that could have the greatest potential for improving healthcare outcomes comes to mind a Waze for Healthcare that harnesses the power multi-stakeholder data intelligence and crowd-sourced peer health support in real-time to help patients make informed decisions about the next steps of their care journeys where they are.

Patients and families are currently experiencing harm from automated or AI-assisted payer systems used to process claims at scale, including batch denials and algorithm-supported downcoding of reimbursements. When coverage or payment determinations are automated but correction pathways remain slow, opaque, or difficult to access, patients can face delayed care, unexpected costs, and administrative barriers that directly affect health outcomes. Oversight, transparency, and clear accountability standards must keep pace with automated decision tools so that efficiency gains do not come at the expense of patient access, provider sustainability, or medically necessary care.

**8. Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.**

The greatest opportunities for AI advancement lie where patient data still remains siloed. Interoperability that connects clinical records, claims, imaging, labs, genomics, wearable data, and patient-reported outcomes would dramatically expand research capability, improve model accuracy, and allow real-world validation across diverse populations as well as enable robust patient-navigation technologies. Standardized formats and terminologies such as FHIR, USCDI, LOINC, and SNOMED CT make it possible for systems to exchange information safely and meaningfully, while benchmarking infrastructure and shared evaluation datasets enable transparent performance comparisons. At the same time, new or expanded standards are needed to support consistent exchange, interpretation, and documentation of AI-derived outputs, such as risk scores, predictions, recommendations, confidence levels, and model provenance, so that these outputs can move across systems, be understood by clinicians and patients, and be evaluated for safety and accuracy over time. Patients experience the consequences of fragmented data every day; when our information cannot move, innovation cannot scale. Interoperability is therefore not only a technical priority but a patient safety and equity imperative.

**9. What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?**

Patients and caregivers want AI in clinical care to address persistent system failures and gaps in care that routinely negatively impact our outcomes, including delayed or missed diagnoses, fragmented care coordination, administrative burden that limits our clinician's time, inequitable access to specialty expertise, preventable medical errors,

and financial toxicity. Tools that accelerate time to diagnosis, close care gaps, identify risk earlier, reduce patient administrative burden, streamline treatment planning and clinical trial navigation are examples of solutions viewed as high-value because they improve both safety and the lived experience of our or our loved one's care. At the same time, patients and caregivers have significant concerns about the adoption of AI, particularly when its role in decision-making is not transparent, when outputs cannot be accessed, reviewed, or discussed, when systems may encode bias or worsen disparities and gaps in care, or when accountability is unclear if harm occurs. Duty of care and do no harm are key. There is also concern that poorly implemented AI could increase cognitive burden on clinicians, generate alert fatigue, or shift decision responsibility onto patients without adequate support. Many patients do not have a primary carepartner or caregiver, may not be digitally savvy or may have language and accessibility need. Patients consistently emphasize that trust depends on clear disclosure of when AI is used, understandable explanations of how it informs care, the ability to correct or challenge outputs, and assurance that human clinical judgment remains responsible for final decisions. Patients do not oppose AI in care; we oppose invisible, unaccountable AI that affects our health, our care, and our lives without our knowledge or voice.

As we critically examine the concerns of the patient stakeholder voice, I strongly encourage all of HHS agencies to read the [AI Rights for Patients](#). This document was created by a cohort convened by The Light Collective, a national community of patient activists, clinicians, and health data experts committed to building tech and digital rights for patients. There are no assumptions of what patients want in this rights framework as this was fully co-created with 100% patient voices.

## **Closing**

I applaud HHS's leadership in advancing responsible AI adoption in patient care. The patient, carepartner, caregiver, and consumer stakeholder perspectives are critically important in the advancement of responsible adoption of AI in patient care. Patient care happens not only within the four walls of medicine, but also in the privacy of our homes, places of work, worship, and community, and play.

I applaud HHS's leadership in advancing responsible AI adoption in clinical care and for recognizing that the voices of patients, carepartners, caregivers, and consumers are essential to getting this right. Thank you for the opportunity to contribute to this important national conversation. Patients are not asking to slow innovation, we are asking to strengthen it. We want to help build an AI-enabled health system that is safer, more transparent, more equitable, and more humane than the one we have today. I urge the Department to build on this moment by ensuring that patient partnership is not symbolic but structural, embedded in policy design, evaluation, governance, and funding priorities. With intentional policy, meaningful partnership, and shared accountability, AI can fulfill its promise not just as a technological advance but as a

public good. Patients, carepartner, and caregivers stand by ready to partner, to collaborate, and lead with our lived experiences and expertise. The question is not whether we are prepared to lead alongside you, but whether the system is ready to support and prioritize authentic partnership.

Respectfully,

**Grace Cordovano**

Grace Cordovano, PhD, BCPA  
Board Certified Patient Advocate & Founder  
Enlightening Results