

NOETED

Public Comment — Request for Information

Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care

Docket No. HHS-ONC-2026-0001 | Submitted February 2026

Submitted by:

Keith “Duv” Williams
Founder & CEO,
NOETED
duv@noeted.com |
noeted.ai
Redding, California

About NOETED:

AI-powered healthcare compliance intelligence platform. Member, NVIDIA Inception Program. Building centralized regulatory intelligence infrastructure for U.S. healthcare organizations.

Preamble

NOETED is an AI-enabled healthcare compliance intelligence platform headquartered in California. We are building what U.S. healthcare lacks today: a reliable regulatory intelligence layer that continuously monitors, aggregates, and connects compliance requirements from federal agencies (CMS, OIG, FDA, SAMHSA, ONC), state health departments, and regional regulators, and translates those signals into actionable guidance for payers and providers.

Our perspective on this RFI reflects two things: (1) direct experience developing AI systems that must operate inside the U.S. healthcare regulatory environment, and (2) hands-on attempts to programmatically consume federal and state regulatory data for large health systems, plans, and safety-net providers.

Across all ten questions, we return to a single, solvable root cause: the lack of modern, standardized regulatory data infrastructure. Today, many APIs, databases, and publication systems for healthcare rules are fragmented, inconsistently maintained, or simply absent. This forces organizations into manual monitoring, slows responsible AI adoption, and ultimately harms patients when preventable compliance failures disrupt care.

We therefore focus our recommendations on a sequenced approach: (1) modernize regulatory data infrastructure; (2) standardize regulatory data formats; (3) embed regulatory intelligence in certification, payment, and safe-harbor frameworks; and (4) fund the research needed to evaluate AI systems that depend on this infrastructure. NOETED is building a complementary private-sector layer on top of such an infrastructure, and we are eager to partner with ONC and HHS so that public and private efforts reinforce each other.

NOETED Response

From NOETED's experience, four barriers consistently slow responsible AI innovation and adoption in clinical care:

Regulatory fragmentation and limited machine-readability.

AI developers and health systems must comply with requirements from CMS, OIG, FDA, ONC, state regulators, and others. These rules are published through heterogeneous channels (Federal Register, PDFs, web pages, listservs) with no unified, machine-readable structure. For many organizations—especially FQHCs, small hospitals, and regional plans—this makes it practically impossible to maintain a current, organization-wide view of “what the rules actually are,” which in turn makes it risky to deploy AI at scale.

Unclear and inconsistent reimbursement for AI-enabled care.

In the absence of durable, technology-neutral payment pathways for AI-assisted services, investments are skewed toward a narrow set of large systems and high-margin use cases. Smaller and rural providers, who could benefit disproportionately from AI support, are often unable to make a business case for adoption.

Liability ambiguity for AI-assisted decisions.

Existing frameworks do not clearly allocate responsibility among AI vendors, deploying organizations, clinicians, and data suppliers when AI contributions are involved in clinical decisions. This ambiguity encourages risk-averse organizations to constrain AI to low-stakes workflows, even where carefully governed AI could materially improve safety, quality, and equity.

Regulatory data infrastructure that lags behind clinical data infrastructure.

Substantial progress has been made on clinical and administrative data interoperability. By contrast, regulatory data—conditions of participation, coverage rules, enforcement guidance—is often published through APIs that are undocumented or deprecated, datasets that lag real-time changes, and formats that vary widely across agencies. As an AI compliance platform, NOETED encounters this daily: we spend disproportionate effort normalizing and validating public regulatory data that should already be available in standardized, machine-readable form. This raises costs, limits innovation, and leaves smaller organizations unable to automate compliance monitoring at all.

HHS can have outsized impact by treating regulatory data infrastructure with the same seriousness as clinical data infrastructure—because AI cannot be safe, equitable, or scalable if health care organizations cannot reliably know the rules AI-assisted workflows must follow.

Q2

What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care? What HHS regulations, policies, or programs could be revisited? Please provide specific changes and applicable CFR citations.

NOETED Response

NOETED recommends a sequenced set of actions that align with HHS's goals for safe, equitable AI in clinical care:

1. *Modernize and standardize regulatory data publication across HHS.*

Direct CMS, OIG, FDA, SAMHSA, ONC and related entities to publish regulatory requirements, guidance, and change notices through maintained, documented, versioned APIs with common schema and metadata (effective date, affected entity type, CFR citations, superseded guidance).

Use the Evidence Act and OMB A-130 authorities to establish minimum publication standards for regulatory data, analogous to existing expectations for clinical data interoperability.

Charge ONC with developing a "Regulatory Data Implementation Guide" aligned with FHIR, so that certified health IT can consume regulatory signals in the same way it consumes clinical data.

2. *Embed regulatory intelligence into certification criteria.*

Extend ONC's Health IT Certification Program to include baseline requirements for AI tools that influence clinical workflows or compliance-sensitive operations: documented regulatory data sources, update mechanisms, and auditability.

Over time, require certified systems to demonstrate an ability to ingest and act on standardized regulatory notifications.

3. *Create durable reimbursement pathways for AI-assisted services.*

Under 42 CFR Part 414, establish technology-neutral codes and coverage criteria for AI-assisted clinical services where evidence supports improved outcomes or efficiency, rather than relying solely on contractor-level decisions.

Prioritize use cases with clear equity impact (e.g., documentation support in safety-net settings, AI-assisted chronic disease management in underserved populations).

4. *Clarify safe harbors and HIPAA guidance for AI.*

Update OIG safe harbors to explicitly address provision of AI tools to affiliated providers under conditions similar to existing EHR safe harbors, enabling health systems to extend AI capabilities without AKS uncertainty.

Issue targeted HIPAA guidance (45 CFR §164.514) clarifying permissible uses of de-identified and synthetic data for AI training and evaluation, with plain-language examples and guardrails.

5. Explicitly recognize regulatory intelligence infrastructure as a funding priority.

Through programs such as ONC's LEAP in Health IT, explicitly invite proposals focused on regulatory intelligence infrastructure—systems that help organizations monitor and respond to evolving requirements in near real time. These platforms are a precondition for safe AI deployment in any compliance-sensitive clinical workflow.

Collectively, these actions would lower compliance uncertainty for innovators, reduce the burden on providers, and make it far more realistic for organizations of all sizes to adopt AI tools with confidence.

Q3

For non-medical devices, what novel legal and implementation issues exist regarding AI in clinical care (liability, indemnification, privacy, security)? What role should HHS play?

NOETED Response

Several novel issues arise when AI is used in clinical care outside the medical-device framework, particularly in compliance-sensitive workflows:

1. Multi-party liability for AI-assisted decisions.

Today, it is unclear how responsibility is allocated among the AI developer, the deploying organization, clinicians, and data suppliers when AI contributes to a clinical or operational error. This uncertainty chills adoption, particularly in high-stakes settings.

HHS can lead by developing, with DOJ and state partners, model liability frameworks that:

- Clarify roles and duties across the AI lifecycle.
- Create safe harbors for organizations that follow defined governance and monitoring practices.
- Encourage transparent incident reporting and learning instead of purely punitive responses.

2. Compliance liability for AI recommendations.

In compliance use cases, organizations remain fully responsible when AI-generated guidance is wrong, yet they currently receive no recognition for using higher-assurance tools or for following best-practice evaluation processes.

HHS could clarify, for example in OIG and CMS enforcement guidance, how reliance on certified or credentialed AI tools factors into penalty calculations and corrective-action expectations. This would incentivize use of higher-quality systems without shifting responsibility away from providers.

3. Source data reliability in regulatory AI.

For AI tools that depend on government-published regulatory data, the quality and timeliness of that data is itself a risk factor. If a government API is outdated or deprecated, an otherwise well-governed AI tool may produce incorrect recommendations.

HHS can address this by:

- Committing to service-level expectations for key regulatory data feeds (timeliness, uptime, change logs).
- Clarifying how documented reliance on official data sources will be treated in enforcement discretion, especially where those sources can be shown to be incomplete or out of date.

4. Privacy in AI processing of unstructured data.

Large-scale processing of notes, images, and voice introduces questions about how HIPAA's minimum necessary standard applies when models need broad context to function.

HHS should issue targeted guidance on AI handling of unstructured PHI, including expectations for minimization, logging, access control, and secondary use (e.g., model training).

Recommended HHS role.

HHS is uniquely positioned to convene multi-stakeholder working groups and convert their outputs into practical tools: model policies, certification criteria, and updated guidance. Extending ONC's certification program to cover AI governance attributes (data provenance, monitoring, explainability, and incident response) would provide a concrete mechanism to operationalize these legal and implementation expectations.

Q4

For non-medical devices, what are the most promising AI evaluation methods (pre- and post-deployment), metrics, robustness testing, and human-centered evaluation methods for clinical care? Should HHS support these? If so, which mechanisms?

NOETED Response

Pre-deployment evaluation.

The most promising pre-deployment approaches for non-device clinical AI share three traits:

- They evaluate performance across demographically and clinically diverse subpopulations, not just aggregate metrics.
- They test robustness under distribution shift (how models behave when “real-world” data differs from training data).
- They incorporate structured expert review of representative AI outputs before deployment.

For compliance and operations AI, pre-deployment evaluation must also include regulatory data source quality: what sources are used, how often they update, and how gaps or inconsistencies are handled. A compliance AI tool cannot be safer than the regulatory data infrastructure it depends on.

Post-deployment monitoring.

Today, systematic post-deployment monitoring is rare. For AI in clinical and compliance workflows, monitoring plans should:

- Track performance drift over time and across subpopulations.
- Detect bias emergence and unintended use cases.
- Incorporate automatic checks for changes in underlying regulatory or policy data that could invalidate model assumptions.

For compliance AI specifically, this means monitoring both: (1) model outputs and (2) the freshness of regulatory data sources—flagging when upstream data lags new rules.

Human-centered evaluation.

Technical metrics alone cannot show whether AI improves care or compliance. HHS should support evaluations that measure:

- How clinicians and compliance staff actually use or override AI recommendations and why.
- Effects on workload, burnout, and error detection.
- Downstream impact on patient outcomes, access, and equity.

Recommended HHS mechanisms.

HHS can:

- Use cooperative agreements to fund multi-site evaluations of AI tools in real clinical and compliance workflows.
- Use prize challenges for well-defined technical benchmarks (e.g., robustness under specified shift scenarios).
- Create a Healthcare AI Evaluation Repository, a public (or tiered access) database where evaluation results for deployed AI tools are reported in a comparable format, helping providers and payers make informed procurement and governance decisions.

NOETED would welcome the opportunity to participate in such evaluations specifically for compliance and regulatory intelligence use cases, contributing both tooling and real-world evaluation data.

Q5

How can HHS best support private sector activities (accreditation, certification, industry-driven testing, credentialing) to promote innovative and effective AI use in clinical care?

NOETED Response

HHS can accelerate safe AI adoption by aligning federal levers with emerging private-sector accreditation and certification efforts:

1. Extend ONC Health IT Certification to non-device AI tools.

Expand 45 CFR Part 170 to include criteria for AI systems that directly influence clinical workflows or compliance-sensitive operations (e.g., documentation assistants, clinical decision support, compliance intelligence tools).

Criteria should address data provenance, bias and robustness testing, interoperability with certified EHRs, update mechanisms for clinical and regulatory changes, and post-deployment monitoring.

2. Recognize industry-driven credentialing frameworks.

Professional bodies (e.g., HIMSS, AHIMA, and AI-focused consortia) are designing voluntary AI credentials. HHS can create a formal “recognized credentialing framework” list, akin to how NCQA accreditation is recognized in Medicare/Medicaid programs.

Such recognition would steer the market toward higher standards without requiring exhaustive federal rulemaking for every AI use case.

3. Link certification/credentialing to enforcement incentives.

In CMS and OIG enforcement guidance, clarify that healthcare organizations using AI tools that meet ONC certification or recognized credentialing can receive defined consideration (e.g., in penalty calculations or corrective-action planning) when AI-assisted errors occur.

This does not shift responsibility away from providers, but rewards organizations that choose higher-assurance tools and follow best-practice governance.

4. Make regulatory intelligence integration a medium-term certification expectation.

Once standardized regulatory data publication infrastructure exists, ONC could require that certified AI tools demonstrate an ability to detect and adapt to relevant regulatory changes via those standardized feeds.

This would ensure that certified tools remain aligned with current rules, not just the rules in effect at the time of their initial evaluation.

NOETED is designing its platform to be “certification-ready” along these dimensions and would welcome collaboration with ONC and credentialing bodies to pilot certification criteria specifically for regulatory and compliance AI.

Q6

Where have AI tools in clinical care met or exceeded performance and cost expectations, and where have they fallen short? What novel AI tools would have the greatest potential to improve health care outcomes, quality insights, and cost reduction?

NOETED Response

Areas where AI has met or exceeded expectations.

- **Ambient clinical documentation:** AI-assisted transcription and note generation has shown clear value in reducing clinician burden without requiring autonomous clinical judgment. Adoption has been strongest where integration with EHR workflows is tight and governance is clear.
- **Imaging and narrow diagnostic tasks:** In radiology and pathology, AI models trained on well-curated datasets have demonstrated high performance on specific classification tasks, with emerging evidence of real-world impact when embedded thoughtfully into workflows.

Areas where AI has fallen short.

- **Generalization and equity:** Performance often degrades for under-represented populations (racial/ethnic minorities, rural patients, rare conditions), especially when training data do not reflect real-world diversity.
- **Dynamic regulatory and policy environments:** AI tools that touch billing, utilization management, or compliance often lag behind rule changes. Here NOETED has seen two failure modes:
- **Update lag:** Developers or implementers do not update logic promptly when rules change.
- **Infrastructure failures:** Even when developers try to update, the necessary regulatory data are not available in timely, machine-readable forms. This is especially problematic for smaller organizations that lack dedicated regulatory intelligence staff.

High-potential AI tools for outcomes, insight, and cost.

- **Regulatory intelligence and compliance monitoring:** AI systems that continuously track regulatory changes, map impacts across an organization's obligations, and surface actionable guidance for compliance, risk, and clinical operations. These tools are "infrastructure AI" that make other AI deployments safer.
- **Predictive compliance and integrity risk scoring:** Models that identify organizations, service lines, or populations at elevated risk of quality or compliance failures before they manifest as

adverse events, denials, or enforcement. When paired with supportive interventions, this can improve safety and reduce avoidable cost.

- **AI-assisted care coordination and prior authorization:** Tools that align clinical documentation, coverage rules, and medical necessity criteria in real time, reducing administrative friction and delays in care.

HHS support for these categories—through pilots, reimbursement innovation, and inclusion in interoperability and certification efforts—would help direct AI innovation toward areas with high patient and system benefit, not just administrative efficiency.

NOETED Response

From NOETED's perspective, the most promising AI uses for improving care, equity, and access share three attributes: they reduce preventable variation, they make complex rules navigable for frontline teams, and they are explicitly designed to work in resource-constrained settings. Examples include:

- **AI-supported guideline and coverage alignment at the point of care**
 - Tools that help clinicians document and order in ways that are simultaneously clinically appropriate and aligned with coverage criteria, reducing avoidable denials and care delays—especially important for safety-net providers and Medicaid populations.
- **Proactive compliance and quality risk detection**
 - AI that identifies emerging patterns (documentation gaps, ordering patterns, access bottlenecks) that put certain populations at higher risk of poor outcomes or coverage disruptions, enabling targeted interventions before harm occurs.
- **Navigation and communication tools tailored for underserved communities**
 - AI-driven patient communication, translation, and navigation tools designed with community input, that respect cultural and linguistic needs and help patients understand and act on care plans and benefits.

HHS can promote these uses by:

- Prioritizing them in demonstration projects and payment innovation models.
- Including equity and access metrics as core evaluation endpoints in AI funding opportunities.
- Ensuring that regulatory and data infrastructure investments explicitly consider the needs of rural, tribal, and safety-net organizations that lack extensive IT and compliance staff.

NOETED's regulatory intelligence work is directly aligned with this: smaller organizations benefit the most when complex rules are translated into clear, actionable guidance that can be operationalized without large legal departments.

Q8

Where would enhanced interoperability widen market opportunities, fuel research, and accelerate AI development for clinical care? Consider specific data types, data standards, and benchmarking tools.

NOETED Response

Interoperability has unlocked new AI possibilities wherever standardized data exist (e.g., FHIR for clinical data). To widen opportunities and accelerate AI, HHS should prioritize a third domain alongside clinical and administrative data: regulatory data interoperability.

Current gaps.

- Regulatory content—conditions of participation, coverage decisions, quality specifications, compliance guidance, enforcement summaries—is often available only as narrative documents or ad hoc web content.
- APIs, where they exist, use inconsistent structures and are not uniformly documented or maintained.
- As a result, AI tools that depend on “knowing the rules” must build custom scrapers and mappers for each source. Smaller organizations cannot realistically do this.

Standards and infrastructure HHS could catalyze:

1. Standardized regulatory change notifications

- A common schema (potentially an implementation guide aligned with FHIR resources) for publishing regulatory and policy change events, with required metadata: effective date, affected entity types, relevant citations, superseded guidance, and unique identifiers.

2. Structured program and quality specifications

- Publishing measure definitions, participation requirements, and program guidance as versioned, machine-readable resources rather than PDFs.

3. Enforcement and oversight data

- Providing de-identified, structured data on enforcement actions and recurring findings to support AI-driven risk modeling and benchmarking.

Recommended HHS actions:

- Direct ONC to work with CMS, OIG, and other agencies on a Regulatory Data Implementation Guide, leveraging FHIR where appropriate.
- Treat regulatory data infrastructure as a core component of interoperability roadmaps and implementation funding.

With such standards in place, private-sector platforms like NOETED can build robust, reusable regulatory intelligence capabilities that every health system, plan, and safety-net provider can leverage—dramatically lowering the barrier to safe, compliant AI deployment.

NOETED Response

Safety in AI-assisted clinical care depends on more than model performance metrics; it requires governance, transparency, and continuous alignment with evolving standards. HHS can strengthen safety by:

- **Requiring risk-appropriate governance frameworks.**
 - For higher-risk uses (diagnosis, treatment selection, utilization management), require documented governance structures: accountable owners, oversight committees, policies for model selection, and procedures for handling incidents and overrides.
- **Mandating documentation of data provenance and limitations.**
 - AI tools should clearly describe training data characteristics, known limitations, intended use, and populations for which performance is uncertain, in language clinicians and compliance teams can understand.
- **Embedding safety checks into certification.**
 - **ONC certification for AI-enabled systems should include:**
 - Support for capturing clinician and staff feedback and overrides.
 - Logging sufficient to reconstruct AI-assisted decisions in the event of an adverse event or audit.
 - Mechanisms to suspend or adjust AI behavior when performance concerns arise.
 - **Aligning AI tools with regulatory intelligence.**
 - For compliance-sensitive workflows, a safe AI tool must be able to detect when relevant rules change and flag potential misalignment. This is where regulatory data interoperability and platforms like NOETED are critical.

HHS's role is to set these expectations through guidance, certification criteria, and funding conditions, while allowing innovation in specific technical approaches. NOETED stands ready to support this by demonstrating how regulatory intelligence can be integrated into AI governance for real-world compliance use cases.

NOETED Response

HHS should prioritize AI research that addresses **foundational infrastructure gaps, equity in real-world settings, and system-level integration**—areas where AI can deliver the greatest clinical and economic impact. Specific priorities include:

1. Regulatory intelligence and compliance monitoring

- Research on AI systems that continuously ingest, normalize, and operationalize regulatory, policy, and program data (CMS rules, OIG guidance, quality specifications, enforcement patterns).
- Focus: mapping obligations to workflows, predicting compliance risk by population/service line, and auto-generating remediation plans.
- Impact: Enables safe AI deployment across compliance-sensitive workflows (documentation, utilization management, quality reporting) by ensuring tools stay aligned with current rules.

2. Equity-focused model robustness and generalization

- Research on techniques to ensure AI performance across under-represented demographics, geographies, and clinical contexts (rural patients, minority populations, rare conditions).
- Focus: adversarial robustness testing, domain adaptation for distribution shift, and synthetic data generation to augment limited real-world datasets.
- Impact: Prevents AI from deepening disparities and supports equitable care in safety-net and rural settings.

3. AI-assisted care coordination and access optimization

- Research on AI for prior authorization, documentation alignment with coverage criteria, and navigation tools tailored to underserved communities.
- Focus: reducing administrative friction, aligning clinical decisions with coverage realities, and supporting culturally/linguistically appropriate communication.
- Impact: Improves access, reduces delays, and lowers cost for Medicaid, Medicare Advantage, and safety-net populations.

4. Post-deployment monitoring and continuous learning

- Research on automated mechanisms for detecting drift, bias emergence, and performance degradation in deployed clinical AI, including clinician override patterns and downstream outcomes.
- Focus: integration with regulatory change detection so models adapt to evolving clinical and policy standards.

Published findings on AI impact and costs/benefits

- **Ambient documentation:** Reduces clinician burden by 20–40% (e.g., Nuance DAX, Suki); studies show 1.5–2.5x ROI from reclaimed time and improved throughput.
- **Imaging diagnostics:** Systematic reviews show AI improves specificity by 5–15% in radiology/pathology, with cost savings from reduced false positives (e.g., reduced downstream procedures); ICERs often < \$10,000/QALY.
- **Chronic disease prediction:** Models for AFib, diabetic retinopathy screening show 14–20% cost reduction and QALY gains (e.g., Xie et al., 2023; ICERs ~\$1,100/QALY).
- **System-level savings:** NBER (2023) estimates 5–10% U.S. healthcare spending reduction (\$200–400B annually) from broader AI adoption, primarily through efficiency and reduced unnecessary care.

HHS can advance these priorities through LEAP funding, cooperative agreements for multi-site pilots, and prize challenges for well-defined benchmarks. NOETED is prepared to contribute regulatory intelligence capabilities to such research efforts.

Conclusion

NOETED thanks HHS and ONC for this opportunity to contribute to the future of AI in clinical care. Our recommendations center on a single, high-leverage opportunity: treating **regulatory data infrastructure** with the same seriousness as clinical data infrastructure. By modernizing publication standards, embedding regulatory intelligence into certification and safe harbors, and funding practical pilots, HHS can unlock AI's potential to improve care, equity, and efficiency across all provider types.

We are eager to partner with ONC and HHS to make this vision real—through pilots, standards co-development, and real-world demonstrations of how regulatory intelligence enables safe, scalable AI. Please contact us at duv@noeted.com to discuss collaboration opportunities.

Respectfully submitted,

Keith “Duv” Williams

Founder, NOETED

noeted.ai

duv@noeted.com