

Fairview

February 23, 2026

The Honorable Thomas Keane, MD, MBA
National Coordinator for Health Information Technology and
Assistant Secretary for Technology Policy
U.S. Department of Health and Human Services
330 C St SW
Floor 7
Washington, DC 20201

Submitted electronically via www.regulations.gov

Re: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care (RIN 0955-AA13)

Dear Dr. Keane,

Fairview Health Services (Fairview) appreciates the opportunity to provide comments to the Assistant Secretary for Technology Policy (ASTP) and Office of the National Coordinator for Health Information Technology (ONC) on the Request for Information (RFI) on artificial intelligence (AI) adoption in clinical care.¹ Fairview appreciates the Administration's ongoing commitment to tech-enabling healthcare to reduce administrative burden while improving access to high-quality care. Fairview supports development of a responsible, innovation-friendly, and transparent regulatory framework for AI in healthcare that improves patient health outcomes and supports clinicians through technology. Clear regulatory guardrails, strong data privacy protections, robust workforce enablement, and industry aligned- standards are essential for safe and effective AI adoption.

Fairview is a nonprofit, integrated health system providing exceptional health care across the full spectrum of medical services through our 10 hospitals and 60 clinics. Fairview's broad care continuum includes academic and community hospitals, primary and specialty care clinics, senior and long-term care facilities, retail and specialty pharmacies, pharmacy benefit management services, rehabilitation centers, counseling, home health care services, and an integrated provider network. Fairview's 34,000+ employees and 5,000+ system providers embrace innovation and new thinking to drive a healthier future through healing, discovery, and education.

AI can deliver meaningful benefits to providers and patients such as improved clinician productivity, reduced documentation burden, lower burnout, more timely clinical decision-making, fewer missed orders, improved throughput, and better patient engagement. These gains slowly translate into

¹Regulatory text available here: <https://www.federalregister.gov/documents/2025/12/23/2025-23641/request-for-information-accelerating-the-adoption-and-use-of-artificial-intelligence-as-part-of>

Fairview

reduced staffing pressure and overtime costs, higher-quality documentation, improved access and satisfaction, and better clinical outcomes through earlier interventions.

However, clinical leadership can be reluctant to adopt and scale AI. Because generative AI is probabilistic, its performance must be validated not only before deployment but also continuously over time, requiring skills and resources that may be scarce in healthcare settings. The costs, which are driven by computing, storage, and model usage, are also highly variable and difficult for organizations to predict, budget, or govern using traditional software pricing models. Realizing the benefits of AI requires healthcare leaders to move past understandable initial hesitation and lean into AI capabilities. The government needs to make sure policies help foster adoption and keep health systems and providers in control of how AI is deployed within healthcare while also protecting patients and health systems from unregulated AI companies and tools that may cause significant harm.

I. Documenting Return on Investment

One barrier to adoption and scaling of AI is determining return on investment (ROI). Healthcare organizations operate today on thinner margins than ever, making it difficult to invest in technologies that will not pay off financially in the near term. Fairview believes many applications of generative AI in healthcare do present sufficient return on investment (ROI).

The Department of Health and Human Services (the Department) could help document ROI data in a way that healthcare providers can both rely on and act on this data.

For example, studies show that AI ambient scribe technology, in which AI records and transcribes clinical interactions and then drafts corresponding notes, can reduce the time clinicians spend working after hours. Less administrative time for providers, however, does not by itself offset the cost of the technology. The true ROI of AI scribes appears downstream. For example, through happier clinicians who are less likely to burn out or leave healthcare, reducing recruitment and onboarding costs. AI summarization tools that help providers and staff find and summarize clinical information, agents that help do lower risk highly standardized workflow like wound care, and even clinical decision support tools that assist providers in making decisions are all potential use cases that could benefit from this kind of ROI standardization.

The Department should address this gap by assembling a repository of AI use cases with demonstrated ROI. By conducting independent research, such a repository would validate vendor claims and highlight metrics relevant to healthcare administrators' bottom-line decisions. Such a repository would be vendor-agnostic, focusing on use cases rather than specific implementations, to not advantage those vendors that get their claims published. Done right, it would serve as a trusted reference for organizations evaluating AI investments.

The Department should also address price controls on AI and the need for lower cost options to make sure IT companies do not increase prices to capture all the ROI that may be gained by health systems.

Fairview

Prioritizing low-cost options or price controls should help health delivery organizations that need these efficiencies to offset the very small margins, financial pressures, and workforce challenges that are coming to all health care delivery systems.

II. Bridging the Adoption-to-ROI Gap and Establishing Payment

Even when an AI use case has a proven ROI, initial adoption may still be financially challenging. The benefits of AI are rarely immediate and often require workflow changes, training, and iteration to be fully realized. Early phases of adoption are typically marked by uneven usage, ongoing configuration, and infrastructure costs that scale before productivity gains materialize. As a result, health systems may incur real costs before seeing offsetting savings, creating a financial mismatch that can slow or stall adoption.

A response to these challenges would be fee-for-service reimbursement for AI use, but this approach could quickly become costly and, for AI with proven ROI, unnecessary. Organizations primarily need support through the initial scaling period; once broad adoption is achieved, AI that delivers value will likely pay for itself assuming it is done well and appropriate price controls are in place for AI companies.

Fairview respectfully asks the Department to consider AI adoption grants that cover infrastructure costs for a fixed period of time with a discreet infrastructure delivery goal rather than a fixed amount of money without a concrete infrastructure deliverable.

Traditional grants incentivize slow scaling to extend funding; time-based grants with clear implementation goals would encourage organizations to scale quickly, assess value, and determine which use cases merit continued investment. After the grant period ends, organizations could discontinue low-value AI while sustaining AI that delivers sufficient benefit using the infrastructure and learnings that were created – without ongoing government spending.

Such programs could be structured as true grants or as fixed-price arrangements. In the latter case, however, pricing should scale with organization size to avoid disadvantaging smaller or lower-resourced providers. If desired, pricing could even be structured disproportionately to give an advantage to organizations that are caring for underserved or rural areas. Properly designed, time-based infrastructure support would help organizations escape the “pilot trap” and reach sustainable AI adoption.

Additionally, performance-based contracting arrangements between providers and healthcare vendors – and data collected on AI and software performance – can help to decrease the go-to-market time and cost for new AI solutions, further extending America’s lead in innovative technology. Performance-based contracts between healthcare vendors and providers can ensure the new technologies purchased by providers deliver the promised benefits and return on investment.

Fairview encourages ASTP/ONC to work with CMS to reduce regulatory barriers to AI and software approval and payment.

Fairview

HHS should explore ways to allow data collected from performance-based contracting arrangements to offer a pathway to CMS coverage and/or reimbursement decisions for novel digital health technologies and Software as a Service (SaaS) solution. Hospitals have already entered performance-based contracts with innovative new healthcare suppliers, including AI technology and SaaS vendors. By making payment dependent on predefined outcomes, these contracts make it easier for hospitals to take a chance on technologies that lack extensive histories of clinical evidence or claims data. That is, if the purported return on investment fails to materialize, hospitals can recoup initial costs.

III. Building AI Fluency

Technology adoption depends not only on technical readiness, but also on human understanding and trust. Considerable attention has been given to training clinicians on how to use AI tools effectively, verify outputs, and avoid automation bias. While this training is essential, a more foundational form of AI fluency is also needed to overcome barriers to adoption—particularly among non-technical stakeholders such as clinical leadership and frontline staff.

In healthcare, AI is unlikely to displace workers and more likely to enable healthcare workers in efficiency and automation if done appropriately. Persistent labor shortages across clinical roles, combined with rising demand driven by an aging population, mean that AI is increasingly necessary to sustain access to care. In practice, most healthcare AI use cases are designed to reduce administrative burden, support clinical judgment, and allow clinicians to spend more time with patients. Nonetheless, the clinical workforce has experienced significant difficulty integrating AI tools within existing clinical workflows, EHR systems, and operational infrastructure. AI tools often must conform to complex and aging infrastructure, requiring expensive adaptation and reconfiguration. Many tools fail to accommodate workflow constraints faced by clinicians, leading to low adoption and poor usability.

These concerns often stem from a lack of shared understanding about what healthcare AI does and does not do well and how to use it safely and efficiently. The most effective way to build AI fluency is through experience within a safe framework. When clinicians and staff see AI reduce documentation burden, improve workflow efficiency, and support—not override—clinical judgment; skepticism gives way to trust. Policies that enable earlier, broader, and more responsible adoption of AI will therefore do more to build AI fluency than education alone. By helping healthcare organizations move past prolonged pilots and into sustained use, the Department can ensure that AI's benefits are visible, tangible, and widely shared.

In combination with AI experience, credible, accessible education, can mitigate the perceived risks that AI adoption is a cost-cutting measure rather than a capacity-multiplying one. This perception can slow implementation, limit usage, or prompt organizations to constrain AI to narrow pilots rather than deploy it broadly where it could deliver the most benefit.

The federal government can play a significant role in improving AI adoption by supporting neutral, trusted education on AI's role in healthcare and promoting this training in education systems and quality improvement incentives.

Fairview

This could include developing materials tailored to different stakeholder groups; developing repositories of the best materials, supporting development of AI certifications in healthcare, supporting peer-to-peer learning networks among healthcare organizations; and incorporating AI fluency into existing workforce development training programs and incentivizing AI development and governance within quality improvement programs. Fairview urges HHS to invest in education programs, establish best practices for human AI- collaboration and promote AI tools that measurably reduce administrative burden.

IV. AI Supporting Clinical Trials

AI can be deployed to address the challenges related to identifying and enrolling patients in clinical trials in a timely manner.

Delays in meeting trial enrollment targets and timelines can increase the cost of the trial. AI tools can analyze the extensive universe of data available to health systems, like Fairview, to identify patients that may be a match for clinical trials that are currently recruiting. This application of natural language processing (NLP) systems can make developing new drugs less expensive and more efficient, while also improving patient and geographical diversity in trials to address generalizability.

To incentivize continued innovation in clinical trials, Fairview also encourages ASTP/ONC to work with the FDA to put into effect clear guidance on the process for properly obtaining consent from patients for the use of their real-world data (RWD) to produce AI-generated synthetic control arms in clinical trials.

AI, once trained on real-world data (RWD), has the capability to generate synthetic data and patient profiles that share characteristics with the target patient population for a clinical trial. Fairview would be able to use synthetic data to simulate clinical trials to optimize trial designs, model the possible effects or range of results of a novel intervention, and predict the statistical significance and magnitude of effects or biases. Ultimately, synthetic patient data can help optimize trial design, improve safety and reduce costs for decentralized clinical trials. Further, synthetic control arms in clinical trials can help increase trial enrollment by easing patients' fears that they will receive a placebo.

V. Making Local Validation Possible While Holding AI Companies and EHR Vendors Accountable

Another concern AI adopters rightly have is safety. Many healthcare AI use cases fall outside the oversight of the Food and Drug Administration (FDA), and even regulated products are poorly served by traditional evaluation methods. Especially for AI, which is probabilistic, pre-deployment testing on generalized datasets has limited predictive value for real-world performance.

The Department needs to assure the method for validating AI performance in healthcare is "local validation," in which the healthcare facility, independent of the developer, evaluates how the AI performs on its own data while embedded in its own workflow while the Department also puts policies in place to hold AI companies and EHR vendors accountable to collect this feedback and build safe, high value AI tools that promote patient care.

Fairview

This can be done both prior to deployment, using local but retrospective data, and on a continual basis post-deployment (“ongoing monitoring”) to make sure performance does not “drift” over time. This way, AI features can be deemed safe no matter what patient acuity mix is being served by a particular healthcare organization.

Local validation also matters because AI performance can be sensitive to the workflow into which the AI is integrated. For example, a popular use case for predictive modelling forecasts the likelihood of imminent medical deterioration. The output of such a model is not binary but instead a probability. When implementing this model, a deployer must decide what probability threshold should trigger an alert for the care team to investigate further. This choice has a significant impact on the “performance” of the model but also has a major impact on the clinical teams and the burden of the alerts. If an organization sets the alert threshold low, there will likely be more false positives; if the organization sets the threshold higher, there will likely be more false negatives. Yet the organization may have good reason to change this setting: organizations with dedicated remote monitoring teams may prefer false positives, while those who must rely only on frontline nurses managing a large caseload of patients may not be able to tolerate them.

The challenge for post-deployment validation, in healthcare at least, is that small or rural healthcare organizations often struggle to afford the necessary data science resources and expertise to perform effective monitoring. Yet because AI performance will vary from organization to organization, post-deployment validation is not something that can be outsourced in practice. Without appropriate support for “local validation” activities, AI could widen, rather than narrow, the digital divide in healthcare.

The Department needs to make sure this local AI validation data is shared with vendors and other systems in a way that allows all people to benefit from this validation and requires AI vendors and EHR vendors to implement this feedback and report safety issues related to their AI products.

VI. Providing Guidance for AI Validation

Even with the tools and resources needed to conduct local validation, healthcare organizations may not know how to use them effectively. Today, there is little consensus on how to evaluate traditional predictive models, and even less for generative AI. Gold-standard guidance such as the National Institute of Standards and Technology’s AI Risk Management Framework asserts the importance of ongoing monitoring but does not specify how to perform it. Vendors providing evaluation services may use proprietary metrics or allow the user to select among evaluation methodologies (as the AI Trust and Assurance Suite does); both limit comparability.

Regulators can add value by defining standard metrics and validation processes for AI performance. Such standards and scorecards should likely be mandatory; but if not mandatory, there should at minimum standard guidance that could increase confidence among healthcare AI adopters and lower the barrier to widescale adoption by setting and updating those standards. Clear, shared methodologies would allow organizations to focus on performing validation rather than designing it, especially if incorporated into existing tools.

Fairview

Importantly, validation guidance should be specific to both the use case and the deployment context, not one-size-fits-all. The risks, benefits, and appropriate metrics for documentation assistance, clinical decision support, predictive modeling, and patient engagement tools differ substantially. Tailored guidance for common classes of AI applications would allow organizations to focus their efforts on meaningful evaluation rather than reinventing validation frameworks from scratch.

Clear federal guidance would also accelerate adoption by lowering the perceived risk of deployment. Many healthcare leaders delay or limit AI use not because they believe AI is unsafe, but because they lack confidence that they are evaluating it correctly. By clarifying expectations and establishing shared norms, the Department could help organizations move from cautious experimentation to responsible scale.

VII. Setting National Privacy Standards

A significant barrier to the large-scale diffusion of innovative AI applications in healthcare is the lack of a single, comprehensive federal privacy law.

A cohesive national privacy framework and patient education about data privacy is necessary to provide a reasonable level of consumer confidence that businesses will protect their data, sensitive or otherwise and be held accountable for data breaches.

The current patchwork, state-driven approach to privacy policy has resulted in inconsistent data privacy practices. Patients must be able to trust the processes for storage, handling and use of their data, particularly as patient data is increasingly used to train AI algorithms. Missing data from patients wary of data sharing could add to data siloing and perpetuate AI algorithmic biases, and ASTP/ONC should work with data stakeholders across HHS to proactively educate patients on these tradeoffs.

Unharmonized and burdensome requirements are a rate-limiting step in unlocking value through AI innovation. A federal privacy law that incorporates AI related data sharing benefits and continued data sharing concerns can empower innovators by placing clear, harmonized and common-sense guardrails around artificial intelligence tools and the data sharing that drives the benefits of holistic patient care across health entities. Healthcare privacy enforcement authorities are currently diffused across a wide swath of the federal government. The Health Insurance Portability and Accountability Act (HIPAA) only delegates privacy enforcement authority to the Department of Health and Human Services Office of Civil Rights (OCR) for a narrow subset of covered entities. With the proliferation of a healthcare data ecosystem outside of the traditional payer-provider exchange, enforcement authority falls to the Federal Trade Commission, the Department of Justice and other agencies for a broad range of privacy- and data-related violations. There is significant, urgent need for a unified baseline set of enforcement priorities to create a harmonized enforcement authority for the portion of healthcare ecosystem outside of the bounds of HIPAA.

Now, as state AI laws follow a similar path toward fragmentation, it is critical that ASTP/ONC and its regulatory partners either 1) offer directional federal guidance or 2) endorse a private sector-developed framework for AI security, transparency, safety and quality as the recognized standard. Fragmented regulatory approaches across agencies and states impede adoption and increase compliance burdens.

Fairview

VIII. Improving AI Transparency

Trust among patients, providers, payers, and suppliers is critical to the development and deployment of AI tools in healthcare settings. To earn trust, AI tools must have an established standard of transparency. Some policy proposals, including those offered by ONC, suggest transparency can be achieved through a “nutrition label” model, which lists the sources and classes of data used to train the algorithm. Unfortunately, some versions of the “nutrition label” approach to AI transparency fail to acknowledge that when an AI tool is trained on a large, complex dataset, and is by design intended to evolve and learn, the initial static inputs captured by a label do not provide accurate insights into an ever-changing AI tool. Further, disclosure requirements around data inputs or algorithmic processes, if not carefully considered, could force AI developers to publicly disclose intellectual property or proprietary technology, which would stifle innovation.

Fairview believes that AI technology in healthcare should be held to a standardized, outcomes-focused set of metrics, such as accuracy, false positives, inference risks and recommended use/applications. Outcomes, rather than inputs, are where AI technologies have potential to drive health or harm. Thus, it is essential to focus transparency efforts on the accuracy, reliability, and overall appropriateness of AI technology outputs in healthcare to ensure that the evolving tool does not produce harm.

Fairview encourages the Department to consider requiring AI developers and manufacturers to manage and maintain model cards, including uses of new technologies in healthcare settings and the associated safety concerns and mitigation practices to provide much-needed guidance to clinicians, providers, and administrators on safe and appropriate use cases with the associated risks and mitigations.

This approach could provide liability protection for the proper use of AI technology for the defined set of use cases where developers have established and reported the appropriate metrics for accuracy and reliability. Transparency about the intended use of AI tools would be the simplest way for regulators to incorporate AI governance into existing regulations. Health systems would also be able to incorporate this information into our own governance structures, putting internal policies in place to prevent misuse of AI in ways that could be detrimental to patient safety or experiences.

Such disclosure would give health systems and patients a complete picture of the safety and security of the AI technologies they use. Rather than limiting or delaying innovation, such guidelines would level the playing field between established market leaders and startups. Alternatively, or in addition to this, HHS could sanction the use of third-party certification organizations or existing market processes to address this challenge while reducing administrative burden. Such a market-driven approach would be flexible and adaptable, capable of adjusting to the latest developments in AI technology without requiring regulators to frequently revisit the framework, although it would need some appropriate oversight.

Fairview

IX. AI Patient Coaches Need Provider Controls

The Department needs to ensure appropriate guardrails which enable oversight by providers or caregivers for all AI patient coach or clinical decision-making tool.

As AI tools increasingly interact directly with patients—through symptom checkers, care navigation tools, medication guidance, and longitudinal “patient coaches”—it is essential that these technologies operate within a framework that preserves the primacy of the clinician–patient relationship. AI patient facing tools must augment, not replace, the role of licensed providers, and should reflect the clinical judgment, care plans, and values of a patient’s personal care team. Without appropriate safeguards, AI patient coaches risk delivering guidance that is misaligned with a patient’s medical history, local standards of care, or clinician intent. Generic or overly confident recommendations may confuse patients, undermine trust, or prompt inappropriate care seeking behavior. For these reasons, AI patient coaches should be required to operate with explicit provider oversight and clearly defined boundaries on autonomy.

At a minimum, policies should ensure that providers maintain control over:

- **Whether and when AI patient coaches are used** for their patients;
- **Scope of advice or guidance** such tools are permitted to provide;
- **Escalation pathways** that route complex, ambiguous, or high-risk situations back to human clinicians; and
- **Alignment with the patient’s established care plan**, including medications, diagnoses, and documented clinical preferences.

AI patient facing tools should be designed to reflect what a patient’s own clinician would reasonably recommend in similar circumstances, rather than offering generic, population level advice. This requires mechanisms for local configuration, integration with the clinical record, and clinician defined constraints on recommendations. Where such alignment is not possible, AI tools should clearly defer to human review rather than attempting to fill gaps with probabilistic output.

Transparency is also critical. Patients should be informed when they are interacting with an AI system, understand its intended role, and know how to reach a human clinician when needed. AI patient coaches should not present themselves as substitutes for professional medical advice, nor should they operate independently of provider governance structures.

Finally, liability and accountability frameworks must reflect these realities. Developers and deployers of AI patient coaches should be required to clearly specify intended use cases, known limitations, and appropriate safeguards and work directly with health systems and providers instead of developing straight to consumer applications. Health systems must retain the authority to govern these tools within their existing clinical oversight, quality, and safety processes.

Establishing clear federal expectations for provider-in-the-loop AI patient tools will protect patients, support clinicians, and enable responsible innovation—ensuring that AI enhances care delivery without eroding clinical judgment or trust.

Fairview

X. Conclusion

Thank you for consideration of Fairview's comments on ASTP/ONC RFI on artificial intelligence (AI) adoption in clinical care. If you have any questions or need more information regarding these comments, please do not hesitate to contact me at Bryan.Jarabek@fairview.org.

Sincerely,

A handwritten signature in black ink, appearing to be 'BJ' with a stylized flourish.

Bryan Jarabek, MD, PhD
Chief Medical Informatics Officer
Fairview Health Services