

The Honorable Robert F.
Kennedy Jr.
Secretary
U.S. Department of Health
and Human Services

February 23, 2026

Re: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as part of Clinical Care (Docket ID HHS-ONC-2026-0001)

Dear Secretary Kennedy,

ServiceNow appreciates the opportunity to respond to the Department's initiative to support the responsible adoption of Artificial Intelligence in clinical care. We share the Department's desire to reduce healthcare costs, eliminate administrative waste, and give time back to healthcare workers.

As a platform company serving thousands of healthcare organizations globally, we have observed a consistent pattern: the barrier to realizing AI's deflationary potential is rarely the algorithm itself. It is the **fragmented workflow environment** into which that algorithm must be deployed. Today's health systems operate dozens - sometimes hundreds - of disconnected applications. Each new AI tool risks becoming yet another system clinicians must learn, log into, and manually reconcile with their existing work.

This fragmentation creates three compounding problems that any federal strategy must address:

1. **The Cognitive Load Problem.** Clinicians already suffer from alert fatigue and application overload. Point-solution AI tools that operate outside the primary workflow add to this burden rather than reducing it. Certain data and workflows, such as automatic work orders, incidents, and tickets, can in some cases be created and resolved before burdening clinical or ancillary personnel -effectively resolving issues before human staff are even aware they exist.
2. **The Governance Gap.** Health systems are deploying AI tools faster than they can track them. Many organizations lack a centralized inventory of which AI models are in use, what data they access, how they were validated, or who is accountable when they fail. This governance vacuum creates liability uncertainty that chills adoption and leaves patient safety to chance. Health systems need a "control tower" that provides visibility and oversight across all AI assets.
3. **The Data Silo Problem.** Current interoperability standards focus almost exclusively on clinical data - patient records, lab results, medications. But the operational inefficiencies that drive healthcare costs live in a different data layer: staffing schedules, bed availability, equipment status, supply chain inventory, IT service queues. Until AI can access and act on this operational data in real time, we will not achieve the systemic efficiency gains the Department seeks.

The solution to these challenges is not more point solutions. It is a **unified workflow platform** that embeds AI directly into the operational fabric of health system delivery, connecting clinical and administrative processes, providing centralized governance, and enabling real-time action on enterprise-wide data, while aligning with HHS/CMS safeguards.

In the responses that follow, we offer concrete recommendations drawn from our experience enabling workflow transformation across healthcare, with the goal of helping HHS design policies that unlock AI's deflationary potential while maintaining the safety and trust that patients deserve.

We believe that workflow-oriented platforms - which embed AI into operational processes, provide centralized governance, and enable real-time action on enterprise-wide data - are essential to realizing these goals. We welcome further dialogue with the Department on these issues.

Sincerely,



Combiz Abdolrahimi
Vice President and Global Head of Government Affairs and Public Policy
ServiceNow

Responses to Specific Questions

1. What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care? [MEDIUM PRIORITY]

Operational Data Fragmentation and the “Last Mile” Problem. The biggest barrier is not a shortage of AI innovation - it is the inability to deliver that innovation into operational workflows where it can create value. While clinical workflow integration is relatively well-understood (most health systems operate on monolithic EMR platforms), the real challenge lies in **operational data fragmentation**. Administrative data remains “trapped” in dozens - sometimes hundreds - of siloed legacy systems that do not communicate with each other. Each application has its own login, interface, and data model.

To take real cost out of the system using AI, operational workflows need to be totally reimagined - but health systems are limited by these legacy systems. When a new AI tool is introduced, it typically arrives as yet another standalone application. Users must leave their primary workflow, open the AI tool, manually input or transfer data, interpret the output, return to their original system, and take action. This “swivel chair” workflow negates the efficiency gains AI promises. Studies consistently show that clinicians and administrators ignore alerts and recommendations that arrive outside their primary workflow, regardless of validity.

The second major barrier is **governance uncertainty**. Health system executives are hesitant to adopt AI tools when liability frameworks are unclear. Questions like “Who is responsible if the AI is wrong?” and “How do we audit AI decisions for compliance?” remain unanswered. Without clear governance standards, legal and compliance teams default to blocking adoption.

The path forward requires platforms that can *embed AI into existing workflows* rather than adding new ones, governance frameworks that provide *clear accountability without stifling iteration*, and policy that forces operational systems to openly share data.

2. What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? [LOW PRIORITY]

Reward Efficiency Outcomes, Not Just Clinical Outcomes. Current payment models largely reward volume and discrete clinical interventions. They do not reward the operational efficiencies that AI is best positioned to deliver - reduced administrative time per encounter, shorter patient wait times, optimized resource utilization, or decreased length of stay through better care coordination.

Reimbursement models are fundamentally misaligned - they pay for procedures, not outcomes or efficiencies. Fee-for-service doesn't reward AI-driven prevention or automation, so tools that could significantly reduce labor costs get shelved. Incentives favor human hours over smart systems. This keeps AI adoption uneven: 71% of hospitals use predictive AI in EHRs, but small and rural facilities lag below 50%, according to 2024 American Hospital Association IT Survey (link [here](#)).

HHS should consider payment adjustments or bonus structures for health systems that demonstrate measurable “deflationary” improvements attributable to technology adoption. Metrics might include: administrative cost per patient encounter, time from order to fulfillment, staff hours spent on non-clinical tasks, or patient throughput without quality degradation. Tying reimbursement incentives (e.g., value-based adjustments) to proven

efficiencies from AI tools would directly reward administrative burden reduction and cost savings.

Affirm the Self-Regulated Status of Administrative AI. Administrative AI solutions - those that automate scheduling, manage supply chains, route service requests, or optimize staffing - are effectively self-regulated today and do not require FDA oversight as medical devices. This is well-established and appropriate. HHS should publicly affirm this status, providing health systems confidence to adopt low-risk, high-efficiency tools without unnecessary compliance hesitation. Clear communication from HHS that administrative operational AI falls outside medical device regulation would remove a source of caution among legal and compliance teams who may otherwise treat these tools with the same scrutiny as clinical diagnostic devices.

3. For non-medical devices, what novel legal and implementation issues exist and what role, if any, should HHS play to help address them? [HIGH PRIORITY]

The Model Drift and Lifecycle Management Challenge. Unlike traditional software, AI models are not static. A model validated today may degrade over weeks or months as the underlying data patterns shift, a phenomenon called "model drift." A scheduling optimization model trained on pre-pandemic patient volumes, for example, may perform poorly when demand patterns change. This creates a novel liability question: if a health system deploys an AI tool that worked at validation but fails six months later, who bears responsibility? This may be the most consequential barrier to broader AI adoption.

The implementation challenge is that most health systems lack the infrastructure to monitor AI performance over time. They deploy tools and assume they continue working. There is no equivalent of a "check engine light" for AI degradation.

Additional implementation challenges compound this problem:

interoperability remains fragmented despite FHIR progress, with 60-70% of AI failures tied to poor quality and siloed data. Cybersecurity risks are also likely to escalate as non-device AI becomes potential attack vectors, with breaches continuing to rise at health organizations handling PHI.

Recommendation: Require Governance Infrastructure, Not Code-Level Regulation.

Rather than attempting to regulate the code or algorithms of AI tools - which would be impractical given the speed of iteration - HHS should require health systems to maintain **governance infrastructure** capable of:

- Maintaining a centralized inventory of all AI tools in use across the organization
- Tracking version history and configuration changes for each tool
- Logging AI-generated recommendations and subsequent human actions (audit trail)
- Monitoring performance metrics and flagging drift or degradation
- Enforcing human-in-the-loop approval workflows for high-stakes decisions

This "control tower" approach to governance creates clear accountability - the health system is responsible for maintaining oversight, and the governance infrastructure provides the evidence needed to demonstrate due diligence. Workflow platforms that unify AI governance with operational execution are well-suited to provide this infrastructure, as they already manage the process layer where AI recommendations become actions.

Additional HHS Roles to Consider: HHS could further support adoption through practical, non-regulatory actions:

- **Issue liability guidance** via ONC/ASTP on voluntary shared accountability frameworks or indemnification templates, drawing from NIST AI RMF principles, to reduce legal fears without new regulations
- **Expand HIPAA clarifications** or safe harbors for secure federated learning and transparency in adaptive systems, funding ONC pilots demonstrating privacy-preserving data flows
- **Enhance TEFCA and HTI rules** to mandate FHIR-based APIs for seamless AI integration, accelerating data standardization

This isn't about more bureaucracy - it's about engineering clarity and infrastructure, consistent with the Administration's push to remove barriers and lead in AI.

4. For non-medical devices, what are the most promising AI evaluation methods for clinical care? [MEDIUM PRIORITY]

Continuous Post-Deployment Monitoring Over One-Time Validation. Traditional medical device evaluation focuses on pre-market validation: prove the device works before it ships. This model is poorly suited to AI, which can degrade over time and whose real-world performance depends heavily on the specific environment in which it is deployed.

The most promising evaluation approach for AI - whether clinical or operational - is **continuous monitoring of workflow-level outcomes**. Rather than asking "Is the model accurate?" the evaluation asks "Is the model delivering the intended improvement?" Metrics should include:

- Time-on-task reduction for the targeted workflow
- Error rates or rework rates before and after AI deployment
- User adoption and sustained engagement over time
- Downstream impact on efficiency metrics (throughput, cost per unit)

Recommendation: Promote Transparency Through Standardized Disclosure. HHS should encourage development of standardized disclosure formats that require AI tool vendors to transparently report: training data sources, intended use cases, known limitations, validation methodology, and recommended monitoring metrics. This empowers health systems to make informed procurement decisions and sets expectations for ongoing performance evaluation. Such transparency standards can be promoted through guidance and best practices without requiring formal regulatory mandates. The goal is to create market expectations for disclosure that allow health systems to evaluate AI tools consistently.

5. How can HHS best support private sector activities to promote innovative and effective AI use in clinical care? [MEDIUM PRIORITY]

Define "Good Governance" and Let the Market Certify It. The most effective role for HHS is to define *what good AI governance looks like* without mandating specific technical implementations. By establishing clear requirements - audit trail retention periods, version control standards, performance monitoring minimums, human oversight protocols - HHS can create a standard that private accreditation bodies can certify against.

This approach has several advantages: it creates market demand for governance-capable platforms, allows for diverse technical approaches to meeting the standard, and avoids the slow pace of direct federal certification. Health systems would seek out platforms and tools that carry recognized governance certifications, creating a virtuous cycle that raises the bar across the industry.

Existing frameworks like HITRUST for security and SOC 2 for controls demonstrate that private-sector certification, anchored in federal expectations, can drive widespread adoption of best practices more quickly than direct regulation.

Accelerate Interoperability as Foundational Infrastructure. HHS should mandate and fund expansion of TEFCA/FHIR-based APIs specifically for AI tools - including ambient scribes, predictive operations, revenue cycle, and workflow automation. ONC-led pilots or grants could demonstrate scalable approaches.

The metrics make the case: where private tools get clean data, clear incentives, and neutral validation, adoption surges. Organizations exploring AI report moderate-to-high ROI, with clinician time savings of 2+ hours per day on documentation tasks and meaningful reductions in system costs. HHS's role isn't to build or dictate—it's to remove barriers, standardize the rails, and incentivize speed.

6. Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short? [HIGH PRIORITY]

Success: Operational and Administrative Automation. AI has consistently exceeded expectations when applied to **operational workflows** rather than clinical decision-making. Examples include:

- **Patient flow optimization:** Predicting admission volumes, optimizing bed assignments, and reducing emergency department boarding times
- **Supply chain automation:** Predicting inventory needs, automating reorder triggers, and reducing stockouts and waste
- **IT and HR service management:** Automating password resets, equipment requests, and routine HR inquiries that consume staff time
- **Security and compliance:** Identifying vulnerabilities, automating patch management, and detecting anomalous access patterns
- **Documentation assistance:** Drafting clinical notes, summarizing patient histories, and reducing after-hours charting burden.

These "deflationary" applications share a common characteristic: they remove friction from existing workflows rather than adding new decision points. They free human workers to focus on higher-value tasks. Health systems using unified workflow platforms have reported measurable improvements including reduced time-to-resolution for IT issues, faster onboarding, significant reductions in security vulnerabilities, and meaningful time savings for clinicians on administrative tasks.

Shortfalls: Interruptive Clinical Decision Support and Unclear ROI. AI has fallen short when deployed as *additive alerts* that interrupt clinical workflow without offering a clear path to action. A system that flags a potential sepsis risk but requires the clinician to navigate three screens to order the appropriate intervention creates friction, not efficiency. Alert fatigue is well-documented: when AI generates too many recommendations, clinicians learn to ignore them all. The lesson is clear - AI must be *embedded in the workflow*, not bolted on top of it.

Cost expectations have also gone unmet in many cases. While ambient scribes and documentation tools show clear value, many clinical AI tools deployed as point solutions lack demonstrated financial ROI - efficiency gains are often offset by new burdens from monitoring, updating, and managing liability exposure. Fee-for-service payment models

compound this problem by failing to reimburse preventive or efficiency-focused interventions, even when they reduce downstream costs. Tools that could prevent issues get shelved because the payment model doesn't reward prevention.

7. Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles? [HIGH PRIORITY]

The Shift from CMIO to CIO/COO. Historically, clinical informatics leaders (CMIOs, CMOs) drove health IT adoption based on clinical validity. For AI, the decision-making center of gravity has shifted. **Chief Information Officers (CIOs) and Chief Operating Officers (COOs)** now exert significant influence because AI adoption decisions increasingly hinge on:

- Integration feasibility: Can this tool connect to our existing systems?
- Total cost of ownership: What is the implementation and maintenance burden?
- Operational ROI: Will this reduce costs or improve throughput measurably?
- Security and compliance: Does this create new risk exposure?

CFOs are also critical gatekeepers, particularly for large investments, evaluating whether AI tools deliver genuine deflationary benefits or simply shift costs. **Emerging roles such as Chief Digital Officers and Chief AI Officers** are gaining influence in some health systems, though CIO and COO remain the primary decision-makers in most organizations today.

Primary Hurdle: The "Integration Tax." The single biggest administrative hurdle is the cost and time required to integrate each new AI tool with core systems, particularly the Electronic Health Record (EHR). Every point solution requires custom interface development, security review, workflow redesign, and staff training. This "integration tax" often exceeds the cost of the AI tool itself and creates a strong bias toward inaction.

Health systems that have adopted unified workflow platforms report significantly lower integration burden for new AI capabilities, because the platform provides a common integration layer and pre-built connectors to major clinical and operational systems. HHS can help by promoting open API standards for operational data (not just clinical data), reducing the custom development required to connect AI tools to health system infrastructure.

8. Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? [HIGH PRIORITY]

Beyond Clinical Data: The Case for Operational Interoperability. Current interoperability frameworks (FHIR, USCDI) focus on clinical data: patient records, medications, lab results, diagnoses. This is essential for care coordination but insufficient for operational efficiency. The data that drives healthcare costs, and therefore the data AI needs to achieve deflationary impact, often lives outside the clinical record:

- Staffing levels and shift schedules
- Bed availability and patient placement status
- Equipment location and maintenance status
- Supply chain inventory and consumption rates
- IT service requests and resolution queues
- Facilities status (room turnover, environmental conditions)

The Opportunity: “Command Center” Intelligence. If AI tools could access a unified, real-time view of both clinical and operational data, a new category of applications becomes possible. Imagine AI that can:

- Predict patient discharge readiness and automatically trigger room turnover and transport
- Match predicted admission volumes against staffing schedules and flag gaps before they create bottlenecks
- Identify supply consumption anomalies and adjust orders proactively
- Route non-clinical requests (IT, facilities, HR) to resolution without human triage

This “command center” model requires data liquidity across the entire enterprise—not just the clinical record.

Recommendation: Extend Information Blocking Rules to Operational Systems. HHS should consider extending information blocking rules to operational data systems. Just as EHRs are now required to share clinical data, legacy operational systems -scheduling, staffing, supply chain, facilities - should be required to share data via open APIs. This transparency-focused approach aligns with the Administration's philosophy of promoting openness over prescriptive requirements. Rather than dictating how health systems must use AI, it simply requires that the data AI needs be accessible. Workflow platforms that connect clinical and operational systems are well-positioned to enable this integration, providing the data fabric that AI requires to deliver systemic efficiency.