

February 20, 2026

Assistant Secretary Thomas Keane, MD

Assistant Secretary for Technology Policy / Office of the National Coordinator for Health Information Technology

Attention: HHS Health Sector AI RFI

Mary E. Switzer Building, Mail Stop: 7033A,
330 C Street SW
Washington, DC 20201.

CC: Steven Posnack, Principal Deputy Assistant Secretary for Technology Policy, ASTP/ONC

Re: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as part of Clinical Care

RIN 0955-AA13

Dear Assistant Secretary Keane,

This comment is submitted by GenHealth.ai, a healthcare technology company that builds and deploys AI-powered automation tools for healthcare administration. We appreciate the ongoing work of ASTP/ONC in exploring how to accelerate the practical use of AI in patient care and applaud the Department's forward-leaning approach to this RFI.

General Points

The central barrier to AI adoption in healthcare is not technological. It is contractual. Across the healthcare ecosystem, providers, patients, payers, and other stakeholders are increasingly prepared to use AI agents to act on their behalf: to submit authorizations, verify eligibility, retrieve clinical documentation, check claim status, and perform the thousands of routine administrative interactions that the healthcare system requires. The technology to do this exists today. What inhibits it is that the terms of service, acceptable use policies, and licensing agreements imposed by counterparties (the systems on the other side of these interactions) attempt to prohibit the use of AI automated means to perform what any authorized human can do manually.

This is a fundamental issue of agency and delegation. When a physician authorizes a staff member to log into a payer portal and submit a prior authorization on the physician's behalf, no one questions the legitimacy of that delegation. When a patient

authorizes a caregiver to call an insurance company and check the status of a claim, that delegation is accepted without controversy. But when those same authorized parties seek to delegate the identical tasks to an AI agent (acting with the same credentials, the same scope of access, and the same purpose) the terms of service governing the counterparty's system attempt to prohibit it. The work itself is unchanged. The authorization is the same. The only difference is the means by which the authorized task is performed.

HHS should establish a clear principle: any party in the healthcare ecosystem, whether a provider, patient, payer, clearinghouse, or other stakeholder, has the right to appoint an AI agent to act on their behalf in performing administrative tasks, and counterparties may not use terms of service, acceptable use policies, or licensing agreements to prevent that delegation. This principle does not require new technology. It does not require new interoperability standards. It requires HHS to recognize that the right to use AI as an authorized agent is an extension of the existing right to delegate administrative tasks to human agents, and that contractual terms which prohibit AI delegation while permitting human delegation are an artificial barrier to the modernization of healthcare administration.

The stakes are significant. Administrative costs consume between 15 and 34 percent of total U.S. healthcare expenditures. Shrank, Rogstad, and Parekh estimated in JAMA (2019) that administrative complexity alone accounts for approximately \$265.6 billion in annual waste, resulting as the single largest category of waste in U.S. healthcare. Himmelstein, Campbell, and Woolhandler found in Annals of Internal Medicine (2020) that \$812 billion (34.2 percent of national health expenditures) went to administration in 2017, compared to 17 percent in Canada. These costs are not abstract. They manifest as the 13 hours per week that the AMA found physicians and staff spend on prior authorization alone, as the 93 percent of physicians who report that prior authorization delays access to necessary care, and as the 29 percent who report it has led to a serious adverse event for a patient in their care.

We address the Department's specific questions below through the lens of this central principle: the right of healthcare stakeholders to delegate administrative tasks to AI agents, and the obligation of counterparties not to contractually prohibit that delegation.

1. What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

The biggest barrier is that counterparties across the healthcare ecosystem use terms of service and licensing agreements to prohibit the very AI-powered interactions that would reduce administrative burden for all parties. This manifests in three ways.

First, payer portals include terms of service that prohibit automated access to their systems. A provider organization with a fully authenticated, credentialed user and legitimate access to a payer's system is contractually barred from delegating the manual data entry, form navigation, and status checking to an AI agent acting on the provider's behalf. The provider's right to interact with the payer's system (a right that exists because the payer requires the provider to use that system) is conditioned on the provider performing the interaction manually or through a human employee. This restriction applies regardless of whether the AI agent uses the same credentials, the same access permissions, and performs the same actions as a human user would.

Second, EHR vendors impose licensing terms that restrict how provider organizations can interact with their own clinical data through automated means. A physician practice that wants to use an AI agent to extract clinical documentation from its own EHR, with documentation that belongs to the practice and its patients, to support an administrative interaction with a payer may find that the EHR vendor's terms prohibit the use of automated tools to interact with the software. This prevents providers from appointing AI agents to act on their behalf even within systems the provider has purchased and operates.

Third, the two dominant medical necessity criteria vendors impose licensing terms that prohibit AI from reasoning over their guidelines. These vendors' criteria are used by the majority of U.S. health plans, thousands of hospitals, and federal agencies including CMS, the VA, and the Department of Defense to govern medical necessity determinations. Notably, these restrictions do not only affect external parties seeking to automate interactions with health plans. The health plans themselves are often prohibited by the same licensing terms from using AI to improve their own internal operations, including their own utilization review and authorization workflows. A health plan that licenses these criteria and wants to deploy AI to accelerate its own medical necessity determinations, reduce its own administrative costs, and issue faster decisions for its own members may find that the criteria vendor's licensing terms forbid it. The restrictions thus impede modernization on both sides of the administrative transaction. We have direct experience with contractual language of the following nature: "I will not ingest, consume, translate, reason or learn over, create, train, test, improve, or otherwise use in any manner the Criteria... in connection with any machine learning algorithms, large language models, or artificial intelligence systems, including any architectures or models." A human utilization review nurse is permitted to read these criteria and apply them to a patient's clinical documentation. An AI agent performing the identical cognitive task of reading the criteria, reviewing the chart, outputting a determination is contractually prohibited from doing so. The task is the same. The authorization is the same. Only the means are different, and the licensing terms forbid it.

On the provider side, ONC's December 19, 2025 FAQ (IB.FAQ54.2025DEC) confirmed that practices interfering with AI and RPA access to electronic health information could implicate the information blocking regulations under 42 CFR 171.103. ONC stated that the information blocking definition "is not specific to any particular means, manners, or mechanisms by which access, exchange, or use of EHI is sought and could be accomplished." This is an important precedent: it establishes that the right to access health information is not limited to manual, human means. However, health plans are not explicitly defined as "actors" under 45 CFR 171.102. The result is a regulatory asymmetry where identical terms-of-service restrictions are potentially unlawful when imposed by providers or EHR vendors, but face no comparable regulatory consequence when imposed by payers.

2. What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why? What HHS regulations, policies, or programs could be revisited to augment your ability to develop or use AI in clinical care? Please provide specific changes and applicable Code of Federal Regulations citations.

HHS should establish, through rulemaking or binding guidance, that any party participating in a federally funded healthcare program has the right to delegate administrative tasks to an AI agent, and that counterparties may not impose terms of service, acceptable use policies, or licensing agreements that prohibit such delegation. This principle should apply symmetrically across all participants in the healthcare ecosystem: providers delegating to AI agents that interact with payer systems, payers delegating to AI agents that interact with provider systems, patients delegating to AI agents that interact with any system holding their health information, and all parties delegating to AI agents that reason over medical policies used in connection with coverage determinations.

To implement this principle, HHS should extend the information blocking protections currently defined in 42 CFR Part 171 to explicitly cover health plans and payer systems. The three regulated actor categories under 45 CFR 171.102 do not include payers by name. ONC has acknowledged that health plans "are not specifically identified within any of these definitions" but also "are not specifically excluded." HHS should resolve this ambiguity by amending the actor definition to include health plans, or by establishing through binding guidance that payer terms of service prohibiting authorized users from employing AI agents constitute information blocking.

HHS should also prohibit licensing terms that prevent AI from reasoning over medical necessity criteria used in federally funded programs. The two dominant clinical criteria vendors control the guidelines that govern the vast majority of medical necessity determinations in the United States, and their licensing terms effectively bar

independent AI agents from applying those criteria. The concern is compounded when a criteria vendor is vertically integrated with a health plan and a healthcare services company, and channels AI access to its criteria exclusively through its own products. When licensing terms permit AI delegation only by the entity that also controls the insurance decisions, the right of providers and patients to appoint their own AI agents to navigate those decisions is nullified. HHS should ensure that any entity licensing medical necessity criteria for use in federally funded programs must permit AI agents, acting on behalf of authorized users, to reason over those criteria in the same manner that human reviewers do.

3. For non-medical devices, we understand that use of AI in clinical care may raise novel legal and implementation issues that challenge existing governance and accountability structures (e.g., relating to liability, indemnification, privacy, and security). What novel legal and implementation issues exist and what role, if any, should HHS play to help address them?

The most significant novel legal issue is that no clear federal framework establishes the right of healthcare stakeholders to delegate administrative tasks to AI agents. In the absence of such a framework, counterparties have filled the vacuum with contractual terms that prohibit AI delegation. The result is a patchwork of terms of service across hundreds of payer portals, dozens of EHR systems, and a handful of medical criteria organizations, each independently restricting the use of AI in ways that collectively make automation of cross-party administrative interactions impractical.

HHS should establish a clear, affirmative framework recognizing that the delegation of administrative tasks to AI agents is a legitimate exercise of an authorized party's existing access rights, and that contractual terms purporting to restrict such delegation in connection with federally funded healthcare programs are contrary to federal policy. This framework does not need to address every aspect of AI governance. It need only establish that the right to perform an administrative task includes the right to delegate that task to an AI agent, and that the counterparty's terms of service cannot override that right. The absence of such a framework today is not a neutral condition, it is an active impediment. Every day that passes without clarity on AI delegation rights is a day in which counterparties continue to enforce contractual prohibitions that preserve manual processes at the expense of patients and the federal programs that serve them.

7. Which role(s), decision maker(s), or governing bodies within health care organizations have the most influence on the adoption of AI for clinical care? What are the primary administrative hurdles to the adoption of AI in clinical care?

The primary hurdles are external, not internal. A healthcare organization's decision to adopt AI for administrative automation is increasingly irrelevant if the counterparties it

must interact with contractually prohibit AI agents from performing those interactions, especially as organizations are implementing AI internally to respond to human submissions. The payer whose portal the provider must use sets the terms. The criteria organization whose guidelines the provider must apply sets the licensing conditions. The EHR vendor whose system holds the provider's clinical data sets the acceptable use policy. The organization's own leadership may be fully committed to AI adoption and still find that the contractual terms imposed by external parties prevent their AI agents from functioning. HHS can have the greatest impact by addressing these external contractual barriers, establishing the right of healthcare organizations to delegate to AI agents regardless of counterparty terms of service.

8. Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care? Please consider specific data types, data standards, and benchmarking tools.

The most impactful interoperability improvement HHS can make is not a new standard, but ensuring that AI agents can operate existing systems on behalf of authorized users. Current standards (FHIR, HL7 v2, X12 278, and the Da Vinci Implementation Guides) are limited in scope and granularity. The CMS Interoperability and Prior Authorization Final Rule (CMS-0057-F) requires payers to implement FHIR-based APIs by January 1, 2027, but this approach has significant limitations.

CMS-0057-F will not fully automate administrative workflows. Many aspects remain outside the scope of API-based approaches. The FHIR-based APIs will in many cases introduce additional burden: providers will be required to answer dozens of structured questions per submission through Documentation Templates and Rules (DTR), adding new data entry requirements that do not exist in current workflows. Not all clinical data necessary for administrative determinations is available through FHIR APIs. Medical charts, clinical notes, imaging reports, and other unstructured documentation are not fully represented in structured FHIR resources. The rule also relies on unproven technologies: Clinical Quality Language (CQL) is a rules-based query language fundamentally incapable of reasoning over medical chart data the way human clinicians do and the way large language models can. CQL evaluates structured data against predefined logic but cannot interpret free-text narratives, synthesize information across chart sections, or exercise clinical judgment.

AI agents that operate existing web-based systems, using the same authentication methods and interfaces available to human users, can bridge this gap today. These agents can navigate payer portals, extract clinical documentation from EHR systems, and complete administrative workflows across the fragmented landscape of healthcare technology without requiring new standards. Unlike traditional automation, modern AI agents powered by large language models adapt to interface changes, interpret

unstructured data, and handle the variability inherent in healthcare workflows. But this approach is only viable if the terms of service governing these systems permit authorized users to delegate interactions to AI agents. Establishing that right is the single most impactful interoperability action HHS can take.

9. What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?

Patients and caregivers want the right to appoint AI agents to navigate the healthcare system on their behalf. Today, patients must manually call insurance companies, navigate payer portals, track prior authorization status, request medical records, and manage billing disputes, or rely on already-overburdened clinical staff to do so. These are administrative tasks that patients have every right to delegate, and AI agents are now capable of performing them. But the terms of service governing the systems patients and their providers must use prohibit automated interaction, effectively requiring patients to perform these tasks manually or not at all. The AMA found that 93 percent of physicians report prior authorization delays access to care, and 82 percent say it leads patients to abandon treatment. AI agents acting on behalf of patients and their providers can address this, if the contractual terms governing counterparty systems permit them to function. Patients should not be required to accept that the systems controlling access to their care are governed by terms of service that prohibit the most efficient means of interacting with those systems.

Conclusion

The principle at the heart of this comment is straightforward. Every participant in the healthcare ecosystem, every provider, patient, payer, and caregiver, currently has the right to delegate administrative tasks to human agents who act on their behalf. A provider delegates to staff. A patient delegates to a caregiver. A payer delegates to a contractor. No counterparty questions the legitimacy of these delegations or imposes terms of service that prohibit them. The same right should extend to AI agents. When an authorized party delegates an administrative task to an AI agent acting with the same credentials, the same scope of access, and the same purpose as a human agent, the counterparty's terms of service should not be permitted to prohibit that delegation.

The administrative burden in U.S. healthcare (\$265 billion or more in annual waste) persists not because automation is technically impossible, but because the contractual environment prohibits it. Terms of service forbid automated portal interactions. Licensing agreements prohibit AI from reasoning over medical policies. Acceptable use policies restrict programmatic access to clinical data. These contractual barriers collectively prevent healthcare stakeholders from exercising their right to appoint AI agents,

preserving manual inefficiency at enormous cost to patients, providers, and the federal programs that finance their care.

HHS has the authority to change this. We urge the Department to establish that the right to delegate administrative tasks to AI agents is an inherent extension of the right to delegate those tasks to human agents, to extend information blocking protections to payer systems, to prohibit licensing terms that prevent AI from reasoning over medical policies used in federally funded programs, and to recognize AI agents operating existing systems as a legitimate and necessary complement to standards-based interoperability. These actions are consistent with the President's AI Action Plan, with HHS's own AI Strategy, and with the interests of every participant in the healthcare ecosystem who is ready to use AI but is contractually prohibited from doing so. The technology is ready. The stakeholders are willing. The only thing standing in the way is contractual language that serves no patient interest, no clinical purpose, and no public policy objective. HHS should remove that barrier.

We appreciate the Department's commitment to accelerating AI adoption in clinical care and welcome the opportunity to provide further detail on any of the issues raised in this comment.

Sincerely,

Ricky Sahu

Ricky Sahu
CEO
GenHealth.ai