

February 23, 2026

Thomas Keane, MD, MBA
Assistant Secretary for Technology Policy
Office of the National Coordinator for Health Information Technology
U.S. Department of Health and Human Services
330 C Street SW
Washington, DC 20201

RE: Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care

Dear Assistant Secretary Keane:

On behalf of the American College of Emergency Physicians (ACEP) and our nearly 40,000 members, thank you for the opportunity to respond to the “Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care” request for information (RFI). Emergency departments serve as a critical entry point to the health care system, operate in high-acuity and time-sensitive environments, and have already become a testing ground for a wide range of artificial intelligence-driven (AI) tools. Accordingly, emergency medicine is an essential lens through which to examine how AI adoption and utilization can supplement the delivery of high-quality, lifesaving emergency care.

ACEP supports responsible acceleration of AI adoption in clinical care, provided that federal policy ensures (1) meaningful clinician oversight, (2) clear accountability frameworks, (3) evidence-based evaluation standards, and (4) transparency in development and deployment.

ACEP's emergency physician members, including those working in academic medical centers, community hospitals, and safety-net settings, are encountering AI tools in their daily practice with increasing frequency, often without having meaningfully participated in the decisions to adopt those tools. Our response to this RFI is informed by input gathered from emergency medicine clinicians and health system informaticists and is intended to describe the landscape of uncertainty, concern, and cautious optimism that characterizes emergency medicine's current relationship with clinical AI. Our comments seek to articulate the key areas where clarity is most urgently needed, and to express ACEP's willingness to serve as a partner to HHS as it works to develop a coherent, equitable, and evidence-based framework for AI in clinical care.

I. Emergency Medicine as a Distinct Clinical Environment

We want to underscore for HHS that emergency medicine presents a set of conditions that make it both a high-priority area for AI development and a setting in which the consequences of poorly designed or poorly deployed AI are particularly acute. Emergency departments manage an unselected patient population, operate continuously, and are frequently under conditions of crowding, resource constraint, and time pressure. These factors affect both the types of AI tools that are most likely to be beneficial and the ways in which AI failures are likely to manifest.

Our members have observed that AI tools which show genuine promise in the emergency department, including documentation assistance, radiology support for critical finding detection, and tools to support patient triage, tend to be those that are designed with a clear understanding of how emergency care is actually delivered.

WASHINGTON, DC OFFICE

901 New York Ave, NW
Suite 515E
Washington DC 20001-4432

202-728-0610
800-320-0610
www.acep.org

BOARD OF DIRECTORS

L. Anthony Cirillo, MD, FACEP
President
Ryan A. Stanton, MD, FACEP
President-Elect
Chadd K. Kraus, DO, DrPH, CPE, FACEP
Chair of the Board
Henry Z. Pitzele, MD, FACEP
Vice President – Communications
Heidi C. Knowles, MS, MD, FACEP
Vice President – Membership
Abhi Mehrotra, MD, MBA, FACEP
Secretary-Treasurer
Alison J. Haddock, MD, FACEP
Immediate Past President
Jennifer Casaletto, MD, FACEP
Dan Freess, MD, FACEP
Steven B. Kailas, MD, MPH, FACEP
C. Ryan Keay, MD, FACEP
Kristin B. McCabe-Kline, MD, FACEP
Diana B. Nordlund, DO, JD, FACEP
Bing S. Pao, MD, FACEP
James L. Shoemaker, Jr., MD, FACEP

COUNCIL OFFICERS

Michael J. McCrea, MD, FACEP
Speaker
Larisa M. Traill, MD, FACEP
Vice Speaker

EXECUTIVE DIRECTOR

Michael Fraser, PhD, MS, CAE

Conversely, tools designed for elective or inpatient settings, or that generate frequent low-value alerts, tend to create more burden than benefit. We are also attentive to a concern raised by members regarding the appropriate role of AI in the emergency department: there is a risk that AI tools will be deployed to manage or obscure systemic problems, such as crowding, boarding, and staffing shortages, rather than to improve the quality of care delivered within those constraints. AI adoption in emergency medicine should be problem-driven and tied to evidence of genuine improvement in patient outcomes or clinician experience, not used as a technological substitute for addressing structural challenges in emergency care delivery.

II. Liability and Accountability: A Critical Area of Uncertainty

Our members have described a troubling bind: following an AI recommendation that proves incorrect may give rise to claims of automation bias or over dependence on technology, while ignoring a recommendation that turns out to have been correct may expose the clinician to allegations of negligence for failing to use available clinical decision support. This tension places the entire burden of risk on the clinician, despite the fact that emergency physicians typically have no role in selecting, procuring, validating, or configuring the AI tools they are expected to use.

Several specific areas of uncertainty warrant attention:

- There is currently no established standard of care that defines when a clinician is expected to follow, document disagreement with, or override an AI-generated recommendation.
- The question of who bears liability when AI is involved including developers, health systems, or individual clinicians has not been resolved, and existing malpractice frameworks were not designed to accommodate this complexity.
- When health systems deploy AI tools that have not been validated against local patient populations, it is unclear whether and to what degree the institution bears responsibility for downstream patient harm.
- The liability question extends well beyond FDA-regulated medical devices to encompass ambient listening tools, natural language processing applications, clinical decision support systems, and other non-device AI categories where governance is especially underdeveloped.

ACEP does not prescribe a particular liability model in this response, but we believe that addressing this uncertainty is prerequisite to responsible AI adoption. We are willing to engage with HHS/ONC and relevant stakeholders to examine what a workable framework might look like, and to share the perspectives of our frontline members who are navigating this uncertainty daily.

III. Governance and Physician Involvement in AI Decision-Making

A recurring and frustrating theme in feedback from our members is that emergency physicians are rarely meaningfully involved in decisions about which AI tools are adopted in their departments. It is ACEP's position that artificial intelligence algorithms used in health care decision-making must be transparent, accountable, and subject to oversight by qualified health care professionals. Yet technology procurement decisions are most commonly driven by institutional IT departments, finance leadership, and hospital administrators, with clinical input typically added late in the process and given limited weight.

This is not simply a matter of professional preference. Emergency medicine's care environment is distinctive: it is characterized by high patient volume, diagnostic uncertainty, time pressure, and a patient population that skews toward complexity and acuity. AI tools that perform well in other settings frequently underperform or create new problems when deployed in the emergency department without ED-specific validation. When the clinicians who will actually use a tool are not involved in evaluating it, the likelihood of poor fit, and the downstream risks to patient safety, increases substantially.

Our members have raised concerns that mirror those expressed about electronic health record adoption: systems are selected and implemented with inadequate clinical input, and frontline physicians are left to adapt to tools that do not reflect how emergency medicine is actually practiced. We believe HHS has an opportunity to address this structural gap through guidance on post deployment monitoring processes.

IV. Reimbursement and Economic Alignment

Emergency physicians and health system leaders have described a structural economic misalignment in how AI adoption is currently financed. The costs and effort associated with implementing, validating, training staff on, and maintaining AI tools are frequently absorbed by clinical departments, while the financial benefits such as reduced downstream costs, billing efficiencies, system-level savings accrue to health systems, payers, or vendors. This imbalance creates a disincentive for departments to invest in thoughtful, resource-intensive adoption and contributes to burnout risk when AI tools become additional unfunded mandates on already-stretched clinical teams.

Current reimbursement models were not designed with AI in mind, and do not adequately account for the full lifecycle of responsible AI use: pre-deployment validation, staff training, integration support, governance overhead, and ongoing performance monitoring. Our members have observed that where resources for these activities are not explicitly allocated, they either do not occur or are performed inadequately thus creating patient safety risks that may not be apparent until problems emerge in practice.

AI systems used in utilization management, prior authorization, and reimbursement determinations also materially affect clinical decision-making and patient access to emergency care. ACEP policy supports prohibiting the use of AI as the sole basis for denial, delay, or downcoding of care and, most importantly, supports requiring review by qualified physicians in the relevant specialty.

V. Evaluation Standards and the Gap Between Vendor Claims and Real-World Performance

Our members have consistently described a gap between the performance metrics presented by AI vendors and the real-world behavior of those tools once deployed in an emergency department. AI tools are typically developed and evaluated on datasets that may not reflect the demographic makeup, data quality, workflow constraints, or clinical complexity of the populations and settings in which they are eventually used. Emergency department data, in particular, is often fragmented, entered under time pressure, and structured differently than the cleaner data used in development environments.

This gap has real consequences. Tools that generate high rates of false-positive alerts, as has been reported with some sepsis detection systems, create alert fatigue that extends to other clinical decision support tools, potentially undermining trust in AI systems more broadly. Tools developed for use in other specialties and repurposed for the emergency department often fail to perform as expected, and the health systems and clinicians bearing the consequences of those failures have limited recourse.

There is currently no standard framework for how AI tools used in clinical care, particularly those that fall outside FDA medical device regulation, should be evaluated before deployment or monitored after. Our members have expressed interest in a more structured approach, though there are differing views about whether emergency medicine requires a specialty-specific evaluation framework or whether existing standards for medical interventions can be adapted. These are questions on which further dialogue would be valuable.

Areas where HHS guidance or investment could be particularly impactful include:

- Development of pre-deployment validation standards or guidelines for non-device clinical AI, with particular attention to the importance of local validation against institution-specific patient populations.
- Support for ongoing post-deployment monitoring infrastructure, including mechanisms to identify and respond to model drift over time.
- Investment in research capacity at community and mid-size health systems that currently lack the resources to conduct independent AI validation.
- Promotion of transparency requirements that would allow health systems and clinicians to understand how AI tools were developed, on what data, and what their known limitations are.

VI. ACEP's Commitment to Partnership

ACEP is committed to ensuring that AI is developed, deployed, and governed in ways that genuinely serve patients and support the clinicians who care for them. We recognize that this is a rapidly evolving policy area and that definitive answers to many of the questions raised in this letter are not yet available. Our goal in this response is not

to foreclose options but to ensure that the perspectives of emergency medicine are part of the conversation as HHS develops its approach.

We are prepared to:

- Participate in technical working groups, advisory panels, or stakeholder convenings convened by HHS on AI in clinical care.
- Provide access to emergency physician members with direct experience developing, evaluating, and using AI tools in clinical settings.
- Collaborate on the development of evaluation frameworks, governance standards, or other guidance documents where emergency medicine input would add value.
- Share ongoing data and feedback from our membership as AI adoption in emergency medicine continues to evolve.

We look forward to an ongoing dialogue and thank the Department for seeking input from the clinical community. If you have any questions, please contact Annalane Miller, ACEP's Manager of Regulatory and External Affairs, at amiller1@acep.org.

Sincerely,

A handwritten signature in black ink that reads "L. Anthony Cirillo, MD, FACEP". The signature is fluid and cursive, with the last name "Cirillo" being the most prominent part.

L. Anthony Cirillo, MD, FACEP
President
American College of Emergency Physicians