



Response to HHS Request for Information: Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care

RIN 0955-AA13

Submitted by: HealthEx 2339 3rd St., Suite 41 San Francisco, CA 94107 info@healthex.io | www.healthex.io

Date: February 23, 2026

1. EXECUTIVE SUMMARY

HealthEx is building the modern stack for consumer-driven care - combining digital identity, real-time records access, and consent management - and is already live powering AI grounded in individuals' personal health history through its partnership with Anthropic, announced in January 2026.

HealthEx appreciates the opportunity to respond to HHS's Request for Information on accelerating the adoption and use of AI as part of clinical care. Our response is grounded in a core conviction: **the single greatest barrier to effective clinical AI is not the model - it is the data**. AI tools can only deliver personalized, clinically meaningful insights when they are grounded in a patient's actual, comprehensive health history. Today, that data remains fragmented, institutionally siloed, and inaccessible to the AI tools that patients and clinicians increasingly rely on.

HealthEx solves this problem and we are the only platform live with both TEFCA Individual Access Services (IAS) and a FHIR endpoint directory for patient-directed records retrieval, enabling individuals to securely access their complete health records across all their places of care and share them with AI tools, providers, payers, and trusted third parties of their choosing.

HealthEx responds to Questions 1, 2, 5, 6, 8, and 9 below. Our recommendations center on three themes: (1) interoperability and patient-directed data access are the foundational infrastructure for clinical AI; (2) HHS should leverage existing frameworks - particularly TEFCA IAS - to accelerate AI adoption rather than creating parallel regimes; and (3) patient empowerment and trust must be embedded in AI deployment, not bolted on after the fact. We also note the uniquely new requirements emerging from conversational AI contexts around more efficient and real-time mechanisms to keep a patient's clinical record updated, and the need for high-quality imaging data to be available to patients to access via IAL2/IAS.

RFI prepared by:

- Priyanka Agarwal, MD, MBA, Co-founder and CEO, HealthEx
- Anand Raghavan, Co-founder and CTO, HealthEx
- Jeremy Schwarz, Chief Commercial Officer, HealthEx

2. ABOUT HEALTHEX

HealthEx is building the modern stack for consumer-driven healthcare. Combining digital identity, unified longitudinal records, and consent and permissions, HealthEx presents to end users as a health data wallet that enables them to securely access and share their health records with trusted third parties across payers, digital apps, and AI agents in a repeatable and efficient manner.

Below are key differentiators of our platform and approach:

- **Dual approach to records retrieval across best-in-class TEFCAs and FHIR endpoint based flows:** HealthEx is the only team enabling comprehensive records access via IAS under TEFCAs and a FHIR endpoint directory. Our record locator service automatically matches individuals' demographic information - including historical addresses - to find where they have health records across thousands of places of care nationwide, eliminating the need for patients to recall every provider they've seen.

HealthEx connects to the TEFCAs network-of-networks infrastructure via our QHIN partners, MedAllies and CommonWell, spanning 12,000+ organizations and 72,000+ unique connections. HealthEx additionally supports direct connections to over 52,000 healthcare organizations via FHIR-based exchange.

Through these efforts, HealthEx supports records retrieval for 250M Americans, supporting records access for ~90% of provider organizations.

- **The HealthEx Digital Wallet for records powers Identity-powered network effects:** The HealthEx Digital Wallet enables users to access their records once and share across all platforms where they've chosen to connect, creating powerful network effects. HealthEx's identity foundation (IAL2/AAL2 via CLEAR) powers the digital wallet capabilities- a Claude user can seamlessly share the same data to other AI tools or providers. This identity layer is also foundational for agentic care experiences like EHR write-back.
- **AI-optimized data delivery:** The HealthEx MCP server is built for performance, speed, and accuracy while giving users privacy and control. It addresses the unique challenges at the intersection of health data and AI: context-window constraints (through proprietary data cleaning and compression achieving ~90% improvement), and privacy (honoring consent on every tool call, returning only query-relevant data, and operating exclusively in contexts where model training is off). AI-ready data can also be delivered via REST APIs.
- **Production-proven at scale:** Live integration with Anthropic's Claude (January 2026). Additional launches in Q1 with two top-5 health insurers, and two additional large-scale AI teams. Across its ecosystem, HealthEx will be enabled for 10M+ Americans in Q1.
- **Deep ecosystem partnerships:** HealthEx is collaborating with every major EHR provider across Epic, athenahealth, Cerner, Veradigm (AllScripts), eClinicalWorks,

MEDITECH, NextGen, and NetSmart. These are not proprietary partnerships, rather they represent a willingness from HealthEx to work together across the ecosystem.

- HealthEx collaborates closely with every major EHR provider including Epic, athenahealth, Cerner, Veradigm (AllScripts), eClinicalWorks, MEDITECH, NextGen, Greenway, and NetSmart. Here are a few recent, relevant public announcements and efforts:

[Epic](#)

[Athena/Medallies](#)

[MEDITECH](#)

[CMS Interoperability Initiatives](#)

Powering AI for Claude & Americans at scale:

In January 2026, HealthEx launched a live integration with Anthropic's Claude through the Model Context Protocol (MCP), enabling consumers to connect their health records and use conversational AI for personalized health insights grounded in their clinical history. This was the first patient-data MCP integration with a major AI platform. HealthEx has forthcoming partner announcements in Q1 2026 that will extend its consumer network to over 10 million US users.

HealthEx is a San Francisco-based company backed by General Catalyst. HealthEx is SOC2 and TX-RAMP certified, and is an approved IAS vendor in TEFCA.

3. RESPONSE TO SPECIFIC QUESTIONS

Question 1: What are the biggest barriers to private sector innovation in AI for health care and its adoption and use in clinical care?

The biggest barrier to clinical AI innovation is the fragmentation and inaccessibility of patient health data to provide context to AI.

AI models are only as useful as the data they can access. Today, a patient's health records are scattered across dozens of providers, health systems, and payer platforms, each with its own portal, credentials, and data formats. This fragmentation means that even the most advanced AI tools operate on incomplete, single-source data – producing generic insights rather than the personalized, longitudinal reasoning that transforms clinical care.

The difference is significant. An AI tool with access to a single EHR system can tell a patient what a lab result means in isolation. An AI tool grounded in a patient's full longitudinal record – spanning conditions, medications, labs, immunizations, clinical notes, and procedures across every provider they've seen – can identify trends over time, surface patterns across specialists, and help patients prepare for more informed conversations with their doctors. This is the kind of complex, contextual reasoning that meaningfully improves care.

HealthEx has demonstrated that this barrier is solvable with existing infrastructure. Through TEFCA IAS and our FHIR endpoint directory, we enable patients to pull their complete records in under five minutes at the median and connect them directly to AI tools. Since launching with Anthropic's Claude in January 2026, approximately half of HealthEx connector users return within five days to have another conversation using their health data, a strong signal that patients find real value in AI grounded in their clinical history.

Beyond data fragmentation, we see three additional barriers:

1. **The evaluation paradigm for clinical AI has not kept pace with the speed of model development.** Models are updated on cycles measured in weeks, but evaluation frameworks still reflect the cadence of traditional medical device review. HHS should support real-time, outcomes-based evaluation methods for a set of templated use cases that can assess AI performance as models evolve - rather than evaluating the efficacy of AI on models that are already out of date by the time the assessment is complete.
2. **The consent and trust infrastructure for AI-health-data interactions remains underdeveloped.** Standard MCP server frameworks do not come with native privacy or consent semantics. HealthEx has built this layer – honoring consent on every tool call, returning only data relevant to a specific query, ensuring users control whether data is saved across sessions, and operating only in environments where model training on PHI is off. HHS should encourage the development and adoption of these trust-enabling

practices across the ecosystem, and HealthEx would be pleased to lend its expertise and experience in the consent arena.

3. **Regulatory uncertainty around non-medical-device AI tools creates hesitancy among both developers and health systems.** Organizations are unsure how AI tools that synthesize patient records, generate care preparation summaries, or support shared decision-making fit within existing regulatory and liability frameworks. HHS can reduce this uncertainty by clarifying that existing frameworks already accommodate these use cases and establishing proportionate guidance - removing ambiguity rather than adding new requirements or regulations.

Question 2: What regulatory, payment policy, or programmatic design changes should HHS prioritize to incentivize the effective use of AI in clinical care and why?

HHS should prioritize policy changes that ensure patient health data can flow to AI tools of a patient's choosing, securely and at scale. The regulatory and payment infrastructure that governs health data exchange will determine whether clinical AI reaches its potential or remains constrained by the same data silos that have limited interoperability for decades.

We recommend the following specific policy actions:

Strengthen TEFCA IAS (or broader CMS-Aligned network IAS) adoption via Incentivation from CMS. TEFCA's Individual Access Services pathway is the most significant federal infrastructure enabling patient-directed data exchange at national scale, and we believe is the most secure and user-friendly mode for patients to access their data. IAS enables patients to retrieve their records from across the TEFCA network and direct them to any trusted third party, including AI tools. HHS should incentivize CMS-regulated entities to participate in TEFCA, and specifically that all TEFCA Participants and Subparticipants respond to IAS requests. While TEFCA adoption has been strong, it is still voluntary and there are some organizations who have not yet pledged to join.

Today, some organizations decline IAS queries due to lingering uncertainty about whether TEFCA's identity and authorization verification meets HIPAA requirements. HHS should issue guidance confirming that IAL2/AAL2 identity verification plus IAS consent satisfies HIPAA identity and authorization requirements under 45 CFR 164.524, removing some of the basis for non-response.

Grant safe harbor to organizations performing patient matching and proposed "double checking" as part of IAS fulfillment. Hesitancy to respond to IAS queries is partly driven by liability concerns around patient matching. Granting qualified immunity to organizations that perform patient matching in response to IAL2-enabled IAS requests would reduce this barrier and unlock broader participation.

Reject additional authorization burdens that impede patient access. Some EHR vendors currently require patients to provide patient portal credentials on top of IAL2 identity verification when fulfilling IAS requests. This added step is not universally required - users of athenahealth and eClinical works EHRs, for example, do not face it - and creates a friction point that disproportionately burdens patients with limited digital literacy, multiple providers, or complex care needs. HHS should clarify if IAS requests using IAL2/AAL2 tokens plus digitally enforceable consent are sufficient for authorization as well, without additional portal credential requirements. We are advocates to converge on the most friction free flows for end users possible, while respecting the applicable legal requirements.

Enable persistent data connections with extended token validity and refresh mechanisms. Nearly all patients expect their health data connections to remain fresh and current without manual intervention after each provider visit or lab result. Under the current IAS framework, OIDC tokens from credential service providers expire within 30 days, requiring patients to re-authenticate biometrically for subsequent data access. This is a significant step backward from the refresh token model available under patient portal-based SMART-on-FHIR workflows. HHS should update guidance to clarify that IAL2 tokens may persist for an extended period – up to 12 months – with appropriate refresh mechanisms, so patients can direct recurring data access to AI tools and other trusted parties without repeated identity verification.

Support push-based data updates to maintain record freshness. Manual refresh workflows place an unreasonable burden on patients and create gaps in the data available to AI tools. HHS should establish standards and incentives for push-based notifications when new clinical data becomes available. This could include webhook-based updates triggered when new encounters, lab results, or imaging studies are documented, enabling AI tools to maintain current, comprehensive patient context without requiring patients to manually pull updates after every healthcare interaction. Push-based updates aligned with patient consent would dramatically improve the clinical utility of AI while reducing patient friction.

Modernize payment policy to accelerate AI adoption. The RFI rightly identifies the limitations of legacy fee-for-service regimes in supporting innovation. HHS should explore multiple payment mechanisms to ensure AI tools that demonstrably reduce burden, lower costs, or improve outcomes can reach patients:

- Medicare Advantage supplemental benefits: CMS should clarify that AI-enabled care navigation, health literacy, and care coordination tools qualify as supplemental benefits under MA plans, enabling plans to offer these capabilities to beneficiaries without additional regulatory approval.
- CMMI innovation models: CMS should establish innovation models that test AI-enabled care pathways - including AI-assisted prior authorization, care gap identification, and care navigation - under value-based payment arrangements that reward outcomes rather than volume.
- Safe harbors for patient-directed AI tools: HHS should consider safe harbor protections under anti-kickback statutes for AI tools that are patient-directed and operate on

patient-authorized data, removing the legal uncertainty that currently deters health systems from recommending AI tools to patients.

Accelerate public availability of a national healthcare endpoint directory that maps OIDs to Patient portals. Effective AI-driven care requires that patients (or platforms acting on their behalf) can discover where their records are held. We were pleased to see CMS's recent commitments toward building a national healthcare endpoint directory, and urge HHS to expedite its development. Consistent with the ONC information blocking regulations at 45 CFR Part 171, providers should be expected to publish and maintain FHIR API endpoints in this directory, ensuring accurate record location at scale.

Question 5: How can HHS best support private sector activities to promote innovative and effective AI use in clinical care?

HHS can best support private sector AI innovation by strengthening the interoperability infrastructure that AI tools depend on and by endorsing trust frameworks for AI-health-data interactions.

The current generation of patient-facing AI tools has focused primarily on information and insights-helping patients understand their health data through natural language conversations. However, significant consumer demand exists for AI that can move beyond explanation to action: booking appointments, messaging providers, requesting prescription refills, and coordinating care. To realize this vision, the healthcare ecosystem needs standardized protocols that enable AI platforms to both access comprehensive health data and interact with clinical systems to fulfill real-world care tasks-all with appropriate patient consent and security controls.

The HealthEx integration with Anthropic's Claude demonstrates how private sector collaboration, built on public infrastructure, can deliver immediate patient value while creating a foundation for more advanced capabilities. HealthEx leverages TEFCA-a federal framework-to retrieve patient records, then delivers them to Claude via the Model Context Protocol with consent honored on every interaction. Today, consumers can ask Claude to explain their lab results, identify patterns across their health history, and prepare questions for their doctors, all grounded in their actual clinical data. HealthEx is actively working to extend this integration beyond information and insights to agentic care fulfillment, creating a comprehensive healthcare MCP that supports the full spectrum from understanding health to acting on it.

HHS should create pathways for patient-directed AI tools to serve Medicare beneficiaries. Approaches already demonstrated in the commercial market provide a model for what these pathways could enable. For example, the HealthEx-Claude integration has shown that consumers will engage deeply with AI grounded in their comprehensive health records. A similar approach for Medicare beneficiaries could deliver three transformative capabilities:

Personalized AI for health understanding: Through the HealthEx x Claude integration or similar AI-powered experiences, Medicare beneficiaries could access conversational AI grounded in their complete longitudinal health record. Members could ask natural language questions about their medications, understand what their latest lab results mean in the context of their conditions, identify patterns in their vitals over time, and prepare informed questions before appointments—all in plain language, with explanations tailored to their actual clinical history rather than generic health information.

Personalized, proactive care enabled by a 360-degree health picture: With a comprehensive view of each beneficiary's health-spanning conditions, medications, labs, procedures, and care gaps across all their providers—AI can surface personalized, proactive recommendations. The system could identify when a diabetic patient is overdue for an eye exam, flag potential medication interactions when new prescriptions are added, or notice declining kidney function trends that warrant earlier intervention. This comprehensive view, made possible by HealthEx's aggregation across the TEFCA network and FHIR endpoints, transforms reactive care into proactive health management.

Agentic care fulfillment for real-world action: Beyond insights, AI can help Medicare beneficiaries act on recommendations—scheduling that overdue eye exam with an in-network ophthalmologist, booking their annual wellness visit, coordinating transportation to appointments, or requesting prescription refills. This vision of AI-assisted care navigation, grounded in comprehensive patient data and executed with patient consent, could meaningfully reduce barriers to care access for Medicare beneficiaries while alleviating administrative burden on providers and plans.

These three capabilities—personalized AI understanding, proactive care recommendations, and agentic fulfillment—represent a fundamentally different member experience: one where Medicare beneficiaries have a trusted AI companion that knows their complete health story, helps them understand it, identifies what care they need, and helps them get it done. HHS should create pathways for patient-directed AI tools to serve Medicare beneficiaries, modeled on approaches already demonstrated in the commercial market.

HHS can promote more such innovation by supporting industry-driven trust frameworks for AI-health-data connectors. These frameworks should address identity verification, consent management, data provenance, and data minimization practices. Establishing clear, adoptable standards for how AI platforms should handle health data will accelerate responsible innovation across the private sector.

HHS should also support pilot programs and cooperative research agreements that test patient-authorized AI use cases in real clinical workflows. Pilot structures could be used to test AI-enabled form filling (reducing administrative burden during e-check-in), AI-assisted prior authorization (using patient-directed clinical data to accelerate approvals), and AI-driven care coordination (surfacing missing screenings and follow-ups from longitudinal records). These are

use cases HealthEx is actively deploying with partners including major national payers and health systems.

HHS should establish mechanisms for real-time evaluation of model performance and patient experience with clinical AI. As AI tools move from advisory roles to agentic care fulfillment, continuous monitoring becomes essential. This includes tracking accuracy of clinical recommendations, patient satisfaction and trust, identification of bias or disparities in AI-generated recommendations, and measurement of clinical outcomes and cost impacts. Public-private partnerships could establish shared evaluation frameworks and reporting standards that enable rapid identification of safety issues while preserving innovation velocity. CMS could leverage Medicare claims data combined with patient-reported outcomes to create feedback loops that improve AI performance over time, similar to post-market surveillance for medical devices but adapted for the unique characteristics of AI systems that learn and evolve.

Question 6: Where have AI tools deployed in clinical care met or exceeded performance and cost expectations and where have they fallen short?

Patient-facing AI grounded in comprehensive health records has exceeded engagement expectations, demonstrating strong demand for personalized health AI.

The HealthEx integration with Anthropic's Claude, launched in January 2026, provides early but compelling evidence that patient-facing AI tools can exceed engagement benchmarks when grounded in comprehensive, longitudinal health data. Key findings from the initial deployment:

- **87%** of users who go through our TEFCA flows have at least 1 site of care automatically identified, with a median retrieval time of under five minutes. The average user retrieves **3.3 records** through our TEFCA + FHIR endpoint flows, compared to **1.4 records** when going through our FHIR endpoint only flow, which requires that the patient provide their portal logins in order to gain access. For the average user, TEFCA IAS retrieval is approximately 43% faster than FHIR endpoint-only retrieval, reflecting the elimination of manual provider search and portal login steps.
- 83% of users with retrieved records initiate a health conversation, with health summaries and lab result interpretation as the most common entry points. This significantly exceeds typical engagement rates for patient-facing health tools.
- Approximately 50% of users return within five days for subsequent conversations, indicating sustained utility rather than one-time curiosity.
- Users engage across the full breadth of their clinical data-conditions, medications, labs, immunizations, procedures, and clinical notes-suggesting that comprehensive longitudinal data, not single-source snapshots, is what drives meaningful AI interactions.

Where AI tools have fallen short, data limitations are almost always the root cause. AI tools operating on data from a single EHR system or institution produce generic insights that patients do not find actionable. The gap between single-source and comprehensive-data AI is not incremental-it is the difference between an AI that can explain a lab result and one that can identify a trend across three years of visits to multiple specialists. Closing this gap requires the interoperability and patient-directed data access reforms outlined in our response to Question 2.

AI tools today also do not have access to imaging data as these sit on PACS systems that are not governed by the 21st Century Cures act in terms of access by patients. This limits their ability to have a 360 degree view of the patient's condition.

Lastly, the current background poll based mechanisms without incremental updates of keeping a user's data up to date result in highly inefficient volumes of queries which return with no new information about a patient. This drives up implementation costs as the traditional business models in this space rely on the number of successful record pulls which are not always incentivized to keep the number of queries to a minimum.

Question 8: Where would enhanced interoperability widen market opportunities, fuel research, and accelerate the development of AI for clinical care?

Enhanced interoperability – specifically, patient-directed interoperability via TEFCA IAS – is the single most important accelerant for clinical AI.

Today, clinical AI tools overwhelmingly operate on data from a single EHR system or a single institution. This constrains the quality, completeness, and generalizability of AI insights and limits the market to use cases that can function within institutional silos. Patient-directed data access, by contrast, unlocks longitudinal, multi-source data that enables fundamentally different AI capabilities: cross-provider trend analysis, comprehensive medication reconciliation, complete care gap identification, and personalized risk assessment grounded in a patient's full clinical history.

Specific areas where enhanced interoperability would widen market opportunities:

TEFCA IAS enables AI tools to access comprehensive patient records without requiring bilateral data-sharing agreements with every provider organization. This dramatically lowers the barrier to entry for AI developers and creates a level playing field where innovation is rewarded based on the quality of insights delivered, not the depth of institutional relationships. HealthEx's integration with Claude demonstrates this: any Claude Pro or Max subscriber can connect their full health record and receive personalized insights, regardless of which health systems they visit.

Production data from the HealthEx-Claude integration illustrates the breadth of interoperability required. The average user has records across multiple provider organizations, requiring connections to health systems, ambulatory practices, and labs. HealthEx retrieves data via both TEFCA IAS and direct FHIR connections, with the dual approach proving essential - TEFCA provides broad network reach while FHIR endpoints enable connections to organizations not yet participating in TEFCA. Neither pathway alone provides sufficient coverage; the combination is what achieves the ~90% provider coverage that makes comprehensive AI possible.

Enhanced interoperability also accelerates research by making patient-authorized clinical data available for outcomes evaluation and implementation science. When patients can direct their

longitudinal records to AI tools and research platforms, it becomes possible to evaluate AI performance against real-world outcomes at a scale and speed that institutional data alone cannot support.

Specific data types, standards, and tools where progress is needed:

USCDI v3 provides a strong baseline, but **imaging** data (particularly from PACS servers) remains difficult to access in lossless formats through standard exchange pathways. Making diagnostic imaging available through TEFCA or FHIR-based exchange would significantly expand AI capabilities in radiology, oncology, and other imaging-intensive specialties. Additionally, establishing a common quality standard for data shared across sources - potentially using novel approaches like open-source LLM-based quality scoring with public weights for transparency - would help AI developers assess data reliability.

FHIR R4 and C-CDA are working well as exchange standards, but the practical challenge is that not all organizations participate in electronic records exchange. As of 2024, roughly a quarter of office-based physicians report never sending electronic records to outside providers. TEFCA's network model can reduce connectivity costs, but HHS policy encouraging broader participation - particularly among small, rural, and independent practices - is essential to closing the remaining gaps that limit AI's access to comprehensive patient data.

Question 9: What challenges within health care do patients and caregivers wish to see addressed by the adoption and use of AI in clinical care? Equally, what concerns do patients and caregivers have related to the adoption and use of AI in clinical care?

Patients want three things from AI in their care: the ability to understand their own health records, access to comprehensive data across all their providers, and control over who sees what.

As a physician, I have seen firsthand how patients struggle to make sense of fragmented health information scattered across multiple portals. The promise of AI in clinical care is that it can serve as a bridge - translating clinical data into plain-language insights, identifying what matters most, and helping patients prepare for conversations with their doctors. Early results from the HealthEx-Claude integration confirm this demand: 83% of users with records start a health conversation, with health summaries and lab results as the most common entry points, and approximately half of users return within five days.

Patients also want AI to help them navigate the complexity of a fragmented system. This includes understanding medication interactions across providers who don't talk to each other, identifying gaps in preventive care, and preparing for specialist visits with a complete picture of their history. These are use cases that only work when AI is grounded in comprehensive, longitudinal data - not data from a single institution.

Patients and caregivers also want AI to move beyond insights and into real-world care action.

AI tools today excel at synthesizing information, but they do not yet extend into care fulfillment—booking appointments, writing back to EHRs, initiating prior authorizations, or coordinating referrals. Consumer demand for these capabilities is strong. HealthEx's identity, consent, and records access capabilities are foundational for enabling these agentic care experiences, and we are actively co-designing EHR write-back capabilities with Epic. HHS should consider how its regulatory and payment frameworks can support the transition from AI as an information tool to AI as a care coordination enabler—clarifying, for example, that patient-directed AI actions taken with appropriate consent and identity verification are consistent with existing HIPAA and HITECH frameworks.

On the concerns side, patients care deeply about privacy, consent, and control.

The sensitivity of health data demands a higher standard than what generic AI platforms provide. HealthEx has built its MCP server to address these concerns directly: consent is honored on every tool call; only data relevant to a specific query is returned to the model; users are given a choice about whether their data is saved across sessions; and the integration is only available in contexts where model training is off. These are not optional features - they are essential to building the trust that enables patients to engage with AI-powered health tools.

HHS can support patient and caregiver trust by encouraging transparency standards for AI-health-data integrations - specifically, requiring that patients be informed about how their data is accessed, what it is used for, and how to revoke access. HealthEx's platform provides this transparency through auditable consent, granular data sharing controls, and the ability to disconnect at any time.

Caregivers present a particularly important use case that current infrastructure does not adequately support.

Family members coordinating care for aging parents, children, or loved ones with complex conditions need consolidated access to health records and the ability to interact with AI tools on behalf of those they care for. HHS should support the development of caregiver proxy access pathways within IAS, so that legally authorized representatives can retrieve and share records - and use AI tools - with appropriate identity verification and consent. Specifically, HHS should work with TEFCA participants and credential service providers to establish standardized proxy identity flows that allow caregivers to connect records on behalf of dependents without requiring the care recipient to perform their own biometric verification.

Finally, patients and caregivers want AI to move beyond insights and into action. AI tools today excel at synthesizing information, but they do not yet extend into real-world care actions like appointment booking, EHR write-back, or prior authorization initiation. HealthEx's identity, consent, and records access capabilities are foundational for enabling these agentic care experiences, and we are actively co-designing EHR write-back capabilities with Epic. HHS

should consider how its regulatory and payment frameworks can support the transition from AI as an information tool to AI as a care coordination enabler.

4. CONCLUSION

HealthEx appreciates the opportunity to contribute to this important discussion. The potential for AI to transform clinical care is immense, but it will only be realized if patients can bring their full health history to the AI tools they choose - securely, with consent, and with control.

The infrastructure exists today. TEFCA IAS provides the regulatory and technical framework for patient-directed data exchange at national scale. HealthEx has demonstrated that this framework works in practice: retrieving records for 94% of users, in under five minutes at the median, and connecting them to AI tools that deliver personalized, clinically grounded insights. What is needed now is policy action that accelerates adoption, removes friction, and ensures that patient empowerment and trust are built into the foundation of AI in healthcare.

We welcome the opportunity to engage further with HHS, ASTP/ONC, and other stakeholders to support these goals.

Contact: Priyanka Agarwal, MD, MBA Co-founder and CEO, HealthEx Priyanka@healthex.io
2339 3rd St., Suite 41, San Francisco, CA 94107 www.healthex.io