

To: Department of Health and Human Services (HHS) Office of the National Coordinator for Health Information Technology (ONC)

Re: *Request for Information (RFI): Accelerating the Adoption and Use of Artificial Intelligence as Part of Clinical Care (Federal Register Doc. 2025-23641)*

From: Tayab Waseem MD PhD, Jefferson Lin JD, Dave Schoolcraft JD, Stephanie Simmons JD, on behalf of the Digital Health Counsel¹

Subject: The "Parallel Cognitive Record" and the Necessity of Privilege for Clinical AI Queries (Response to Questions 3 & 9)

To the Coordinator:

The Digital Health Counsel (DHC) is a multidisciplinary collaboration among clinicians, technologists, legal experts, and policy professionals focused on the responsible development, deployment, and governance of digital health technologies. Our work centers on identifying real-world barriers to adoption, aligning innovation with patient safety, and ensuring that regulatory frameworks reflect how technology is actually used in clinical practice.

We write to urge the Department to address a critical, emerging liability vector that creates a massive, structural barrier to the safe adoption of AI in medicine: the creation of a **"Parallel Cognitive Record."**

While the RFI correctly identifies "liability" as a potential barrier for non-device AI (Question 3), current policy discussions often miss the granular reality of how this liability manifests in the daily workflow of a clinician. The risk is not merely that an AI tool might give a wrong answer; the risk is that the *physician's interaction with the tool* creates a discoverable, permanent record of their internal thought process. This includes their doubts, their differentials, and their double-checks, which can be weaponized in litigation.

Defining the Problem: The Parallel Cognitive Record (Response to Question 3)

Historically, the physician's progress note served as the definitive legal account of a clinical encounter. It is a "Performative Record", a sanitized, professional conclusion designed for communication and billing.

However, the rapid adoption of Large Language Models (LLMs) and decision-support tools (e.g., OpenEvidence, ChatGPT, UpToDate AI, DoxGPT) has created a "Cognitive Record" which includes the raw, timestamped logs of the queries a doctor makes. Under current discovery rules, these logs are treated as standard Electronically Stored Information (ESI).

¹ The Digital Health Counsel is a multidisciplinary collaboration among individuals with shared interest in digital health policy. Each signatory participates in their personal capacity only and not on behalf of any employer, client, or other institution

This creates a dangerous divergence between the official medical record and the digital exhaust of the physician's workflow. These logs present a foreseeable litigation risk that plaintiff attorneys may seek to use.

1. **Impeach Clinical Expertise:** Using a physician's prudent verification of basic facts ("What is the dosage for X?") as evidence of incompetence or lack of training.
2. **Prove "Scienter" (Knowledge of Risk):** Using a physician's exploration of alternative treatments or "what-if" scenarios to prove they were subjectively aware of a risk that was not documented in the final note, elevating negligence claims to recklessness.
3. **Undermine Clinical Judgment:** Highlighting the timestamped gap between a query and a decision to suggest that the physician "shopped" for an answer that fit a desired outcome rather than following the patient's presentation.

The Chilling Effect on Adoption (Response to Question 9)

The RFI asks what concerns exist regarding the adoption of AI. The examples above illustrate the primary barrier: **Radical Transparency creates Defensive Medicine.**

If physicians understand that every question, they ask an AI tool (every admission of uncertainty) is being logged and can be subpoenaed to impeach their expertise, **they will stop asking questions.**

This creates a perverse incentive structure:

- **Safe Behavior (for the Patient):** The doctor uses AI to verify their knowledge, catching potential errors before they happen.
- **Safe Behavior (for the Doctor):** The doctor relies solely on memory to avoid creating a digital paper trail of "doubt."

We cannot expect clinicians to embrace these powerful diagnostic aids while the responsible use of such tools— verifying clinical judgments through AI queries— simultaneously exposes them to heightened litigation risk.

A Potential Solution: A Federal Standard of "Cognitive Privilege"

To accelerate the adoption of AI in clinical care, HHS must clarify that **clinical queries to AI decision-support tools constitute "Patient Safety Work Product" (PSWP).**

Just as hospital Morbidity & Mortality conferences and Peer Review deliberations are privileged to encourage open discussion of errors, **AI query logs must be privileged to encourage the open verification of knowledge.**

We respectfully recommend that HHS, in coordination with the Agency for Healthcare Research and Quality (AHRQ), issue guidance clarifying that:

1. **Protective Scope:** Prompt histories generated during the course of patient care are protected under the framework of the **Patient Safety and Quality Improvement Act of 2005 (PSQIA)**.
2. **Discovery Limits:** These logs should be discoverable only in cases of criminal misconduct, not standard civil liability, mirroring the protections afforded to peer review files.

Without this "Cognitive Privilege," the adoption of AI will be stifled by fear of legal liability, and the potential for these tools to reduce diagnostic error will be impeded.

Respectfully,

Tayab Waseem MD PhD, Jefferson Lin JD, Dave Schoolcraft JD, Stephanie Simmons JD on behalf of the Digital Health Counsel.